



# Apache Kylin: Speed up Cubing with Spark

Luke Han, [luke@kyligence.io](mailto:luke@kyligence.io)

Shaofeng Shi, [shaofeng@kyligence.io](mailto:shaofeng@kyligence.io)

# About Us

## Luke Han (韩卿)

- VP of Apache Kylin
- Co-founder & CEO, Kyligence Inc.

## Shaofeng Shi (史少锋)

- Apache Kylin PMC
- Sr. Architect, Kyligence Inc.



formed by the team who created Apache Kylin

<http://kyligence.io>

# Agenda

- What is Apache Kylin
- Challenges with MapReduce
- Speed up Cubing with Spark
- Q&A

# About Apache Kylin

**Kylin** / 'ki:'lin / 麒麟

-- n. (in Chinese art) a mythical animal of composite form

Apache Kylin

Extreme OLAP Engine for Big Data

- ✓ Leading open source OLAP on Hadoop
- ✓ Fast growing open source community
- ✓ Adopted by 200+ global organizations
- ✓ First born in China Apache Top Level Project
- ✓ InfoWorld Bossie Award:
  - Best Open Source Big Data Tool (2015)
  - Best Open Source Big Data Tool (2016)



<http://kylin.apache.org>

# Global Users

200+ use cases in production





# Use Case



*Apache Kylin is the best OLAP on Hadoop solution,  
we can't do analytics on trillion level data without  
Apache Kylin*

*-- Chaozhong Yang, Bytedance (Toutiao.com)*



- Top 1 News Feed App
- 600+M Users, 8000+ DAU
- 70+ minutes /user/day

## Pain Points

Huge Volume Data, 100+B rows per day

Second query latency by business requirement

## Solution

Apache Kylin as OLAP platform with 60+ HBase region servers

## Scenarios

#1: Video impression data, 4+ TB Cube (expansion rate < 3%)

(**Raw records: 3,000+B**, 100+TB)

#2: APP performance monitoring, **1000+ Analysts**

#3: News Log Insight, **100+ Cubes** for different products

## Benefits

1. Super fast query speed: 90%+ queries in ms level

(10,000+ times fast than hive)

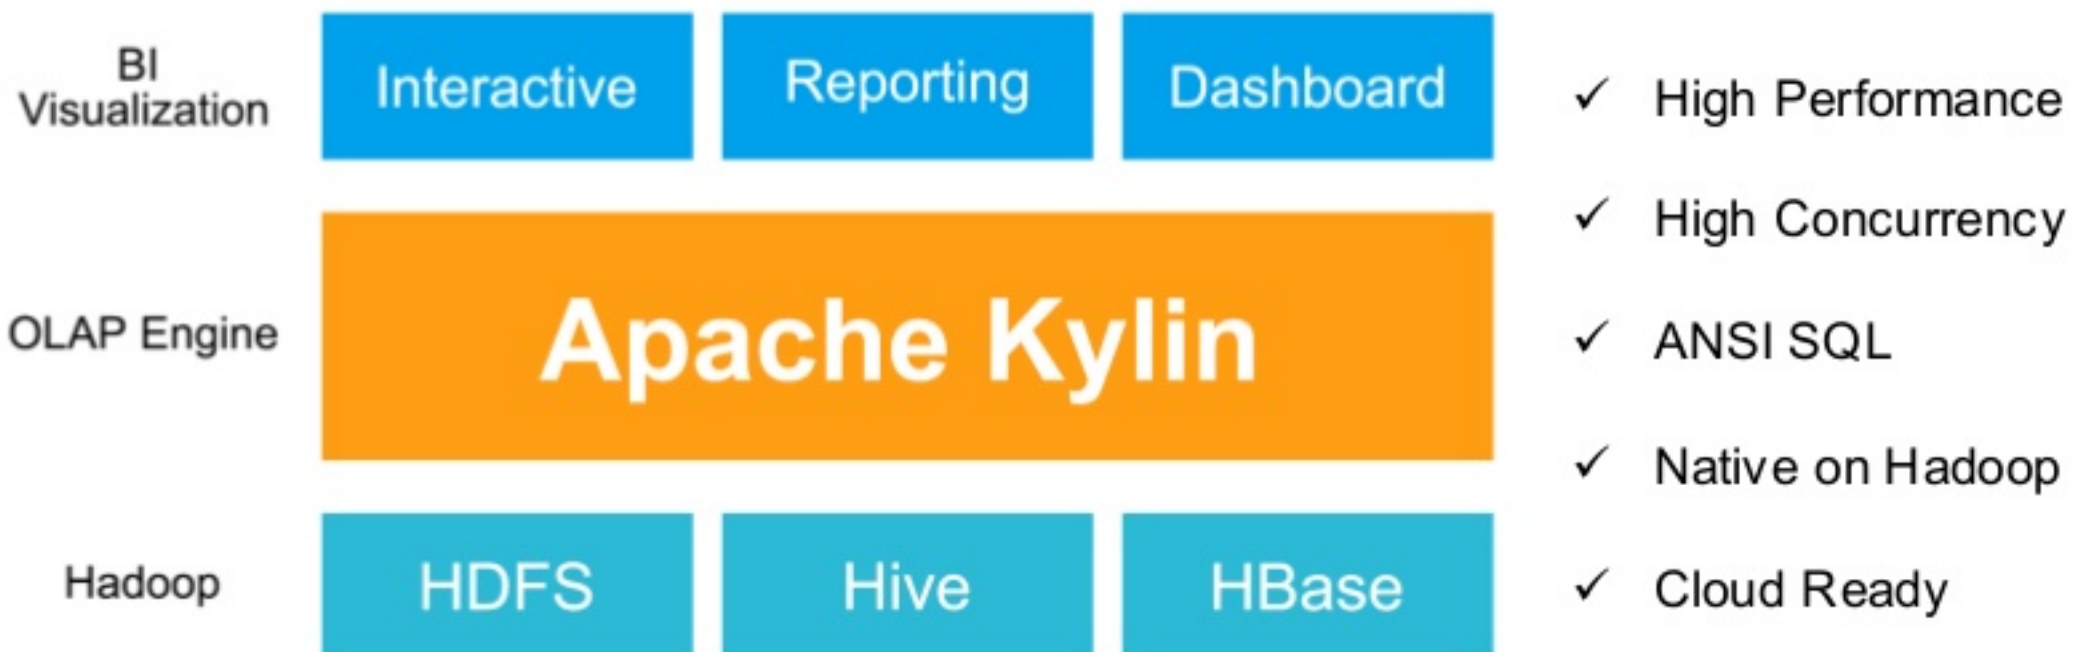
2. Saving Cluster Computing Resources: **One Build for all queries**

3. Rich API and drives, managing 100+Cube automatically

# About Apache Kylin

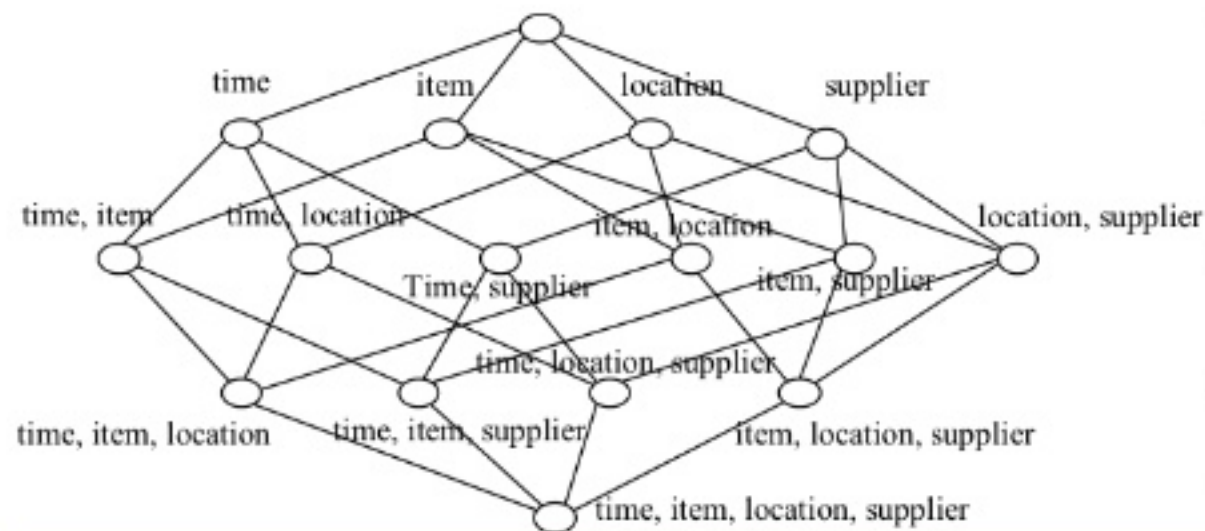


- Extreme OLAP Engine on Hadoop



# The Basic: OLAP Cube

- Pre-calculate metrics by dimensions



- Base vs. aggregate cells, ancestor vs. descendant cells, parent vs. child cells
  - (9/15, milk, Urbana, Dairy\_land) - <time, item, location, supplier>
  - (9/15, milk, Urbana, \*) - <time, item, location>
  - (\*, milk, Urbana, \*) - <item, location>
  - (\*, milk, Chicago, \*) - <item, location>
  - (\*, milk, \*, \*) - <item>

## OLAP Cube

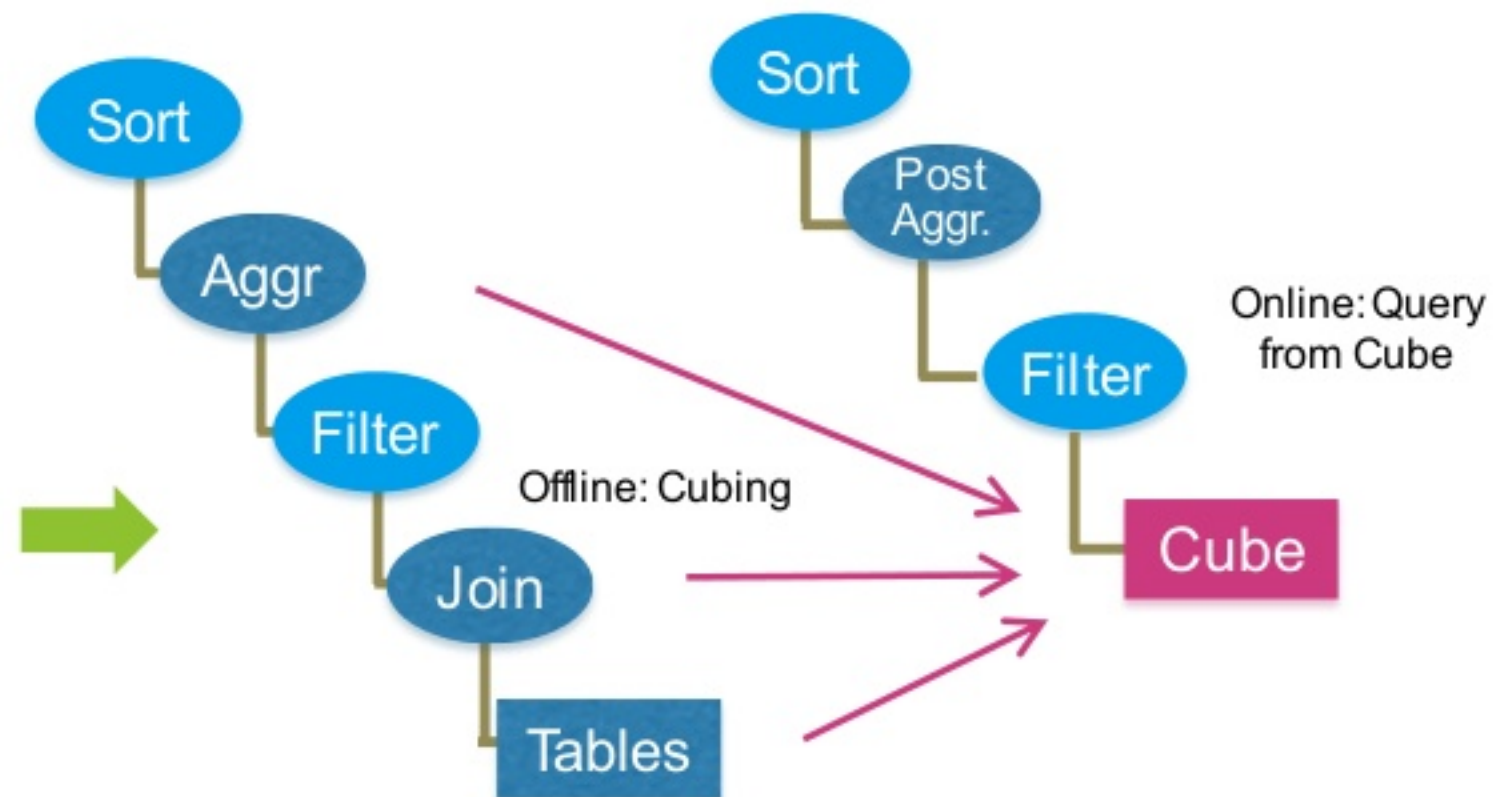
- Cuboid = one combination of dimensions
- Cube = all combination of dimensions (all cuboids)



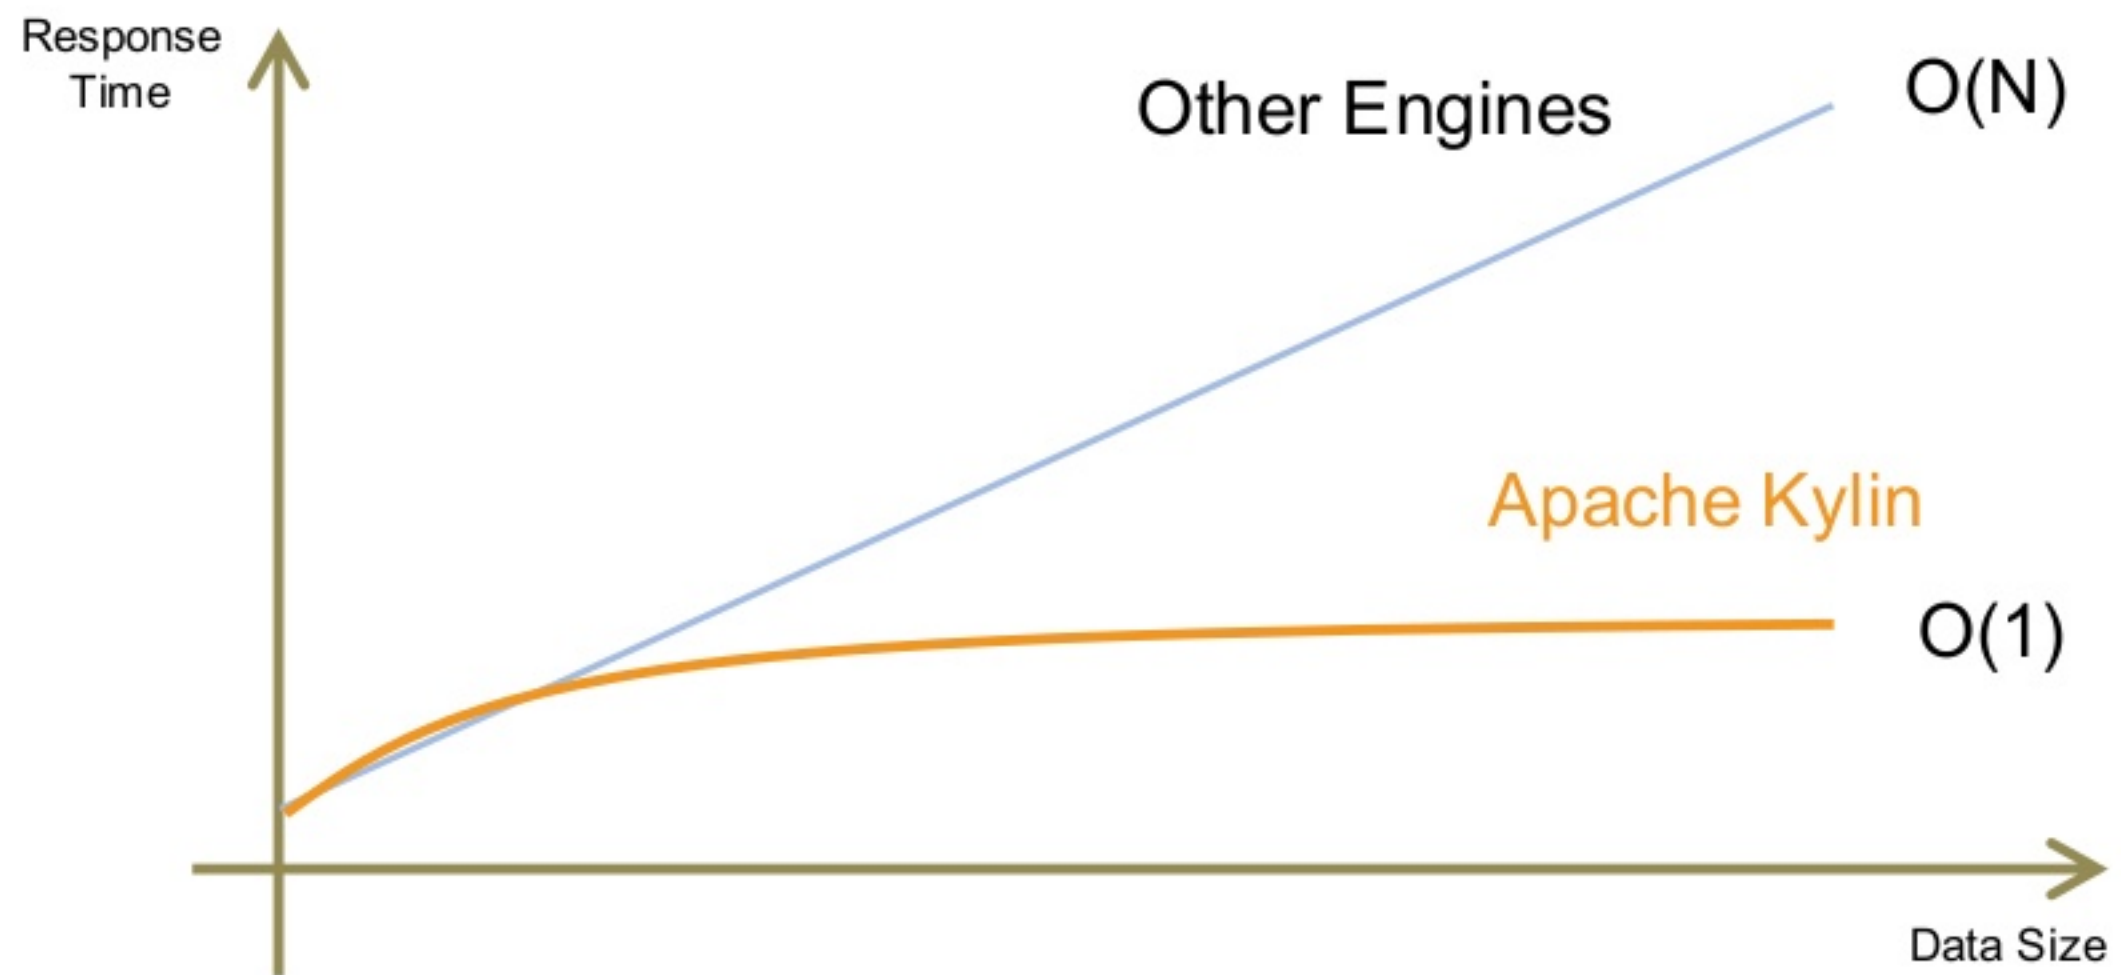


# The Magic: Pre-Calculation

```
SELECT
  returned,
  status,
  sum(quantity),
  sum(price)
FROM
  lineitem INNER JOIN orders
    on l_orderkey = o_orderkey
WHERE
  shipdate = '2016-09-16'
GROUP BY
  returned, status
ORDER BY
  returned, status;
```

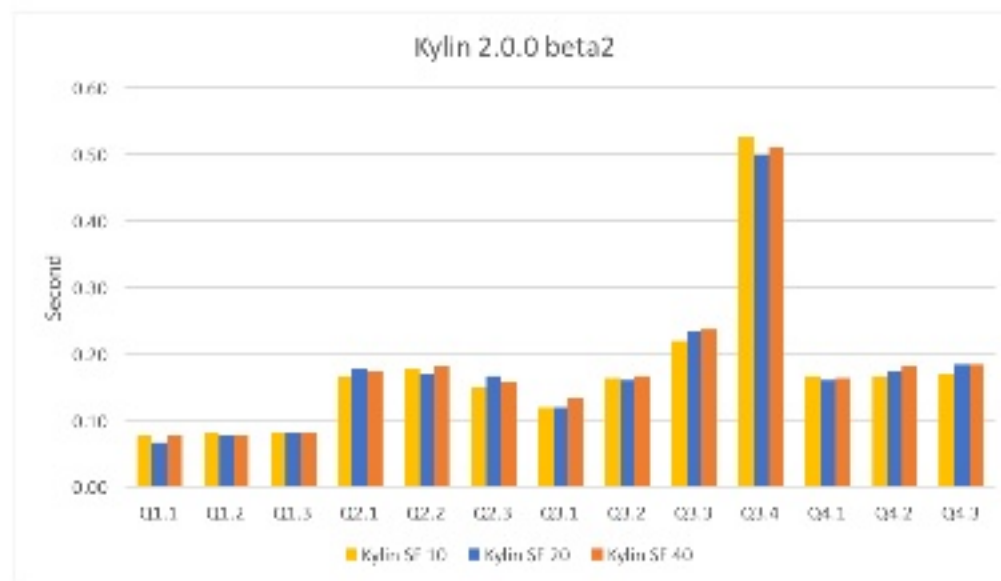


# Apache Kylin: $O(1)$

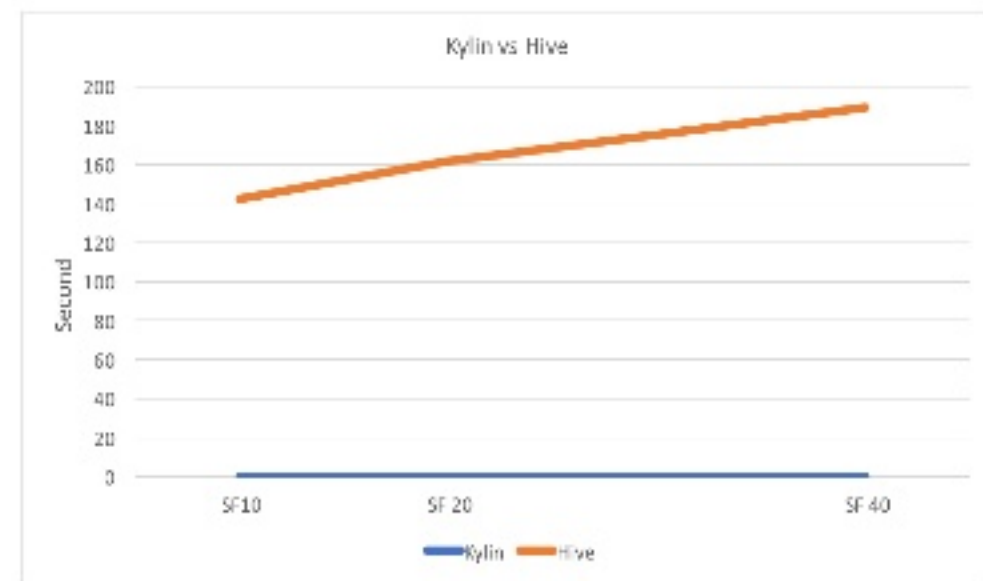


# Star-Schema Benchmark

Run SSB at 10, 20 and 40 million-row scales,  
Kylin's response time keep stable

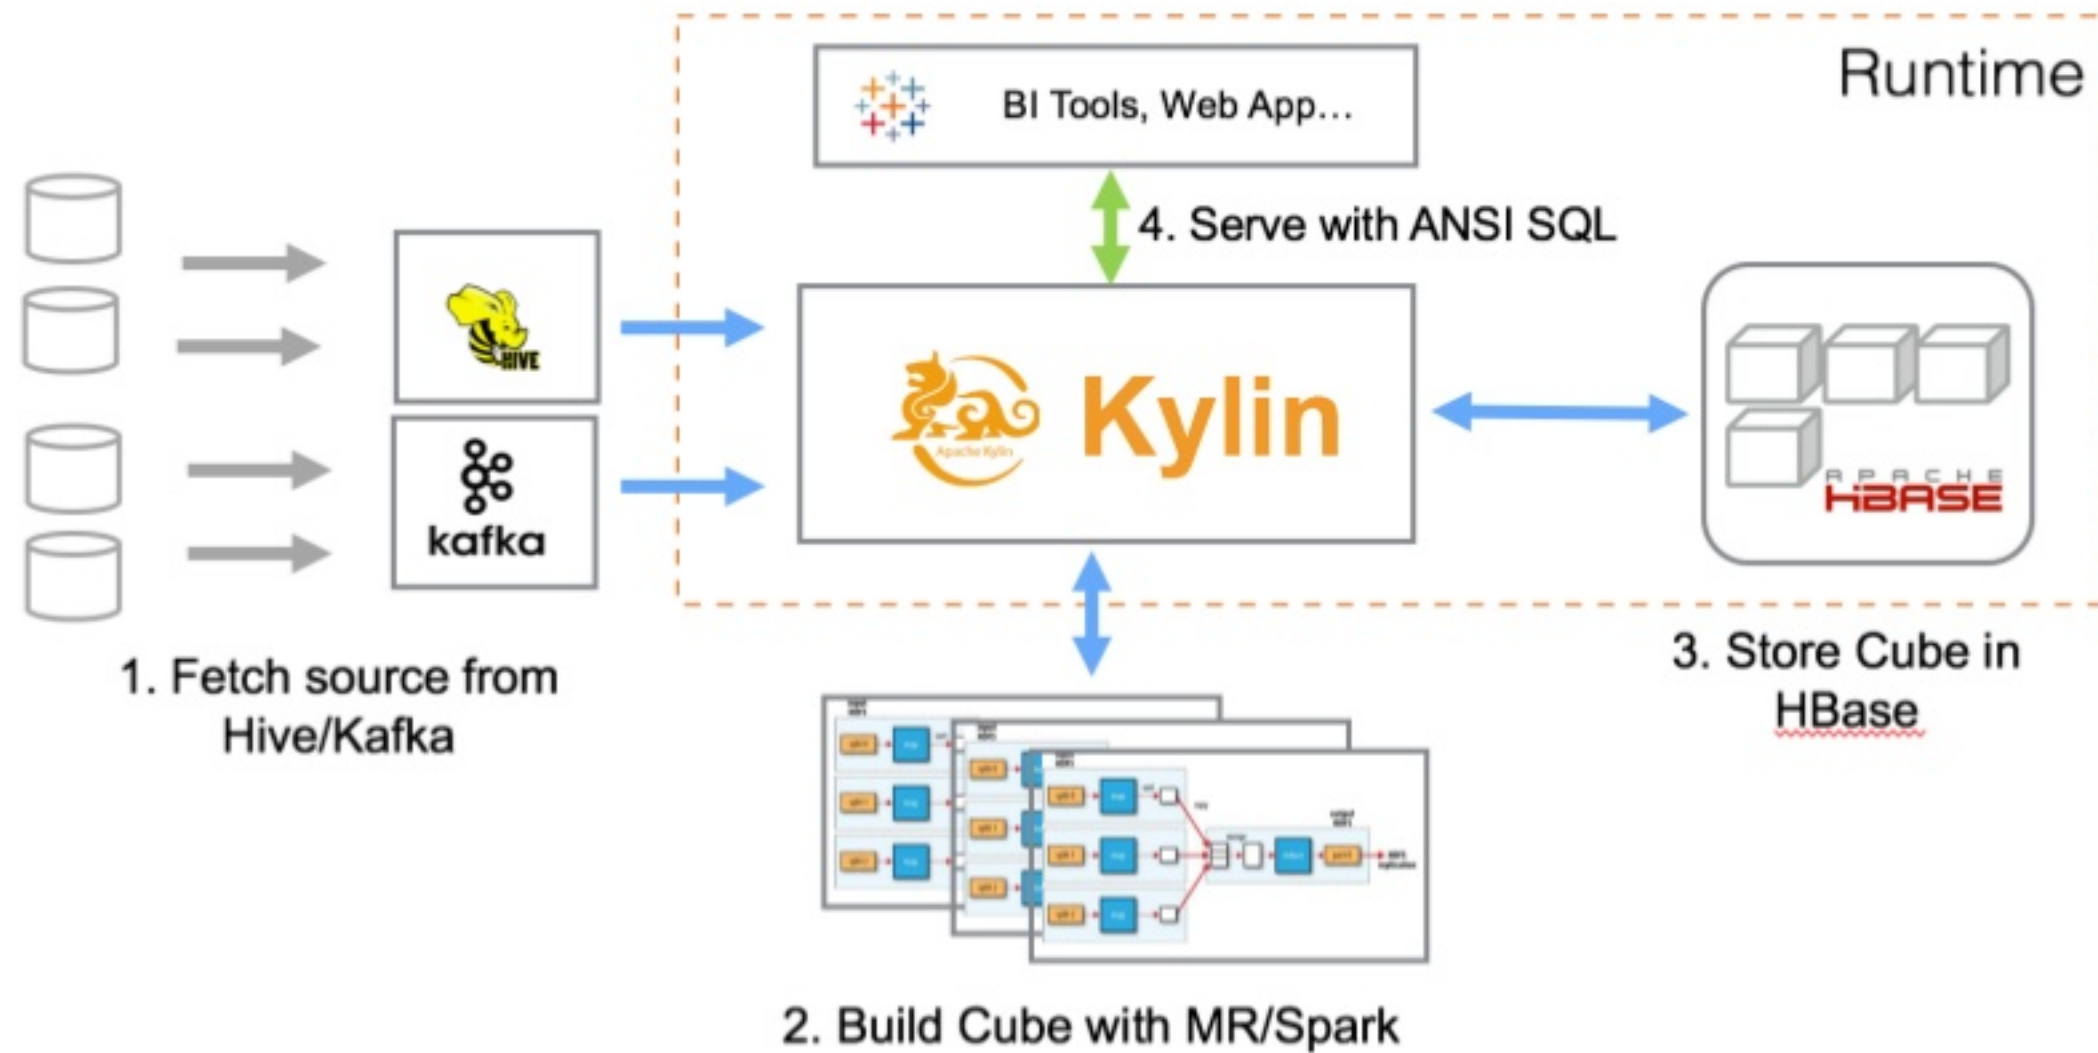


Kylin vs Hive,  $O(1) : O(N)$



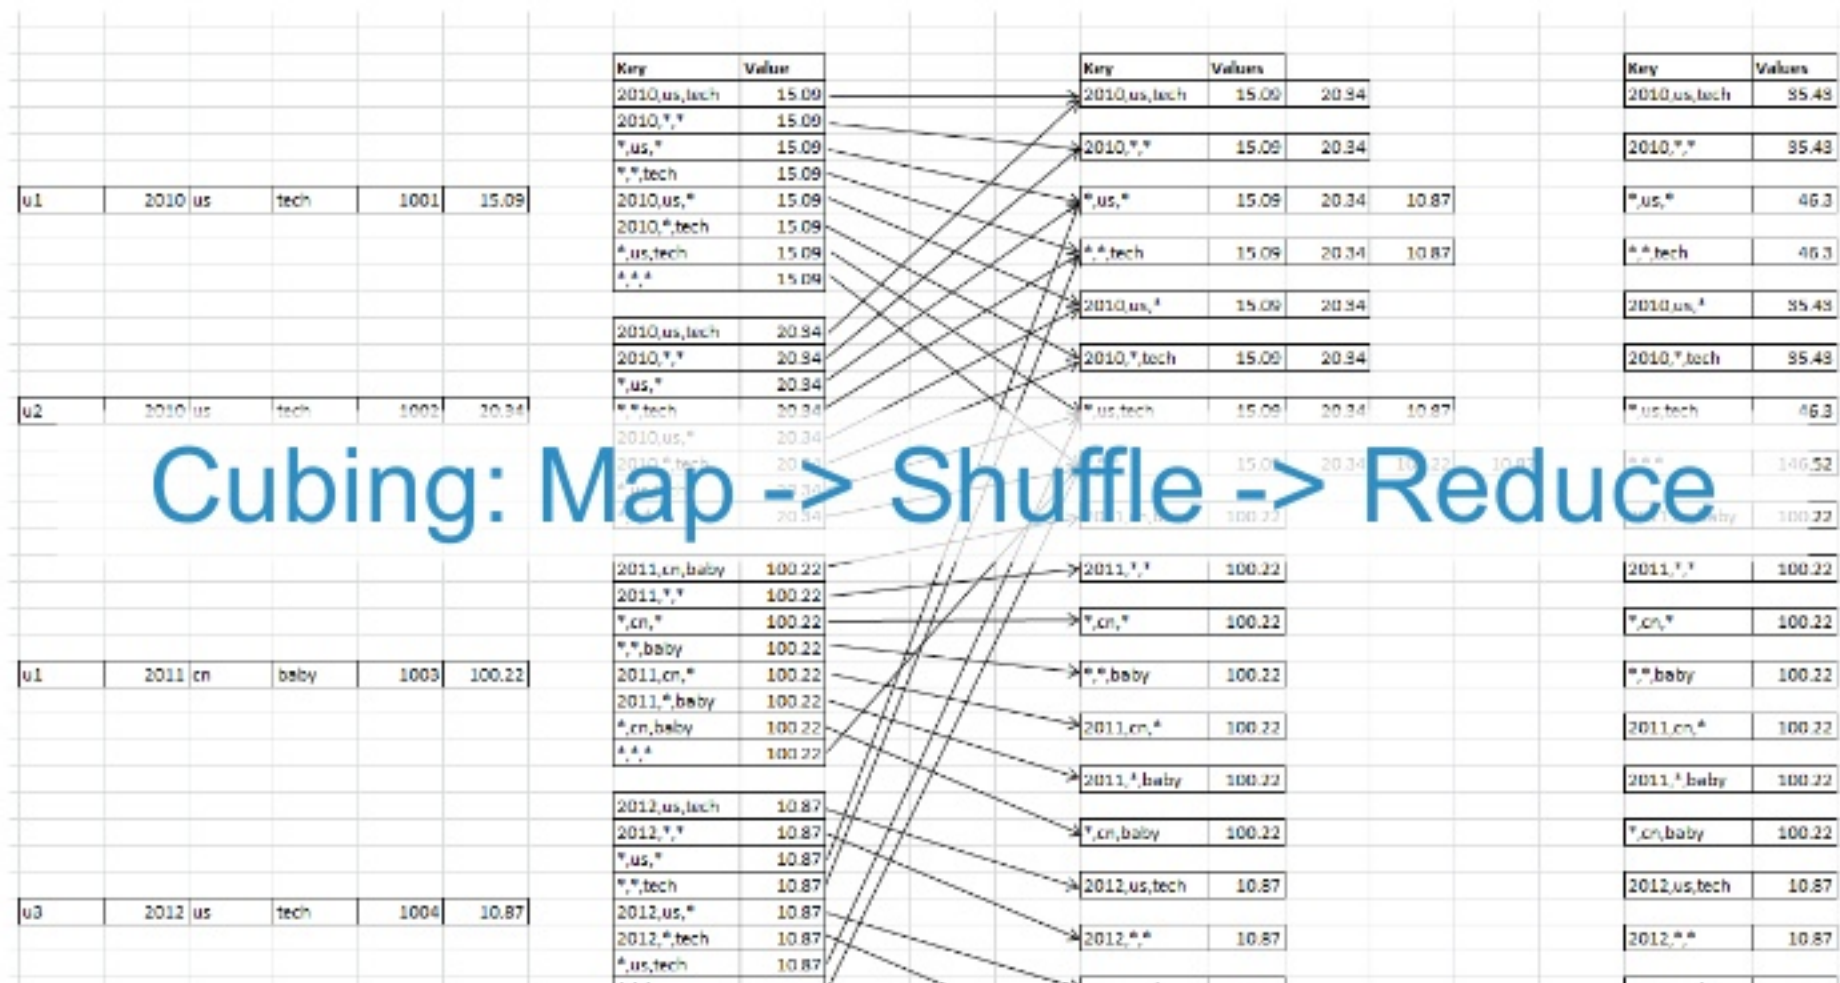
Read more at: <https://github.com/kyligence/ssb-kylin>

# Architecture



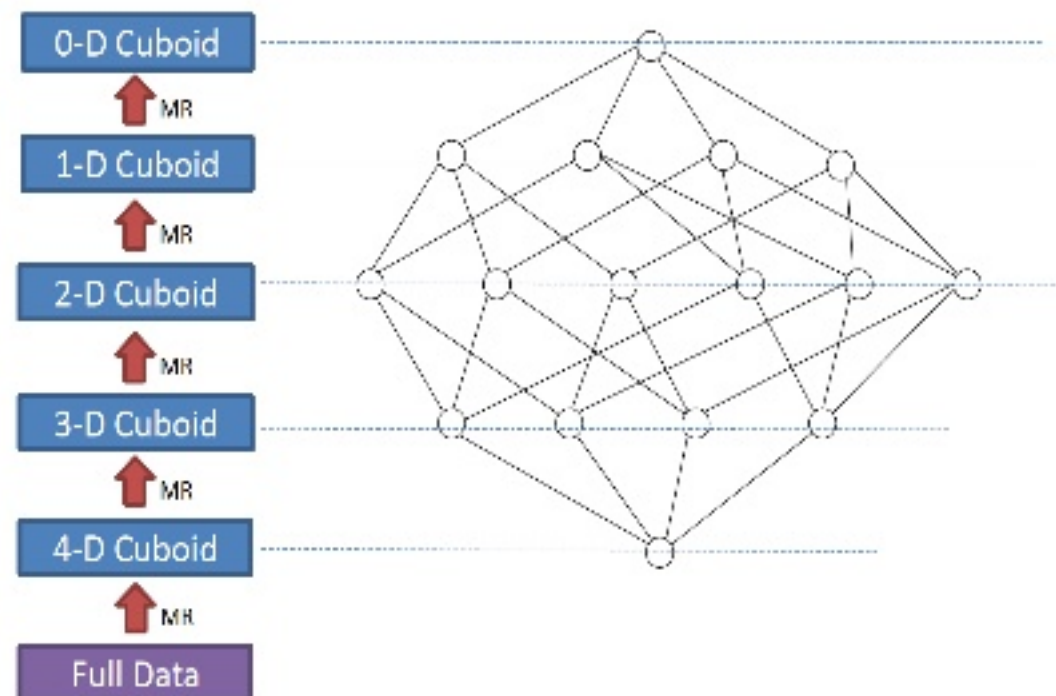


# Cubing Process



# Build Cube with MapReduce

- Calculate Cuboids by layer :N dim (Base cuboid), N-1 dim, N-2..., 1, 0
- Reuse previous layer's result
- HDFS used for data sharing
- Totally need N round MR;



# Challenges with MapReduce

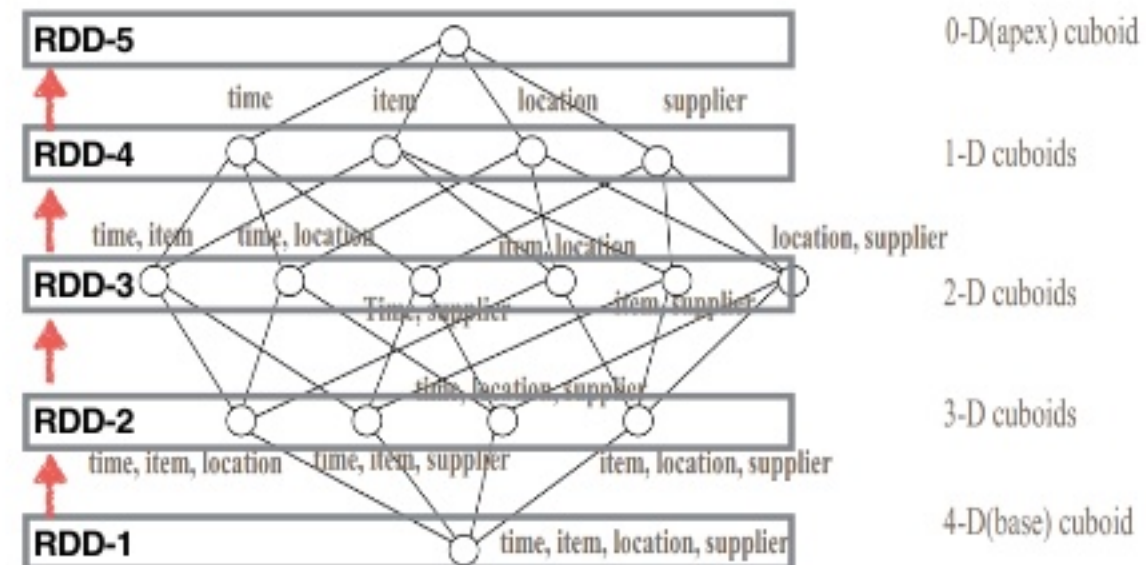
- Slow data sharing
  - ✓ Serialization, Replication, Deserialization...
- Repeated job submission
  - ✓ Submit dependent jar/files repeatedly
  - ✓ Re-queue when cluster is busy
- Short of streaming support
  - ✓ Couldn't support low latency data processing
- Limited storage support
  - ✓ Couldn't ingest data from non-HDFS





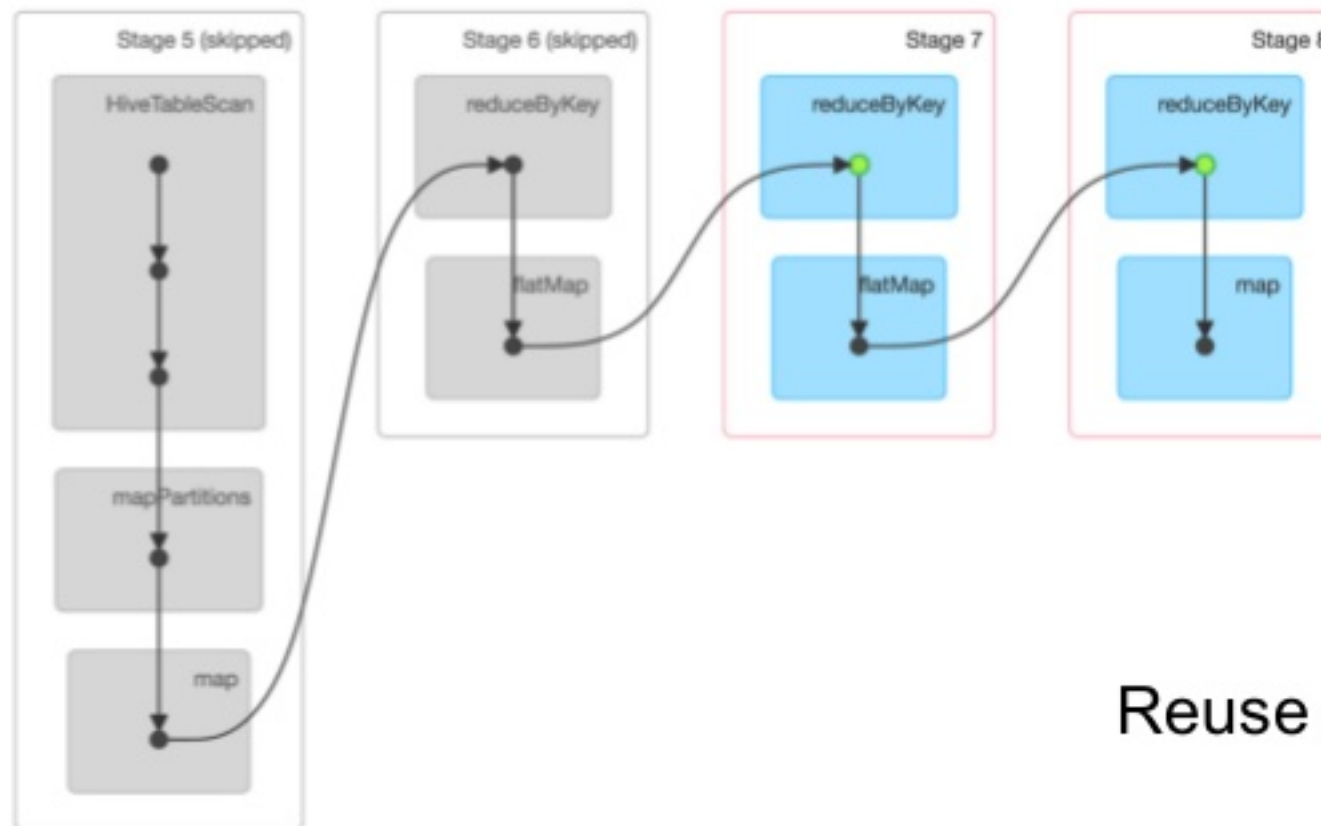
# Speed up Cubing with Spark

- Abstract each layer cuboids as a RDD
- Cache parent RDD to generate Child RDD
- Export RDD when Child be generated
- As-is measure aggregators can be reused with Spark Java API



# Speed up Cubing with Spark

▼ DAG Visualization



Reuse data in memory

# Speed up Cubing with Spark

## Spark Jobs <sup>(?)</sup>

Total Uptime: 5.0 min  
Scheduling Mode: FIFO  
Active Jobs: 1  
Completed Jobs: 6

► Event Timeline

### Active Jobs (1)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
6	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:13:57	47 s	1/8	<div><div></div></div> 502/1603

### Completed Jobs (6)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:12:32	1.4 min	2/2 (5 skipped)	<div><div></div></div> 816/816 (345 skipped)
4	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:11:23	1.1 min	2/2 (4 skipped)	<div><div></div></div> 602/602 (116 skipped)
3	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:10:40	41 s	2/2 (3 skipped)	<div><div></div></div> 315/315 (30 skipped)
2	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:10:28	12 s	2/2 (2 skipped)	<div><div></div></div> 100/100 (16 skipped)
1	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:10:22	5 s	2/2 (1 skipped)	<div><div></div></div> 28/28 (2 skipped)
0	<a href="#">saveAsNewAPIHadoopFile at SparkCubingByLayer.java:287</a>	2017/03/06 14:10:08	14 s	2/2	<div><div></div></div> 16/16

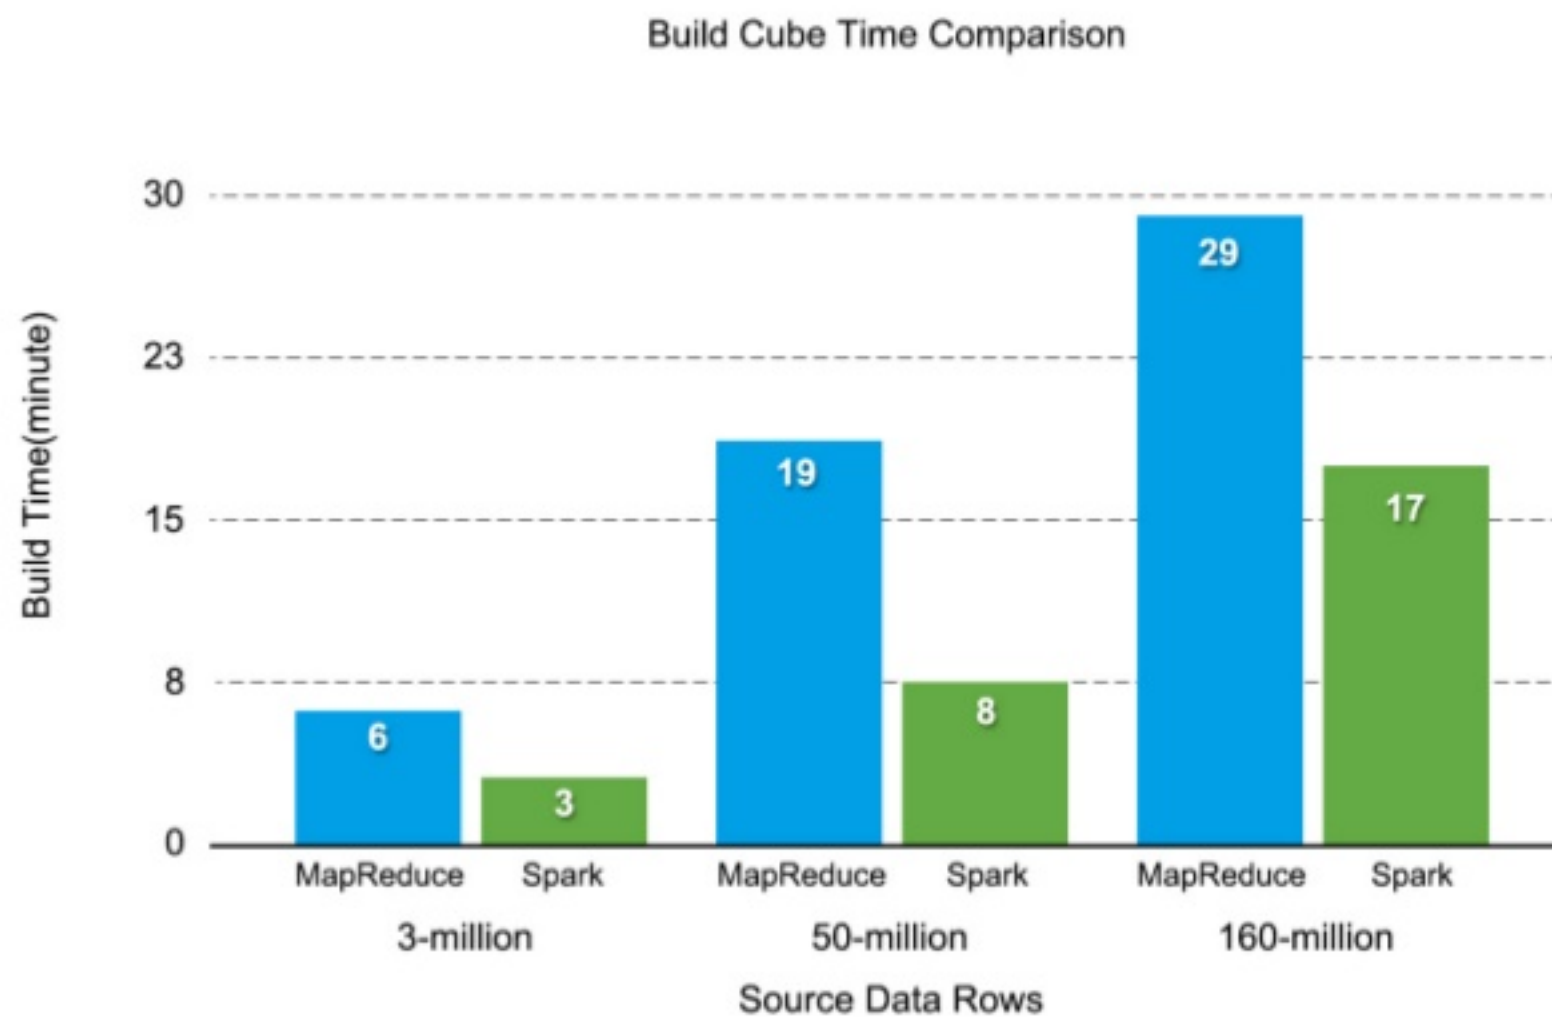
One job finishes all layers aggregation!

# Performance Benchmark

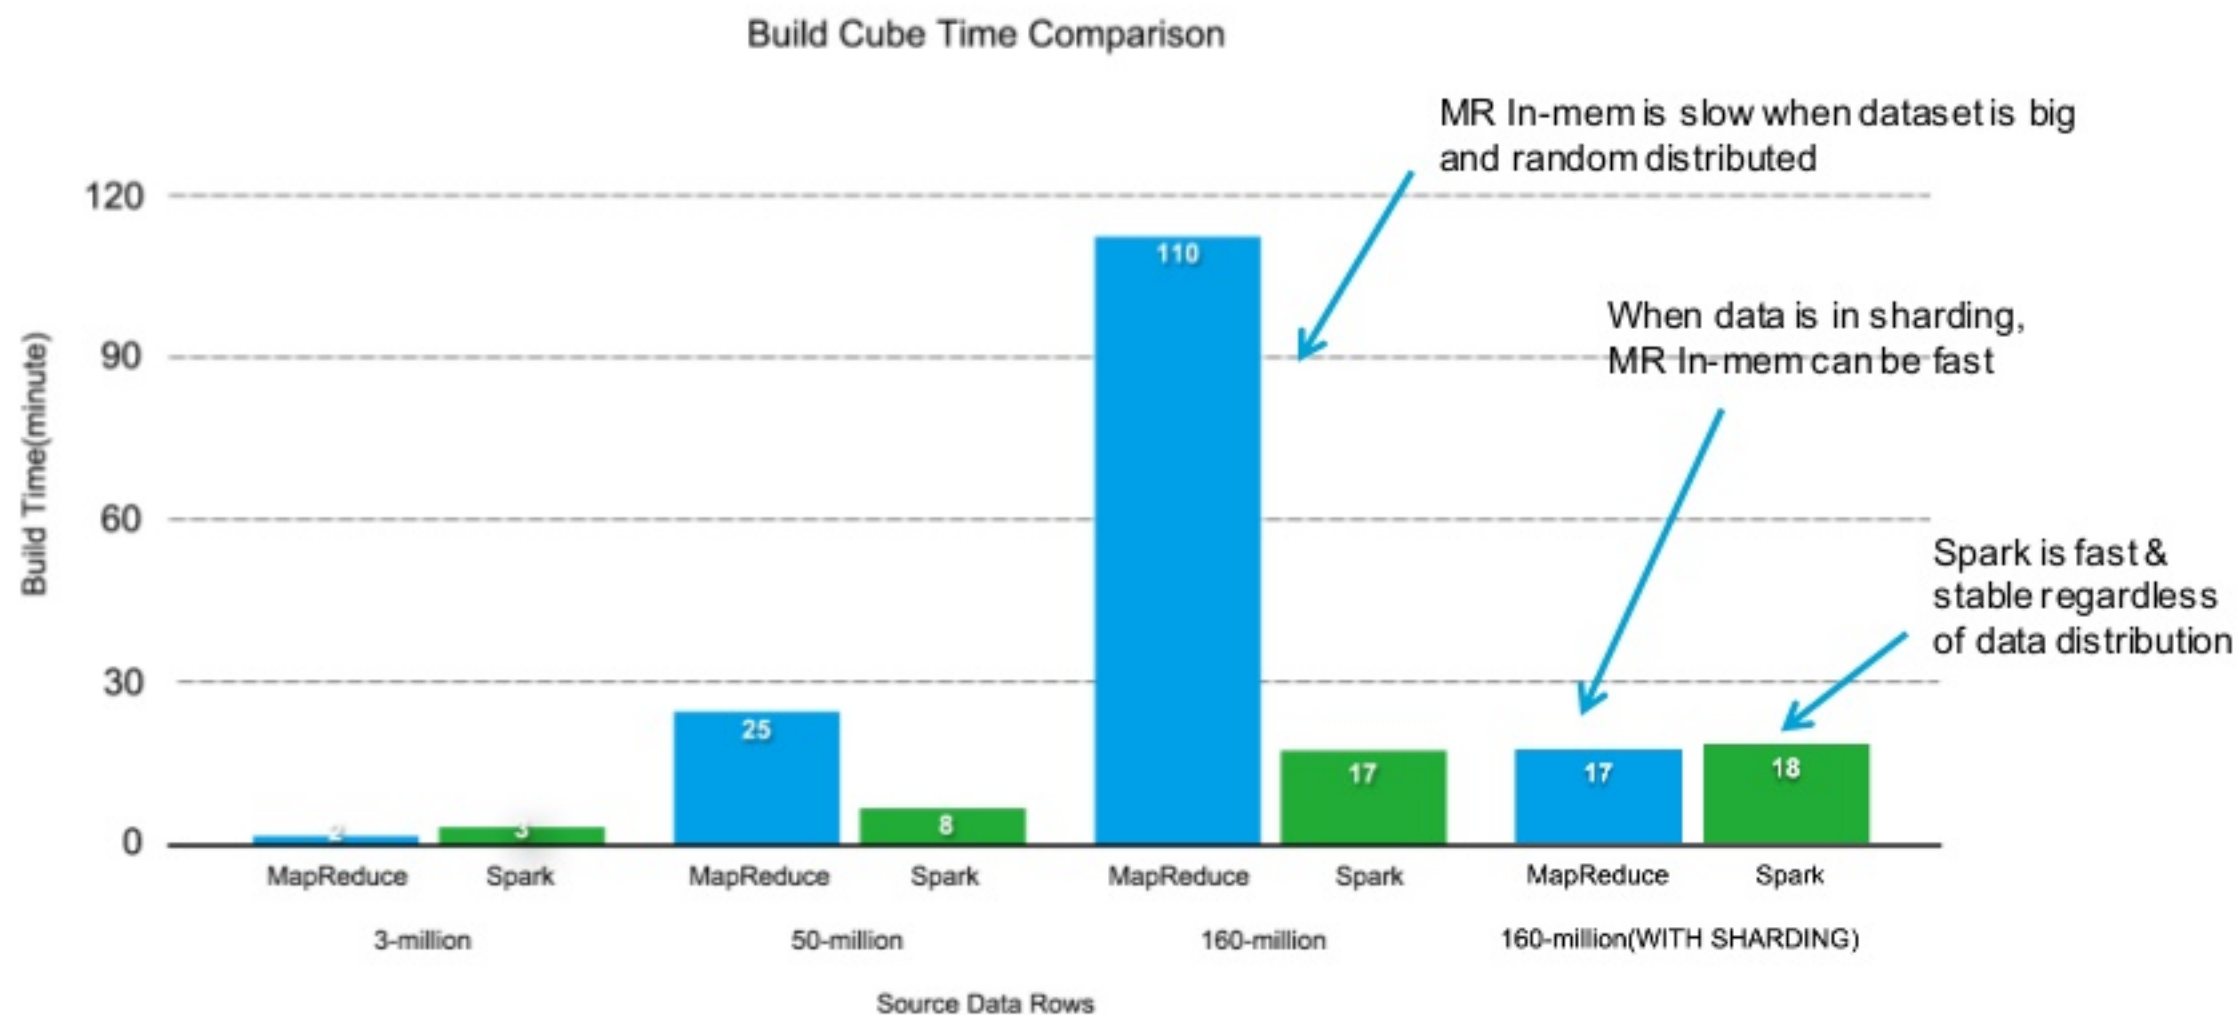
- Environment
  - ✓ 4 nodes Hadoop cluster; each has 28 GB RAM and 12 cores
  - ✓ CDH 5.8, Apache Kylin 2.0
- Spark
  - ✓ Spark 1.6.3 on YARN
  - ✓ 6 executors, each has 4 cores, 5GB memory
- Test Data
  - ✓ Airline data of US DoT, totally 160 million rows
  - ✓ Cube: 10 dimensions, 5 measures (SUM)
- Test Scenarios
  - ✓ Build the cube at different scale: 3 million, 50 million and 160 million rows



# Spark Cubing vs. MR Cubing



# Spark Cubing vs. MR In-Mem Cubing



# Benefits with Spark

- Spark speeds up Cubing at 1x
  - ✓ Half time be saved
- Spark simplifies Kylin's development
  - ✓ More functions, less codes
- Spark brings Kylin to a new era
  - ✓ Real-time OLAP, Ad hoc query, Cloud integration



# Thank You.

Join Apache Kylin community

[dev@kylin.apache.org](mailto:dev@kylin.apache.org)

Visit Kylin & Kyligence home

<http://kylin.apache.org>

<http://kyligence.io>