

Best Practices for Using Alluxio with Spark

Gene Pang, Alluxio, Inc.

Cheng Chang, Alluxio, Inc.

Spark Summit SF - June 2017



Outline

- 1 Alluxio Overview
- 2 Alluxio + Spark Use Cases
- 3 Using Spark with Alluxio
- 4 Performance Evaluation
- 5 Demo

• Data Ecosystem Yesterday



- One Compute Framework
- Single Storage System
- Co-located

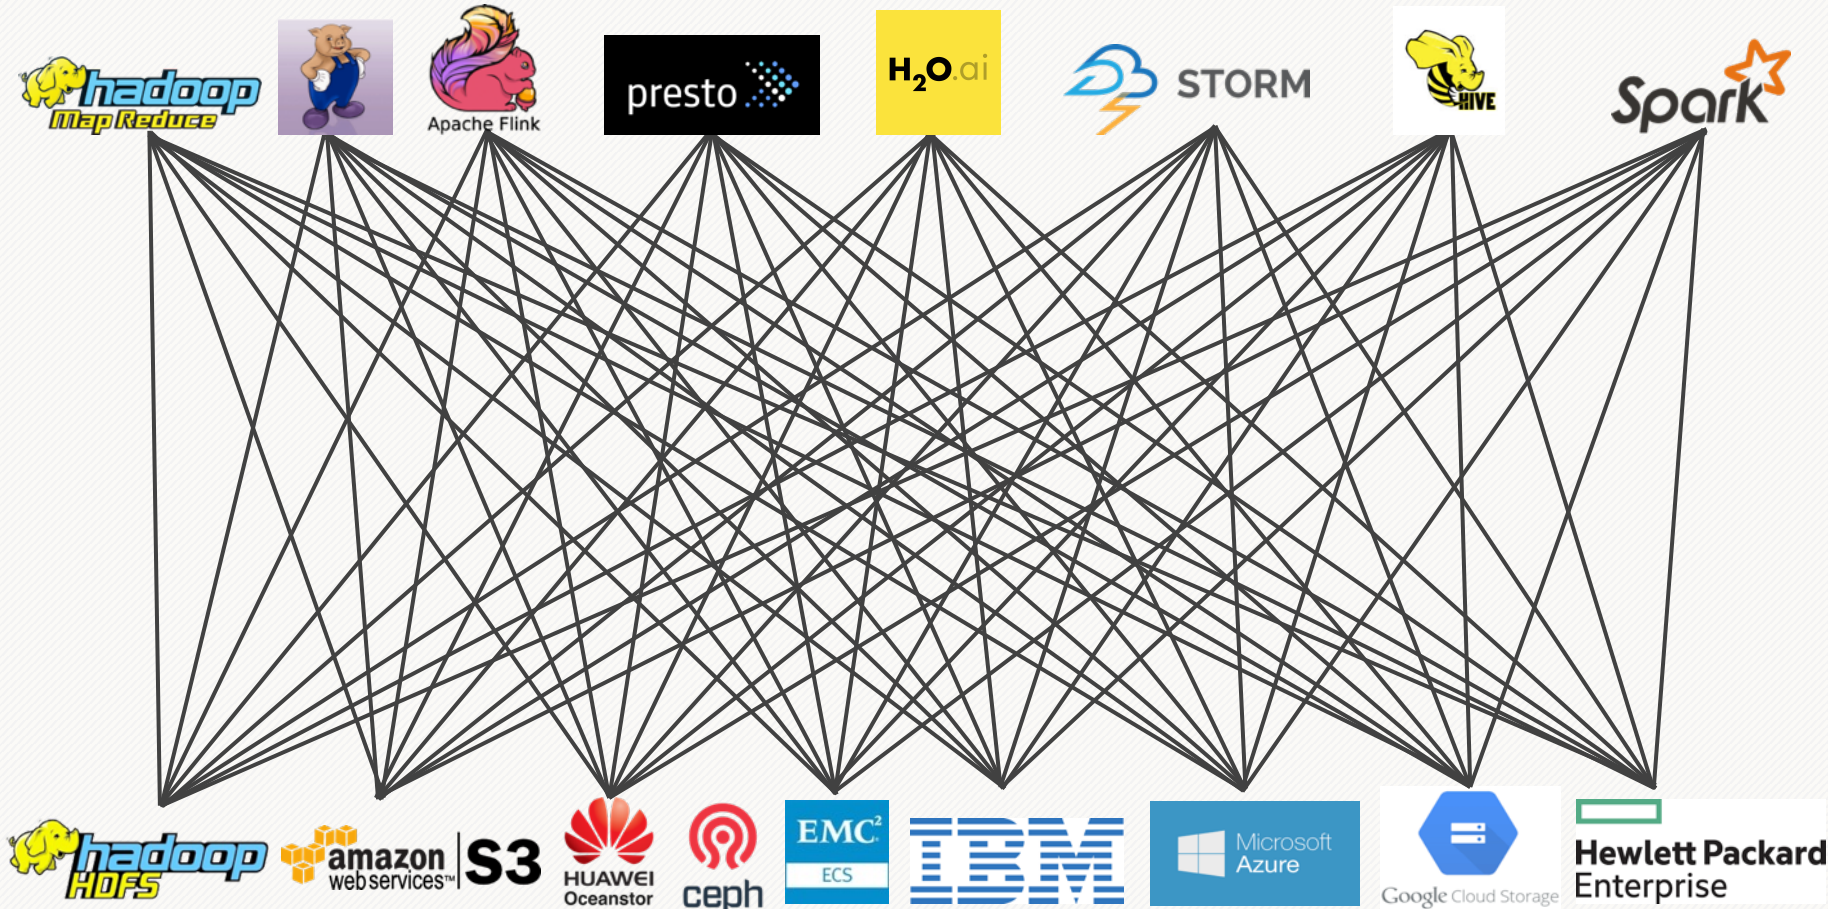
• Data Ecosystem Today



- Many Compute Frameworks
- Multiple Storage Systems
- Most not co-located

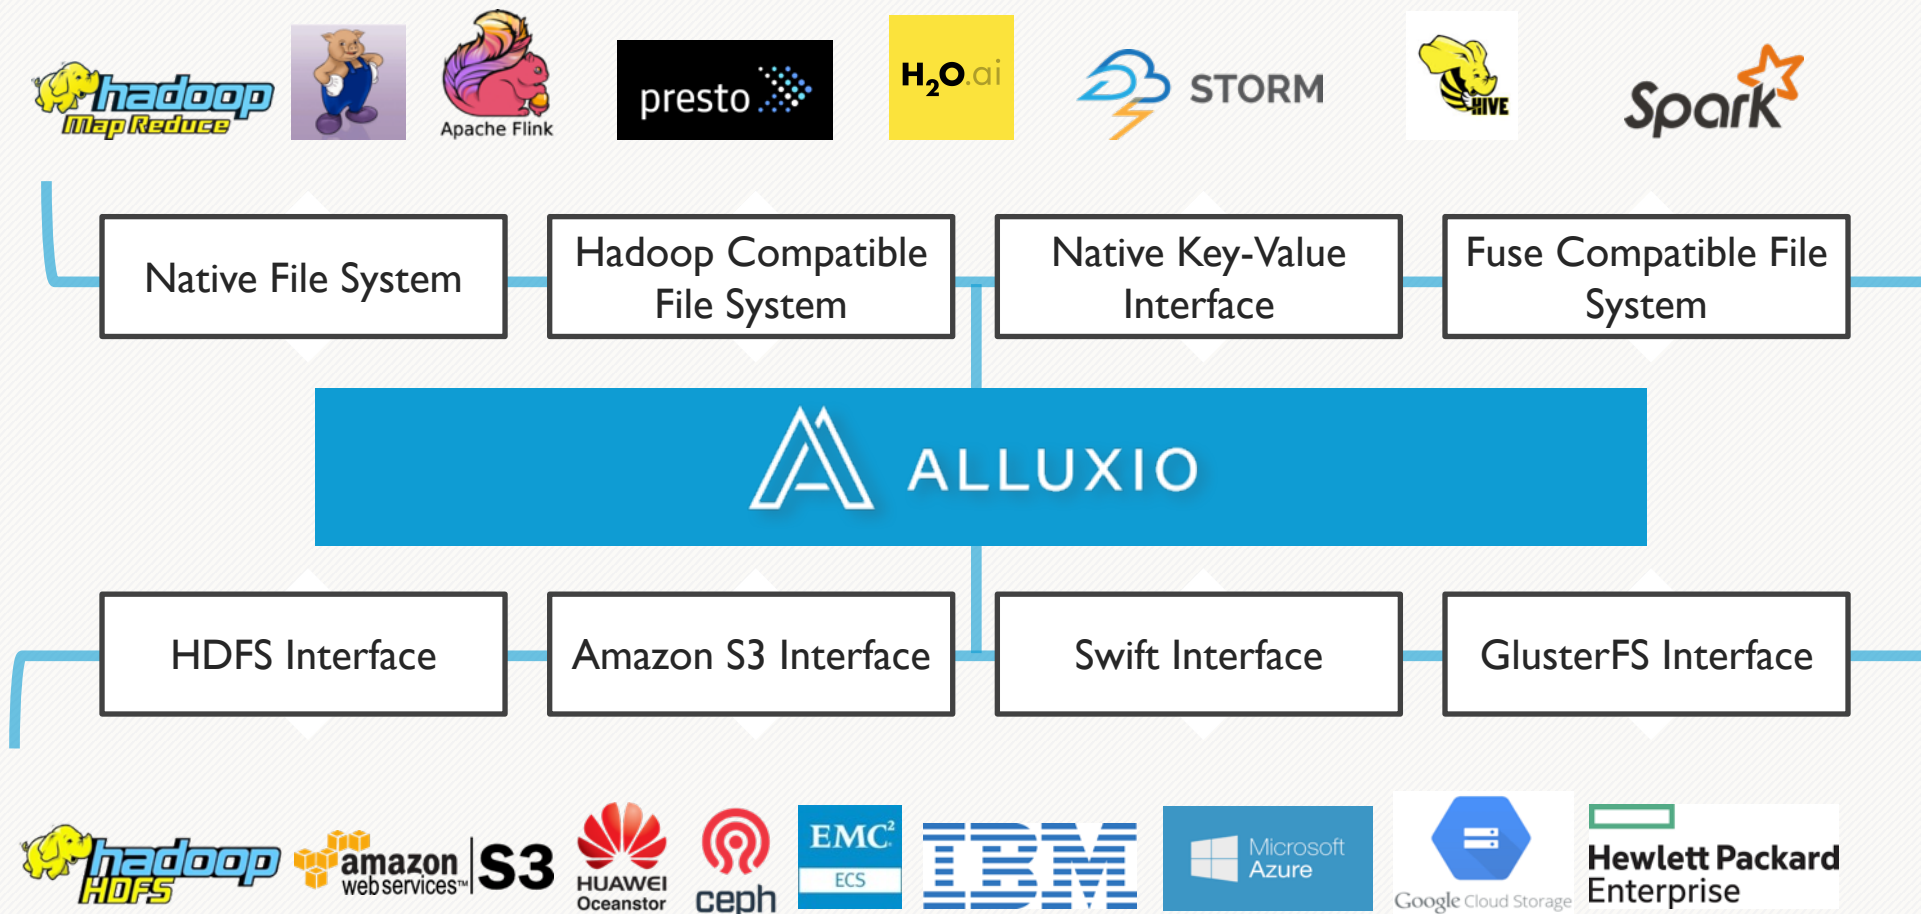


Data Ecosystem Issues



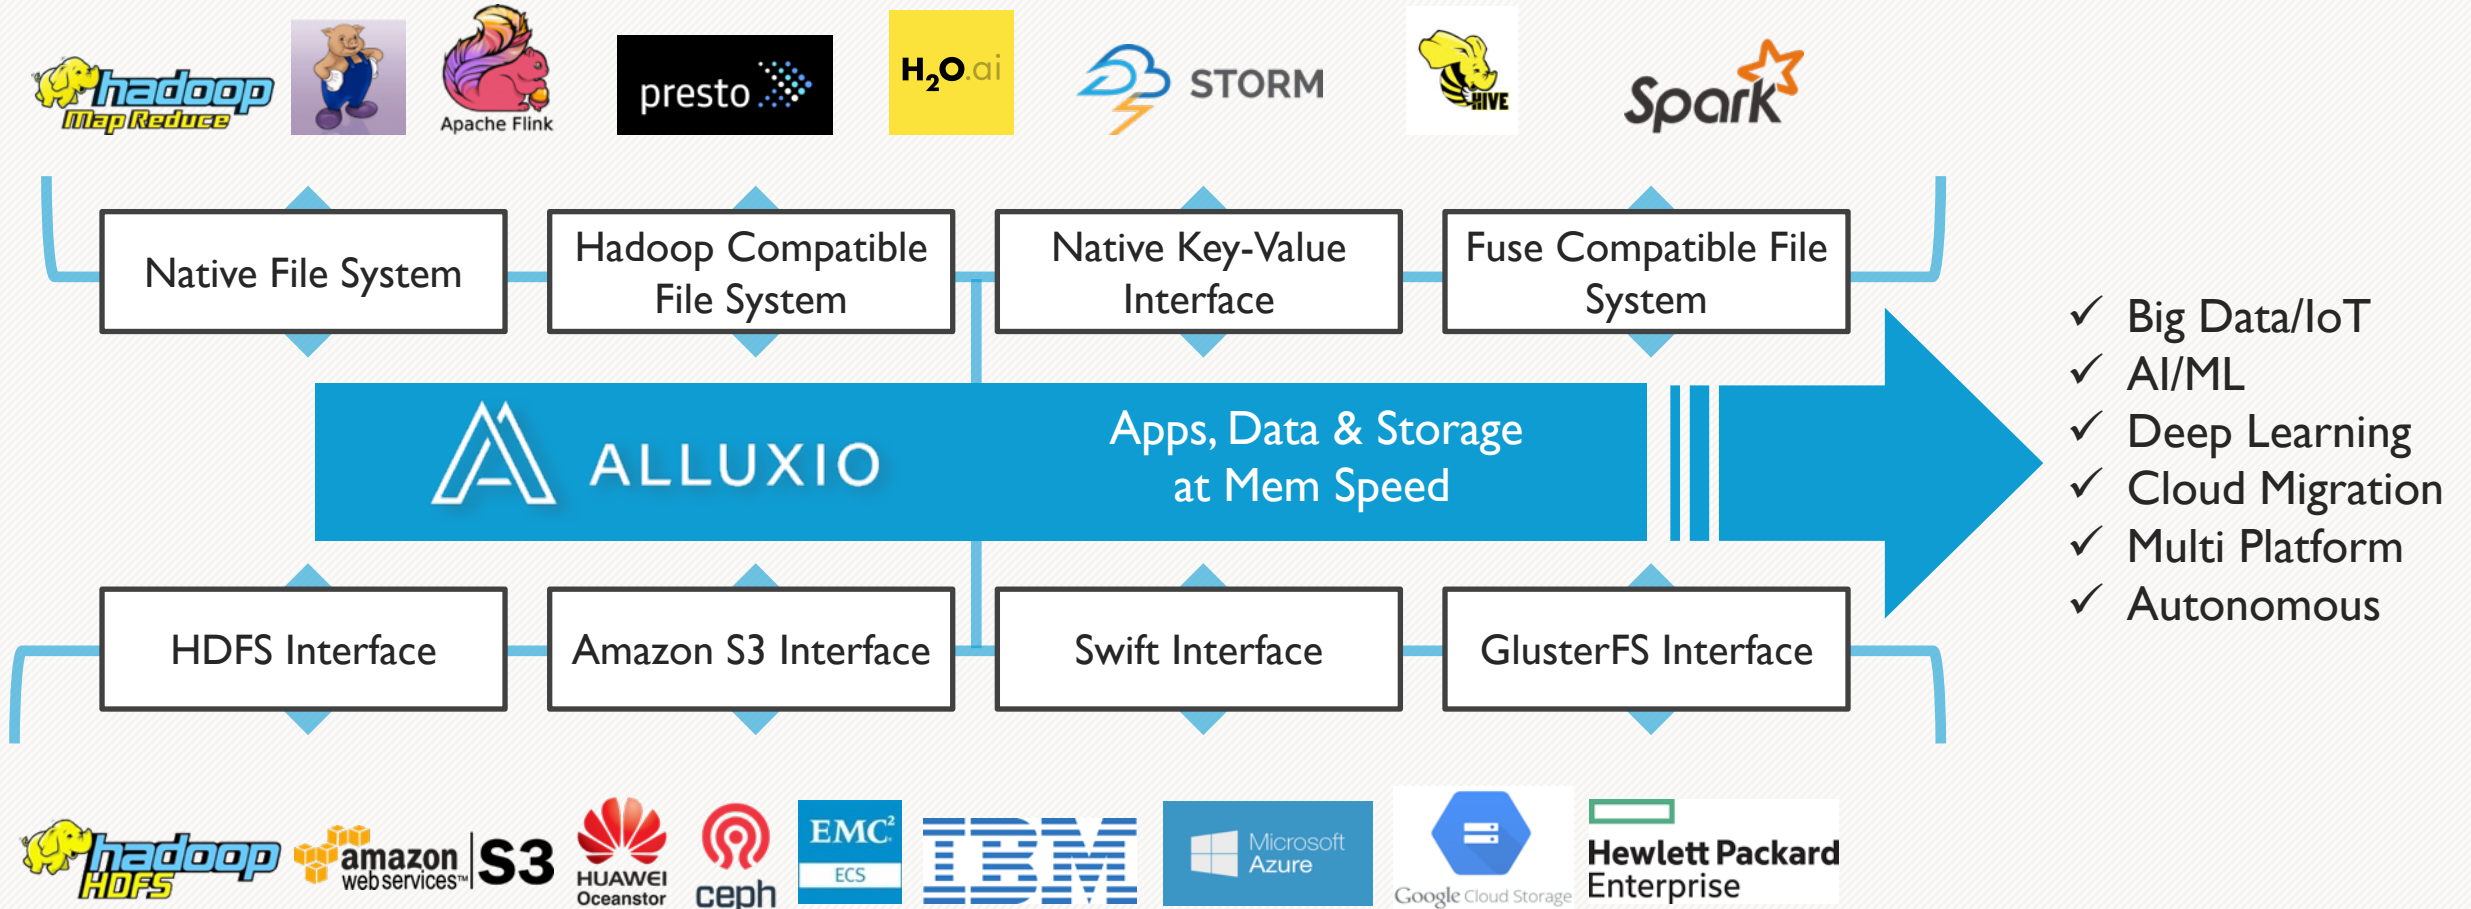
- Each application manage multiple data sources
- Add/Removing data sources require application changes
- Storage optimizations requires application change
- Lower performance due to lack of locality

Data Ecosystem with Alluxio

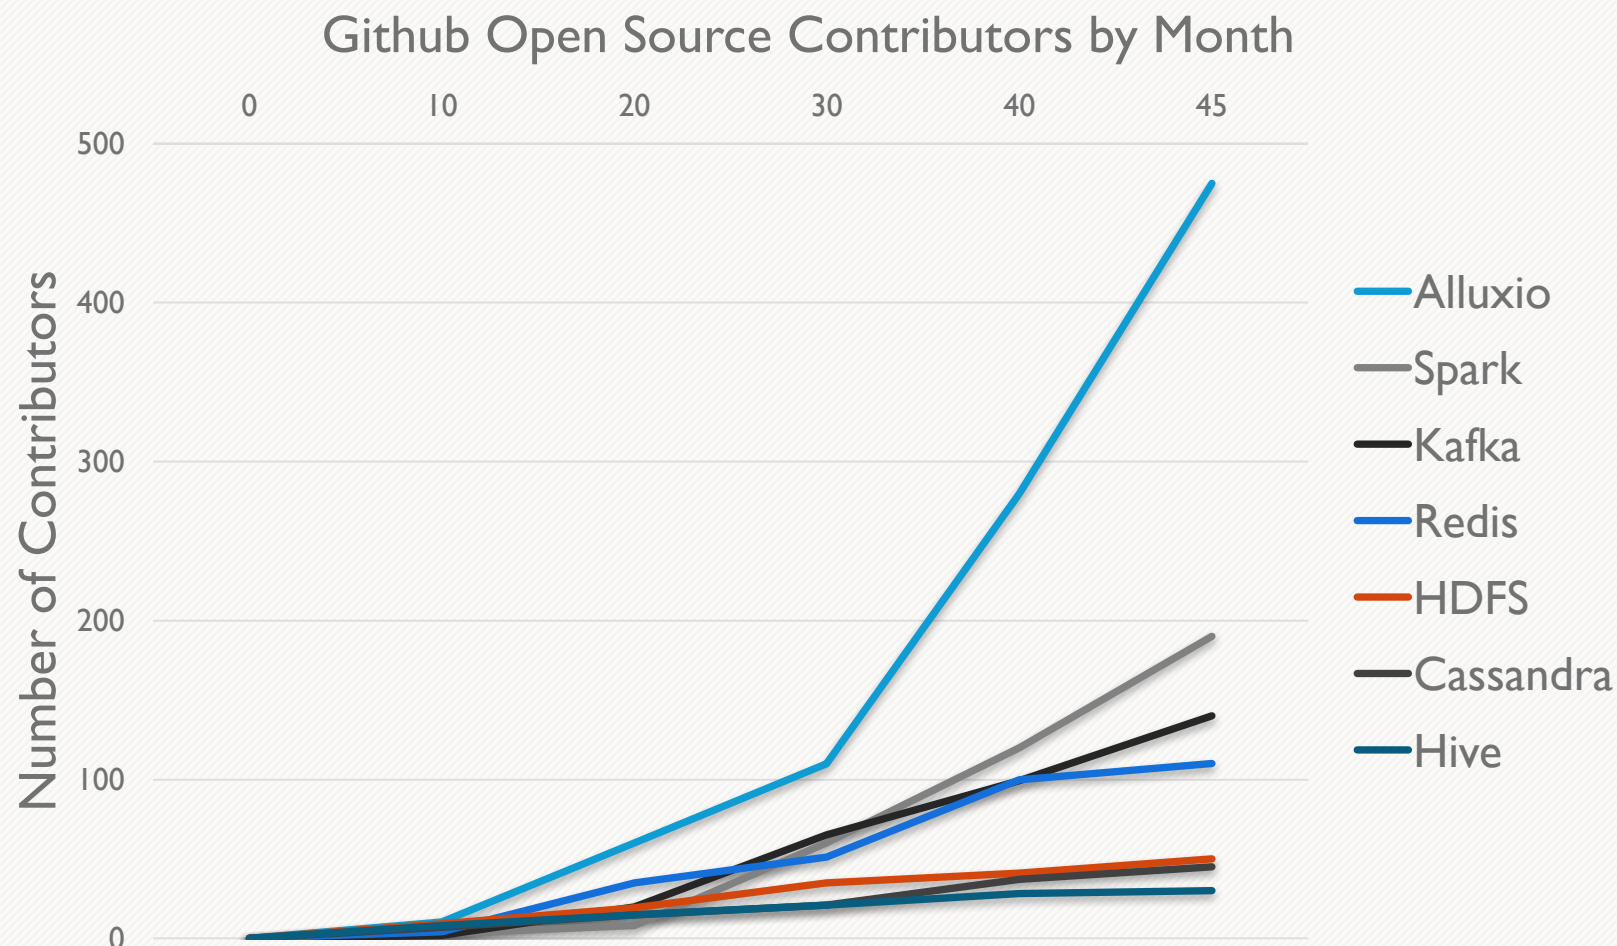


- Apps only talk to Alluxio
- Simple Add/Remove
- No App Changes
- Highest performance in Memory
- No Lock in

Next Gen Analytics with Alluxio



Fastest Growing Big Data Open Source Projects



Fastest Growing open-source project in the big data ecosystem

Running in large production clusters

500+ Contributors from 100+ organizations



Outline

- 1 Alluxio Overview
- 2 Alluxio + Spark Use Cases
- 3 Using Spark with Alluxio
- 4 Performance Evaluation
- 5 Demo

Big Data Case Study – **BARCLAYS**



SPARK

TERADATA

Challenge –

Gain end to end view of business with large volume of data

Queries were slow / not interactive, resulting in operational inefficiency

SPARK

 **ALLUXIO**

TERADATA

Solution –

ETL Data from Teradata to Alluxio

Impact –

Faster Time to Market – “Now we don’t have to work Sundays”

<http://bit.ly/2oMx95W>

Big Data Case Study –



SPARK

Baidu File System

Challenge –

Gain end to end view of business with large volume of data

Queries were slow / not interactive, resulting in operational inefficiency

SPARK

 ALLUXIO

Baidu File System

Solution –

With Alluxio, data queries are 30X faster

Impact –

Higher operational efficiency

<http://bit.ly/2pDHS3O>

Big Data Case Study –



SPARK

FLINK

HDFS

CEPH

Challenge –

Gain end to end view of business with large volume of data for \$5B Travel Site

Queries were slow / not interactive, resulting in operational inefficiency

SPARK

FLINK



HDFS

CEPH

Solution –

With Alluxio, 300x improvement in performance

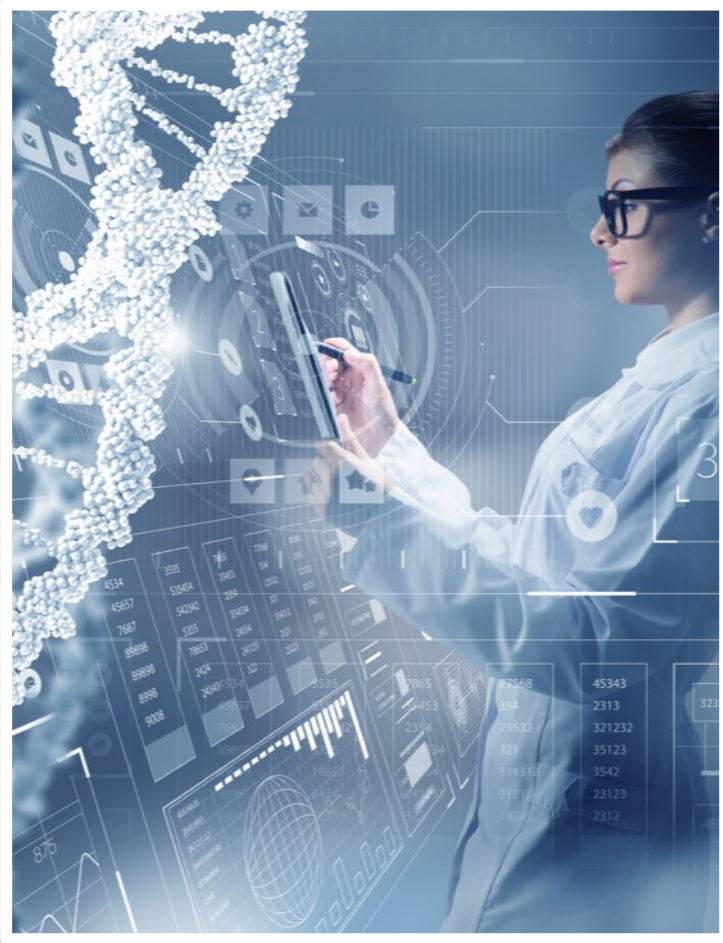
Impact –

Increased revenue from immediate response to user behavior

Use case: <http://bit.ly/2pDJdrq>



Machine Learning Case Study –



SPARK

HDFS

Challenge –

Disparate Data both on-prem and Cloud. Heterogeneous types of data.

Scaling of Exabyte size data.
Slow due to disk based approach.

SPARK



MINIO

MESOS

Solution –

Using Alluxio to prevent I/O bottlenecks

Impact –

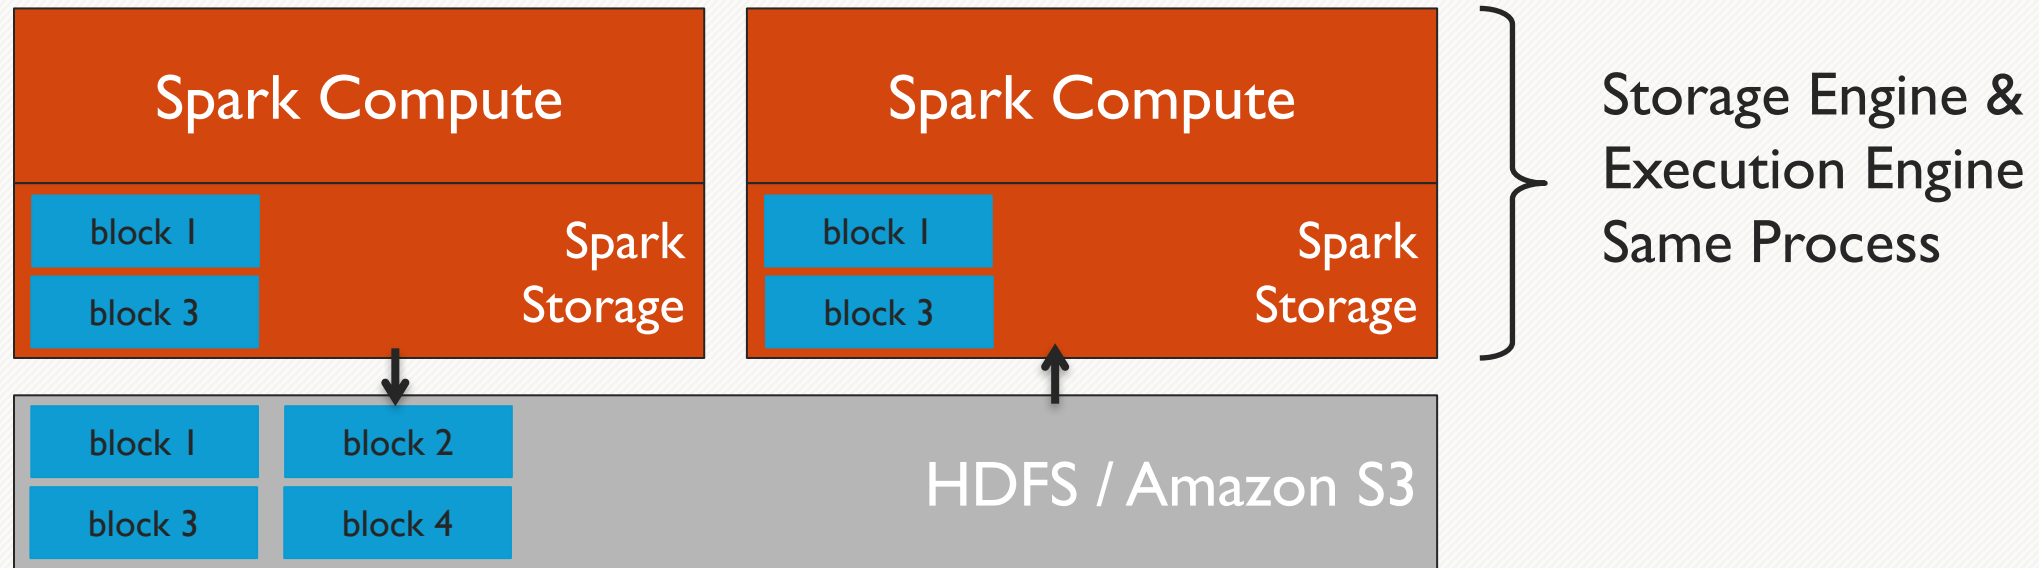
Orders of magnitude higher performance than before.
<http://bit.ly/2pI8ds3>



Outline

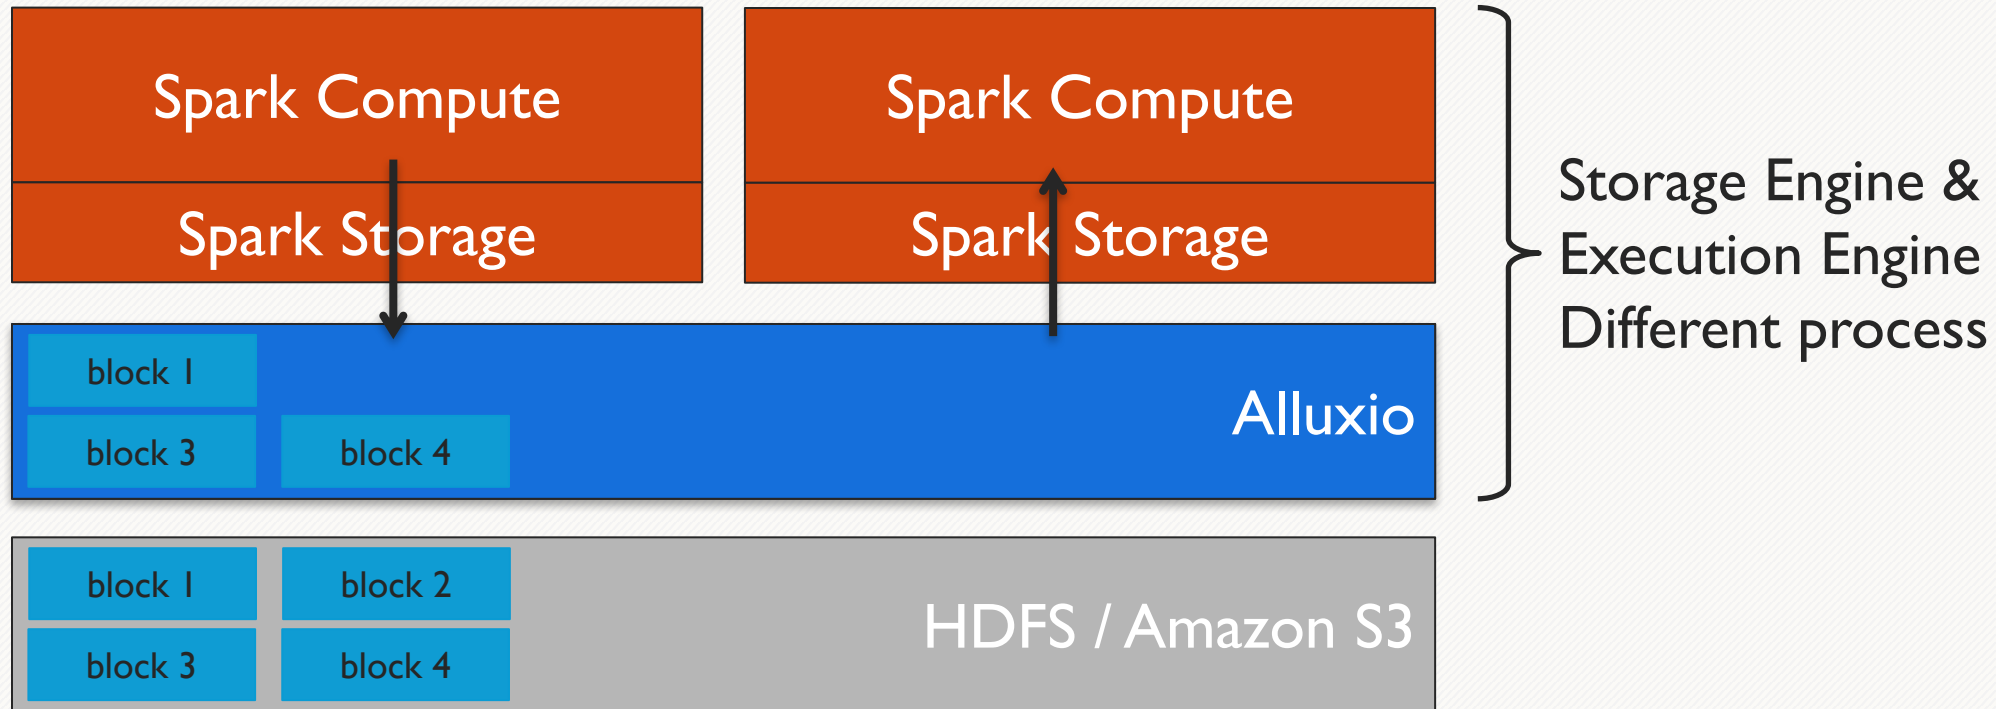
- 
- 1 Alluxio Overview
 - 2 Alluxio + Spark Use Cases
 - 3 Using Spark with Alluxio
 - 4 Performance Evaluation
 - 5 Demo

Consolidating Memory



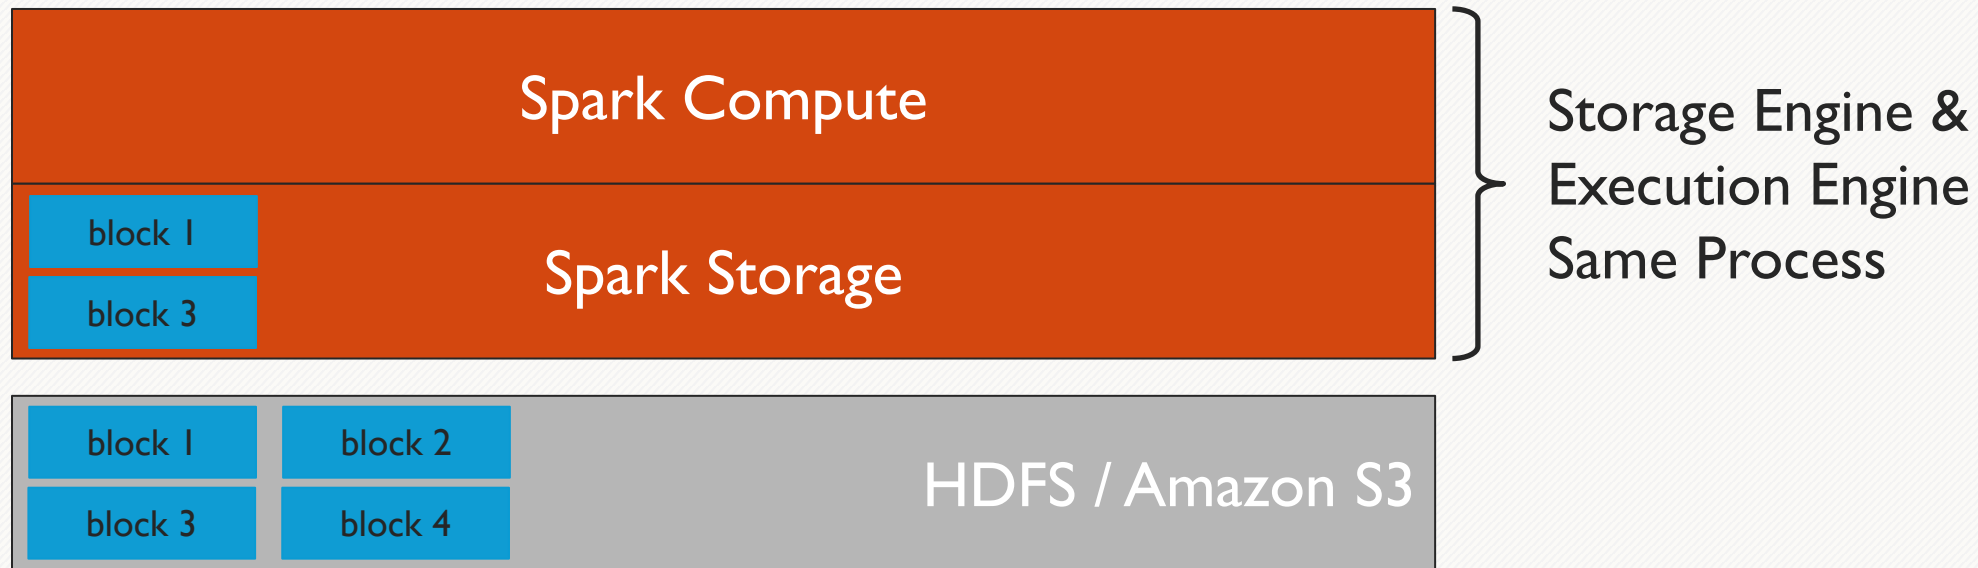
- Two copies of data in memory – double the memory used
- Inter-process Sharing Slowed Down by Network / Disk I/O

Consolidating Memory

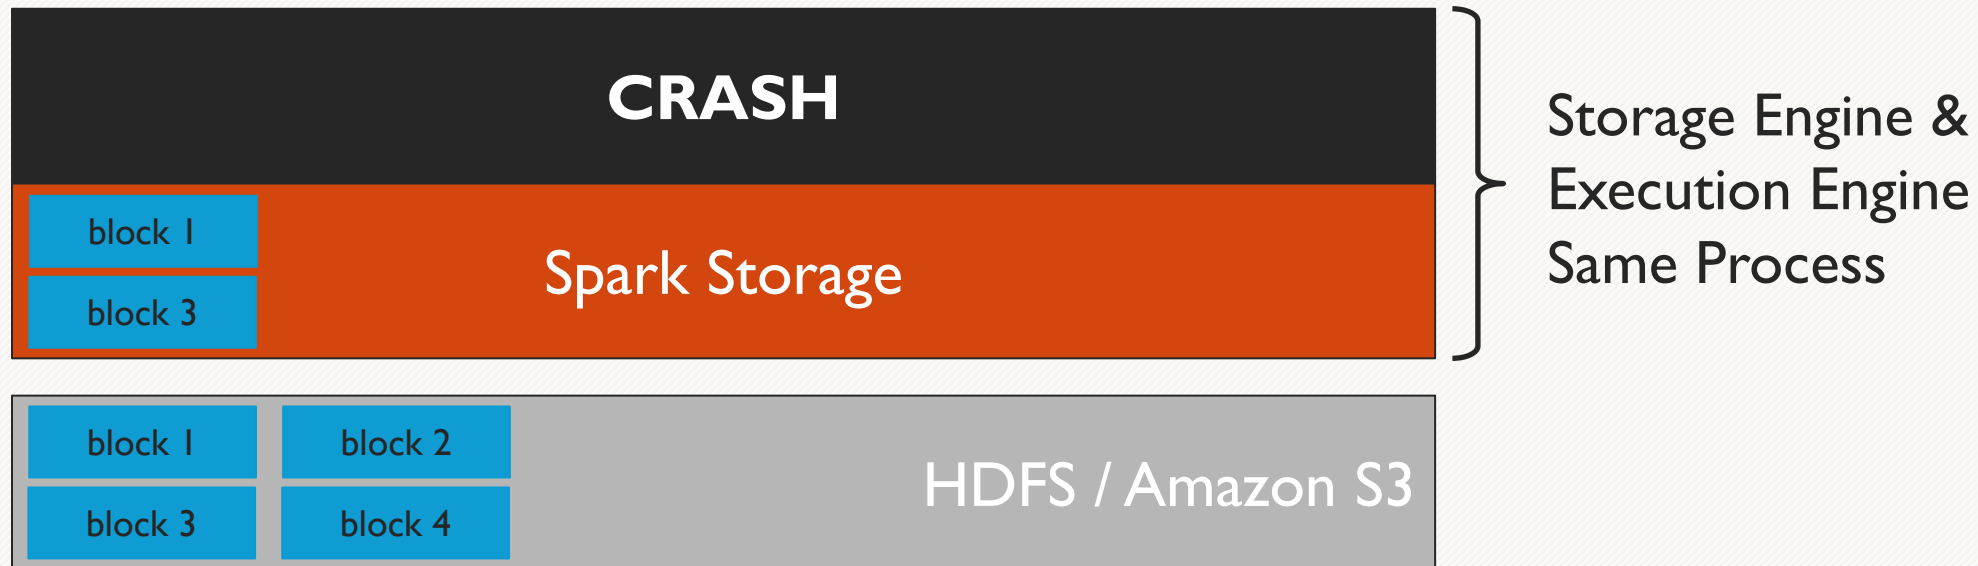


- Half the memory used
- Inter-process Sharing Happens at Memory Speed

• Data Resilience During Crash

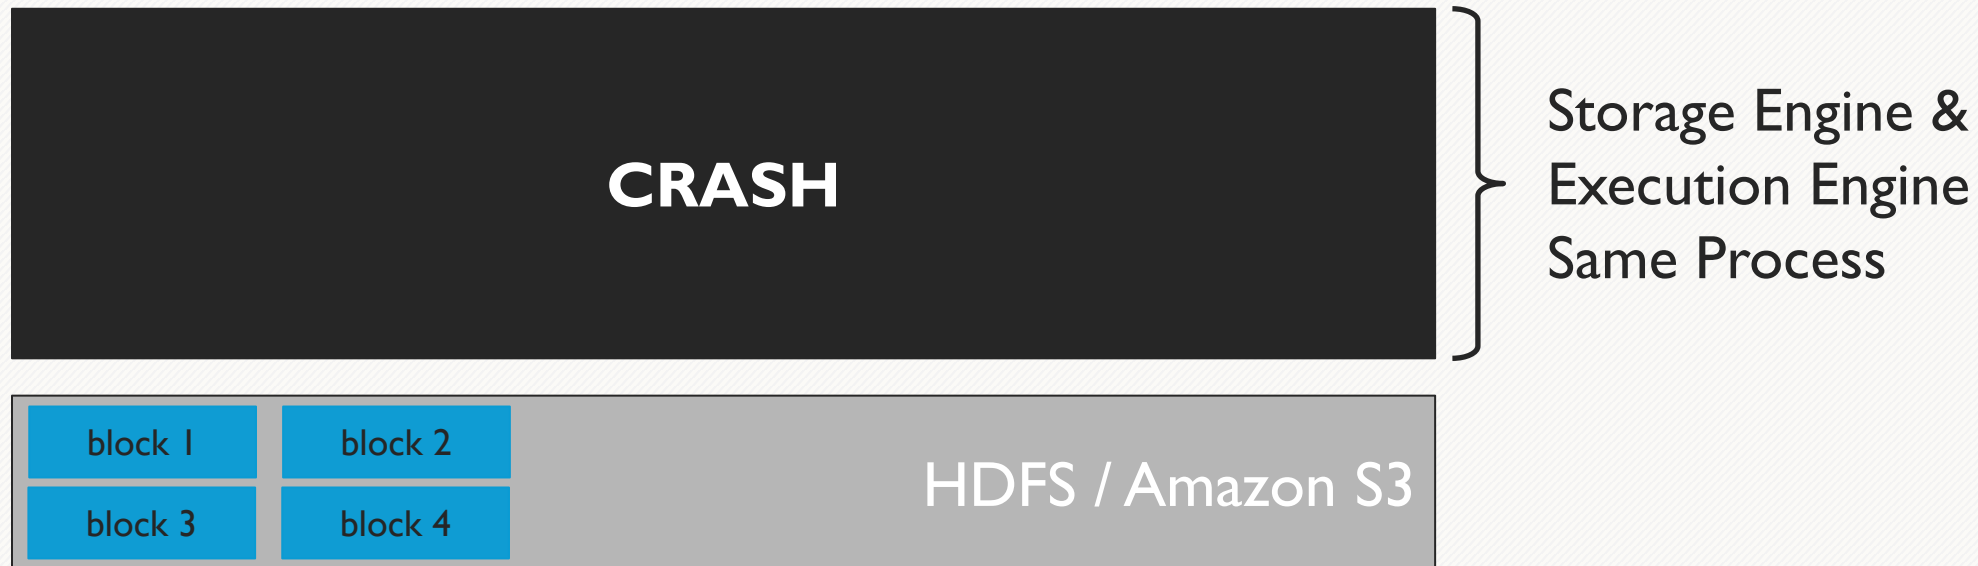


• Data Resilience During Crash



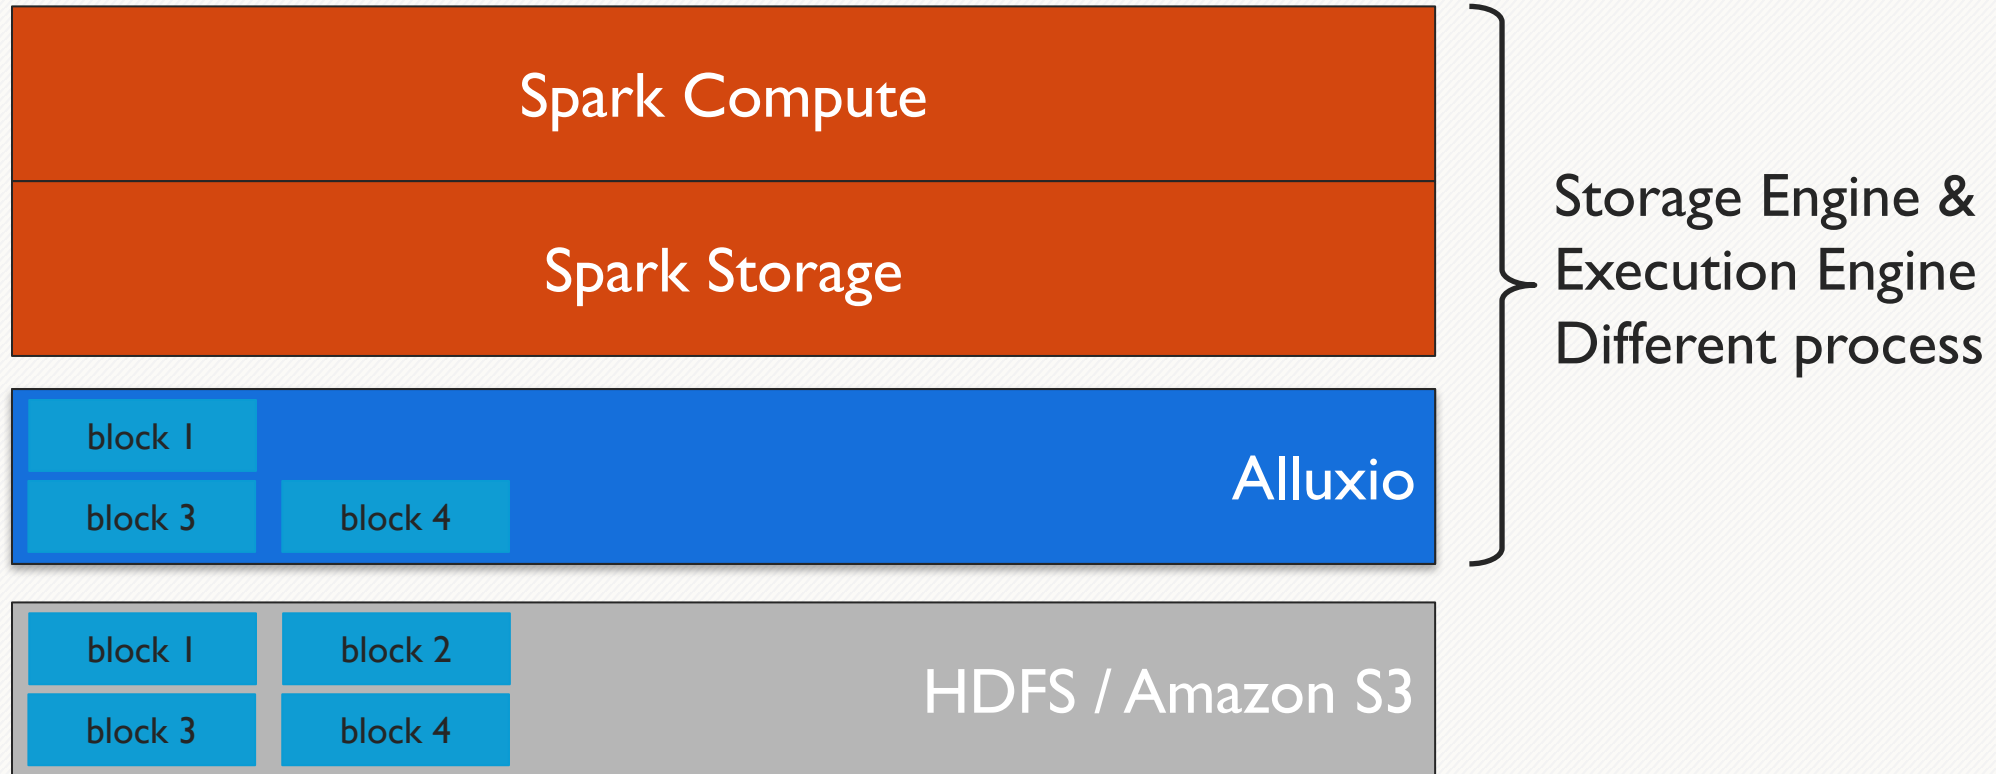
- Process Crash Requires Network and/or Disk I/O to Re-read Data

• Data Resilience During Crash

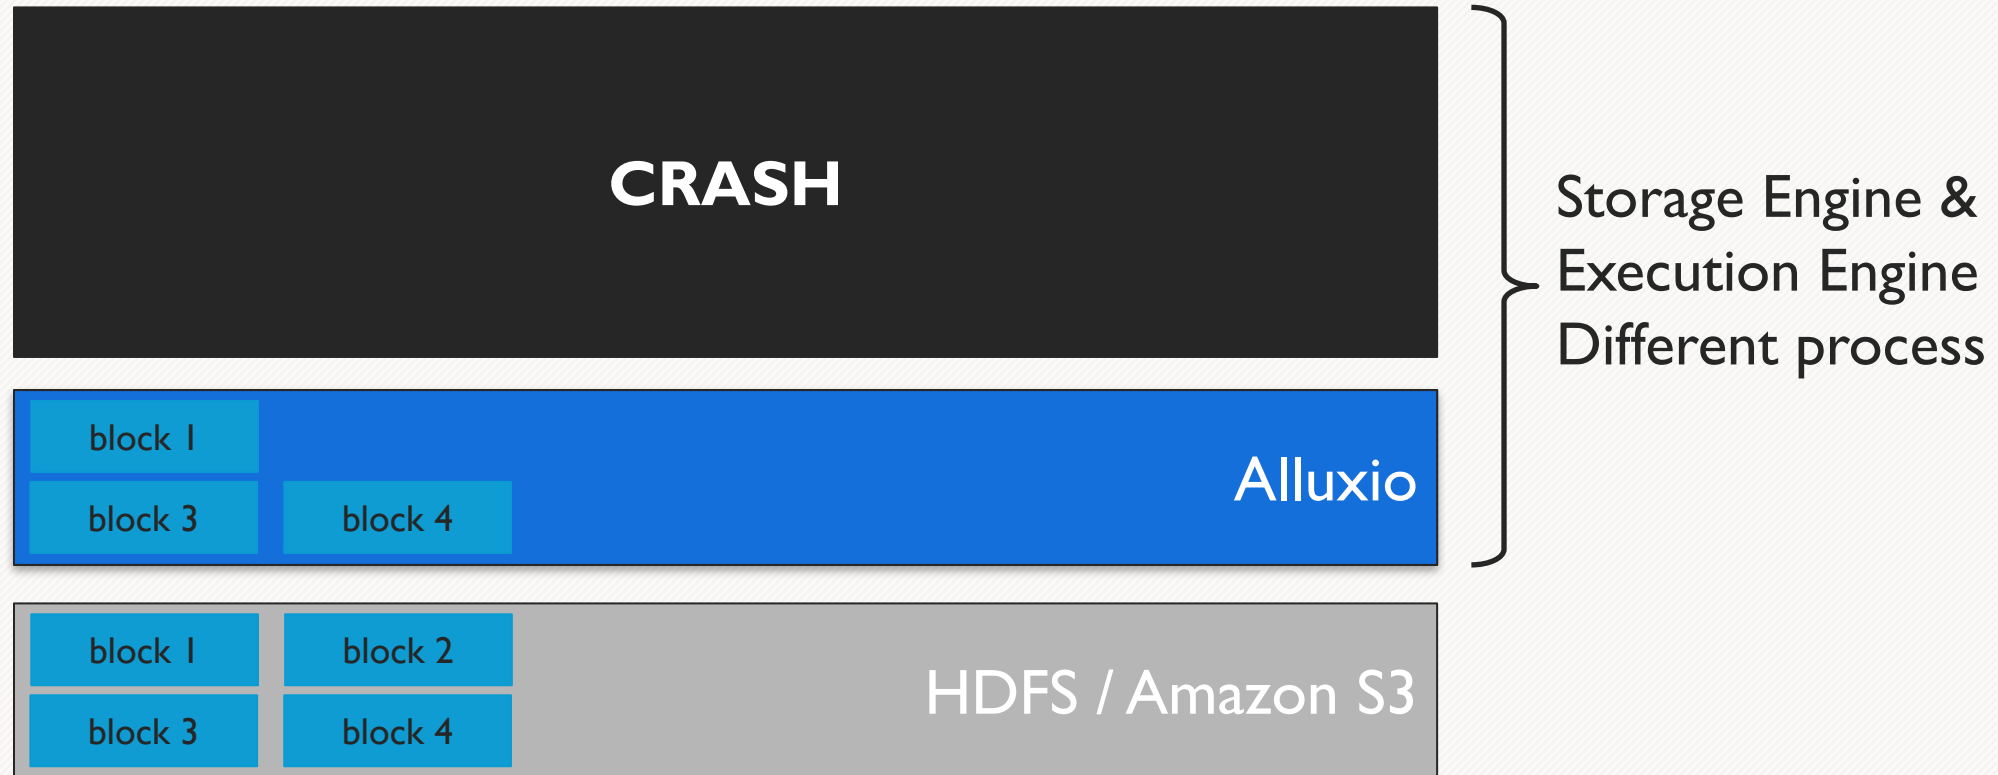


- Process Crash Requires Network and/or Disk I/O to Re-read Data

• Data Resilience During Crash



• Data Resilience During Crash



- Process Crash – Data is Re-read at Memory Speed

• Accessing Alluxio Data From Spark

Writing Data

Write to an Alluxio file

Reading Data

Read from an Alluxio file

• Code Example for Spark RDDs

Writing RDD to Alluxio

```
rdd.saveAsTextFile(alluxioPath)  
rdd.saveAsObjectFile(alluxioPath)
```

Reading RDD from Alluxio

```
rdd = sc.textFile(alluxioPath)  
rdd = sc.objectFile(alluxioPath)
```

• Code Example for Spark DataFrames

Writing to Alluxio

```
df.write.parquet(alluxioPath)
```

Reading from Alluxio

```
df = sc.read.parquet(alluxioPath)
```



Outline

- 1 Alluxio Overview
- 2 Alluxio + Spark Use Cases
- 3 Using Spark with Alluxio
- 4 Performance Evaluation
- 5 Demo

Experiments

Spark 2.0.0 + Alluxio 1.2.0

Single worker: Amazon r3.2xlarge

Comparisons:

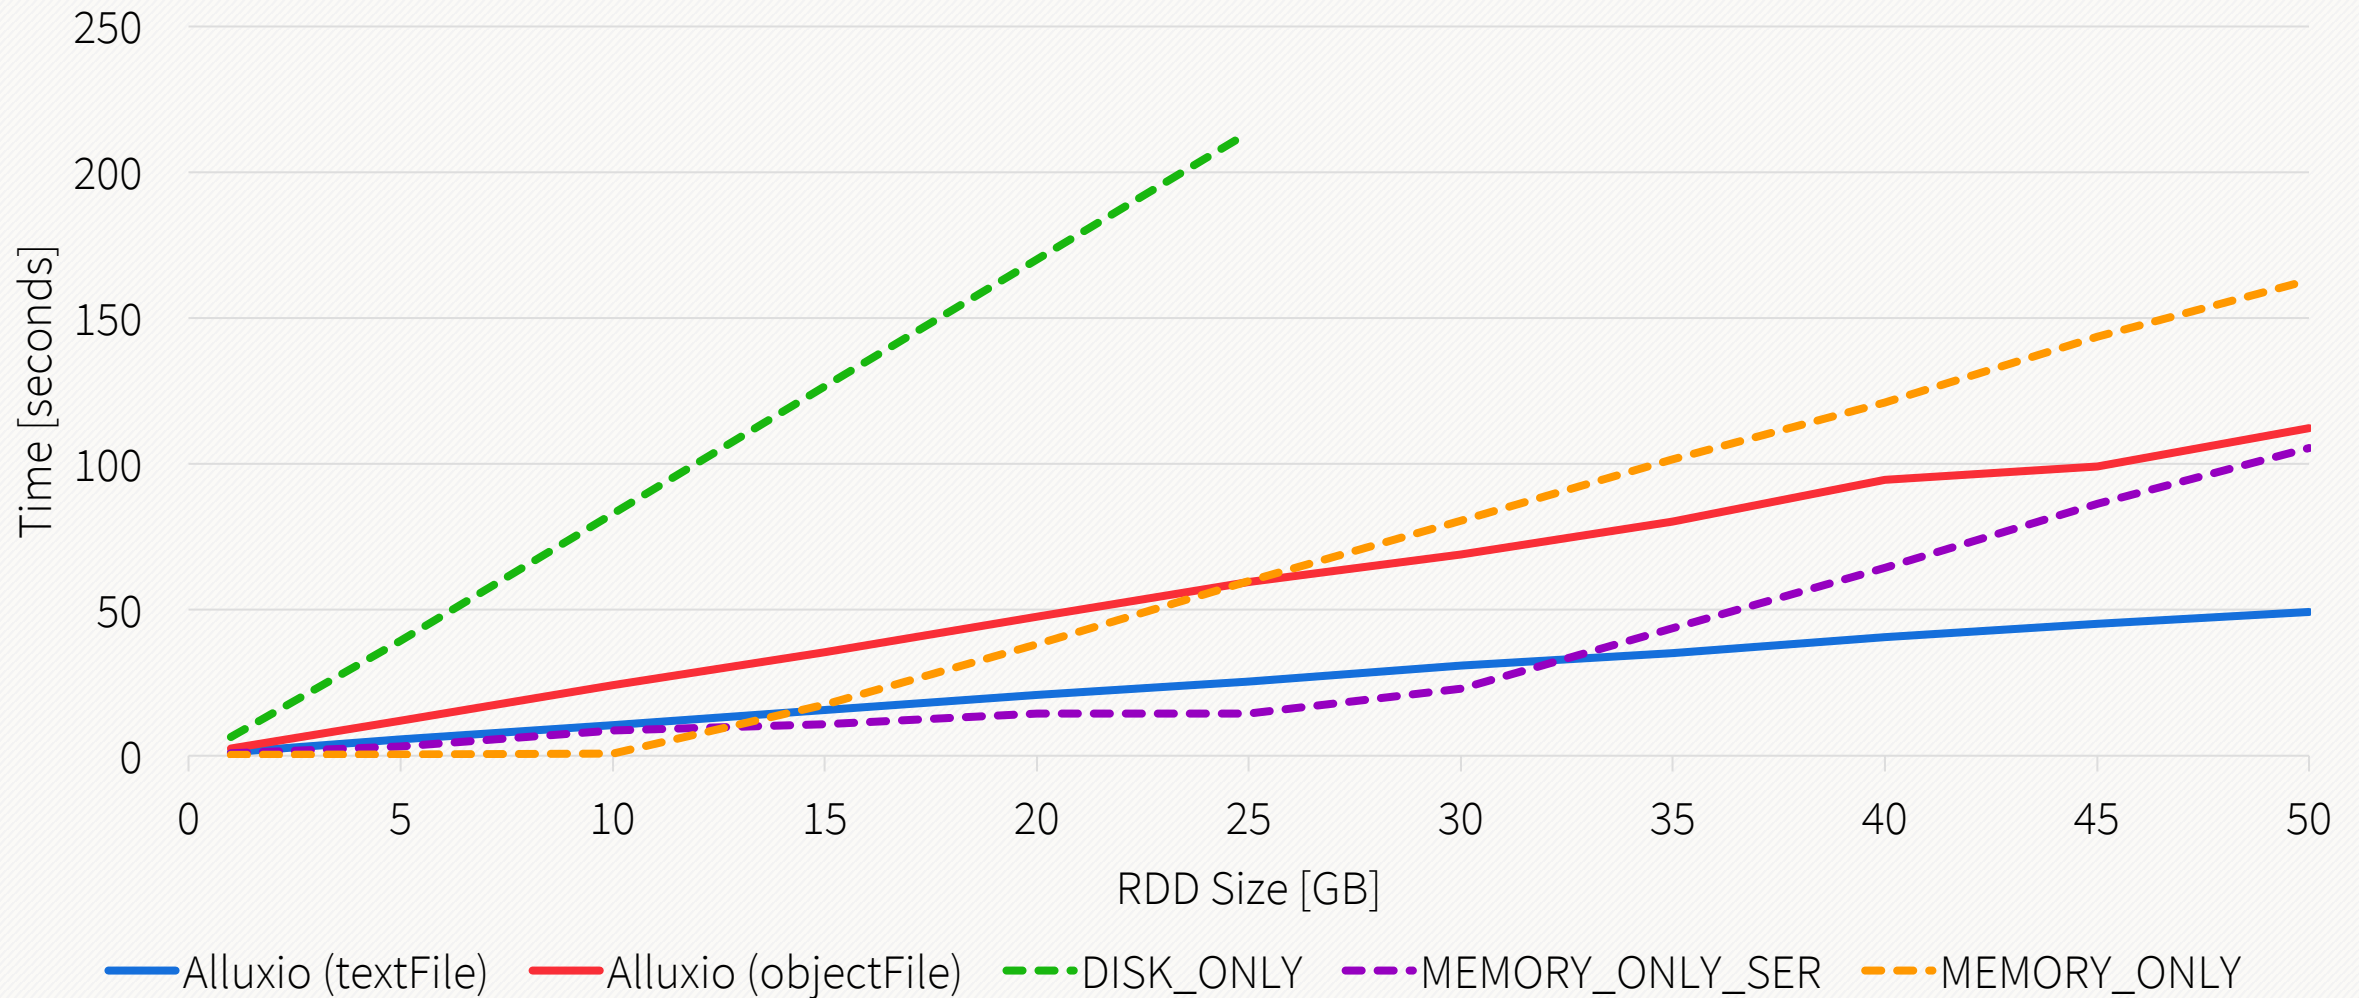
Alluxio

Spark Storage Level: MEMORY_ONLY

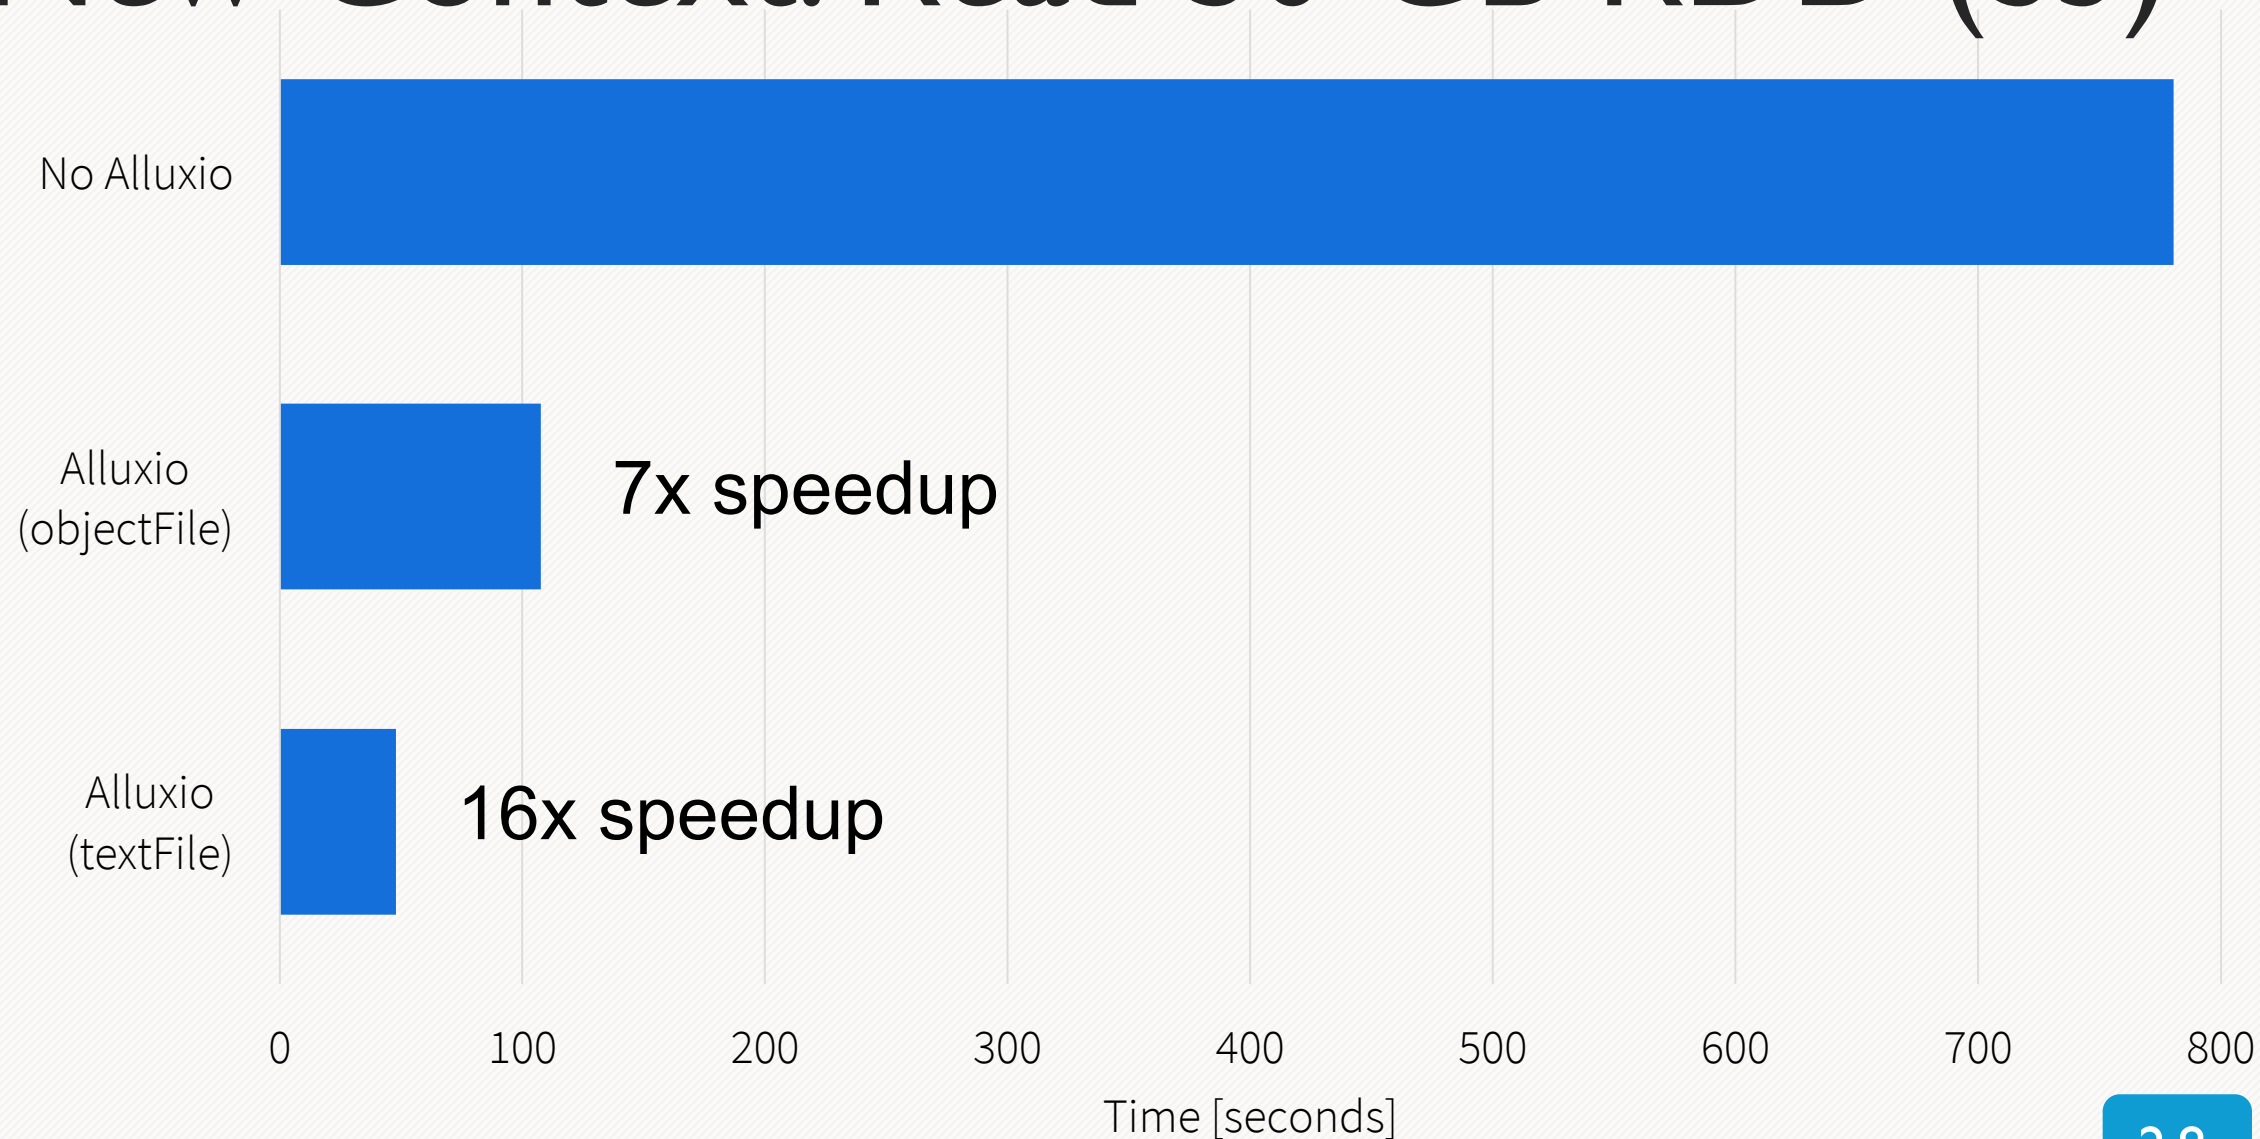
Spark Storage Level: MEMORY_ONLY_SER

Spark Storage Level: DISK_ONLY

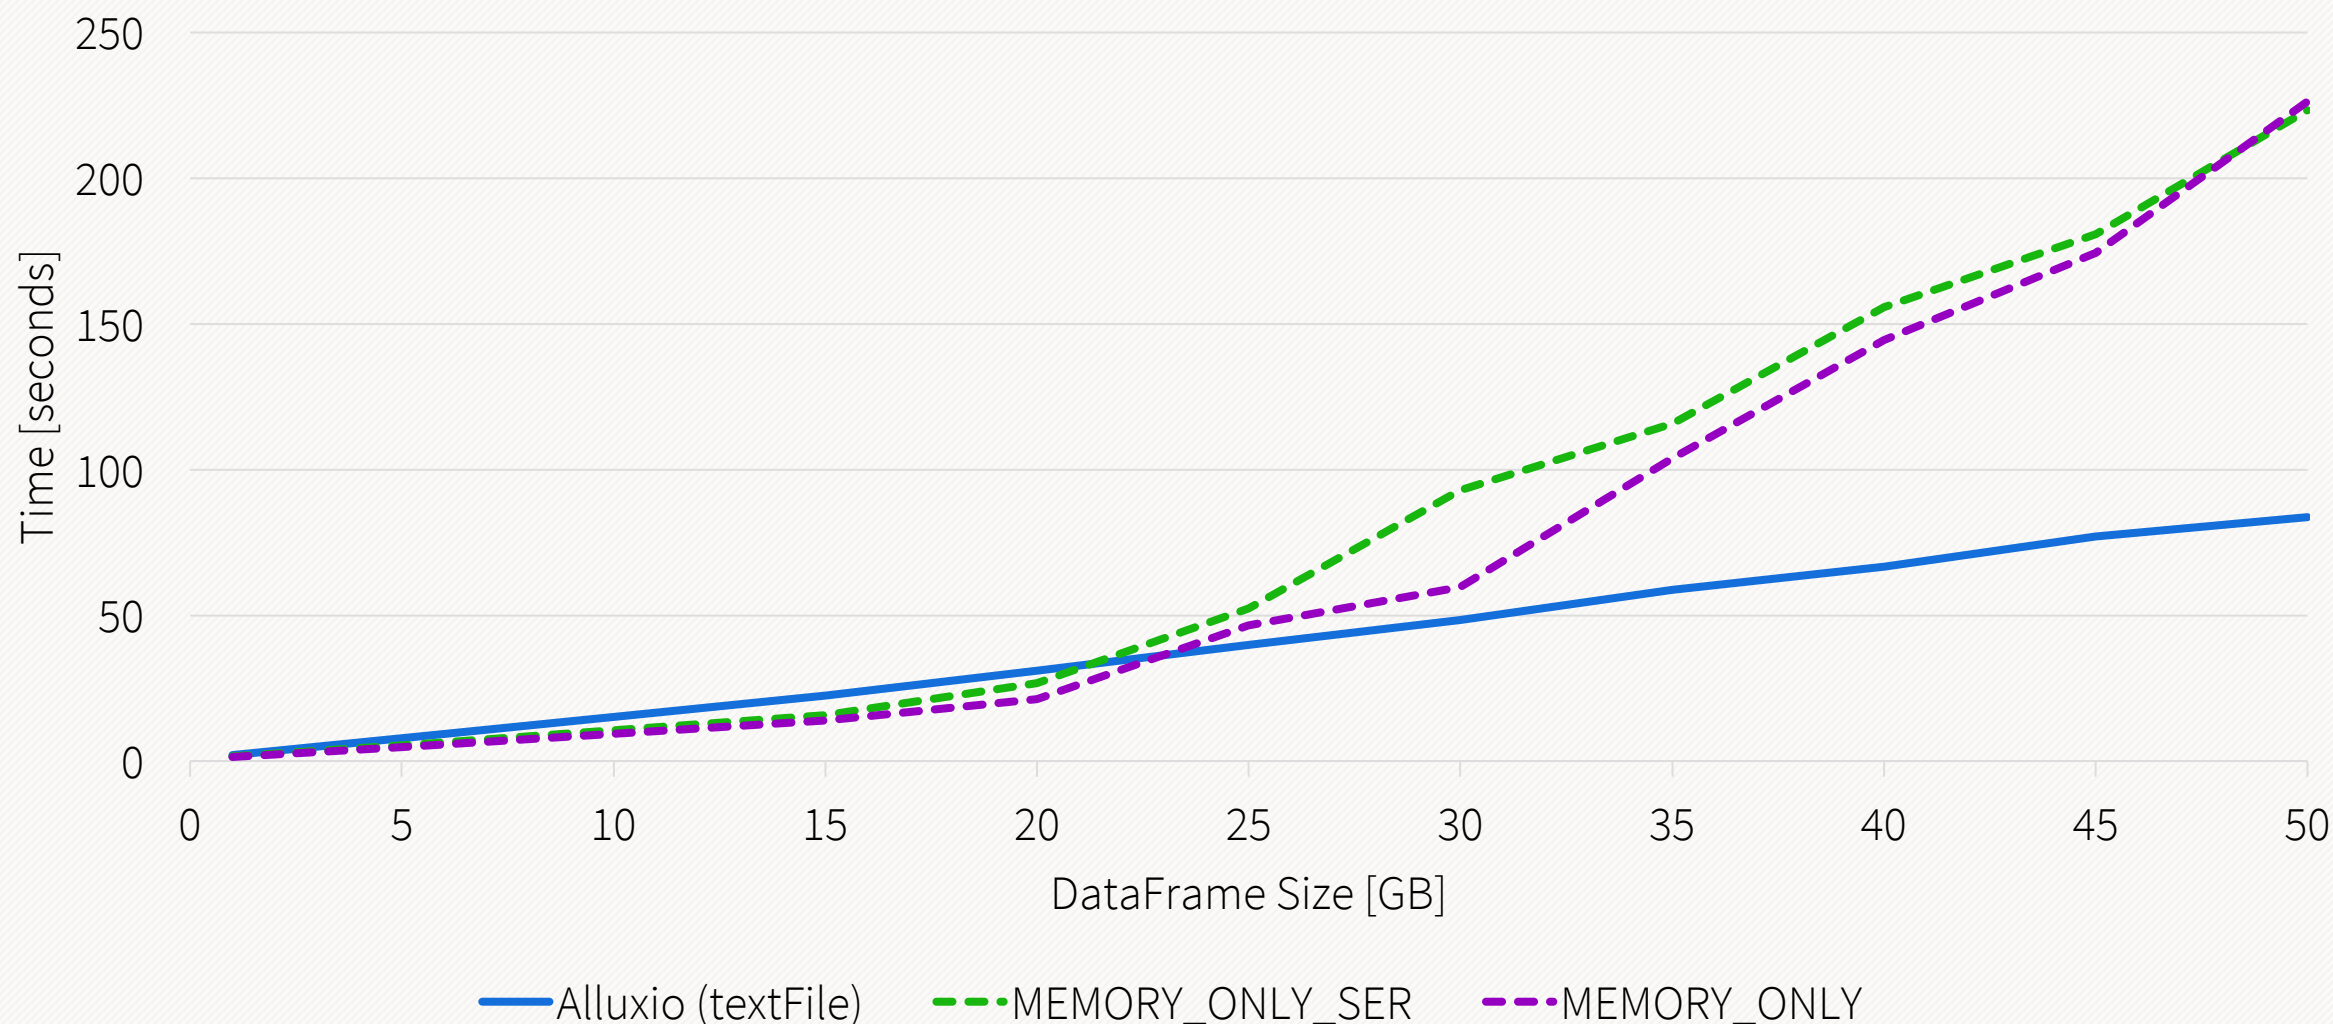
Reading Cached RDD



New Context: Read 50 GB RDD (S3)

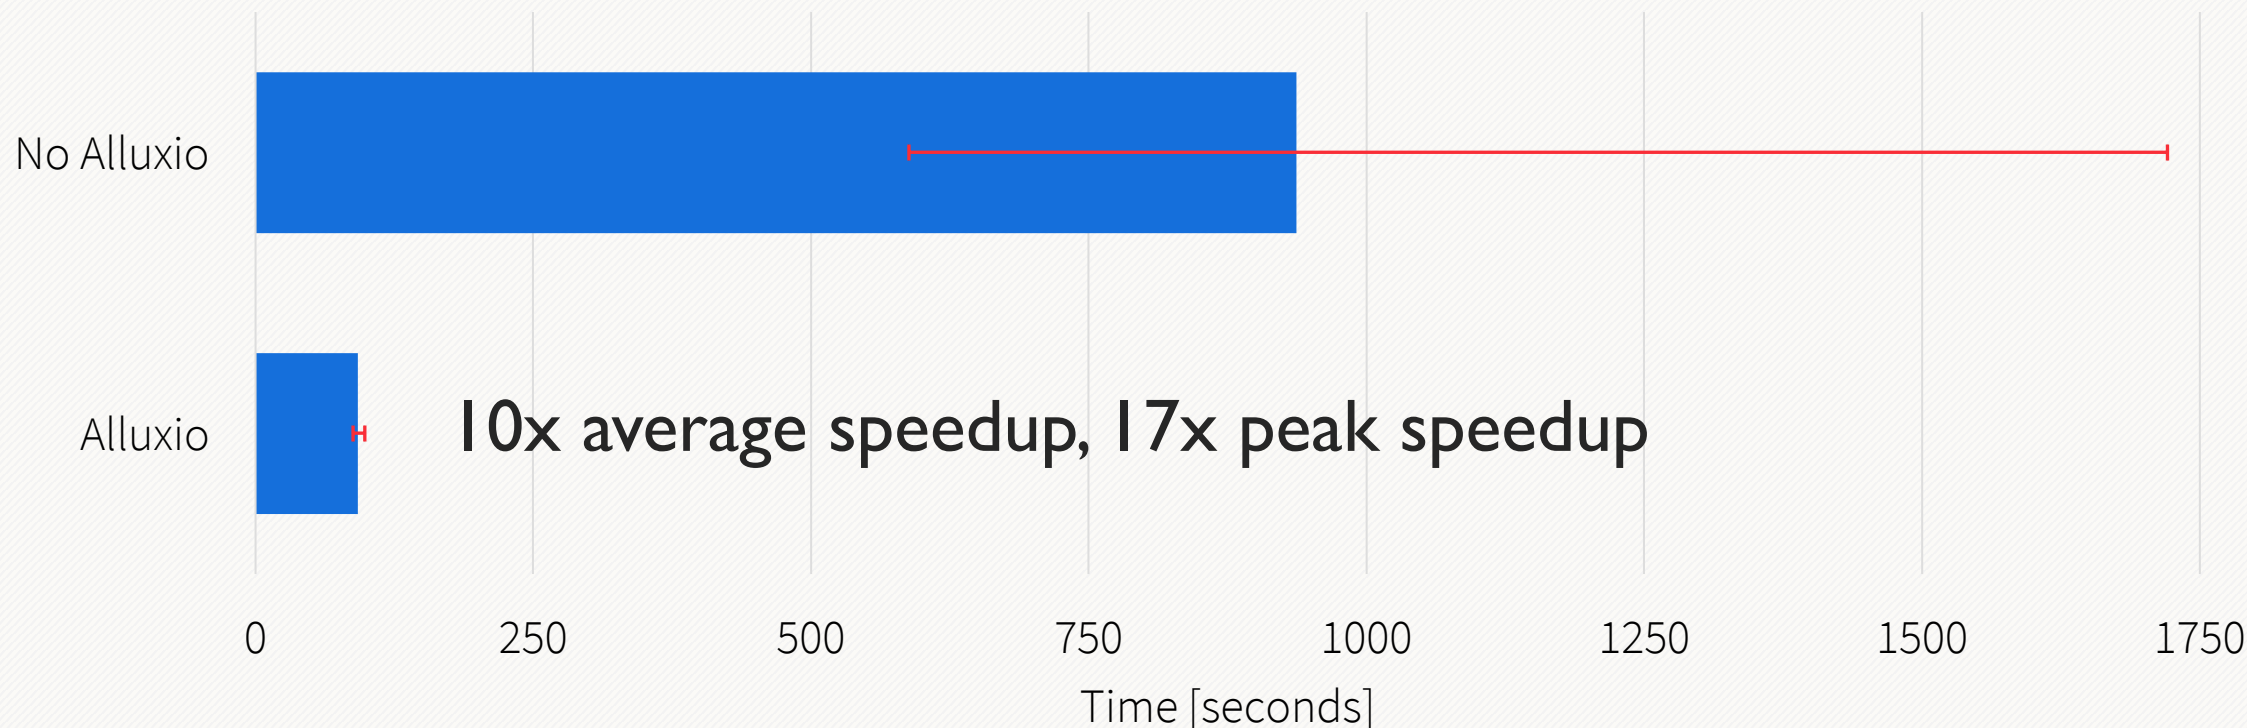


Reading Cached DataFrame (parquet)





New Context: Read 50 GB DataFrame (S3)

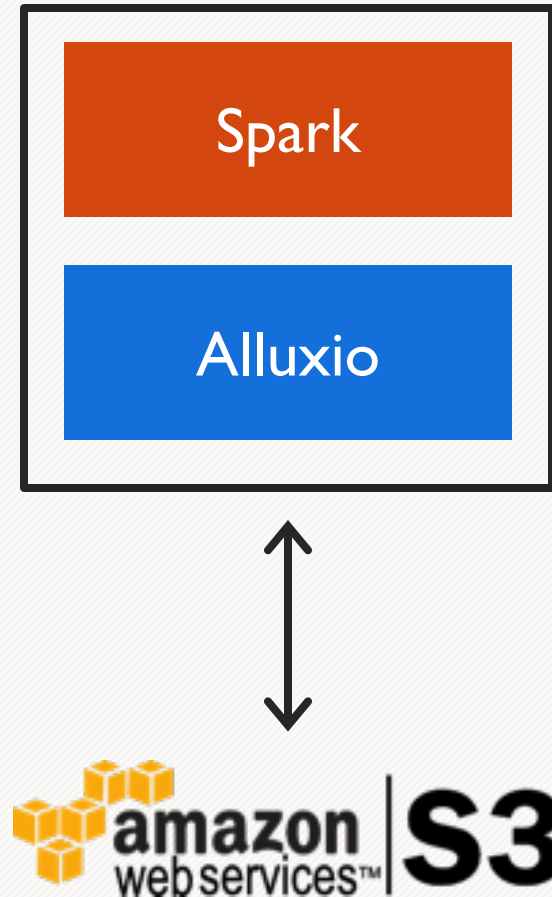




Outline

- 1 Alluxio Overview
- 2 Alluxio + Spark Use Cases
- 3 Using Spark with Alluxio
- 4 Performance Evaluation
- 5 Demo

• Demo Environment



• Conclusion

Easy to use Alluxio with Spark

Predictable and improved performance

Easily connect to various storages

Thank you!

Gene Pang
gene@alluxio.com
Twitter: @unityxx

Cheng Chang
cc@alluxio.com
Twitter: @uronce



Website

www.alluxio.com



E-mail

info@alluxio.com



Social Media

[Twitter.com/alluxio](https://twitter.com/alluxio)

[Linkedin.com/alluxio](https://www.linkedin.com/company/alluxio)