SPARK SUMMIT 2017

Scaling Spark applications by connecting code to resource consumption

Vinod Nair

Director of Product Management @ Pepperdata

Pepperdata does performance (for Big Data)

18 Thousand Production

50 Million Jobs/Year

20 0 Trillion

SPARK SUMMIT 2017

pepperdata

Today's talk will cover…

- Why debugging performance problems is hard
- Data elements needed for a complete view of application performance from separate tools
- Bringing these elements together in a single tool

2 reasons why debugging performance problems is hard

SPARK SUMMIT 2017

Reason # 1

Same external symptoms, but many possible causes

- code
- data
- configuration
- cluster weather

Reason #2

Existing tools provide limited visibility

- Spark Web UI is the most popular
  - Execution plan with some aggregate performance data
- Ganglia, Ambari, CM etc
  - Time series data about cluster, not specific to Spark apps
- Code execution not connected to resource consumption
- Unhealthy hardware or load from other apps unaccounted

3 data elements form a complete picture of Spark application performance

1. Code execution plan
   - Indicates which block of code is being executed
2. Time series view
   - Visual of resource consumption of application
   - Outliers in resource use very easy to detect
3. Cluster weather
   - A view of all applications that runs on the cluster
   - A view of the health of all the nodes in the cluster

First half of the solution

Spark Web UI

SPARK SUMMIT 2017

Logical code execution plan from Spark: Jobs / Stages / DAG

Spark 1.6.0 | Jobs | Stages | Storage | Environment | Executors | **ScalaPageRank** application UI

Stages for All Jobs

Completed Stages: 6

Completed Stages (6)

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
5	saveAsHadoopFile at IOCommon.scala:63	+details	2017/05/29 12:36:53	23 s	8/8		176.4 MB	547.3 MB	
4	flatMap at SparkPageRank.scala:62	+details	2017/05/29 12:29:45	7.1 min	8/8	10.4 GB		879.6 MB	547.3 MB
3	flatMap at SparkPageRank.scala:62	+details	2017/05/29 12:25:32	4.2 min	8/8	10.4 GB		829.8 MB	547.3 MB
2	flatMap at SparkPageRank.scala:62	+details	2017/05/29 12:18:59	6.5 min	8/8	10.4 GB		2.9 GB	497.5 MB
1	distinct at SparkPageRank.scala:58	+details	2017/05/29 12:13:00	6.0 min	24/24			2.7 GB	2.6 GB
0	distinct at SparkPageRank.scala:58	+details	2017/05/29 12:06:55	6.1 min	24/24	2.8 GB			2.7 GB

Physical execution plan from Spark: Executors / Tasks

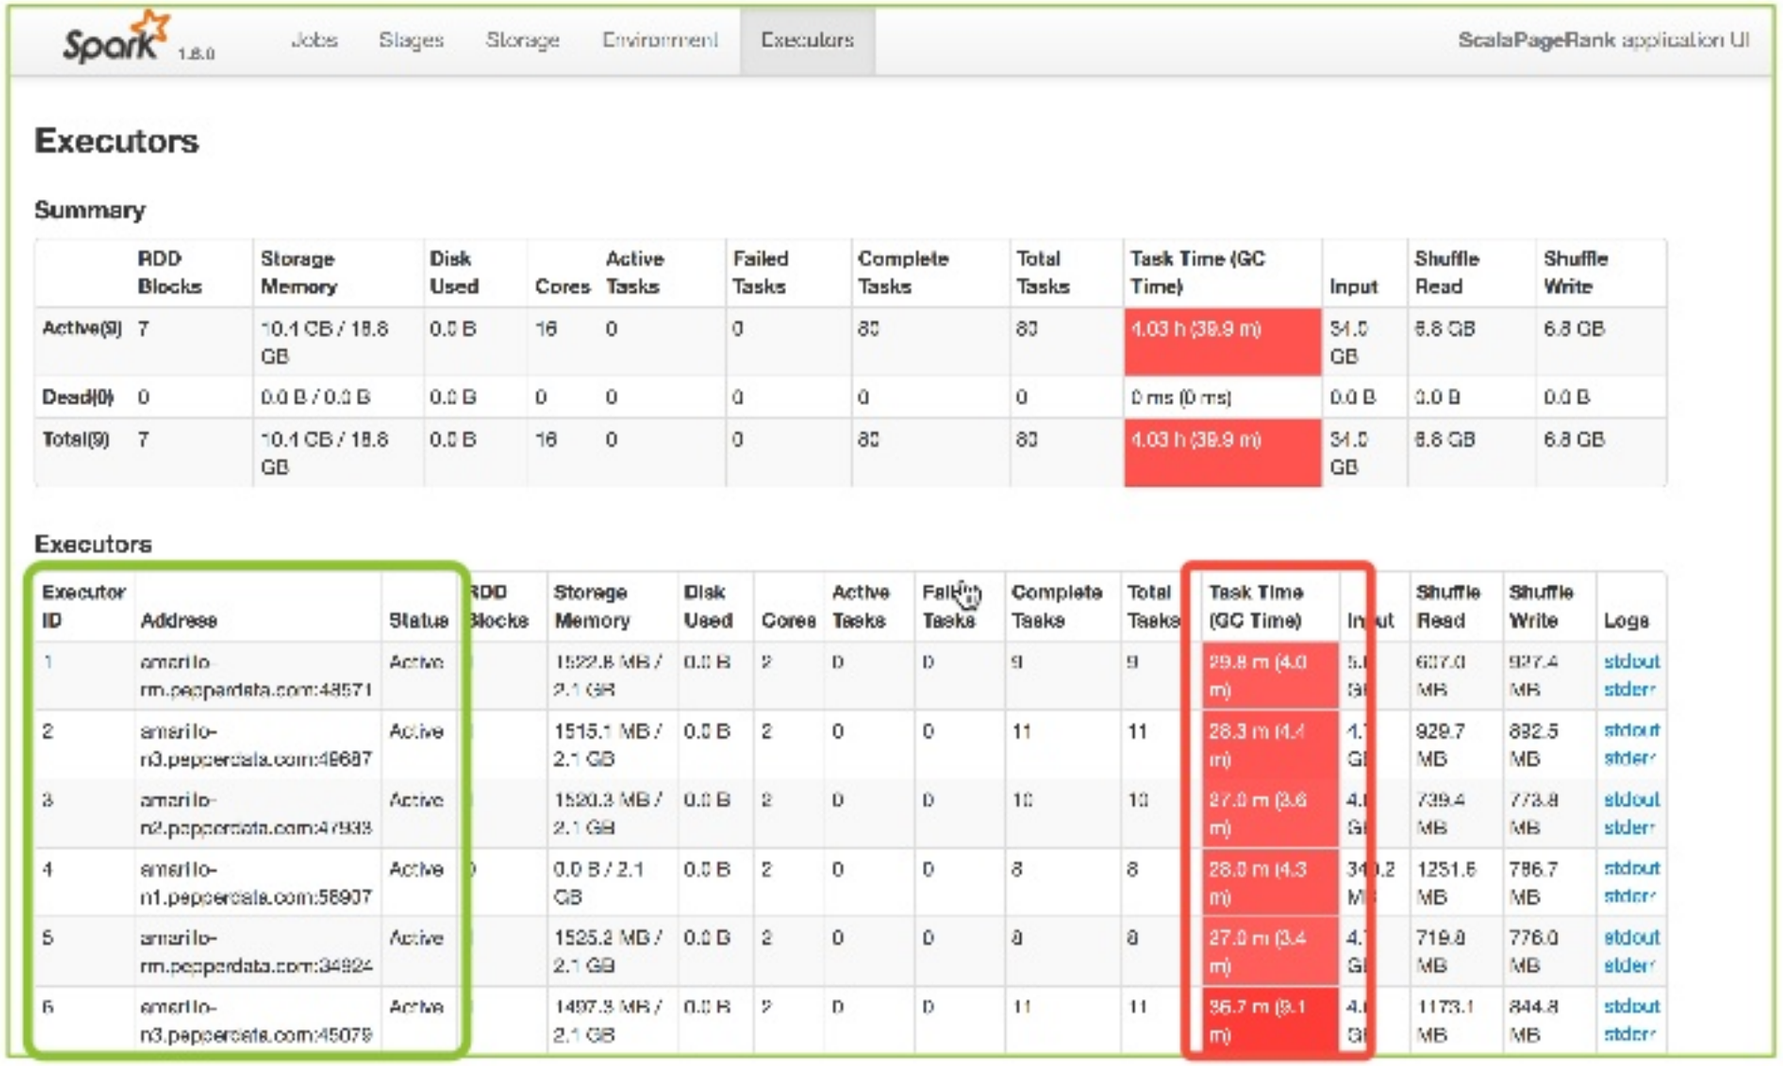

SPARK SUMMIT 2017

Second half of solution

Time series view

SPARK SUMMIT 2017

Time series view of resource consumption for the App

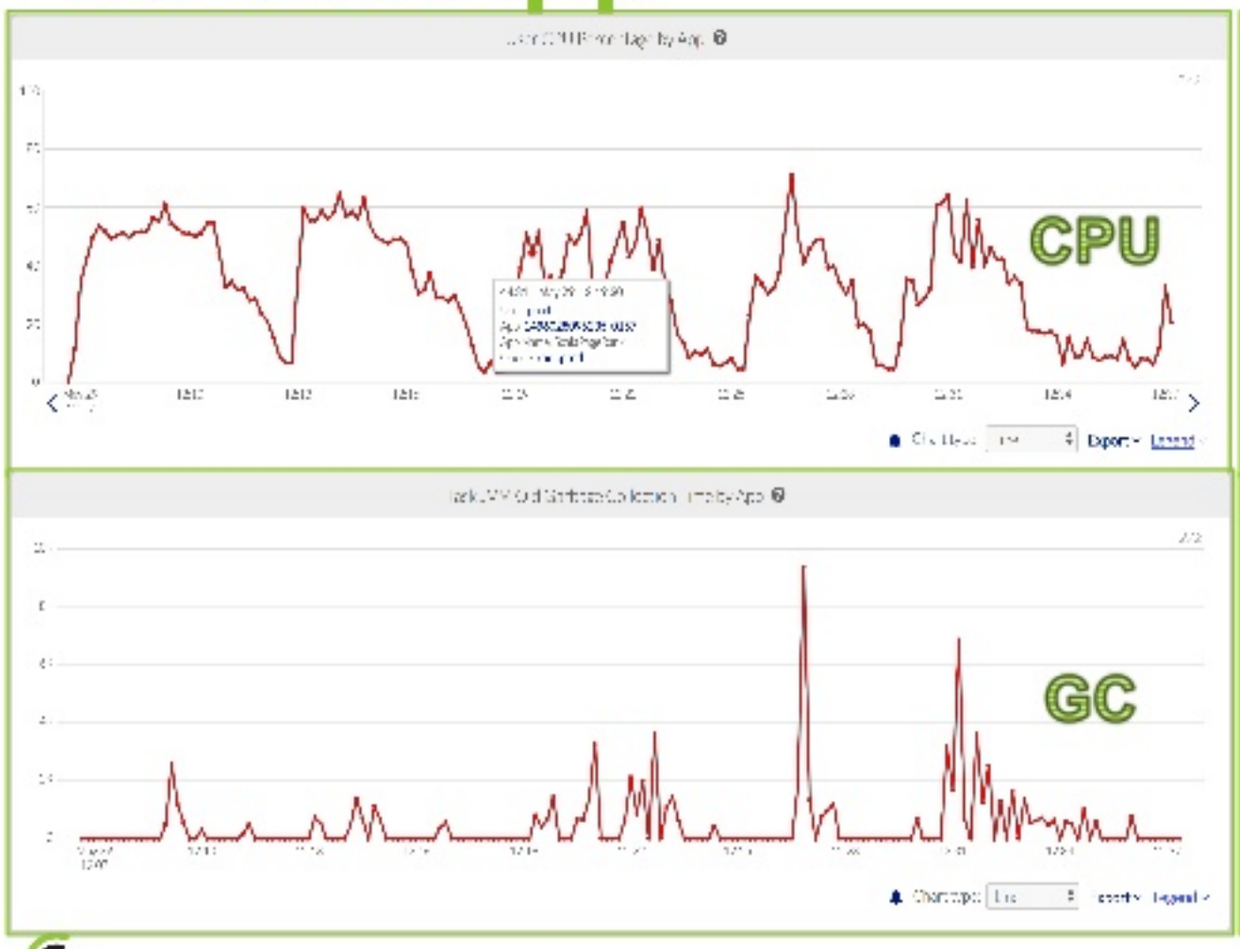

SPARK SUMMIT 2017

Best of both worlds

Bring them together

SPARK SUMMIT 2017

Bringing it all together: CPU across Stages

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

CPU | Memory | Non Heap Memory | GC | HDFS Bytes Read

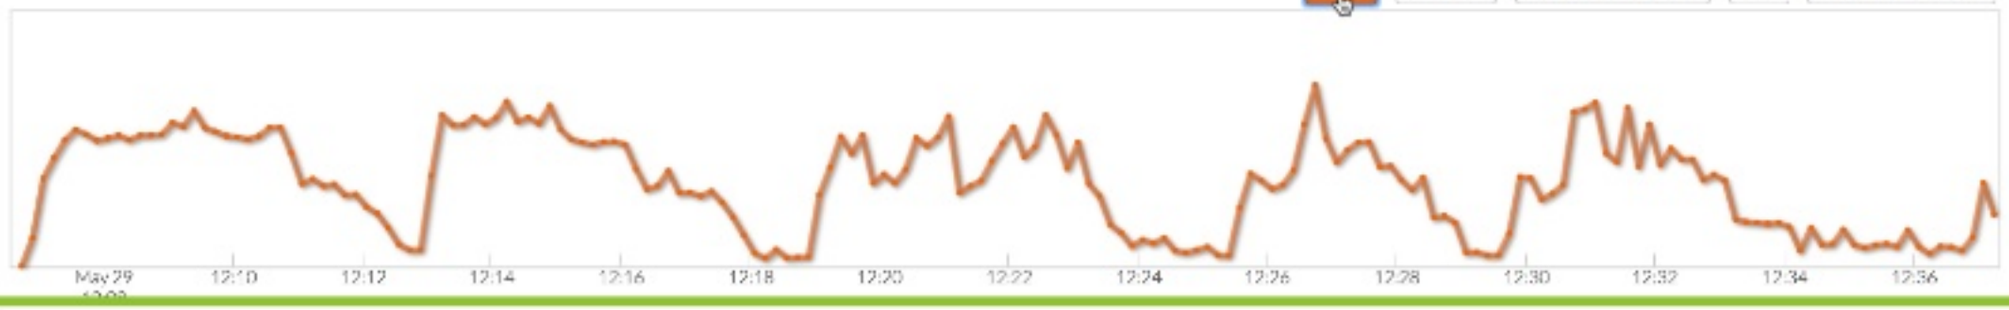

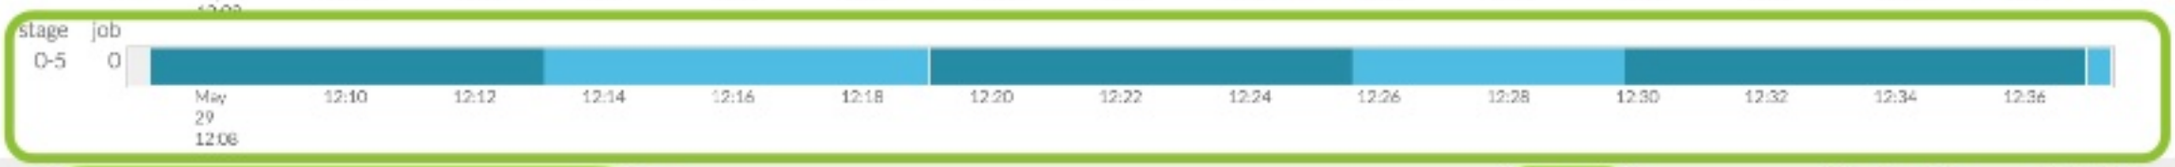

Stage	Job	File Name	Stage Start Time	Stage Duration (hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:05	41.5	23,327	497	5.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

Memory across all Stages of App

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

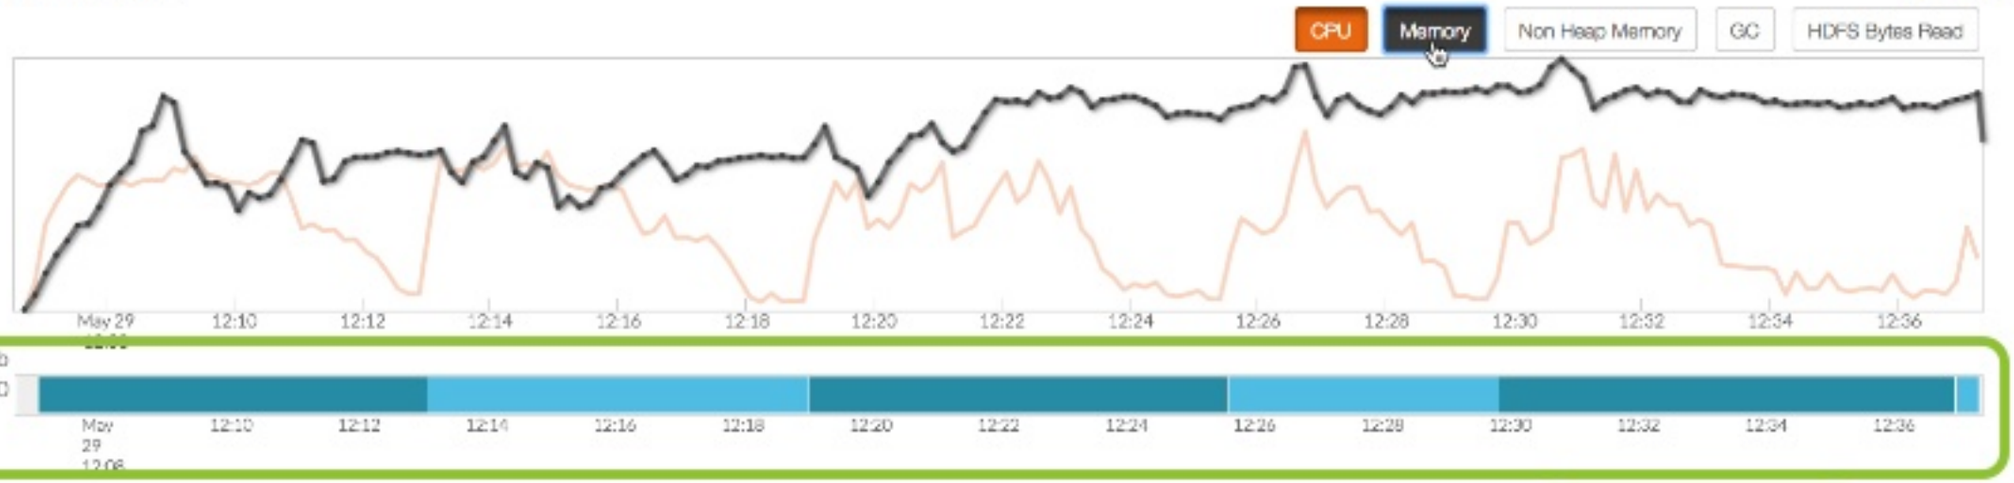

Stage	Job	File Name	Stage Start Time	Stage Duration ((hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:05	41.5	23,327	497	5.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

GC across all Stages of App

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

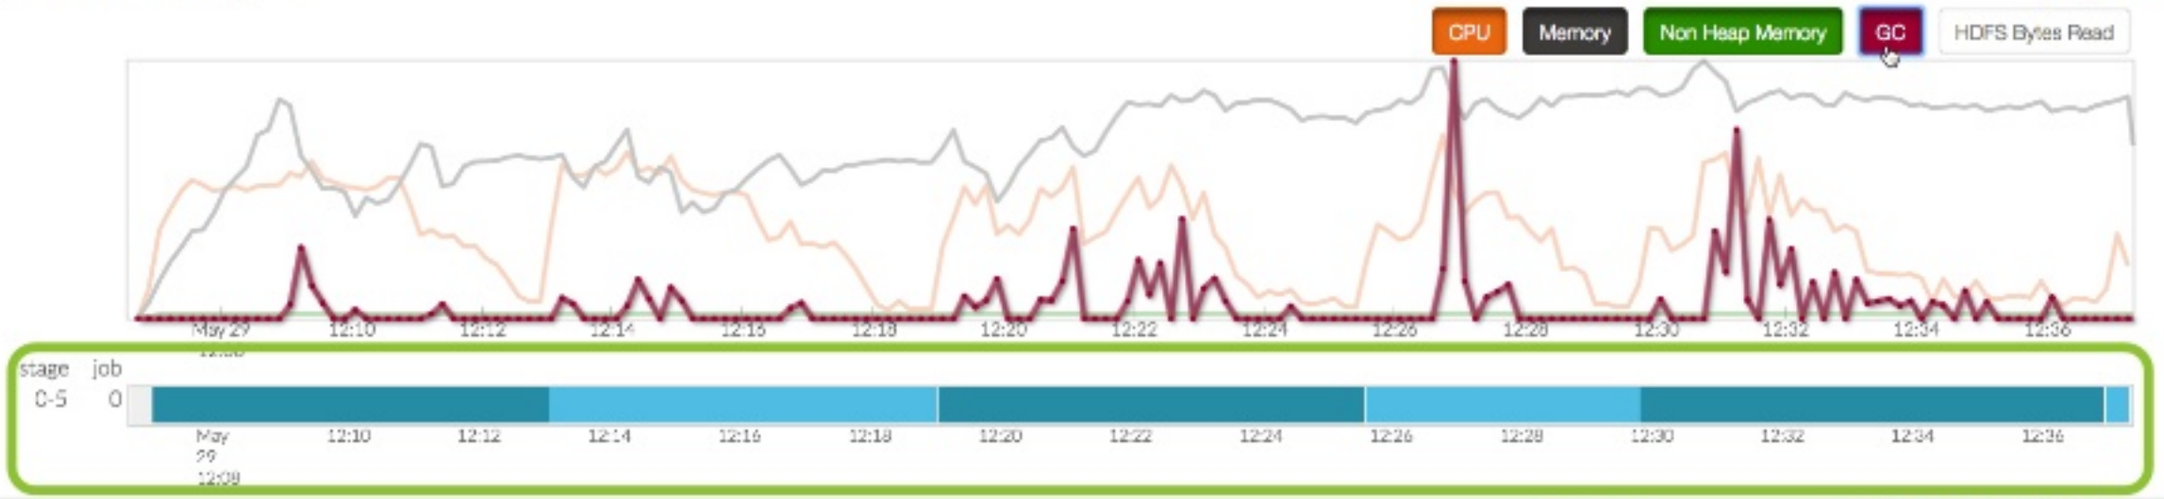

Stage	Job	File Name	Stage Start Time	Stage Duration ((hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:05	41.5	23,327	497	5.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

HDFS Reads across all Stages of App

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

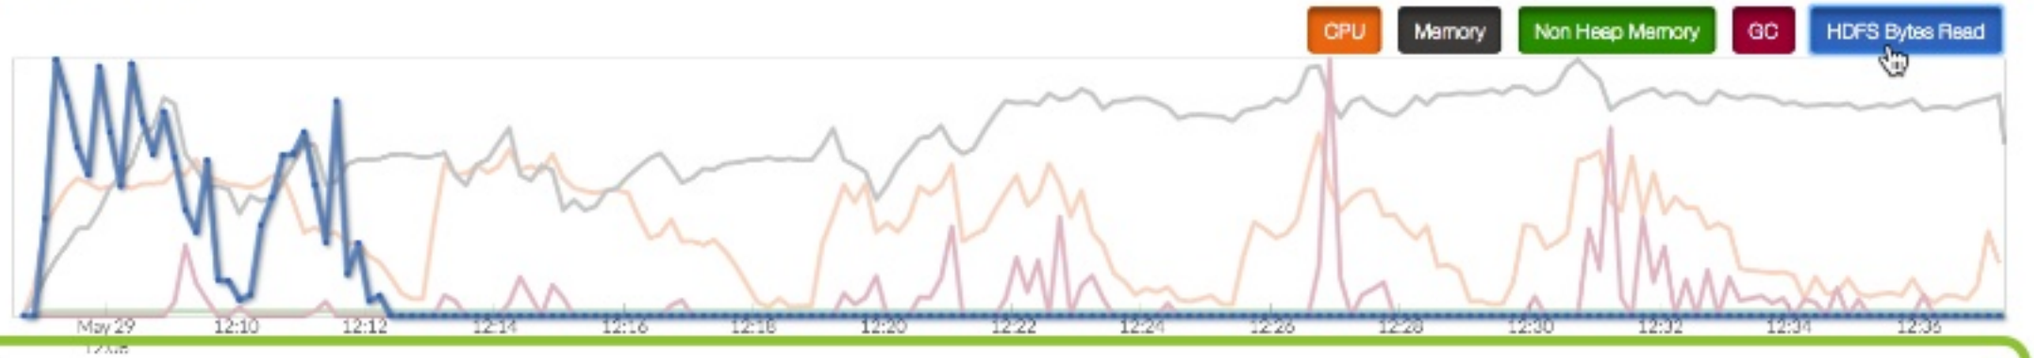

Stage	Job	File Name	Stage Start Time	Stage Duration ((hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:05	41.5	23,327	497	5.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

Bringing it all together

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

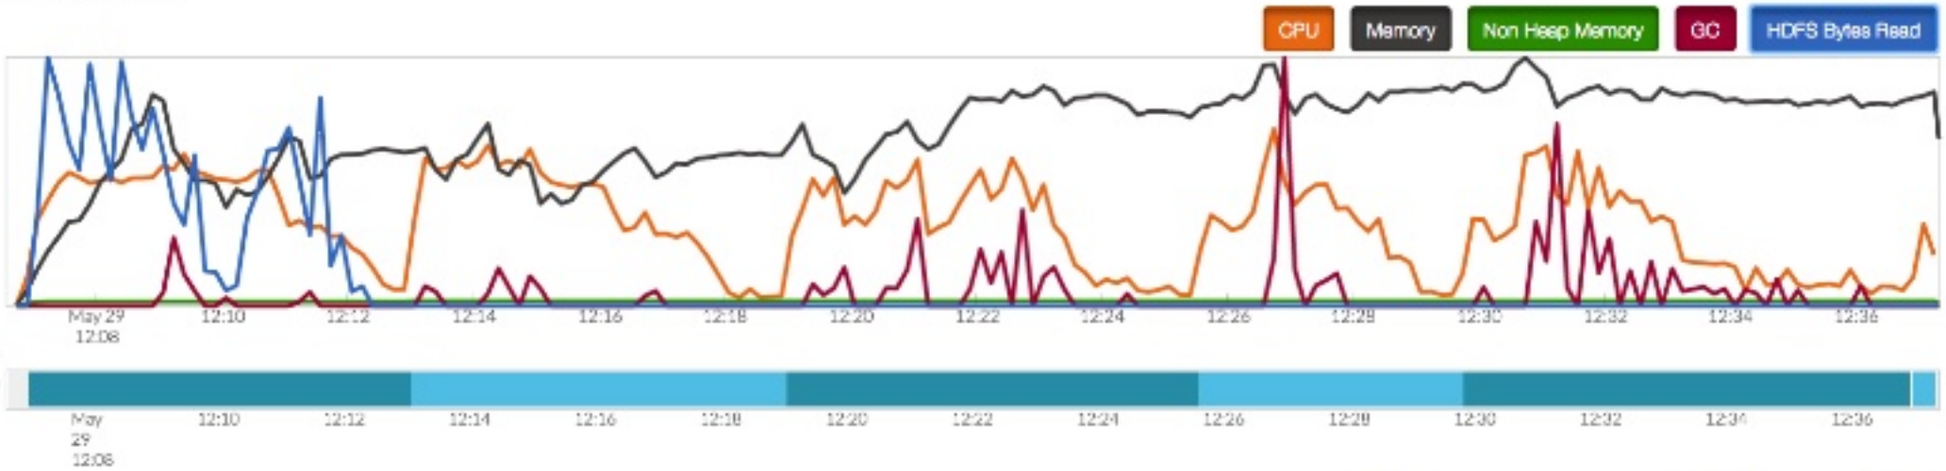

Stage	Job	File Name	Stage Start Time	Stage Duration (hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:05	41.5	23,327	497	5.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

Let's examine GC activity in Stage 4

ScalaPageRank

App Id: 1496028096235_0157

Cluster Load View

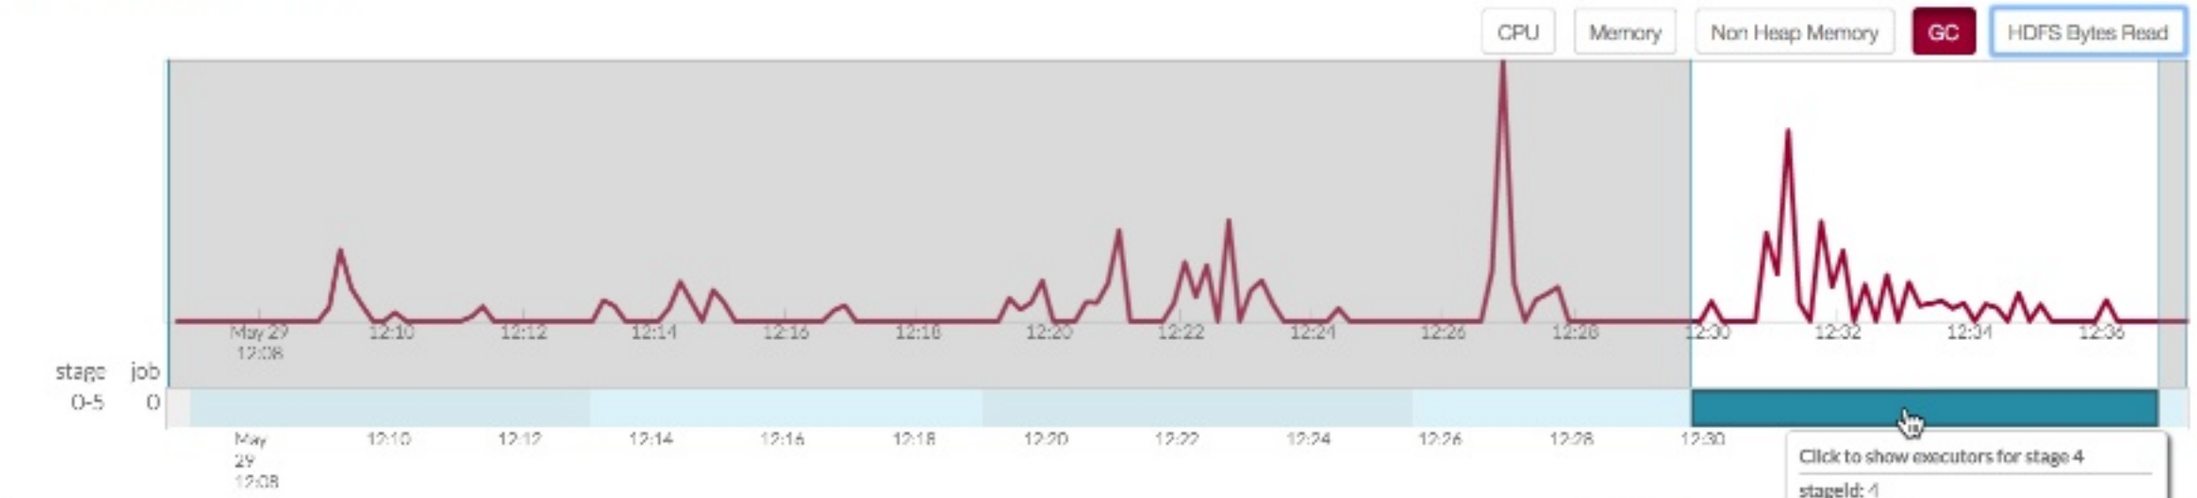

Stage	Job	File Name	Stage Start Time	Stage Duration ((hh:]mm:ss)	Average Executor CPU (%)	Stage Pe: Memc		GC ime nds)	Parallel Stages
0	0	distinct at SparkPageRank.scala:58	2017/05/29-12:06	06:06	41.5			6.82	none
1	0	distinct at SparkPageRank.scala:58	2017/05/29-12:13	05:58	37.8	20,106	507	6.6	none
2	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:18	06:32	32.1	24,256	514	22.2	none
3	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:25	04:12	31.6	26,676	517	15.5	none
4	0	flatMap at SparkPageRank.scala:62	2017/05/29-12:29	07:08	26	27,414	518	31.3	none
5	0	saveAsHadoopFile at IOCommon.scala:63	2017/05/29-12:36	00:22	26.7	23,639	524	0	none

Executor skew increased Stage duration 2x

SPARK SUMMIT 2017

Possible solution: increase number of partitions

Cluster weather

What if it’s not your fault?

How does cluster weather impact your app ?

Spark 1.6.0 **History Server**

Event log directory: hdfs:///user/spark/applicationHistory

Showing 101-120 of 2045 < 1 ... 4 5 6 7 8 ... 103 >

App ID	App Name	Started	Completed	Duration	Spark User	Last Updated
application_1495423352820_0377	ScalaTeraSort	2017/05/23 18:08:26	2017/05/23 18:32:48	24 min	prod	2017/05/23 18:32:48
[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]
application_1495423352820_0373	ScalaTeraSort	2017/05/23 16:50:49	2017/05/23 17:10:59	20 min	prod	2017/05/23 17:11:00
[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]	[illegible]
application_1495423352820_0369	ScalaTeraSort	2017/05/23 14:47:59	2017/05/23 15:53:38	1.1 h	prod	2017/05/23 15:53:38

SPARK SUMMIT 2017

No apparent reason for delay from Spark Web UI

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	saveAsNewAPIHadoopFile at ScalaTeraSort.scala:60	+details	2017/05/23 16:55:02	16 min	8/8		29.8 GB	13.1 GB	
1	map at ScalaTeraSort.scala:49	+details	2017/05/23 16:51:46	3.3 min	240/240	29.1 GB			13.1 GB
0	BaseRangePartitioner at ScalaTeraSort.scala:56	+details	2017/05/23 16:51:30	15 s	8/8				

Stage Id	Description		Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	saveAsNewAPIHadoopFile at ScalaTeraSort.scala:60	+details	2017/05/23 15:12:33	41 min	8/8		29.8 GB	13.1 GB	
1	map at ScalaTeraSort.scala:49	+details	2017/05/23 14:50:46	22 min	240/240	29.8 GB			13.1 GB
0	BaseRangePartitioner at ScalaTeraSort.scala:56	+details	2017/05/23 14:49:35	32 s	8/8	768.0 MB			

Time series shows slower run of App with much lower resources

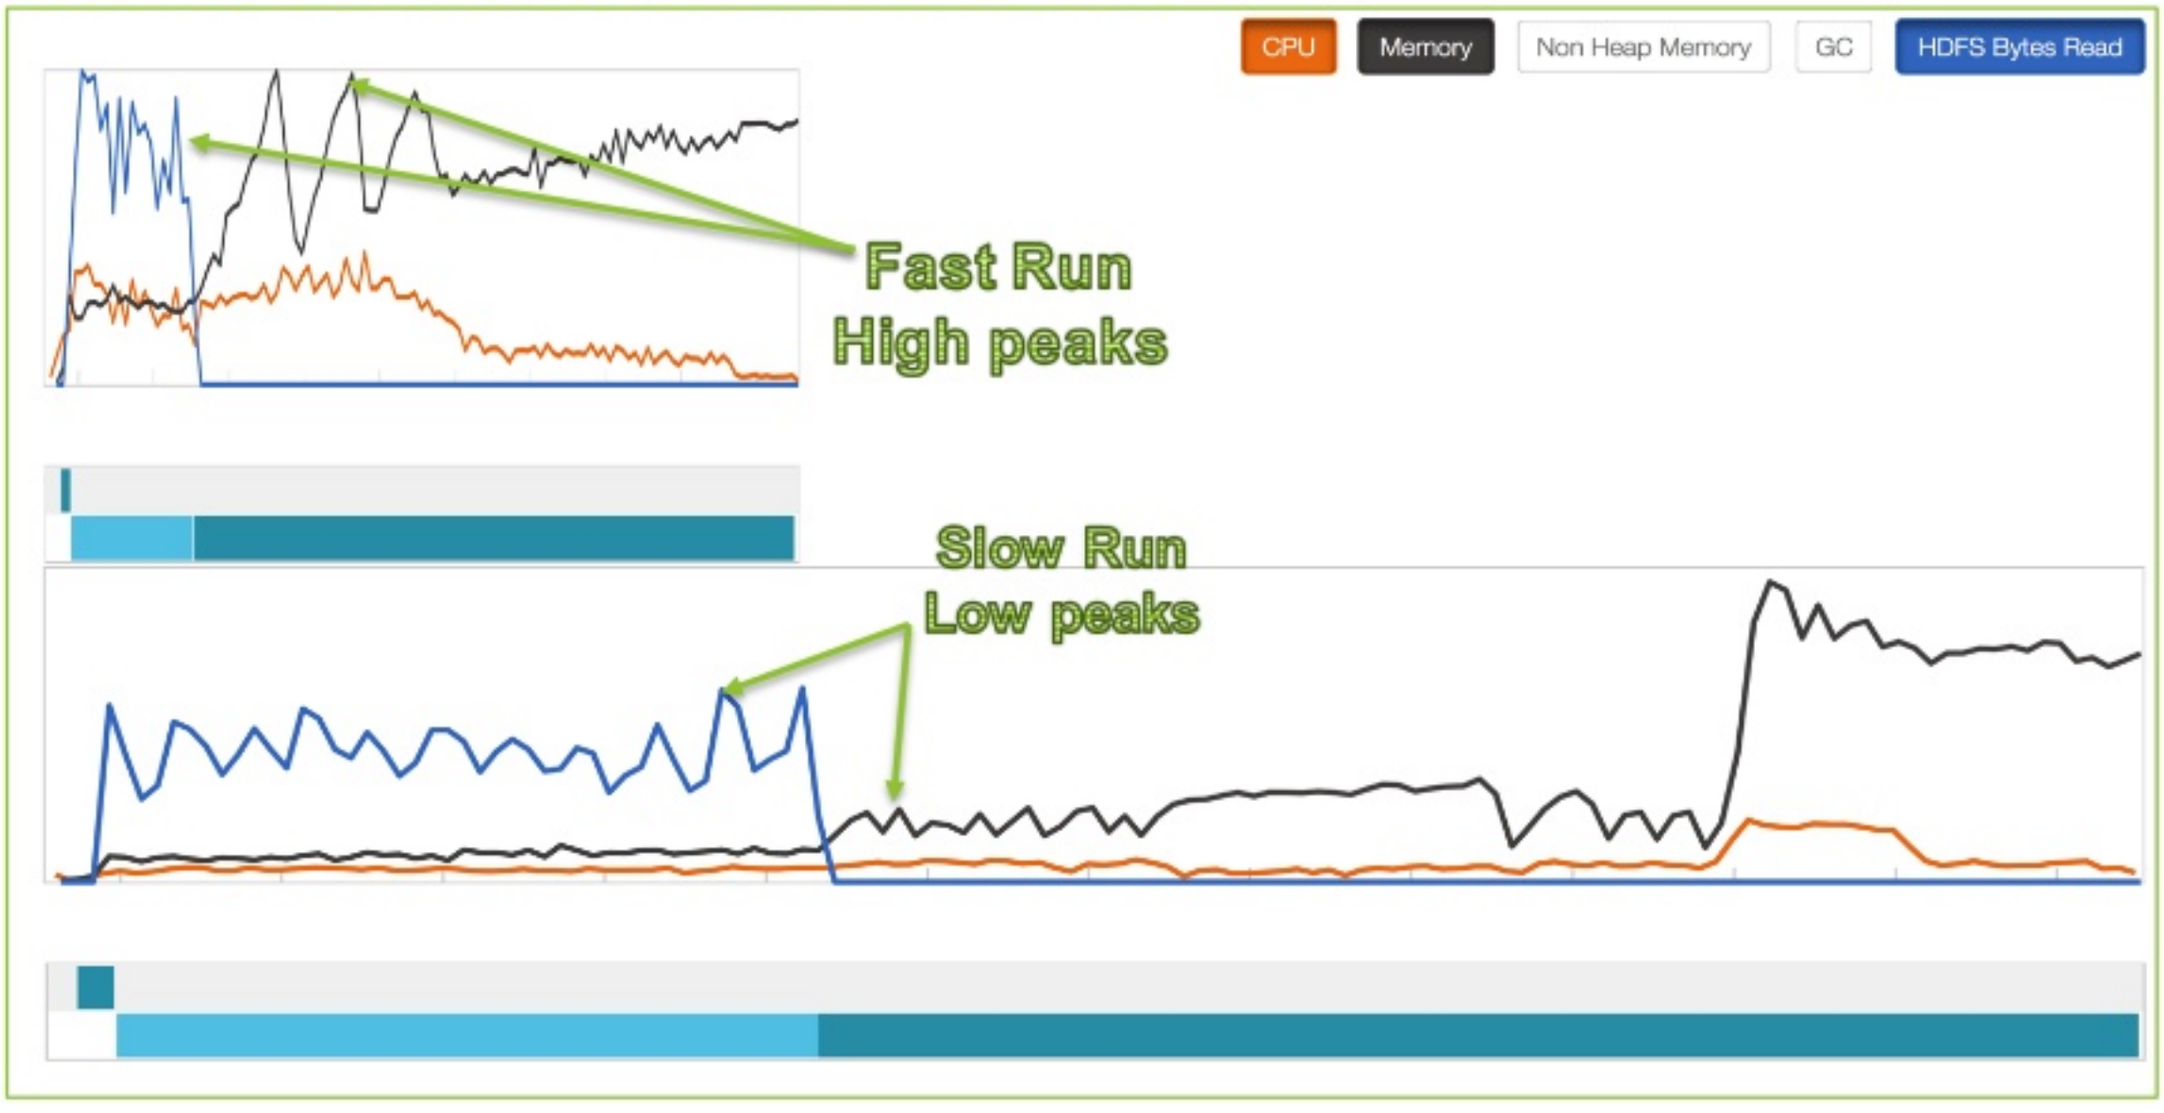

View cluster weather conditions for slower run of App

ScalaTeraSort

App Id: 1495423352820_0369

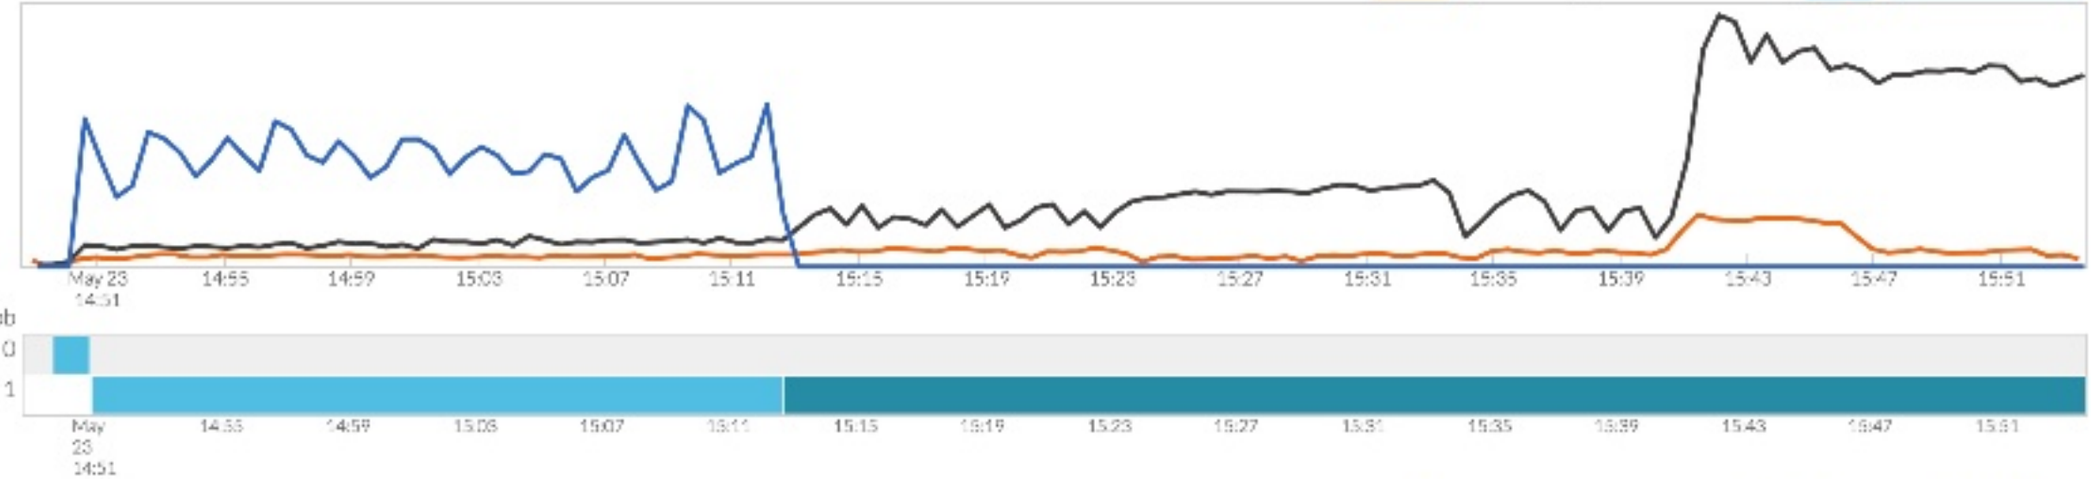

Stage	Job	File Name	Stage Start Time	Stage Duration ((hh:)mm:ss)	Average Executor CPU (%)	Stage Peak Heap Memory(MB)	Stage Peak Non Heap Memory(MB)	Total GC Time (seconds)	Parallel Stages
0	0	BaseRangePartitioner at ScalaTeraSort.scala:56	2017/05/23-14:49	01:08	2.02	1,030	95.1	0	none
1	1	map at ScalaTeraSort.scala:49	2017/05/23-14:50	21:47	3.97	1,479	97	0	none
2	1	saveAsNewAPIHadoopFile at ScalaTeraSort.scala:60	2017/05/23-15:12	41:02	6.74	10,098	411	5.49	none

Cluster weather reveals reason for CPU constraints on slower app

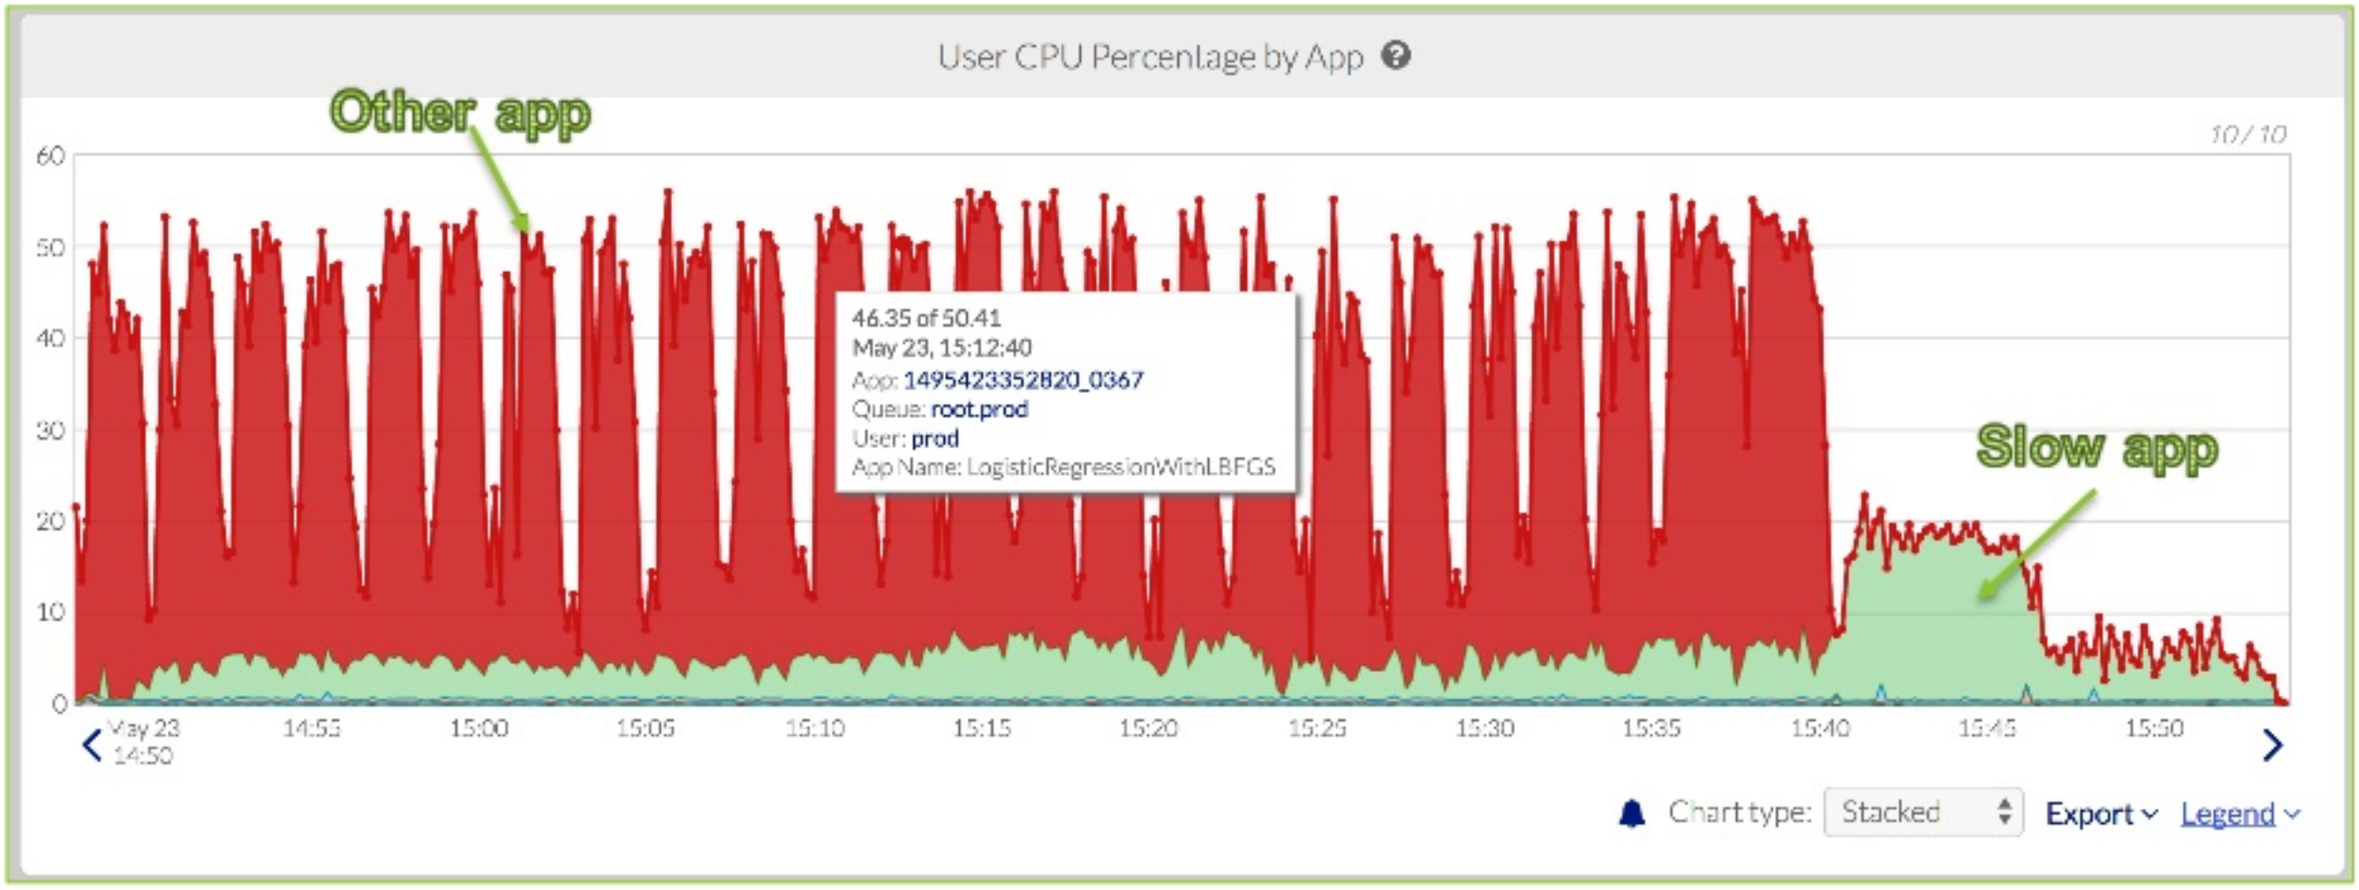

Cluster weather reveals reason for memory constraints on slower app

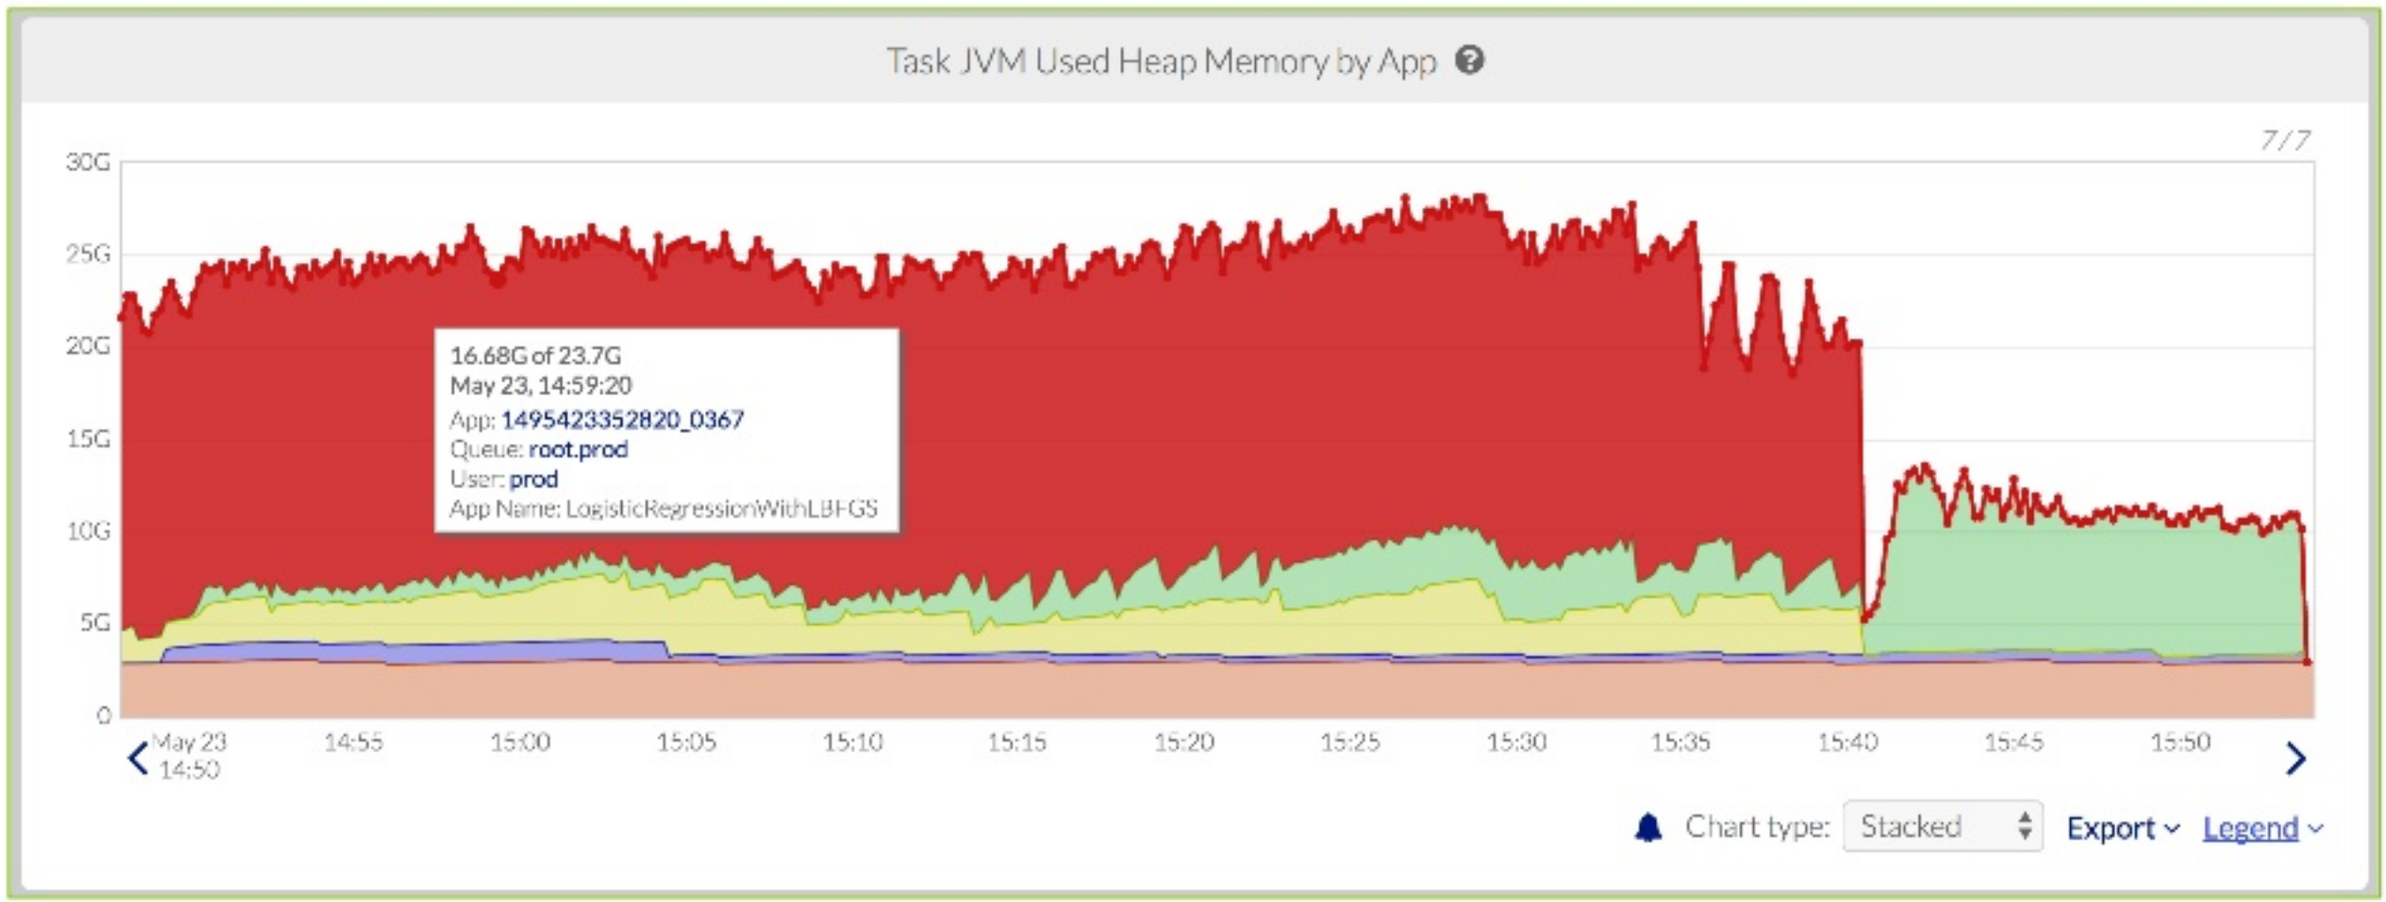

SPARK SUMMIT 2017

Cluster weather reveals reason for HDFS constraints on slower app

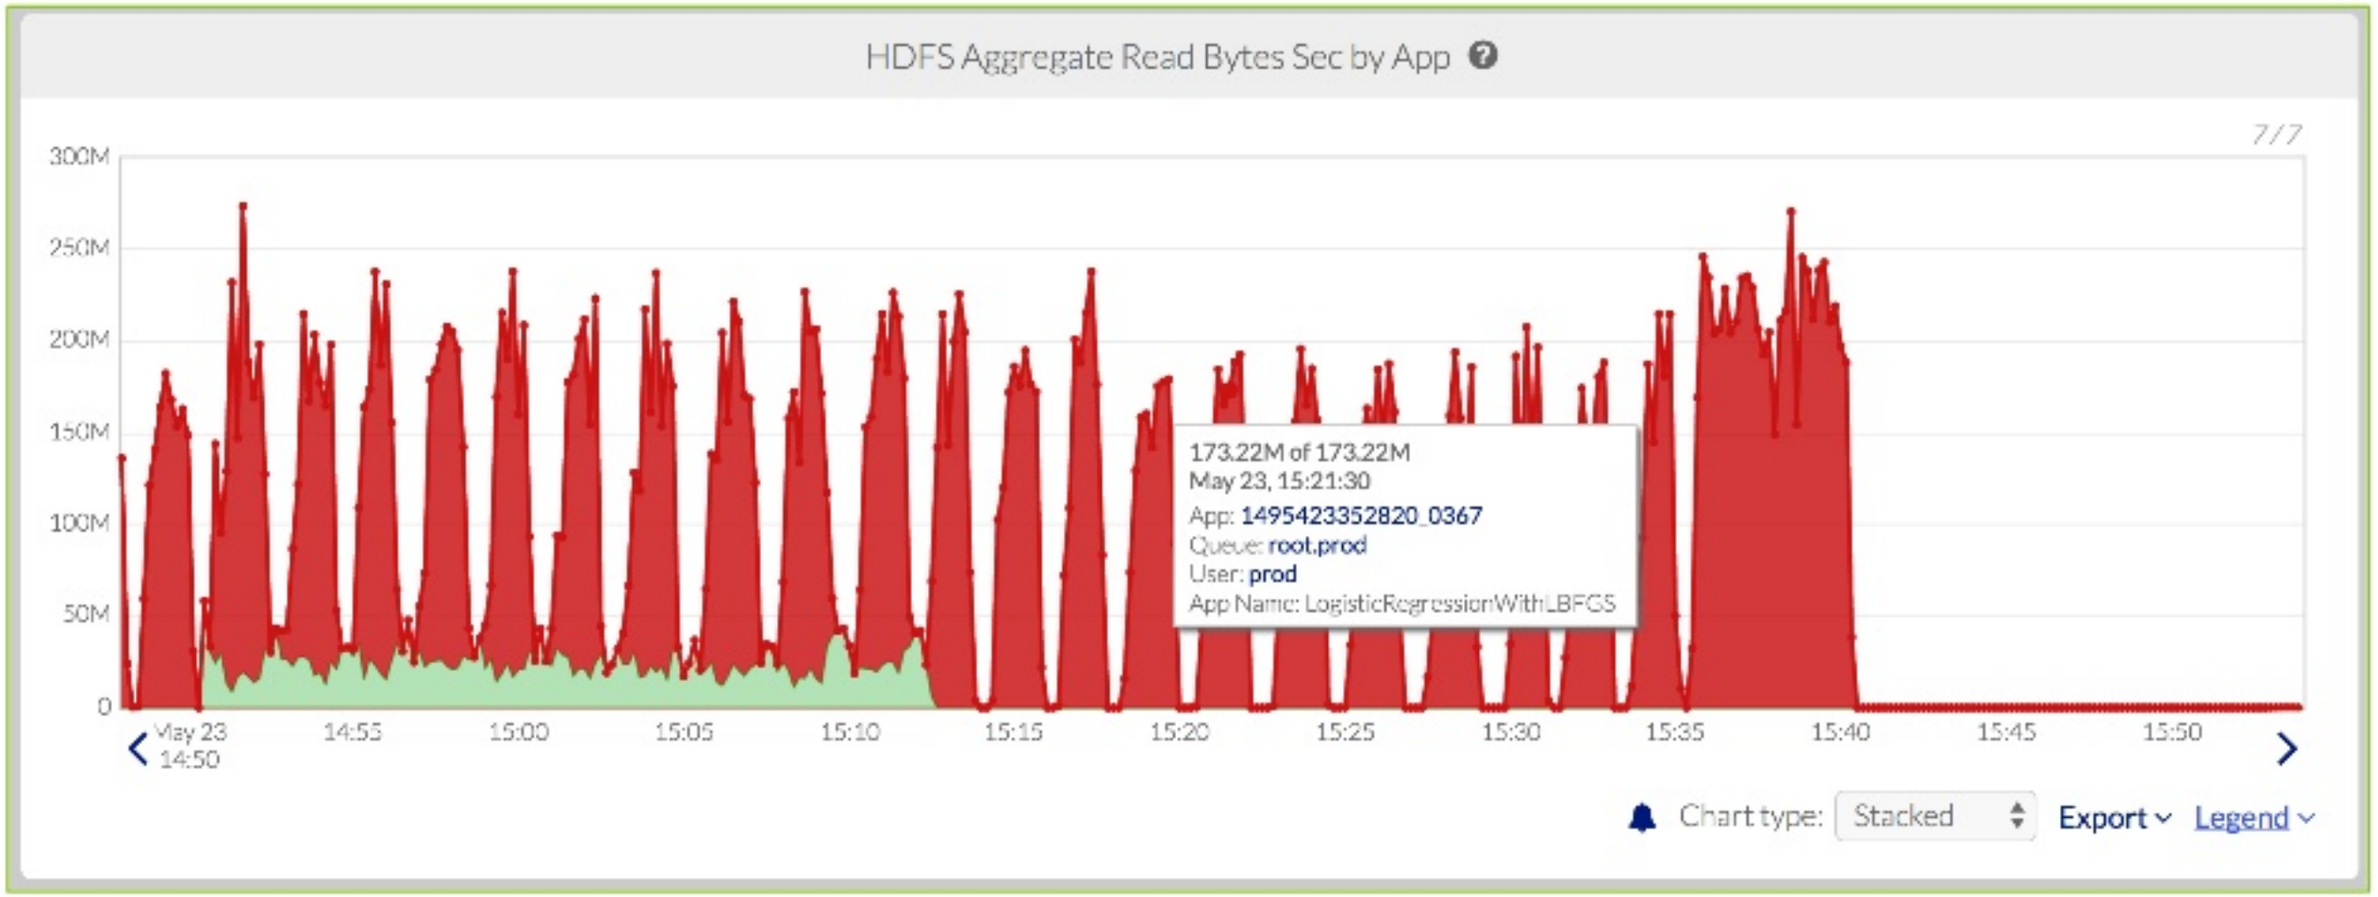

SPARK SUMMIT 2017

To recap

- Execution plan integrated with time series shines a spot light on problems
- Stage and code section integrated with resources consumed enables focus on most impactful areas for optimization
- Knowing cluster weather can prevent time wasted debugging non-existent performance issues

Code Analyzer for Apache Spark

- Free during Early Access starting June 5th, 2017

- Early Access is for development teams

- To learn more visit pepperdata booth #101

- My contact vinod@pepperdata.com

pepperdata.com/products/code-analyzer

SPARK SUMMIT 2017

SPARK SUMMIT 2017

Thank You.

www.pepperdata.com/products/code-analyzer/