



Apache Spark™ MLlib 2.x: How to Productionize your Machine Learning Models

Richard Garris (Principal Solutions Architect)



VISION

Empower anyone to innovate faster with big data.

PRODUCT

A fully-managed data processing platform for the enterprise, powered by Apache Spark.

WHO WE ARE

Founded by the creators of Apache Spark.
Contributes **75%** of the open source code,
10x more than any other company.

YOUR DATA



databricks[™]

VIRTUAL ANALYTICS PLATFORM

CLUSTER TUNING &
MANAGEMENT

INTERACTIVE
WORKSPACE

OPTIMIZED DATA
ACCESS



PRODUCTION
PIPELINE
AUTOMATION

DATABRICKS ENTERPRISE SECURITY

YOUR TEAMS



Data Science



Data Engineering



BI Analysts

Many others...

About Me



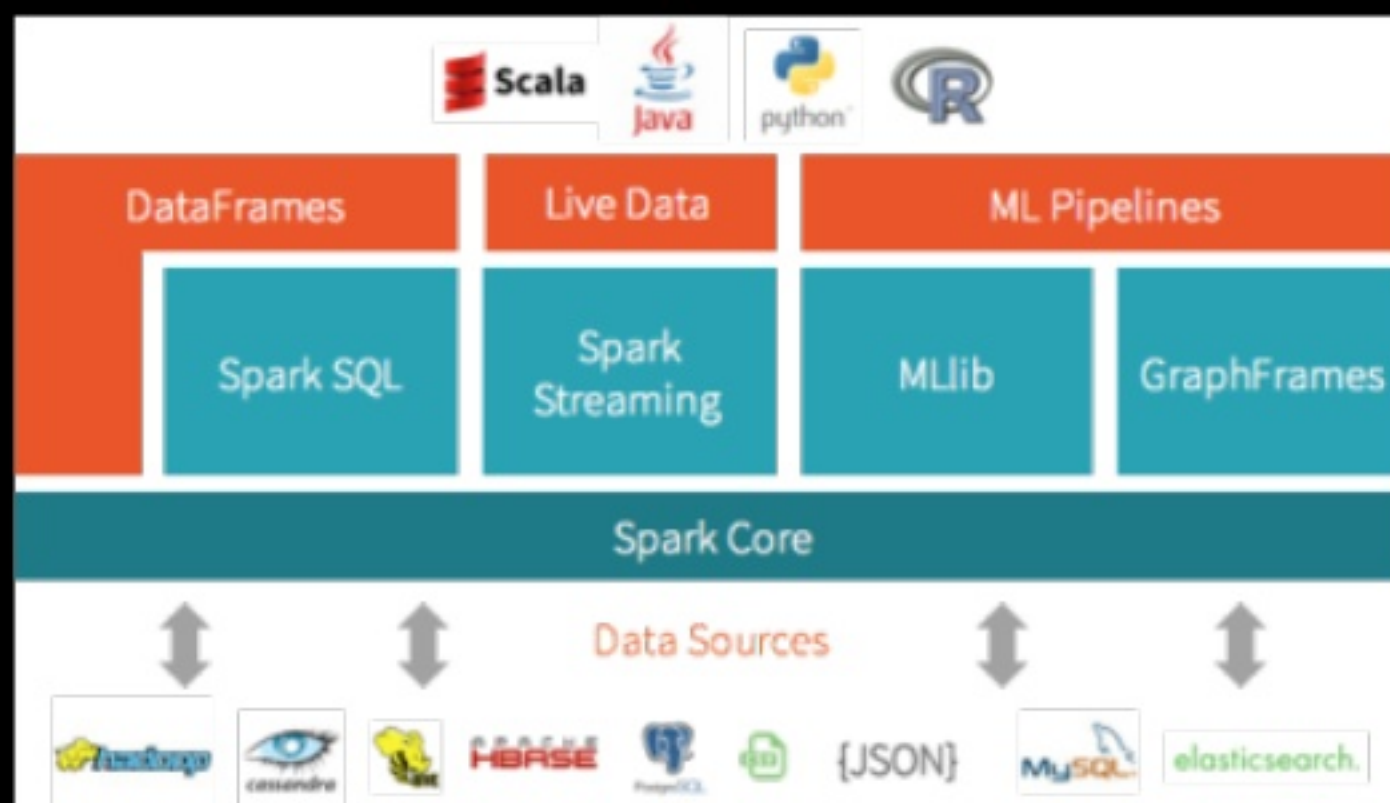
- Richard L Garris
 - rlgarris@databricks.com
 - Twitter @rlgarris
- Principal Data Solutions Architect @ Databricks
- 12+ years designing Enterprise Data Solutions for everyone from startups to Global 2000
- Prior Work Experience PwC, Google and Skytree – the Machine Learning Company
- Ohio State Buckeye and Masters from CMU

Outline

- Spark Mllib 2.X
- Model Serialization
- Model Scoring System Requirements
- Model Scoring Architectures
- Databricks Model Scoring

About Apache Spark™ MLlib

- Started with Spark 0.8 in the AMPLab in 2014
- Migration to Spark DataFrames started with Spark 1.3 with feature parity within 2.X
- Contributions by 75+ orgs, ~250 individuals
- Distributed algorithms that scale linearly with the data

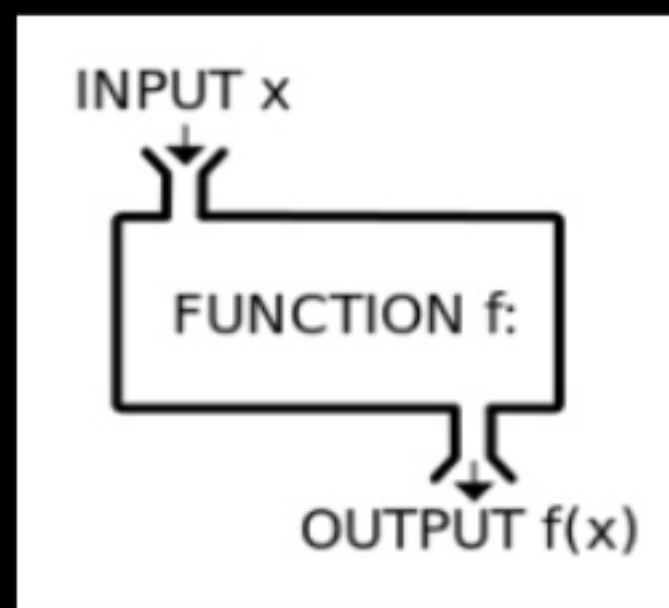
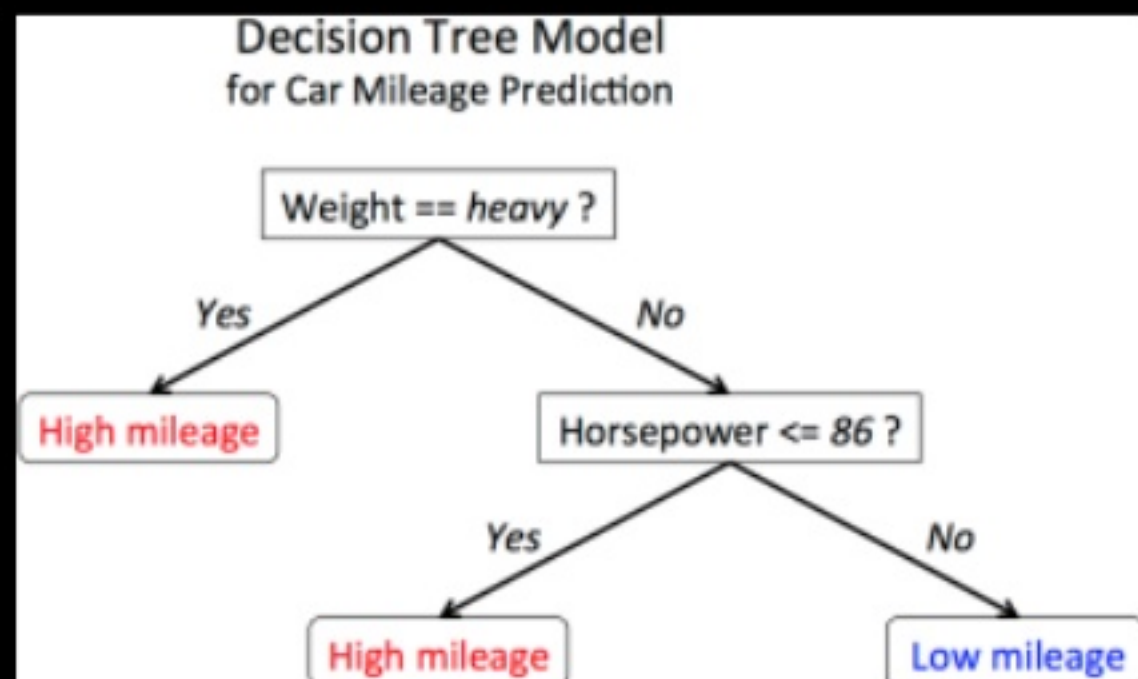


MLlib's Goals

- General purpose machine learning library optimized for big data
 - Linearly scalable = 2x more machines , runtime theoretically cut in half
 - Fault tolerant = resilient to the failure of nodes
 - Covers the most common algorithms with distributed implementations
- Built around the concept of a Data Science Pipeline (scikit-learn)
- Written entirely using Apache Spark™
- Integrates well with the Agile Modeling Process

A Model is a Mathematical Function

- A model is a function: $f(x)$
- Linear regression $y = b_0 + b_1x_1 + b_2x_2$



ML Pipelines



A very simple pipeline

ML Pipelines

A real pipeline!



Productionizing Models Today

Data Science

Develop Prototype
Model using Python/R

Data Engineering

Re-implement model for
production (Java)

Problems with Productionizing Models

Data Science

Develop Prototype
Model using Python/R

Data Engineering

Re-implement model for
production (Java)

- Extra work
- Different code paths
- Data science does not translate to production
- Slow to update models

MLlib 2.X Model Serialization

Data Science

Develop Prototype
Model using Python/R



Persist model or Pipeline:

```
model.save("s3n://...")
```

Data Engineering

Load Pipeline (Scala/Java)

```
Model.load("s3n://...")
```

Deploy in production



MLlib 2.X Model Serialization Snippet

Scala

```
val lrModel = lrPipeline.fit(dataset)
```

```
// Save the Model
```

```
lrModel.write.save("/models/lr")
```

Python

```
lrModel = lrPipeline.fit(dataset)
```

```
# Save the Model
```

```
lrModel.write.save("/models/lr")
```


Model Serialization Output

Code

```
// List Contents of the Model Dir  
dbutils.fs.ls("/models/lr")
```

Output

path	name
dbfs:/models/lr/metadata/	metadata/
dbfs:/models/lr/stages/	stages/

Remember this is a pipeline model and these are the stages!

Transformer Stage (StringIndexer)

Code

```
// Cat the contents of the Metadata dir
dbutils.fs.head("/models/lr/stages/00_strIdx_bb9728f85745/metadata/part-00000")

// Display the Parquet File in the Data dir
display(spark.read.parquet("/models/lr/stages/00_strIdx_bb9728f85745/data/"))
```

Output

```
{
  "class": "org.apache.spark.ml.feature.StringIndexerModel",
  "timestamp": 1488120411719,
  "sparkVersion": "2.1.0",
  "uid": "strIdx_bb9728f85745",
  "paramMap": {
    "outputCol": "workclassIdx",
    "inputCol": "workclass",
    "handleInvalid": "error"
  }
}
```

Metadata and params

Data (Hashmap)

labels	
▼ array	
0:	Private
1:	Self-emp-not-inc
2:	Local-gov
3:	?
4:	State-gov
5:	Self-emp-inc
6:	Federal-gov
7:	Without-pay
8:	Never-worked

Estimator Stage (LogisticRegression)

Code

```
// Cat the contents of the Metadata dir
dbutils.fs.head("/models/lr/stages/18_logreg_325fa760f925/metadata/part-00000")
```

```
// Display the Parquet File in the Data dir
display(spark.read.parquet("/models/lr/stages/18_logreg_325fa760f925/data/"))
```

Output

Model params

```
{"class": "org.apache.spark.ml.classification.LogisticRegressionModel",
 "timestamp": 1488120446324,
 "sparkVersion": "2.1.0",
 "uid": "logreg_325fa760f925",
 "paramMap": {
   "predictionCol": "prediction",
   "standardization": true,
   "probabilityCol": "probability",
   "maxIter": 100,
   "elasticNetParam": 0.0,
   "family": "auto",
   "regParam": 0.0,
   "threshold": 0.5,
   "fitIntercept": true,
   "labelCol": "label"
 }}
```

numClasses	numFeatures	InterceptVector	coefficientMatrix
2	100	<div>▼ array</div> <div>0: 1</div> <div>1: 1</div> <div>2: []</div> <div>▼ 3:</div> <div>0:</div> <div>-2.0412260644120477</div>	<div>▼ object</div> <div>numRows: 1</div> <div>numCols: 100</div> <div>▼ values:</div> <div>0: -1.6536519269918135</div> <div>1: -2.1437754918551044</div> <div>2: -1.8372425650595294</div>

Intercept + Coefficients

Estimator Stage (DecisionTree)

Code

```
// Display the Parquet File in the Data dir
display(spark.read.parquet("/models/dt/stages/18_dtc_3d614bcb3ff825/data/"))

// Re-save as JSON
spark.read.parquet("/models/dt/stages/18_dtc_3d614bcb3ff825/data/").json("/models/json/dt").
```

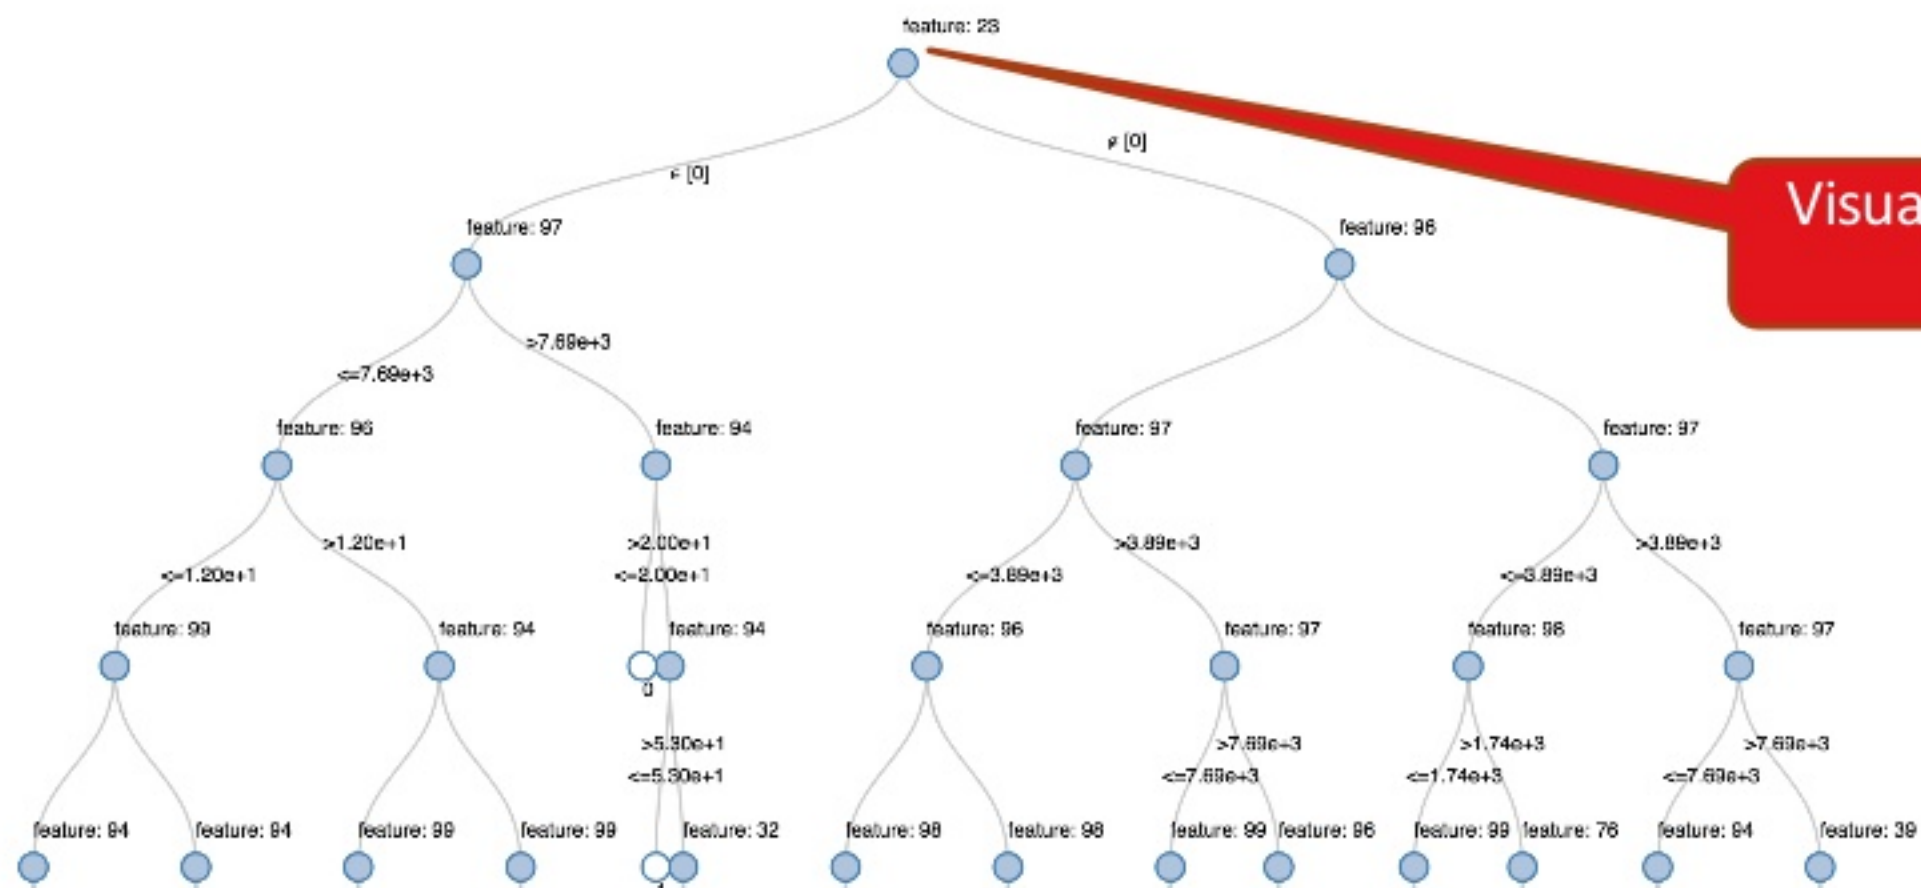
Output

id	prediction	impurity	impurityStats	gain	leftChild	rightChild	split
99	1	0	▶ [0,2]	-1	-1	-1	▶ {"featureIndex":-1,"leftCategoriesOrThreshold":[],"numCategories":-1}
100	1	0.005154604757994008	▶ [1,386]	0.00013864488567297793	101	102	▶ {"featureIndex":94,"leftCategoriesOrThreshold": [66],"numCategories":-1}
101	1	0	▶ [0,353]	-1	-1	-1	▶ {"featureIndex":-1,"leftCategoriesOrThreshold":[],"numCategories":-1}
102	1	0.007000105000550010	▶ [1,386]	0.00013864488567297793	101	102	▶ {"featureIndex":94,"leftCategoriesOrThreshold": [66],"numCategories":-1}

Decision Tree Splits

Visualize Stage (DecisionTree)

```
display(dtModel.stages.last.asInstanceOf[DecisionTreeClassificationModel])
```



Visualization of the Tree In Databricks

What are the Requirements for a Robust Model Deployment System?

Model Scoring Environment Examples

- In Web Applications / Ecommerce Portals
- Mainframe / Batch Processing Systems
- Real-Time Processing Systems / Middleware
- Via API / Microservice
- Embedded in Devices (Mobile Phones, Medical Devices, Autos)



Hidden Technical Debt in ML Systems

“Hidden Technical Debt in Machine Learning Systems”, Google NIPS 2015

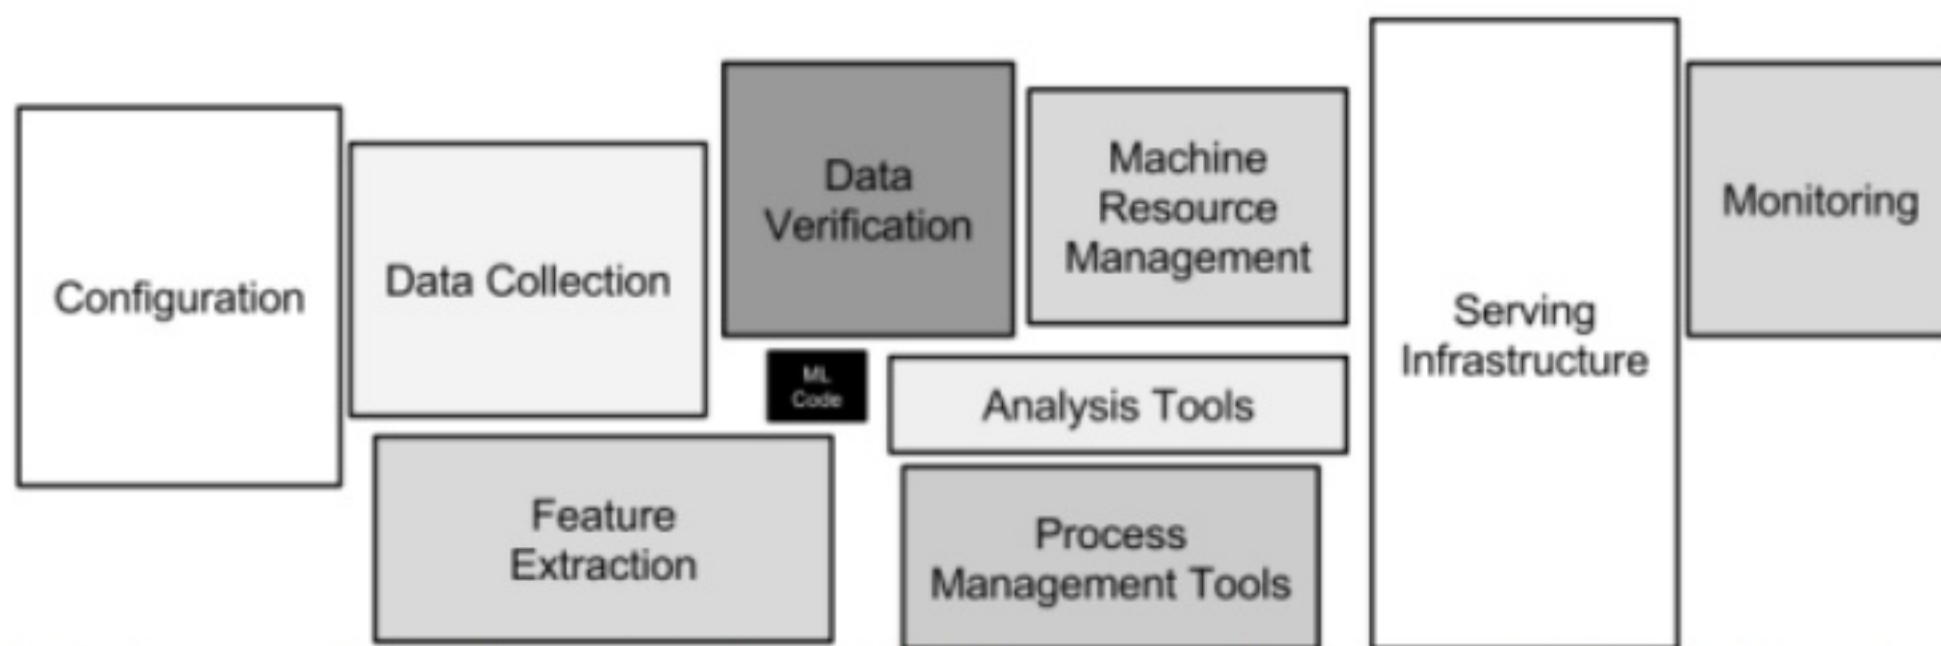
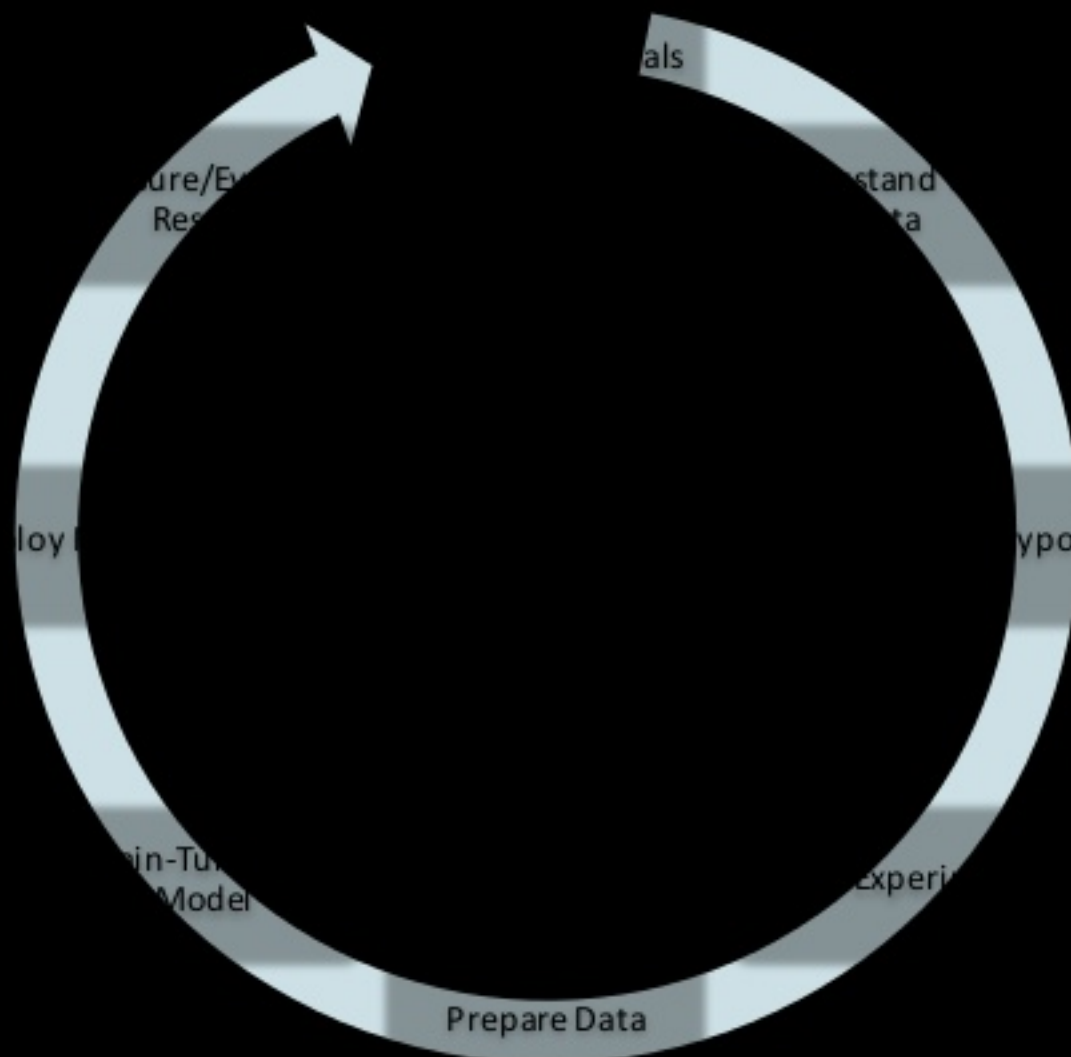
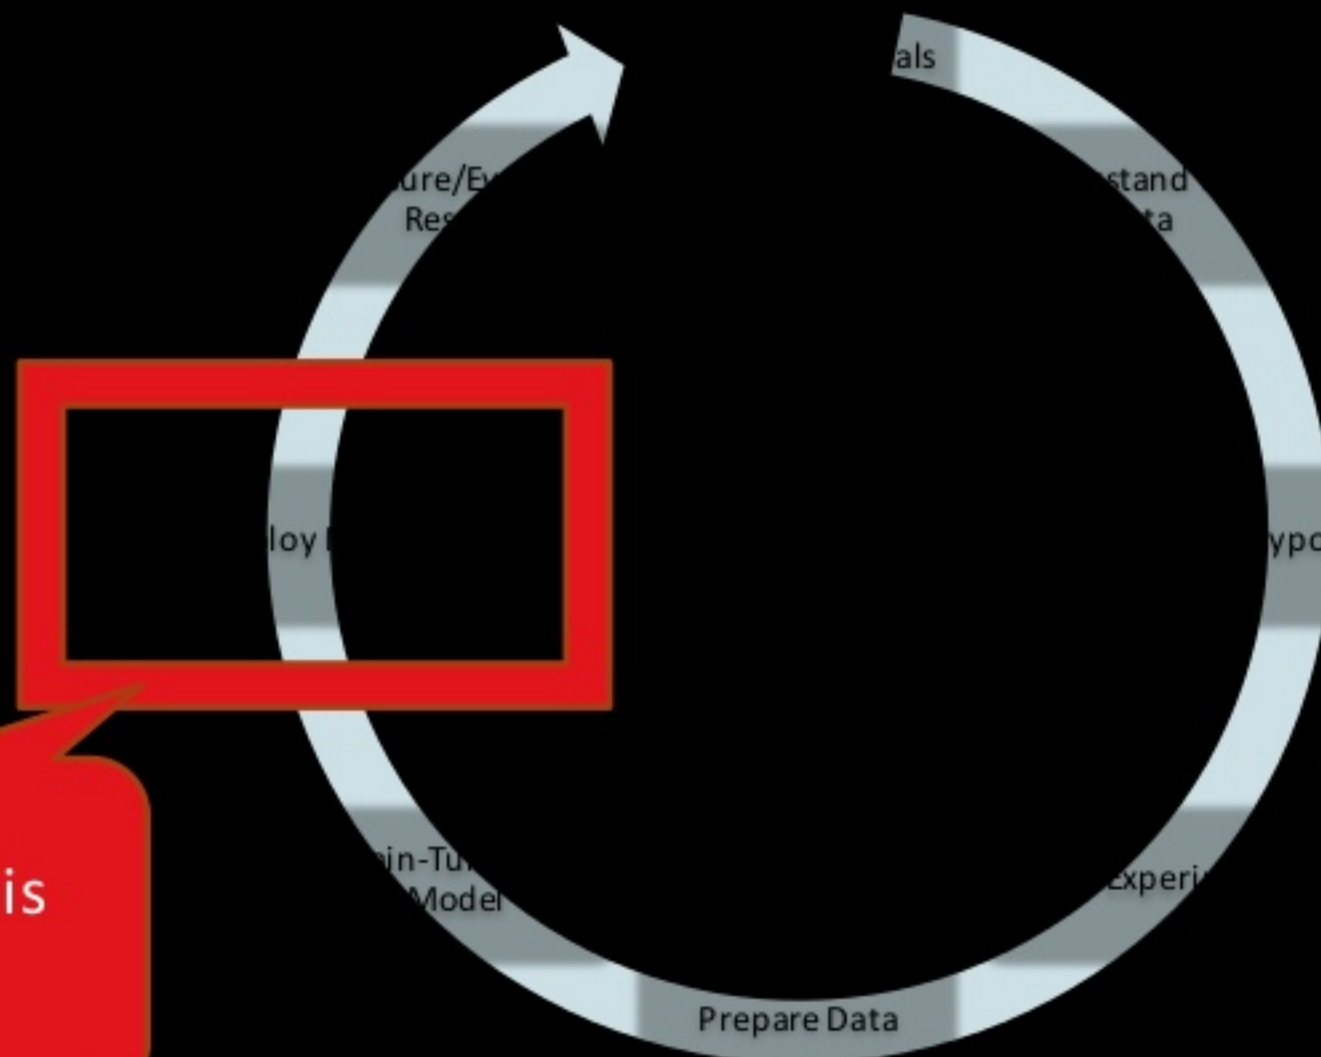


Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

Agile Modeling Process

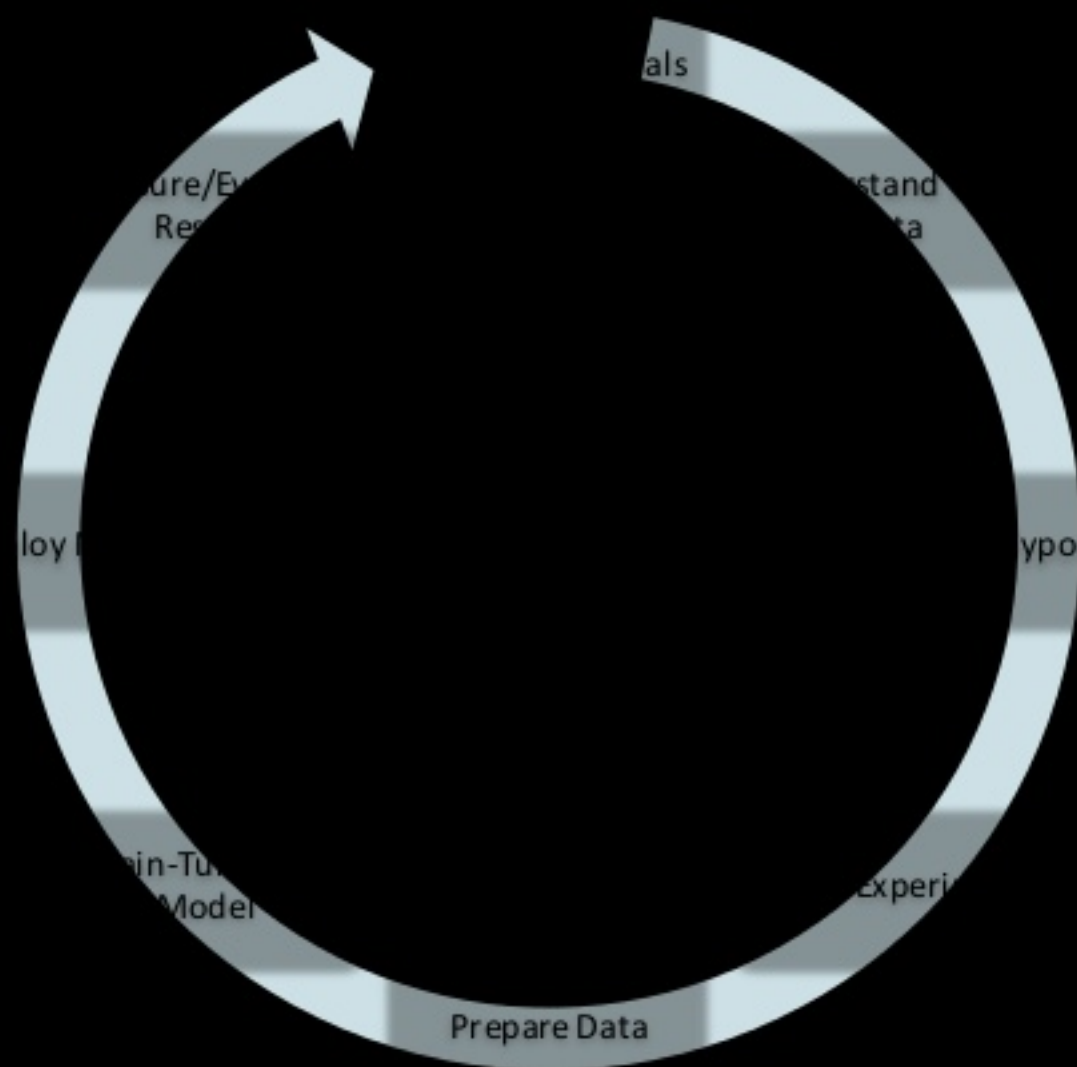


Agile Modeling Process



Focus of this
talk

Deployment Should be Agile



- Deployment needs to support A/B testing and experiments
- Deployment should support measuring and evaluating model performance
- Deployment should be fast and adaptive to business needs

Model A/B Testing, Monitoring, Updates

- A/B testing – comparing two versions to see what performs better
- Monitoring is the process of observing the model's performance, logging it's behavior and alerting when the model degrades
- Logging should log exactly the data feed into the model at the time of scoring
- Model update process
 - Benchmark (or Shadow Models)
 - Phase-In (20% traffic)
 - Avoid Big Bang



Consider the Scoring Environment

Customer SLAs

- Response time
- Throughput
(predictions per second)
- Uptime / Reliability

Tech Stack

- C / C++
- Legacy (mainframe)
- Java

Scoring in Batch vs Real-Time

Batch

- Asynchronous
- Internal Use
- Triggers can be event based on time based
- Used for Email Campaigns, Notifications

Real-Time

- Synchronous
- Could be Seconds:
 - Customer is waiting (human real-time)
- Subsecond:
 - High Frequency Trading
 - Fraud Detection on the Swipe

Online Learning and Open / Closed Loop

Open / Closed Loop

Open Loop – human being involved

Closed Loop – no human involved

- Model Scoring – almost always closed loop, some open loop e.g. alert agents or customer service
- Model Training – usually open loop with a data scientist in the loop to update the model

Online Learning

- Online is closed loop, entirely machine driven but modeling is risky
- need to have proper model monitoring and safeguards to prevent abuse / sensitivity to noise
- MLlib supports online through streaming models (k-means, logistic regression support online)
- Alternative – use a more complex model to better fit new data rather than using online learning

Model Scoring – Bot Detection

Not All Models Return Boolean – e.g. a Yes / No

Example: Login Bot Detector

Different behavior depending on probability that user is a bot

0.0-0.4 ➡ Allow login

0.4-0.6 ➡ Send Challenge Question

0.6 to 0.75 ➡ Send SMS Code

0.75 to 0.9 ➡ Refer to Agent

0.9 - 1.0 ➡ Block



Model Scoring – Recommendations

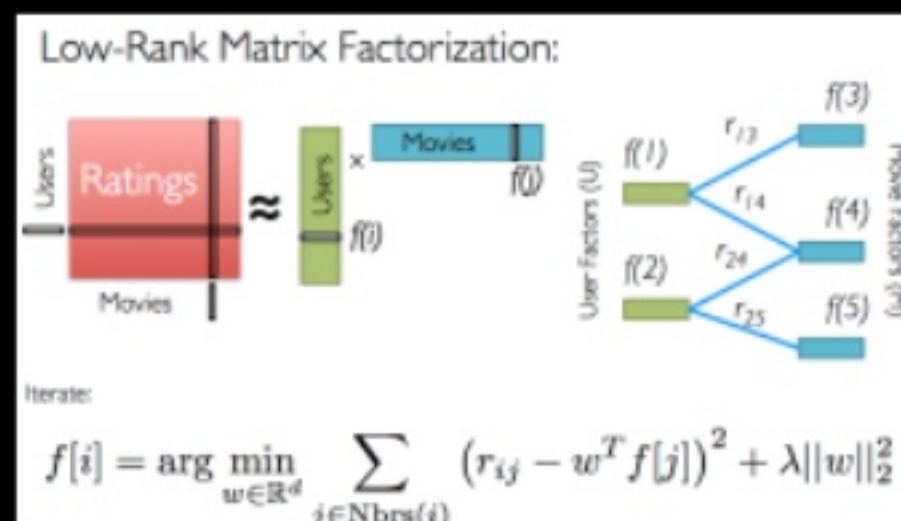
Output is a ranking of the top n items

API – send user ID + number of items

Return sorted set of items to recommend

Optional –

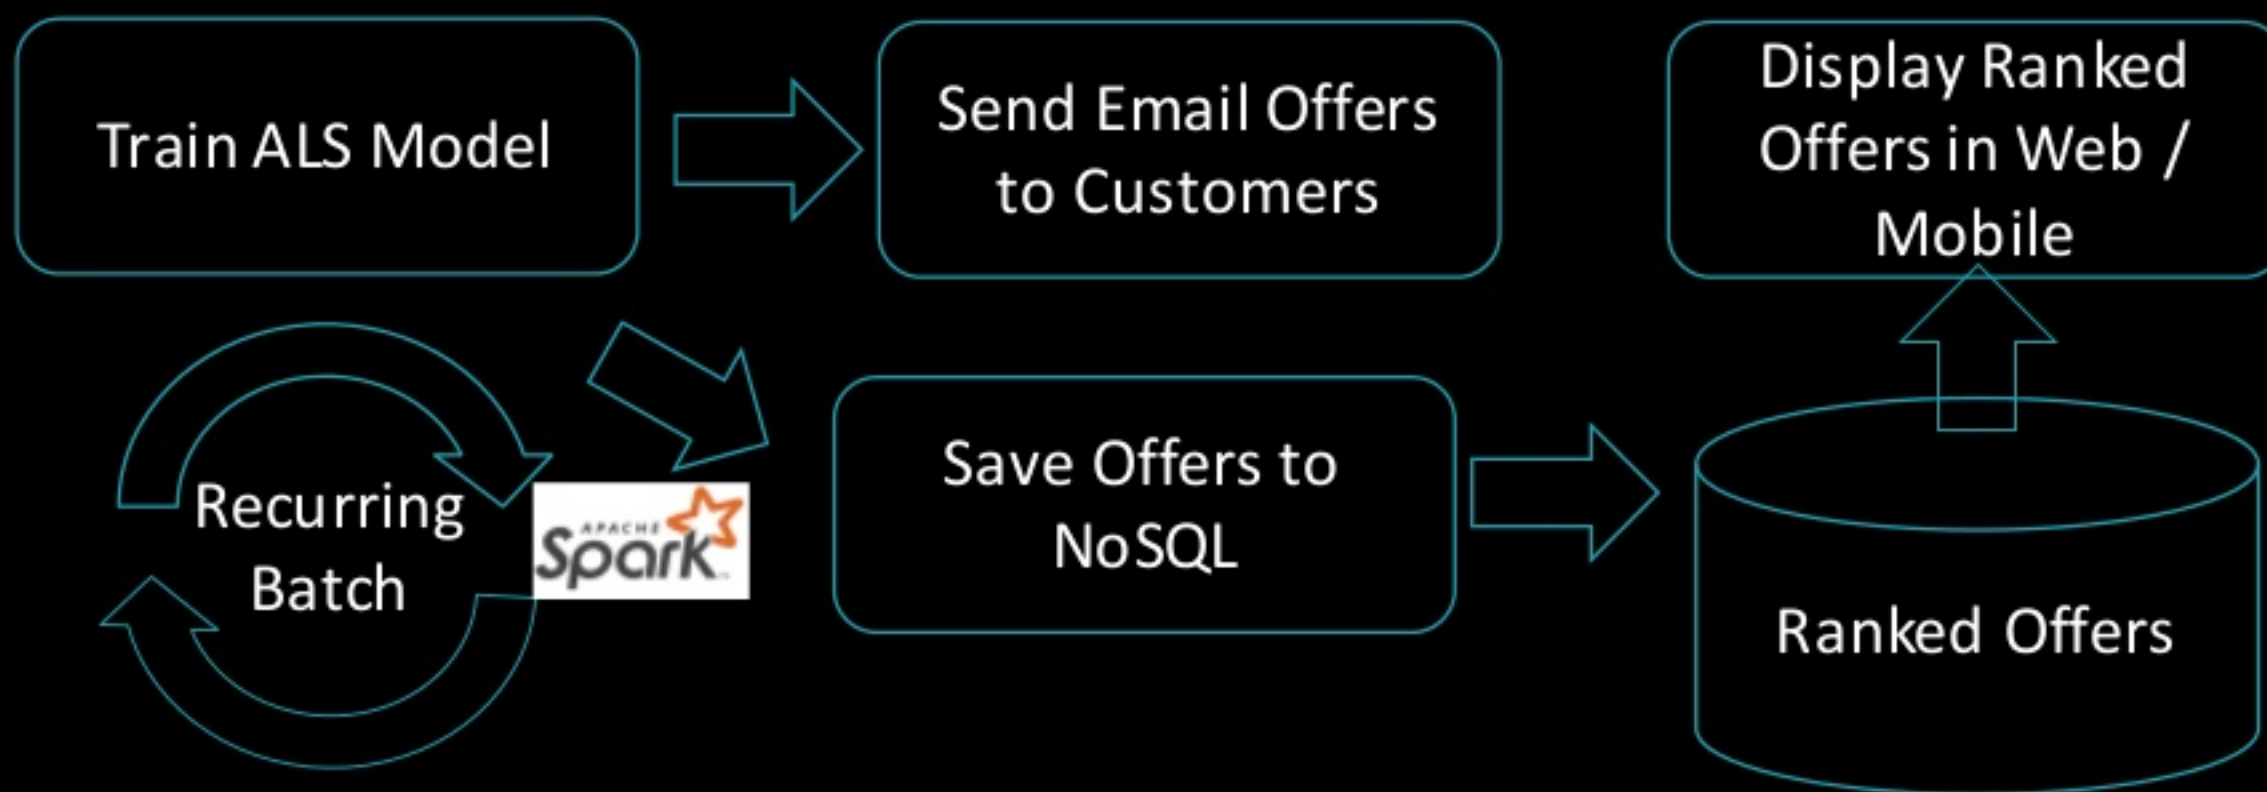
pass context sensitive information to tailor results



Model Scoring Architectures

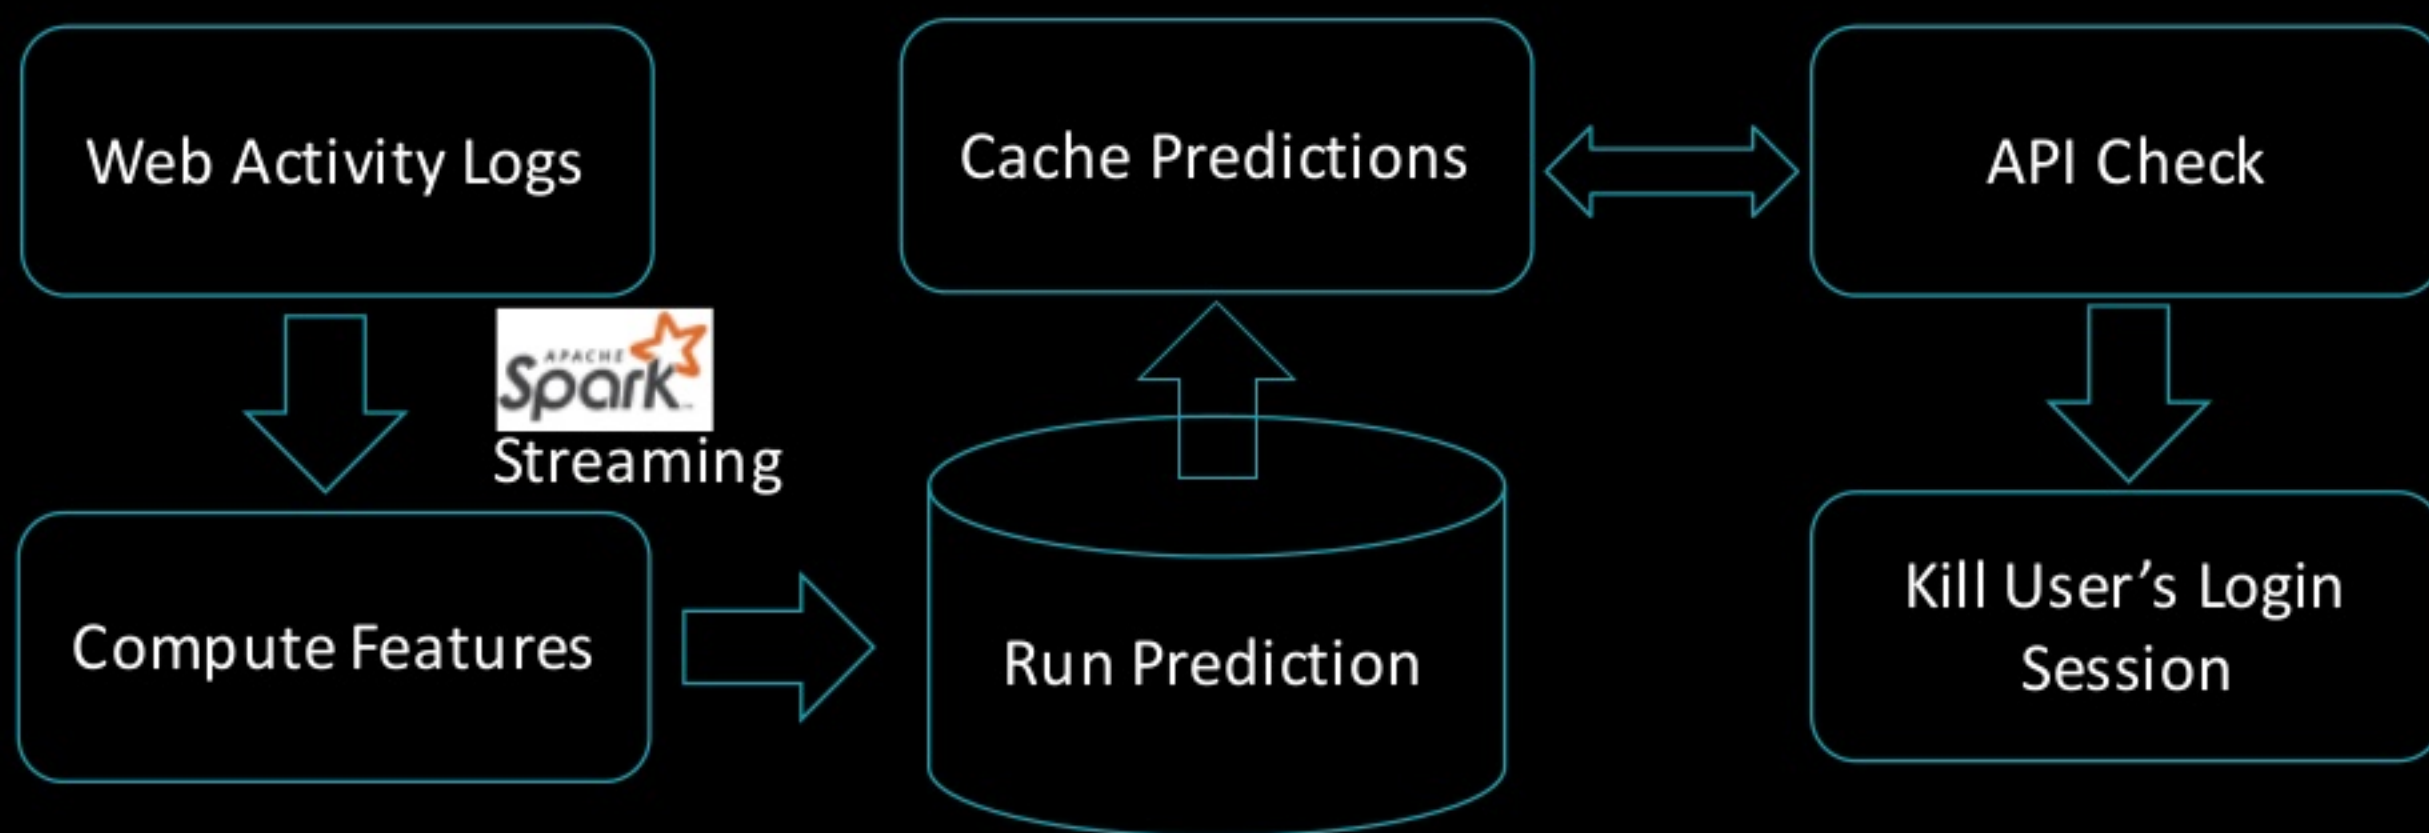
Architecture Option A

Precompute Predictions using Spark and Serve from Database



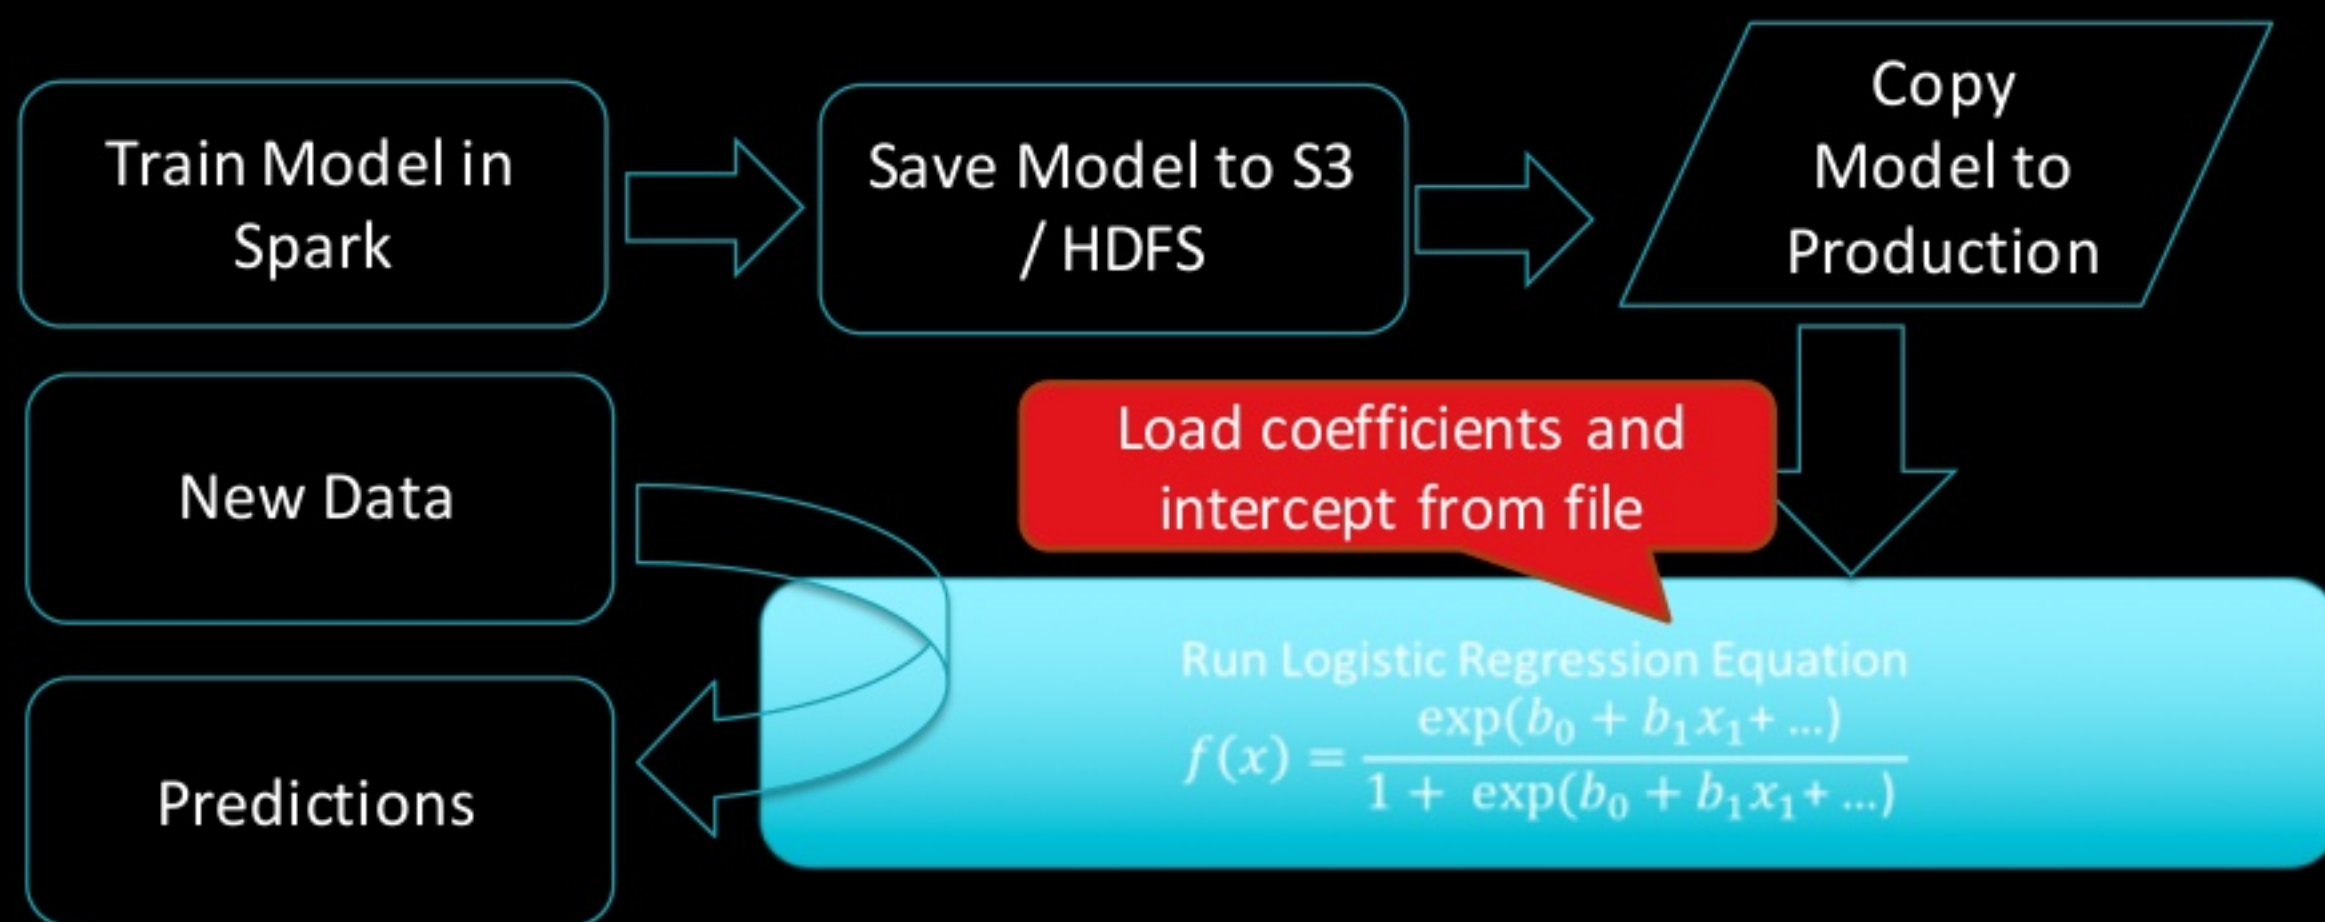
Architecture Option B

Spark Stream and Score using an API with Cached Predictions



Architecture Option C

Train with Spark and Score Outside of Spark



Databricks Model Scoring

Databricks Model Scoring

- Based on Architecture Option C
- Goal: Deploy MLlib model outside of Apache Spark and Databricks.
 - Easy to Embed in Existing Environments
 - Low Latency and Complexity
 - Low Overhead

Databricks Model Scoring

- Train Model in Databricks
 - Call Fit on Pipeline
 - Save Model as JSON
- Deploy model in external system
 - Add dependency on “dbml-local” package (without Spark)
 - Load model from JSON at startup
 - Make predictions in real time

Code

```
// Fit and Export the Model in Databricks
```

```
val lrModel = lrPipeline.fit(dataset)
```

```
ModelExporter.export(lrModel, "/models/db ")
```

```
// In Your Application (Scala)
```

```
import com.databricks.ml.local.ModelImport
```

```
val lrModel = ModelImport.import("s3a:/...")
```

```
val jsonInput = ...
```

```
val jsonOutput = lrModel.transform(jsonInput)
```


Databricks Model Scoring Private Beta

- Private Beta Available for Databricks Customers
- Available on Databricks using Apache Spark 2.1
- Only logistic regression available now
- Additional Estimators and Transformers in Progress

Demo Model Scoring

<https://community.cloud.databricks.com/?o=1526931011080774#notebook/1904316851197504>



Thank You.

Questions?

Happy Sparking
richard@databricks.com