

@superguenter @dcos #smack

POWERING PREDICTIVE MAPPING AT SCALE WITH SPARK, KAFKA & ELASTIC



Spark Summit West

June 07, 2017



Tobi Knaup

Co-Founder & CTO, Mesosphere



@superguenter

HYPERSCALE MEANS VOLUME *AND* VELOCITY

Days

Hours

Minutes

Seconds

Microseconds

Batch

Micro-Batch

Event Processing

Reports what has happened using descriptive analytics

Solves problems using predictive and prescriptive analytics

Billing, Chargeback



Product Recommendations



Real-time Pricing and Routing



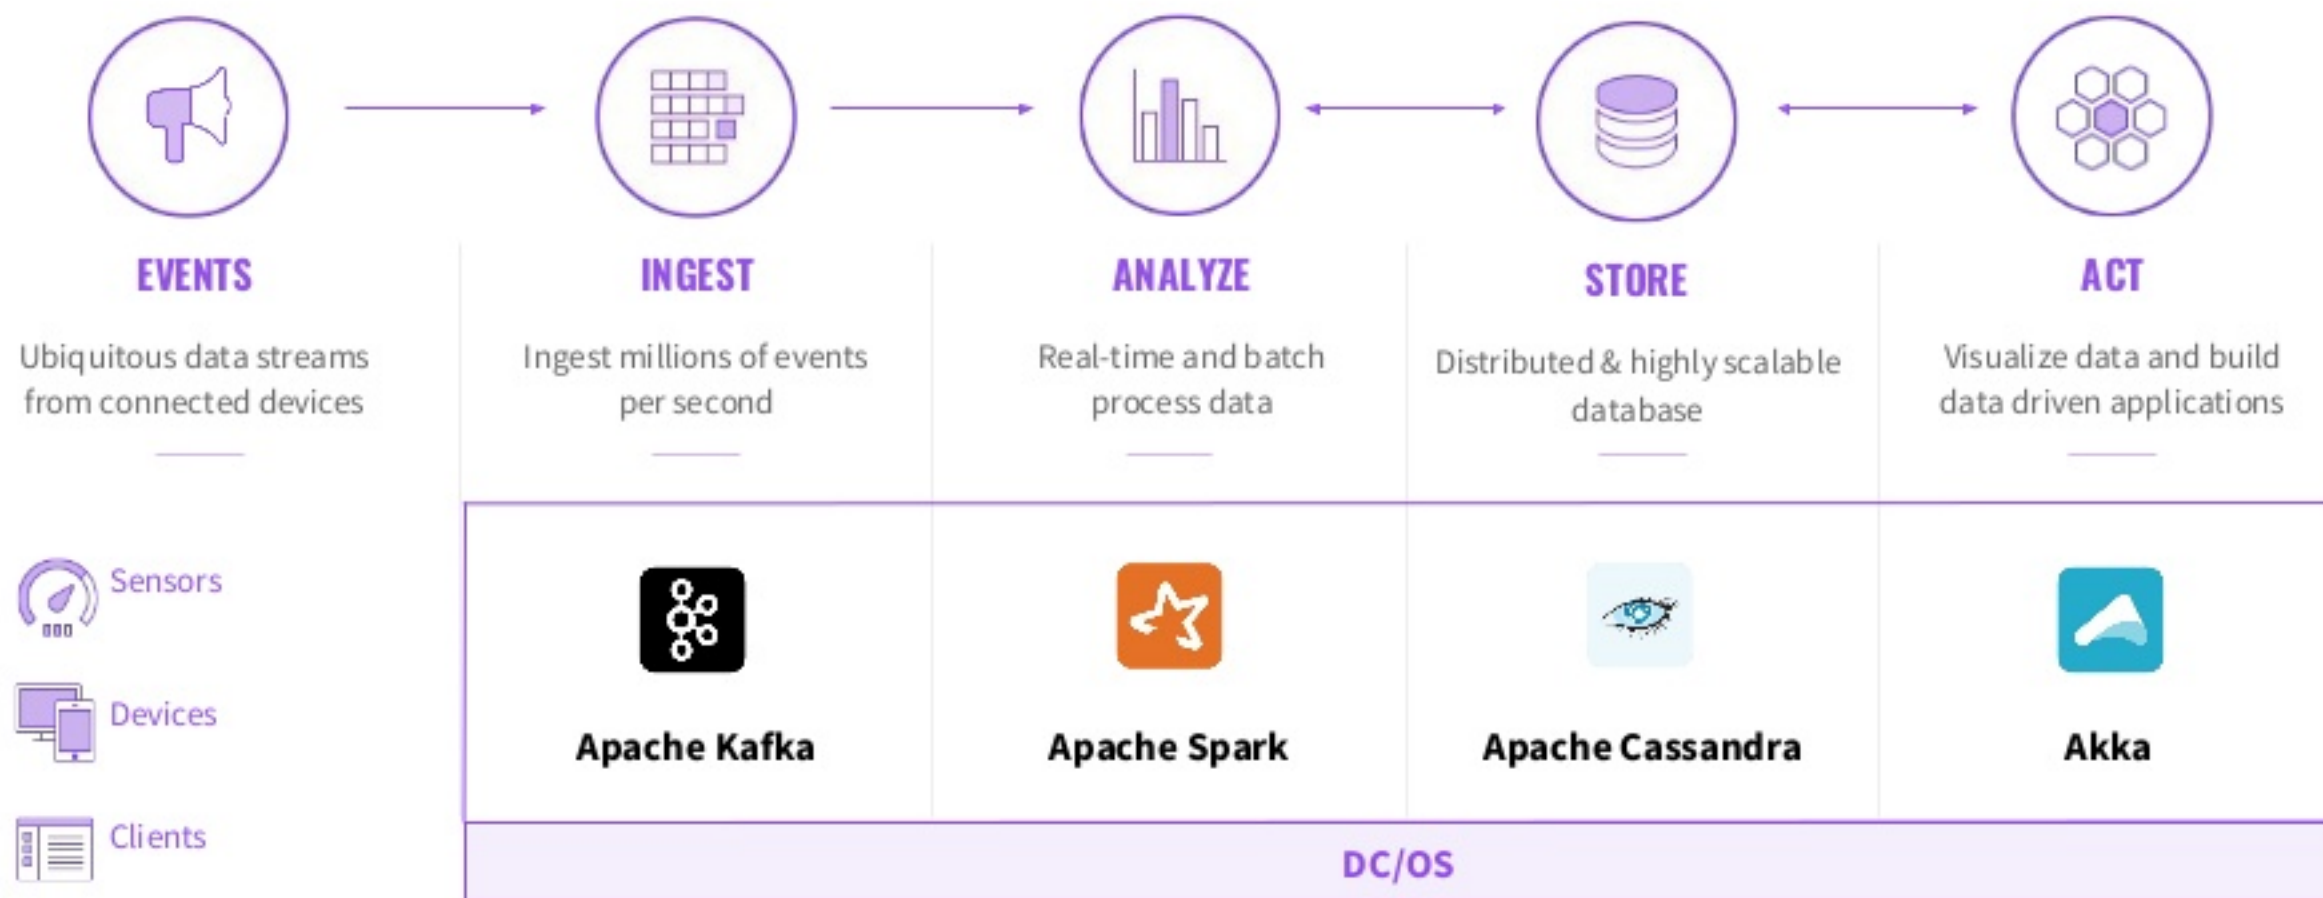
Real-time Advertising



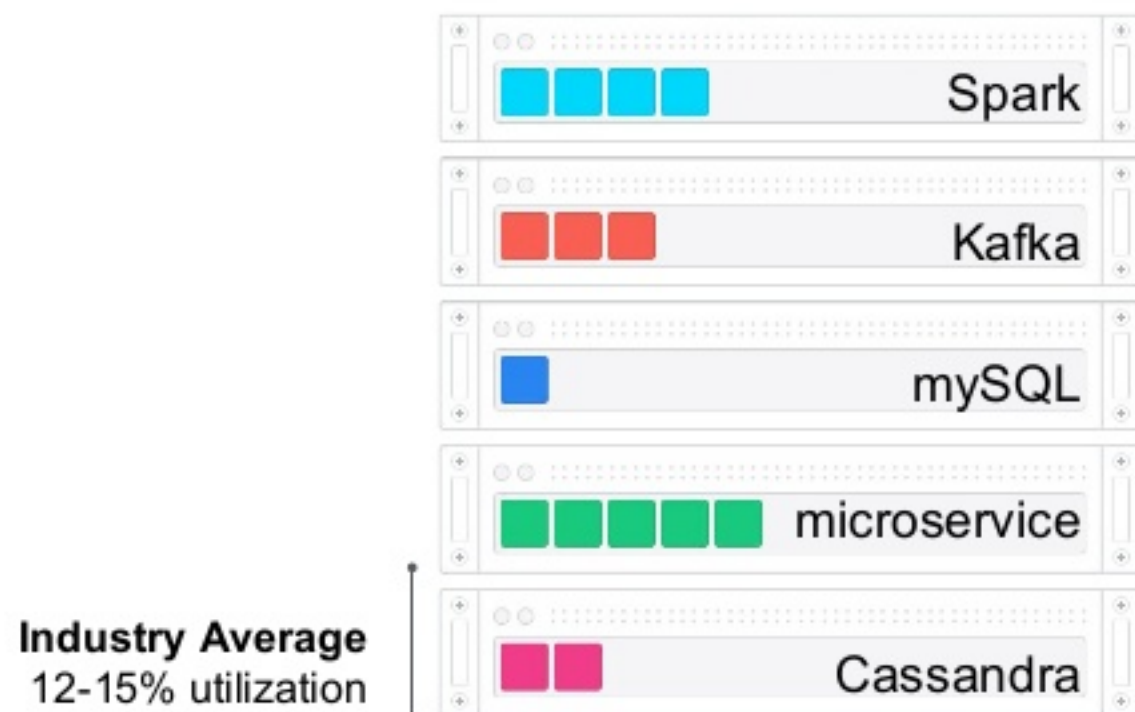
Predictive User Interface



S.M.A.C.K. STACK



NAIVE APPROACH



Typical Datacenter
siloed, over-provisioned servers,
low utilization

DC/OS APPROACH




DC/OS

MULTIPLEXING OF DATA, SERVICES, USERS, ENVIRONMENTS



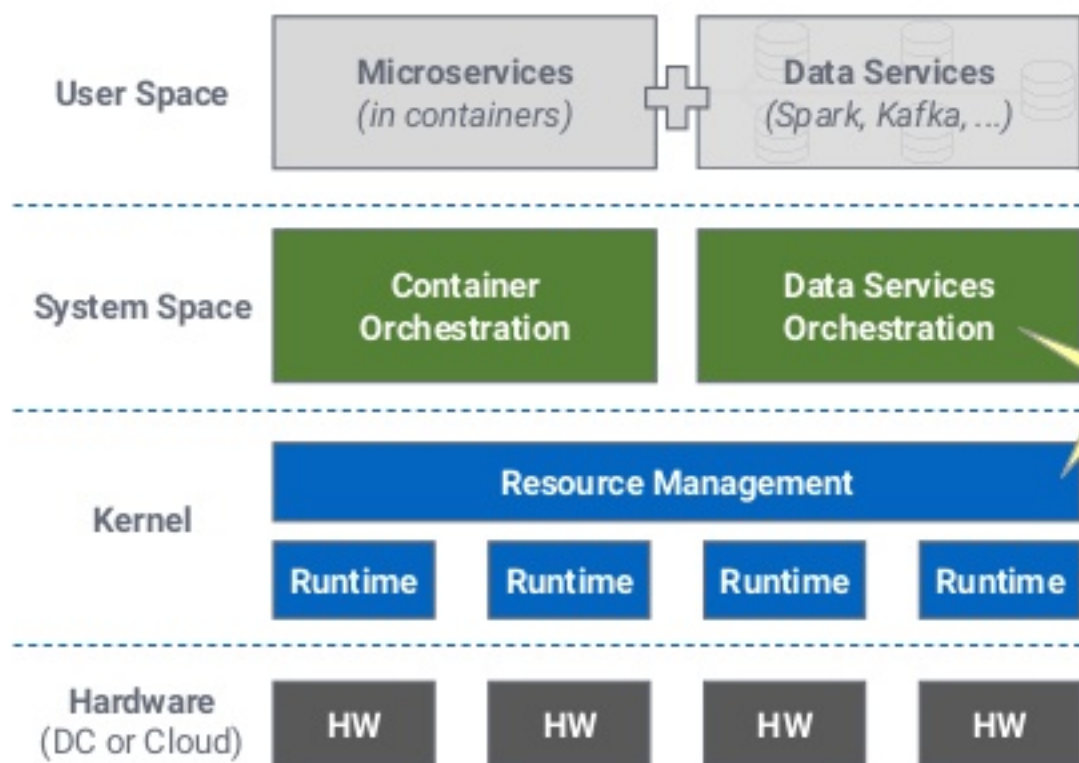
Typical Datacenter
siload, over-provisioned servers,
low utilization



DC/OS
automated schedulers, workload multiplexing onto the
same machines

DC/OS APPROACH TO CONTAINER ORCHESTRATION

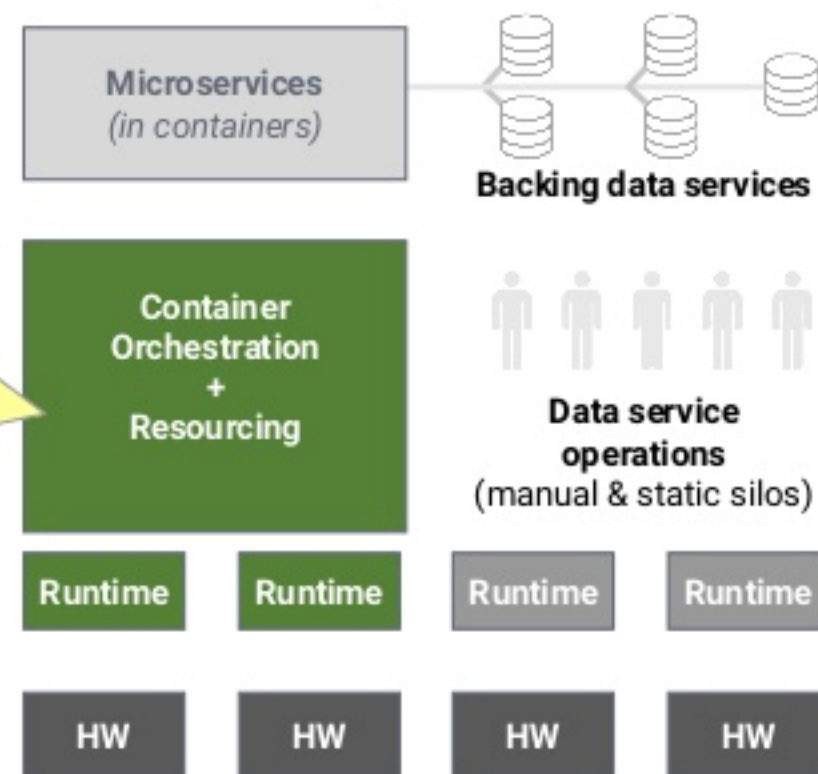
DC/OS Approach



Two-level scheduling

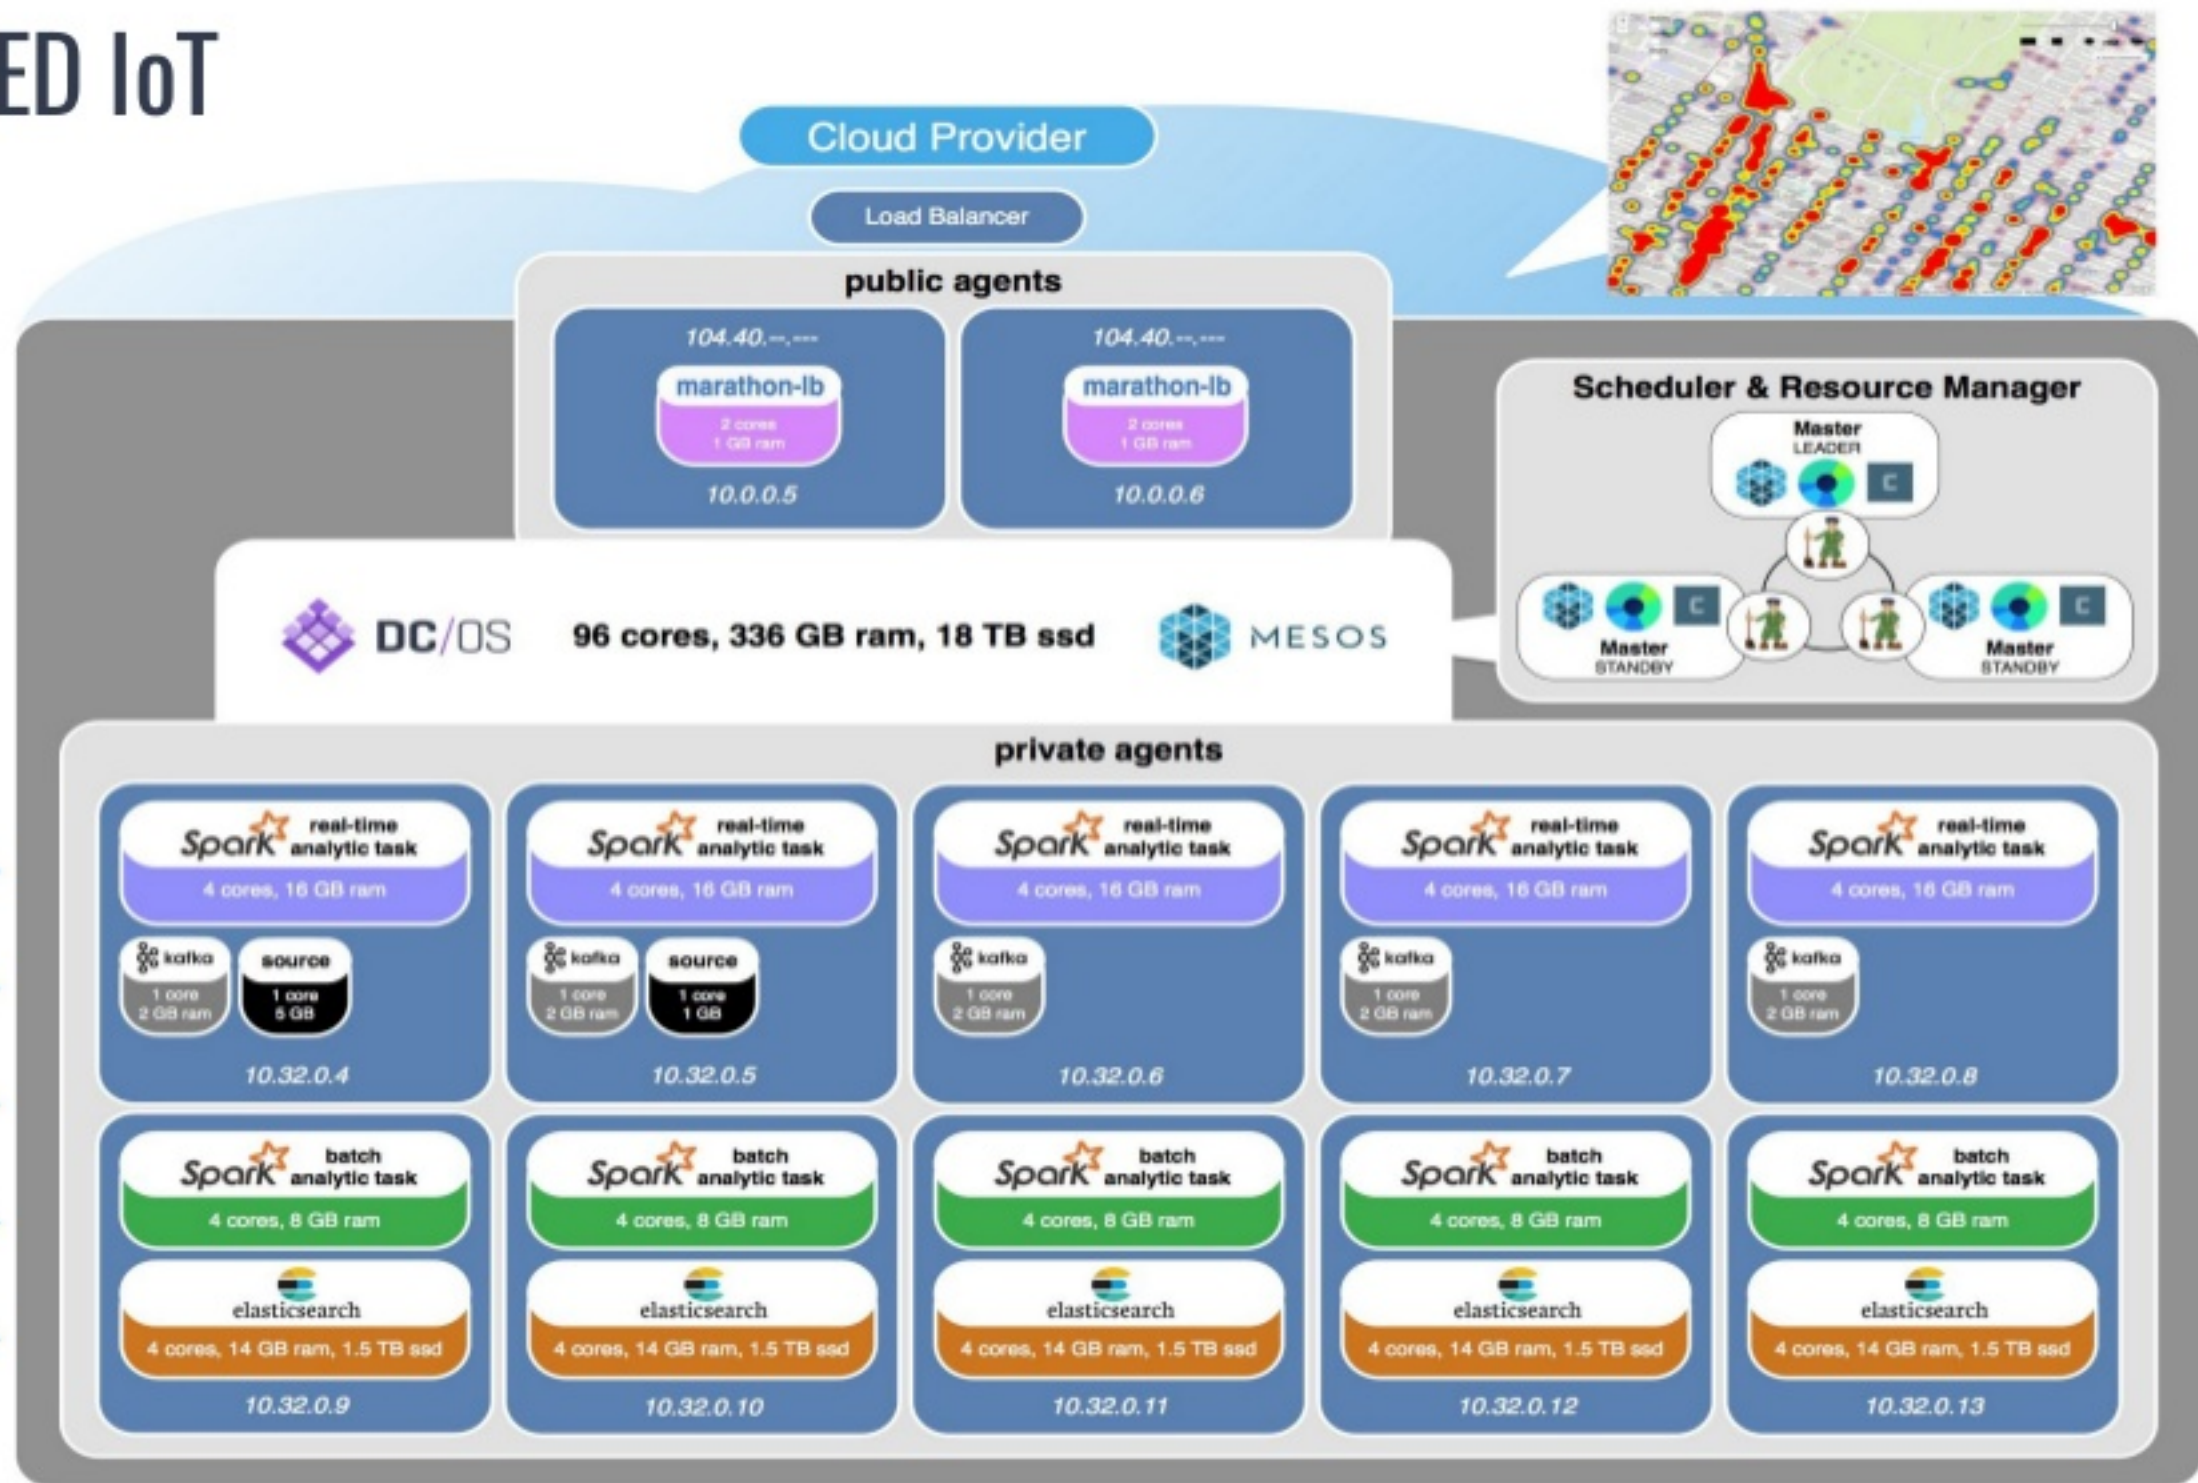
Monolithic scheduling

Container Orchestrator Approach

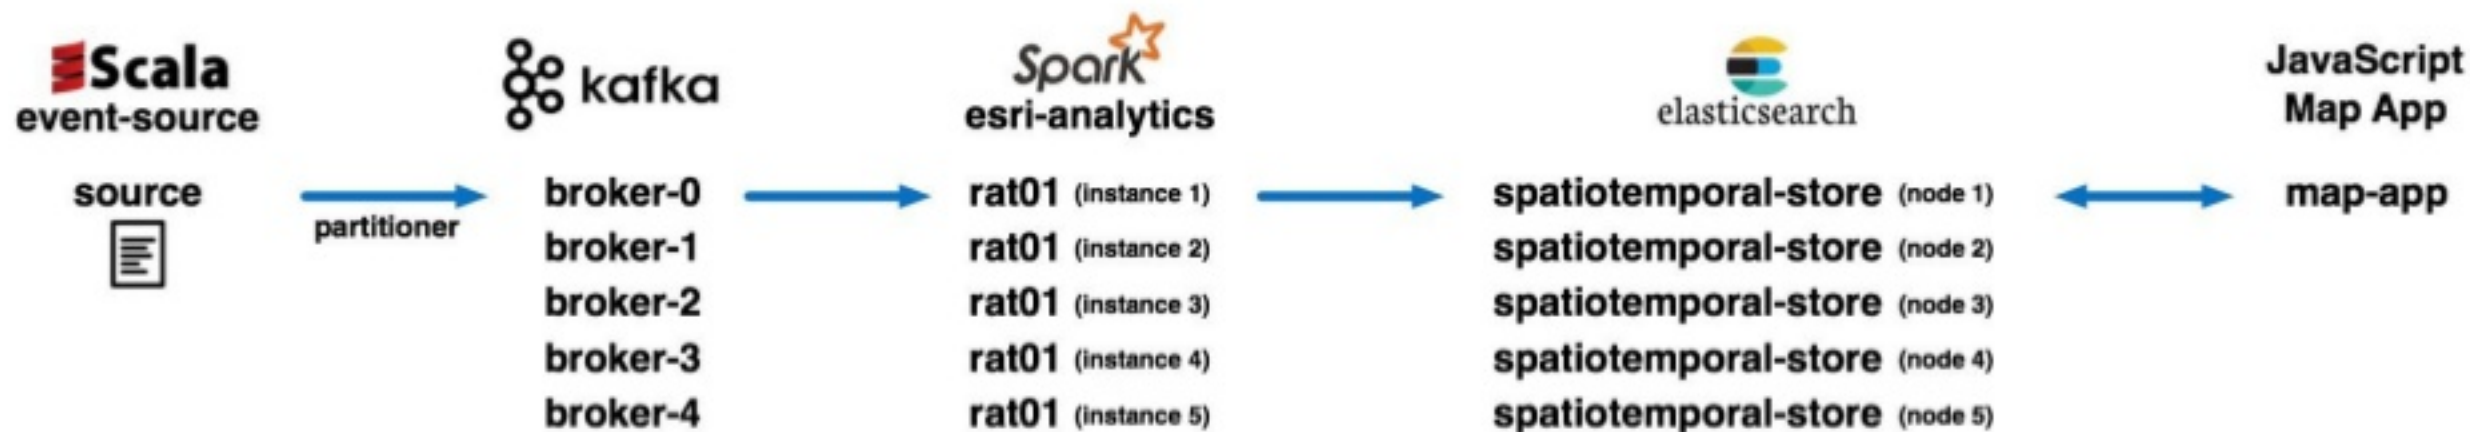


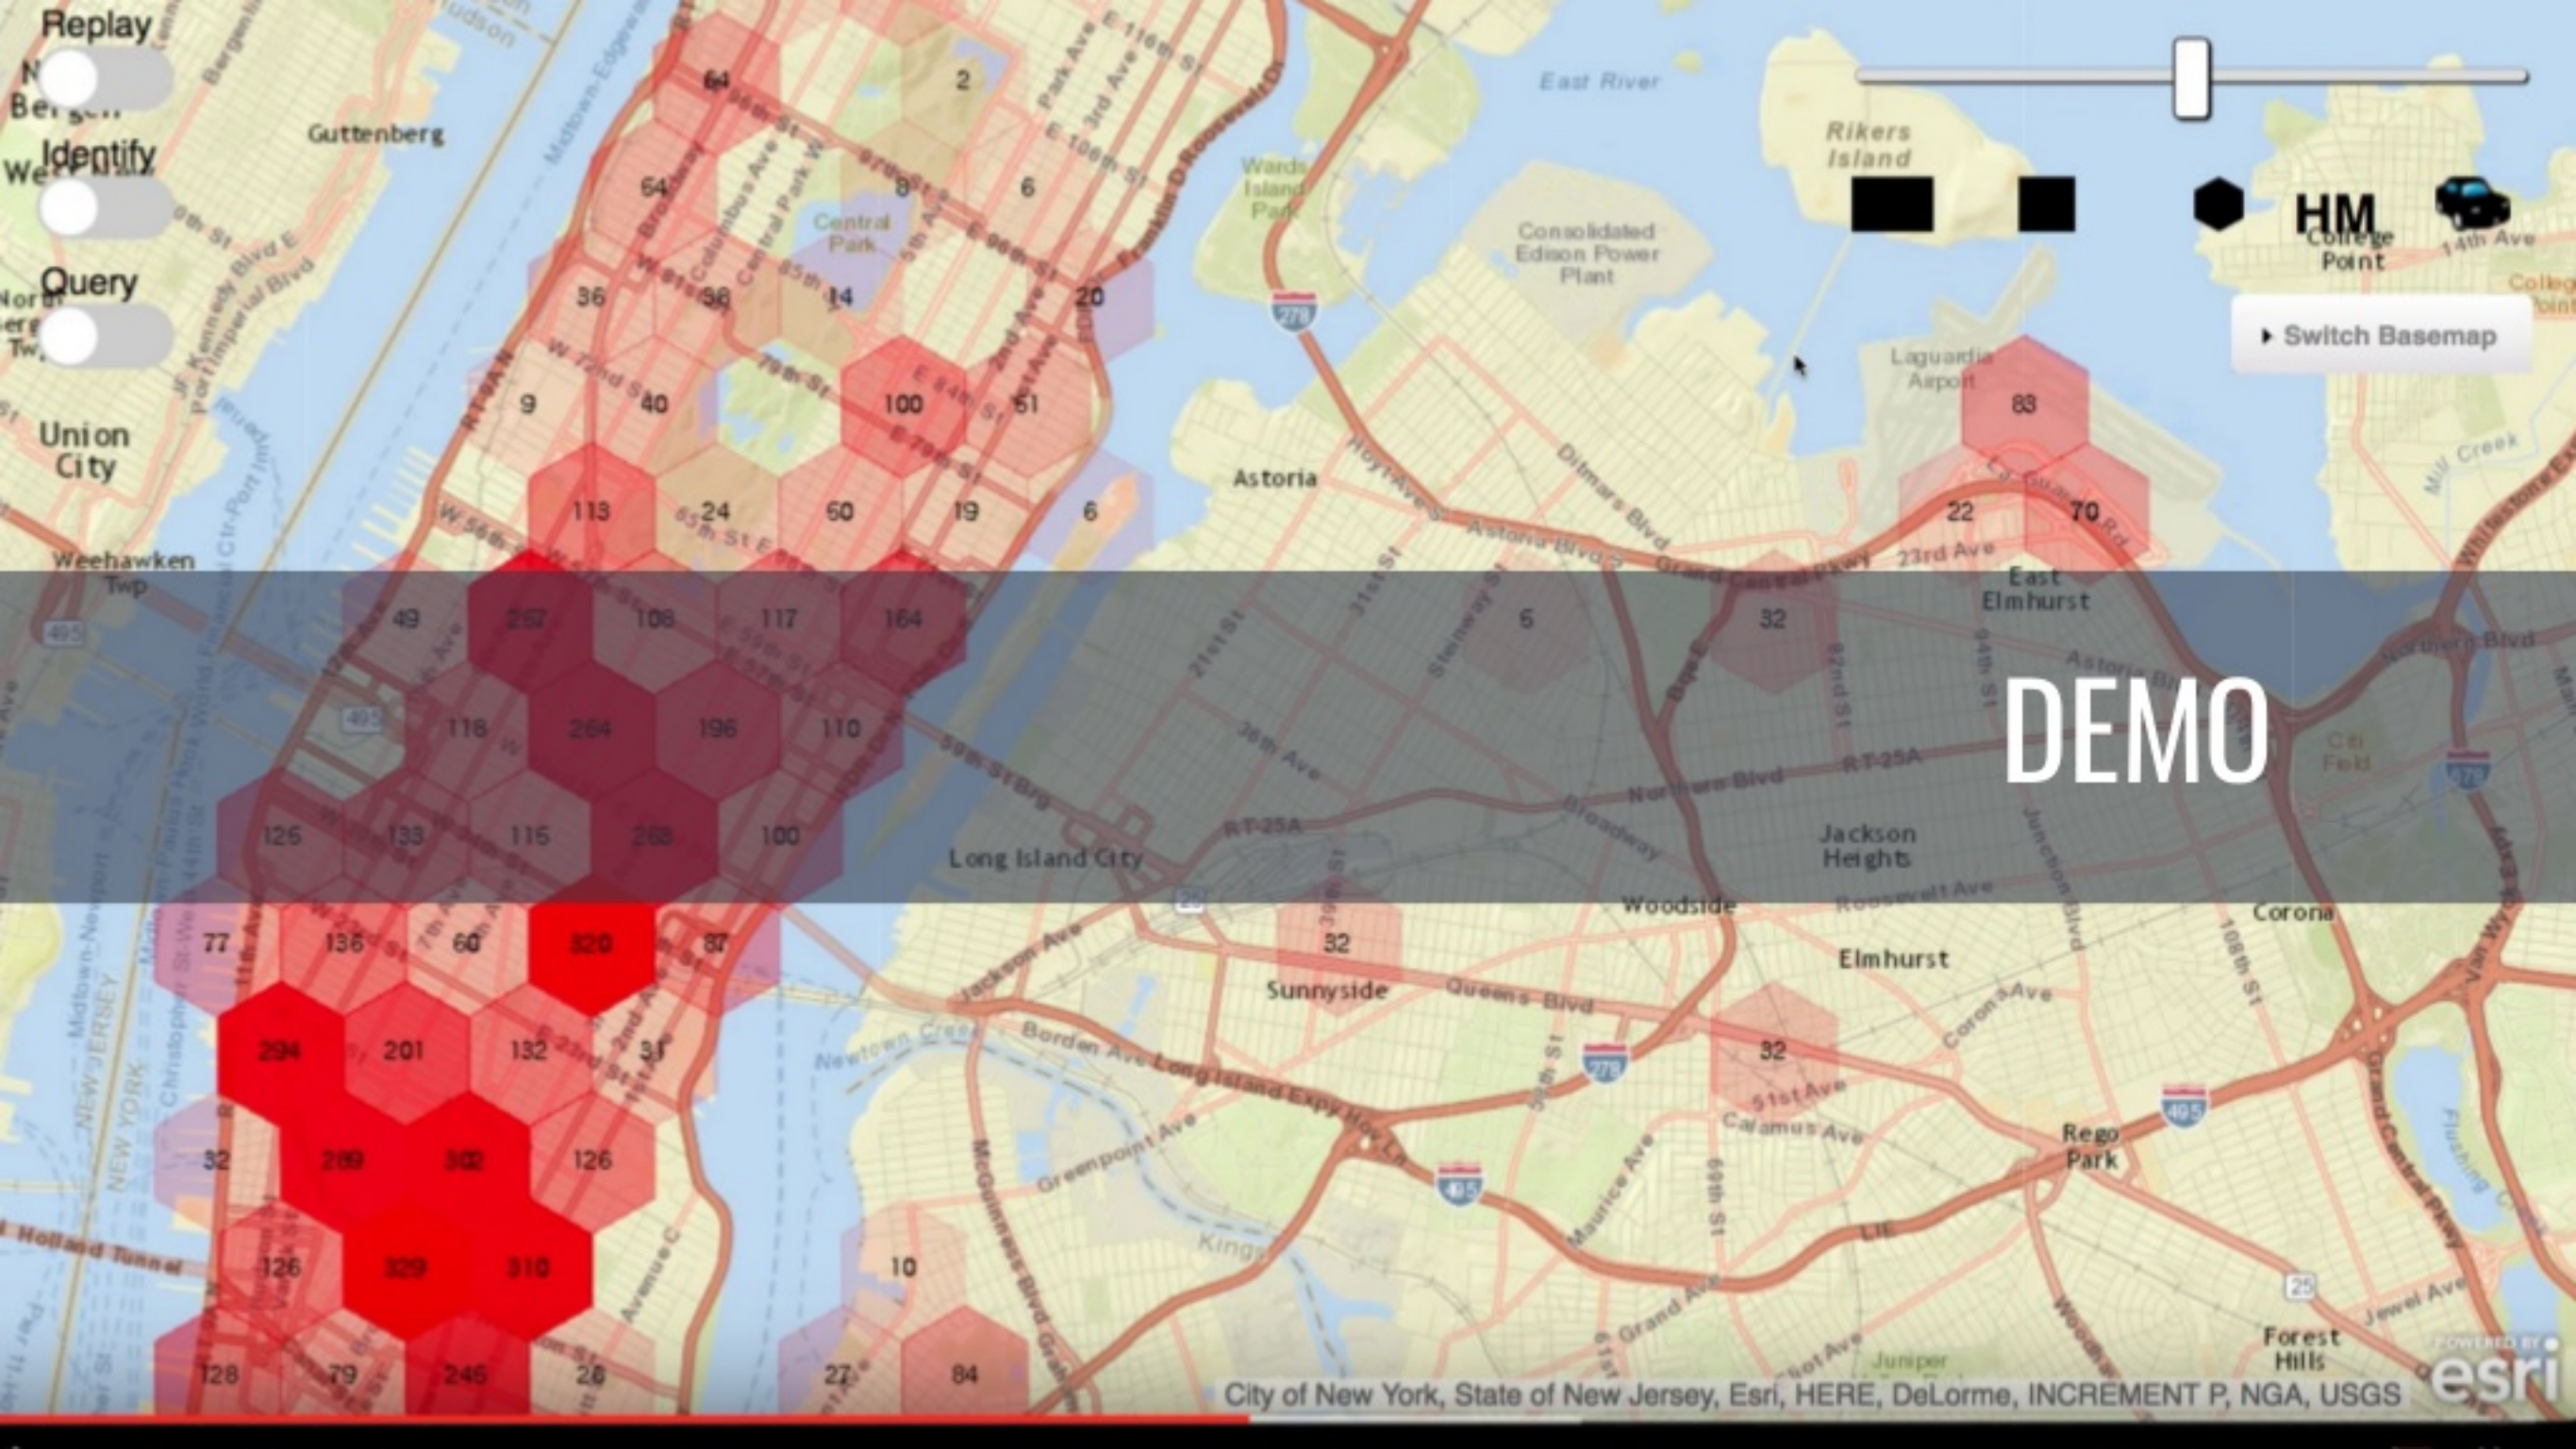
EXAMPLE: REAL-TIME CAR TRACKING

GEO-ENABLED IoT



DATA FLOW







@dcos



chat.dcos.io



users@dcos.io



/groups/8295652



/dcos
/dcos/examples
/dcos/demos

THANK YOU!

ANY QUESTIONS?