

Apache Spark – Apache HBase Connector

**Feature Rich and Efficient Access to HBase
through Spark SQL**

Weiqing Yang

Mingjie Tang

June, 2017



About Authors

◆ Weiqing Yang

- Contribute to Apache Spark, Apache Hadoop, Apache HBase, Apache Ambari
- Software Engineer at Hortonworks

◆ Mingjie Tang

- SparkSQL, Spark Mllib, Spark Streaming, Data Mining, Machine Learning
- Software Engineer at Hortonworks

◆ ... All Other SHC Contributors

Agenda

Motivation

Overview

Architecture & Implementation

Usage & Demo

Motivation

- ◆ Limited Spark Support in HBase Upstream
 - RDD level
 - But Spark Is Moving to DataFrame/Dataset
- ◆ Existing Connectors in DataFrame Level
 - Complicated Design
 - Embedding Optimization Plan inside Catalyst Engine
 - Stability Impact with Coprocessor
 - Serialized RDD Lineage to HBase
 - Heavy Maintenance Overhead

Overview

Apache Spark– Apache HBase Connector (SHC)

◆ Combine Spark and HBase

- Spark Catalyst Engine for Query Plan and Optimization
- HBase as Fast Access KV Store
- Implement Standard External Data Source with Build-in Filter, Maintain Easily

◆ Full Fledged DataFrame Support

- Spark SQL
- Integrated Language Query

◆ High Performance

- Partition Pruning, Data Locality, Column Pruning, Predicate Pushdown
- Use Spark UnhandledFilters API
- Cache Spark HBase Connections

Data Coder & Data Schema

Support Different Data Coders

- PrimitiveType: Native Support Java Primitive Types
- Avro: Native Support Avro Encoding/Decoding
- Phoenix: Phoenix Encoding/Decoding
- Plug-In Data Coder
- Can Run on the Top of Existing HBase Tables


Support Composite Key

- ```
def cat = s"""{
 "table":{"namespace":"default", "name":"shcExampleTable", "tableCoder":"Phoenix"},
 "rowkey":"key1:key2",
 "columns":{
 "col00":{"cf":"rowkey", "col":"key1", "type":"string"},
 "col01":{"cf":"rowkey", "col":"key2", "type":"int"},
 ...
 }
}
```

# Architecture & Implementation



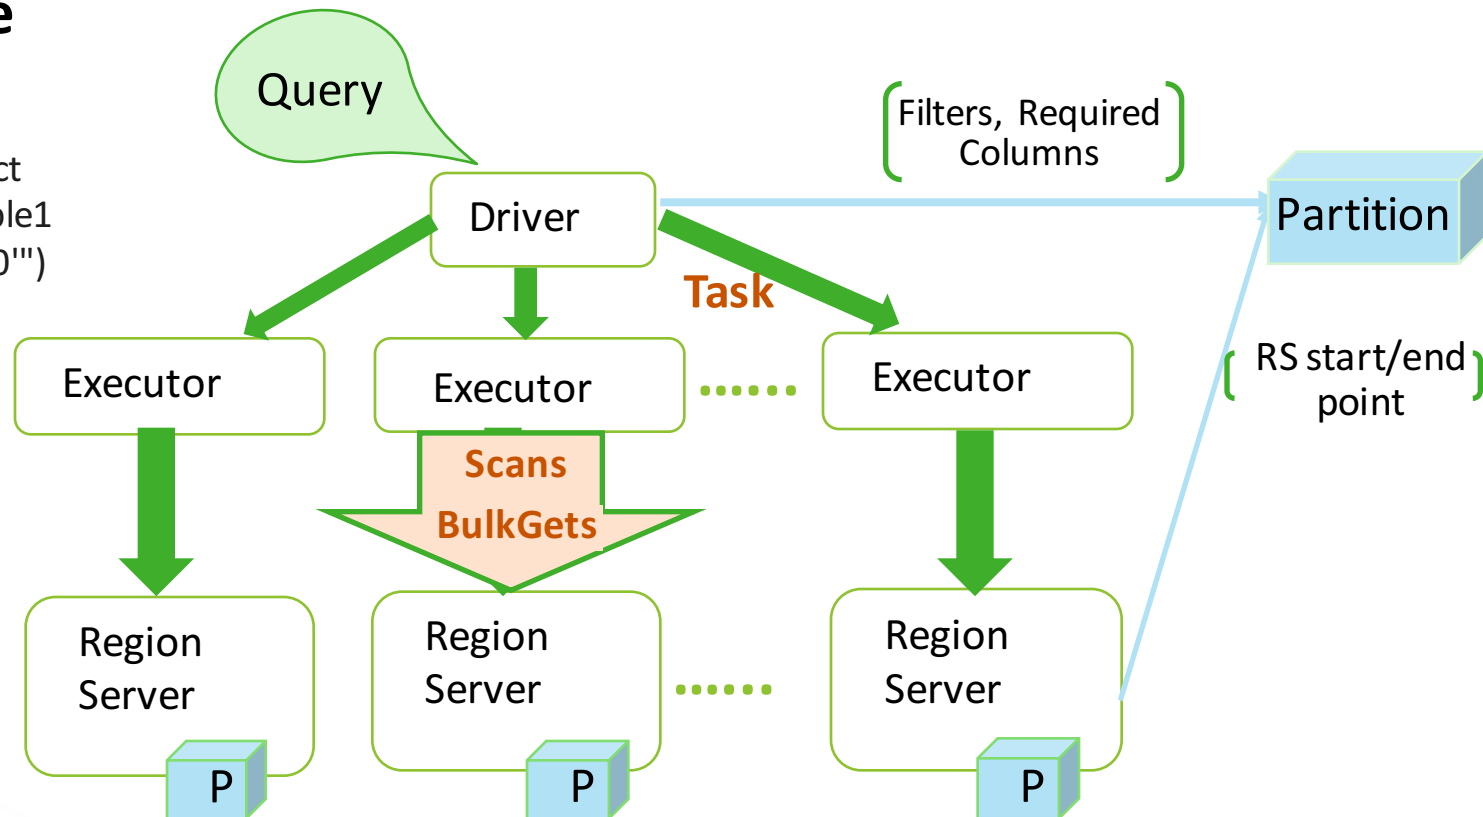
# Architecture



Picture 1. SHC architecture

# Architecture

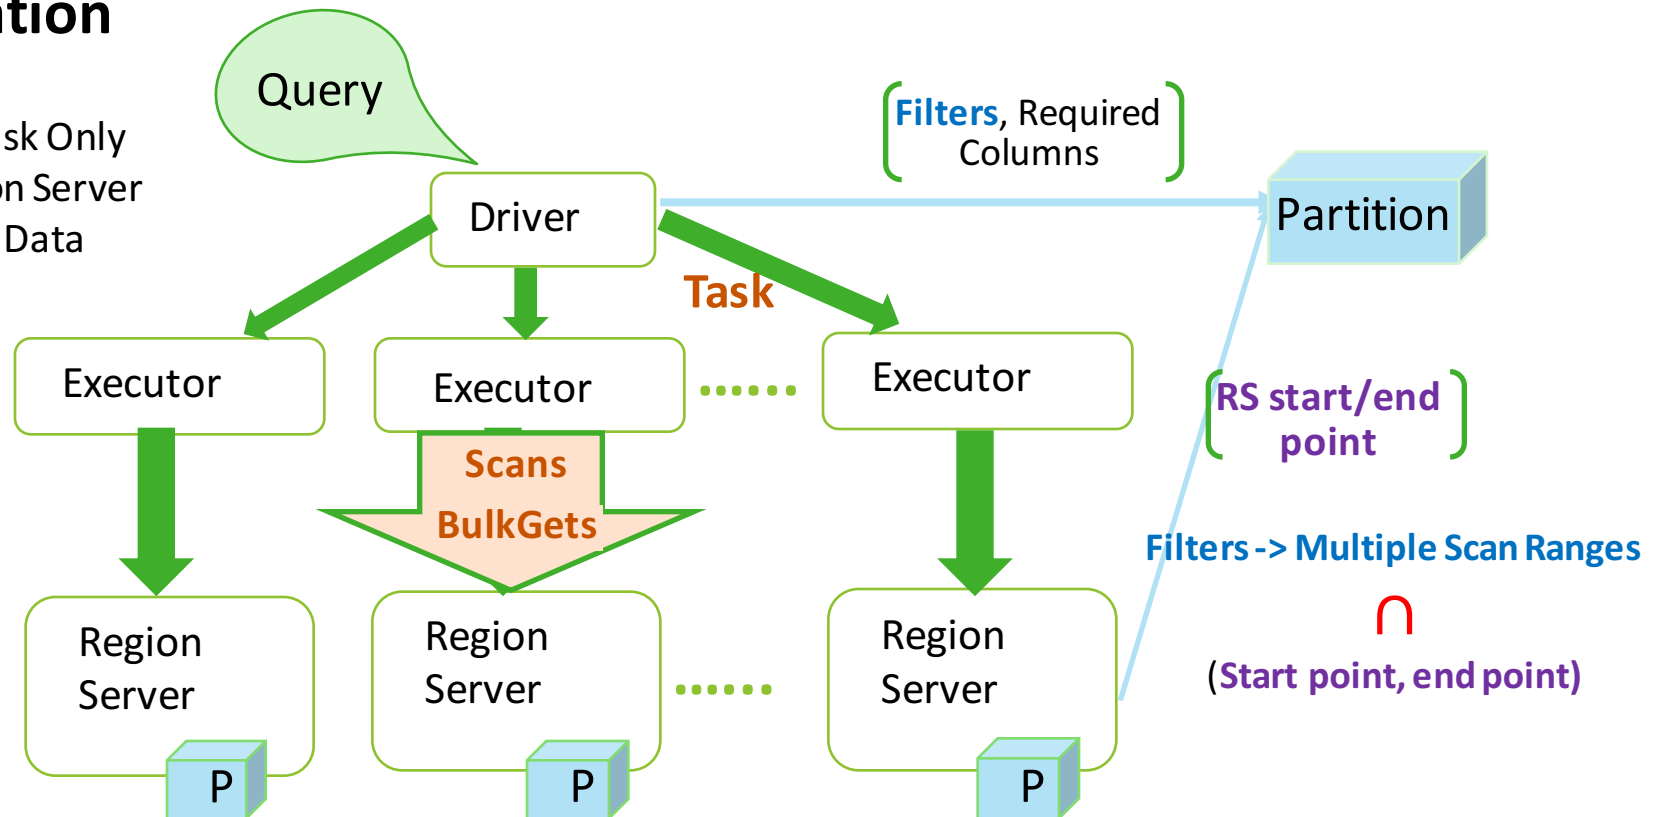
```
sqlContext.sql("select
count(col1) from table1
where key < 'row050'")
```



Picture 1. SHC architecture

# Implementation

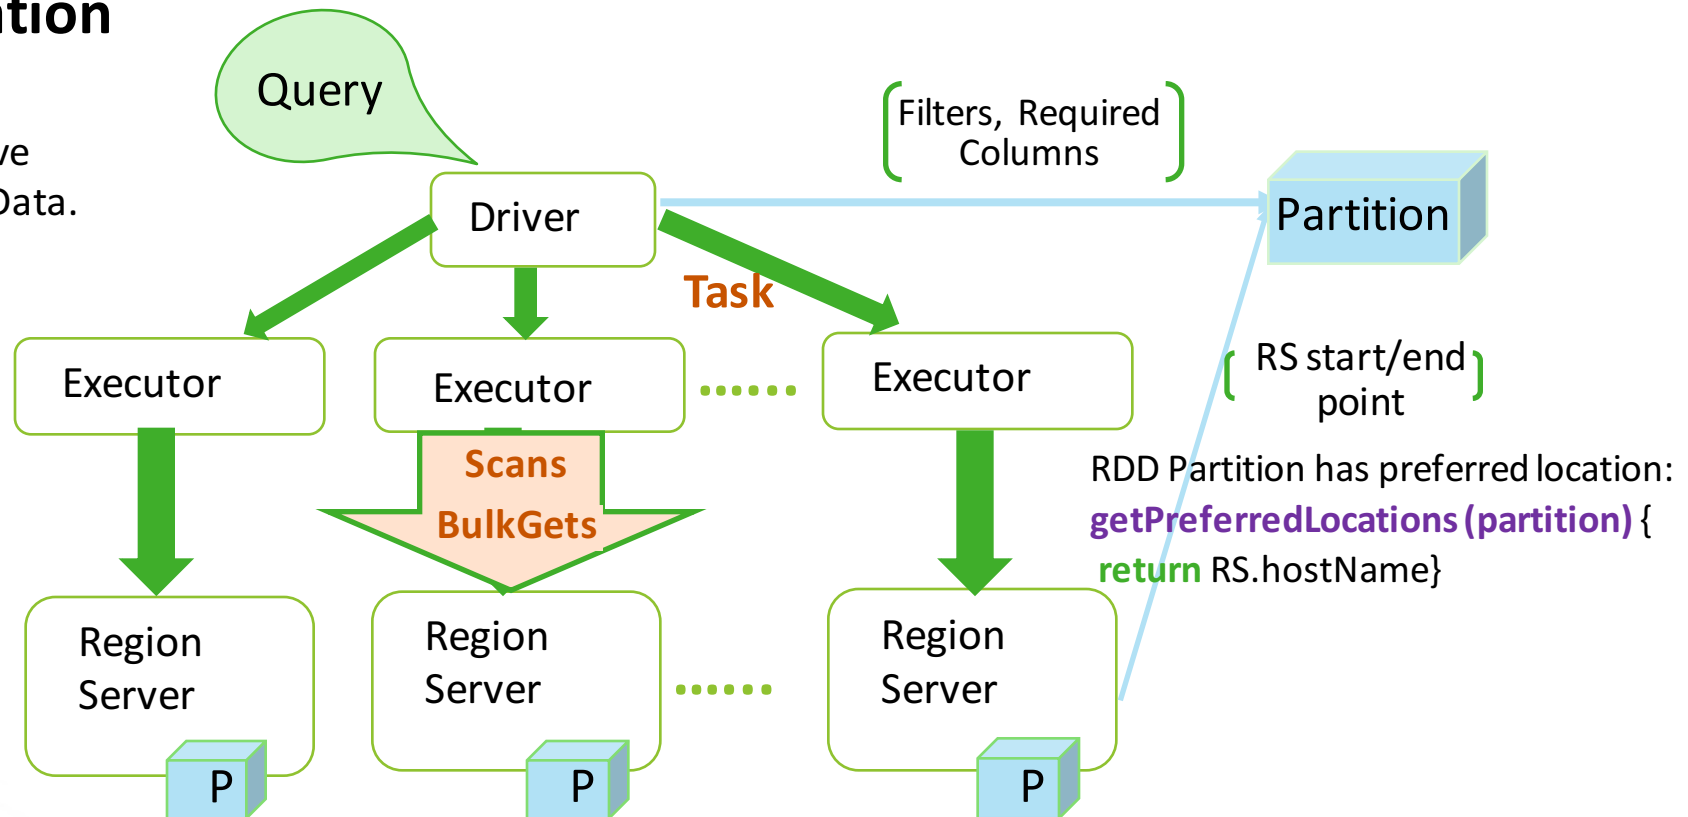
**Partition Pruning:** Task Only Performed in Region Server Holding Requested Data



Picture 1. SHC architecture

## Implementation

## Data Locality: Move Computation to Data.

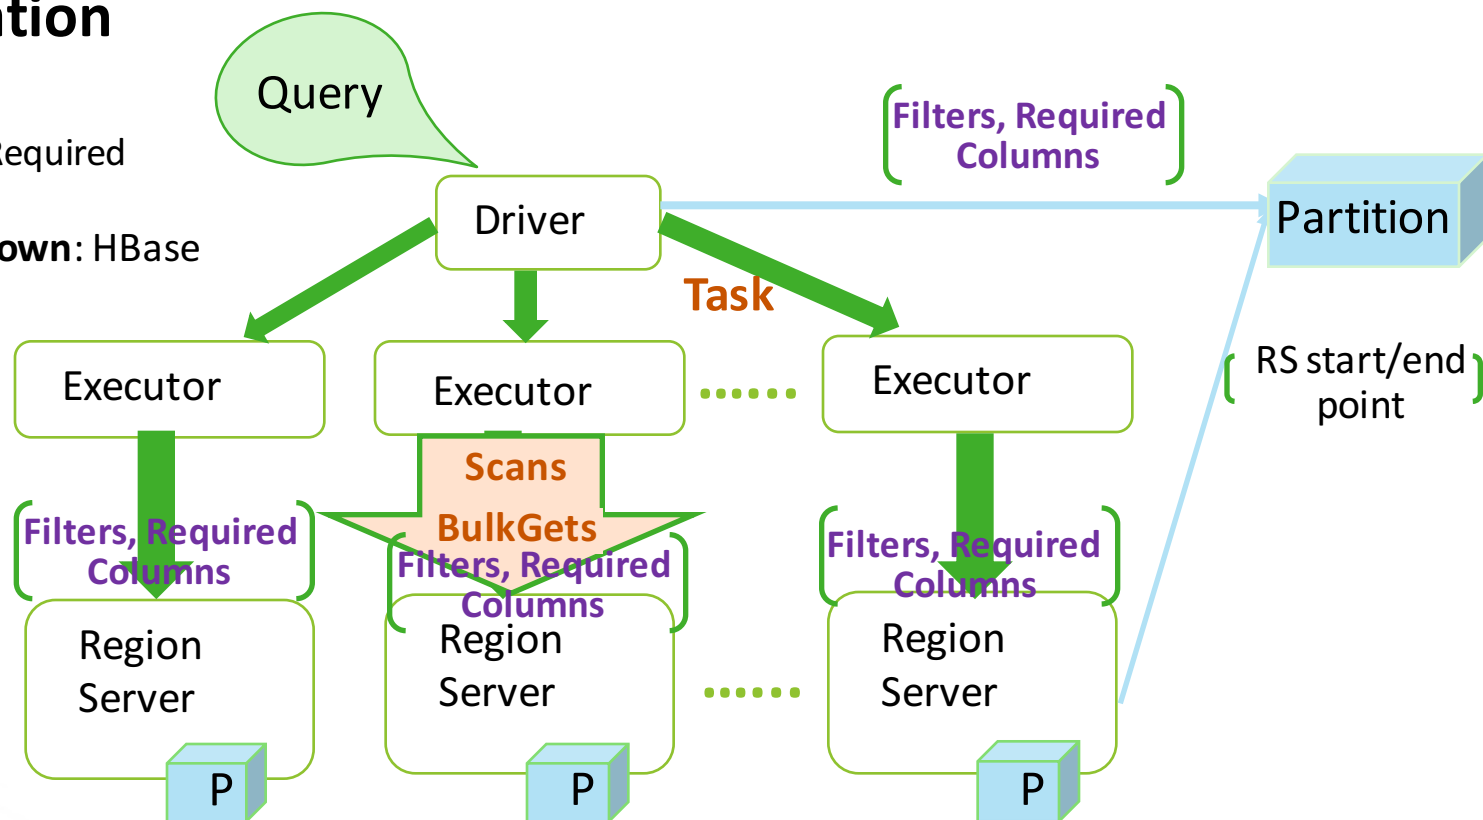


### Picture 1. SHC architecture

# Implementation

**Column Pruning:** Required Column

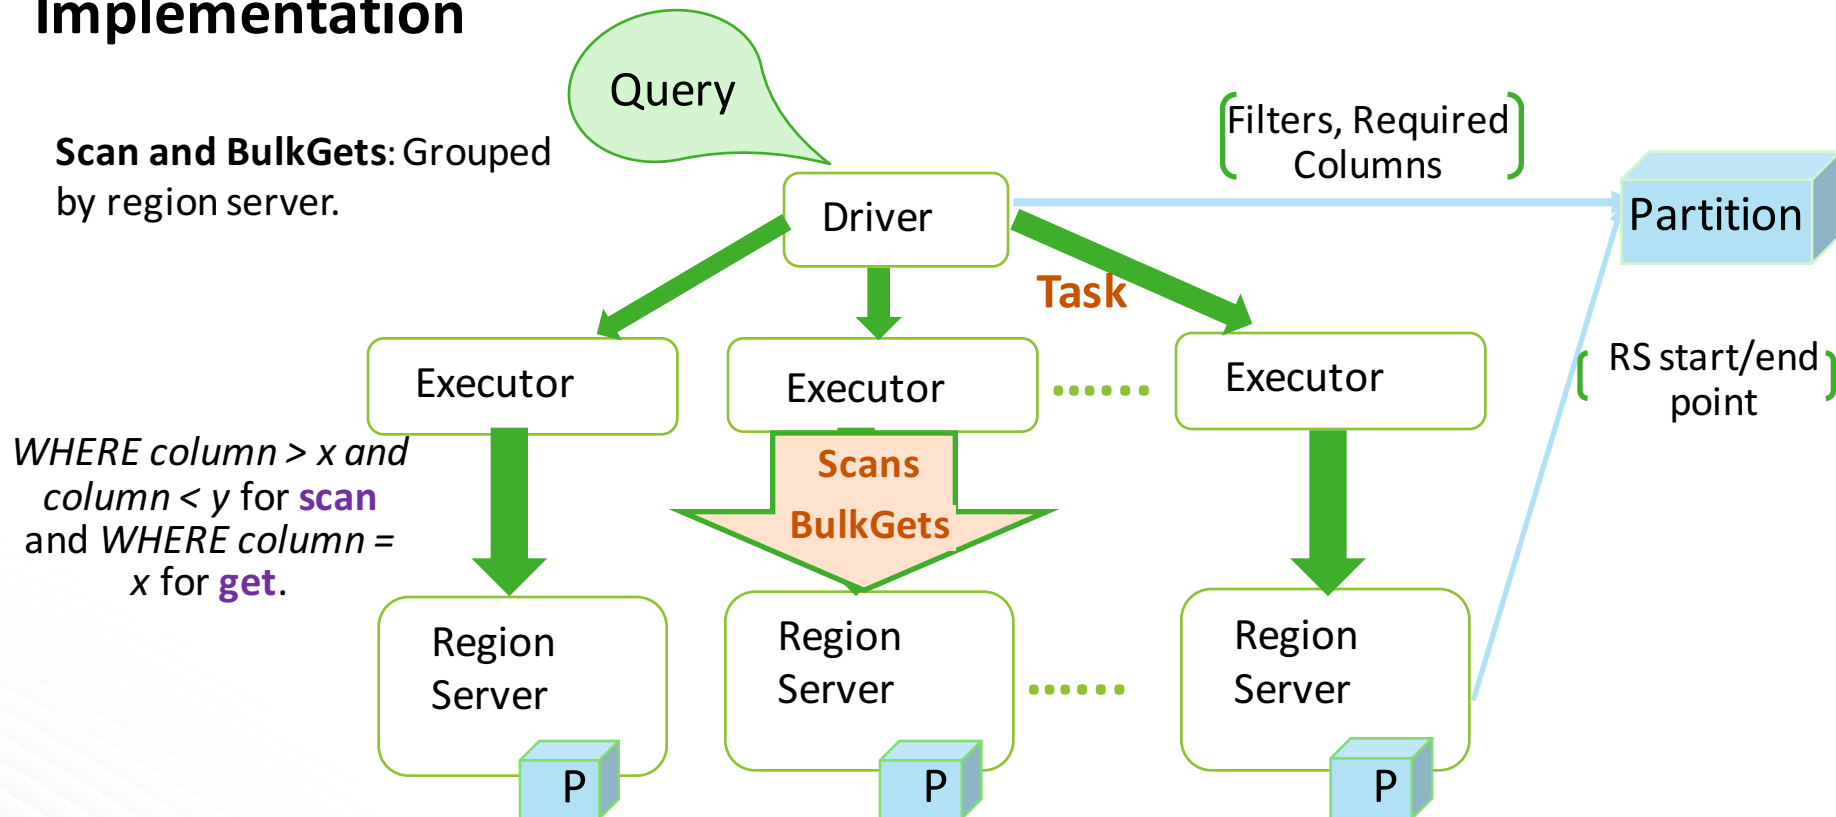
**Predicate Pushdown:** HBase built-in Filters



Picture 1. SHC architecture

# Implementation

**Scan and BulkGets:** Grouped by region server.



Picture 1. SHC architecture

# Usage & Demo

# How to Use SHC?

## ◆ Github

- <https://github.com/hortonworks-spark/shc>

## ◆ SHC Examples

- <https://github.com/hortonworks-spark/shc/tree/master/examples>

## ◆ Apache HBase Jira

- <https://issues.apache.org/jira/browse/HBASE-14789>



# Demo

- ◆ Interactive Jobs through Spark Shell
- ◆ Batch Jobs

# Acknowledgement

- ◆ HBase Community & Spark Community
- ◆ All Spark-HBase Contributors, Zhan Zhang

# Reference

## ◆ Hortonworks Public Repo

- <http://repo.hortonworks.com/content/repositories/releases/com/hortonworks/>

## ◆ Apache Spark

- <http://spark.apache.org/>

## ◆ Apache HBase

- <https://hbase.apache.org/>

# Thanks

## Q & A

Emails:  
[wyang@hortonworks.com](mailto:wyang@hortonworks.com)

# BACKUP

# Kerberos Cluster

- ◆ Kerberos Ticket
  - `kinit -kt foo.keytab foouser` or Principle/Keytab
- ◆ Long Running Service
  - `--principal, --keytab`
- ◆ Multiple Secure HBase Clusters
  - Spark only Supports Single Secure HBase Cluster
  - Use SHC Credential Manager
  - Refer [LRJobAccessing2Clusters Example in github](#)

# Usage

Define the catalog for the schema mapping:

```
def catalog = s"""{
 |"table":{"namespace":"default", "name":"phoenixTable",
 | "tableCoder":"Phoenix", "version":"2.0"},
 |"rowkey":"key",
 |"columns":{
 |"col0":{"cf":"rowkey", "col":"key", "type":"string"},
 |"col1":{"cf":"cf1", "col":"col1", "type":"boolean"},
 |"col2":{"cf":"cf2", "col":"col2", "type":"double"},
 |"col3":{"cf":"cf3", "col":"col3", "type":"float"},
 |"col4":{"cf":"cf4", "col":"col4", "type":"int"},
 |"col5":{"cf":"cf5", "col":"col5", "type":"bigint"},
 |"col6":{"cf":"cf6", "col":"col6", "type":"smallint"},
 |"col7":{"cf":"cf7", "col":"col7", "type":"string"},
 |"col8":{"cf":"cf8", "col":"col8", "type":"tinyint"}
 |}
 |}""".stripMargin
```

## Usage

### ◆ Prepare the data and populate the HBase table

```
val data = (0 to 255).map { i => HBaseRecord(i, "extra") }
```

```
sc.parallelize(data).toDF.write.options(
 Map(HBaseTableCatalog.tableCatalog -> catalog, HBaseTableCatalog.newTable -> "5"))
 .format("org.apache.spark.sql.execution.datasources.hbase")
 .save()
```



# Usage

## ◆ Load the DataFrame

```
def withCatalog(cat: String): DataFrame = {
 sqlContext
 .read
 .options(Map(HBaseTableCatalog.tableCatalog->cat))
 .format("org.apache.spark.sql.execution.datasources.hbase")
 .load()
}
```

```
val df = withCatalog(catalog)
```

# Usage

## ◆ Query

### Language integrated query:

```
val s = df.filter((($"col0" <= "row050" && $"col0" > "row040") ||
 $"col0" === "row005" && ($"col4" === 1 || $"col4" === 42))
 .select("col0", "col1", "col4")
```

### SQL:

```
val s = df.filter((($"col0" <= "row050" && $"col0" > "row040")
df.registerTempTable("table")
sqlContext.sql("select count(col1) from table").show
```

# Usage

## ◆ Work with different data sources

*// Part 1: write data into Hive table and read data from it*

```
val df1 = sql("SELECT * FROM shcHiveTable")
```

*// Part 2: read data from Hbase table*

```
val df2 = withCatalog(cat)
```

*// Part 3: join the two dataframes*

```
val s1 = df1.filter($"key" <= "40").select("key", "col1")
```

```
val s2 = df2.filter($"key" <= "20" && $"key" >= "1").select("key", "col2")
```

```
val result = s1.join(s2, Seq("key"))
```

```
result.show()
```