



Spark on Supercomputers: A Tale of the Storage Hierarchy

Costin Iancu, Khaled Ibrahim – LBNL

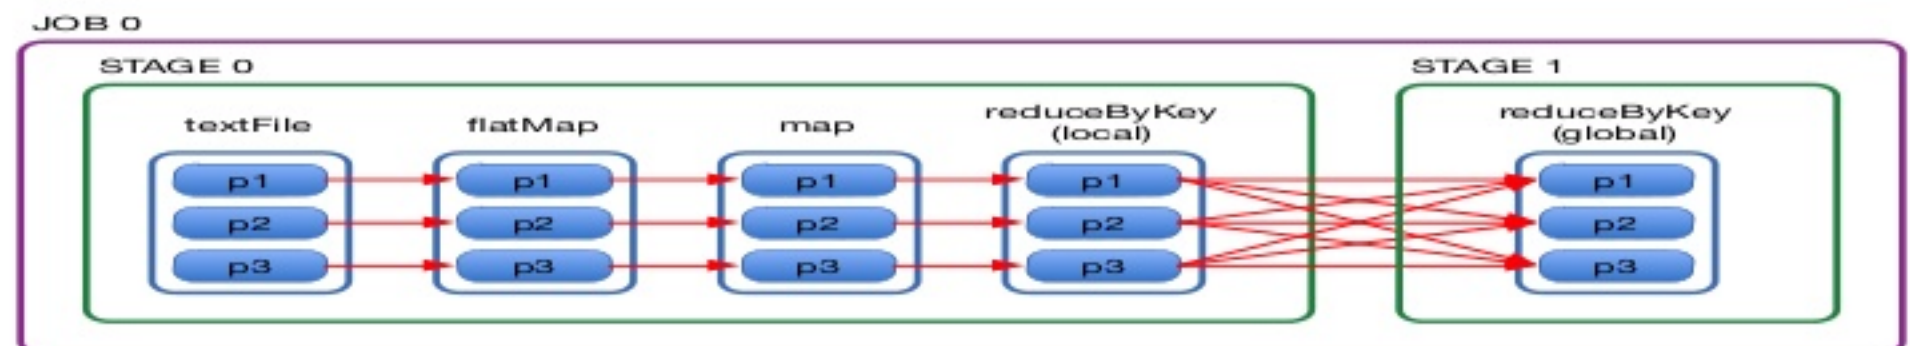
Nicholas Chaimov – U. Oregon

Apache Spark

- Developed for cloud environments
- Specialized runtime provides for
 - Performance 😊, Elastic parallelism, Resilience
- Programming productivity through
 - HLL front-ends (Scala, R, SQL), multiple domain-specific libraries: Streaming, SparkSQL, SparkR, GraphX, Splash, MLlib, Velox
- We have huge datasets but little penetration in HPC

Apache Spark

- In-memory Map-Reduce framework
- Central abstraction is the **Resilient Distributed Dataset**.
- **Data movement is important**
 - Lazy, on-demand
 - Horizontal (node-to-node) – shuffle/Reduce
 - Vertical (node-to-storage) - Map/Reduce



Data Centers/Clouds

Node local storage, assumes all disk operations are equal

Disk I/O optimized for latency

Network optimized for bandwidth




HPC

Global file system, asymmetry expected


Disk I/O optimized for bandwidth

Network optimized for latency



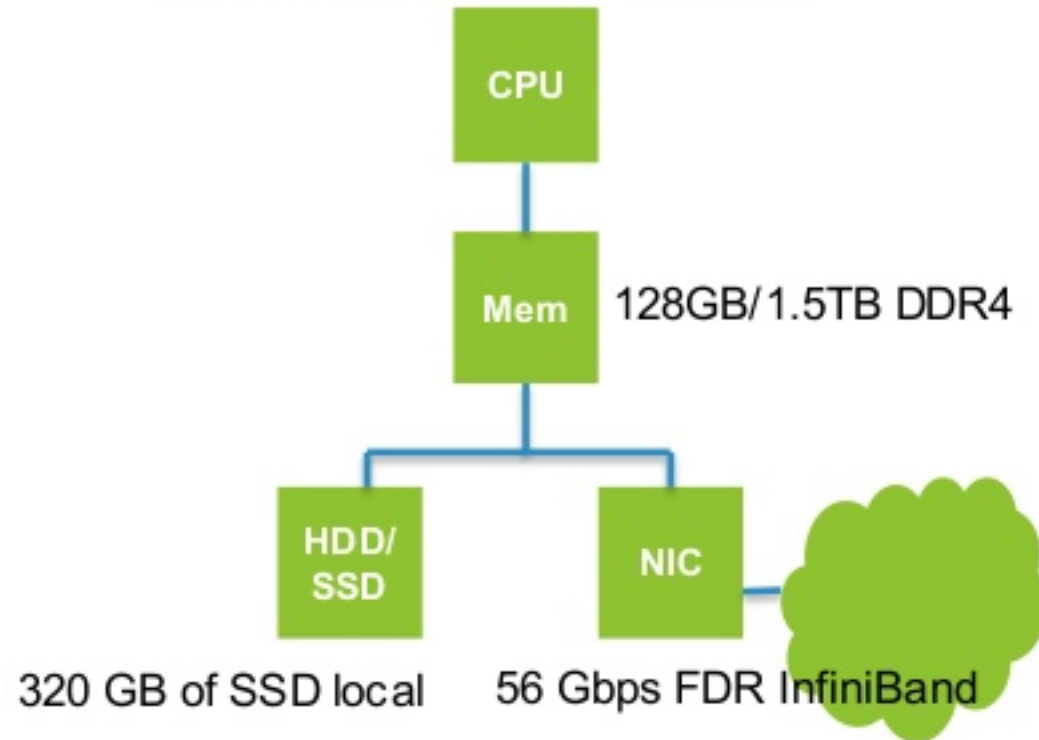
Cloud: commodity CPU,
memory, HDD/SSD NIC

Data appliance: server CPU,
large fast memory, fast SSD



HPC: server CPU, fast memory,
combo of fast and slower storage

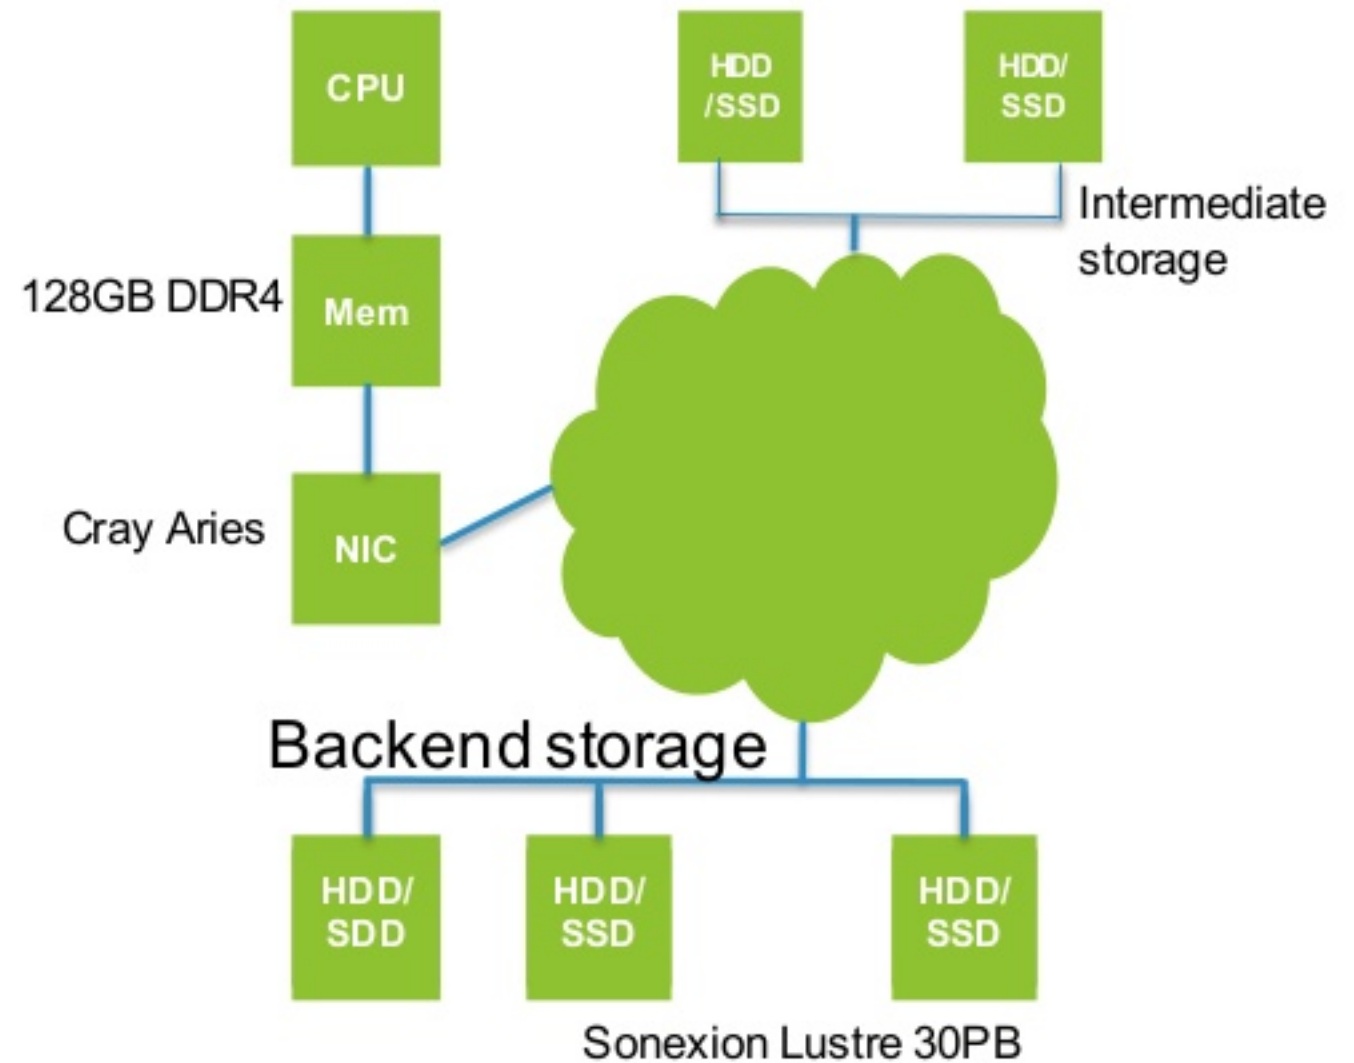
2.5 GHz Intel Haswell - 24 cores



Comet (DELL)

2.3 GHz Intel Haswell – 32 cores

Cray Data Warp
1.8PB at 1.7TB/s




Cori (Cray XC40)

Scaling Spark on Cray XC40

(It's all about file system metadata)


Not ALL I/O is Created Equal



Shuffle opens = # Shuffle reads $\rightarrow O(\text{cores}^2)$

Time per open increases with scale, unlike read/write

I/O Variability is HIGH



fopen is a problem:

- Mean time is 23X larger than SSD
- Variability is 14,000X

Improving I/O Performance

Eliminate file metadata operations

1. Keep files open (cache fopen)

- Surprising 10%-20% improvement on data appliance
- Argues for user level file systems, gets rid of serialized system calls

2. Use file system backed by single Lustre file for shuffle

- This should also help on systems with local SSDs


3. Use containers

- Speeds up startup, up to 20% end-to-end performance improvement

- Solutions need to be used in conjunction
 - E.g. fopen from Parquet reader

Plenty of details in “Scaling Spark on HPC Systems”. HPDC 2016

Scalability



File-Backed Filesystems

- **NERSC Shifter** (container infrastructure for HPC)
 - Compatible with Docker images
 - Integrated with Slurm scheduler
 - Can control mounting of filesystems within container
- **Per-Node Cache**
 - File-backed filesystem mounted within each node's container instance at common path (/mnt)
 - `--volume=$SCRATCH/backingFile:/mnt:perNodeCache=size=100G`
 - File for each node is created stored on backend Lustre filesystem
 - Single file open — intermediate data file opens are kept local



Now the fun part 😊

Architectural Performance Considerations

Cori

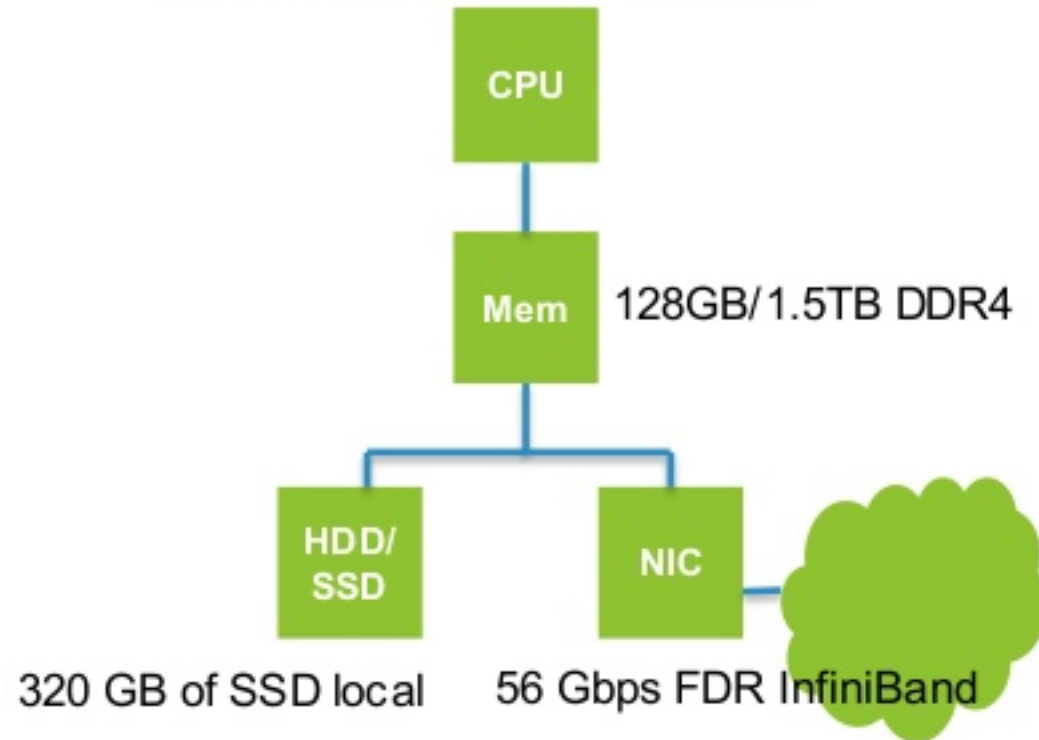
The Supercomputer

vs

Comet

The Data Appliance

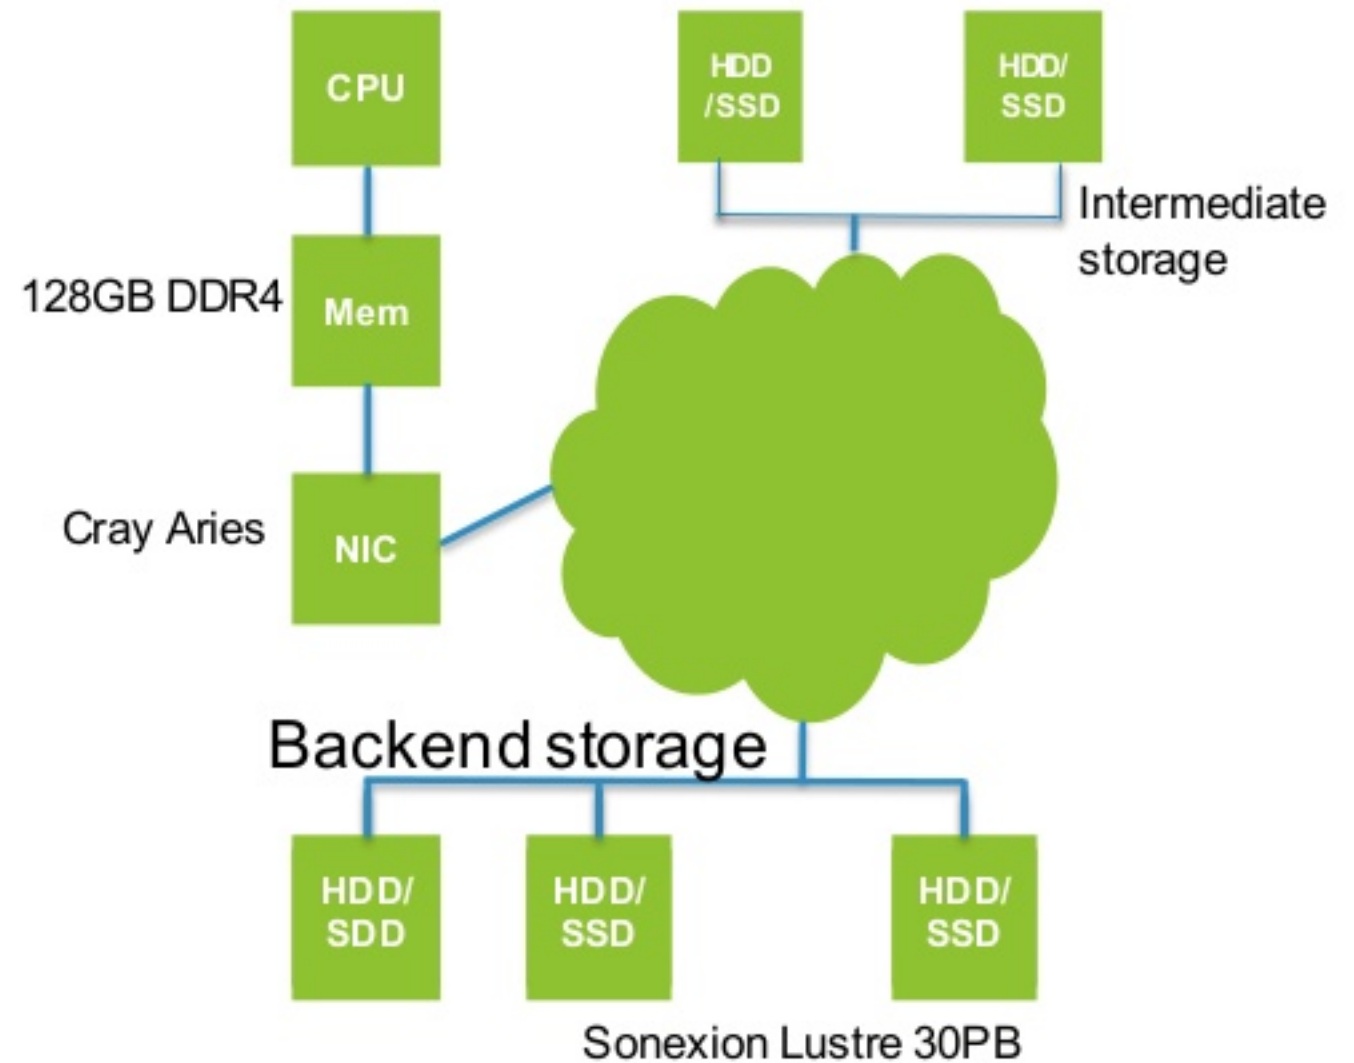
2.5 GHz Intel Haswell - 24 cores



Comet (DELL)

2.3 GHz Intel Haswell – 32 cores

Cray Data Warp
1.8PB at 1.7TB/s



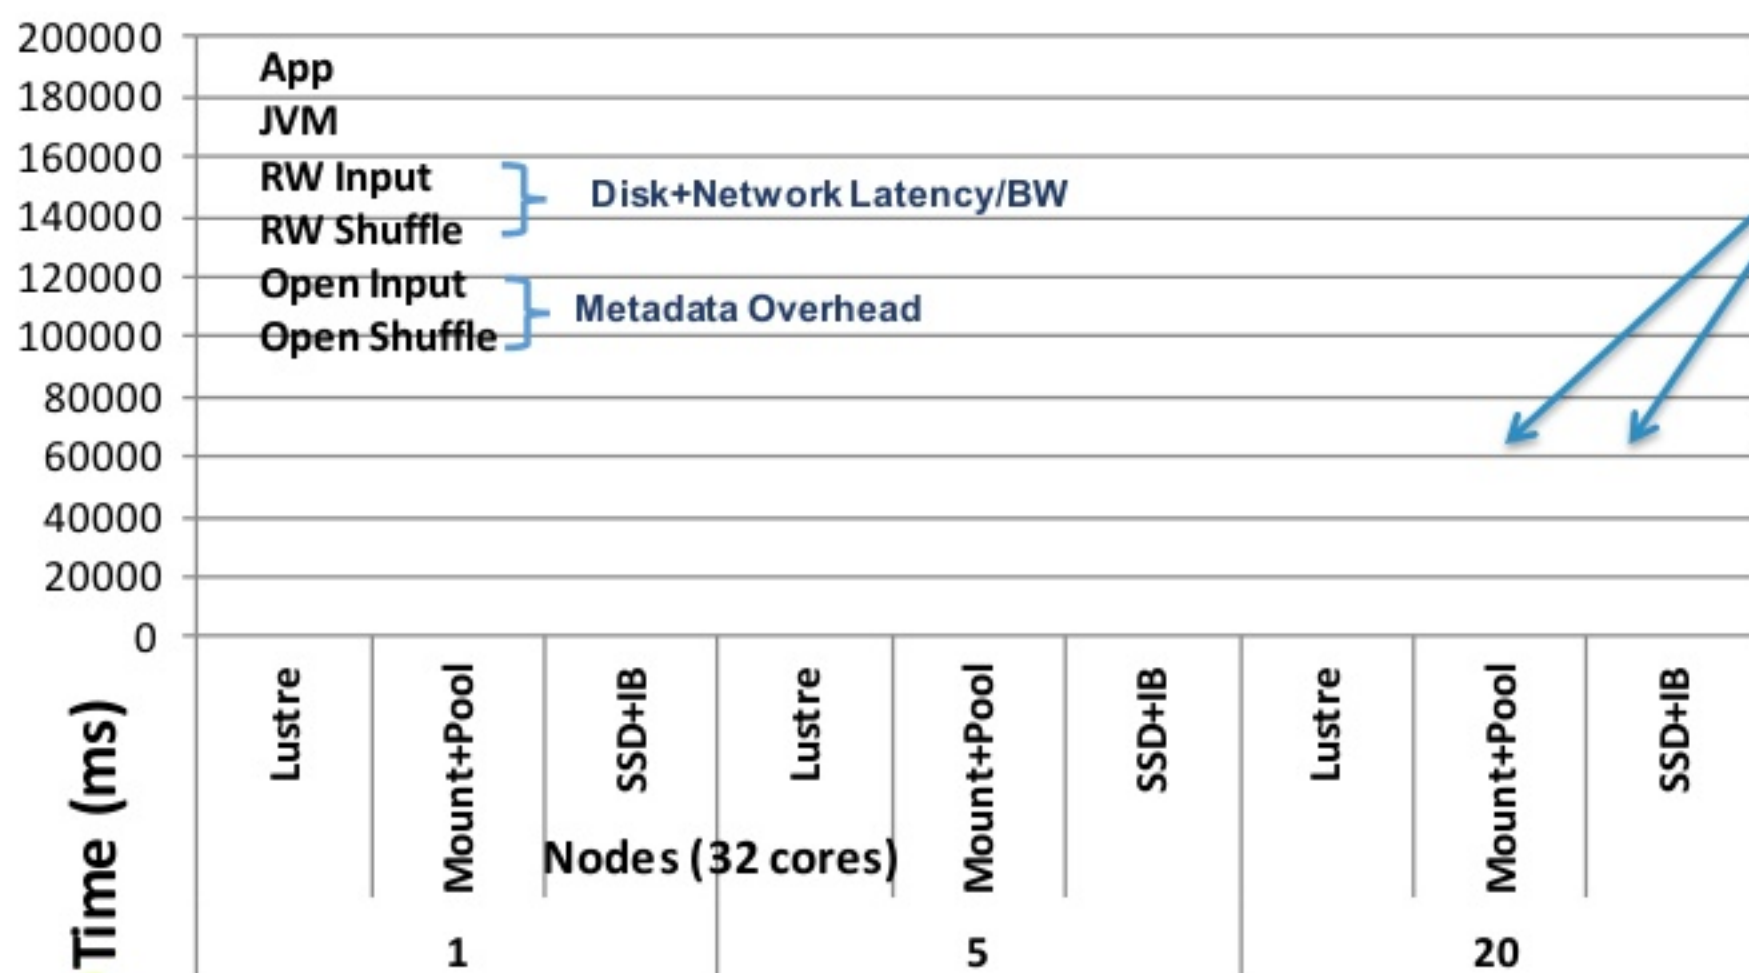
Cori (Cray XC40)

CPU, Memory, Network, Disk?

- Multiple extensions to Blocked Time Analysis (Ousterhout, 2015)
- BTA indicated that CPU dominates
 - Network 2%, disk 19%
- Concentrate on scaling out, weak scaling studies
 - Spark-perf, BigDataBenchmark, TPC-DS, TeraSort
- Interested in determining right ratio, machine balance for
 - CPU, memory, network, disk ...
- Spark 2.0.2 & Spark-RDMA 0.9.4 from Ohio State University, Hadoop 2.6


Storage hierarchy and performance

Global Storage Matches Local Storage




- Variability matters more than advertised latency and bandwidth number
- Storage performance obscured/mitigated by network due to client/server in BlockManager
 - Small scale local is slightly faster
 - Large scale global is faster

Global Storage Matches Local Storage




Intermediate Storage Hurts Performance



(Without our optimizations, intermediate storage scaled better)

Networking performance

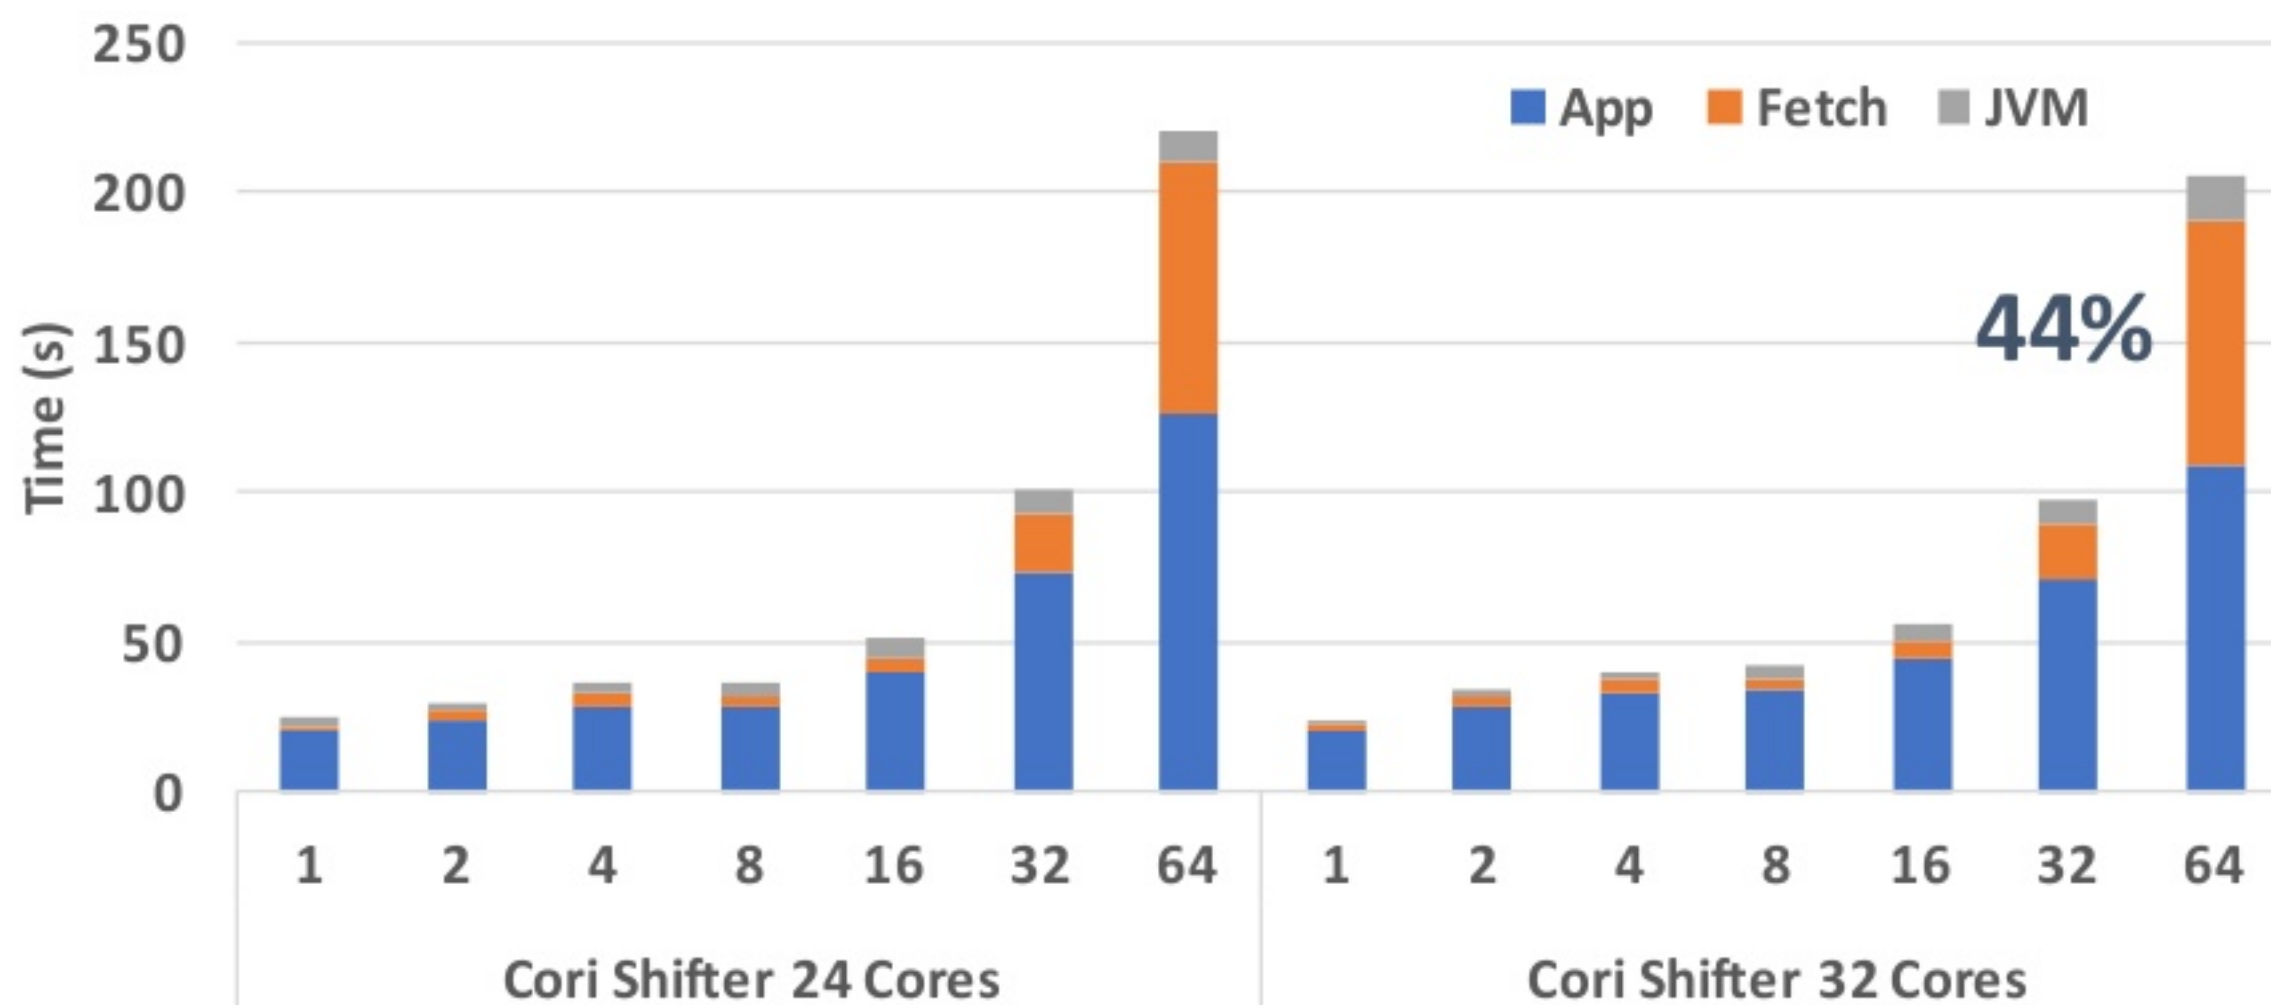
Latency or Bandwidth?



Average message size for spark-perf is 43B


Network Matters at Scale

Average Across Benchmarks



CPU

More cores or better memory?



- Need more cores to hide disk and network latency at scale.
- Preliminary experiences with Intel KNL are bad
 - Too much concurrency
 - Not enough integer throughput
- Execution does not seem to be memory bandwidth limited

Summary/Conclusions

- Latency and bandwidth are important, but not dominant
 - Variability more important than marketing numbers
- Network time dominates at scale
 - Network, disk is mis-attributed as CPU
- Comet matches Cori up to 512 cores, Cori twice as fast at 2048 cores
 - Spark can run well on global storage
- Global storage opens the possibility of global name space, no more client-server

Acknowledgement

Work partially supported by



Intel Parallel Computing Center: Big Data Support
for HPC



Thank You.

Questions, collaborations, free software


cciancu@lbl.gov

kzibrahim@lbl.gov

nchaimov@uoregon.edu

Setup

- Cray XC30 at NERSC (Edison): 2.4 GHz IvyBridge - Global
- Cray XC40 at NERSC (Cori): 2.3 GHz Haswell + Cray DataWarp



- Comet at SDSC: 2.5GHz Haswell, InfiniBand FDR, 320 GB SSD, 1.5TB memory - LOCAL