



# Real-time Machine Learning with Redis-ML and Apache Spark

Shay Nativ - Redis Labs

# Agenda

- Intro to Redis and Redis Labs - 5 min
- Using Redis-ML for Model Serving - why and how - 10 min
- Building a recommendation system using Spark-ML and Redis-ML - 10 min
- QA

# Redis Labs – Home of Redis



The commercial company behind Open Source Redis




Provider of the **Redis Enterprise (Redis<sup>e</sup>)** technology, platform and products

*Founded in 2011*

*HQ in Mountain View CA, R&D center in Tel-Aviv IL*

# Redis Labs Products



# A Brief Overview of Redis

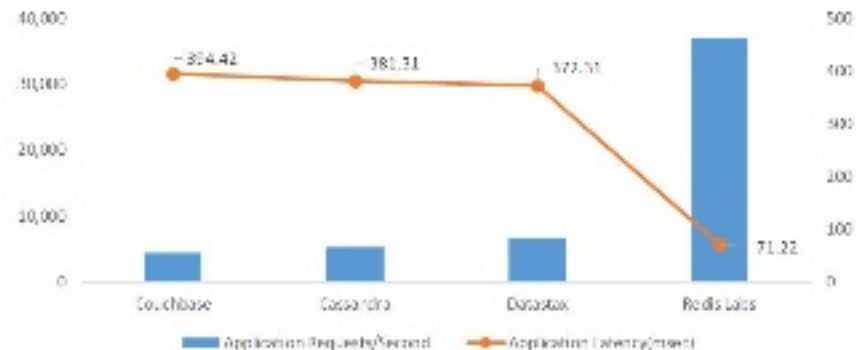
- Started in 2009 by Salvatore Sanfilippo
- Most popular KV store
- In memory - disk backed
- Notable Users:
  - Twitter, Netflix, Uber, Groupon, Twitch
  - Many, many more...



redis

# Redis Main Differentiations


NoSQL Performance Benchmark



**Performance**



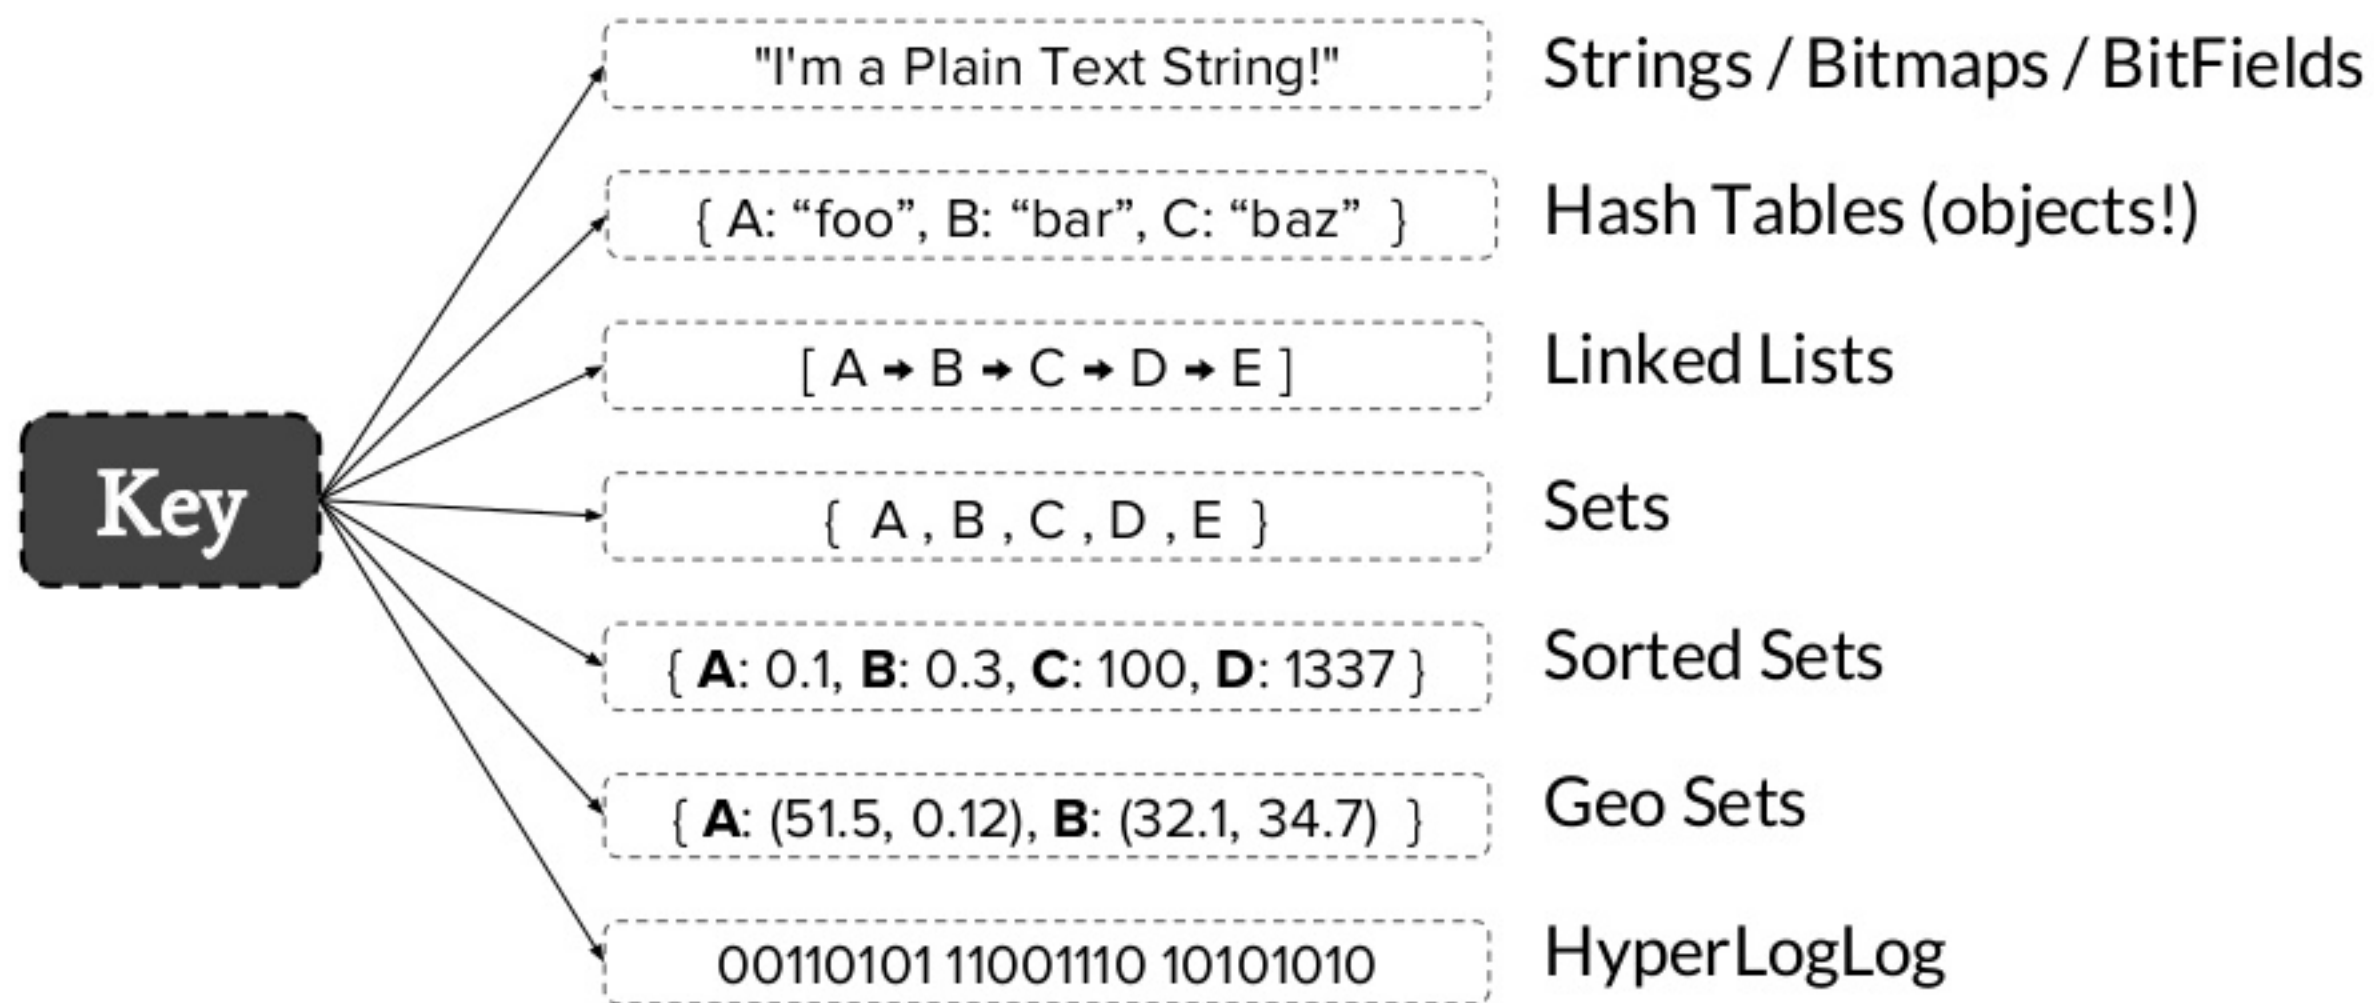
**Simplicity**  
(through Data Structures)



**Extensibility**  
(through Redis Modules)



# A Quick Recap of Redis



# Simple Redis Example (string, hash)

```
127.0.0.1:6379> SET spark summit
OK
127.0.0.1:6379> GET spark
"summit"
127.0.0.1:6379> HMSET spark_hash org apache version 2.1.1
OK
127.0.0.1:6379> HGET spark_hash version
"2.1.1"
127.0.0.1:6379> HGETALL spark_hash
1) "org"
2) "apache"
3) "version"
4) "2.1.1"
```



# Another Simple Redis Example (sorted set)

```
127.0.0.1:6379> zadd my_sorted_set 1 foo
```

```
(integer) 1
```

```
127.0.0.1:6379> zadd my_sorted_set 5 bar
```

```
(integer) 1
```

```
127.0.0.1:6379> zadd my_sorted_set 3 baz
```

```
(integer) 1
```

```
127.0.0.1:6379> ZRANGE my_sorted_set 0 2
```

```
1) "foo"
```

```
2) "baz"
```

```
3) "bar"
```

```
127.0.0.1:6379>
```

# What Modules Actually Are

- Dynamic libraries loaded to redis
- Written in C/C++
- Use a C ABI/API isolating redis internals
- Use existing or add new data-structures
- Near Zero latency access to data



New Data Types

New Commands

New Capabilities

# Modules : A Revolutionary Approach

Adapt your database to your data, not the other way around

## Neural Redis

Simple Neural Network Native  
to Redis

## Redis-ML

Machine Learning Model  
Serving

## RediSearch

Full Text Search Engine in Redis

## ReJSON

JSON Engine on Redis.  
Pre-released

## Time Series

Time series values aggregation  
in Redis

## Graph

Graph database on Redis based  
on Cypher language

## Rate Limiter

Based on Generic Cell Rate  
Algorithm (GCRA)

## Crypto Engine Wrapper

Secure way to store data in  
Redis via encrypt/decrypt with  
various [Themis](#) primitives

## Secondary Index/RQL

Indexing + SQL-like syntax for  
querying indexes.  
Pre-released

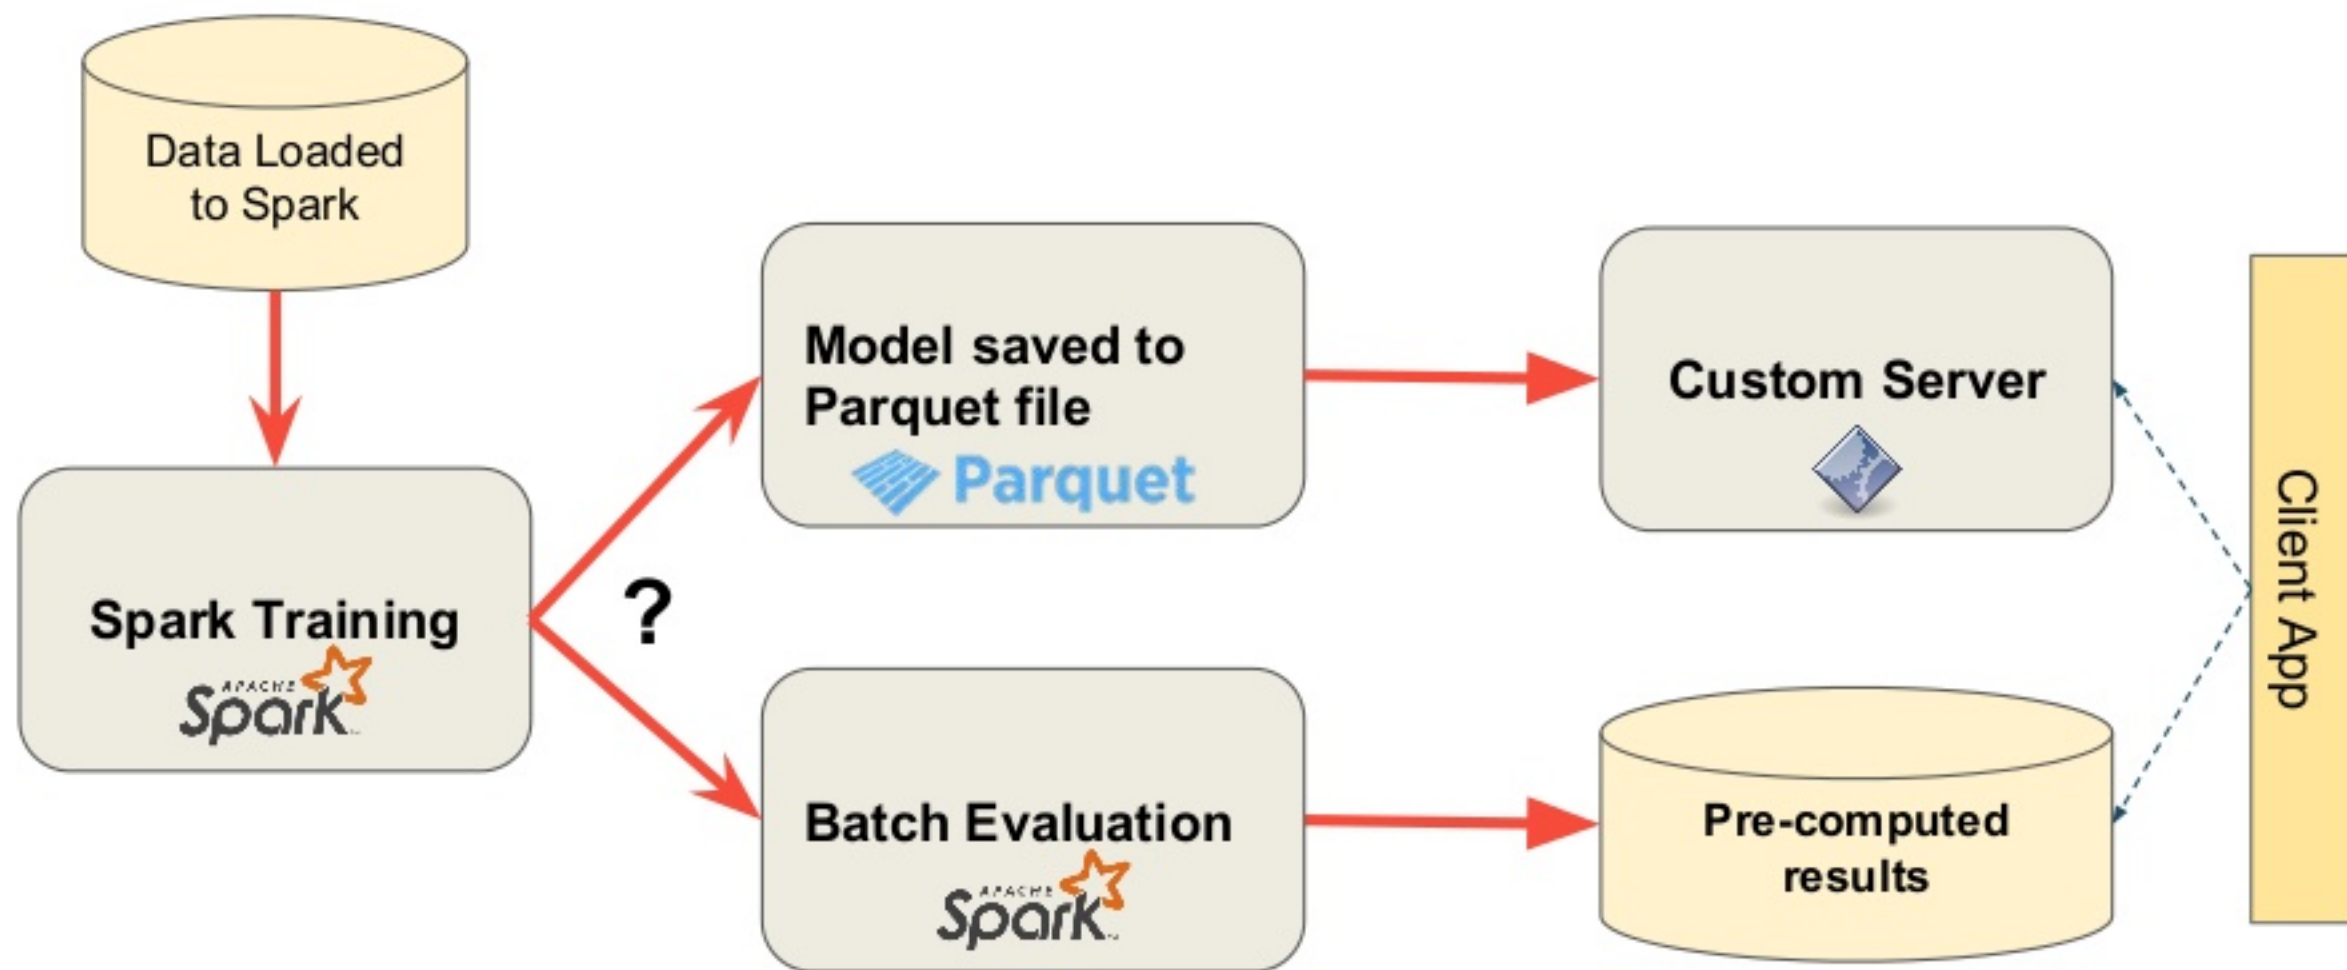




# Redis ML

**Machine Learning Model Server**

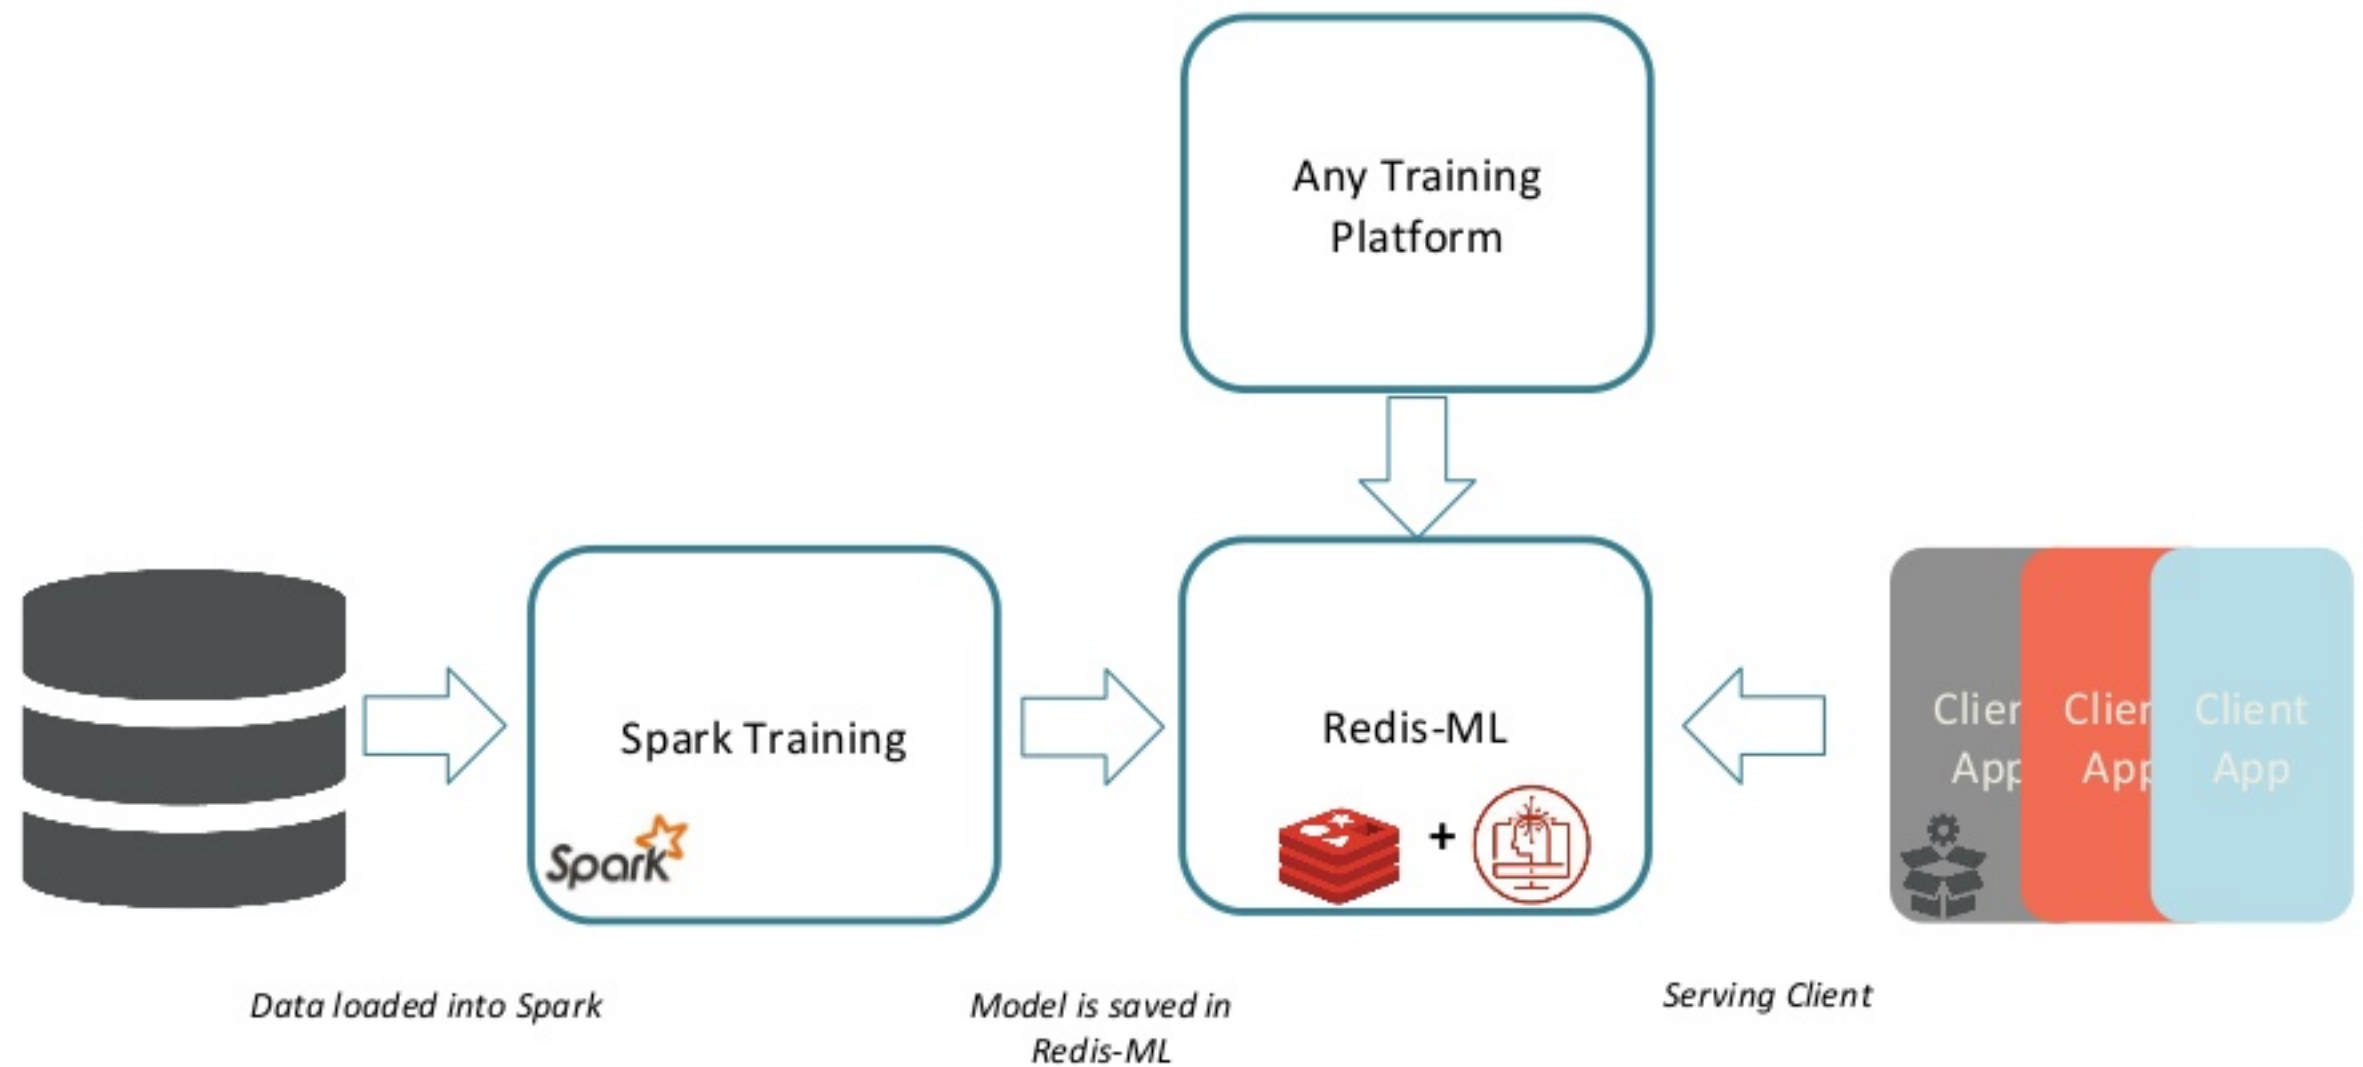
# Spark-ML End-to-End Flow



# ML Models Serving Challenges

- Models are becoming bigger and more complex
- Can be challenging to deploy & serve
- Do not scale well, speed and size
- Can be very expensive

# A Simpler Machine Learning Lifecycle







# Redis-ML – ML Serving Engine

- Store training output as “hot model”
- Perform evaluation directly in Redis
- Enjoy the performance, scalability and HA of Redis



# Redis-ML



## ML Models

Tree Ensembles

Linear Regression


Logistic Regression

Matrix + Vector Operations

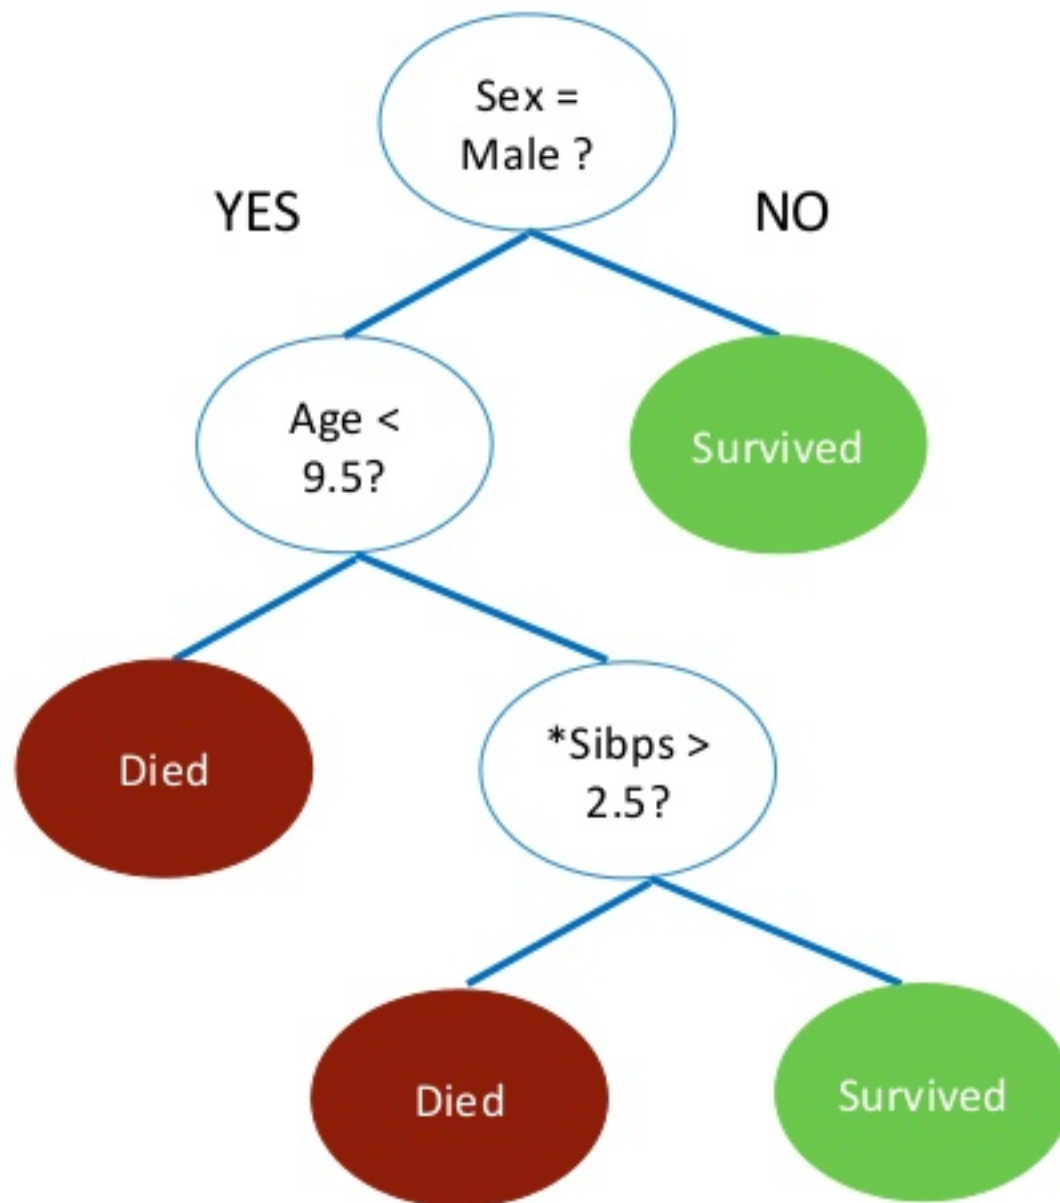
More to come...

# Random Forest Model

- A collection of decision trees
- Supports classification & regression
- Splitter Node can be:
  - Categorical (e.g. day == "Sunday")
  - Numerical (e.g. age < 43)
- Decision is taken by the majority of decision trees

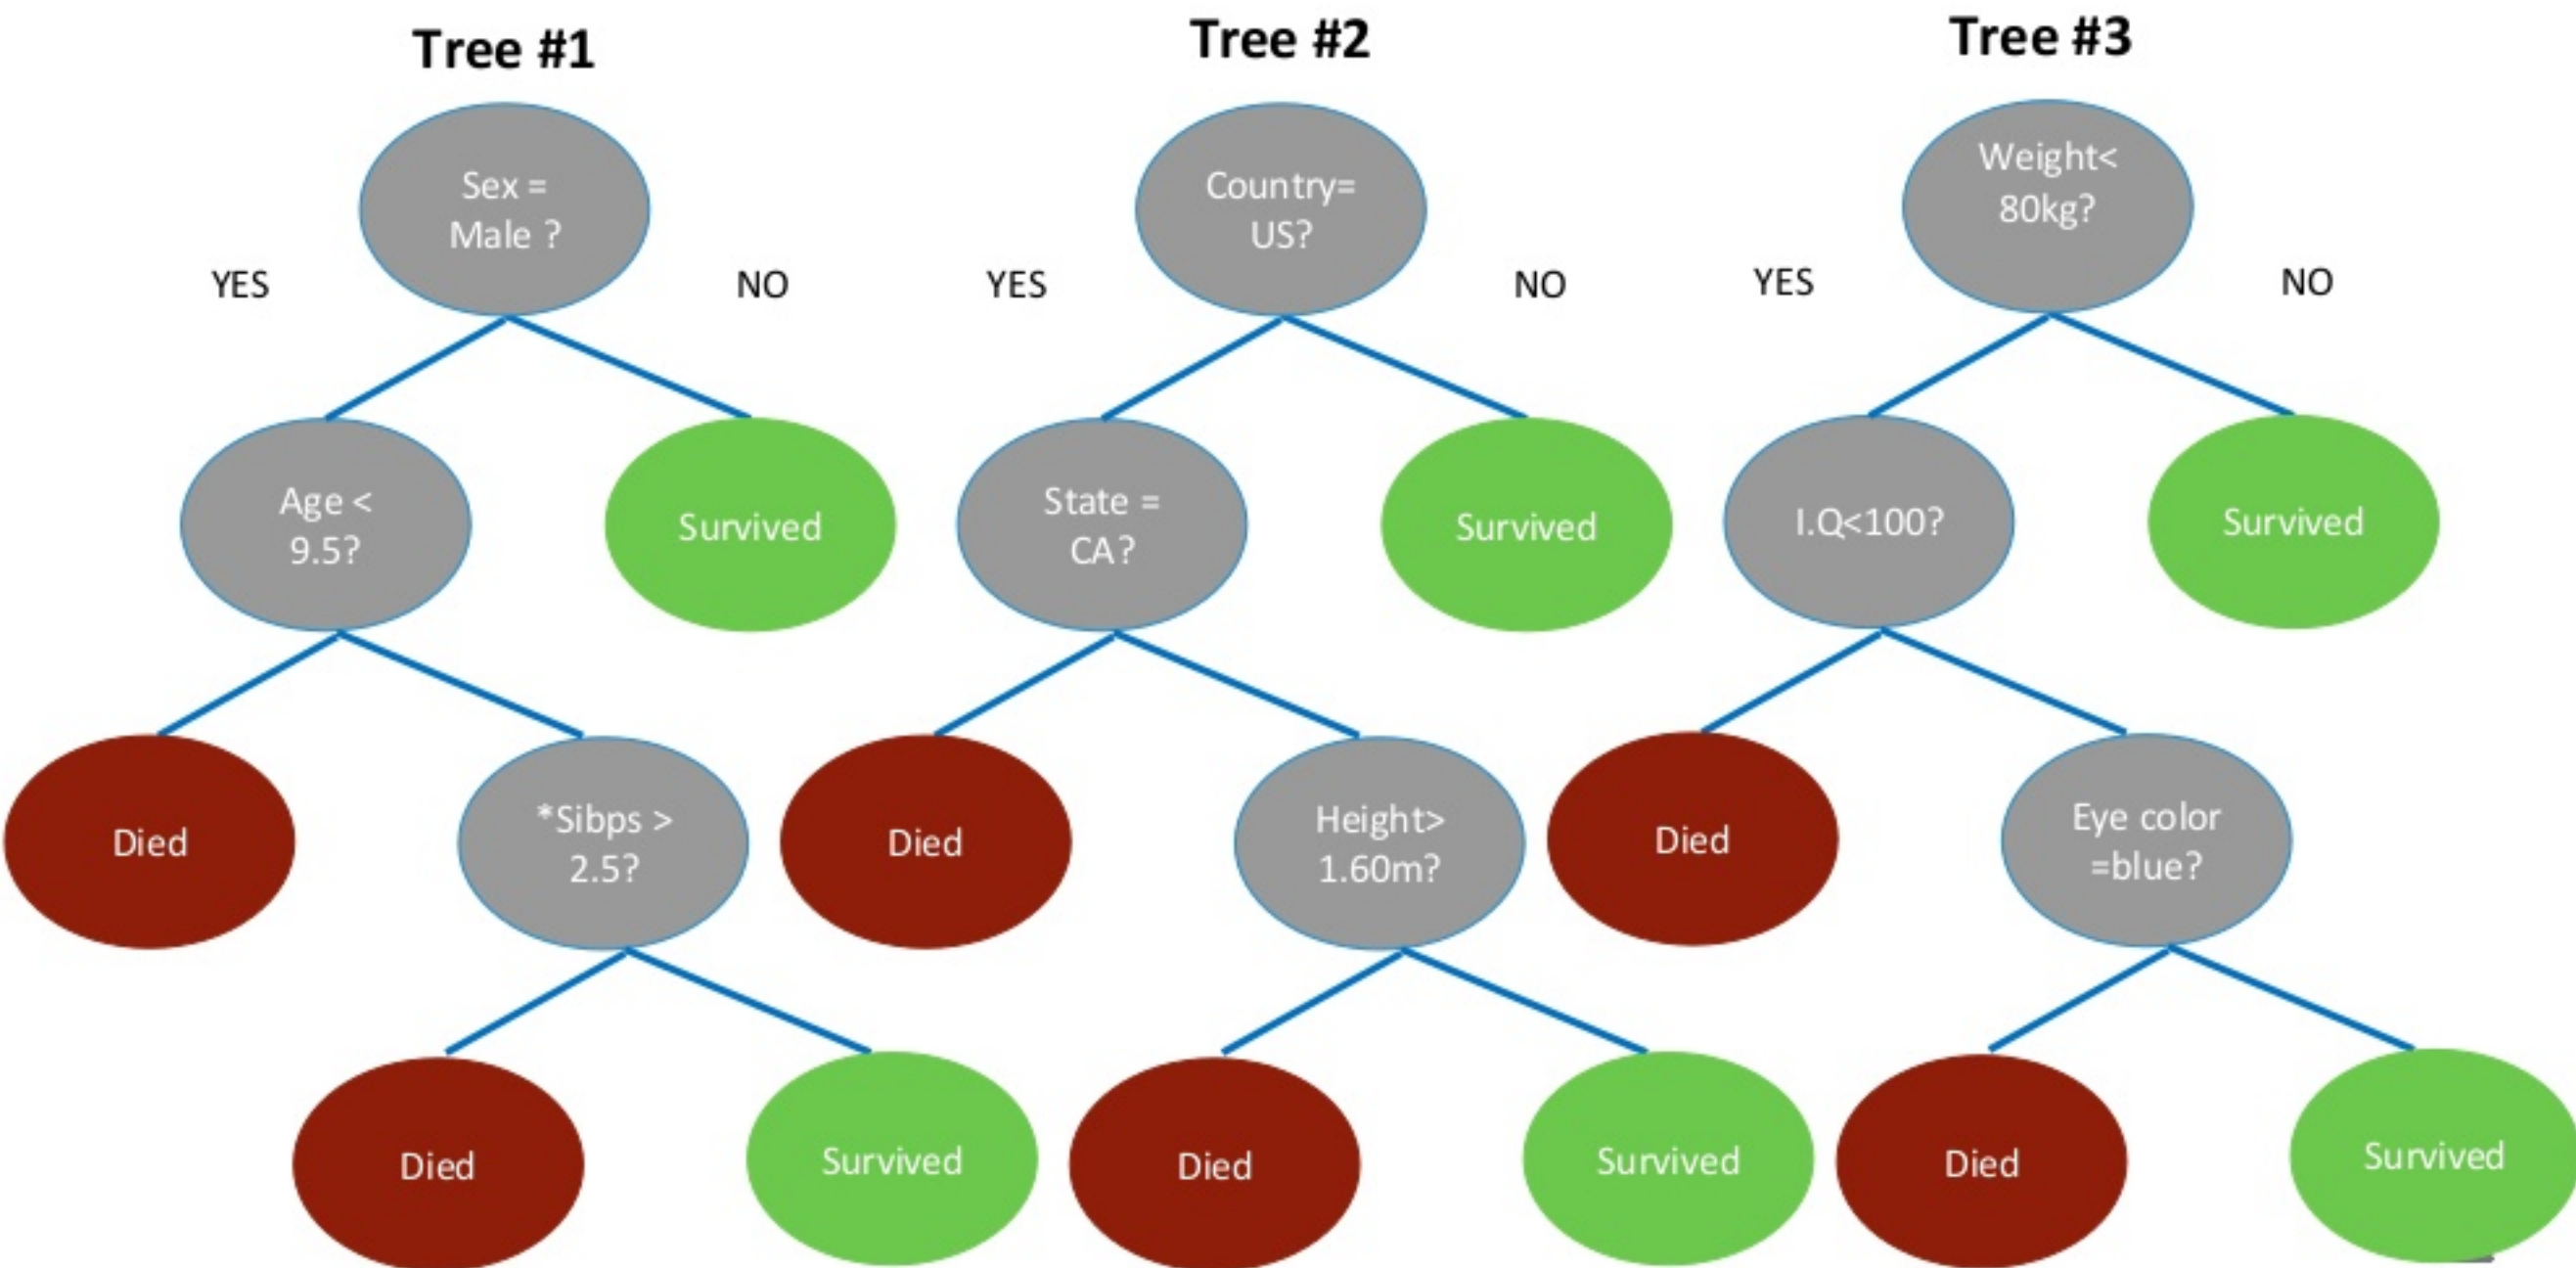


# Titanic Survival Predictor on a Decision Tree



\*Sibps = siblings + spouses

# Titanic Survival Predictor on a Random Forest





# Would John Survive The Titanic

- John's features:  
{male, 34, married + 2, US, CA, 1.78m, 78kg, 110iq, blue eyes}
- Tree#1 – Survived
- Tree#2 – Failed
- Tree#3 – Survived
- Random forest decision - Survived

# Forest Data Type Example

```
> MODULE LOAD "./redis-ml.so"
OK
> ML.FOREST.ADD myforest 0 . CATEGORIC sex "male" .L LEAF 1 .R LEAF 0
OK
> ML.FOREST.RUN myforest sex:male
"1"
> ML.FOREST.RUN myforest sex:no_thankx
"0"
```



# Using Redis-ML With Spark

```
scala> import com.redislabs.client.redism1.MLClient
scala> import com.redislabs.provider.redis.ml.Forest

scala> val jedis = new Jedis("localhost")
scala> val rfModel = pipelineModel.stages.last.asInstanceOf[RandomForest]

// Create a new forest instance
scala> val f = new Forest(rfModel.trees)

// Load the model to redis
scala> f.loadToRedis("forest-test", "localhost")

// Classify a feature vector
scala> jedis.getClient().sendCommand(MLClient.ModuleCommand.FOREST_RUN,
"forest-test", makeInputString(0))

scala> jedis.getClient().getStatusCodeReply
res53: String = 1
```


# Real World Challenge

- Ad serving company
- Need to serve 20,000 ads/sec @ 50msec data-center latency
- Runs 1k campaigns → 1K random forest
- Each forest has 15K trees
- On average each tree has 7 levels (depth)
- Would require < 1000 x c4.8xlarge

# Redis ML with Spark ML


## 40x Faster

Classification Time Over Spark



# **Real World Example: Movie Recommendation System**

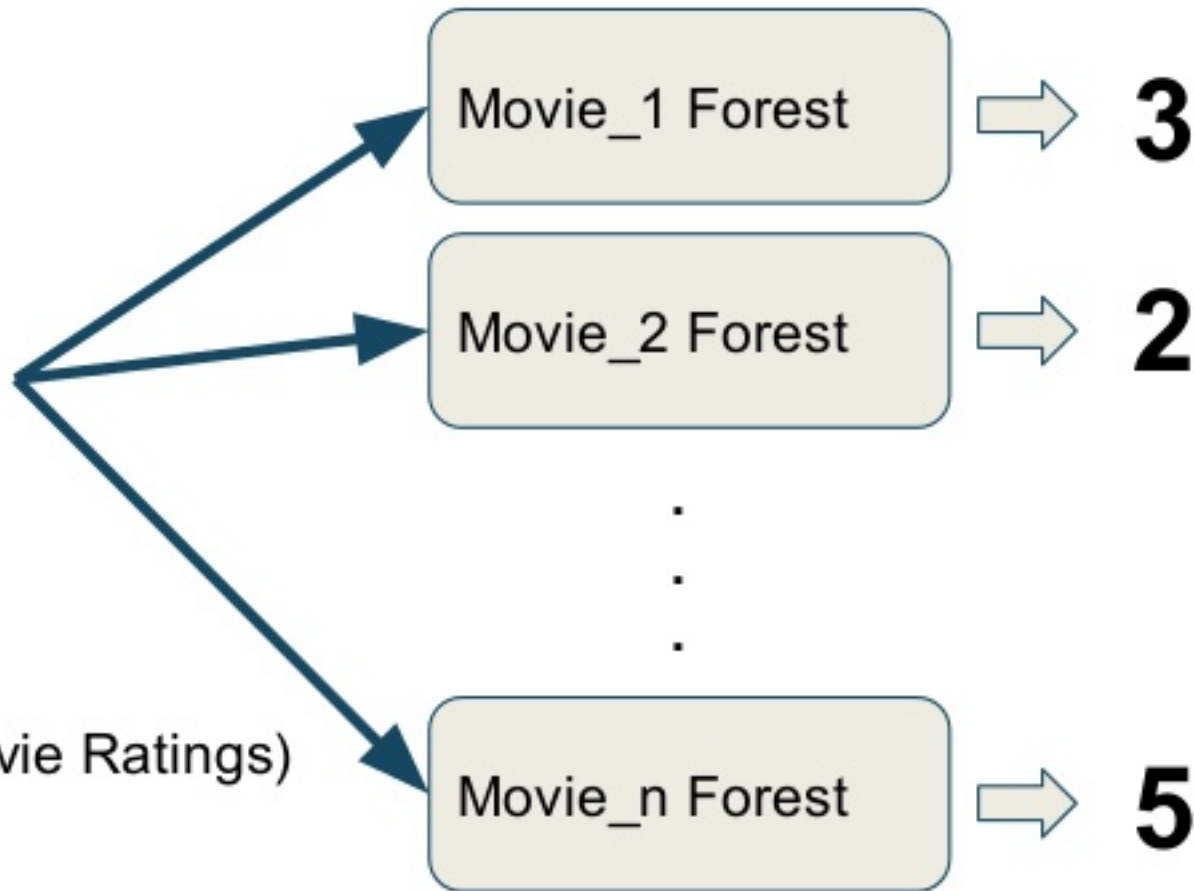
# Overview



# Concept: One Forest For Each Movie



User Features:  
(Age, Gender, Movie Ratings)



# The Tools

Transform: 

Train: 

Classify:  + 

Containers:   
docker



## Using the Dockers

```
$ docker pull shaynativ/redis-m1
$ docker run --net=host shaynativ/redis-m1 &
$
$ docker pull shaynativ/spark-redis-m1
$ docker run --net=host shaynativ/spark-redis-m1
```

# Step 1: Get The Data

- Download and extract the [MovieLens 100K Dataset](#)
- The data is organized in separate files:
  - Ratings: user id | item id | rating (1-5) | timestamp
  - Item (movie) info: movie id | genre info fields (1/0)
  - User info: user id | age | gender | occupation
- Our classifier should return the expected rating (from 1 to 5) a user would give the movie in question

## Step 2: Transform

- The training data for each movie should contain 1 line per user:
  - class (rating from 1 to 5 the user gave to this movie)
  - user info (age, gender, occupation)
  - user ratings of other movies (movie\_id:rating ...)
  - user genre rating averages (genre:avg\_score ...)
- Run `gen_data.py` to transform the files to the desired format

## Step3: Train and Load to Redis

// Create a new forest instance

val rf = new

RandomForestClassifier().setFeatureSubsetStrategy("auto").setLabelCol("indexedLabel").setFeaturesCol("indexedFeatures").setNumTrees(500)

.....

// Train model

val model = pipeline.fit(trainingData)

.....

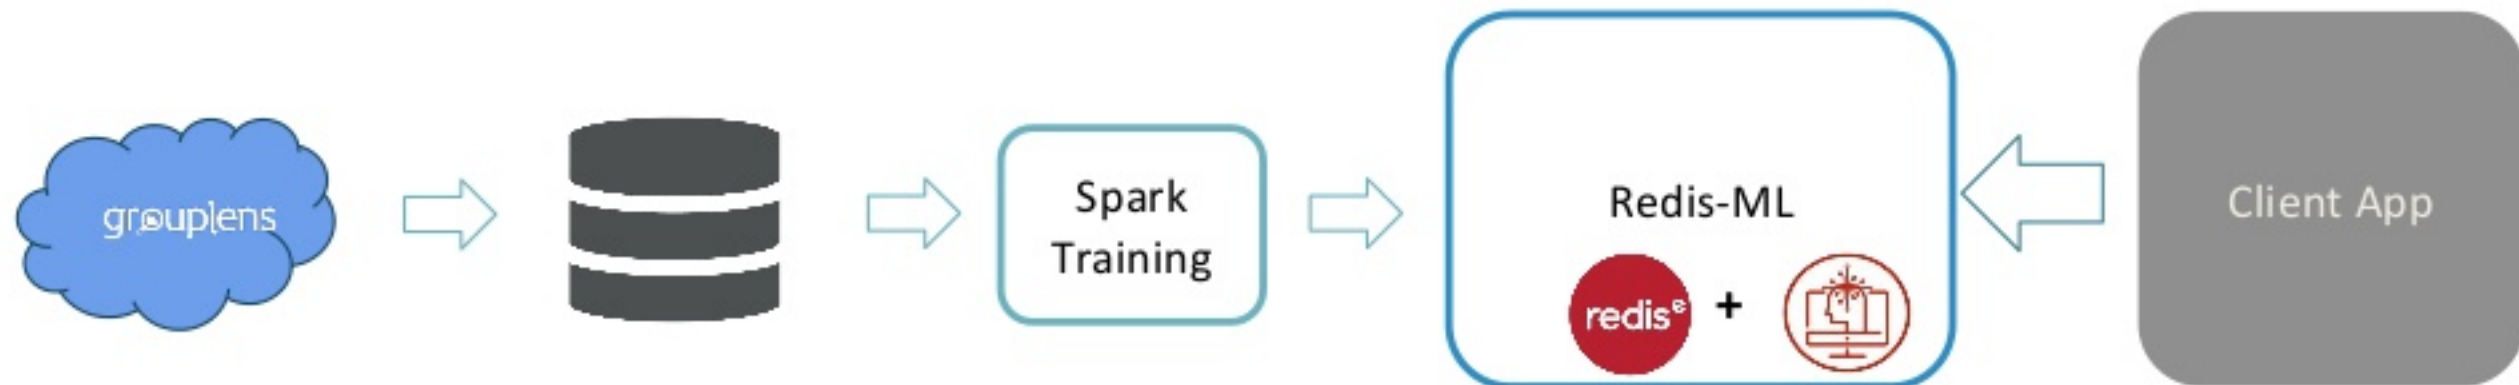
val rfModel = model.stages(2).asInstanceOf[RandomForestClassificationModel]

// Load the model to redis

val f = new Forest(rfModel.trees)

f.loadToRedis("movie-10", "127.0.0.1")

## Step 3: Execute in Redis



# Python Client Example

```
>> import redis
>> config = {"host": "localhost", "port": 6379}
>> r = redis.StrictRedis(**config)
>> user_profile = r.get("user_shay_profile")
>> print(user_profile)
12:1.0,13:1.0,14:3.0,15:1.0,17:1.0,18:1.0,19:1.0,20:1.0,23:1.0,24:5.0,1.0,115:1.0,116:2.0,
0,117:2.0,119:1.0,120:4.0,121:2.0,122:2.0,
.....
1360:1.0,1361:1.0,1362:1.0,
1701:6.0,1799:435.0,1801:0.2,1802:0.11,1803:0.04,1812:0.04,1813:0.07,1814:0.24,1815:0.09
,1816:0.32,1817:0.06
>> r.execute_command("ML.FOREST.RUN", "movie-10", user_profile)
'3'
```

# Redis CLI Example

```
>keys *
```

```
127.0.0.1:6379> KEYS *
```

```
1) "movie-5"
```

```
2) "movie-1"
```

```
.....
```

```
8) "movie-6"
```

```
9) "movie-4"
```

```
10) "movie-10"
```

```
11) "user_1_profile"
```

```
>ML.FOREST.RUN movie-10
```

```
12:1.0,13:1.0,,332:3.0,333:1.0,334:1.0,335:2.0,336:1.0,357:2.0,358:1.0,359:1.0,362:1.0,367:1.
```

```
.....
```

```
,410:3.0,411:2.0,412:2.0,423:1.0,454:1.0,455:1.0,456:1.0,457:3.0,458:1.0,459:1.0,470:1"
```

```
"3"
```

```
>
```



# Performance

```
Redis time: 0.635129ms, res=3
```

```
Spark time: 46.657662ms, res=3.0
```

```
-----
```

```
Redis time: 0.644444ms, res=3
```

```
Spark time: 49.028983ms, res=3.0
```

```
-----
```

```
Classification averages:
```

```
redis: 0.9401250000000001 ms
```

```
spark: 58.01970206666667 ms
```

```
ratio: 61.71488053893542
```

```
diffs: 0.0
```

## Getting Actual Recommendations - Python Script

```
#!/usr/bin/python

import operator
import redis
config = {"host": "localhost", "port": 6379}
r = redis.StrictRedis(**config)

user_profile = r.get("user-1-profile")
results = {}

for i in range(1, 11):
    results[i] = r.execute_command("ML.FOREST.RUN", "movie-{}".format(i), user_profile)

print "Movies sorted by scores:"
sorted_results = sorted(results.items(), key=operator.itemgetter(1), reverse=True)
for k,v in sorted_results:
    print "movie-{}:{}".format(k,v)

print ""
print "Recommended movie: movie-{}".format(sorted_results[0][0])
```

## Getting Actual Recommendations - Results

```
$ ./classify_user.py 1
```

```
Movies sorted by scores:
```

```
movie-4:3
```

```
movie-3:2
```

```
movie-6:2
```

```
movie-7:2
```

```
movie-8:2
```

```
movie-9:2
```

```
movie-1:1
```

```
movie-2:1
```

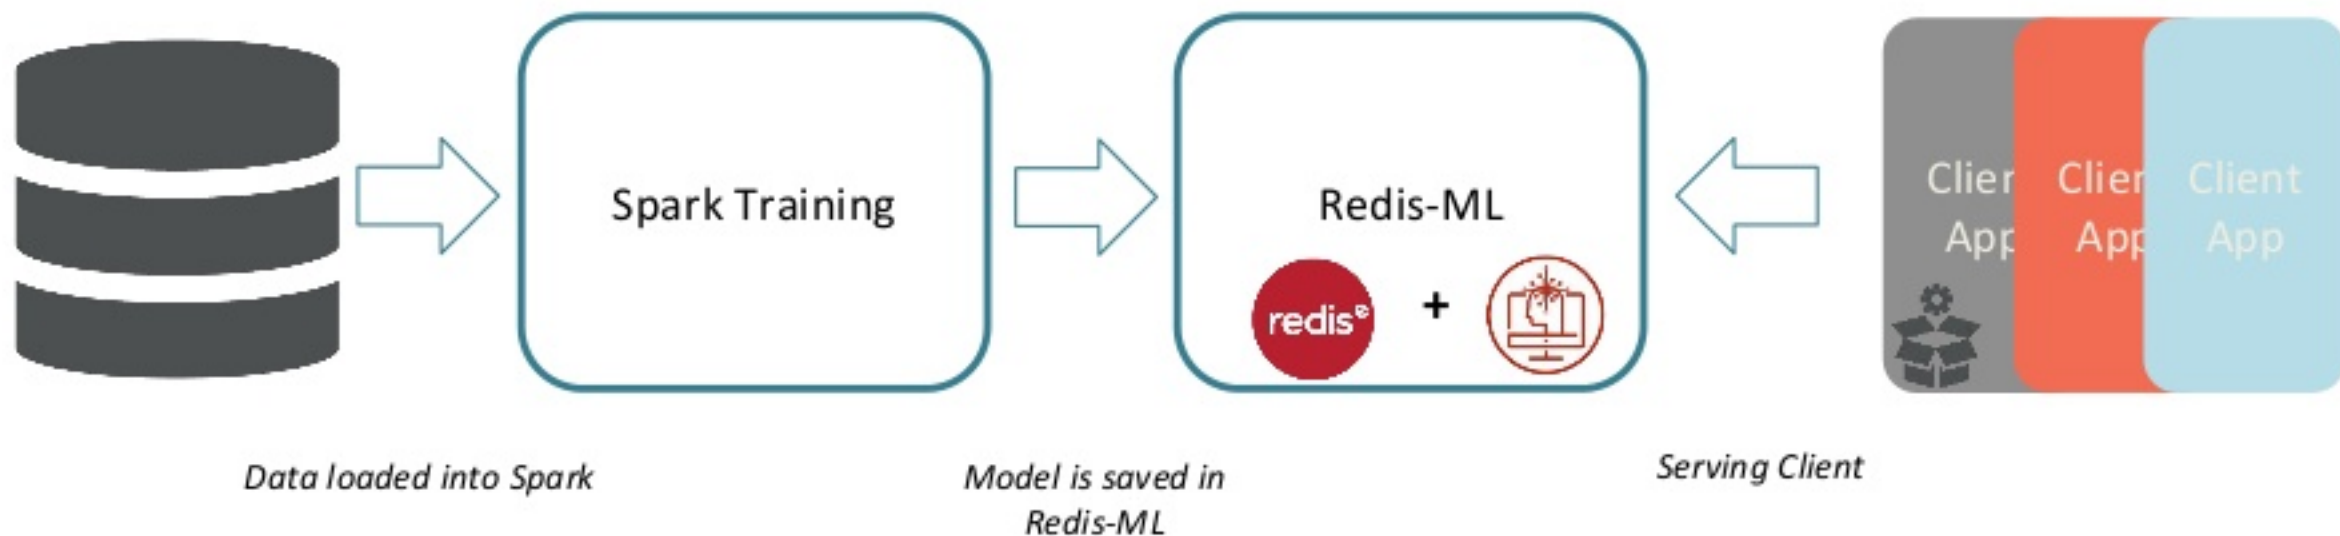
```
movie-5:1
```

```
movie-10:0
```

```
Recommended movie for user 1: movie-4
```

# Summary

- Train with Spark, Serve with Redis
- 97% resource cost serving
- Simplify ML lifecycle
- Redis<sup>e</sup> (Cloud or Pack):
  - Scaling, HA, Performance
  - PAYG – cost optimized
  - Ease of use
  - Supported by the teams who created Spark and Redis



# Resources

- Redis-ML: <https://github.com/RedisLabsModules/redis-ml>
- Spark-Redis-ML: <https://github.com/RedisLabs/spark-redis-ml>
- Databricks Notebook: <http://bit.ly/sparkredism1>
- Dockers: <https://hub.docker.com/r/shaynativ/redis-ml/>  
<https://hub.docker.com/r/shaynativ/spark-redis-ml/>

# Q&A



# Thank You.

[shay@redislabs.com](mailto:shay@redislabs.com)