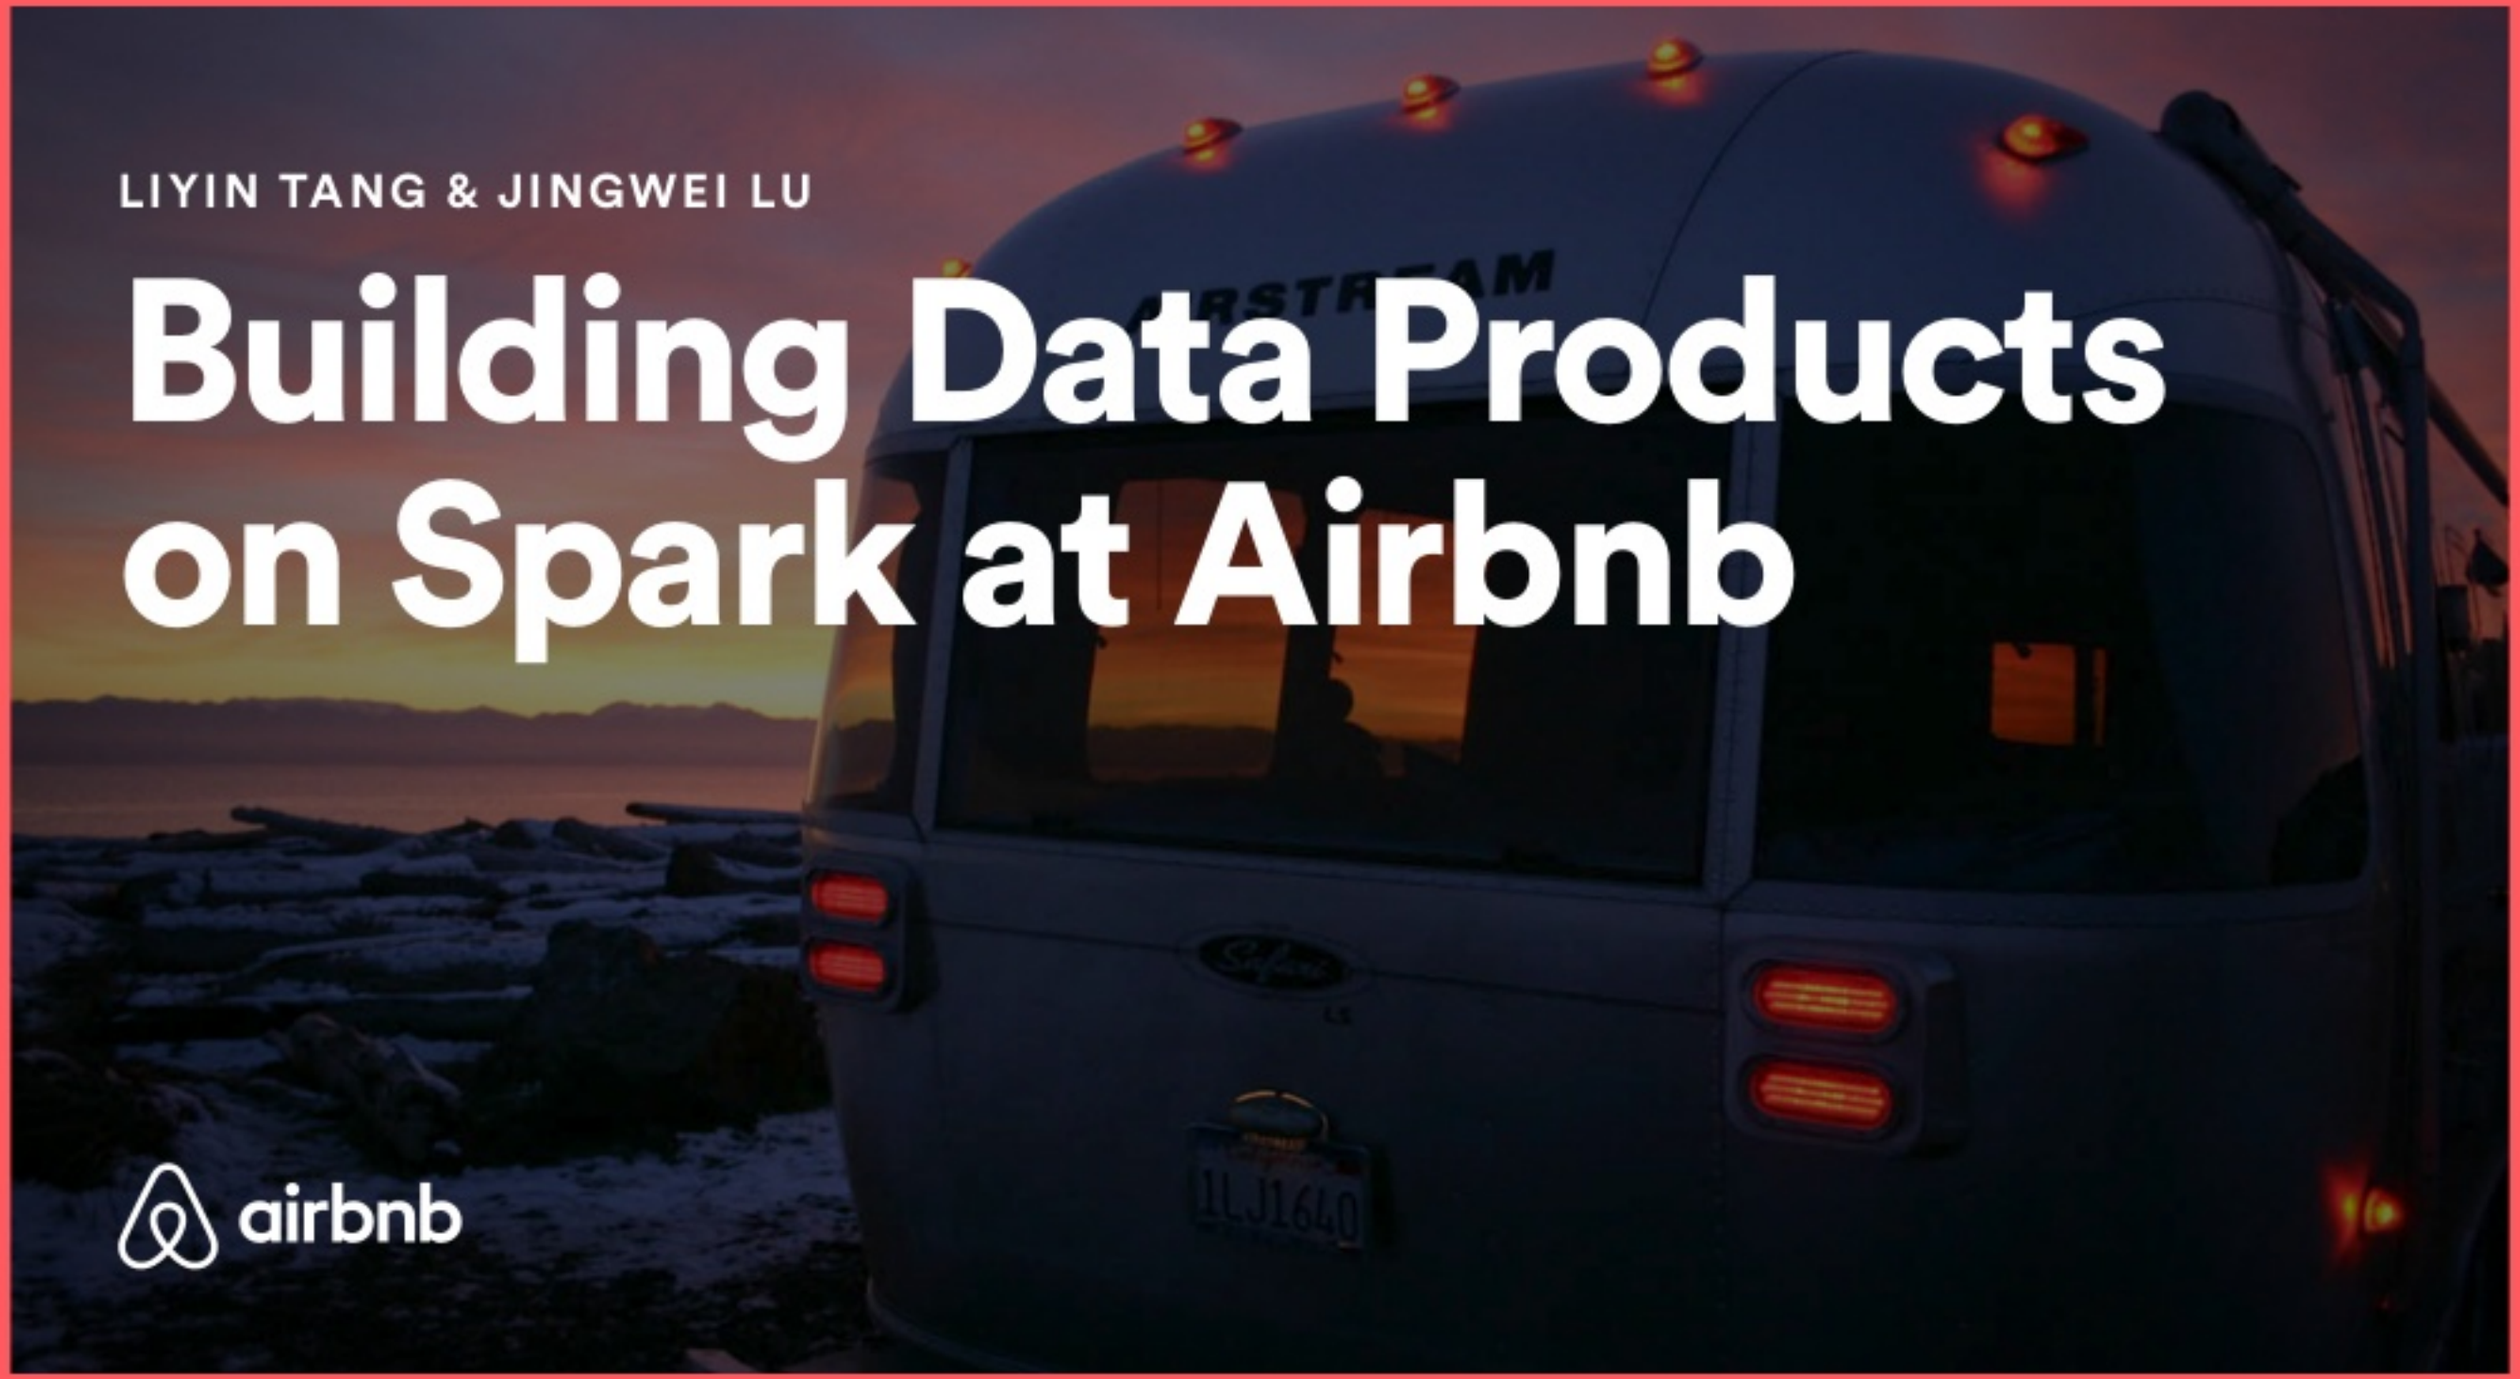
LIYIN TANG & JINGWEI LU
Building Data Products on Spark at Airbnb
airbnb

Data Infrastructure at Airbnb

Batch Infrastructure

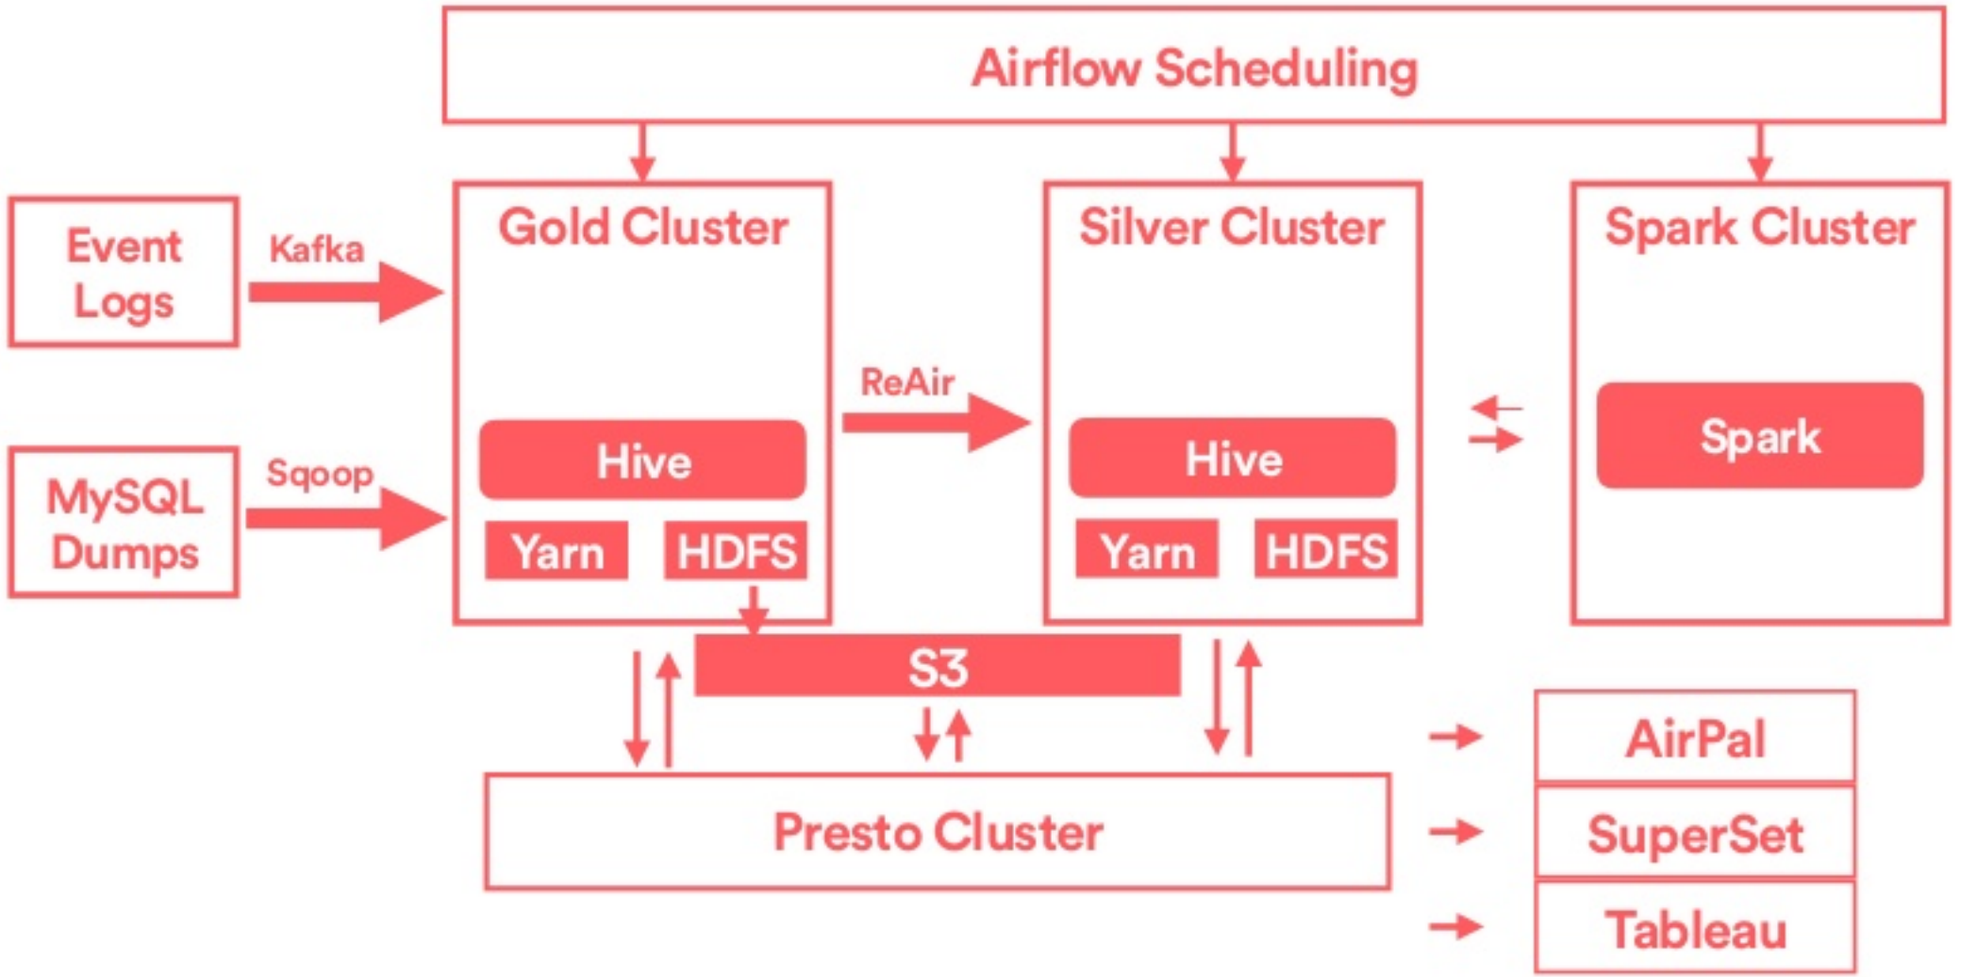

Streaming at Airbnb

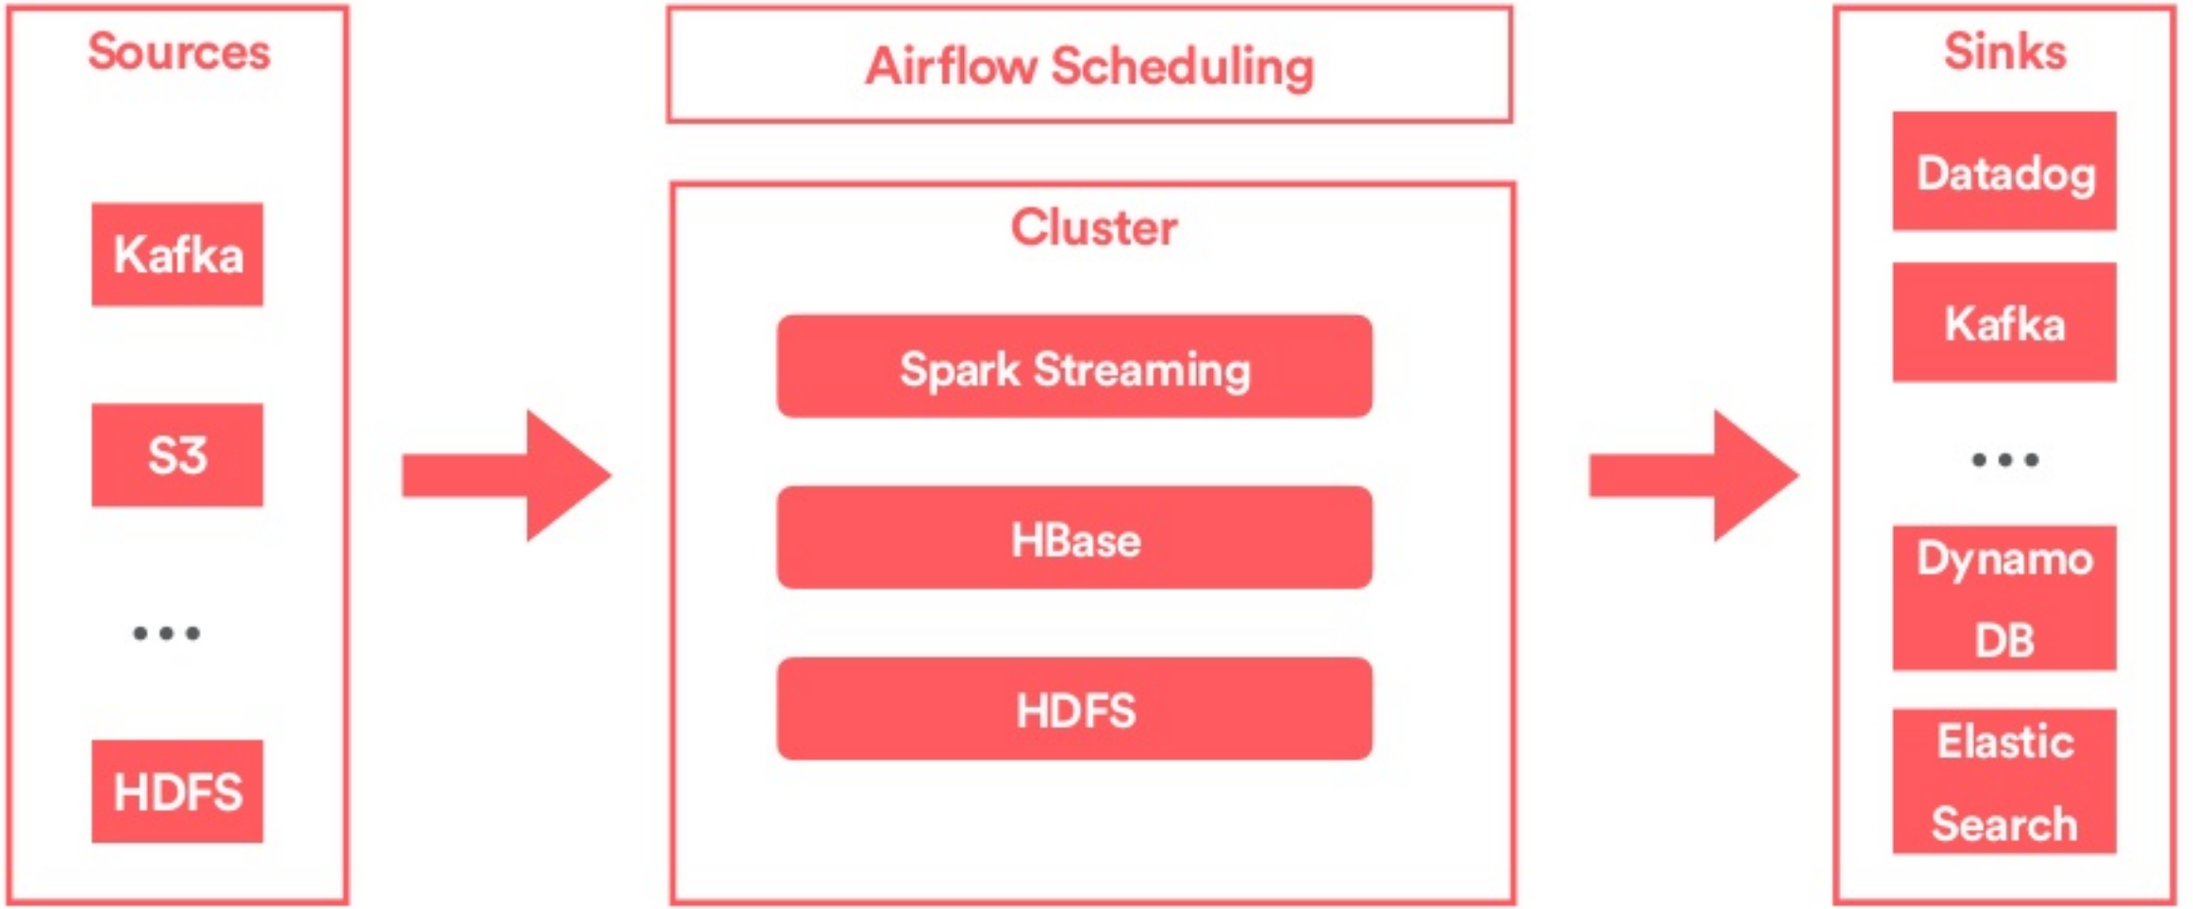

Lambda Architecture

Lambda Architecture

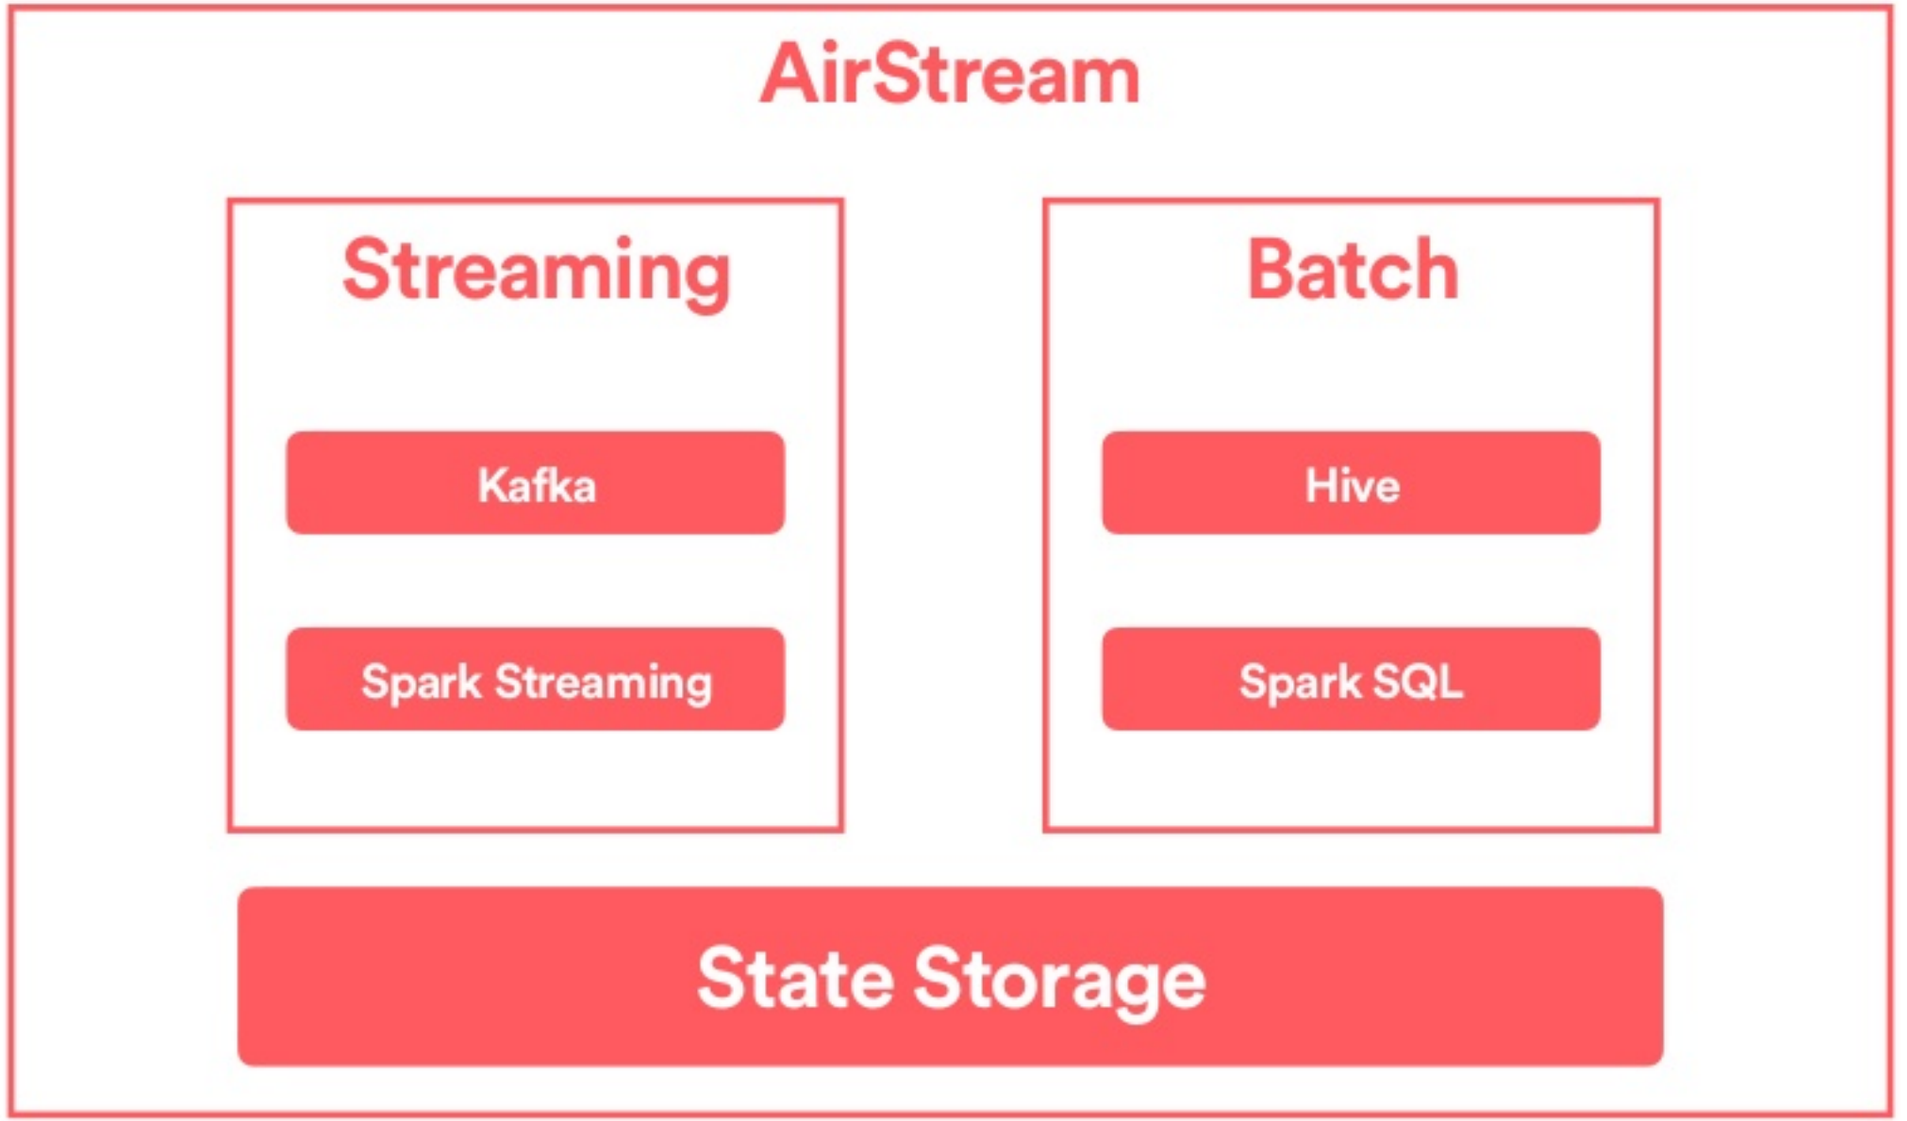

Combine Streaming and Batch Processing

Sources

Streaming

```
source: [
 {
  name: source_example,
  type: kafka,
  config: {
   topic: "example_topic",
  }
 }
]
```

Batch

```
source: [
 {
  name: source_example,
  type: hive,
  sql: {
   select * from db.table where
ds='2017-06-05';
  }
 }
]
```

Computation

Streaming/Batch

```
process: [{
  name = process_example,
  type = sql,
  sql = """
   SELECT  listing_id, checkin_date, context.source as source
   FROM  source_example
   WHERE  user_id IS NOT NULL  """
 }]
```

Sinks

Streaming

```
sink: [
 {
  name = sink_example
  input = process_example
  type = hbase_update
  hbase_table_name = test_table
  bulk_upload = false
 }
]
```

Batch

```
sink: [
 {
  name = sink_example
  input = process_example
  type = hbase_update
  hbase_table_name = test_table
  bulk_upload = true
 }
]
```

Computation Flow

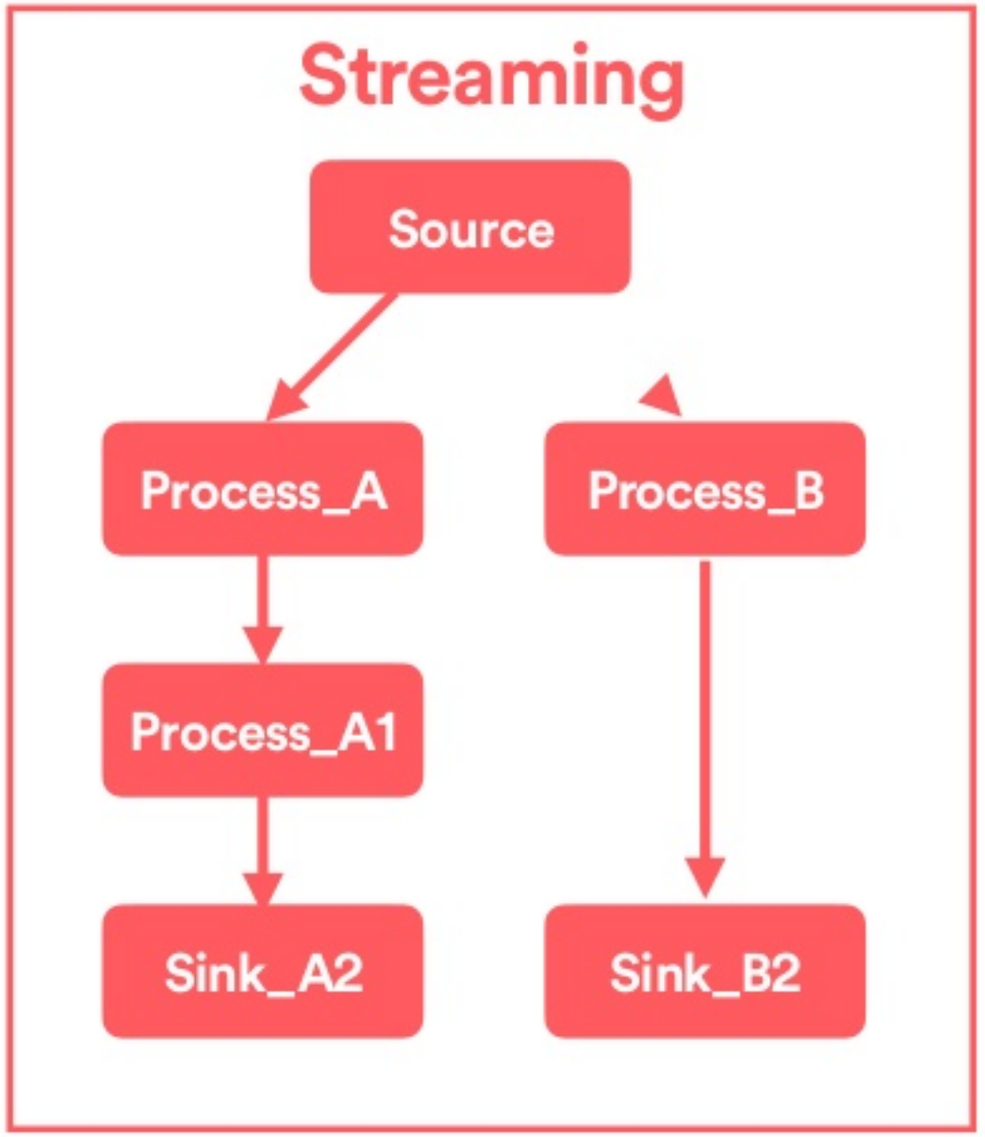

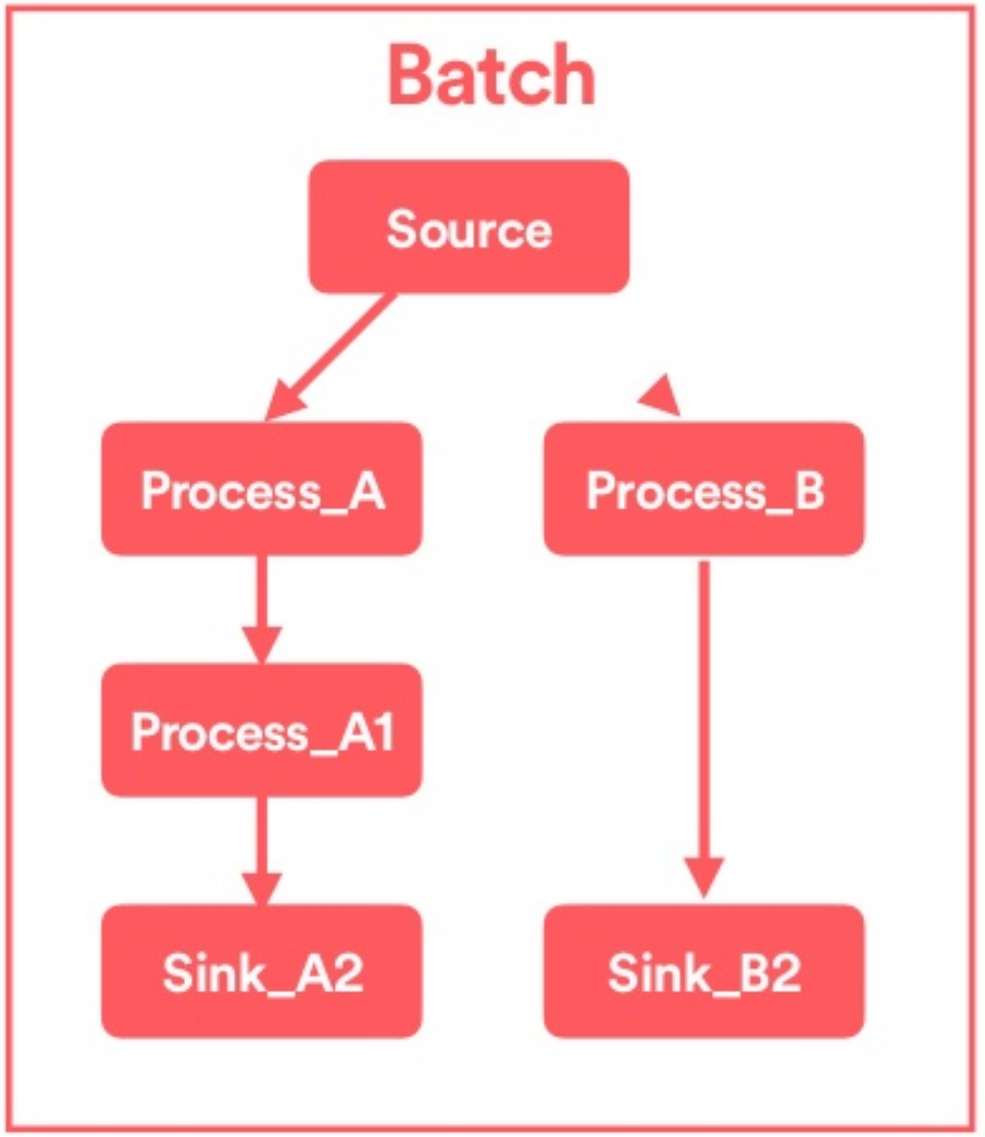

Unified API through AirStream

- Declarative job configuration
- Streaming source vs static source
- Computation operator or sink can be shared by streaming and batch job.
- Computation flow is shared by streaming and batch
- Single driver executes in both streaming and batch mode job

Shared State Storage

Shared Global State Store

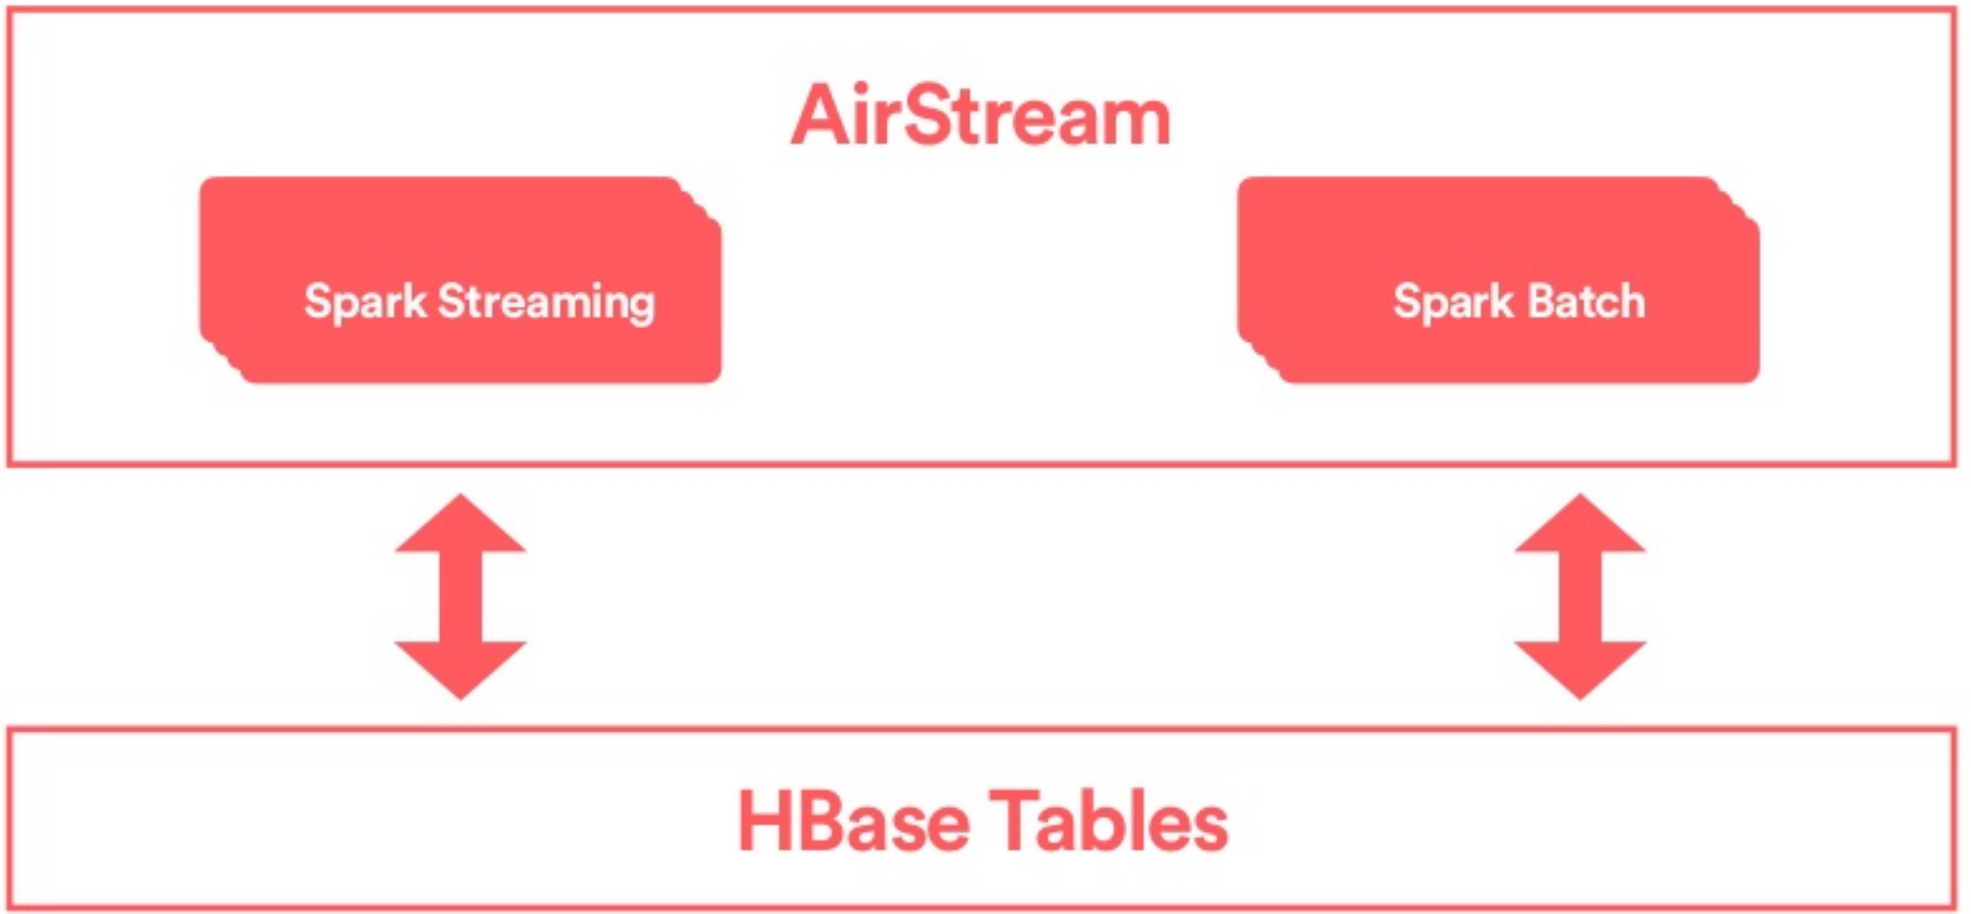

Why HBase

- Well integrated with Hadoop eco system
- Efficient API for streaming writes and bulk uploads
- Rich API for sequential scan and point-lookups
- Merged view based on version

Unified Write API

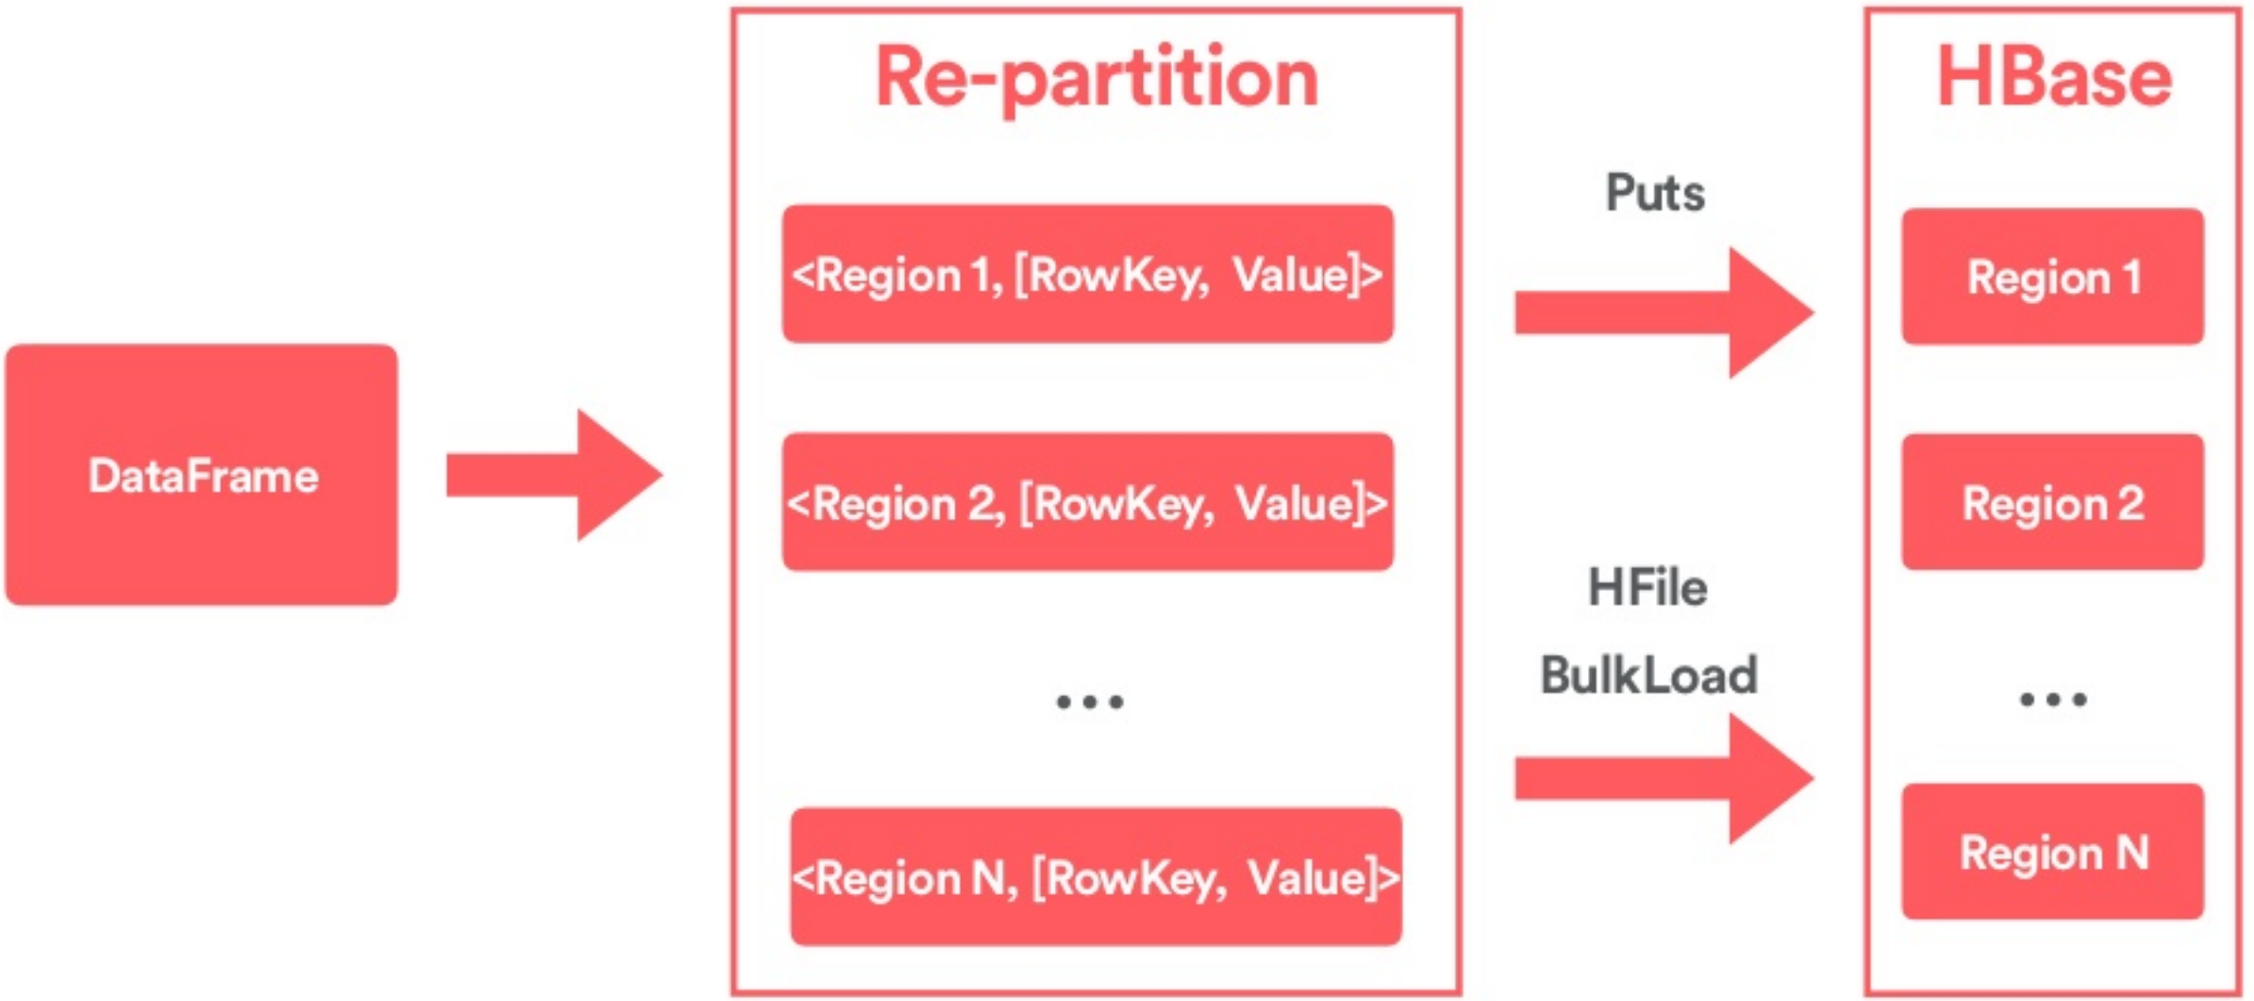

Rich Read API

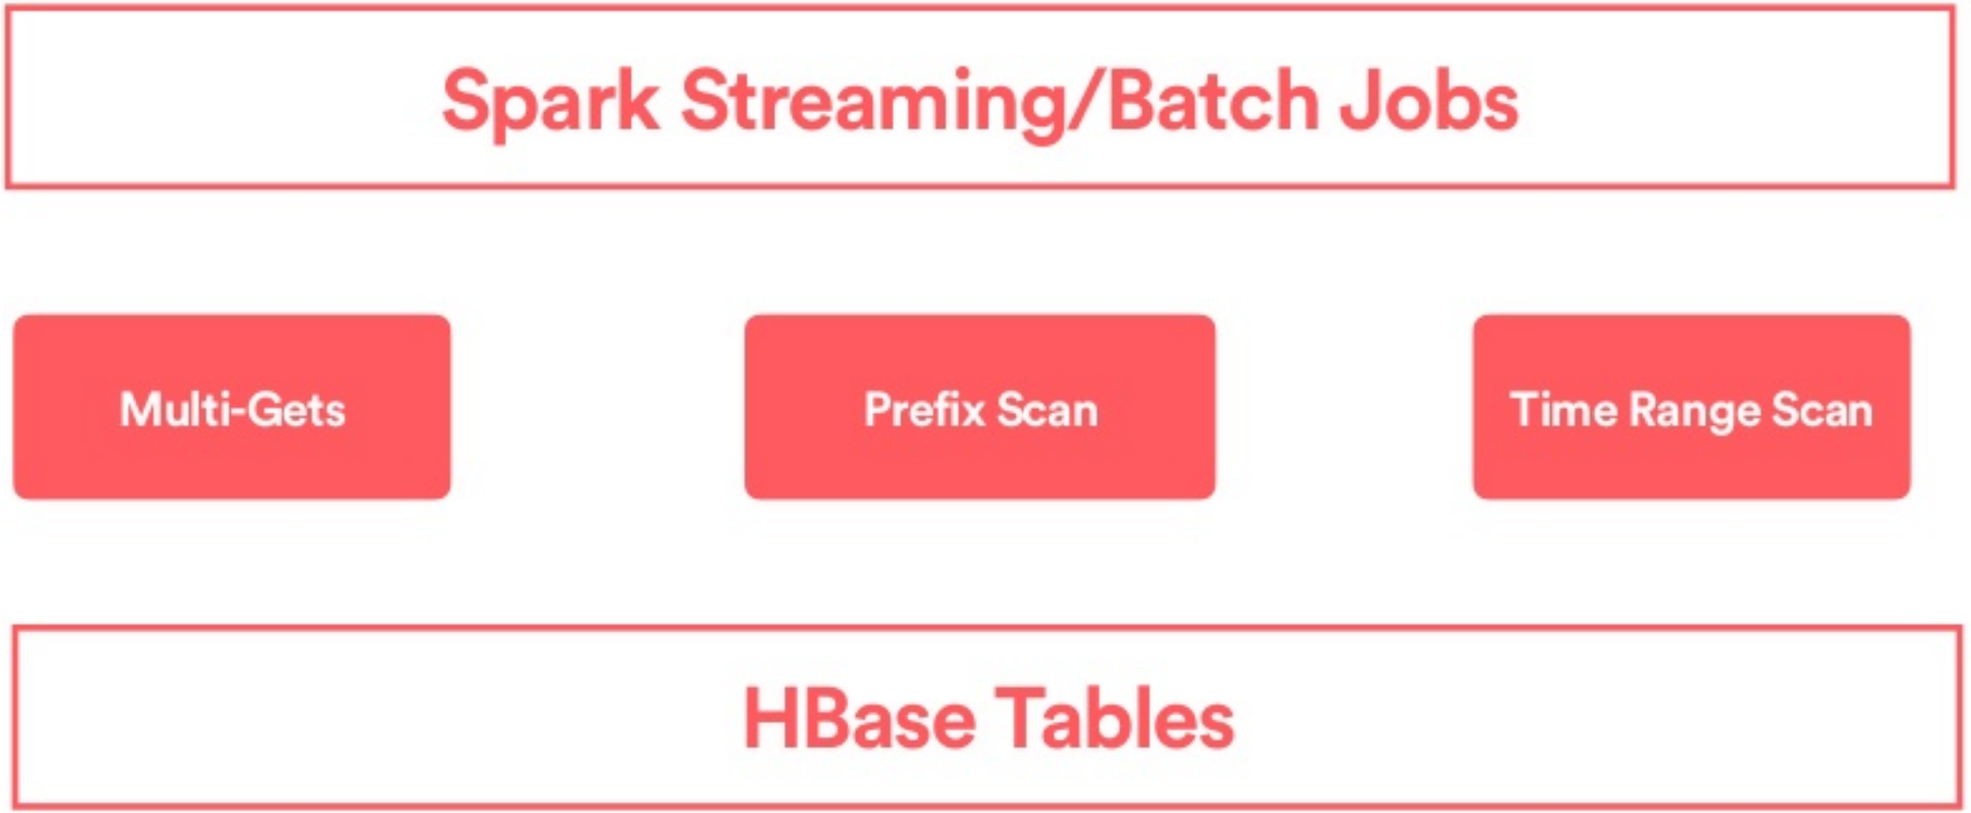

Merged Views

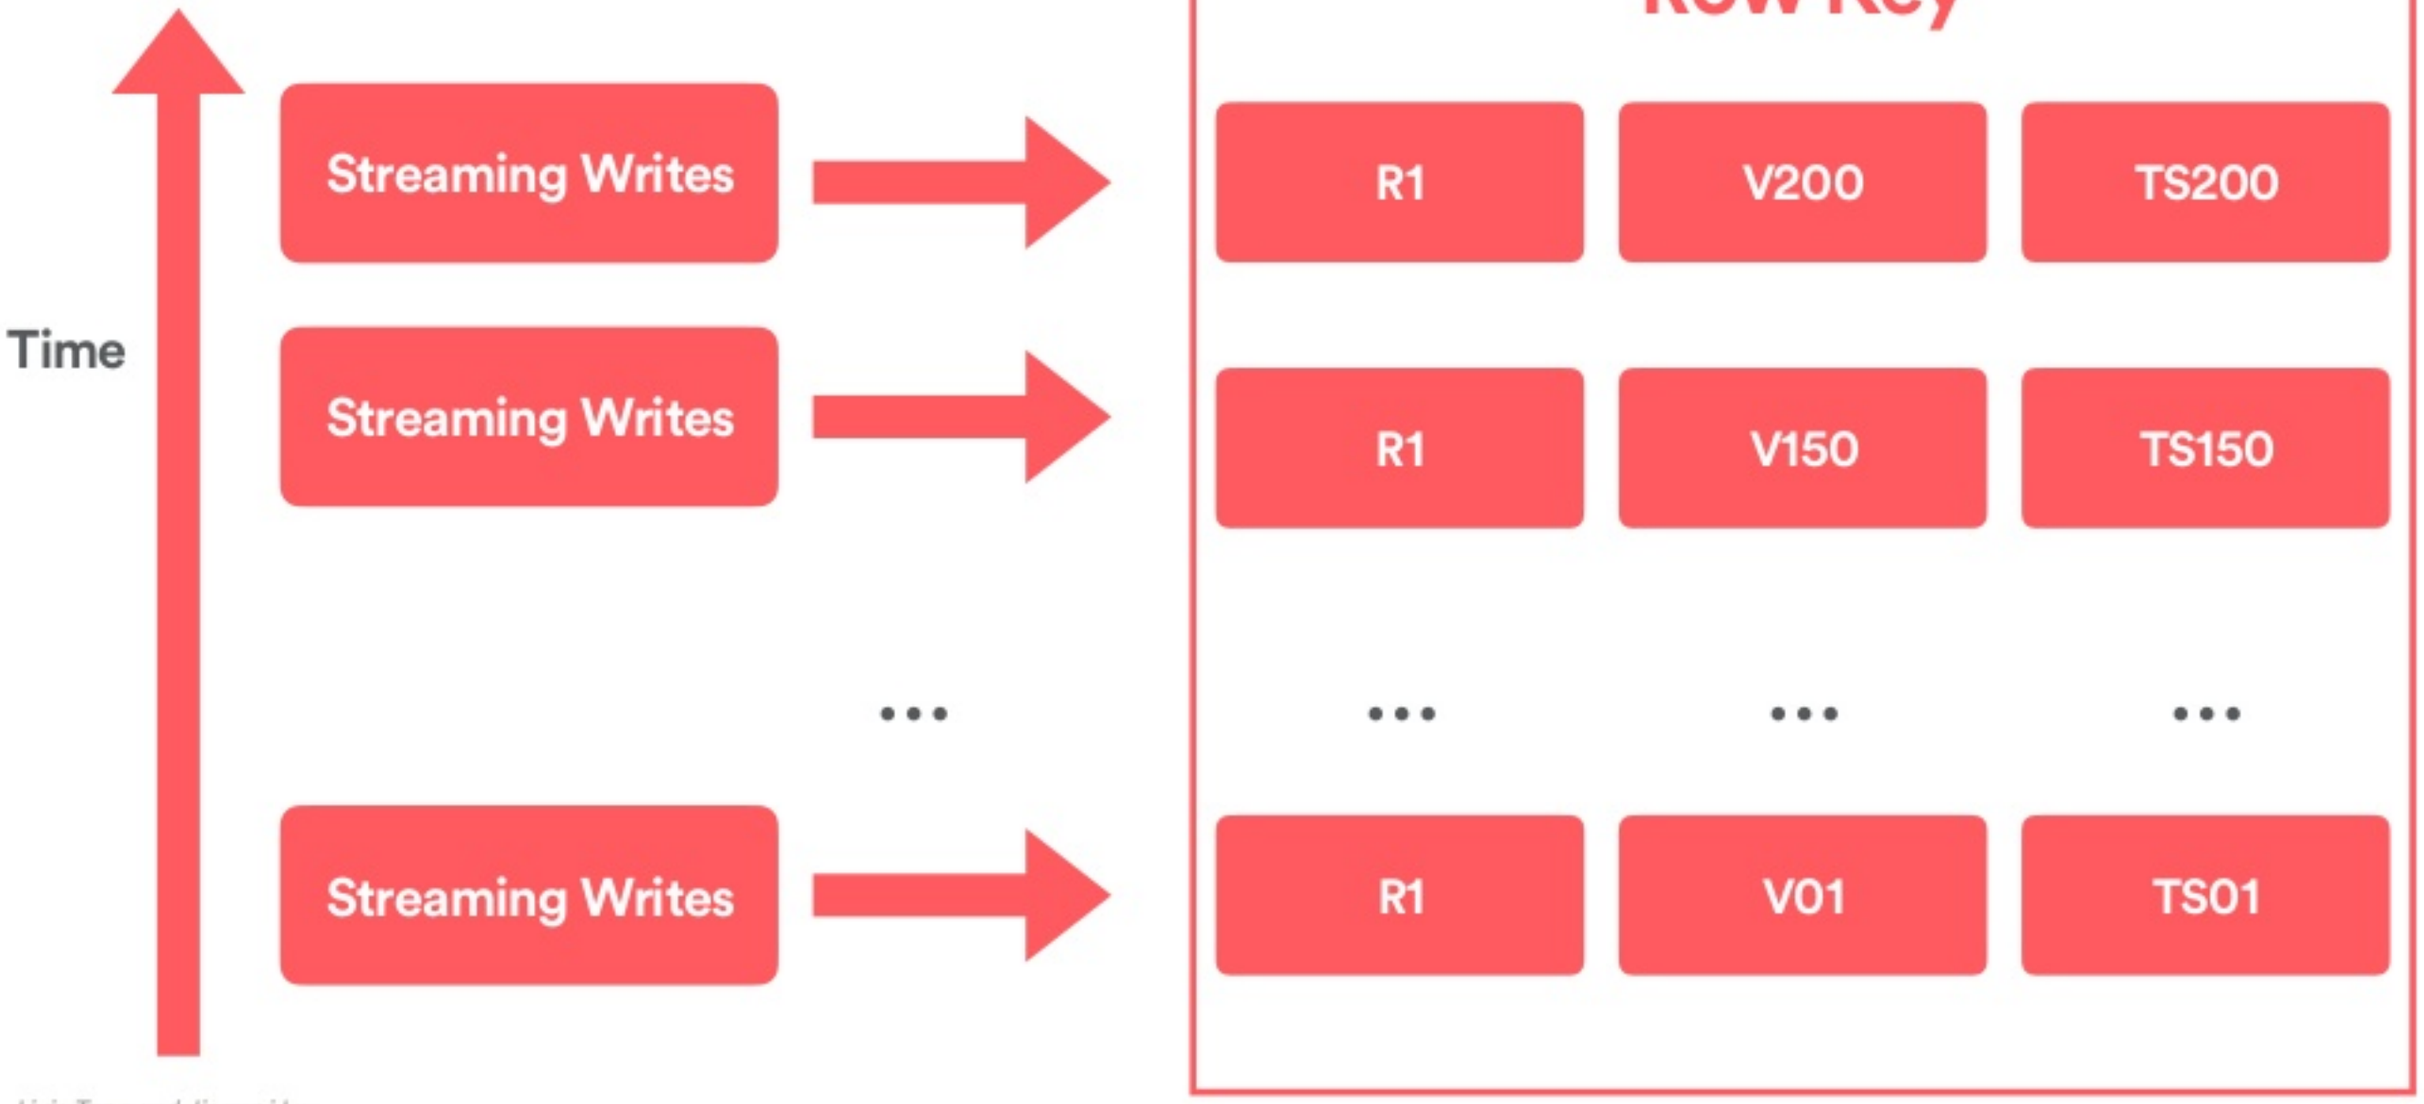

Liyin Tang and Jingwei Lu

Merged Views

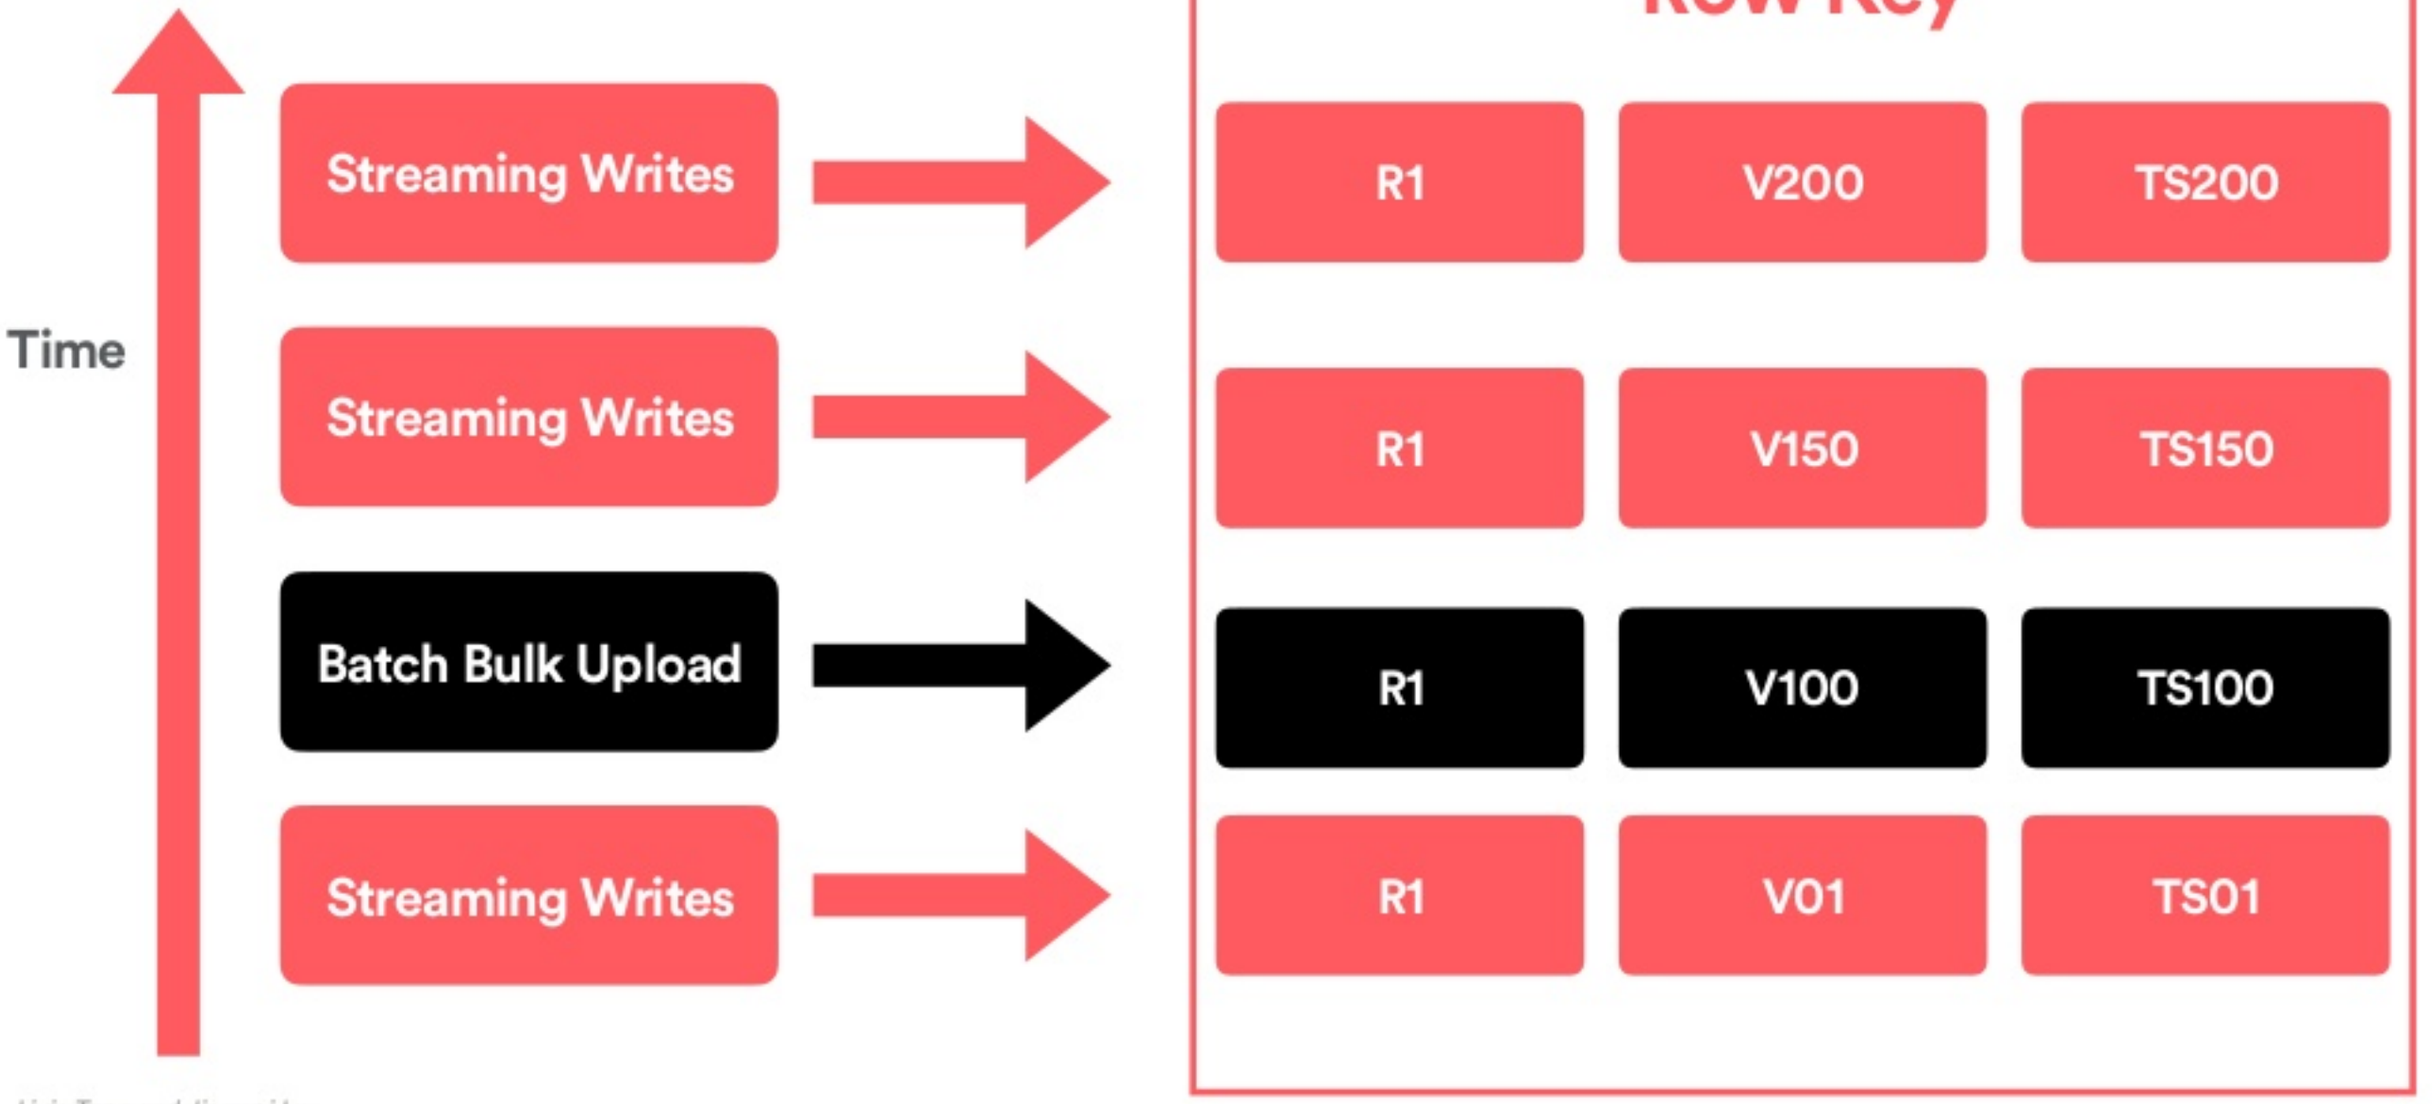

Liyin Tang and Jingwei Lu

Our Foundations

- Unify streaming with batch process
- Shared global state store

Use Cases

MySQL DB Snapshot Using Binlog Replay

Move Elephant

Database Snapshot

- **Large amount of data:** Multiple large mysql DBs
- **Realtime-ness:** minutes delay/ hours delay
- **Transaction :** Need to keep transaction across different tables
- **Schema change:** Table schema evolves

Binlog Replay on Spark

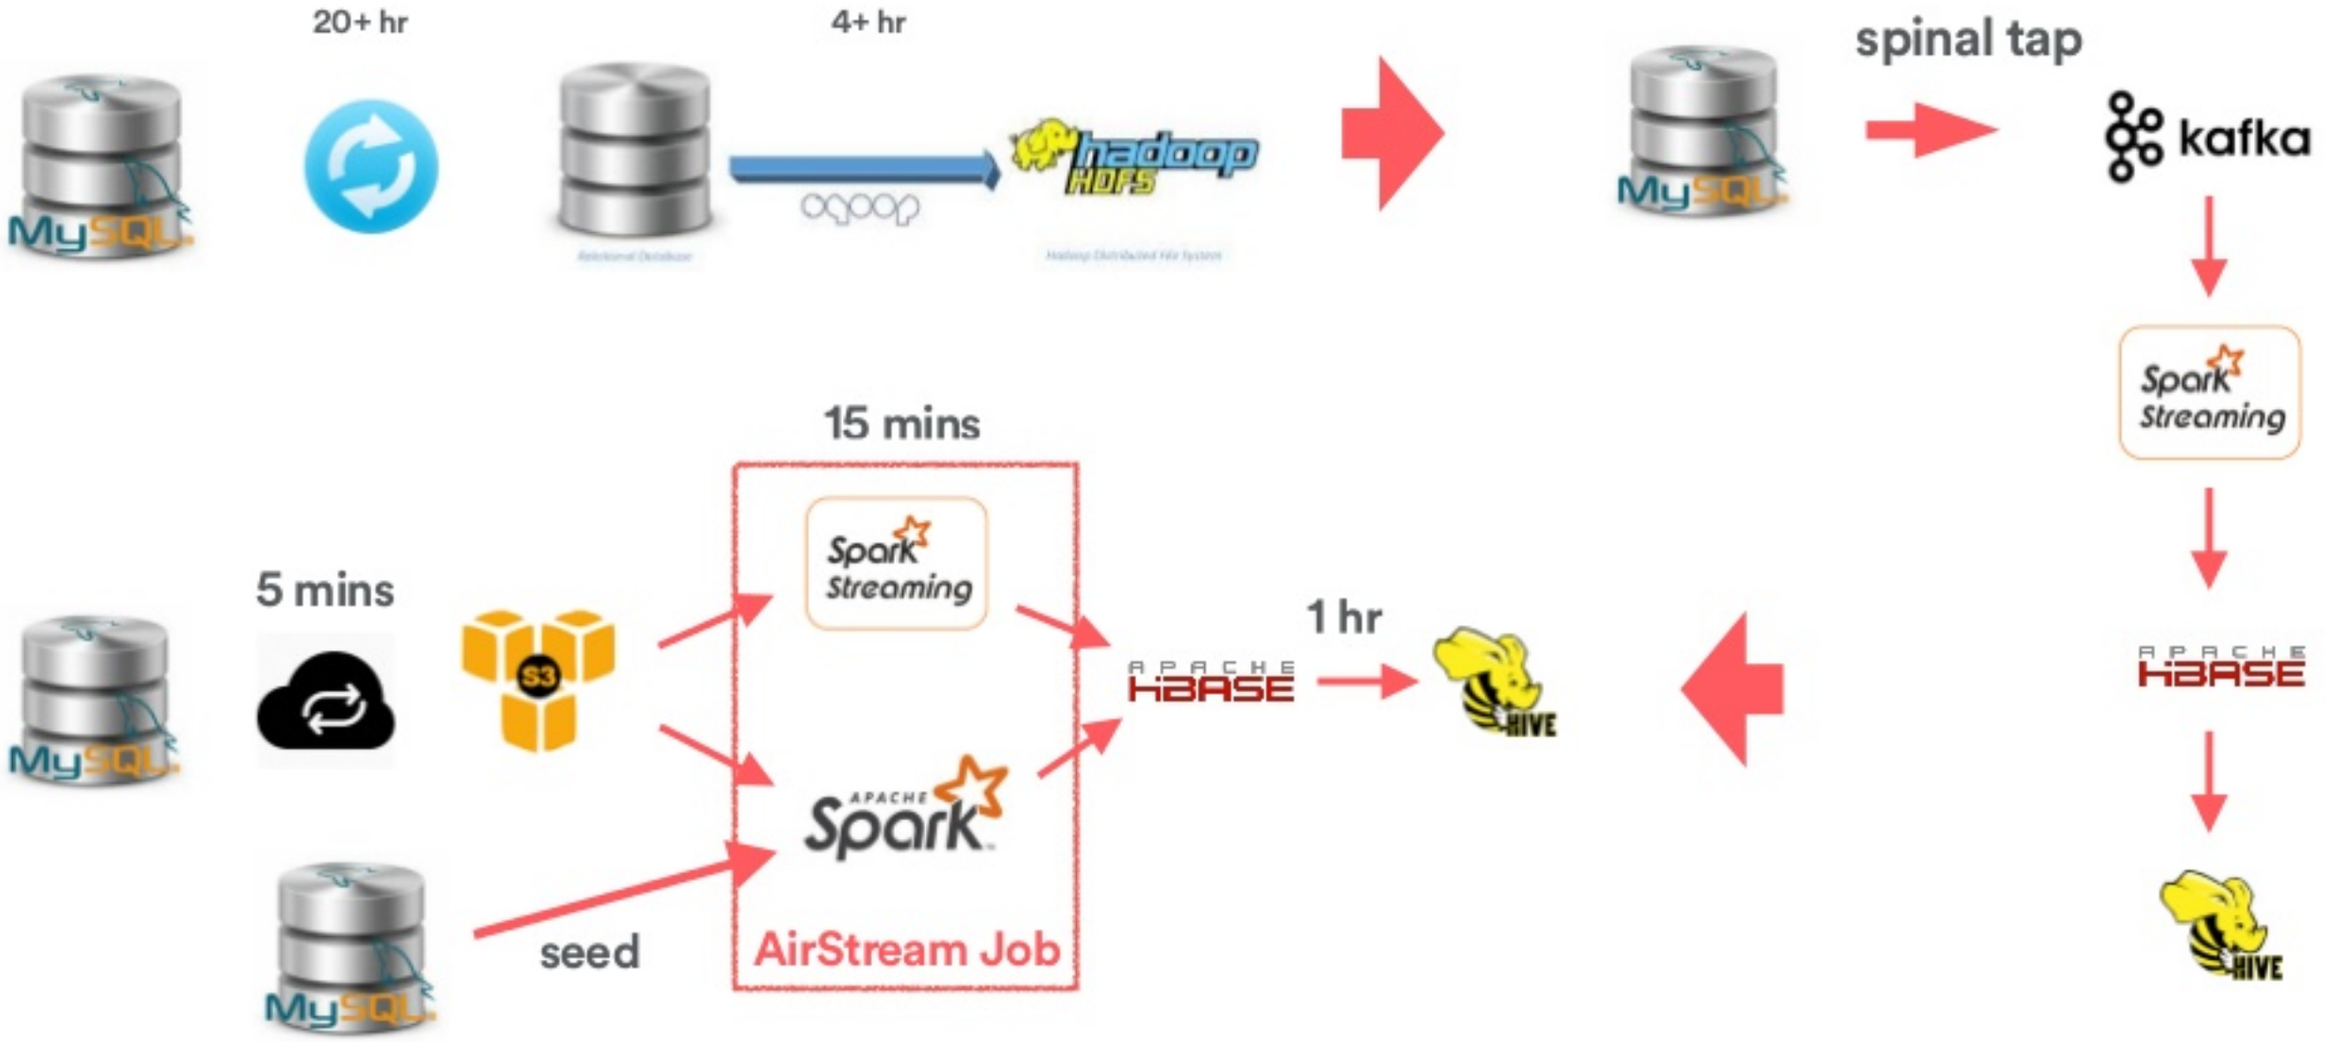

Lambda Architecture

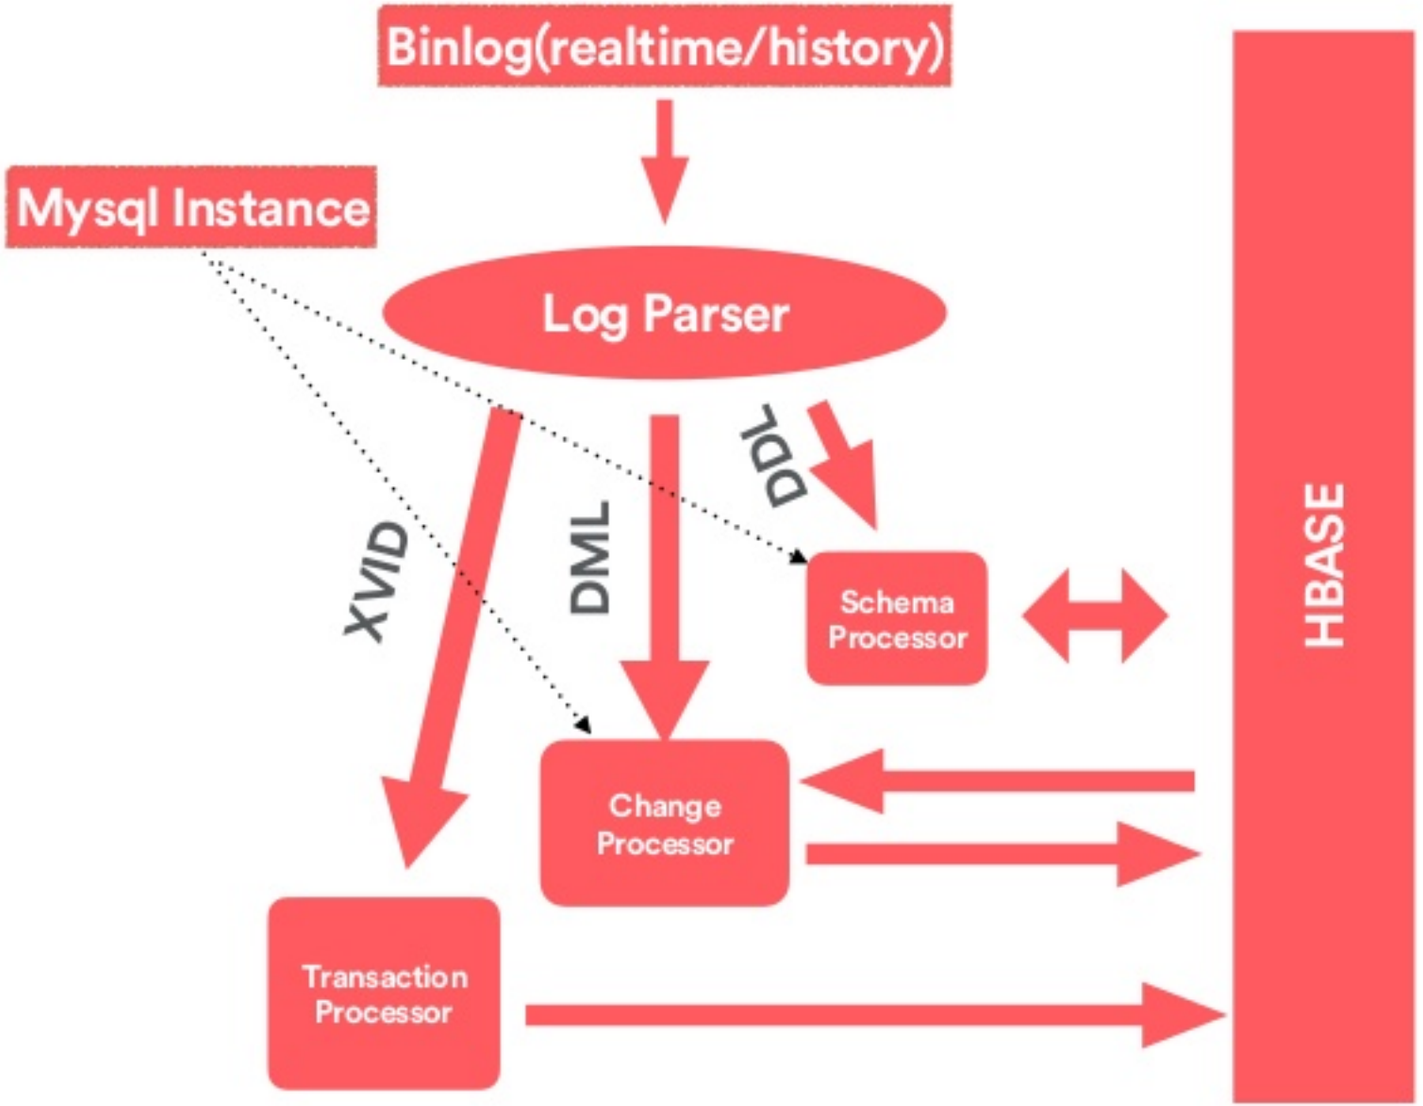

- **Streaming and Batch shares Logic:** Binlog file reader, DDL processor, transaction processor, DML processor.
- **Merged by binlog position:** <filenum, offset>
- **Idempotent:** Log can be replayed multiple times.
- **Schema changes:** Full schema change history.

Realtime Indexing

Realtime Indexing

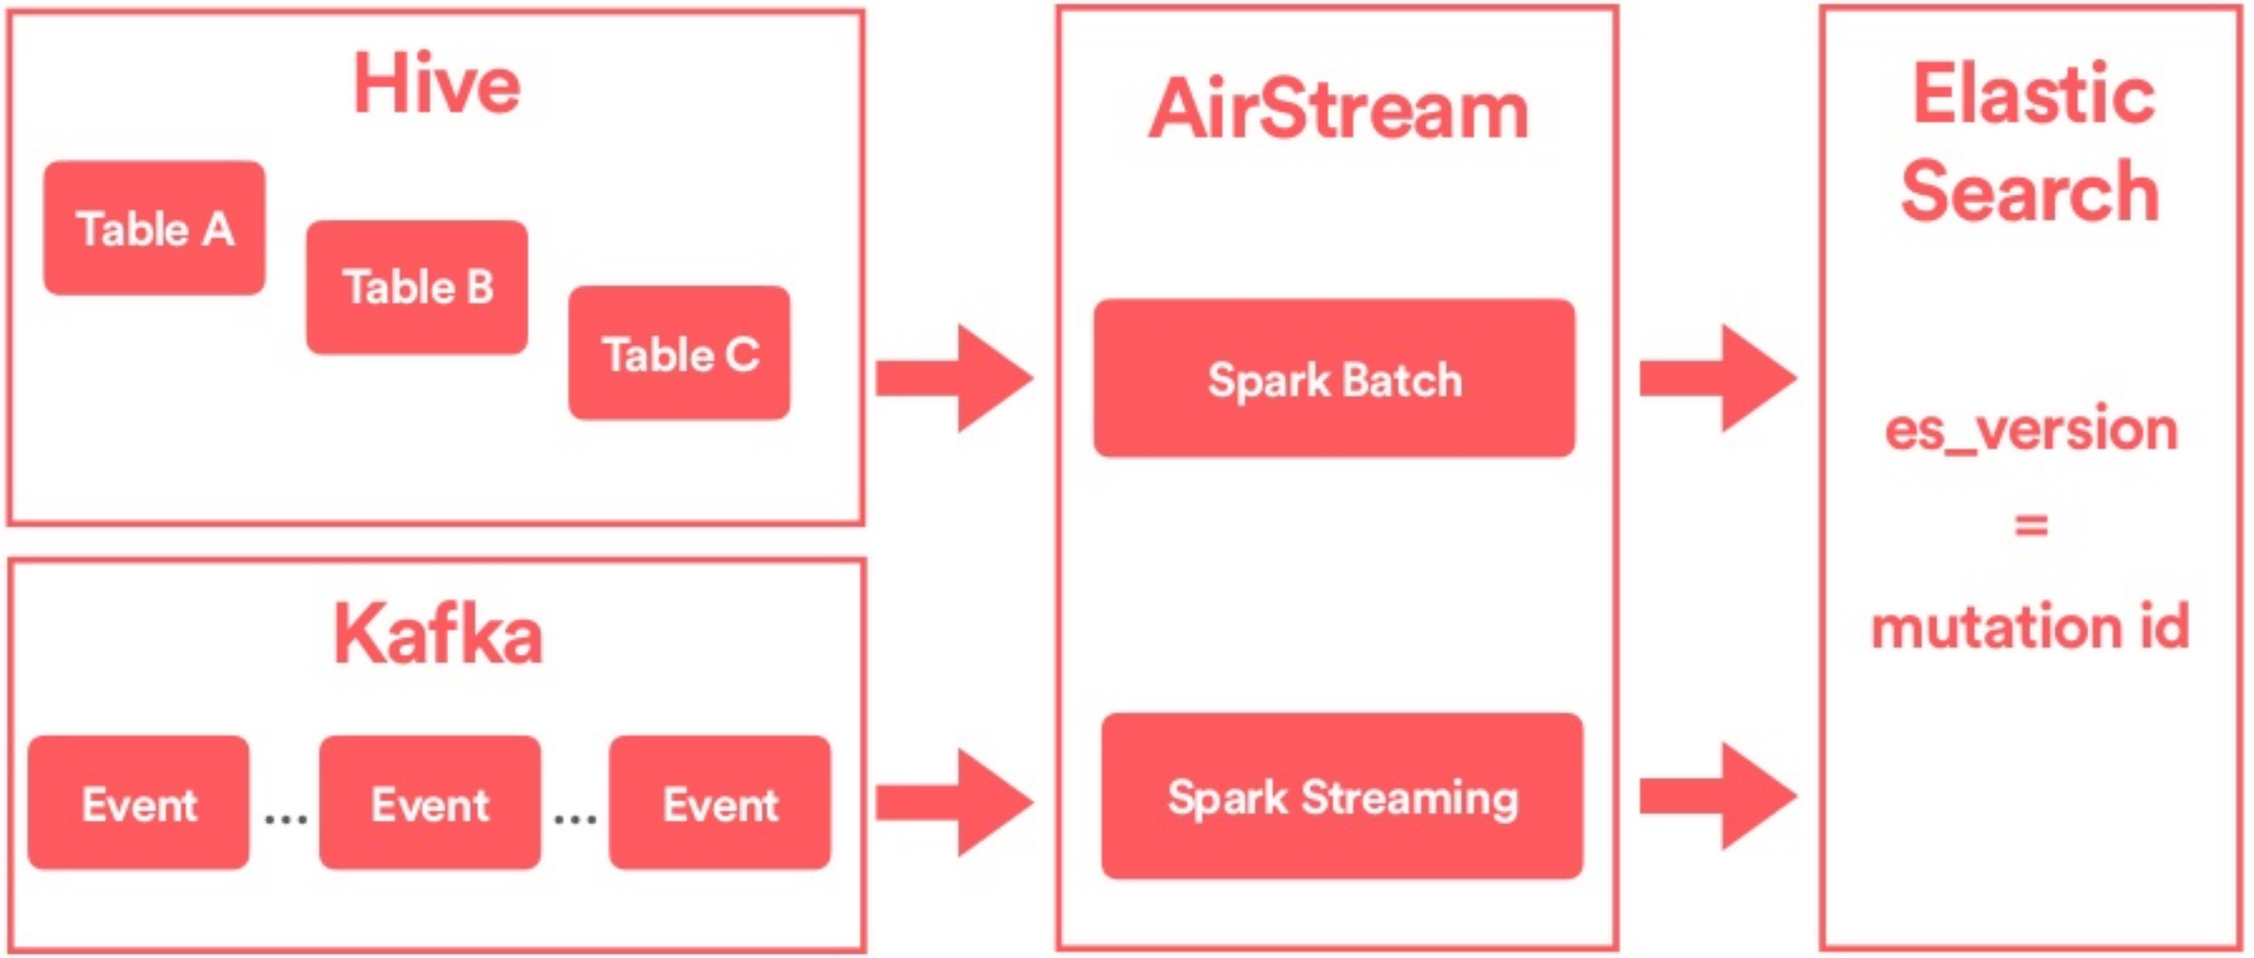

Realtime OLAP with Druid

Druid Ingestion

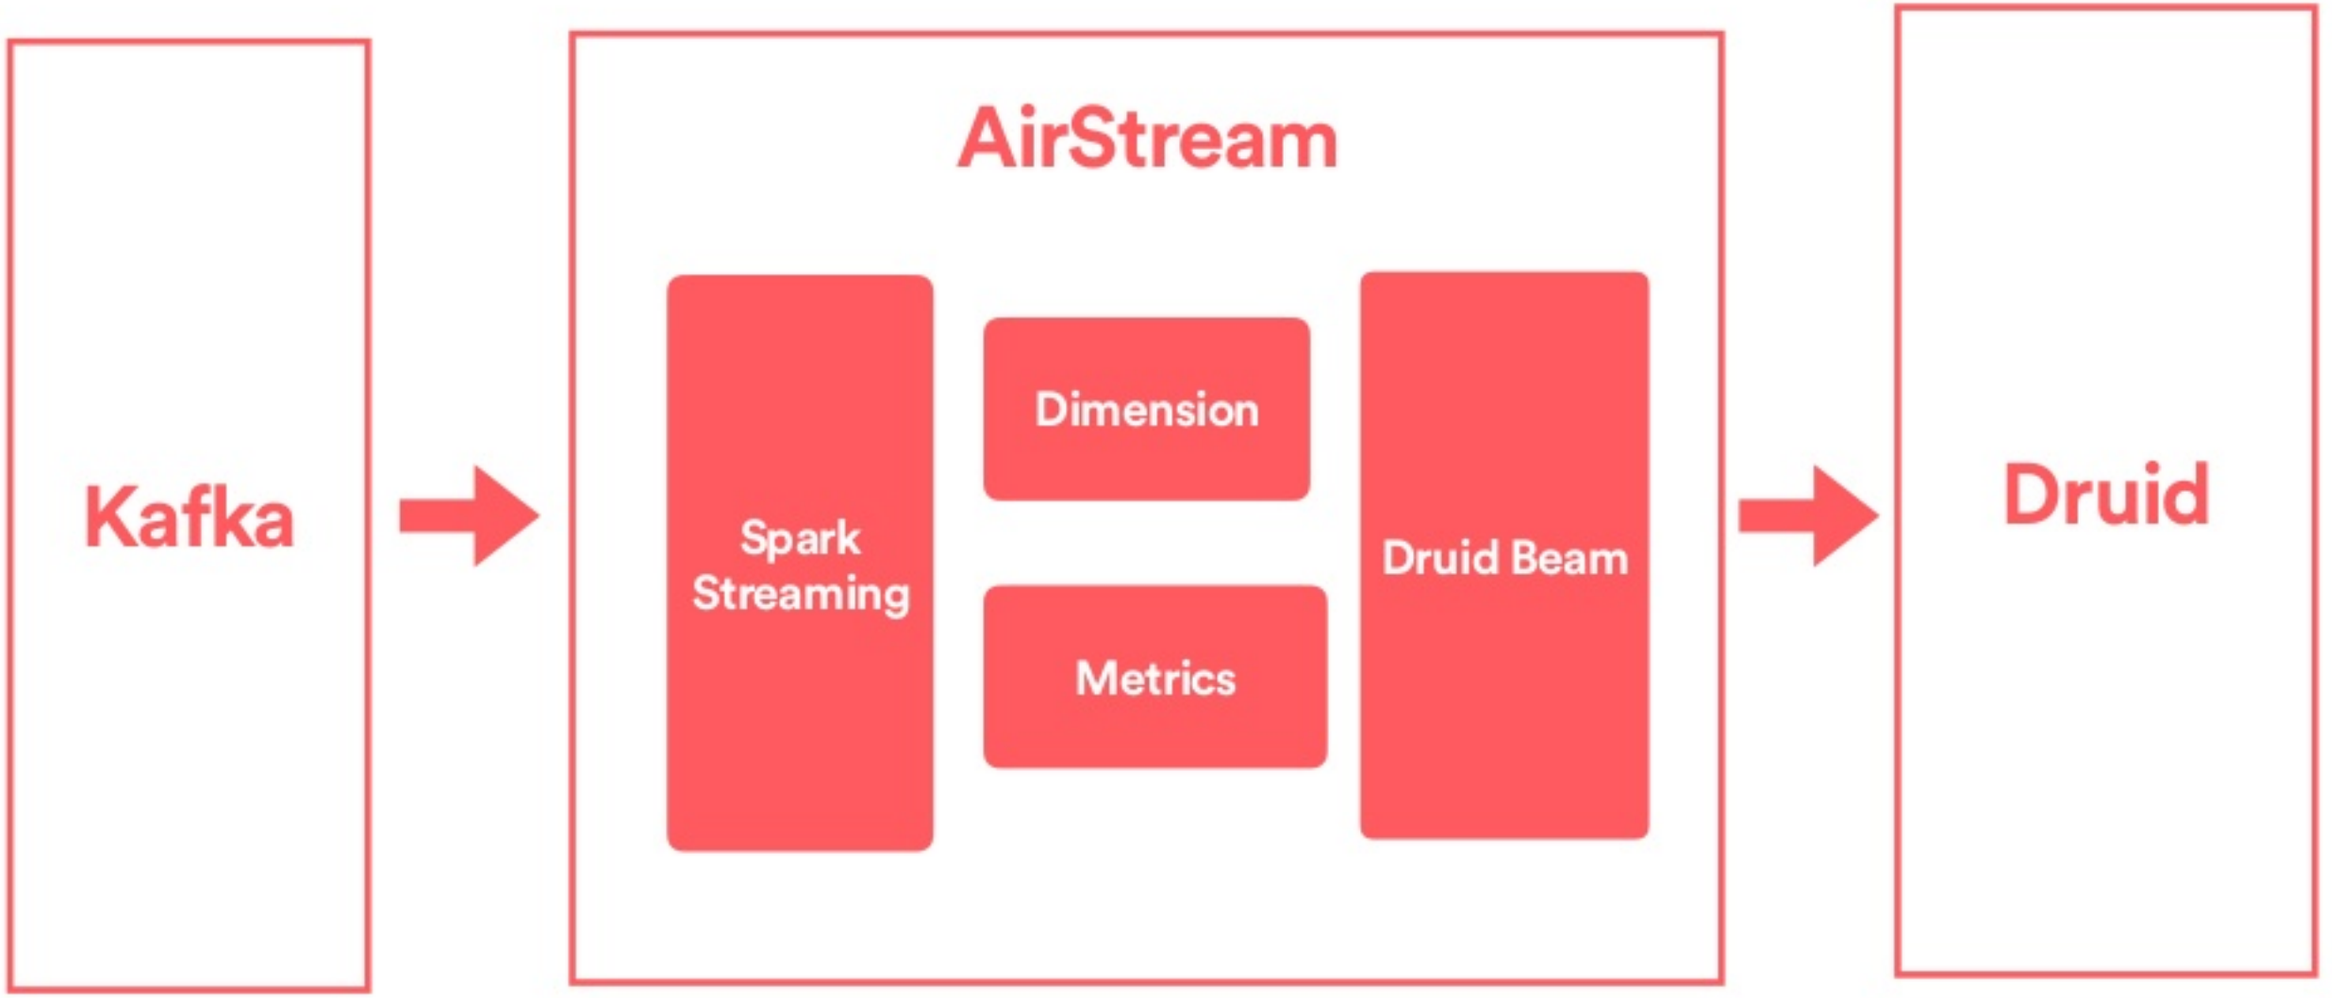

Liyin Tang and Jingwei Lu

Superset Powered by Druid

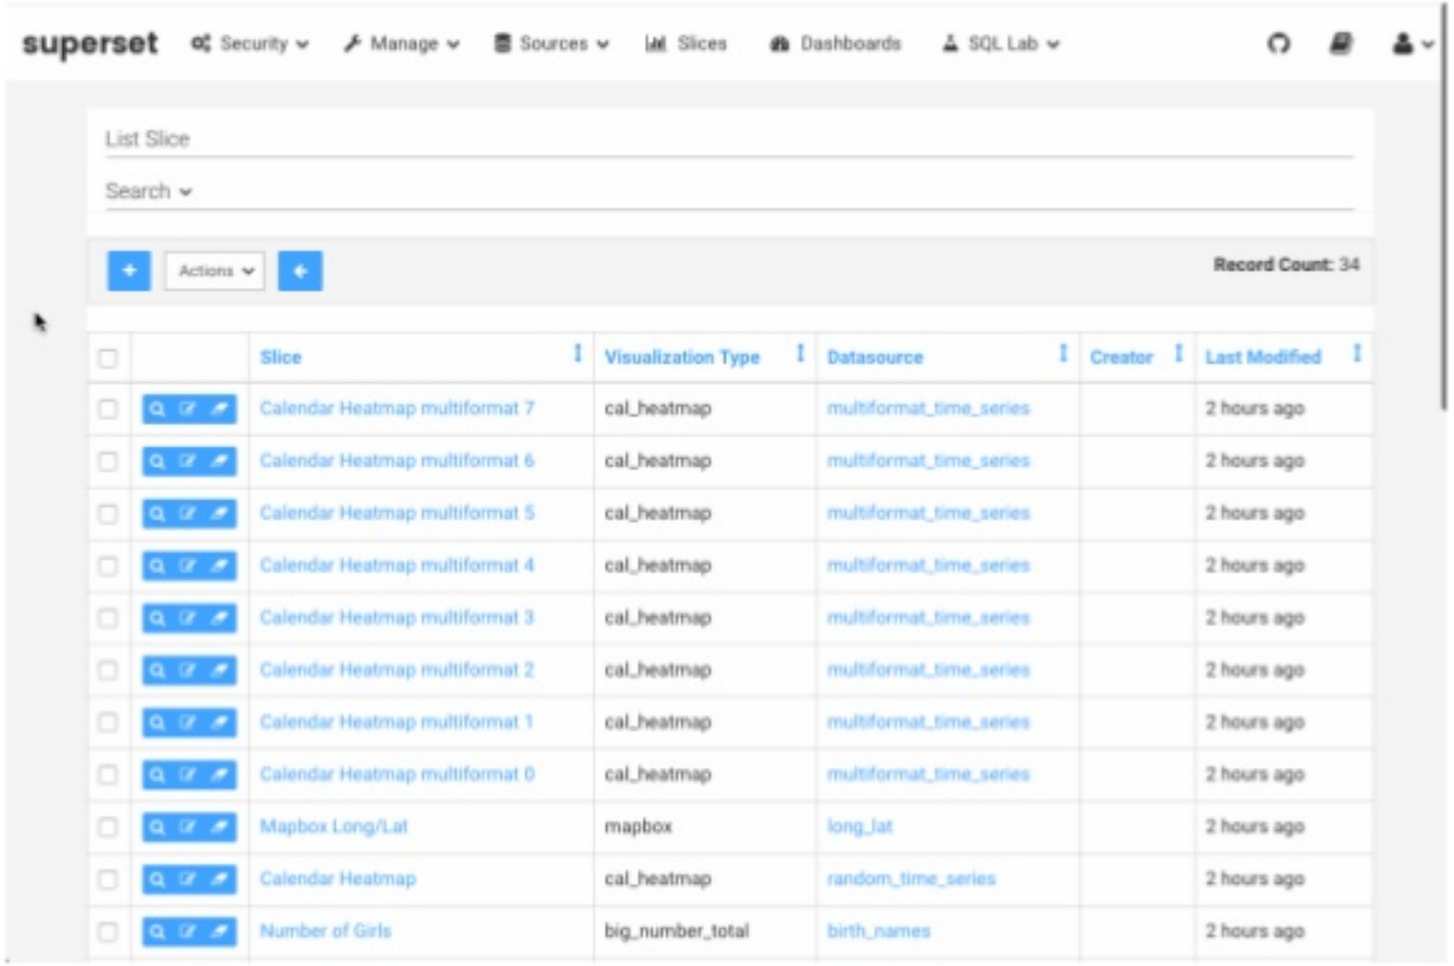

Tips

Moving Window Computation

Long Window Computation

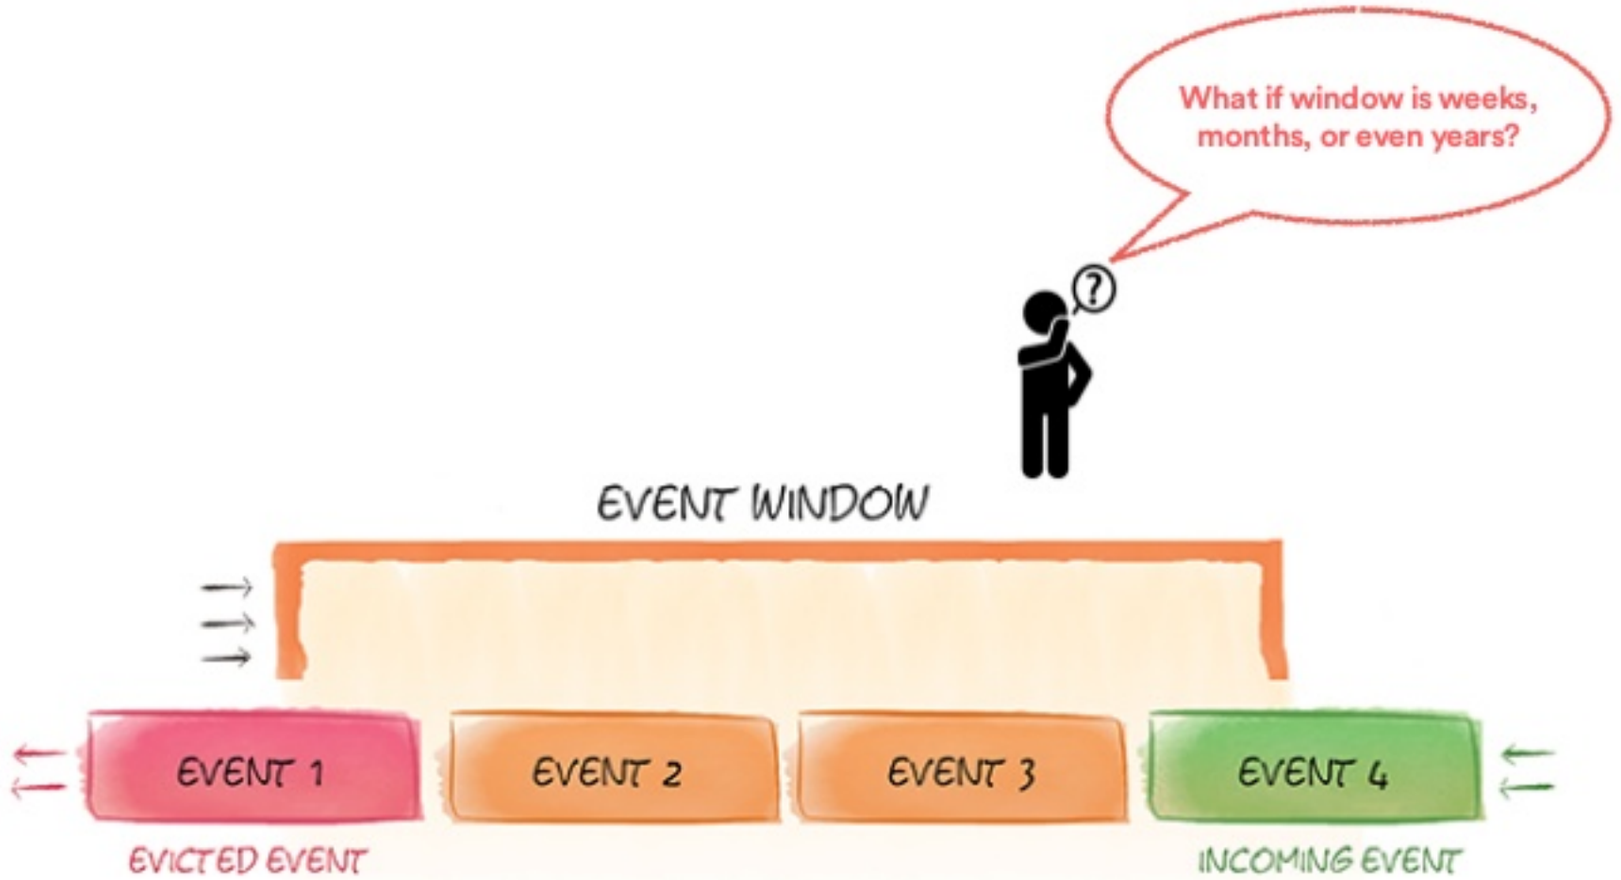

Distinct in a Large Window

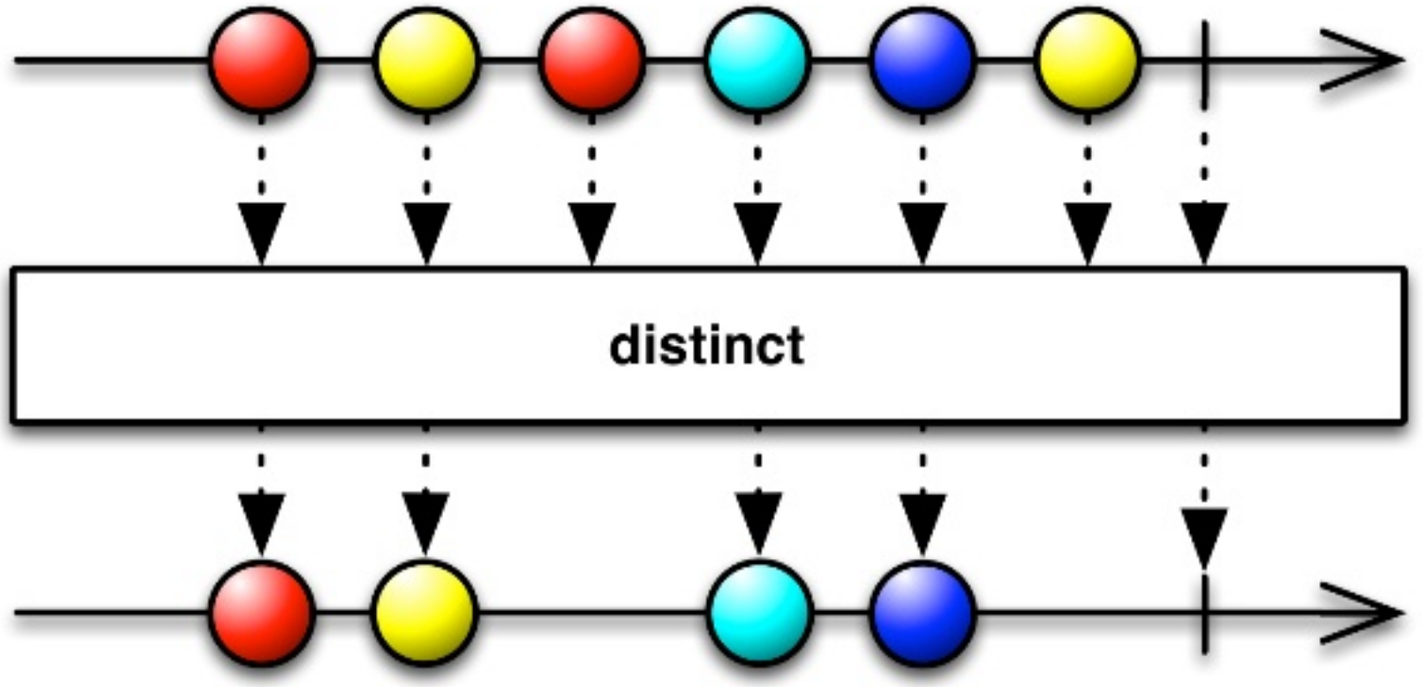

Distinct Count

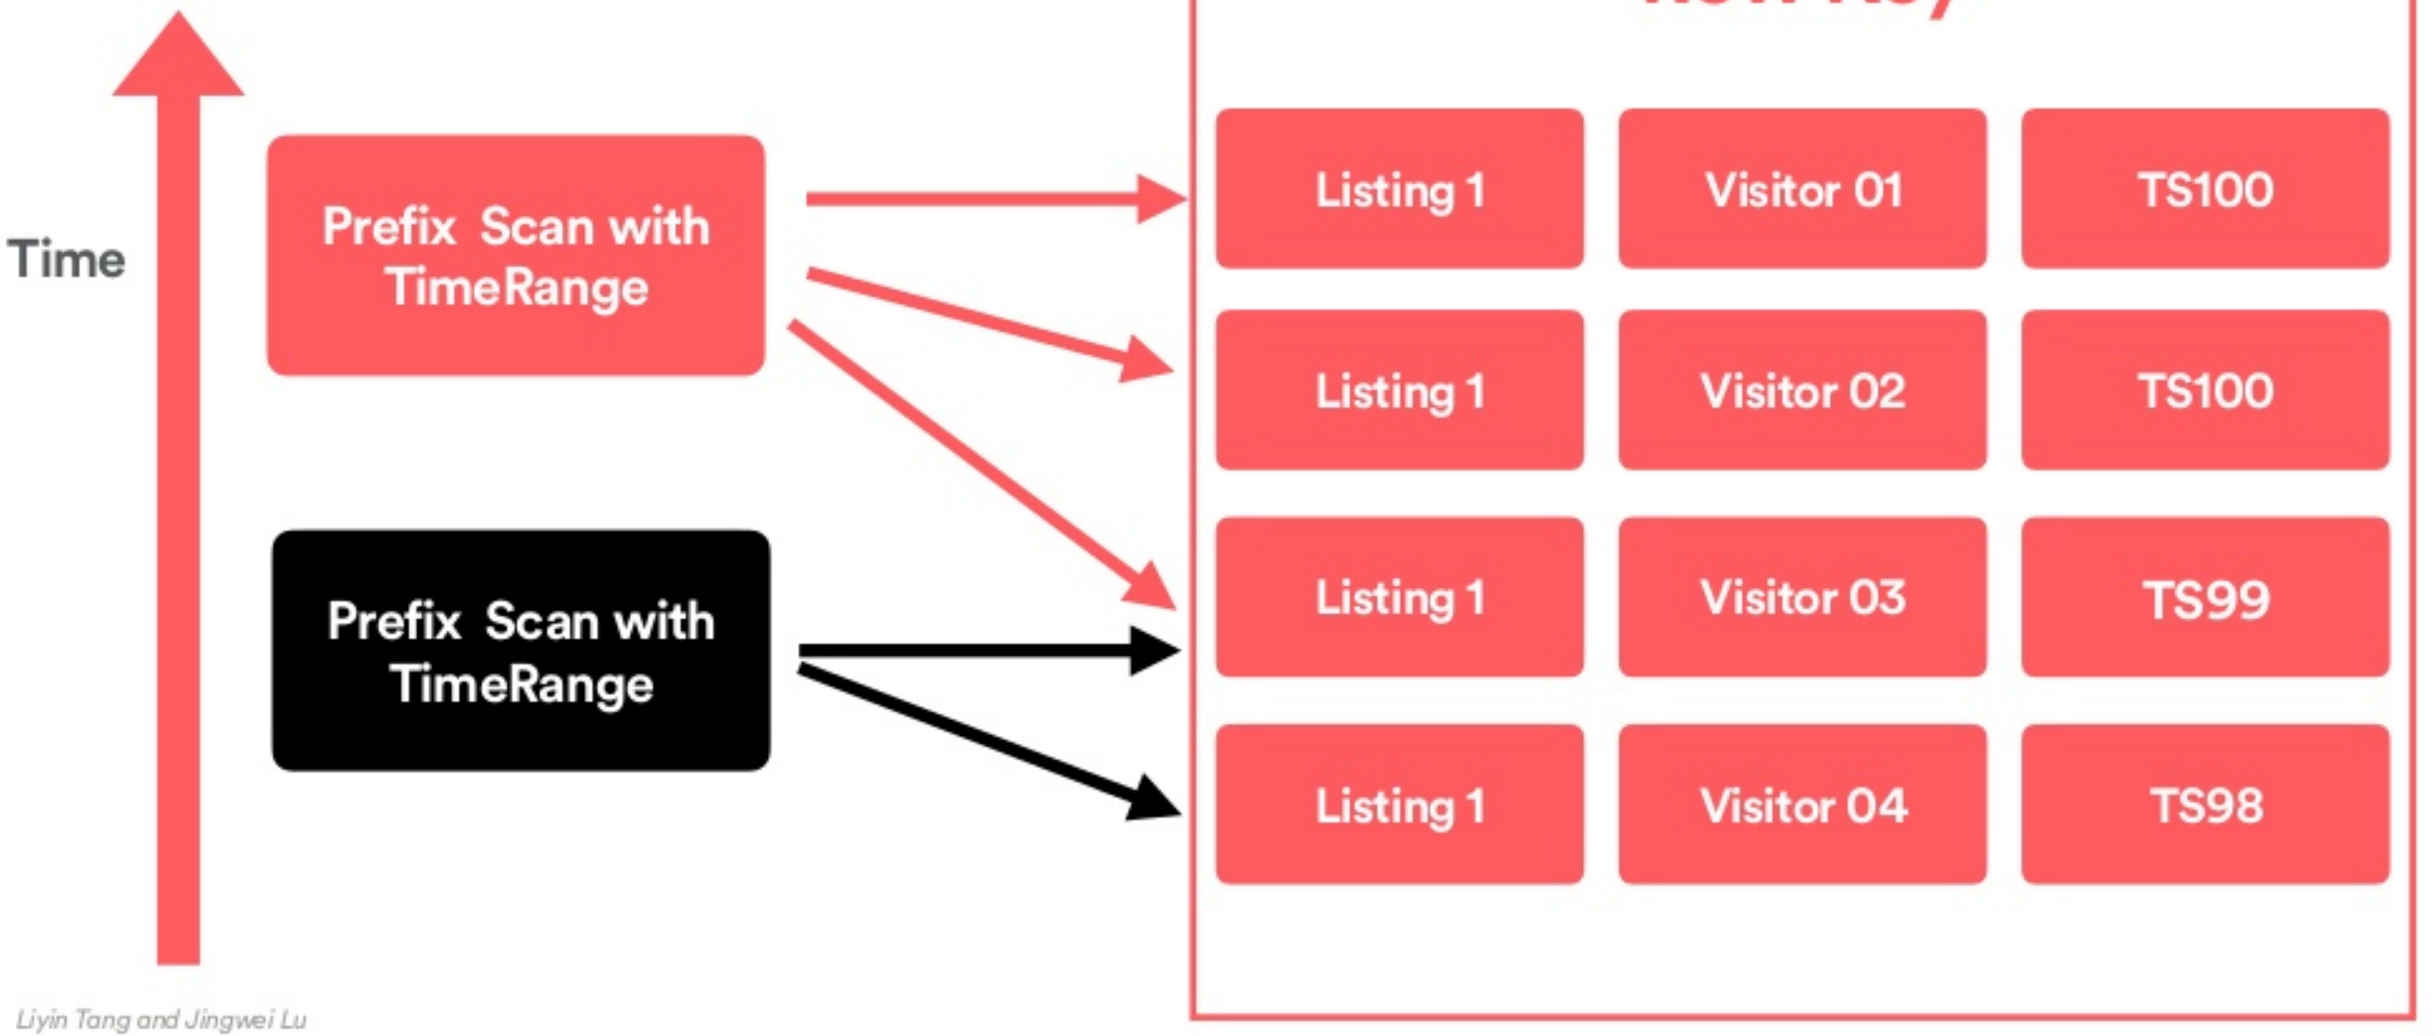

Liyin Tang and Jingwei Lu

Moving Average

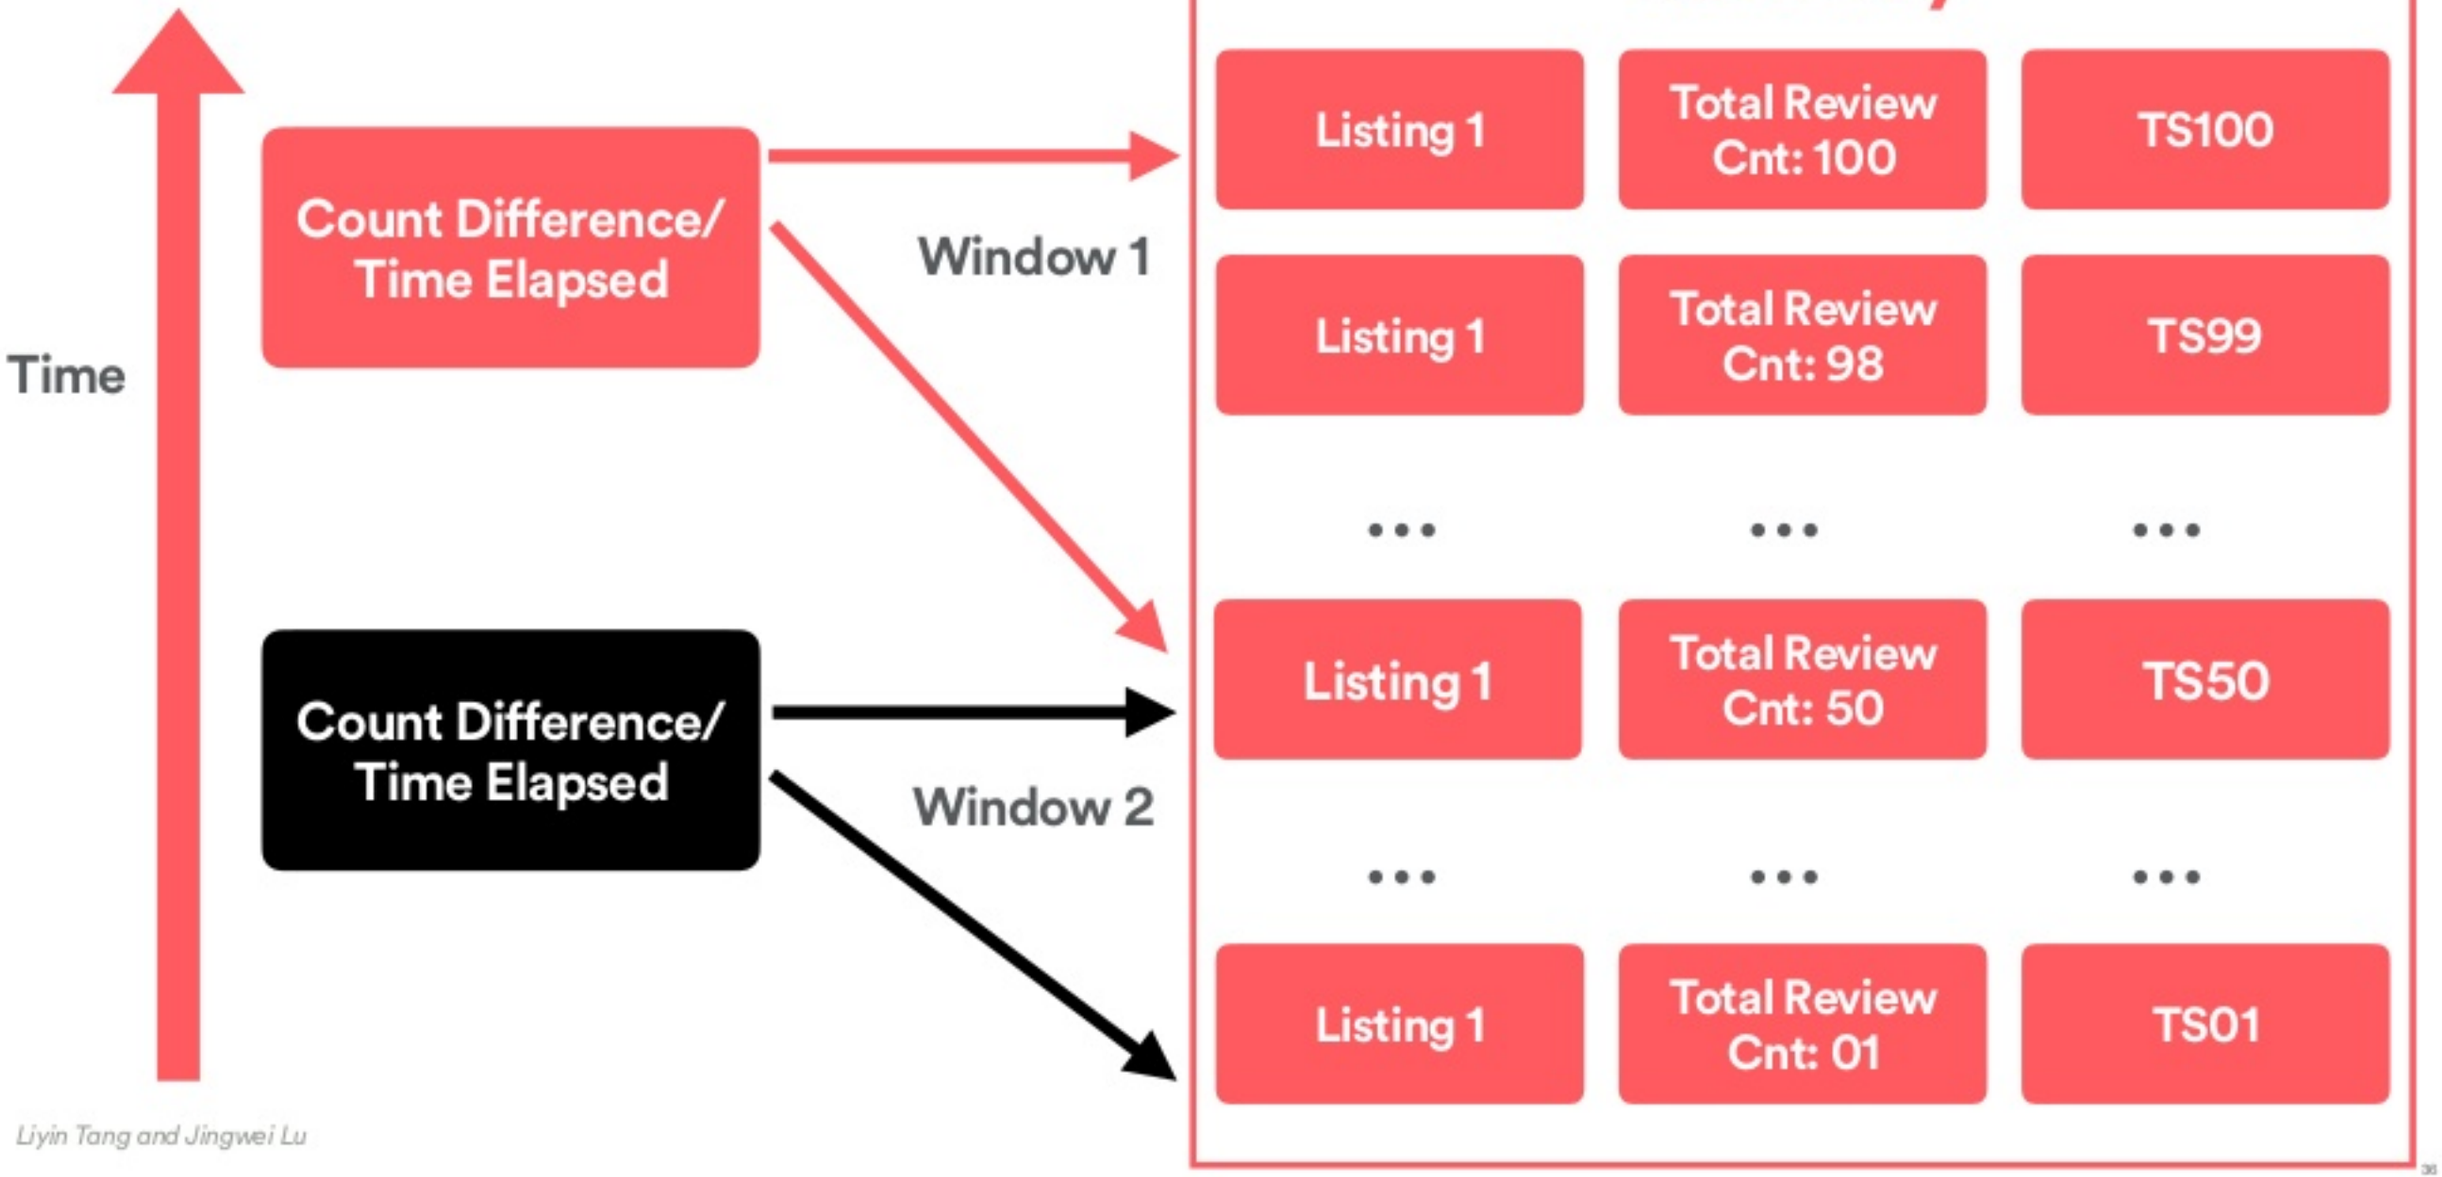

Streaming Ingestion & Realtime Interactive Query

Realtime Ingestion and Interactive Query

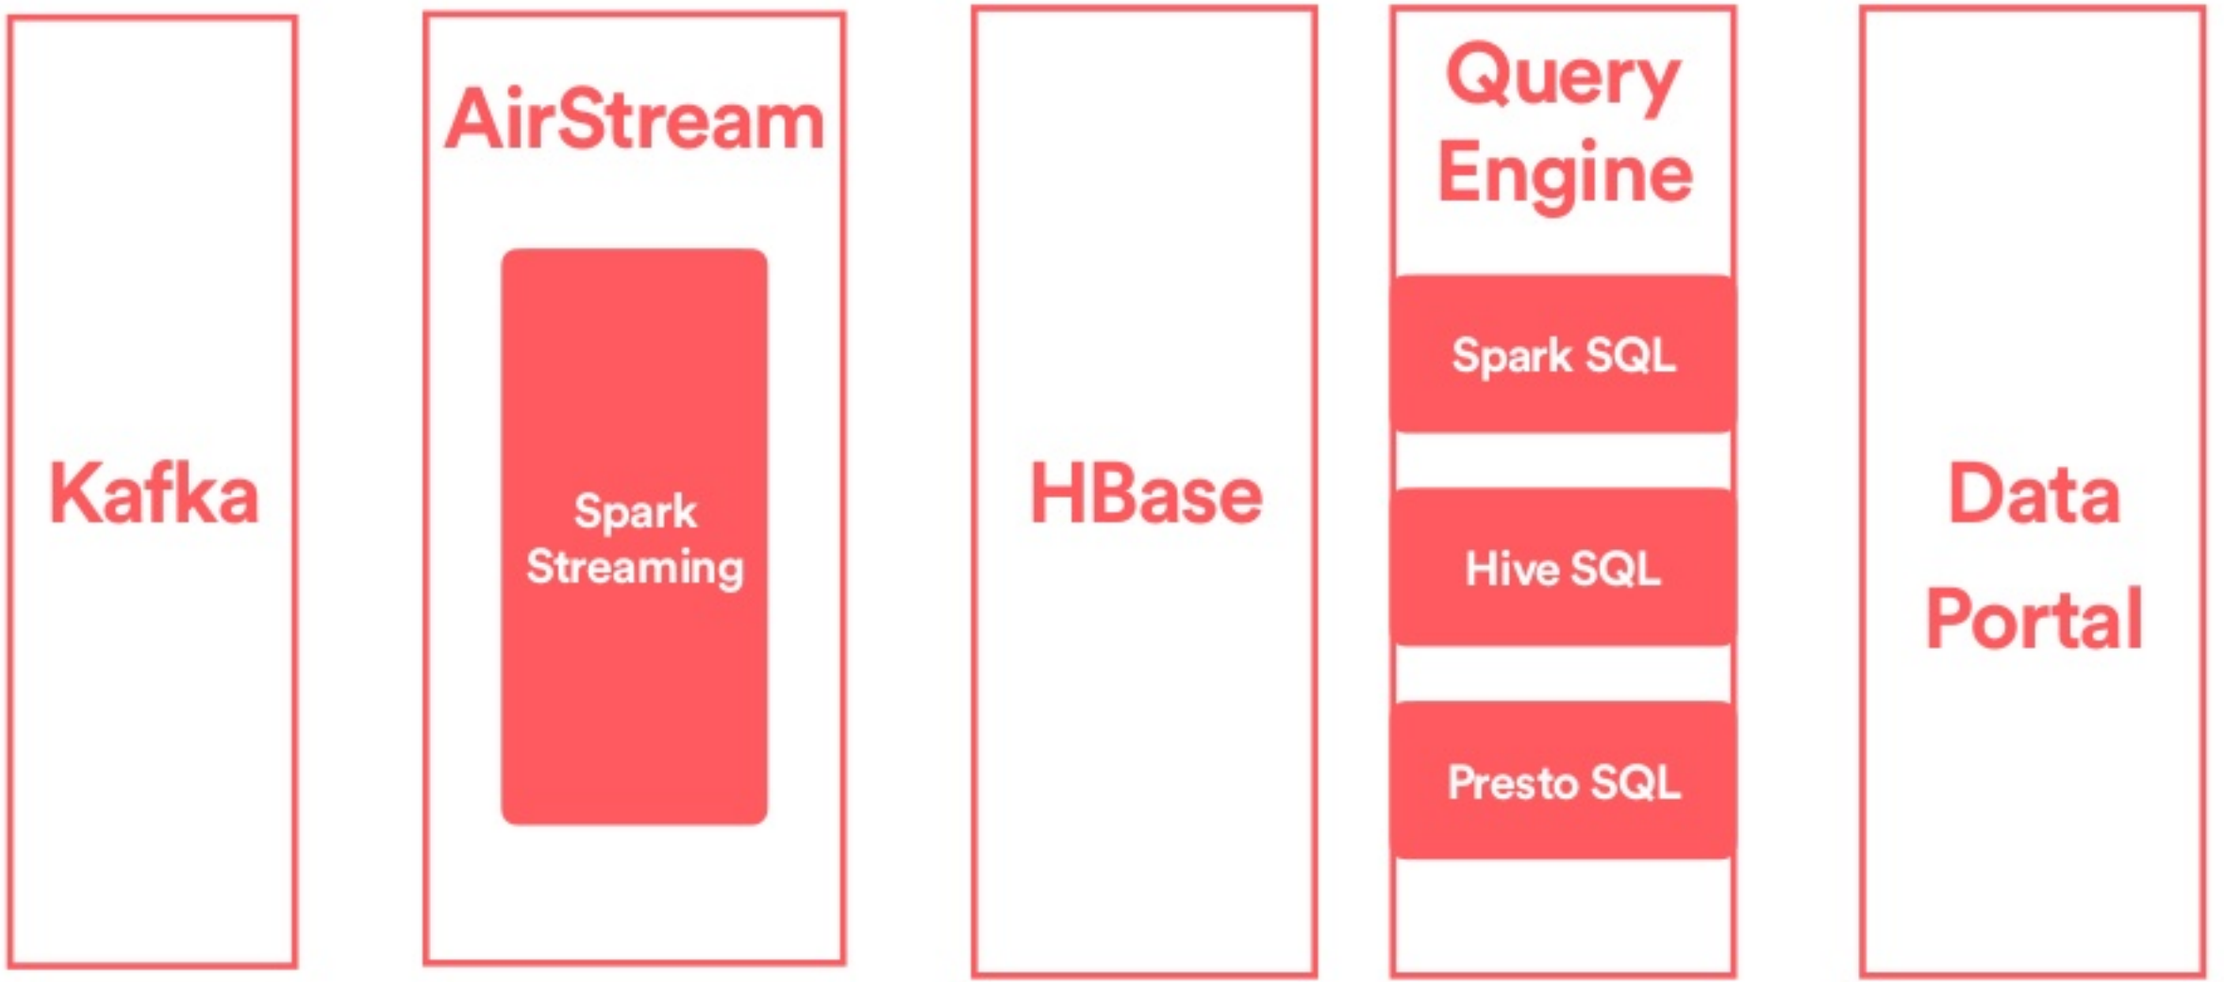

Liyin Tang and Jingwei Lu

Interactive Query in SqlLab

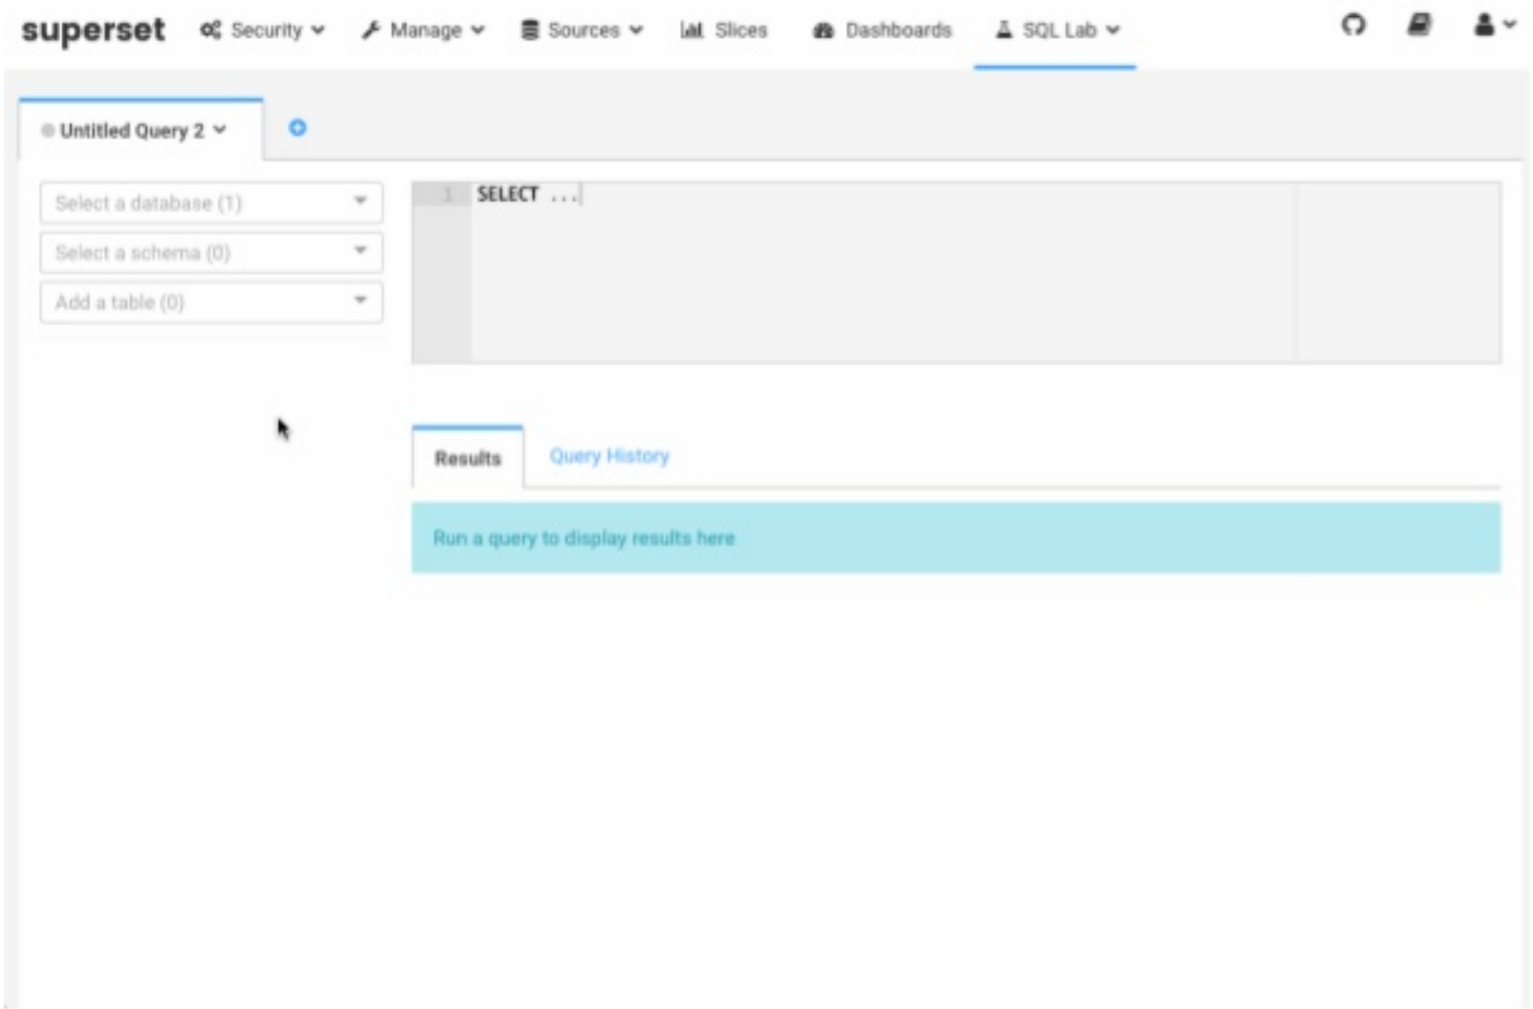

Schema Enforcement Streaming Events

Thrift -> DataFrame

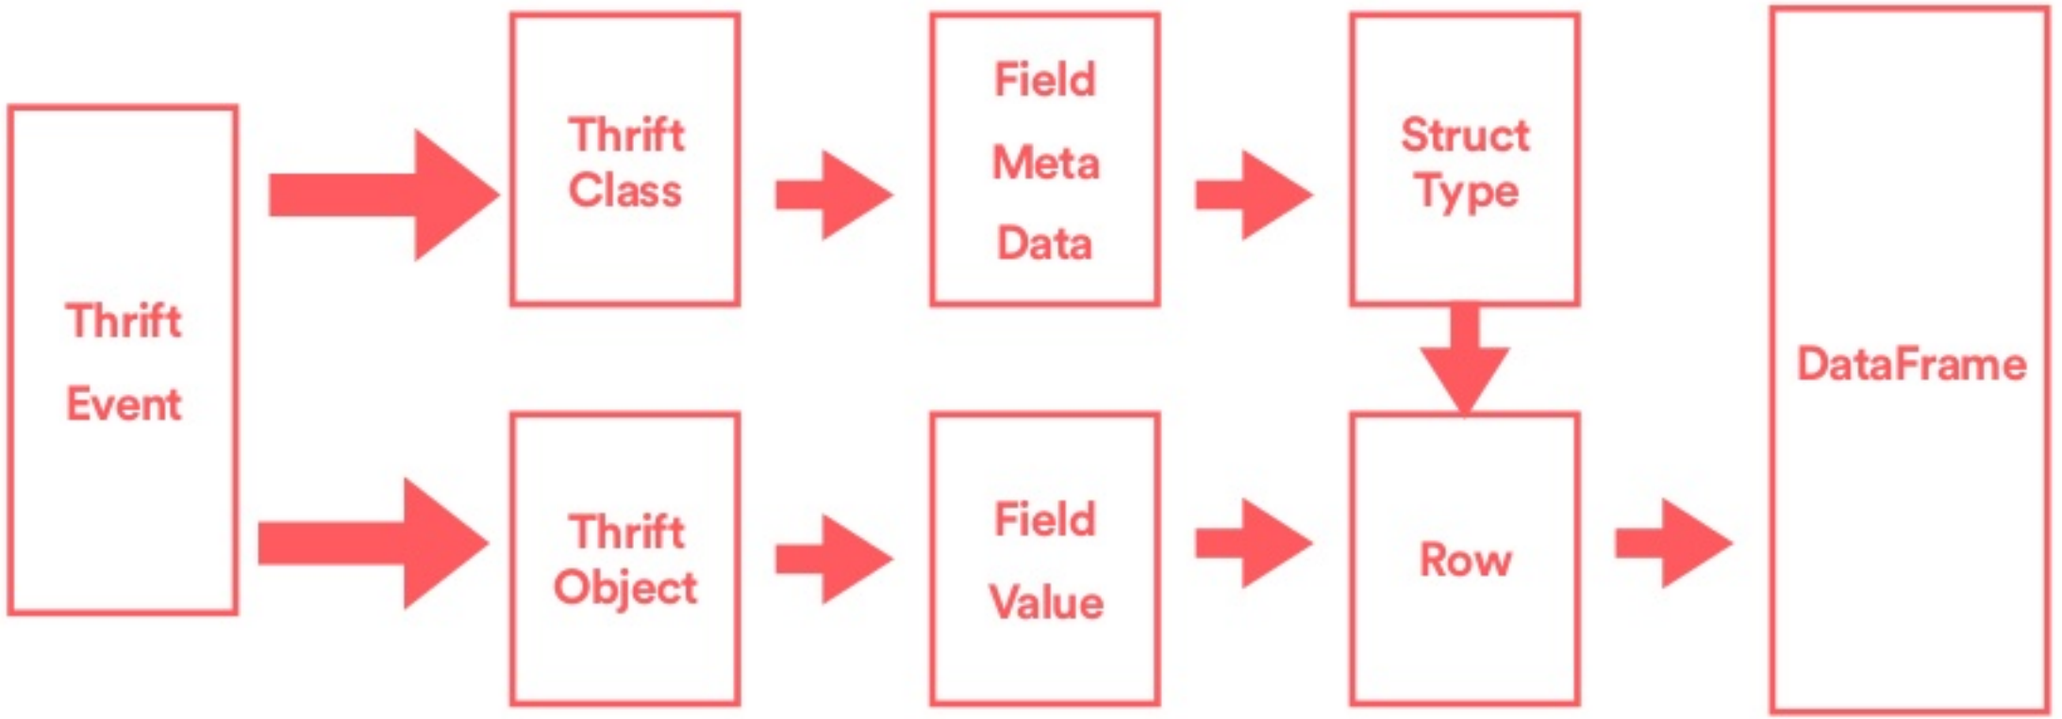

https://github.com/airbnb/airbnb-spark-thrift

Summary

Unify Batch and Streaming Computation

Global State Store Using HBase

airbnb

We are hiring

Happy Hour:
6pm, B Restaurant & Bar, 720 Howard St, SF