




Challenging Web-Scale Graph Analytics with Apache Spark

Xiangrui Meng

Spark Summit 2017



About me

- Software Engineer at Databricks
 - machine learning and data science/engineering
- Committer and PMC member of Apache Spark
 - MLlib, SparkR, PySpark, Spark Packages, etc

GraphFrames

GraphFrames

- A Spark package introduced in 2016 (graphframes.github.io)
 - collaboration between Databricks, UC Berkeley, and MIT
- GraphX to RDDs as GraphFrames are to DataFrames
 - Python, Java, and Scala APIs,
 - expressive graph queries,
 - query plan optimizers from Spark SQL,
 - graph algorithms.

Quick examples

Launch a Spark shell with GraphFrames:

```
spark-shell --packages graphframes:graphframes:0.5.0-spark2.1-s_2.11
```

Or [try it on Databricks Community Edition \(databricks.com/try\)](https://databricks.com/try).

Quick examples

Find 2nd-degree followers:

```
g.find("(A)-[]->(B); (B)-[]->(C); !(A)-[]->(C)")  
  .filter("A.id != C.id")  
  .select("A", "C")
```

Quick examples

Compute PageRank:

```
g.pageRank(resetProbability=0.15, maxIter=20)
```

Supported graph algorithms

- breath-first search (BFS)
- connected components
 - strongly connected components
- label propagation algorithm (LPA)
- PageRank and personalized PageRank
- shortest paths
- triangle count

Moving implementations to DataFrames

- Several algorithms in GraphFrames are simple wrappers over GraphX RDD implementations, which do not scale very well.
- DataFrames are optimized for a huge number of small records.
 - columnar storage
 - code generation
 - query optimization

Assigning integral vertex IDs

... lessons learned

Pros of having integral vertex IDs

GraphFrames take string vertex identifiers, whose values are not used in graph algorithms. Having integral vertex IDs can help

- optimize in-memory storage,
- save communication.

So the task is to map unique vertex identifiers to unique (long) integers.

The hashing trick?

- It is easy to hash the vertex identifier to a long integer.
- What is the chance of collision?
 - $1 - (k-1)/N * (k-2)/N * \dots$
 - seems unlikely with long range $N=2^{64}$
 - with 1 billion nodes, the chance is ~5.4%
- And collisions change graph topology.

Name	Hash
Tim	84088
Joseph	-2070372689
Xiangrui	264245405
Felix	67762524

Generating unique IDs

Spark has builtin methods to generate unique IDs.

- RDD: `zipWithUniqueId()` / `zipWithIndex()`
- DataFrame: `monotonically_increasing_id()`

So given a DataFrame of distinct vertex identifiers, we can add a new column with generated unique long IDs. Simple?

How it works?

Partition 1		Partition 2		Partition 3	
Vertex	ID	Vertex	ID	Vertex	ID
Tim	0	Xiangrui	$100 + 0$...	$200 + 0$
Joseph	1	Felix	$100 + 1$...	$200 + 1$

... but not always work

- DataFrames/RDDs are immutable and reproducible by design.
- However, records do not always have stable order.
 - distinct
 - repartition
- And cache doesn't help.

Partition 1		re-compute →	Partition 1	
Vertex	ID		Vertex	ID
Tim	0		Joseph	0
Joseph	1		Tim	1

Our implementation


We implemented (v0.5.0) an expensive but correct version:

1. (hash) re-partition + distinct vertex identifiers,
2. sort vertex identifiers within each partition,
3. generate unique integral IDs

Connected Components

Connected Components

- Assign each vertex a component ID such that vertices receive the same component ID iff they are connected.
- Applications:
 - fraud detection
 - [Spark Summit 2016 keynote from Capital One](#)
 - clustering



A naive implementation

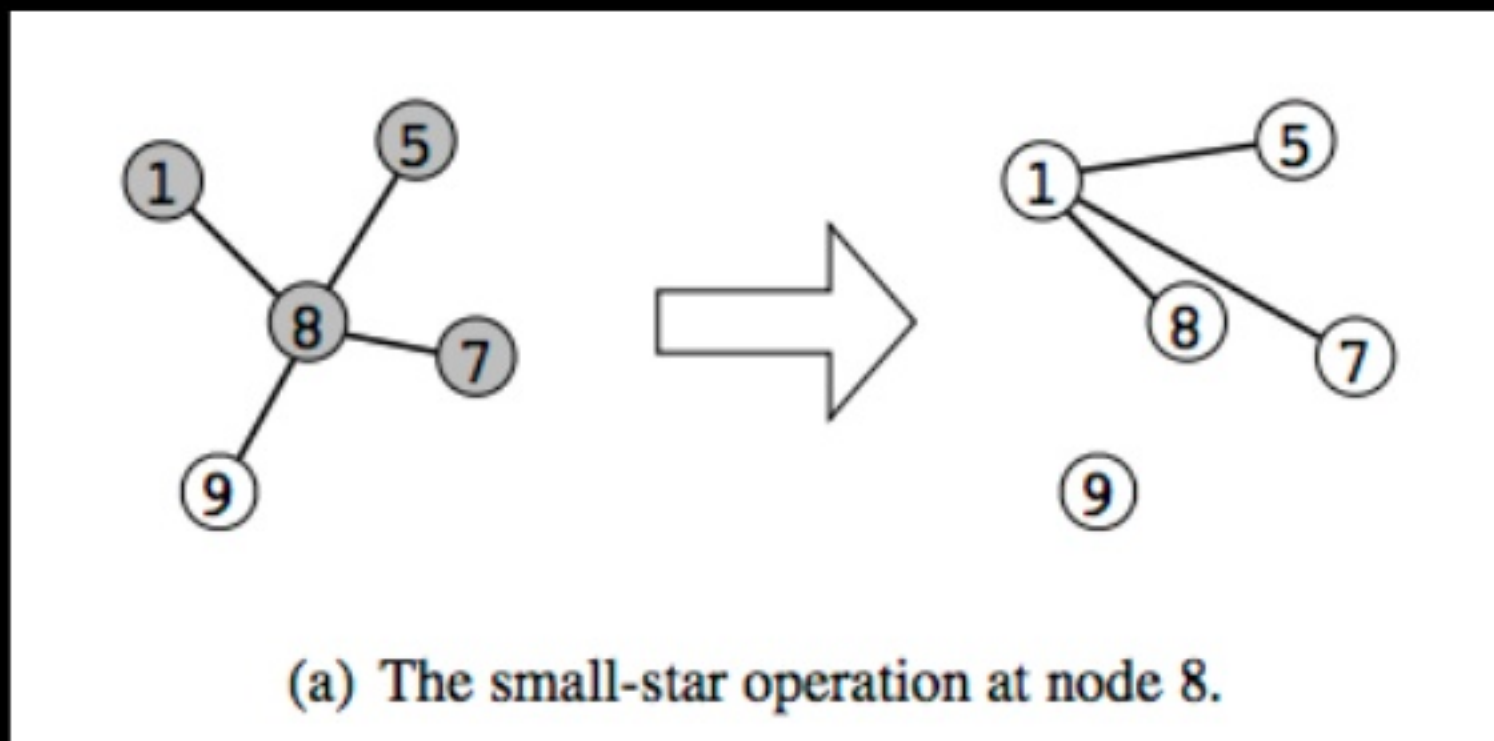
1. Assign each vertex a unique component ID.
 2. Run in batches until convergence:
 - For each vertex v , update its component ID to the smallest component ID among its neighbors' and its own.
- easy to implement
 - slow convergence on large-diameter graphs

Small-/large-star algorithm [Kiveris14]

Kiveris et al., Connected Components in MapReduce and Beyond.

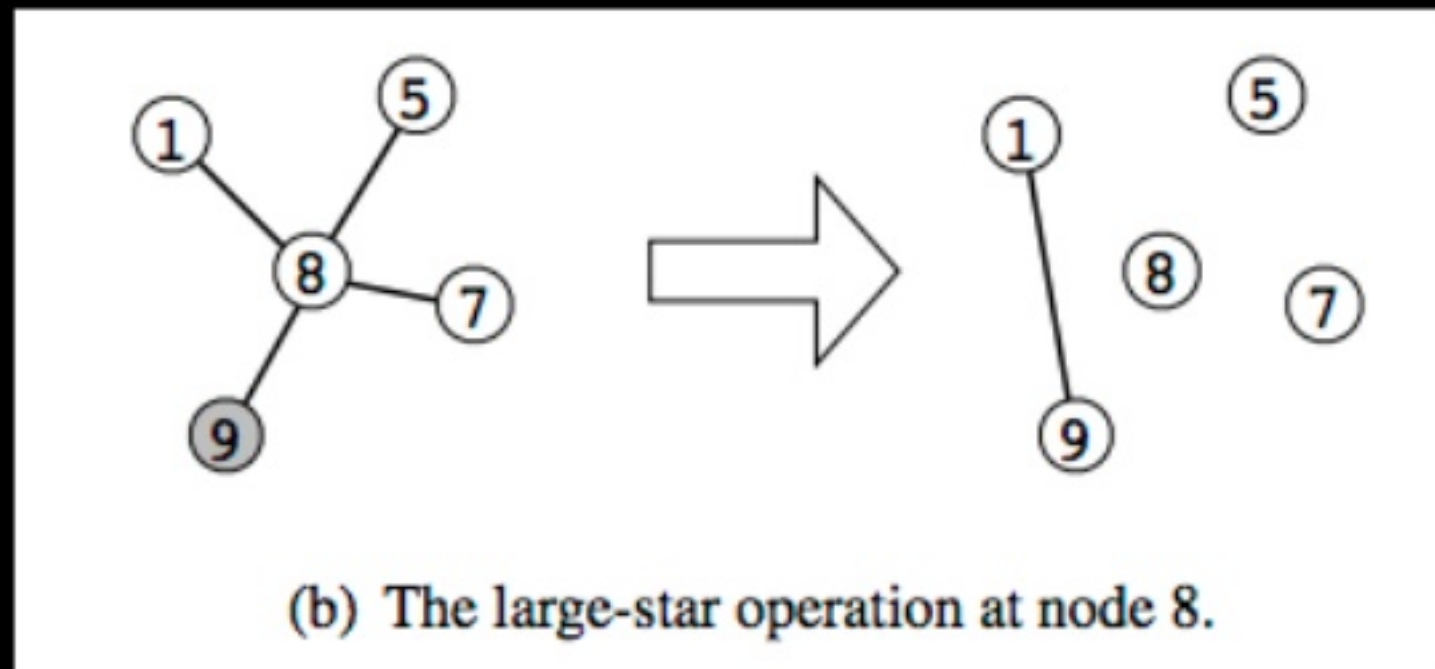
1. Assign each vertex a unique ID.
2. Alternatively update edges in batches until convergence:
 - (small-star) for each vertex, connect its smaller neighbors to the smallest neighbor vertex
 - (big-star) for each vertex, connect its bigger neighbors to the smallest neighbor vertex (or itself)

Small-star operation




Kiveris et al., Connected Components in MapReduce and Beyond.

Big-star operation



Kiveris et al., Connected Components in MapReduce and Beyond.

Another interpretation



adjacency matrix

	1	5	7	8	9
1				x	
5				x	
7				x	
8					x
9					

Small-star operation

	1	5	7	8	9
1				x	
5				x	
7				x	
8					x
9					

rotate & lift

	1	5	7	8	9
1		x	x	x	
5					
7					
8					x
9					

Big-star operation

	1	5	7	8	9
1				x	
5				x	
7				x	
8					x
9					

lift

	1	5	7	8	9
1				x	x
5				x	
7				x	
8					
9					

Convergence

	1	5	7	8	9
1	x	x	x	x	x
5					
7					
8					
9					

Small-/big-star algorithm

- Small-/big-star operations do not change graph connectivity.
- Extra edges are pruned during iterations.
- Each connected component converges to a star graph.

Kiveris et al. proved one variation of the algorithm converges in $\log^2(\text{\#nodes})$ iterations. We chose a variation that alternates small-/big-star operations in GraphFrames.

Implementation

Essentially the small-/big-star operations map to a sequence of filters and self joins with DataFrames. So we need to handle the following operations at scale:

- joins
- iterations

Skewed joins

A real-world graph usually contains big component, which leads to data skewness at connected components iterations.

src	id	nbrs
0	0	2,000,000
1	0	10
2	3	5

join

src	dst
0	1
0	2
0	3
0	4
...	...
0	2,000,000
1	3
2	5

Skewed joins

(#nbrs > 1,000,000)

src	id	nbrs
0	0	2,000,000

broadcast join

src	dst
0	1
0	2
0	3
0	4
...	...
0	2,000,000

union

1	0	10
2	3	5

hash join

1	3
2	5

Checkpointing

We do checkpoint at every 2 iterations to avoid:

- query plan getting too big (exponential growth)
- optimizer taking too long
- disk out of shuffle space
- unexpected node failures

Experiments

- twitter-2010 from [WebGraph datasets](#) (small diameter)
 - 42 million vertices, 1.5 billion edges
- 16 r3.4xlarge workers on Databricks
 - GraphX: 4 minutes
 - GraphFrames: 6 minutes
 - algorithm difference, checkpointing, checking skewness

Experiments

- uk-2007-05 from [WebGraph datasets](#)
 - 105 million vertices, 3.7 billion edges
- 16 r3.4xlarge workers on Databricks
 - GraphX: 25 minutes
 - slow convergence
 - GraphFrames: 4.5 minutes

Experiments

- regular grid 32,000 x 32,000 (large diameter)
 - 1 billion nodes, 4 billion edges
- 32 r3.8xlarge workers on Databricks
 - GraphX: failed
 - GraphFrames: 1 hour

Experiments

- regular grid 50,000 x 50,000 (large diameter)
 - 2.5 billion nodes, 10 billion edges
- 32 r3.8xlarge workers on Databricks
 - GraphX: failed
 - GraphFrames: 1.6 hours

Future improvements

- update inefficient code (due to Spark 1.6 compatibility)
- better graph partitioning
- local iterations
- node pruning and better stop criterion
- letting Spark SQL handle skewed joins and iterations
- graph compression
- prove $\log(N)$ iterations or maybe a better algorithm?



Thank You

- graphframes.github.io
- docs.databricks.com