



CONTINUOUS APPLICATION WITH FAIR SCHEDULER

Robert Xue

Software Engineer, Groupon

About me

Software engineer, Behavioral Modeling team



Groupon Online Defense

- Malicious behavior detection
- A log stream processing engine

Tech stack

- Kafka
- Storm
- Spark on Yarn
- In-house hardware

Glossary

Application

- Created by spark-submit

Job

- A group of tasks
- Unit of work to be submitted

Task

- Unit of work to be scheduled

Scheduler

- Code name: SchedulableBuilder
- [FIFO](#) and [FAIR](#).
- Update [spark.scheduler.mode](#) to switch

Job pool scheduling mode

- Code name SchedulingAlgorithm
- [FIFO](#) and [FAIR](#), applies to FAIR scheduler only
- Update [fairscheduler.xml](#) to decide

What does a scheduler do?

- Determine who can use the resources
- Tweak to optimize
 - ▶ Total execution time
 - ▶ Resource utilization
 - ▶ Total CPU seconds

When does your scheduler work?

- Only when you have more than one job submitted to spark context
- Writing your rdd operations line by line != multiple jobs submitted at the same time


```
Range(0, 24).foreach { i =>  
  sc.textFile(s"file-$i").count()  
}
```


Run - 1

Sequentially submitting 24 file reading jobs

```
Range(0, 24).foreach { i =>
  sc.textFile(s"file-$i").count()
}
```

- 4 parts, ~500MB files * 24
- 8 cores
- 110 seconds
- $8 * 110 * 11.33\% = 99.68$ cpu seconds
- ~20 individual jobs clearly visible



Node_local - Rack_local - Any - Unused

Submit your job in parallel

- Java Runnable
- Python threading
- Scala Future
 - Scalaz Task
 - Monix Task

```
Range(0, 24).foreach { i =>  
  sc.textFile(s"file-$i").count()  
}
```





```
Range(0, 24).par.foreach { i =>  
  sc.textFile(s"file-$i").count()  
}
```

Run - 2

Parallel submit & FIFO(24)

```
Range(0, 24).par.foreach { i =>
  sc.textFile(s"file-$i").count()
}
```

- 8 cores
- ~~110~~22 seconds
- $8 * 22 * 82.04\% =$
~~99.68~~144.32 cpu
seconds
- ~~0~~15 seconds 100%
utilization period



Node_local - Rack_local - Any - Unused

Turn on FAIR scheduler

Provide `--conf spark.scheduler.mode=FAIR` in your spark submit...




... it is that simple

Run - 3

Parallel submit & FAIR(24)

```
Range(0, 24).par.foreach { i =>
  sc.textFile(s"file-$i").count()
}
```

- 8 cores
- 22 **15** seconds
- $8 * 15 * 73.74\% =$
~~144.32~~**88.48**^{record} cpu
seconds
- ~~Bad~~**Great** locality



Node_local - Rack_local - Any - Unused

Tweak your locality config

Provide `--conf spark.locality.wait=1s` in your spark submit...

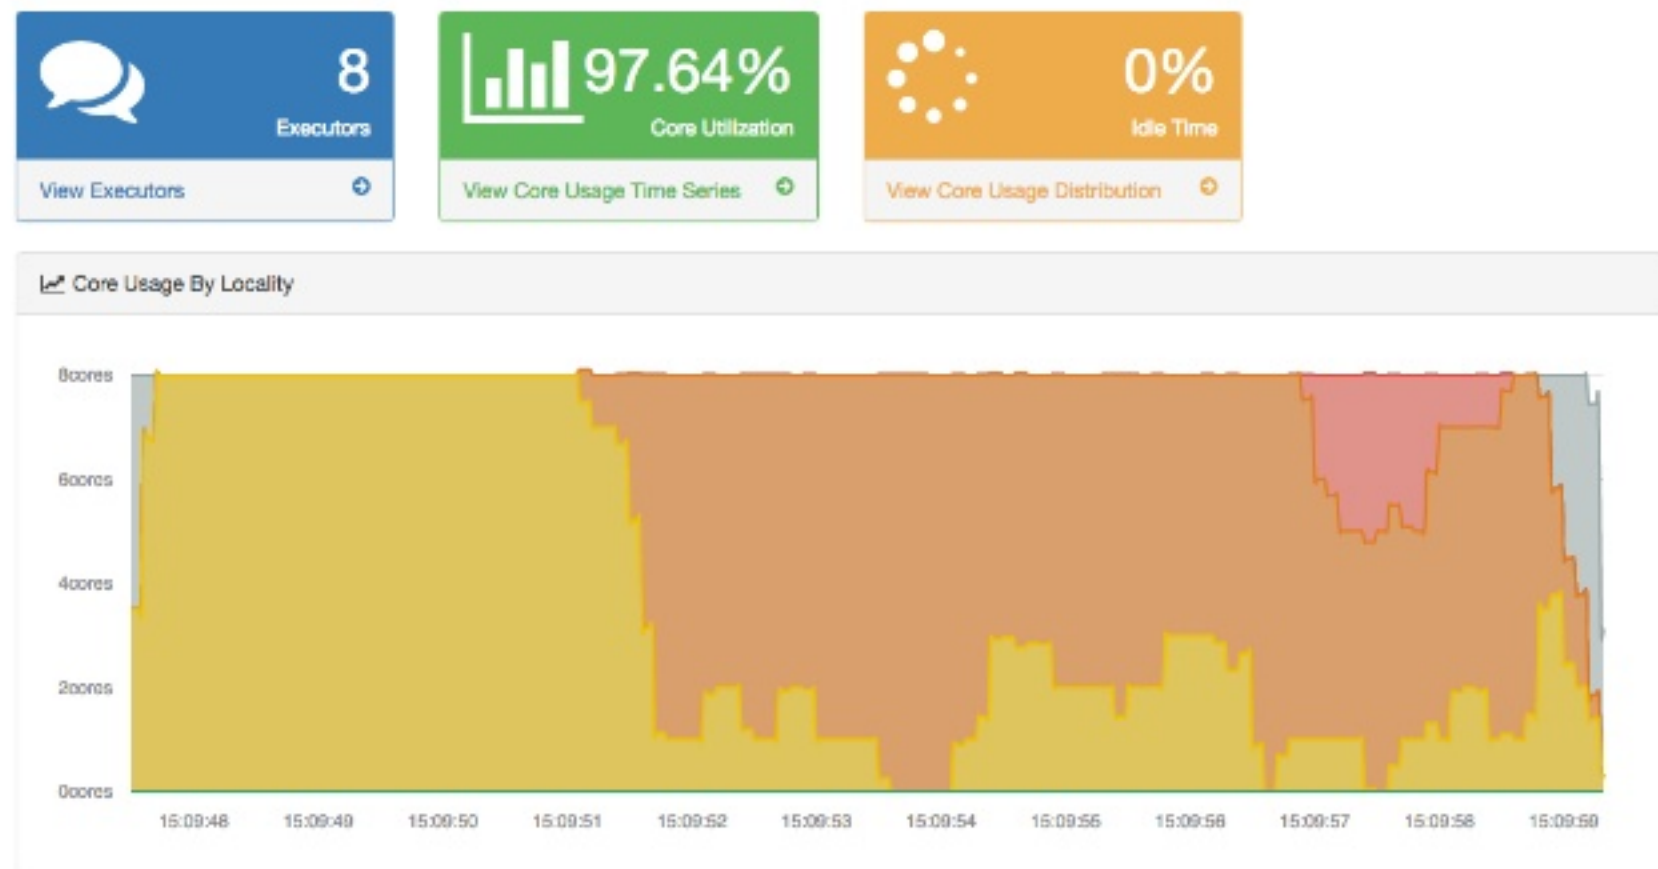
- `spark.locality.wait.process/node/rack` available as well
- default is 3s

Run - 4

Parallel submit & FAIR(24) & locality.wait=1s

```
Range(0, 24).par.foreach { i =>
  sc.textFile(s"file-$i").count()
}
```

- 8 cores
- ~~15~~ **12^{record}** seconds
- $8 * 12 * 97.64\% =$
~~88.48~~ **93.68** cpu
seconds




Node_local - Rack_local - Any - Unused

Summary

	Cluster utilization	Execution time	Locality tweaking	Miscellaneous
Sequential	Poor	Long	Hard	Default behavior in Spark-shell
Parallel {FIFO}	Good	Short	Hard	Default behavior in Notebooks
Parallel {FAIR}	Good	Short	Easy	

Standalone applications

Stream, Batch, Adhoc query, submitted from multiple apps



Continuous application


Stream, Batch, Adhoc query, parallel submitted in one app

Read data ■

Model
Execution ■

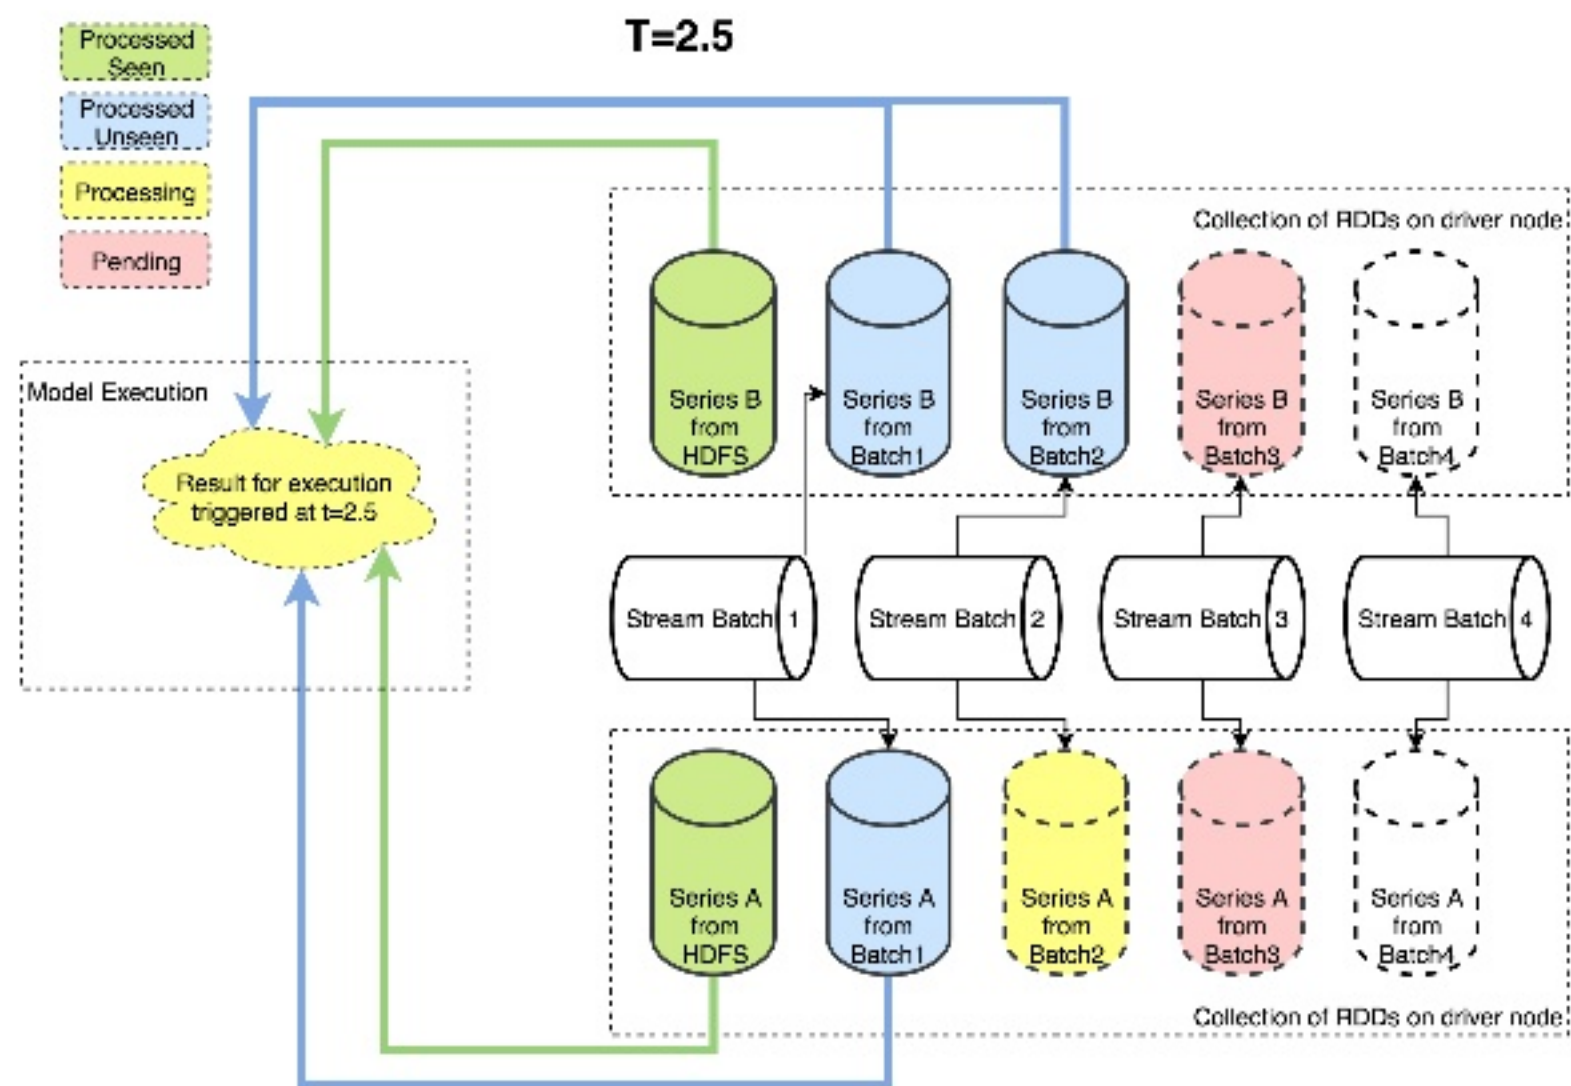
Stream
Process ■

Reporting ■

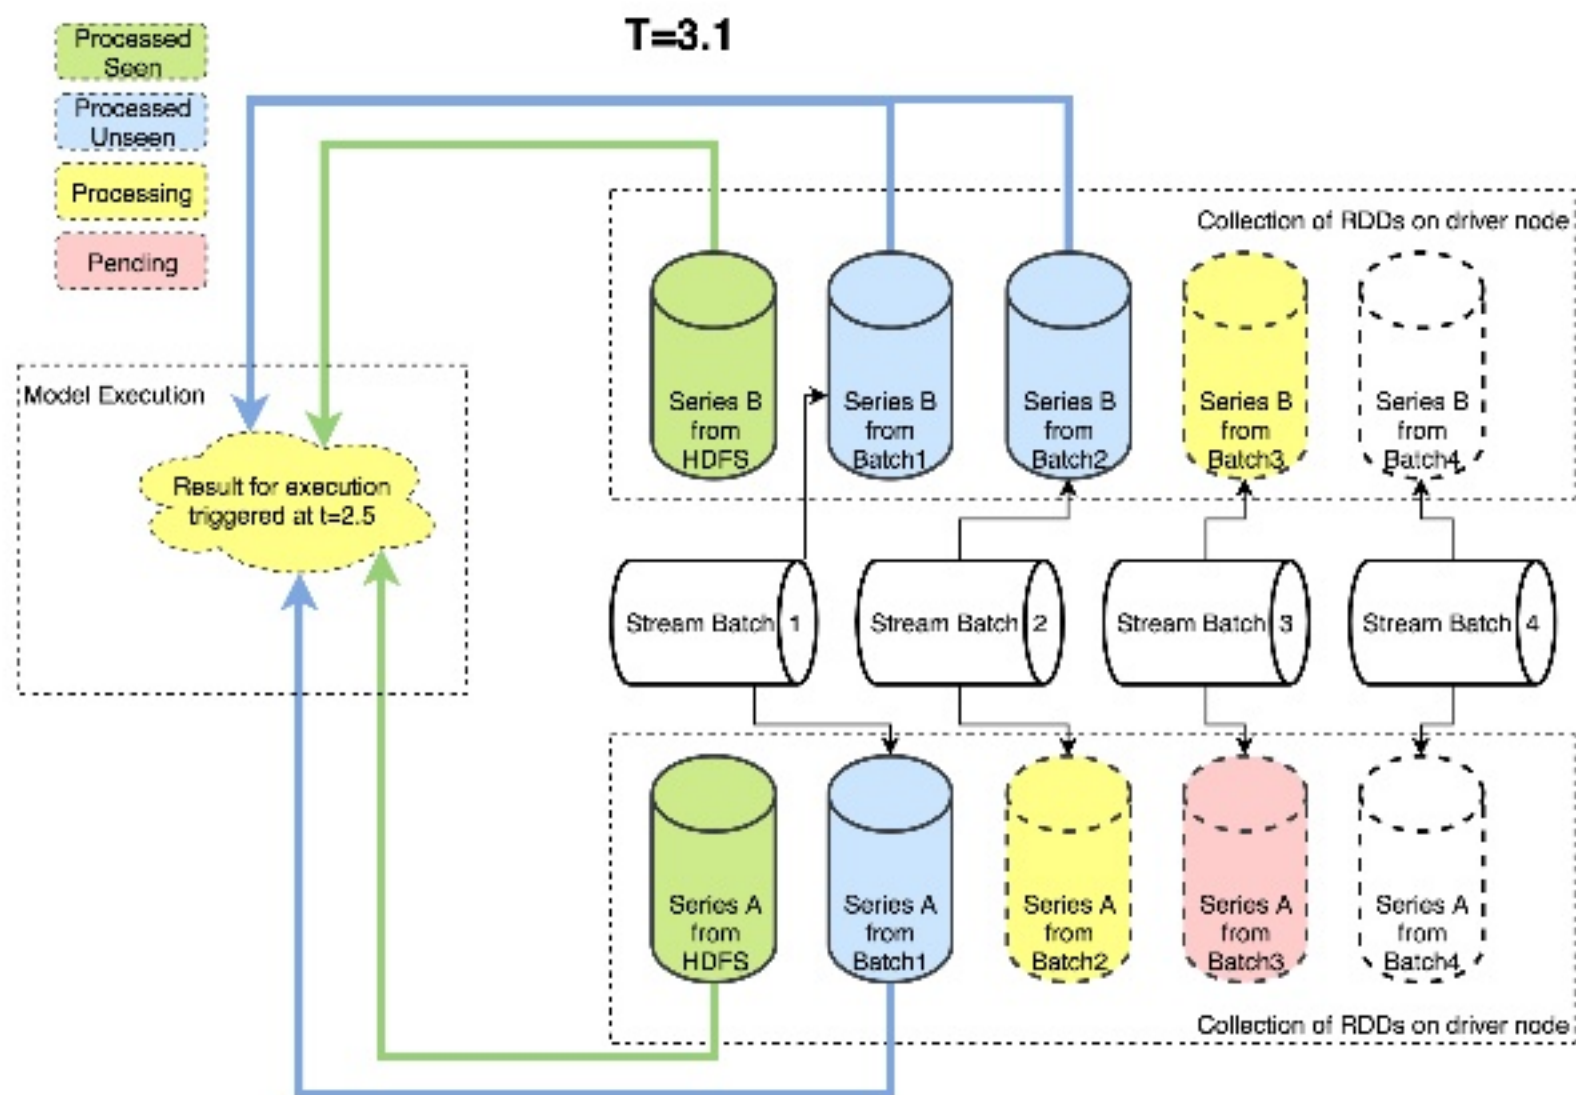


Continuous Application

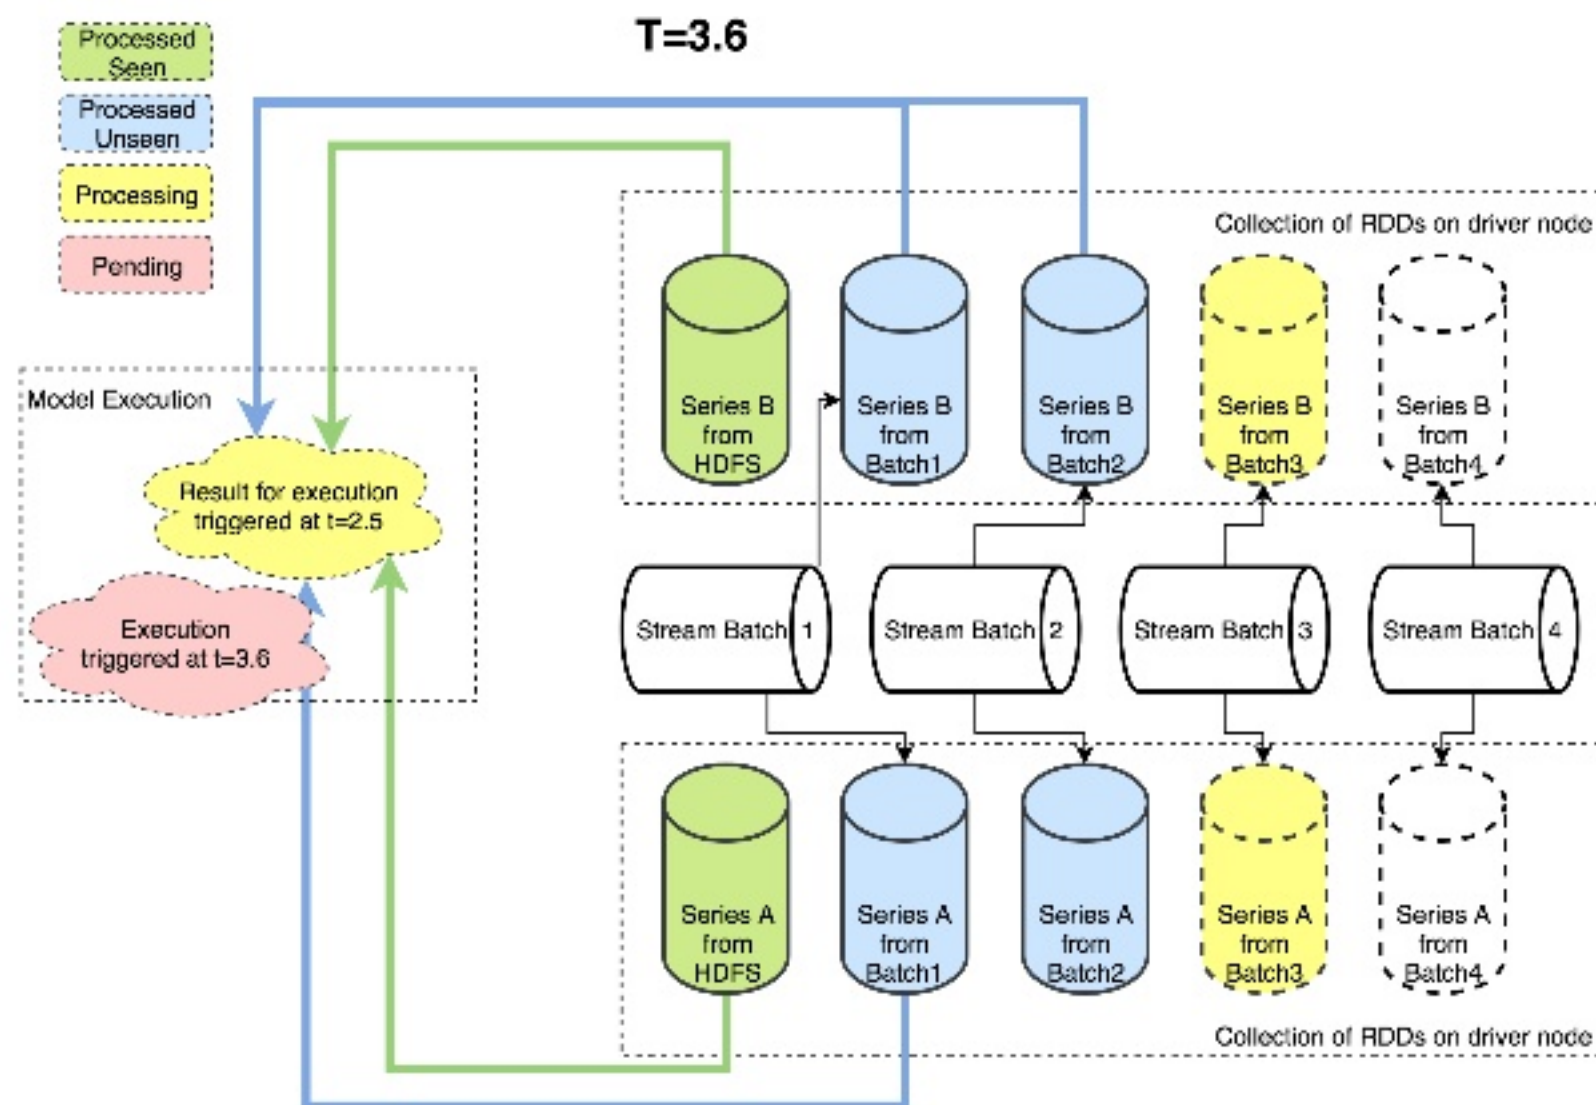
	Multiple standalone applications	Continuous application
Sharing context	Share file content using database, file system, network, etc	Simply pass RDD reference around
Resource allocation	Static; configure more cores than typical load to handle peak traffic	Dynamic; less important tasks yield CPU critical tasks
Job scheduling	Crontab, oozie, airflow... — all approaches leads to spark-submit, a typically 20s - 120s overhead	Spark-submit once and only once



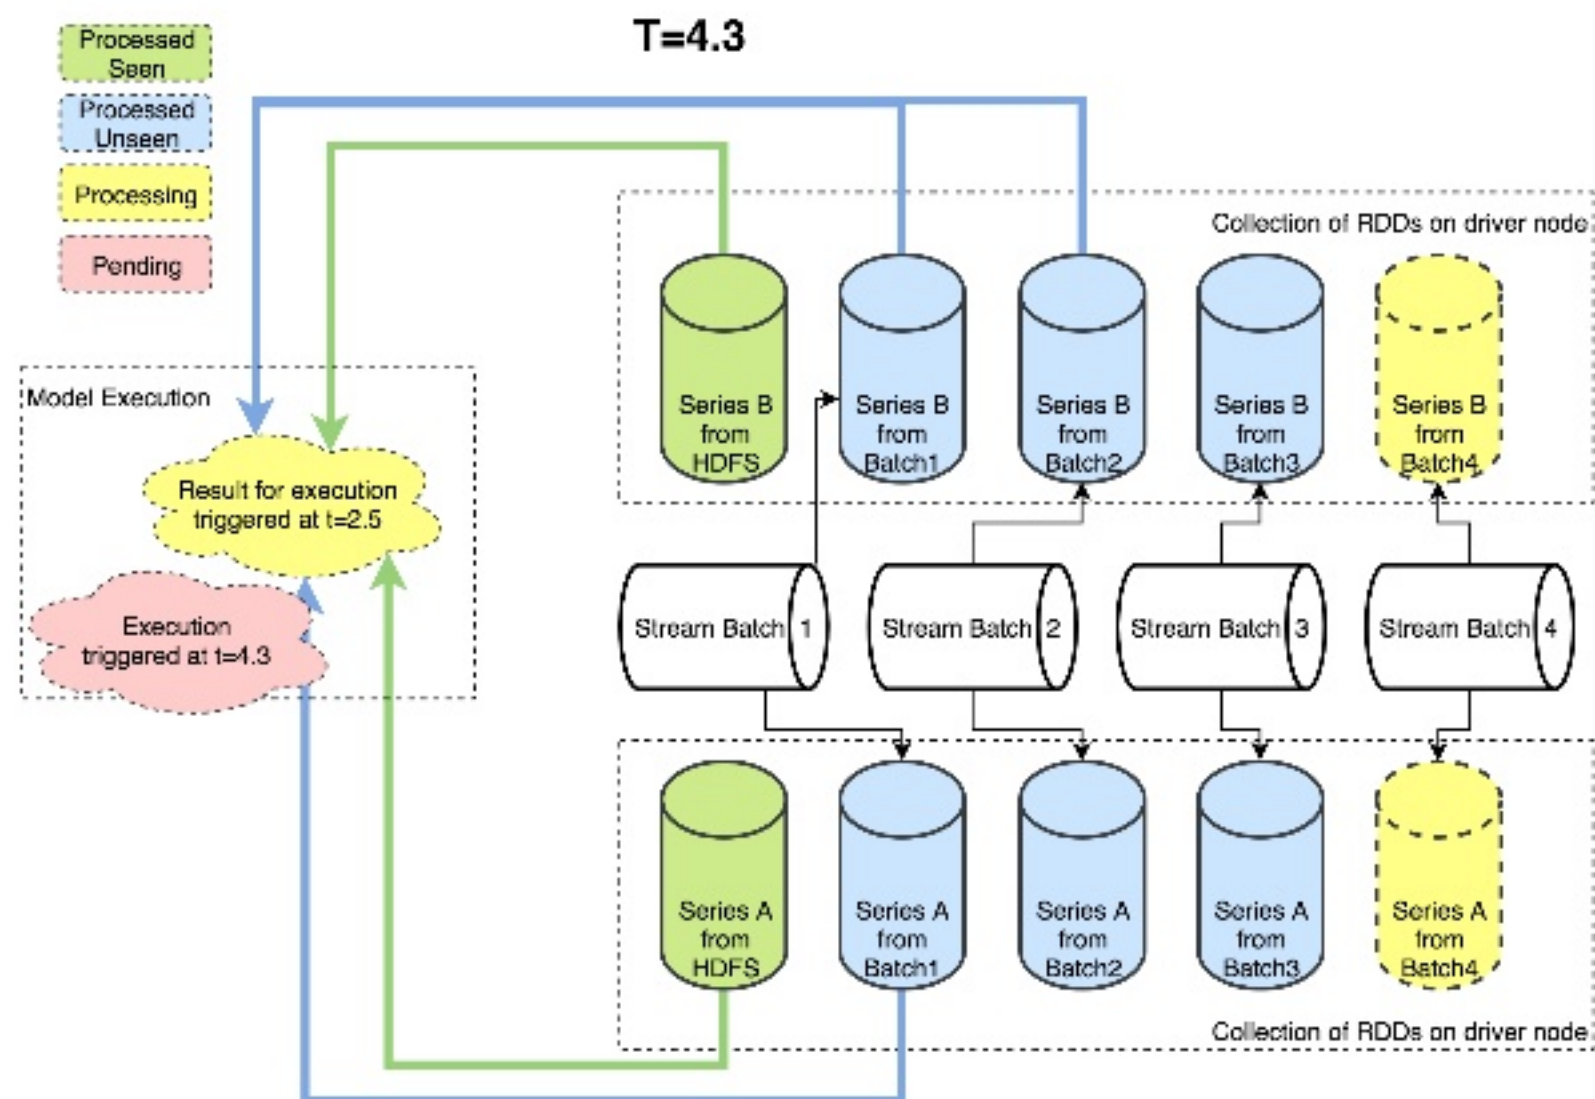
At $t=2.5$, Series A and B finished processing stream input at $t=1.0$
 A round of model execution had been triggered at $t=2.5$
 Model was triggered using all available data at that moment



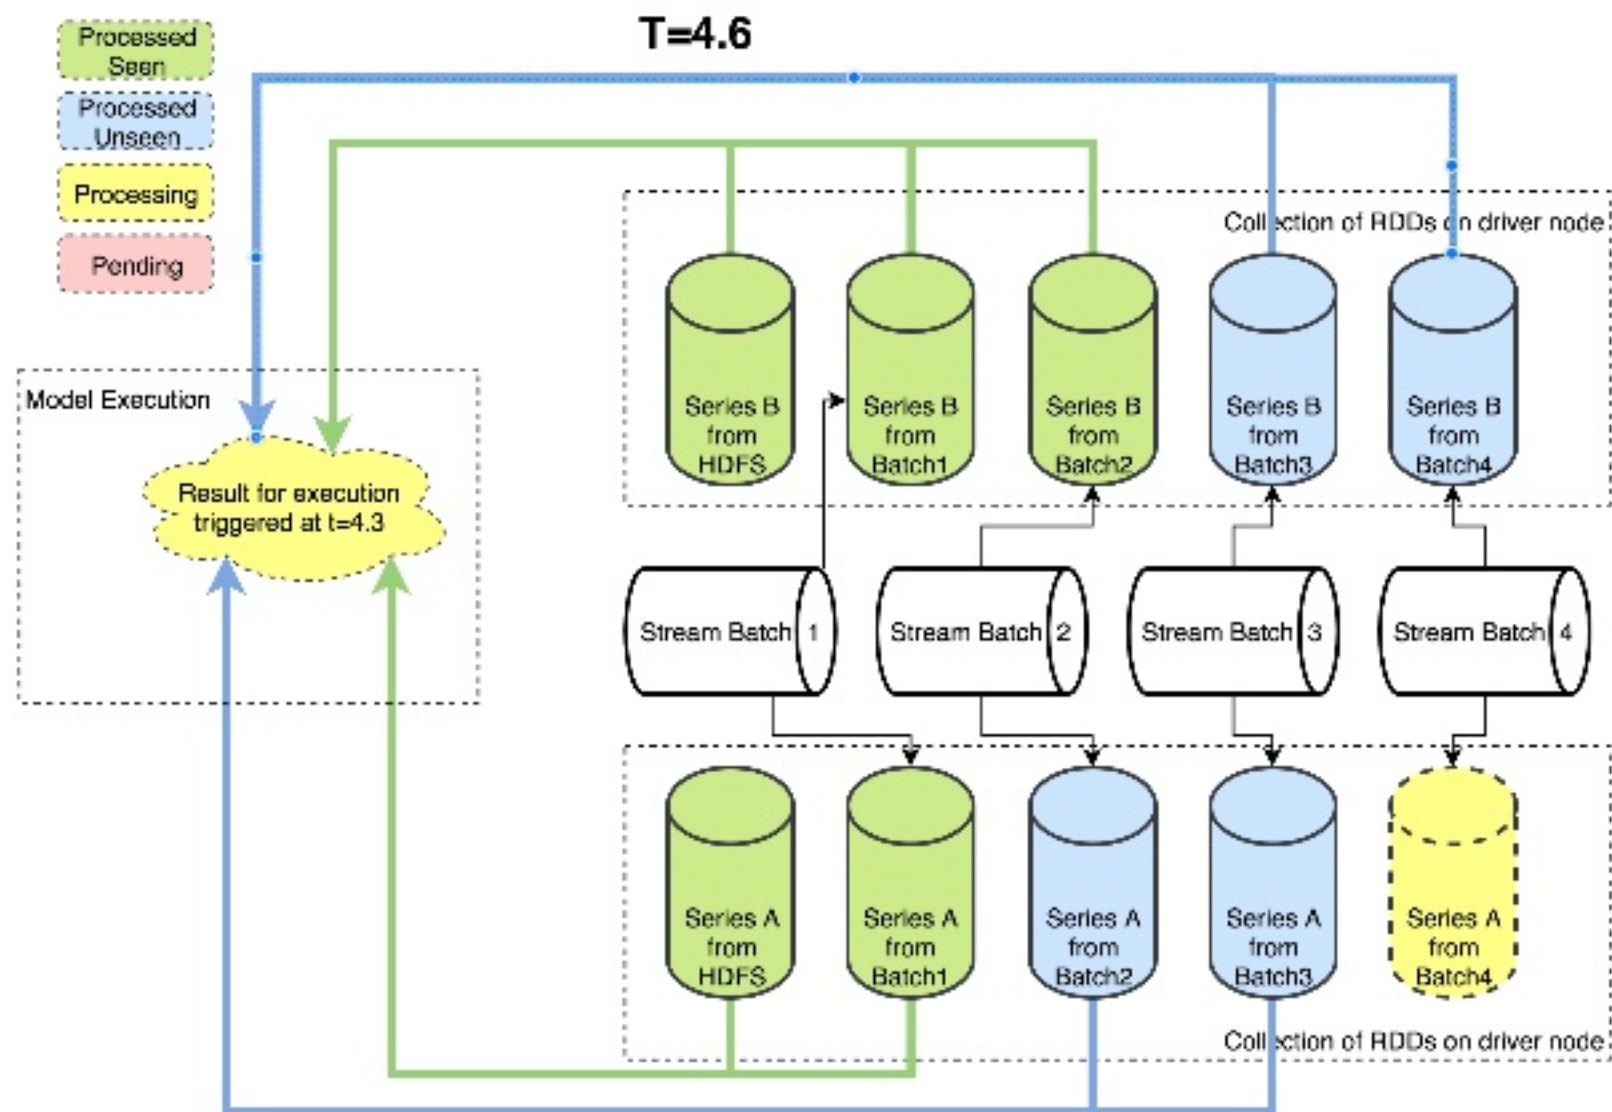
At $t=3.1$, app received stream input at $t=3.0$
 Series A was still processing input at $t=2.0$
 Series B started processing the new input.



At t=3.6, SeriesA finished processing stream input at t=2.0
 SeriesA started processing stream input at t=3.0
 Model had a pending execution triggered at t=3.6



At $t=4.3$, Series A and B finished processing stream input at $t=3.0$
 Model had a pending execution triggered at $t=4.3$
 Model's execution at $t=3.6$ was cancelled



At $t=4.6$, model finished its execution at $t=2.5$
 Series B finished processing stream input at $t=4.0$
 Model started execution that was triggered at $t=4.3$
 Model used data that was available at $t=4.6$ as input

Example:

4 models executing in parallel from a 16 cores app

6 Fair Scheduler Pools

Pool Name	Minimum Share	Pool Weight	Active Stages	Running Tasks	SchedulingMode
default	4	1	4	2	FAIR
Model	0	2	1	1	FAIR
cache	0	1	0	0	FIFO
reporting	0	1	0	8	FIFO
spotCheck	0	1	1	5	FIFO

Active Stages (6)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
177202	default	merging Aggregator buffer count at Utils.scala:129 +details (kill)	2017/06/06 00:57:21	Unknown	0/8				
177199	default	merging Aggregator buffer count at Utils.scala:129 +details (kill)	2017/06/06 00:57:21	0.4 s	<div><div>5/8</div></div>	19.8 KB		12.6 KB	
177193	default	flatMap at Aggregator.scala:73 +details (kill)	2017/06/06 00:57:20	2 s	<div><div>24/25</div></div>				3.0 KB
177189	Model	wrap violators in ActorData, without whitelist info map at AggregatorNamedResourcesSettings.scala:16 +details (kill)	2017/06/06 00:57:19	2 s	<div><div>0/1</div></div>				
177108	default	merging Aggregator buffer count at Utils.scala:129 +details (kill)	2017/06/06 00:57:16	4 s	<div><div>6/8</div></div>	33.3 MB			
177096	spotCheck	perform a spot check for 74.88.68.136 map at AggregatorNamedResourcesSettings.scala:16 +details (kill)	2017/06/06 00:57:12	8 s	<div><div>3/8</div></div>	541.9 MB			41.3 MB

Practices for Continuous Application

Decouple stream processing from batch processing

- Batch interval is merely your minimum resolution

Emphasis tuning for streaming part

- Assign to a scheduler pool with minimum core guarantee

Execute only the latest batch invocation

- Assign to a scheduler pool with high weight

Reporting and query onsite

- Assign to a scheduler pool with low weight

Summary

Coding

- Share data by passing RDD/DStream/DF/DS
- Always launch jobs in a separate thread
- No complex logic in the streaming operation
- Push batch job invocation into a queue
- Execute only the latest batch job

App submission

- Turn on FAIR scheduler mode
- Provide fairscheduler.xml
- Turn off stream back pressure for Streaming app
- Turn off dynamic allocation for Streaming app
- Turn on dynamic allocation for long-lived batch app

Job bootstrapping

- `sc.setJobGroup("group", "description")`
- `sc.setLocalProperty("spark.scheduler.pool", "pool")`
- `rdd.setName("my data at t=0").persist()`

Packaging

- Multiple logic, one repo, one jar
- Batch app can be long-lived as well
- Replace crontab with continuous app + REST

Tuning

- Understand your job's opportunity cost
- Tweak `spark.locality.wait` parameters
- Config cores that yields acceptable SLA with good resource utilization



Thank You.

Robxue Xue, Software Engineer, Groupon
rxue@groupon.com | [@roboxue](https://twitter.com/roboxue)

Tool used in this deck: [groupon/sparklint](https://github.com/groupon/sparklint)

Open sourced on github since [Spark Summit EU 2016](#)

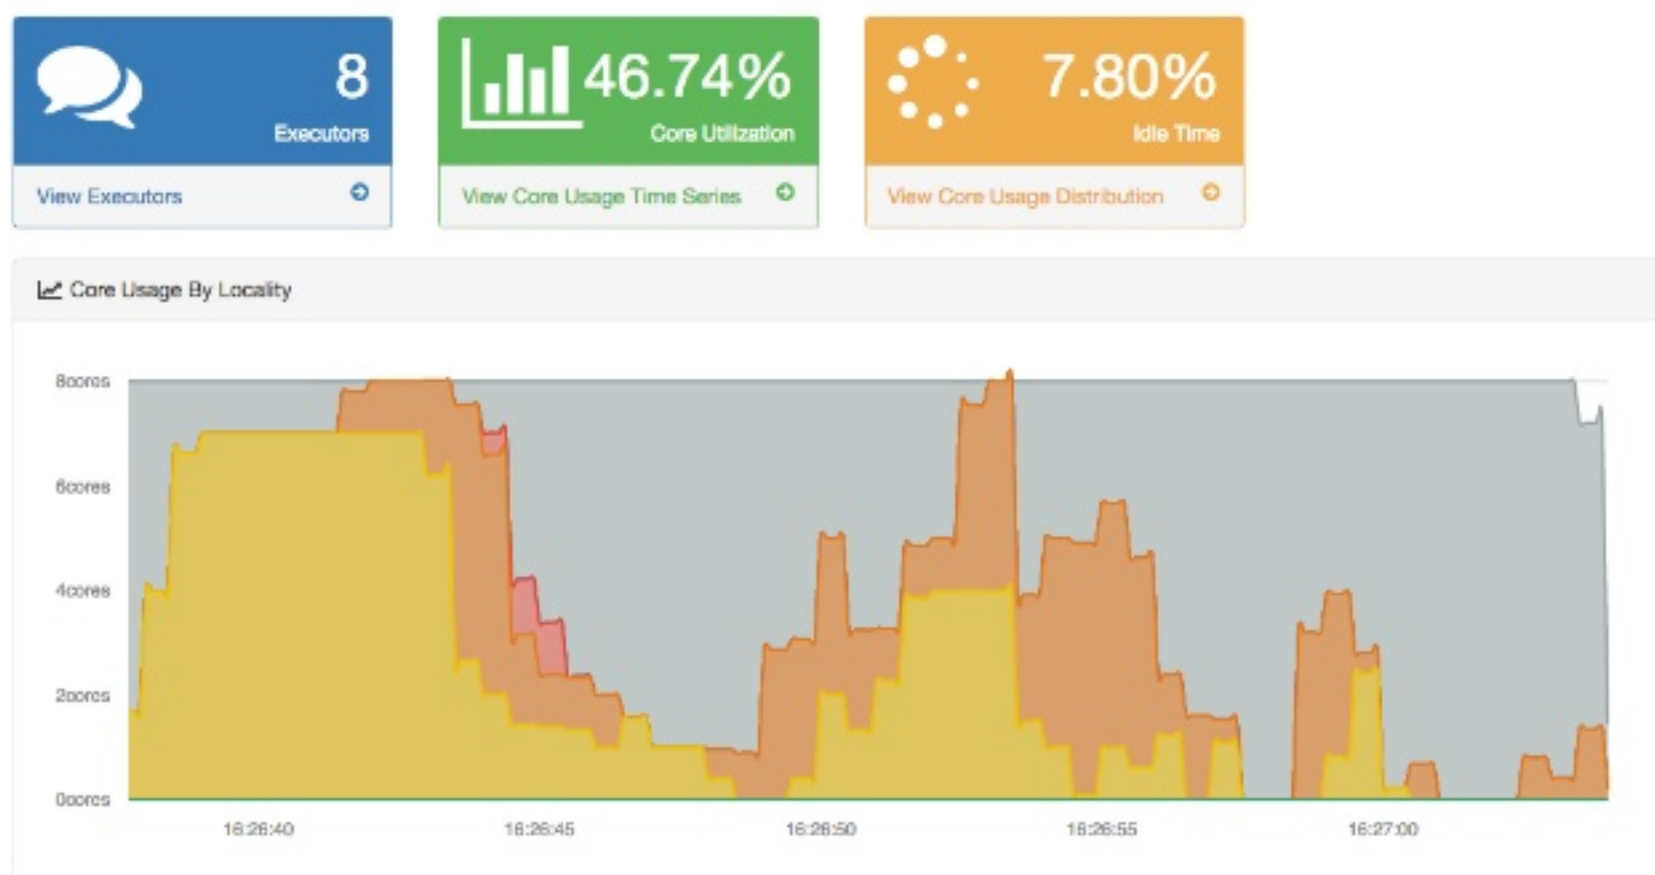
Q & A

Why do you need a scheduler - 3

Parallelly submitting 24 file reading jobs - FIFO(4)

- 8 cores
- ~~22~~27 seconds
- $8 * 27 * 46.74\% =$
~~99.68~~100.88 cpu
seconds
- ~~Bad~~Improved locality

```
implicit val ec =  
  ExecutionContext.fromExecutor(  
    new ForkJoinPool(4))  
Range(0, 24).foreach(i=> Future{  
  sc.textFile(s"file-$i").count()  
})
```





Node_local - Rack_local - Any - Unused

Why do you need a scheduler - 4

Parallelly submitting 24 file reading jobs - FAIR(4)

- 8 cores
- ~~22~~ 32 seconds
- $8 * 32 * 40.92\% = 104.72$ cpu seconds
- ~~Bad~~ Great locality


```
// --conf spark.scheduler.mode=FAIR
implicit val ec =
  ExecutionContext.fromExecutor(
    new ForkJoinPool(4))
Range(0, 24).foreach(i=> Future{
  sc.textFile(s"file-$i").count()
})
```




Node_local - Rack_local - Any - Unused

Back up: locality settings matters

spark.scheduler.mode=FAIR
spark.locality.wait=500ms

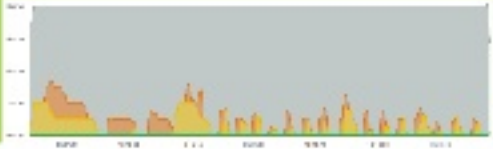








spark.scheduler.mode=FAIR
spark.locality.wait=20s



Node_local - Rack_local - Any - Unused

Backup: Scheduler summary

	Sequential job submit	Parallel job submit Under Parallelized	Parallel job submit Perfect Parallelized	Parallel job submit Over Parallelized / Poor locality settings
				
FIFO scheduler	Under-utilized cluster Good locality Low core usage Long execution time	Under-utilized cluster Poor locality High core usage Short execution time	Well-utilized cluster Poor locality High core usage Short execution time	Well-utilized cluster Poor locality High core usage Long execution time
FAIR scheduler	N/A	Under-utilized cluster Good locality Low core usage Short execution time	Well-utilized cluster Good locality Low core usage Short execution time	
				

Node_local - Rack_local - Any - Unused

Backup:

What's wrong with Dynamic Allocation in streaming

... if you are analyzing time series data / “tensor” ...

- Streaming app, always up. Workloads comes periodically and sequentially.
- Core usage graph has a saw-tooth pattern
 - Less likely to return executors, if your batch interval < executorIdleTimeout
- Dynamic allocation is off if executor has cache on it
 - “by default executors containing cached data are never removed”
- Dynamic allocation == “Resource blackhole” ++ “poor utilization”