



# Dr. Elephant for Monitoring & Tuning Apache Spark Jobs on Hadoop

Carl Steinbach (LinkedIn)  
Simon King (Pepperdata)

# Hadoop @ LinkedIn c. 2015

- > 10 clusters
- > 10,000 nodes
- > 1,000 users
- Thousands of queries and flows in development
- Spark, Pig, Hive, Scalding, Gobblin, Cubert, ...

# What we learned along the way

Scaling Hadoop Infrastructure is Hard  
Scaling User Productivity is Harder






# Tuning Hadoop and Spark

# Some things we tried

- Training
  - doesn't scale
  - interferes with productivity
- Expert Review
  - doesn't scale
  - long wait times



# Birth of Dr. Elephant!



# What does Dr. Elephant do?

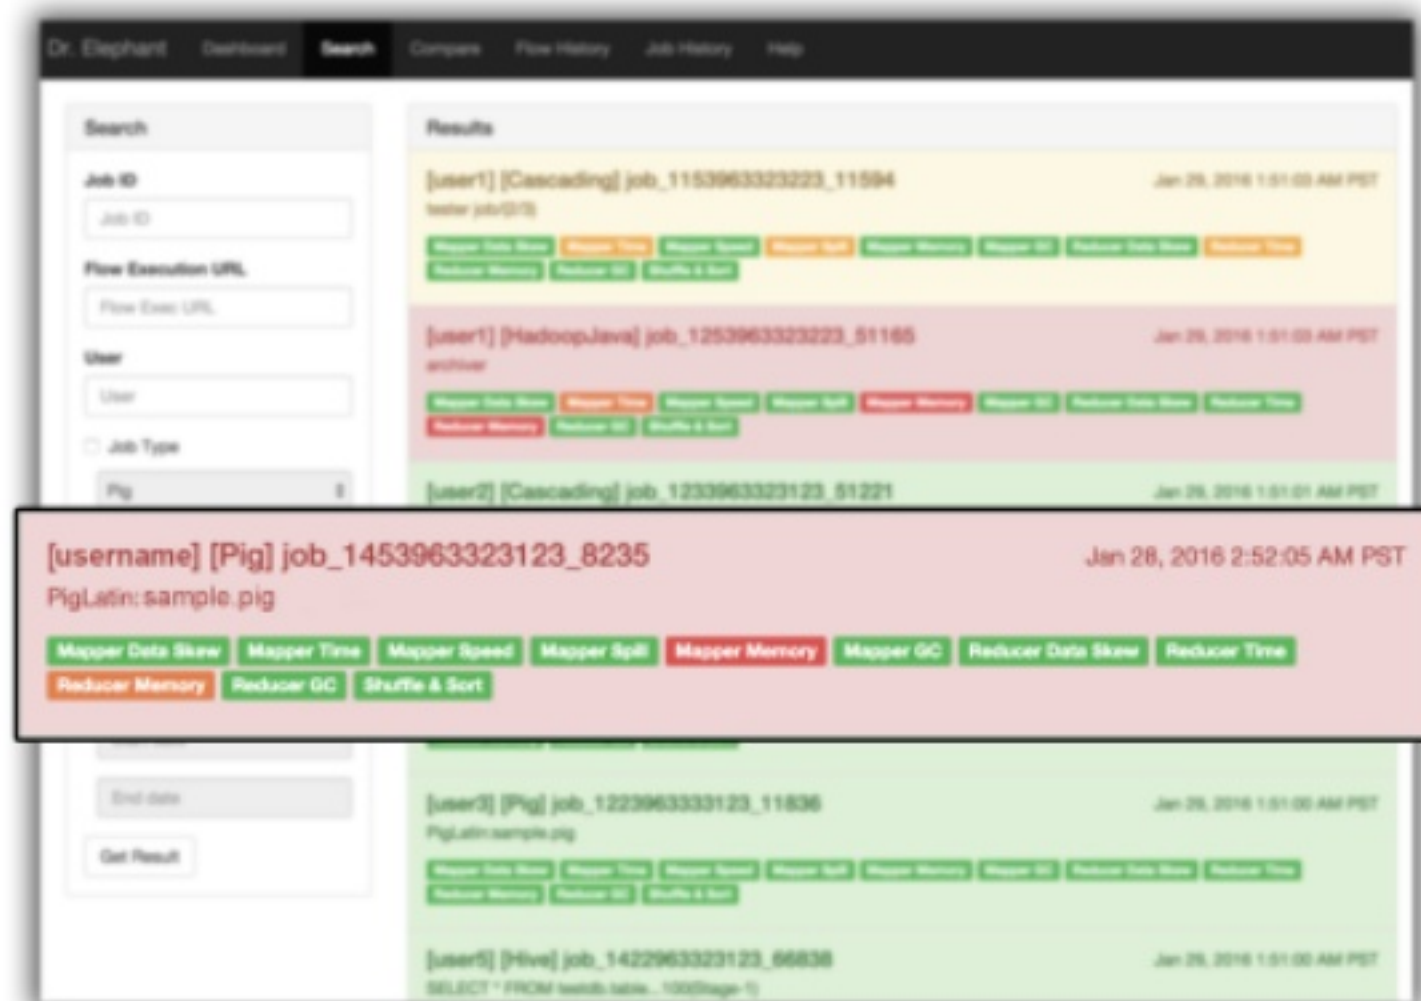
- Performance monitoring and tuning service
- Finds common mistakes
- Provide actionable advice
- Compare performance changes over time

# Architecture

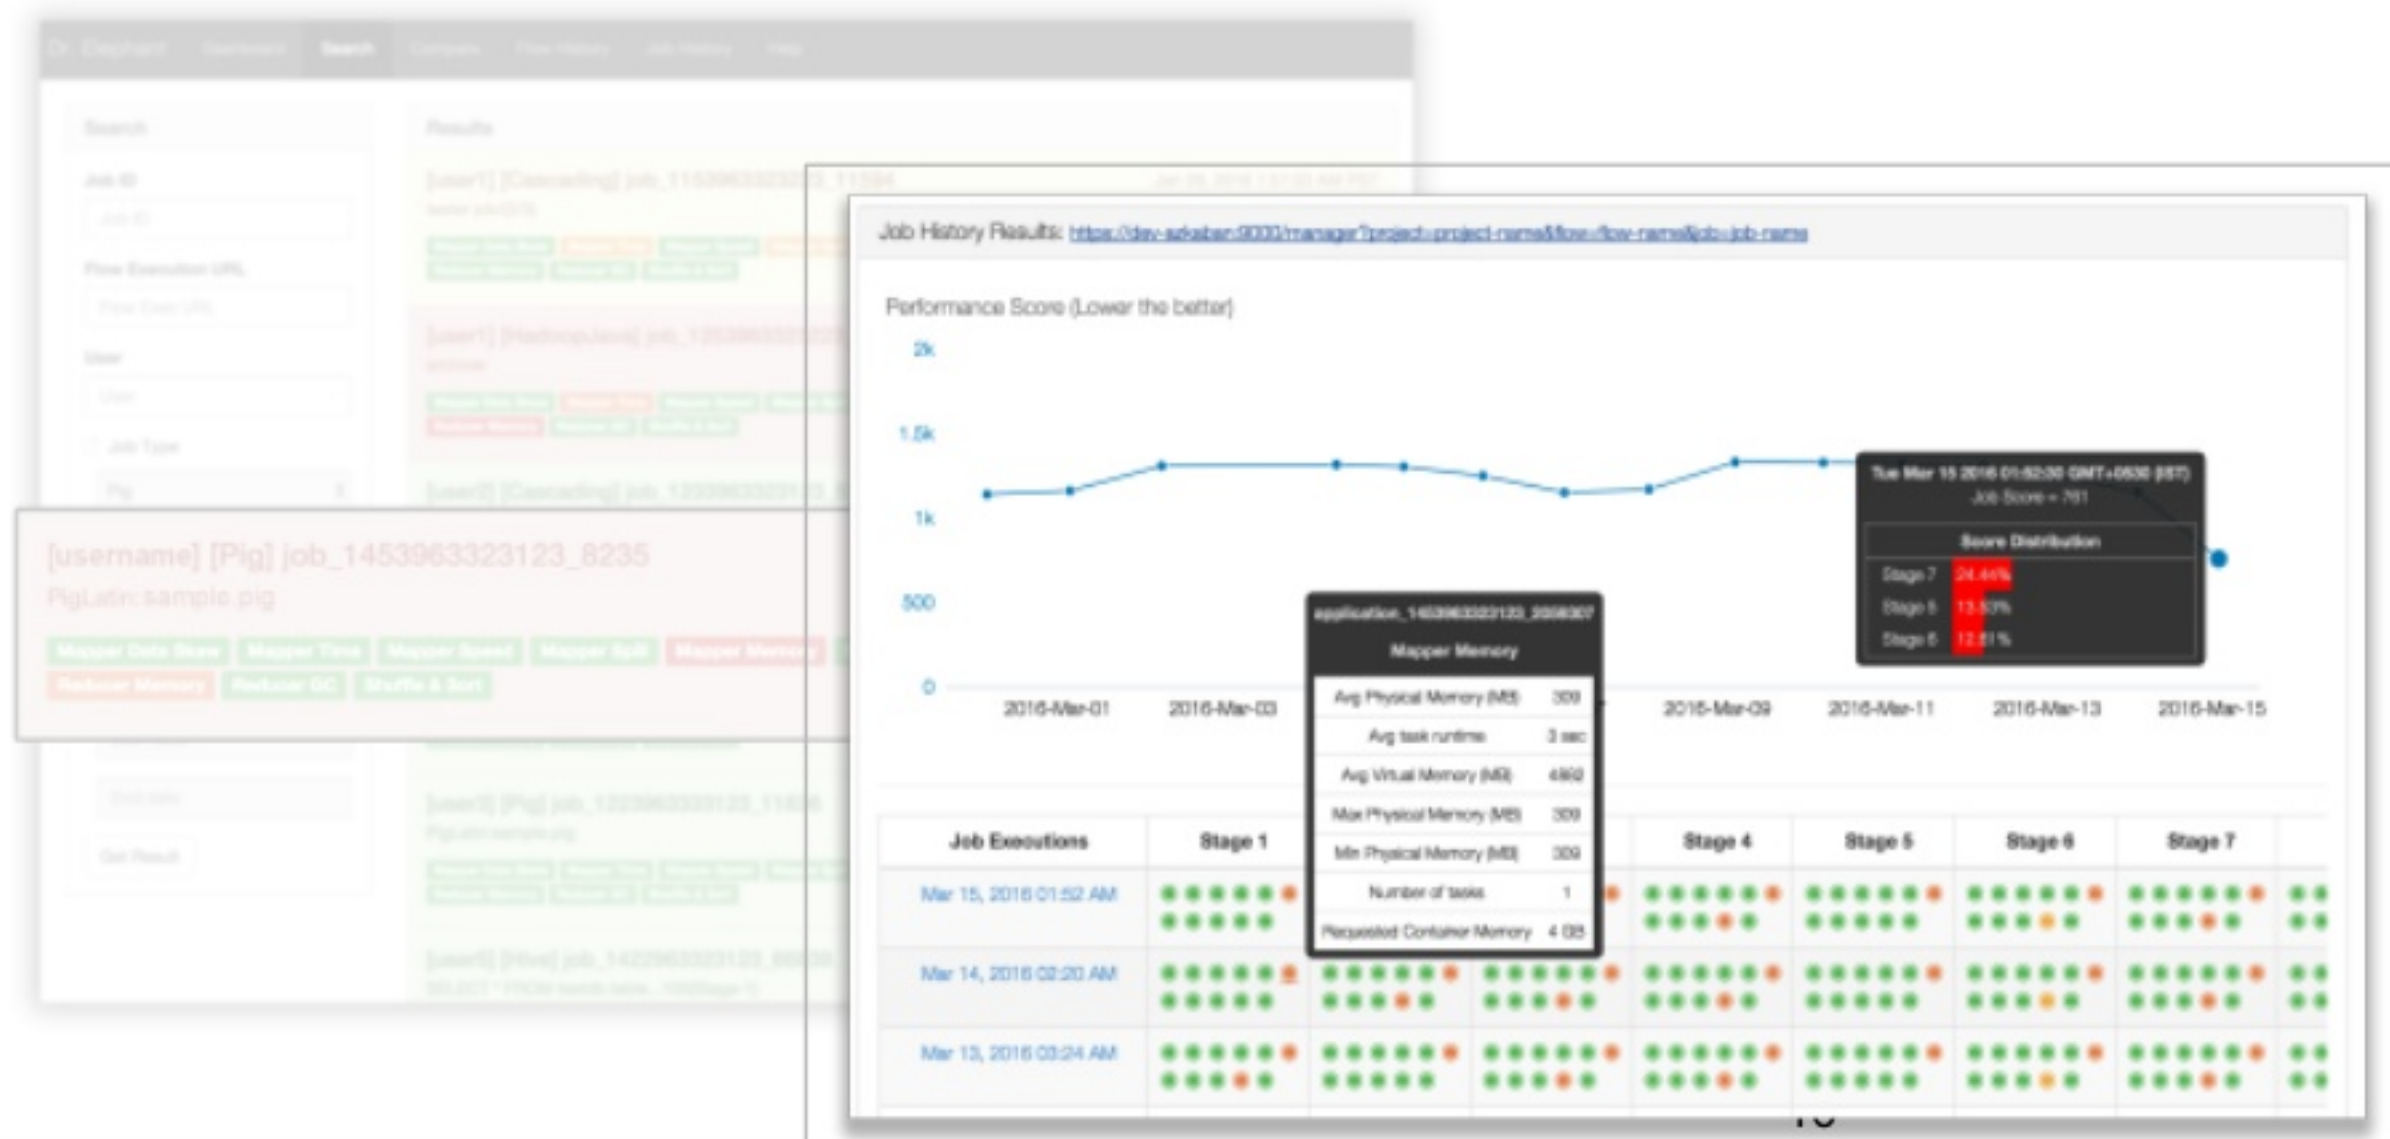




# Dr. Elephant User Interface



# Dr. Elephant User Interface



# Dr. Elephant Community

LinkedIn



hulu



Flipkart



FOURSQUARE



Hops

inMOBI™

kakao

# Outline


- Spark Event Logs and Spark History Server
- Dr. Elephant for Spark
- Pepperdata's Application Profiler



[simon@pepperdata.com](mailto:simon@pepperdata.com)



# Spark History Server

<div> <b>History Server</b></div> <div>Event log directory: <code>hdfs:///user/spark/applicationHistory</code></div> <div>Showing 1-20 of 1925 <span>1 2 3 ... 97 &gt;</span></div>						
App ID	App Name	Started	Completed	Duration	Spark User	Last Updated
<a href="#">application_1495423352820_0363</a>	PepperdataLogAnalyzer	2017/05/23 12:28:33	2017/05/23 12:36:53	8.3 min	root	2017/05/23 12:36:57
<a href="#">application_1495423352820_0361</a>	LogisticRegressionDataGenerator	2017/05/23 12:14:24	2017/05/23 12:18:34	4.2 min	prod	2017/05/23 12:18:36
<a href="#">application_1495423352820_0360</a>	ScalaTeraSort	2017/05/23 11:52:22	2017/05/23 12:13:56	22 min	prod	2017/05/23 12:13:56
<a href="#">application_1495423352820_0358</a>	LogisticRegressionWithLBFGS	2017/05/23 11:43:13	2017/05/23 11:54:56	12 min	prod	2017/05/23 11:54:57
<a href="#">application_1495423352820_0357</a>	LogisticRegressionDataGenerator	2017/05/23 11:41:20	2017/05/23 11:43:00	1.7 min	prod	2017/05/23 11:43:02
<a href="#">application_1495423352820_0356</a>	ScalaTeraSort	2017/05/23 11:38:55	2017/05/23 11:41:04	2.2 min	prod	2017/05/23 11:41:04
<a href="#">application_1495423352820_0354</a>	LogisticRegressionWithLBFGS	2017/05/23 11:29:09	2017/05/23 11:37:34	8.4 min	prod	2017/05/23 11:37:34

# Spark History Server

## Summary Metrics for 8 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	7 s	8 s	9 s	10 s	10 s
GC Time	43 ms	78 ms	0.2 s	0.3 s	1 s
Result Serialization Time	1 ms	1 ms	1 ms	2 ms	2 ms
Peak Execution Memory	64.0 MB	64.0 MB	64.0 MB	64.0 MB	64.0 MB
Shuffle Read Blocked Time	0 ms	0 ms	0 ms	1 ms	1 ms
Shuffle Read Size / Records	27.7 MB / 1247593	27.7 MB / 1249493	27.7 MB / 1250163	27.7 MB / 1251184	27.8 MB / 1252117

# Spark Event Logs

# Spark Event Logs

```
{
  "Event": "SparkListenerTaskEnd",
  "Stage ID": 9,
  "Stage Attempt ID": 0,
  "Task Type": "ResultTask",
  "Task End Reason": {
    "Reason": "Success"
  },
  "Task Info": {
    "Task ID": 775,
    "Index": 55,
    "Attempt": 0,
    "Launch Time": 1495496382885,
    "Executor ID": "9",
    "Host": "amarillo-rm.pepperdata.com",
    "Locality": "PROCESS_LOCAL",
    "Speculative": false,
    "Getting Result Time": 0,
    "Finish Time": 1495496481595,
    "Failed": false,
    "Accumulables": {
      {
        "ID": 7,
        "Name": "peakExecutionMemory",
        "Update": "76154696",
        "Value": "601560113",
        "Internal": true
      }
    },
    "Task Metrics": {
      "Host Name": "amarillo-rm.pepperdata.com",
      "Executor Deserialize Time": 11,
      "Executor Run Time": 98690,
      "Result Size": 1366,
      "JVM GC Time": 51928,
      "Result Serialization Time": 0,
      "Memory Bytes Spilled": 0,
      "Disk Bytes Spilled": 0,
      "Shuffle Read Metrics": {
        "Remote Blocks Fetched": 114,
        "Local Blocks Fetched": 6,
        "Fetch Wait Time": 5,
        "Remote Bytes Read": 743674,
        "Local Bytes Read": 41686,
        "Total Records Read": 120
      }
    }
  }
},
{
  "Event": "SparkListenerTaskEnd",
  "Stage ID": 9,
  "Stage Attempt ID": 0,
  "Task Type": "ResultTask",
  "Task End Reason": {
    "Reason": "Success"
  },
  "Task Info": {
    "Task ID": 770,
    "Index": 50,
    "Attempt": 0,
    "Launch Time": 1495496382879,
    "Executor ID": "8",
    "Host": "amarillo-n1.pepperdata.com",
    "Locality": "PROCESS_LOCAL",
    "Speculative": false,
    "Getting Result Time": 0,
    "Finish Time": 1495496487808,
    "Failed": false,
    "Accumulables": {
      {
        "ID": 7,
        "Name": "peakExecutionMemory",
        "Update": "96536946",
        "Value": "698097059",
        "Internal": true
      }
    },
    "Task Metrics": {
      "Host Name": "amarillo-n1.pepperdata.com",
      "Executor Deserialize Time": 4,
      "Executor Run Time": 104915,
      "Result Size": 1366,
      "JVM GC Time": 68939,
      "Result Serialization Time": 0,
      "Memory Bytes Spilled": 0,
      "Disk Bytes Spilled": 0,
      "Shuffle Read Metrics": {
        "Remote Blocks Fetched": 111,
        "Local Blocks Fetched": 9,
        "Fetch Wait Time": 10,
        "Remote Bytes Read": 921999,
        "Local Bytes Read": 74622,
        "Total Records Read": 120
      }
    }
  }
},
{
  "Event": "SparkListenerTaskEnd",
  "Stage ID": 9,
  "Stage Attempt ID": 0,
  "Task Type": "ResultTask",
  "Task End Reason": {
    "Reason": "Success"
  },
  "Task Info": {
    "Task ID": 769,
    "Index": 49,
    "Attempt": 0,
    "Launch Time": 1495496382874,
    "Executor ID": "5",
    "Host": "amarillo-rm.pepperdata.com",
    "Locality": "PROCESS_LOCAL",
    "Speculative": false,
    "Getting Result Time": 0,
    "Finish Time": 1495496507584,
    "Failed": false,
    "Accumulables": {
      {
        "ID": 7,
        "Name": "peakExecutionMemory",
        "Update": "105946616",
        "Value": "804043675",
        "Internal": true
      }
    },
    "Task Metrics": {
      "Host Name": "amarillo-rm.pepperdata.com",
      "Executor Deserialize Time": 9,
      "Executor Run Time": 124690,
      "Result Size": 1366,
      "JVM GC Time": 81294,
      "Result Serialization Time": 0,
      "Memory Bytes Spilled": 0,
      "Disk Bytes Spilled": 0,
      "Shuffle Read Metrics": {
        "Remote Blocks Fetched": 108,
        "Local Blocks Fetched": 12,
        "Fetch Wait Time": 2,
        "Remote Bytes Read": 972911,
        "Local Bytes Read": 113196,
        "Total Records Read": 120
      }
    }
  }
}
```



# Dr. Elephant

The screenshot displays the Dr. Elephant web application interface. At the top is a navigation bar with links: Dr. Elephant, Dashboard, Search, Compare, Flow History, Job History, and Help. The 'Search' tab is active.

On the left side, there is a search filter panel with the following sections:

- Search**: A search bar.
- Job/App ID**: A text input field containing 'Job ID'.
- Flow Execution URL/ID**: A text input field containing 'Flow Exec URL/ID'.
- User**: A text input field containing 'prod'.
- Queue**: A text input field containing 'Queue'.
- Job Type**: A dropdown menu with 'Spark' selected.
- Severity**: A dropdown menu with 'Critical' selected.
- All Analysis**: A dropdown menu.

The main content area shows search results. At the top, it says 'Results: Showing 20 of 878'. Below this, three job entries are visible:

Job/App ID	Job Name	Timestamp	Spark Configuration	Spark Executor Metrics	Spark Job Metrics	Spark Stage Metrics	GB Hours	Completion %	Duration
[prod] [Spark] application_1495423352820_0435	LogisticRegressionWithLBFGS	Wed May 24 2017 00:05:29 GMT-0700 (PDT)					4,550	0 %	0:08:16
[prod] [Spark] application_1495423352820_0434	LogisticRegressionDataGenerator	Tue May 23 2017 23:56:47 GMT-0700 (PDT)					0.871	6.77 %	0:01:32
[prod] [Spark] application_1495423352820_0433	DeriveKMeans with Params(hdfs://amarillo-n3.pepperdata.com:8020/user/prod/HIBench/Kmeans/Input/sam...	Tue May 23 2017 23:54:51 GMT-0700 (PDT)					1,312	0 %	0:02:19

At the bottom left, a URL bar shows: 220:8069/search?application\_1495423352820\_0435

# Spark Application Heuristics

[prod] [Spark] application_1495423352820_0410		Tue May 23 2017 23:17:17 GMT-0700 (PDT)
NWeightGraphX		
Spark Configuration Spark Executor Metrics Spark Job Metrics Spark Stage Metrics		
18.583 GB Hours 78.18 % 0:34:42		
[prod] [Spark] application_1495423352820_0409		Tue May 23 2017 22:42:16 GMT-0700 (PDT)
NWeight-DataGeneration		
Spark Configuration Spark Executor Metrics Spark Job Metrics Spark Stage Metrics		
1.651 GB Hours 0 % 0:02:58		
[prod] [Spark] application_1495423352820_0408		Tue May 23 2017 22:38:44 GMT-0700 (PDT)
LogisticRegressionWithLBFGS		
Spark Configuration Spark Executor Metrics Spark Job Metrics Spark Stage Metrics		
4.581 GB Hours 0 % 0:08:21		
[prod] [Spark] application_1495423352820_0407		Tue May 23 2017 22:29:58 GMT-0700 (PDT)
LogisticRegressionDataGenerator		
Spark Configuration Spark Executor Metrics Spark Job Metrics Spark Stage Metrics		
0.882 GB Hours 7.71 % 0:01:33		

# Spark Application Heuristics

[prod] [Spark] app  
NWeightGraphX

Spark Configuration

18.583 GB Hours

[prod] [Spark] app  
NWeight-DataGenerat

Spark Configuration

1.851 GB Hours

[prod] [Spark] app  
LogisticRegressionW

Spark Configuration

4.581 GB Hours

[prod] [Spark] app  
LogisticRegressionDa

Spark Configuration

0.882 GB Hours

Job Name  
NWeightGraphX

Job ID  
1495423352820\_0410

Job Type  
Spark

User  
prod

Queue  
root.prod

Severity  
CRITICAL

Duration  
34m

Wait Time  
0m

Start Time  
2017/05/23-22:42

Finish Time  
2017/05/23-23:17

**Spark Executor Metrics SEVERE** [explain](#)

Executor input bytes distribution	min: 0 B, p25: 119.71 MB, median: 1.38 GB, p75: 1.56 GB, max: 3.13 GB
Executor shuffle read bytes distribution	min: 0 B, p25: 359.37 MB, median: 594.32 MB, p75: 989.13 MB, max: 1.2 GB
Executor shuffle write bytes distribution	min: 0 B, p25: 137.55 MB, median: 689.56 MB, p75: 845.84 MB, max: 1.5 GB
Executor storage memory used distribution	min: 0 B, p25: 828.14 MB, median: 1.08 GB, p75: 1.47 GB, max: 1.86 GB
Executor storage memory utilization rate	0.480
Executor task time distribution	min: 0 sec, p25: 1 min 27 sec, median: 3 min 37 sec, p75: 6 min 36 sec, max: 11 min 54 sec
Total executor storage memory allocated	18.84 GB
Total executor storage memory used	9.03 GB

**Spark Job Metrics LOW** [explain](#)

**Spark Stage Metrics CRITICAL** [explain](#)

Spark completed stages count	42
Spark failed stages count	5
Spark stage failure rate	0.106
Spark stages with high task failure rates	stage 6, attempt 2 (task failure rate: 0.800) stage 6, attempt 3 (task failure rate: 0.600) stage 8, attempt 0 (task failure rate: 0.545) stage 8, attempt 1 (task failure rate: 0.571)
Spark stages with long average executor runtimes	

# Spark Application Heuristics

## Spark Executor Metrics

This heuristic concerns the distribution (min, 25p, median, 75p, max) of key executor metrics including input bytes, shuffle read bytes, shuffle write bytes, storage memory used, and task time. The max-to-median ratio determines the severity of any particular metric.

Spark application get resources from YARN allocated all at once, and don't release these until the application completes. Thus, it's important to balance load on the executors to avoid wasting resources.

To achieve better load balancing:

- use an appropriate number of partitions (some small multiple of the # of executors) so that there are enough tasks handling those partitions to keep the executors busy
- try avoiding key skew; you should know which partitioner you are using and what is the distribution of your keys
- consider enabling `spark.speculation`, so that straggler tasks can be re-launched



# 1: Configuration Heuristics

- Display some basic config settings for your app
- Complain if some settings not explicitly set
- Recommend configuring an external shuffle service (especially if dynamic allocation is enabled)
- These recommendations won't change over multiple runs of an application

## 2: Stages and Jobs Heuristics

- Simple alarms showing stage and job failure rates
- Good for seeing when there's a problem

## 3: Executor Heuristics

- Looks at the distribution across executors of several different metrics
- Outliers in these distributions probably indicate:
  - Suboptimal partitioning.
  - One or more slow executors due to external circumstances (cluster weather)

## 3: Partitions Heuristic

- Ideally data for each task will fit into the RAM available to that task.
- Cloudera has an excellent blog on Spark tuning:

$$\frac{(\text{observed shuffle write}) * (\text{observed shuffle spill memory}) * (\text{spark.executor.cores})}{(\text{observed shuffle spill disk}) * (\text{spark.executor.memory}) * (\text{spark.shuffle.memoryFraction}) * (\text{spark.shuffle.safetyFraction})}$$

---

<http://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-2/>




# More Heuristics?

Yes, please! Dr. Elephant is open source.

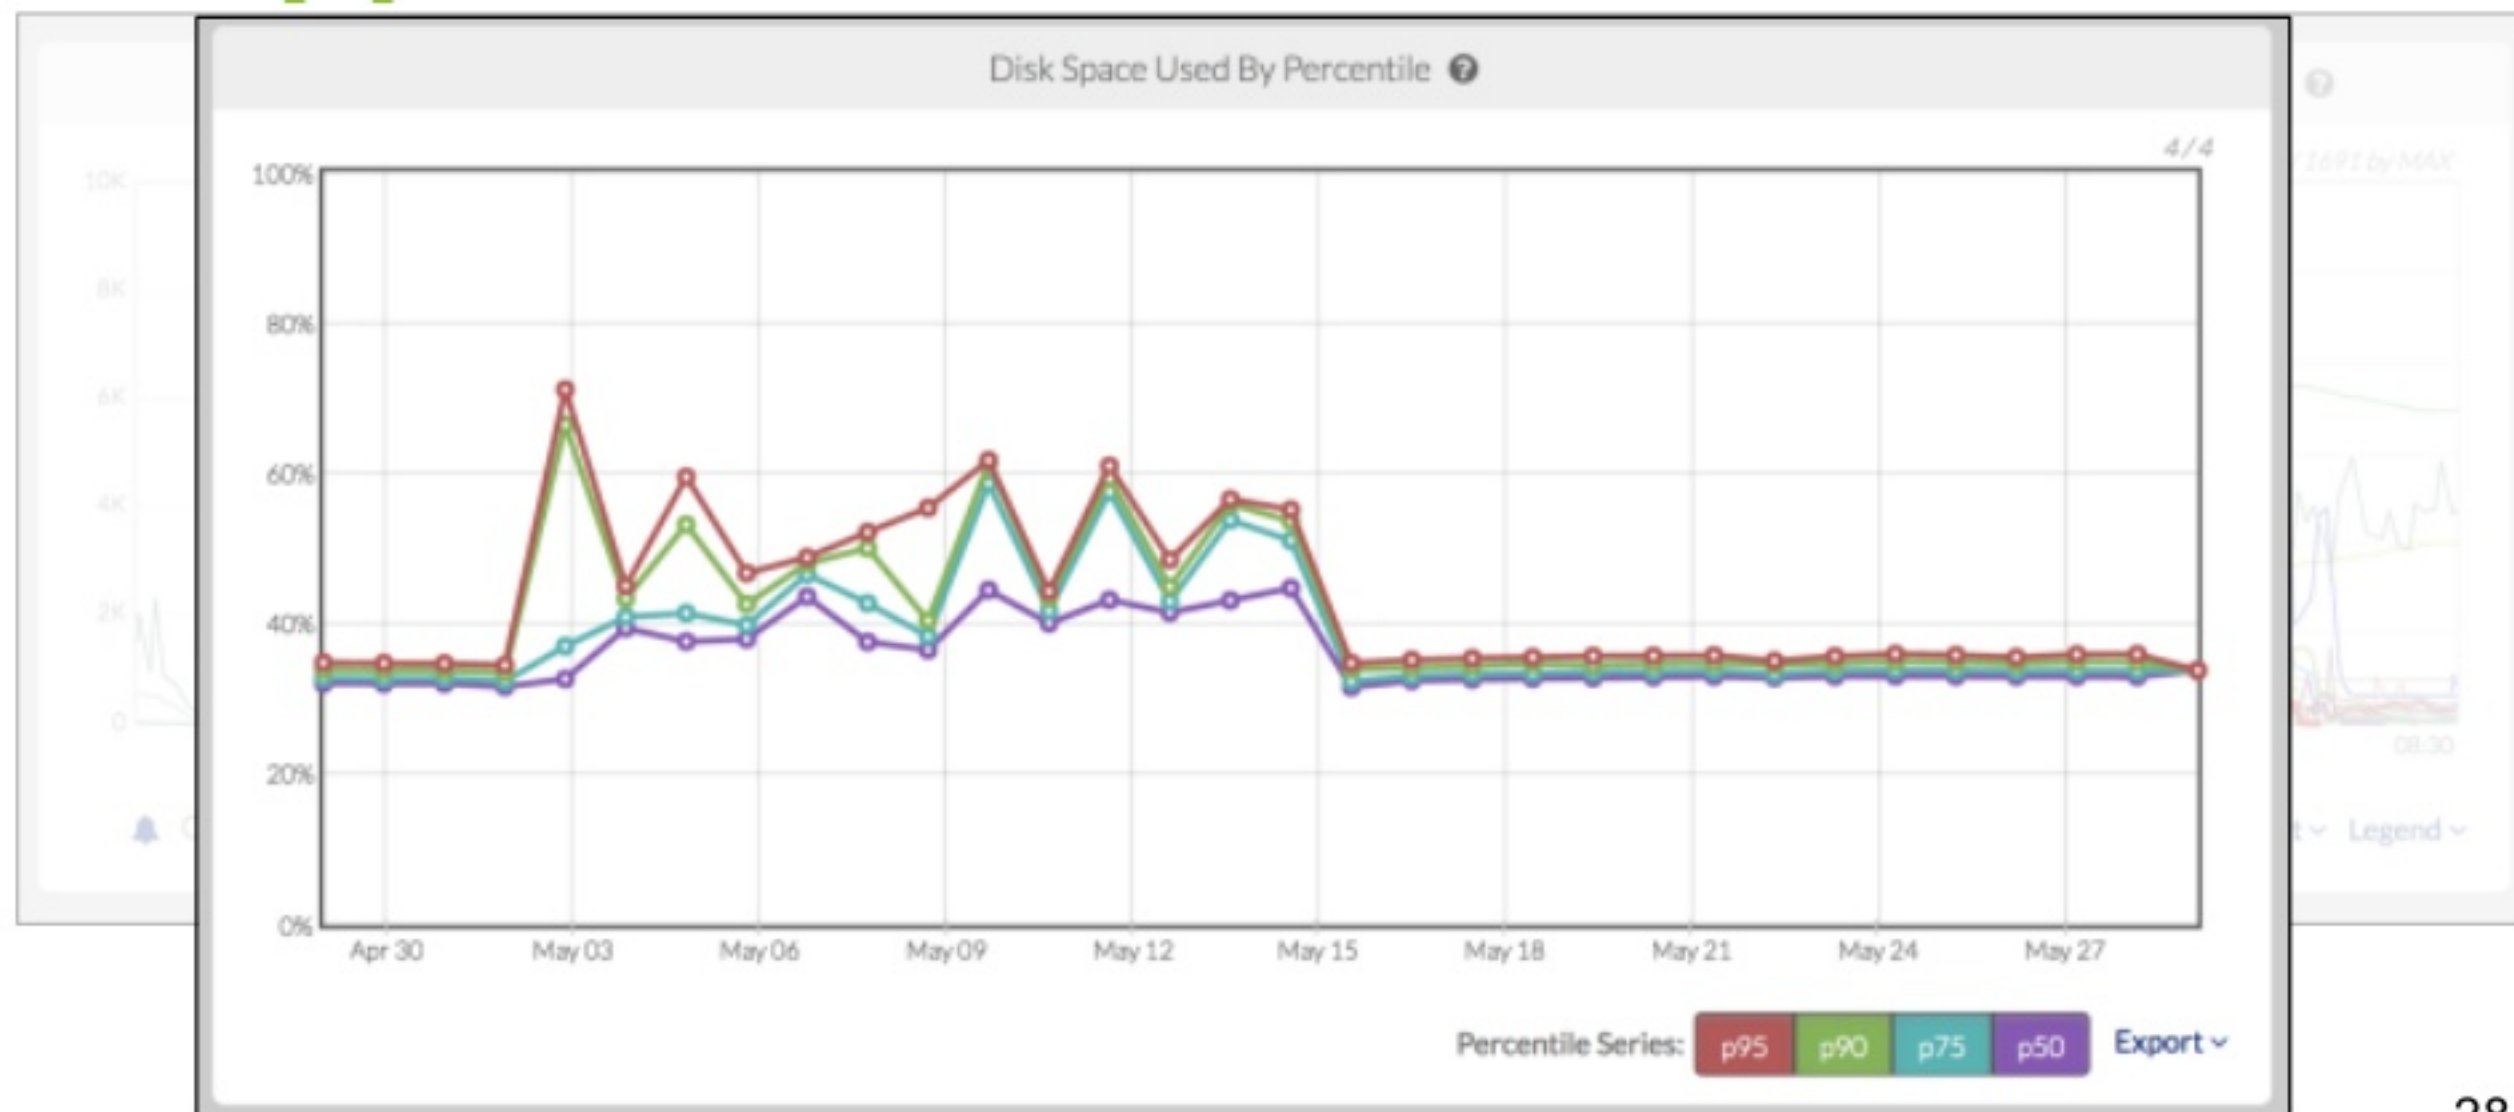
# Pepperdata

- Capacity Optimizer
- Policy Enforcer
- Cluster Analyzer

# Pepperdata



# Pepperdata





# Pepperdata




- Capacity Optimizer
- Policy Enforcer
- Cluster Analyzer
- Application Profiler

} Mostly for Operators  
} For Developers


# Application Profiler


- Benefits to our users:
  - Provide simple answers to simple questions
  - Combination of metrics for experts, simple actionable insights for users
  - Pepperdata support
- Why stay close to open source?
  - Heuristics

# Application Profiler, Hardware and Cluster Weather


  amarillo - Application Profiler 

Show Last: 24h 6h 1h

Start Time: 2017/05/23-10:32 

End Time: 2017/05/24-10:32 

Range: 1 day

Times are: UTC-07:00 (local) 

Expand Time Range

Apply Time Range

App Type: MapReduce Spark

Columns to show

Job Info: Job Name Job ID Job Type User Queue Severity

Duration Wait Time Start Time Finish Time








Heuristics: Spark (4/4)

Compare these jobs:

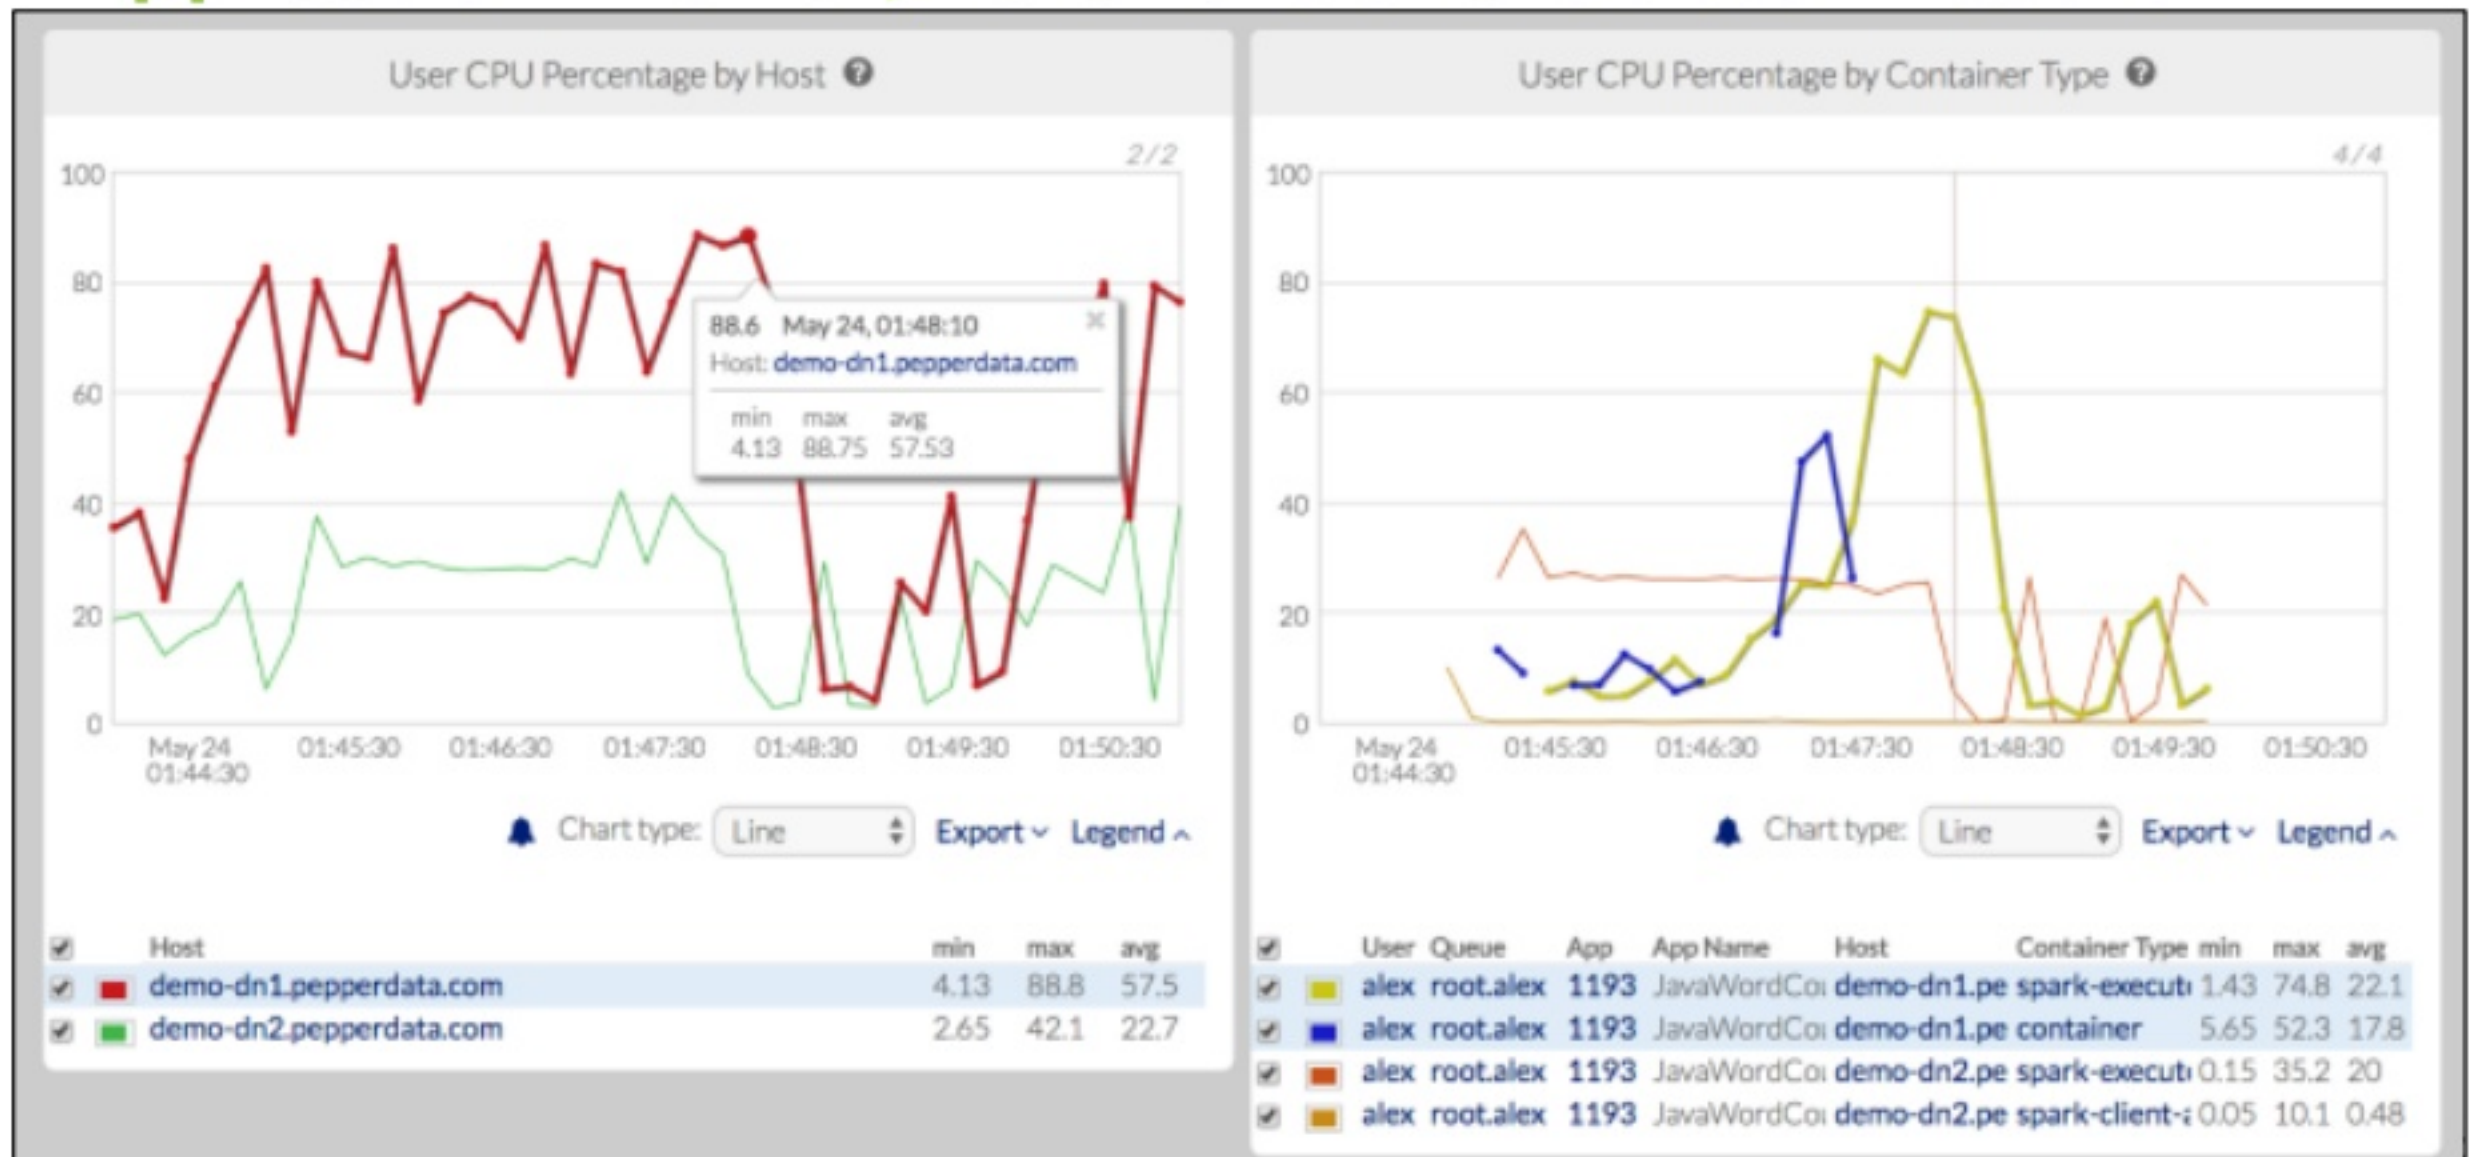
select jobs using checkboxes in table

Compare

60 rows

compare	chart	Job Name	Job ID	Severity	Configuration	Executor Metrics	Job Metrics	Stage Metrics	Duration	Wait Time	Start Time
<input type="checkbox"/>		LogisticRegressio...	1495423352820_0...	CRITICAL	MODERATE	CRITICAL			4m	0m	2017/05/23-12:14
<input type="checkbox"/>		NWeightGraphX	1495423352820_0...	CRITICAL	MODERATE	SEVERE	LOW	CRITICAL	34m	0m	2017/05/23-22:42
<input type="checkbox"/>		NWeight-DataGe...	1495423352820_0...	SEVERE	MODERATE	SEVERE			2m	0m	2017/05/23-22:39
<input type="checkbox"/>		LogisticRegressio...	1495423352820_0...	MODERATE	MODERATE	MODERATE			8m	0m	2017/05/23-23:57
<input type="checkbox"/>		SparseNaiveBaye...	1495423352820_0...	MODERATE	MODERATE	MODERATE			1m	0m	2017/05/23-23:43
<input type="checkbox"/>		ScalaWordCount	1495423352820_0...	MODERATE	MODERATE	MODERATE			< 1m	0m	2017/05/23-23:26
<input type="checkbox"/>		LogisticRegressio...	1495423352820_0...	MODERATE	MODERATE	MODERATE			8m	0m	2017/05/23-22:30

# Application Profiler, Hardware and Cluster Weather







# Thanks!

Stop by the Pepperdata booth (#101)

Come to the Dr Elephant Meetup:  
6:00 PM Wednesday, June 7, 2017  
LinkedIn San Francisco Office  
222 2nd Street, San Francisco

Get involved with Dr. Elephant:  
<https://github.com/linkedin/dr-elephant>

Contact us:  
[simon@pepperdata.com](mailto:simon@pepperdata.com), [csteinbach@linkedin.com](mailto:csteinbach@linkedin.com)