SPARK SUMMIT 2017

JUST-IN-TIME ANALYTICS AND THE NEED FOR AUTONOMOUS DATABASE ADMINISTRATION

Wes Holler, *Chief Architect*
Algebraix Data

Data Algebra

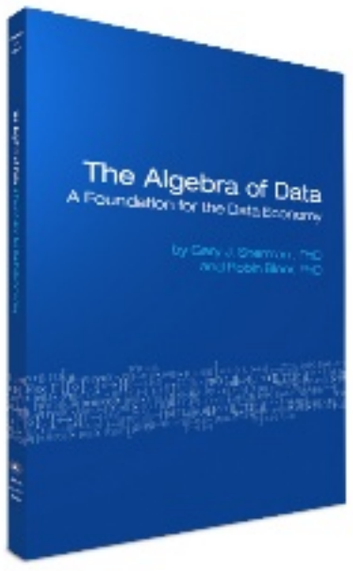

- A **novel application of Set Theory** to Data Processing
- **Applicable to many data models** including SQL

Just-in-Time Analytics

needs…

Autonomous Data Management

JIT Analytics and the Life of a Modern Analyst

This shouldn't require intimate knowledge of how underlying systems work.

Spark for JIT Analytics: *The Good*

- Unified API
- Schema-on-read and Heterogeneous Data Sources
- Declarative Languages/APIs and Catalyst
- Elastic Compute

Spark for JIT Analytics: *The Bad*

- Challenges for interactivity, efficiency, and scalability
- Cost of creating and maintaining "glue code"
- ***Data scientists and engineers are doing DBA work***

Database Management Responsibilities

Capacity planning

Configuration

We will focus on the performance and tuning aspects

Performance tuning

A billion other things

Improving and Maintaining Performance

- Indexes
- Materialize views
- Pre-aggregate data
- *Lots* of configuration

Performance Tuning Strategies in Spark

- Segment, cache, and checkpoint
- Configure cluster parameters
- `spark.sql.shuffle.partitions`

What is the Problem with Manual Tuning?

- Varies with the data (skew and scale), queries, and hardware
- Often done through trial and error
- Problems are exacerbated with JIT analytics case
- Shared resources

What is the Problem with Manual Tuning?

It is hard and time-consuming.

A Motivating Example for
Autonomous Data Management

Query 1

```
SELECT *
FROM A JOIN B ON A.a = B.b
WHERE B.b2 < 100
```

Query 2

```
SELECT A.a, A.foo
FROM A JOIN B ON A.a = B.b
JOIN C ON B.c = C.c
WHERE B.b2 < 100 AND C.bar = "bar"
```

Query 3

```
SELECT (A.foo + D.baz) AS foo_or_fu_baz
FROM A JOIN D ON A.d = D.d
JOIN B ON A.a = B.b
WHERE (D.baz LIKE "baz%" OR A.foo in ("foo", "fu"))
AND B.b2 < 100
```

Query 1 w/ FPP

```
1  SELECT *
2  FROM (
3      SELECT *
4      FROM A JOIN B ON A.a = B.b
5      WHERE B.b2 < 100
6  ) as ChunkExample
```

Query 2 w/ FPP

```
1  SELECT A.a, A.foo
2  FROM (
3      SELECT *
4      FROM A JOIN B ON A.a = B.b
5      WHERE B.b2 < 100
6  ) as ChunkExample
7  JOIN C ON B.c = C.c
8  WHERE C.bar = "bar"
```

Query 3 w/ FPP

```
1  SELECT (A.foo + D.baz) AS foo_or_fu_baz
2  FROM #ChunkExample
3  JOIN D ON A.d = D.d
4  WHERE (D.baz LIKE "baz%" OR A.foo in ("foo", "fu"))
```

A Motivating Example for Autonomous Data Management

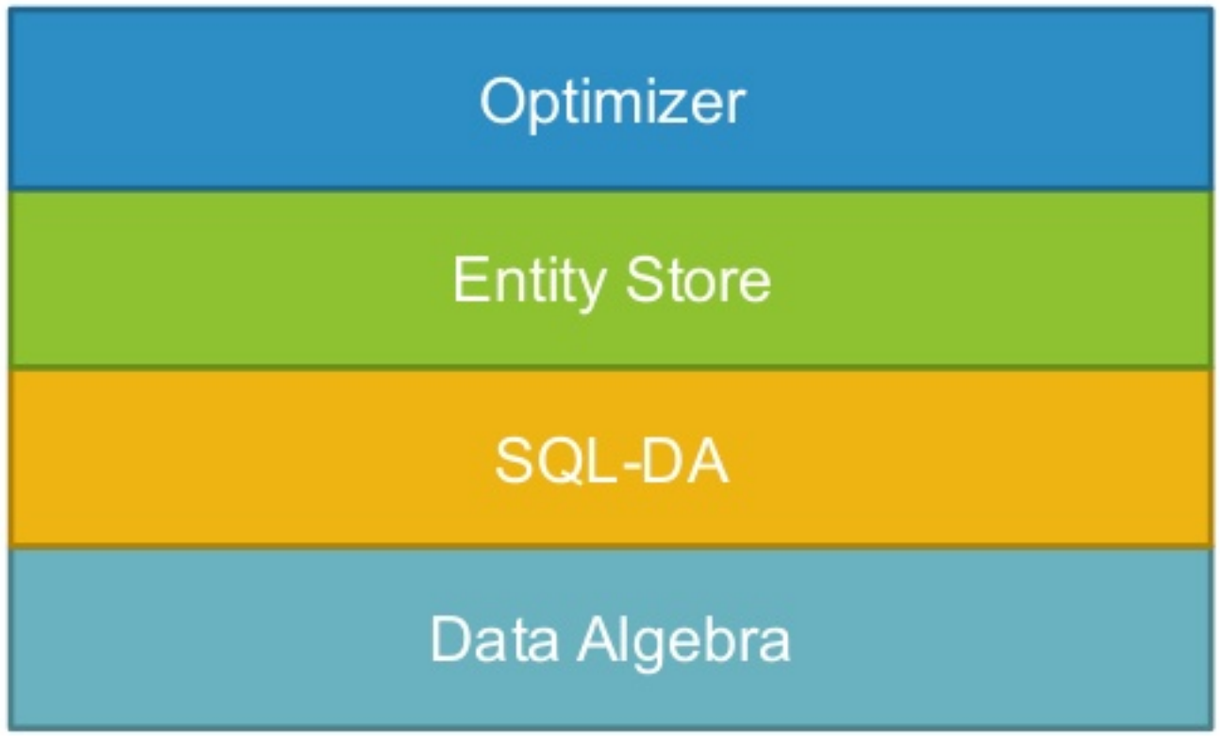

```
SELECT *
FROM A JOIN B ON A.a = B.b
WHERE B.b2 < 100
```

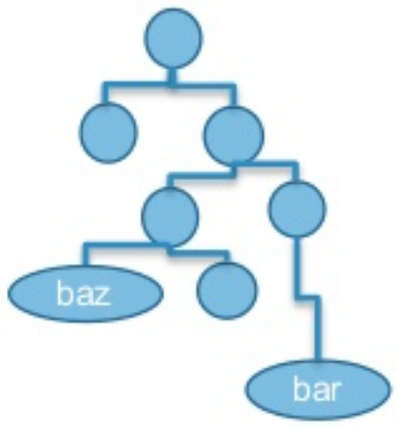

$$A \nabla B = \left(\begin{Bmatrix} \{0 \mapsto \alpha, \ldots, 3 \mapsto 42.0\}: 3, \\ \{0 \mapsto \beta\}: 1, \\ \ldots \end{Bmatrix}, \begin{pmatrix} 0 & \cdots & 3 & 4 \\ a & \cdots & b & b2 \\ int & \ldots & float & int \\ A & A & B & B \end{pmatrix} \right)$$

$$Q := filter_{[4]<100}(filter_{[0]<[3]}(A \nabla B))$$

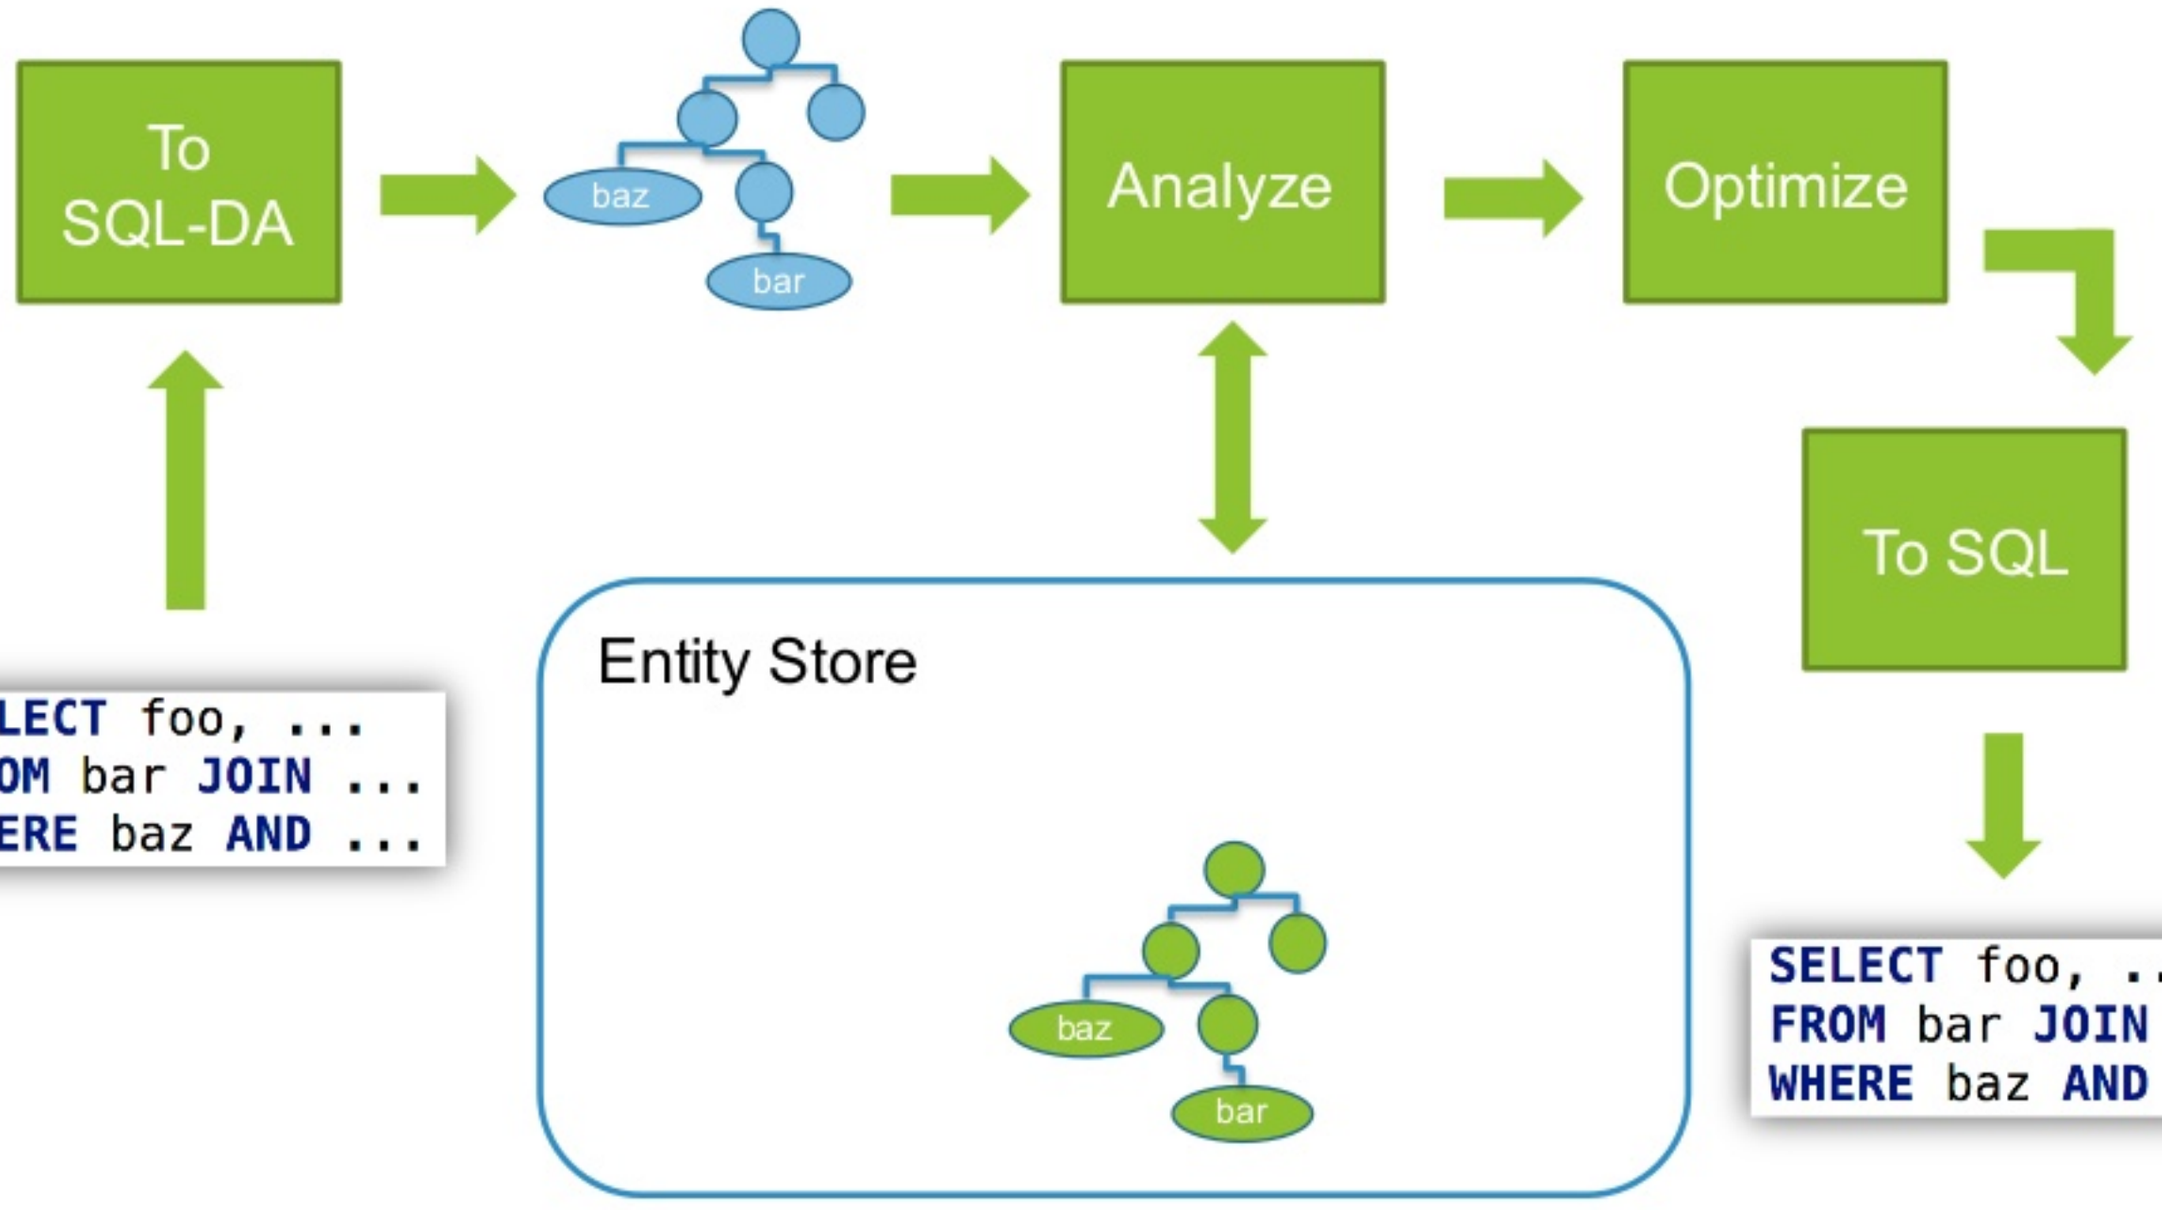
To
SQL-DA
baz
bar
Analyze
Optimize
To SQL
SELECT foo, ...
FROM bar JOIN ...
WHERE baz AND ...
Entity Store
baz
bar
SELECT foo, ...
FROM bar JOIN ...
WHERE baz AND ...

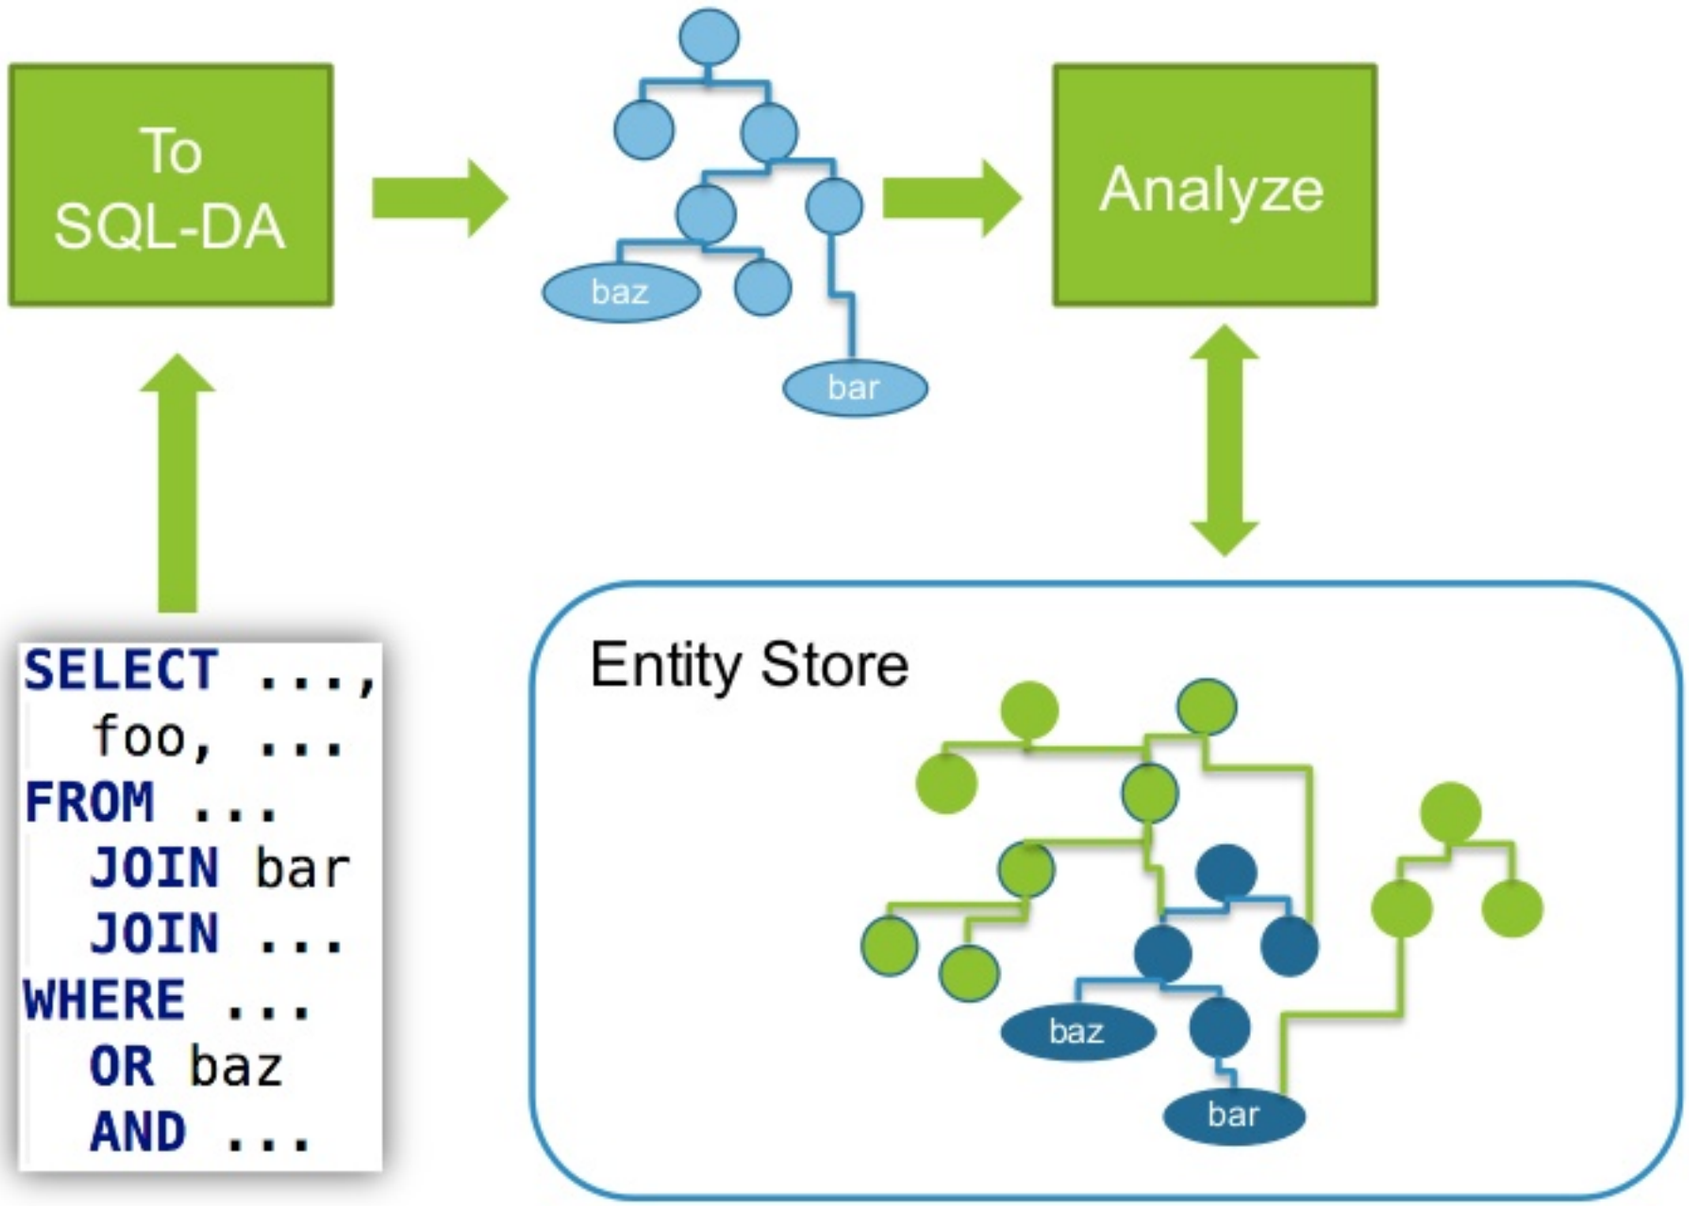
To
SQL-DA
baz
bar
Analyze
SELECT ...,
foo, ...
FROM ...
JOIN bar
JOIN ...
WHERE ...
OR baz
AND ...
Entity Store
baz
bar


To SQL-DA
Analyze
baz
bar
Optimize
Materialize View
SELECT ...,
foo, ...
FROM ...
JOIN bar
JOIN ...
WHERE ...
OR baz
AND ...
Entity Store
baz
bar

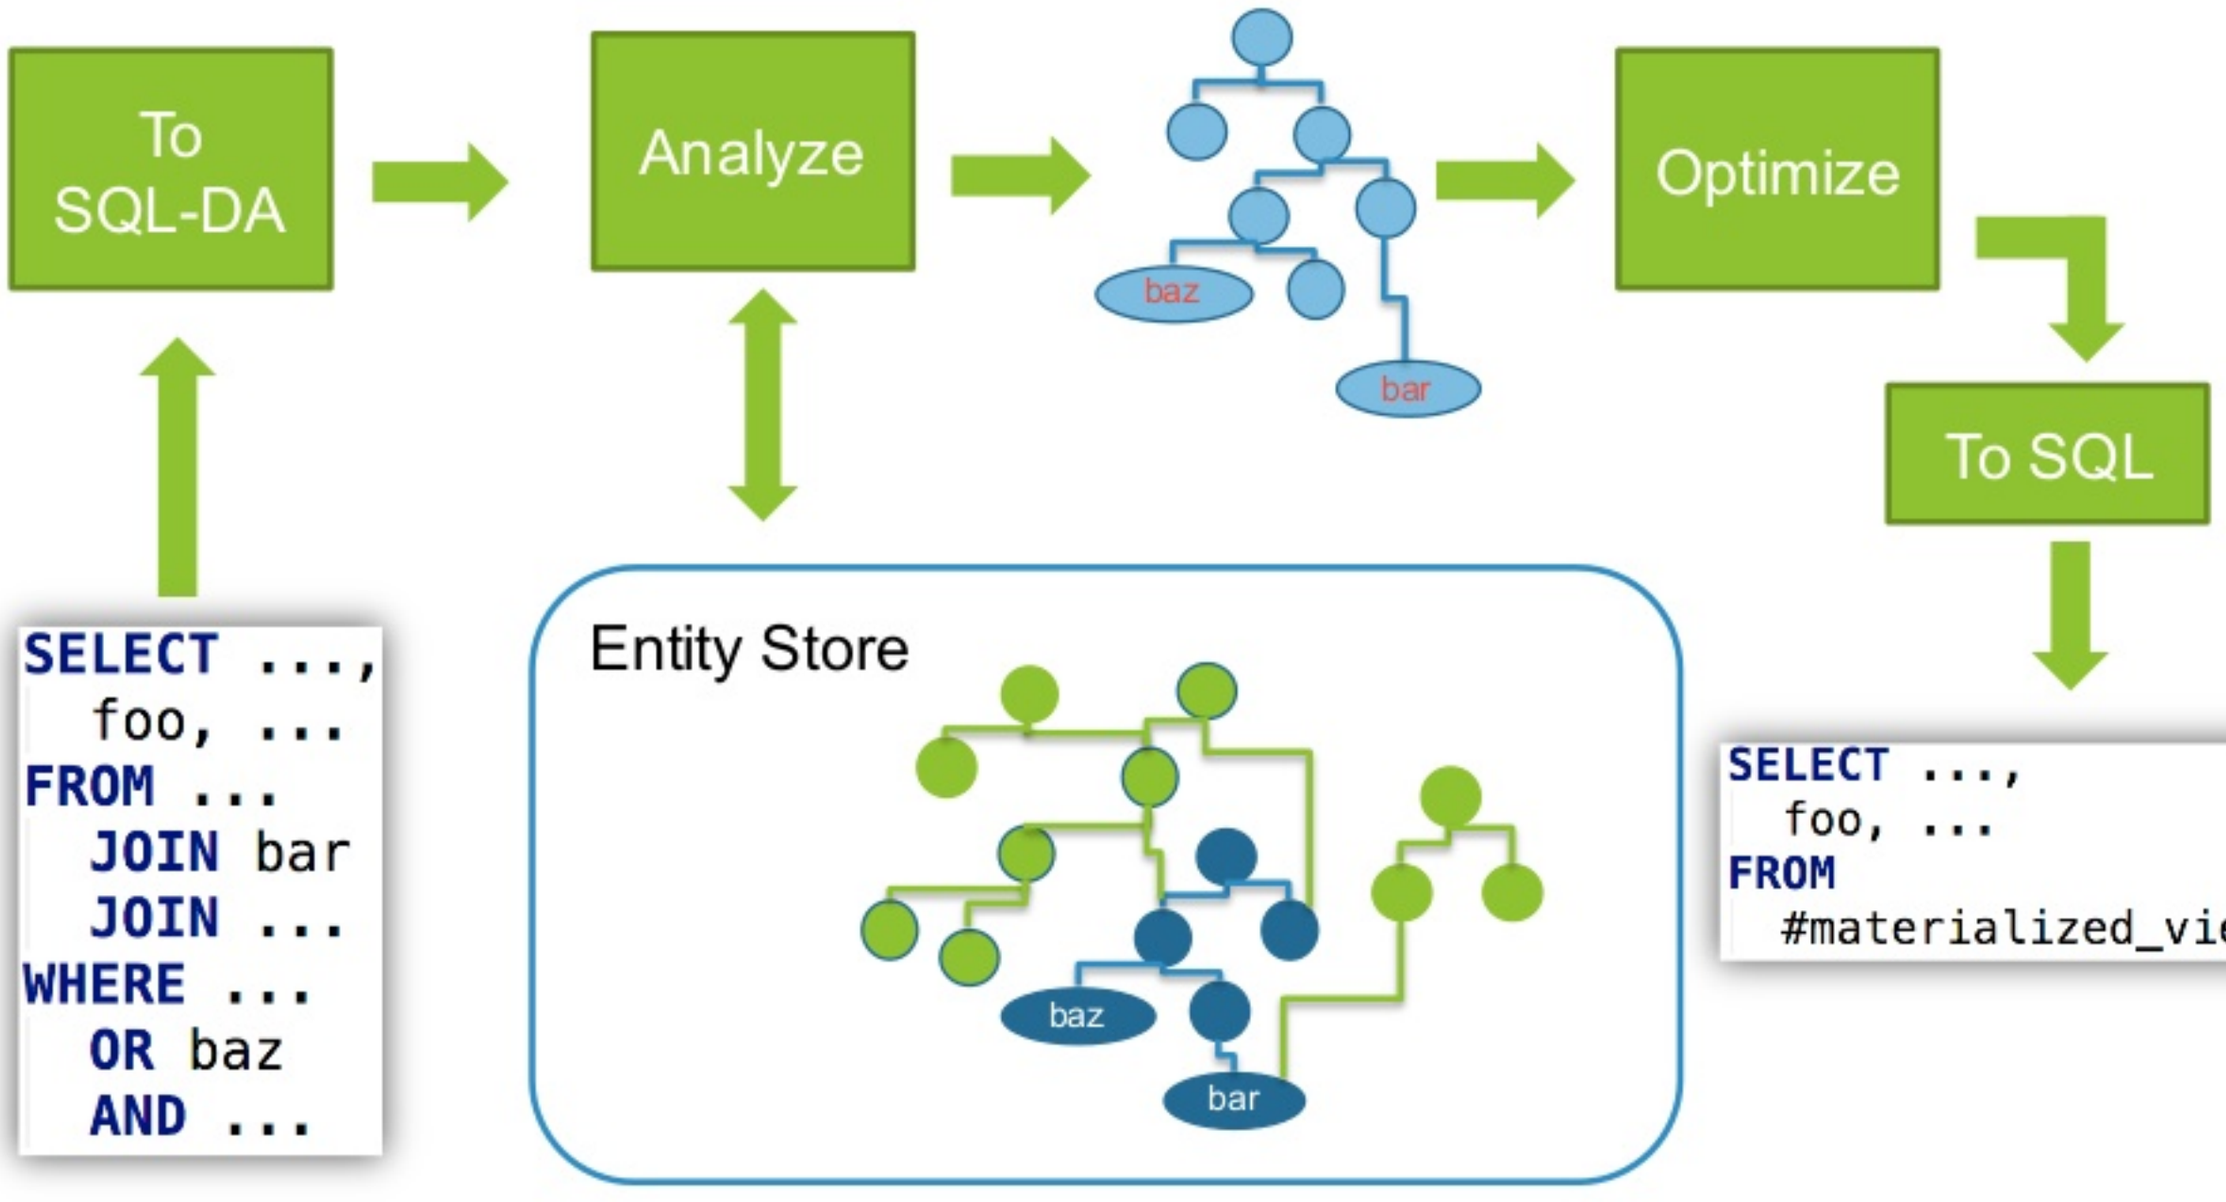
To SQL-DA
Analyze
baz
bar
Optimize
To SQL
SELECT ...,
foo, ...
FROM ...
JOIN bar
JOIN ...
WHERE ...
OR baz
AND ...
Entity Store
baz
bar
SELECT ...,
foo, ...
FROM
#materialized_view

Complex Query Expressions are Turned Into Look-ups

Benefits of Autonomous Data Management

- Reduce query time
- Reduce computation resources required
- ***Allow the analyst to focus on problem solving, not data management***

Algebraix Inside:
An Implementation of ADM

The PySpark API (DataFrames and SQL) is shimmed.

Before

```
1  from pyspark import *
2  from pyspark.sql import SQLContext
3
4  conf = SparkConf()
5  sc = SparkContext(conf=conf)
6  sqlContext = SQLContext(sc)
7
8  names = sc.readText("people.txt")
9
10 namesDF = sc.createDataFrame(names)
11
12 namesDF.registerTempTable("names")
13
14 sqlContext.sql("""
15     SELECT * FROM names
16 """).show()
```

After

```
1  from aqaspark import *
2
3
4  conf = SparkConf()
5  sc = SparkContext(conf=conf)
6  sqlContext = SQLContext(sc)
7
8  names = sc.readText("people.txt")
9
10 namesDF = sc.createDataFrame(names)
11
12 namesDF.registerTempTable("names")
13
14 sqlContext.sql("""
15     SELECT * FROM names
16 """).show()
```

Wrap Up

Autonomous Data Management makes Spark great for SQL analytics.

Thank You.

@wes_holler
wholler@algebraixdata.com

algebraix data
www.algebraixdata.com
tstraub@algebraixdata.com