Smart Scalable Feature Reduction with Random Forests

Erik Erlandson

Red Hat, Inc.

Erik Erlandson

- Software Engineer
- Radanalytics.io community
- Apache Spark on OpenShift
- Intelligent Applications in the cloud

Talk

- Motivate Feature Reduction
- Random Forest Clustering
- T-Digest Feature Sketching
- RF Feature Reduction
- Example: Tox21 Assay Data

Features

Measurable Properties!

Model Training

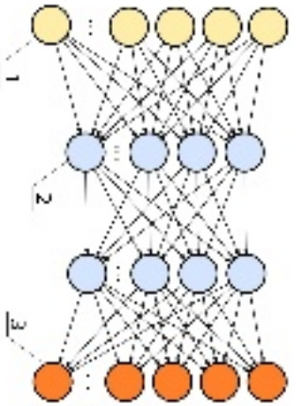

Evaluation

Feature Reduction

Full Feature Set

Identify Useful Features

Reduced Feature Set

Feature Sets Can Be Very Large

hundreds

thousands

...

millions

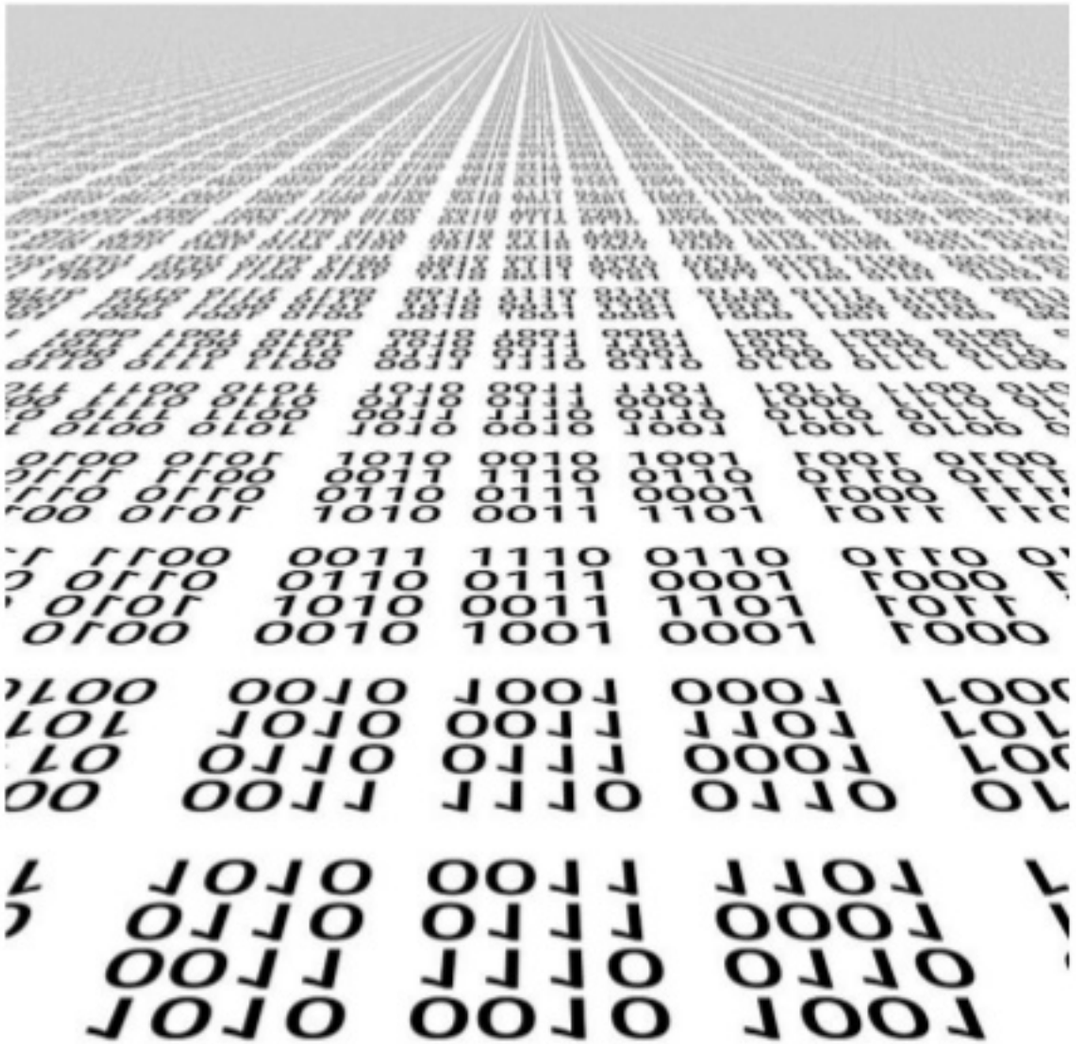

Features Cost Resources

Memory

Network

Time

Disk

Energy

Features Inject Noise

Model Training Without Reduction

With Feature Reduction

Features Impact Model Size

Model Without Reduction

With Reduction

Representation & Transfer Learning

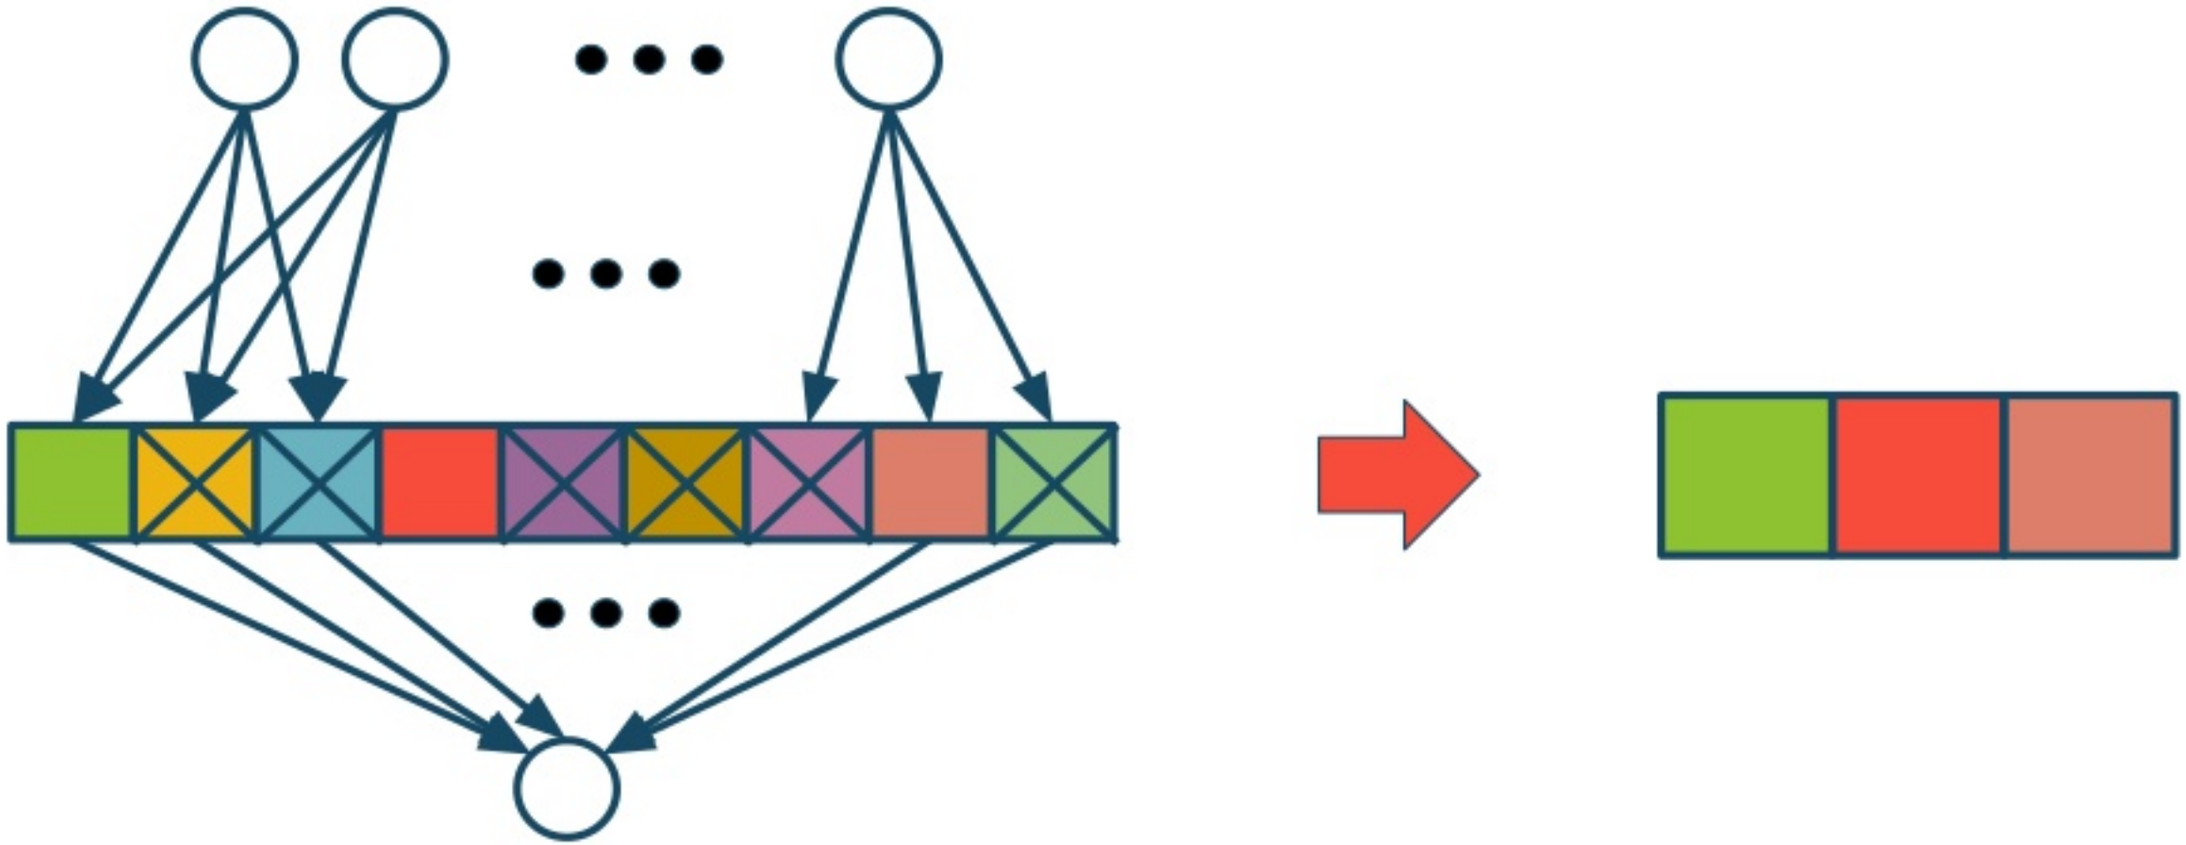

Random Forests

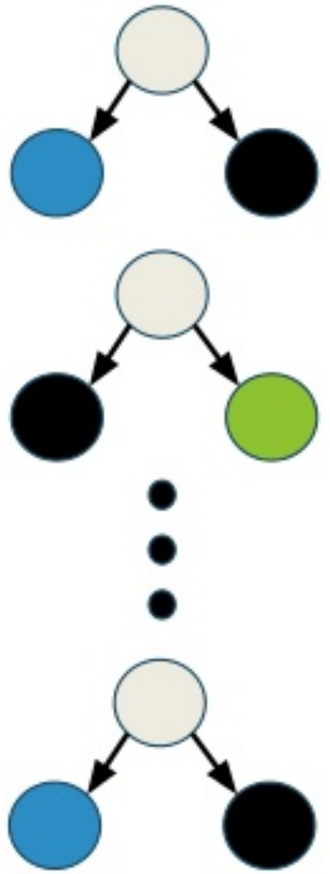

Leo Breiman (2001)

Ensemble of Decision Tree Models

Each tree trains on random subset of data

Each split considers random subset of features

Random Forest Clustering

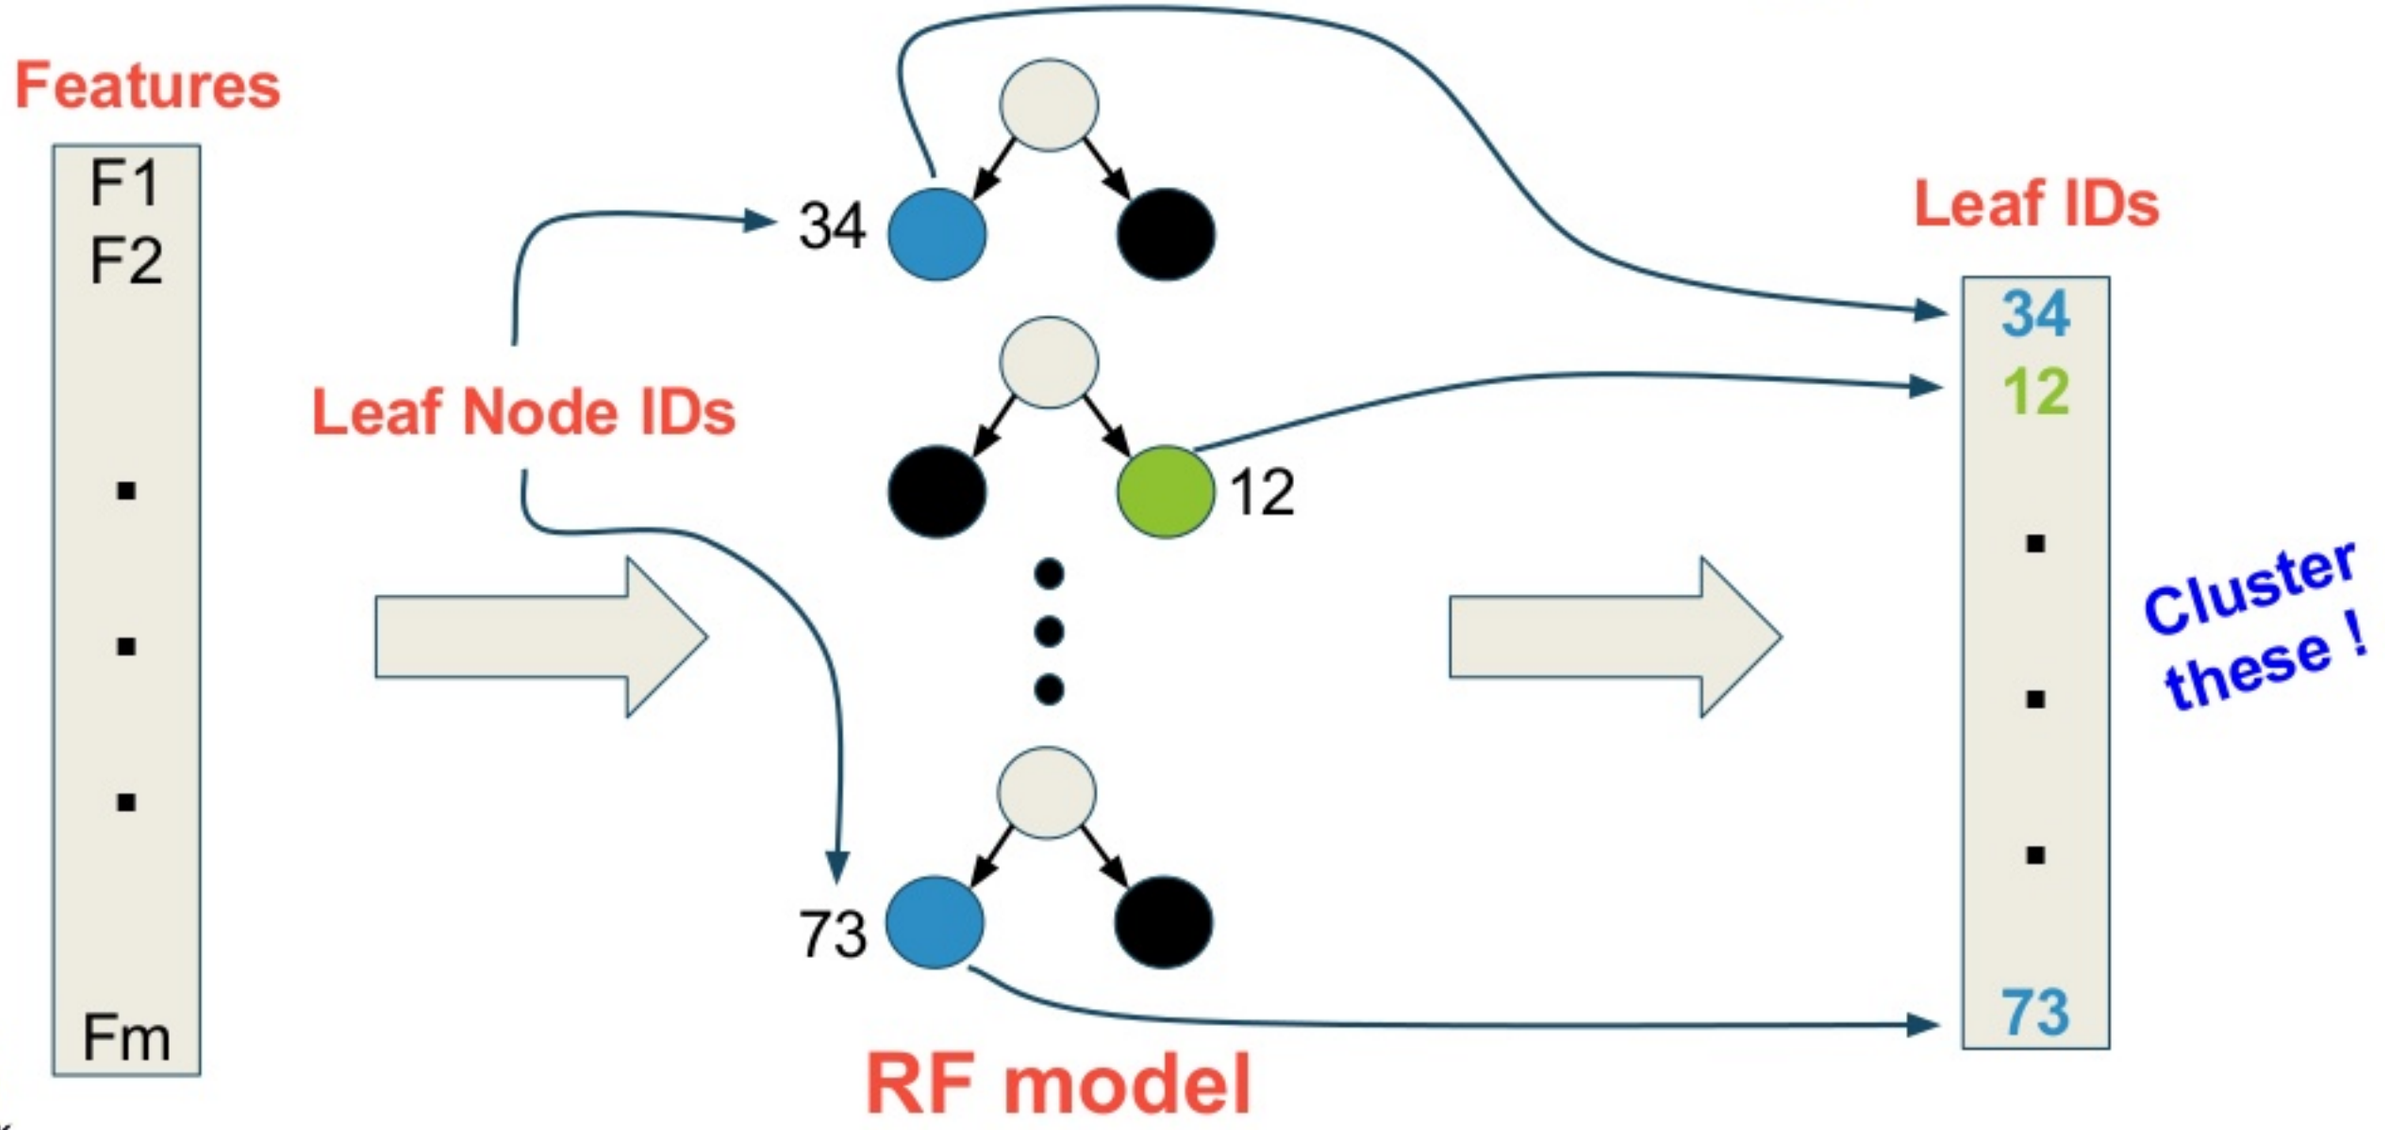

2 Key Benefits of RF Clustering

RF Training ignores unhelpful features

Features Used by RF Model << Full Feature Set

Data with a Joint Distribution in R^2

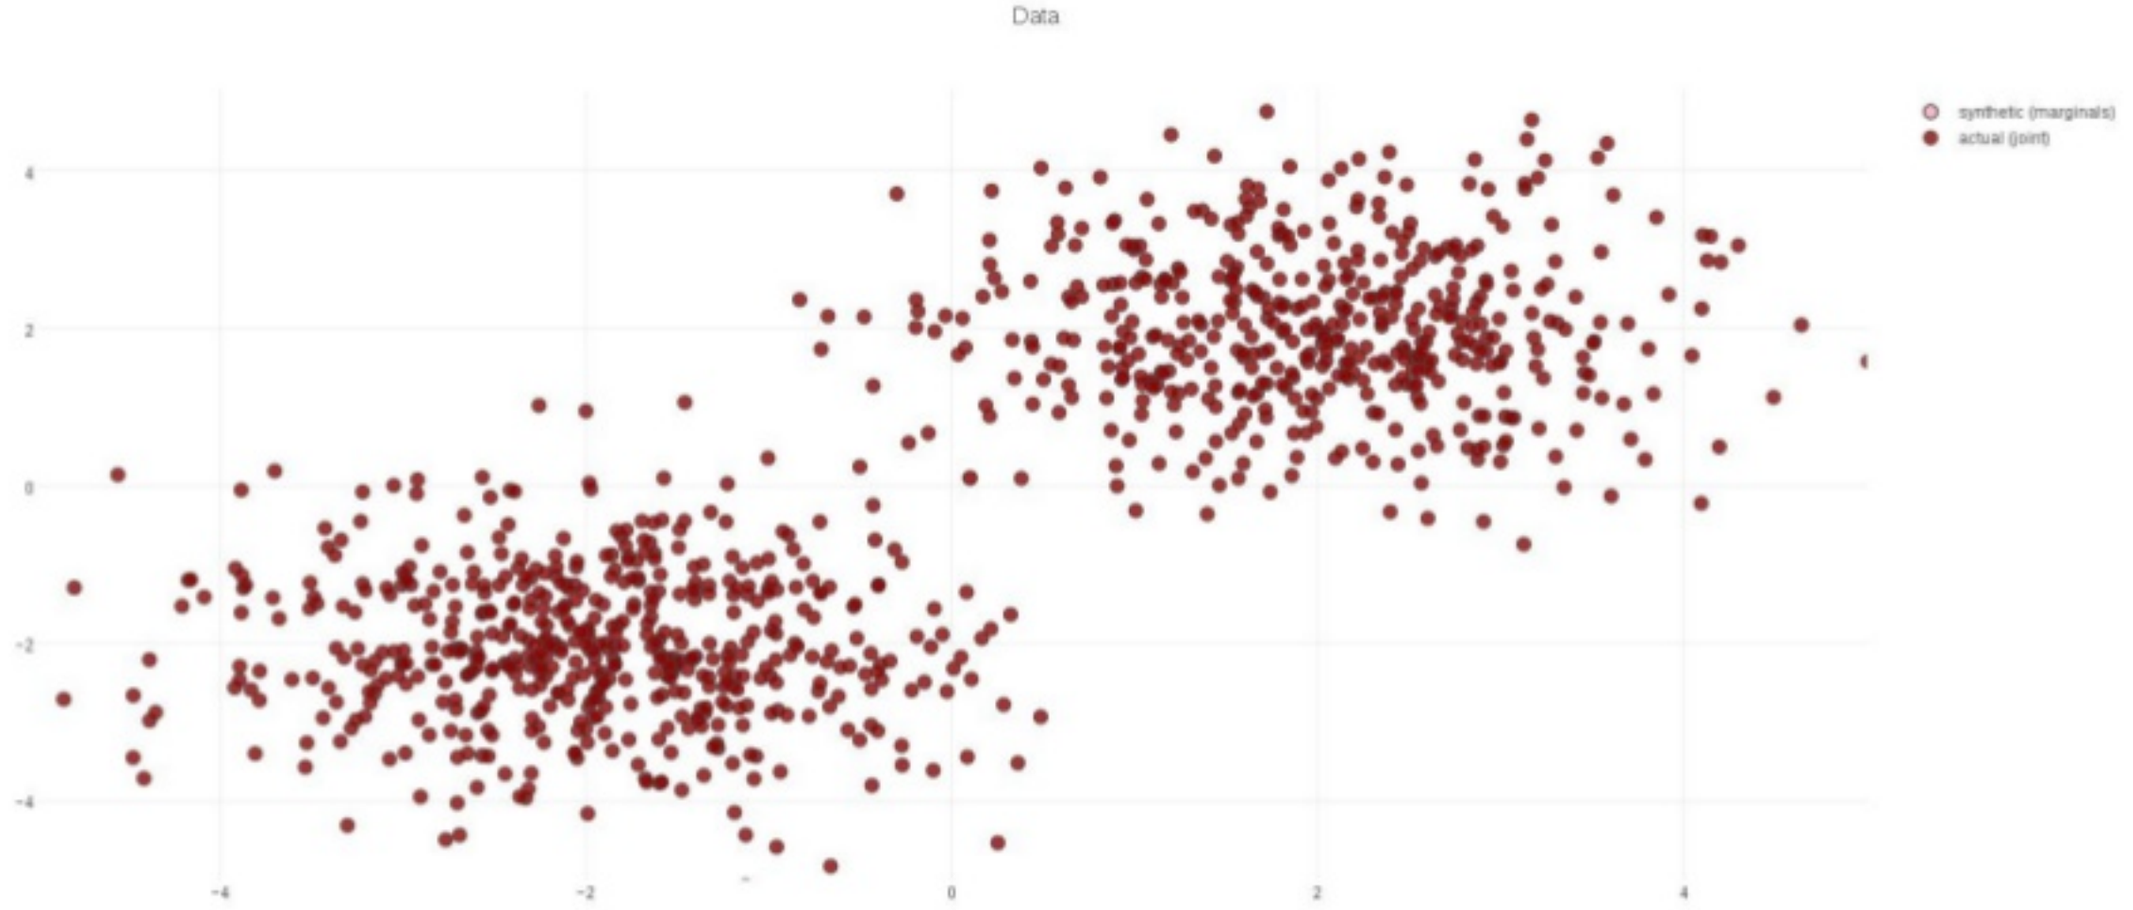

Data with Synthetic

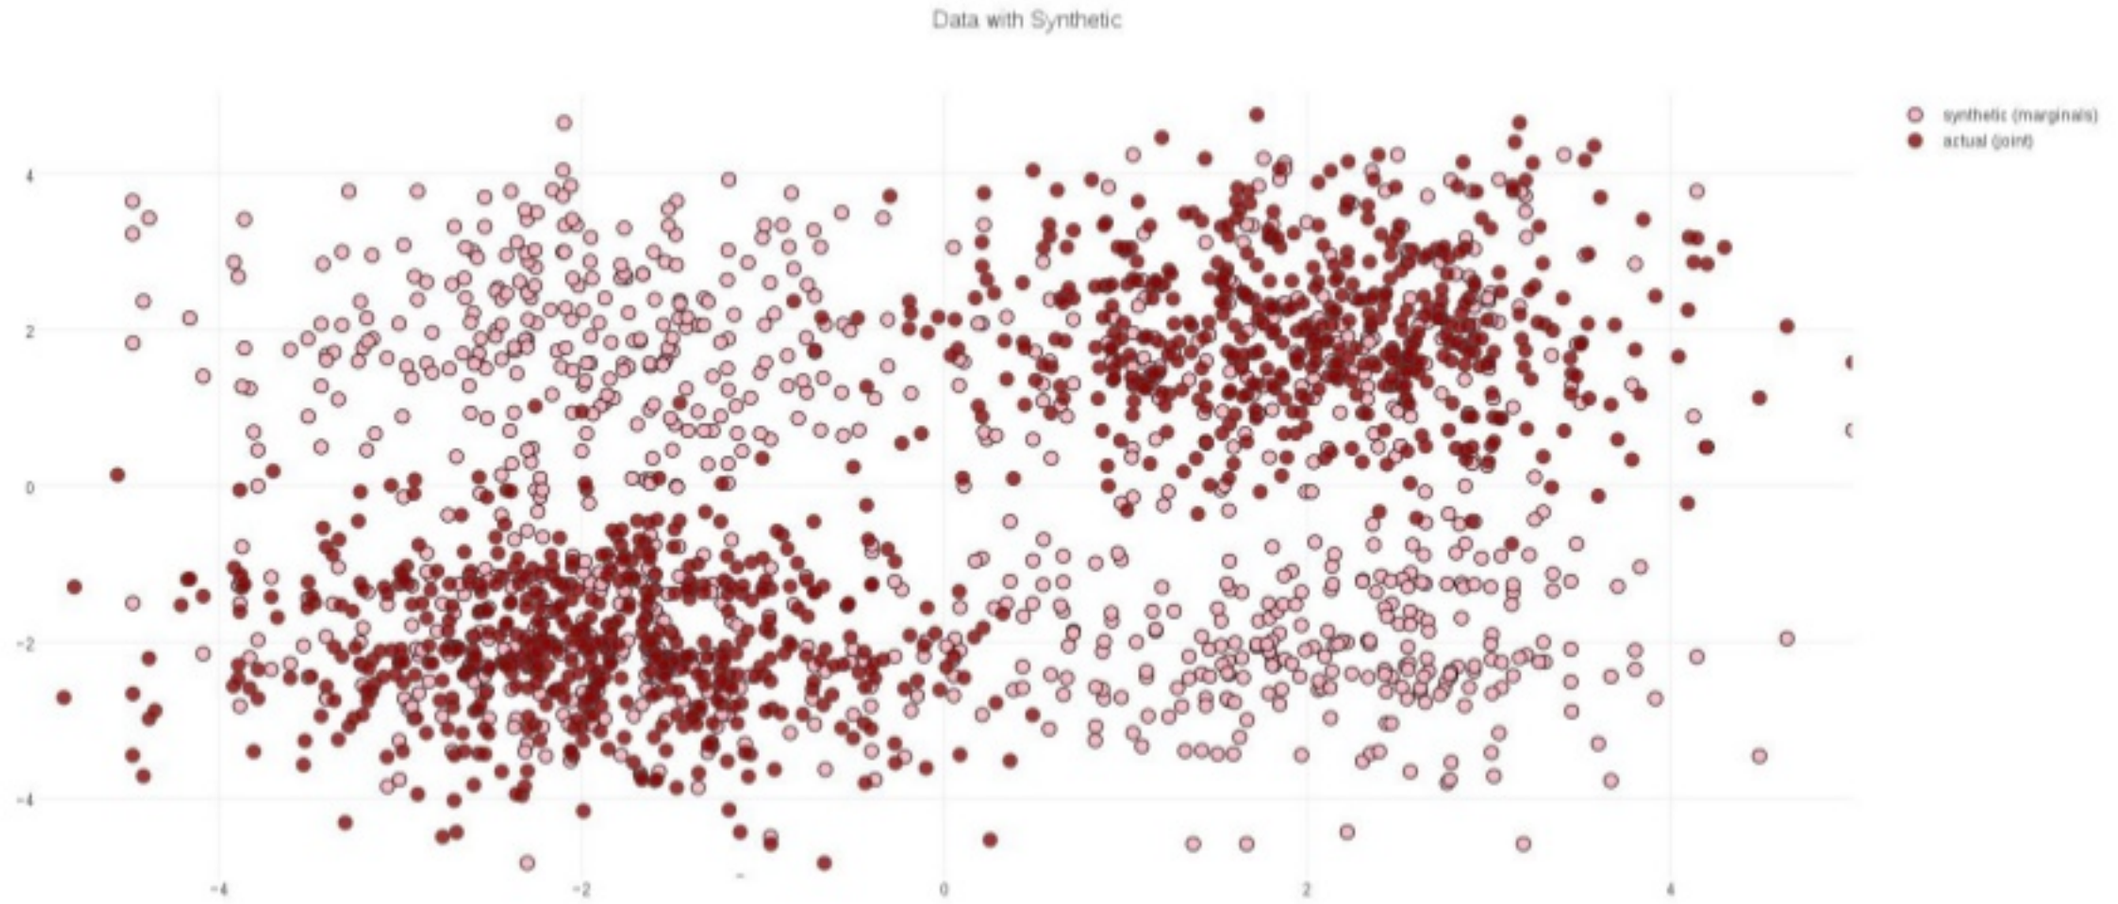

RF Rules for Data (non-synthetic)

```
List((x2 <= -1.32), (x1 <=  0.87))
List((x1  > -1.37), (x2  >  1.03))
List((x2 <=  2.09), (x1 <=  0.87))
List((x1 <=  2.13), (x2 <= -1.32))
List((x2 <= -2.31), (x1 <=  0.87))
```

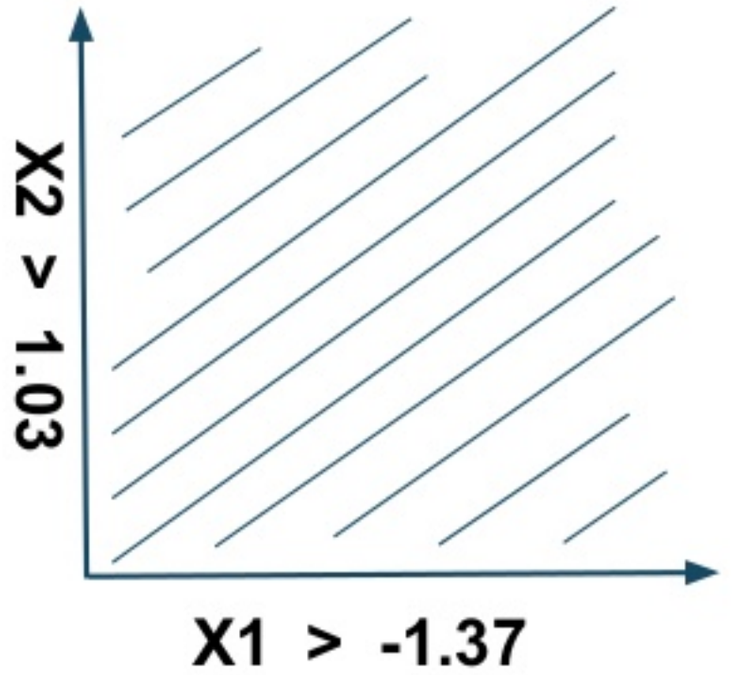

RF Rules in Feature Space

What Features Did the RF Use?

```
List((x2 <= -1.32), (x1 <=  0.87))
List((x1  > -1.37), (x2  >  1.03))
List((x2 <=  2.09), (x1 <=  0.87))
List((x1 <=  2.13), (x2 <= -1.32))
List((x2 <= -2.31), (x1 <=  0.87))
```

reduced = {"x1", "x2"}

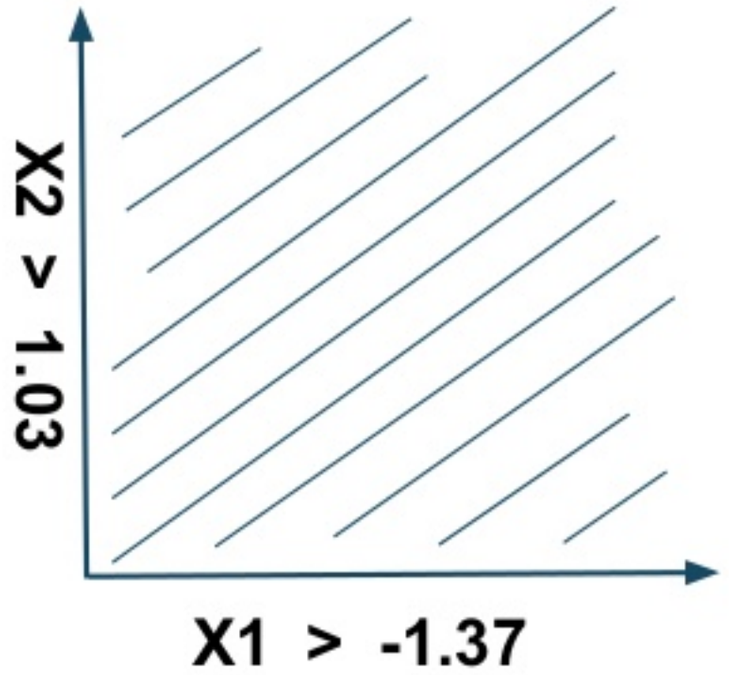

T-Digest Sketches a Distribution

3.4

6.0

2.5

⋮

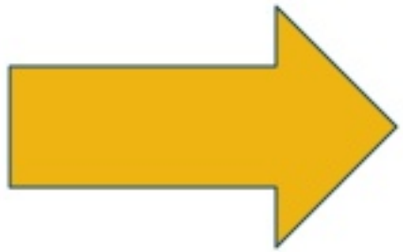

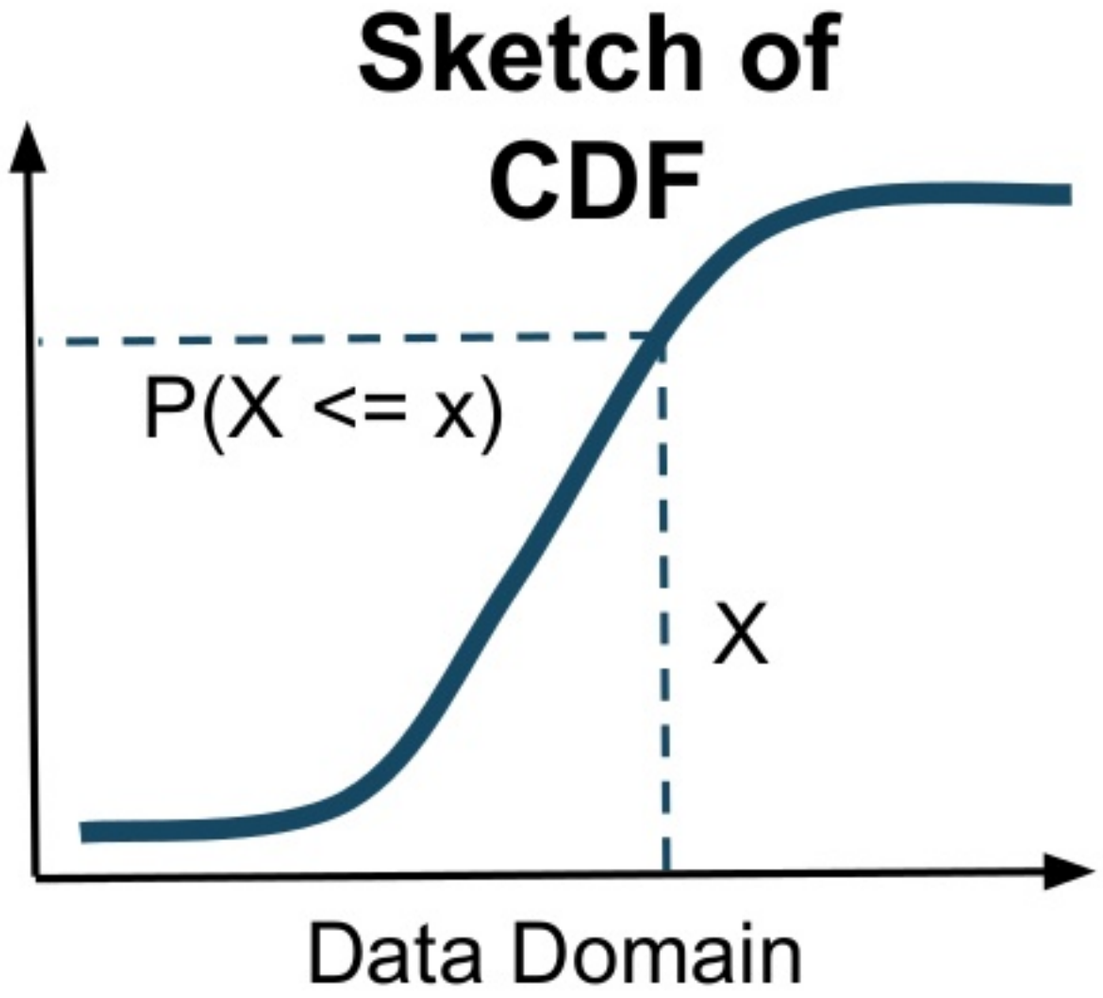

Inverse Transform Sampling

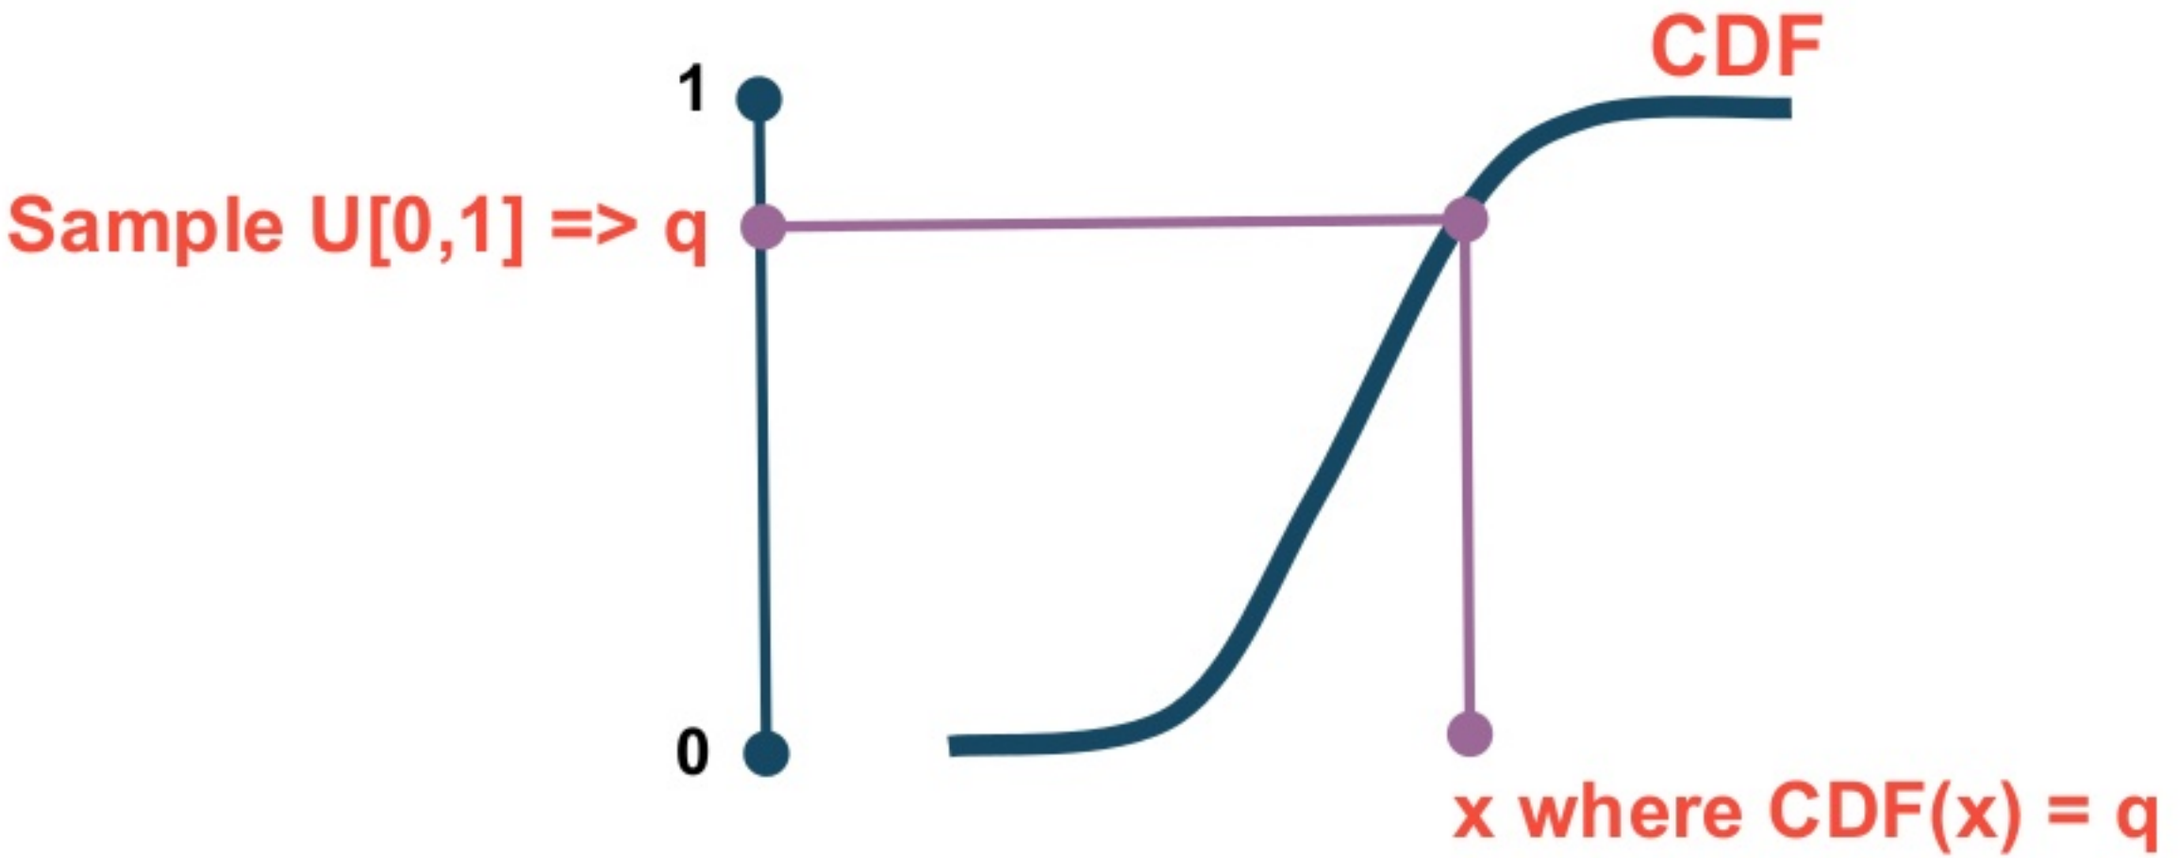

T-Digests Can Aggregate

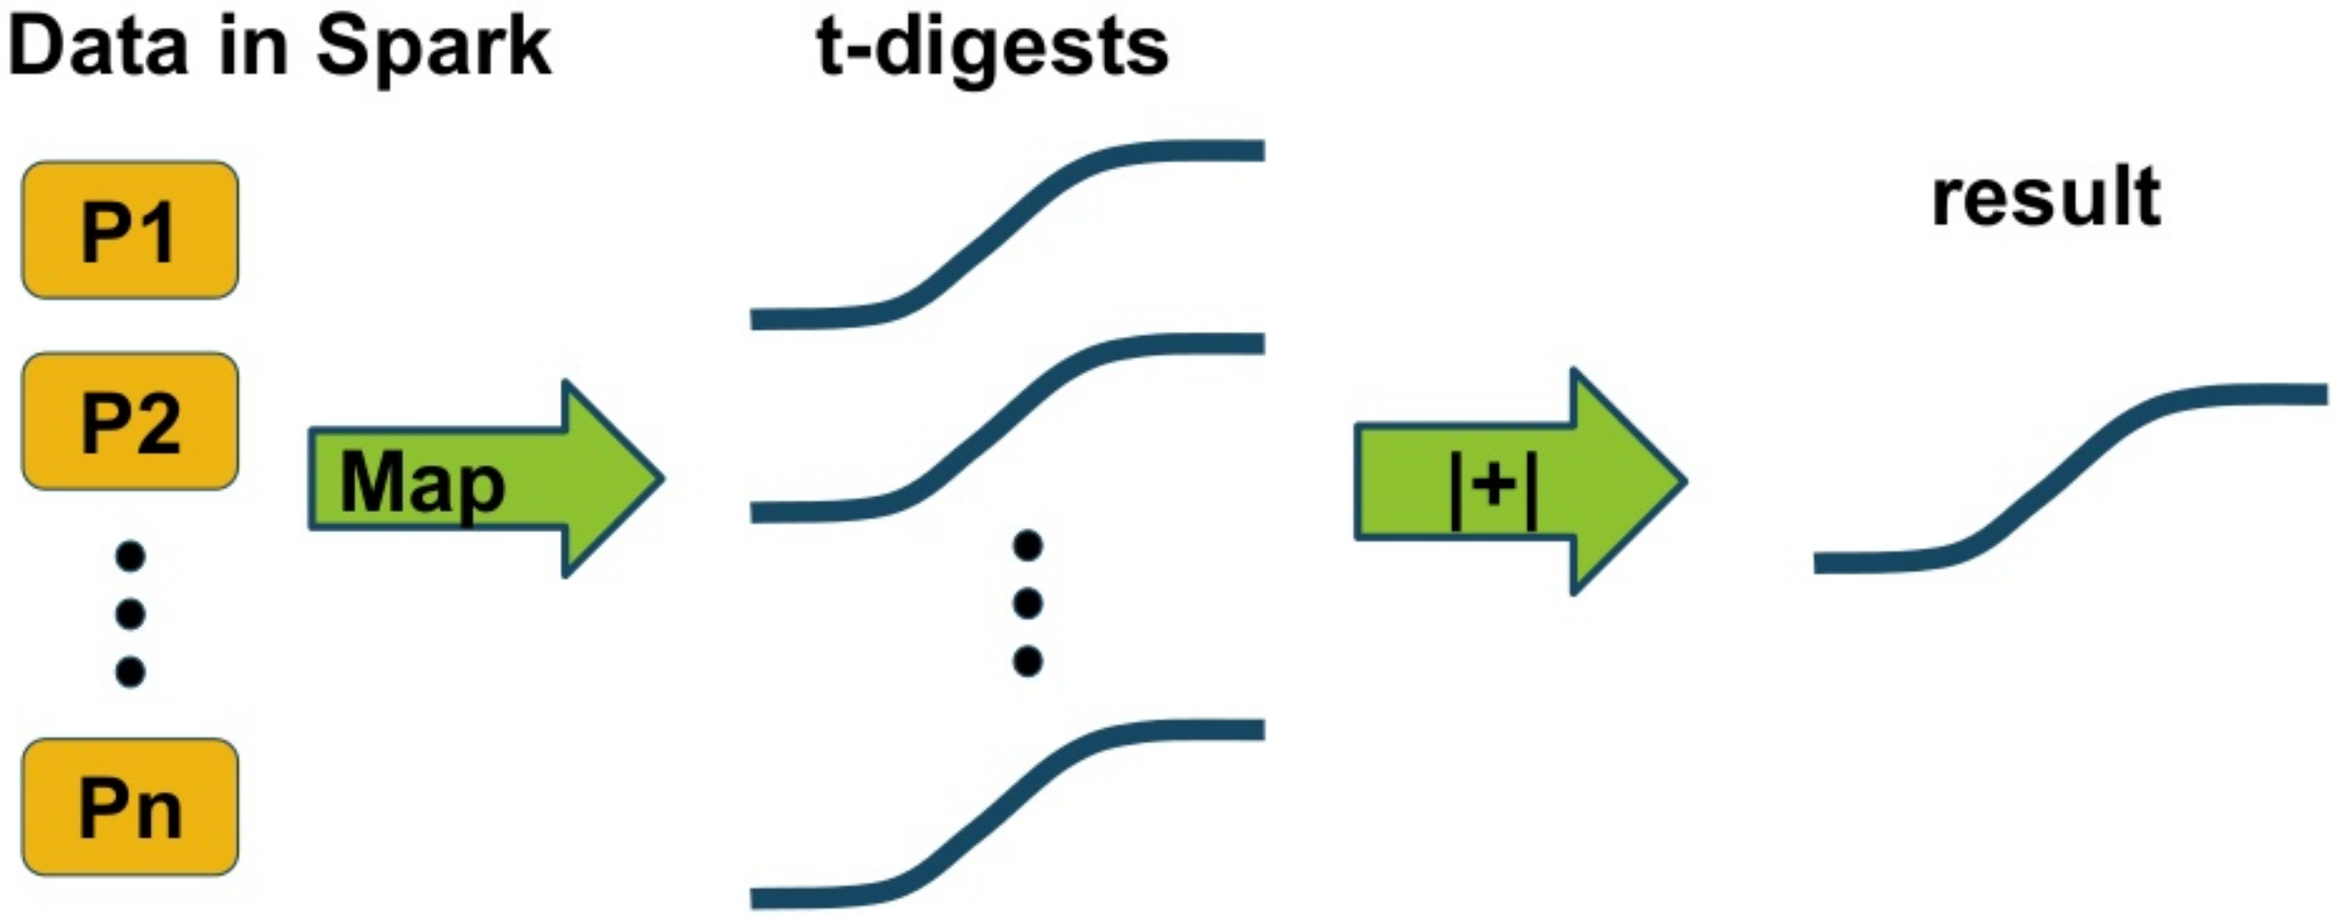

Sketching a Feature

```
feature.aggregate(TDigest.empty())(
    (td, x) => td + x,
    (td1, td2) => td1 ++ td2
)
```

Synthesizing Data from TDigests

```
def synthesize(tdVec: Vector[TDigest],
                    n: Int) = {
   val tdVecBC = sc.broadcast(tdVec)
   sc.parallelize(1 to n).map { _ =>
     tdVecBC.value.map(_.sample)
   }
}
```

Random Forest Training Data

```
val fvSketches = sketchFV(trainFV)
val synthFV = synthesize(fvSketches, 48000)
val trainLab = trainFV.map(_.toLabeledPoint(1.0))
val synthLab = synthFV.map(_.toLabeledPoint(0.0))
val trainFR = trainLab ++ synthLab
```

Random Forest Feature Reduction

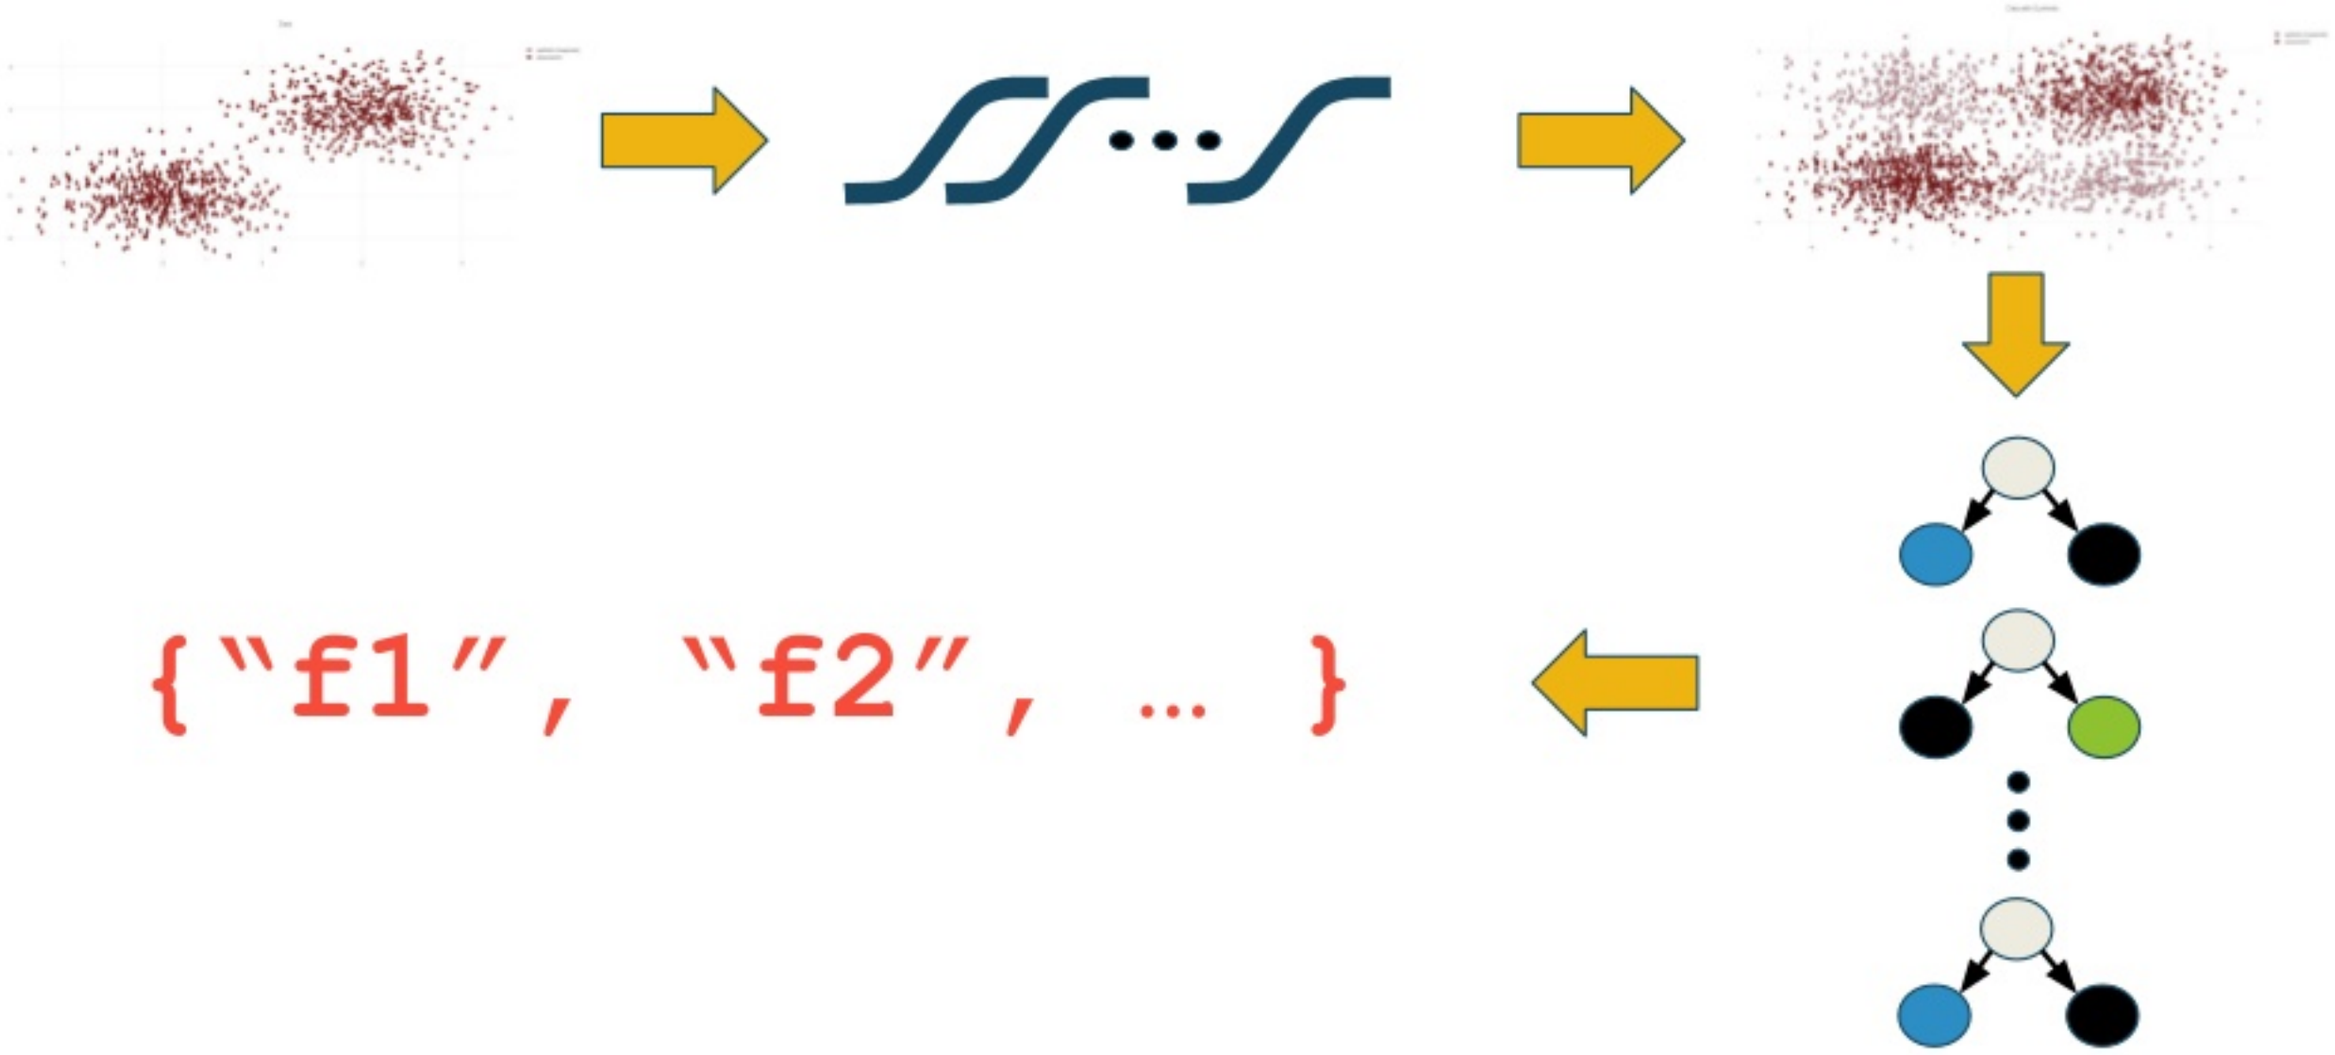

Tox21 Data Challenge

National Institute of Health (2014)

12 Toxicity Assays

12060 compounds + 647 hold-out

https://tripod.nih.gov/tox21/challenge/index.jsp

DeepTox

Johannes Kepler University Linz

Institute of Bioinformatics

http://bioinf.jku.at/research/DeepTox/tox21.html

[Mayr2016] Mayr, A., Klambauer, G., Unterthiner, T., & Hochreiter, S. (2016). DeepTox: Toxicity Prediction using Deep Learning. *Frontiers in Environmental Science*, **3**:80.

[Huang2016] Huang, R., Xia, M., Nguyen, D. T., Zhao, T., Sakamuru, S., Zhao, J., Shahane, S., Rossoshek, A., & Simeonov, A. (2016). Tox21Challenge to build predictive models of nuclear receptor and stress response pathways as mediated by exposure to environmental chemicals and drugs. *Frontiers in Environmental Science*, **3**:85.

Tox21 Data

801 Dense Features

272K Sparse Features

Each assay represented on a different subset

I used these

compound	NR.AhR	NR.AR	NR.AR.LBD	NR.Aromatase	...
NCGC00261900-01	0	1	NA	0	
NCGC00260869-01	0	1	NA	NA	.
NCGC00261776-01	1	1	0	NA	.
NCGC00261380-01	NA	0	NA	1	.
NCGC00261842-01	0	0	0	NA	
NCGC00261662-01	1	0	0	NA	
NCGC00261190-01	NA	0	0	NA	

Experiment

Train models on all 12 assays

Perform Random Forest Feature Reduction

Train similar models on reduced feature set

Compare models on each assay

85 of 801 Features Were Used

Features	Number trees used
RNCS	21
MRVSA7	20
VSAEstate2	19
VSAEstate3	18
slogPVSA8	18
VSAEstate0	17
slogPVSA6	16
RDFM29	12
slogPVSA3	12
RDFM30	12

Full vs Reduced (Logistic Reg)

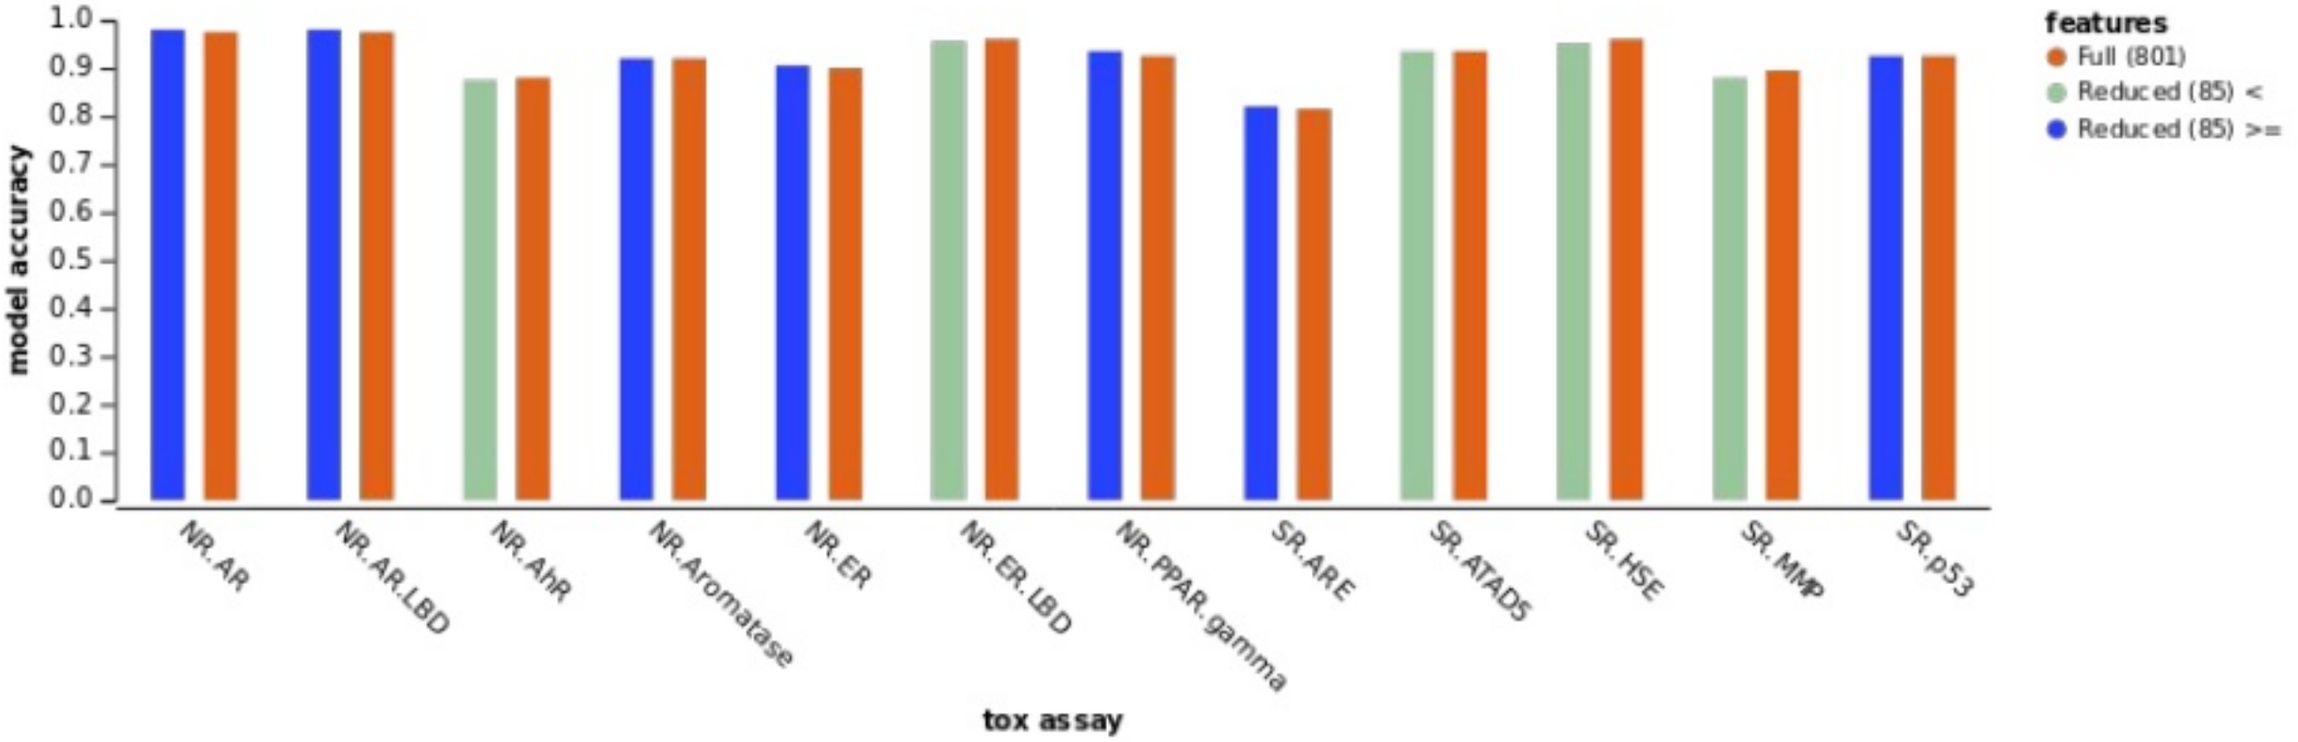

Full vs Reduced (Boosted DTE)

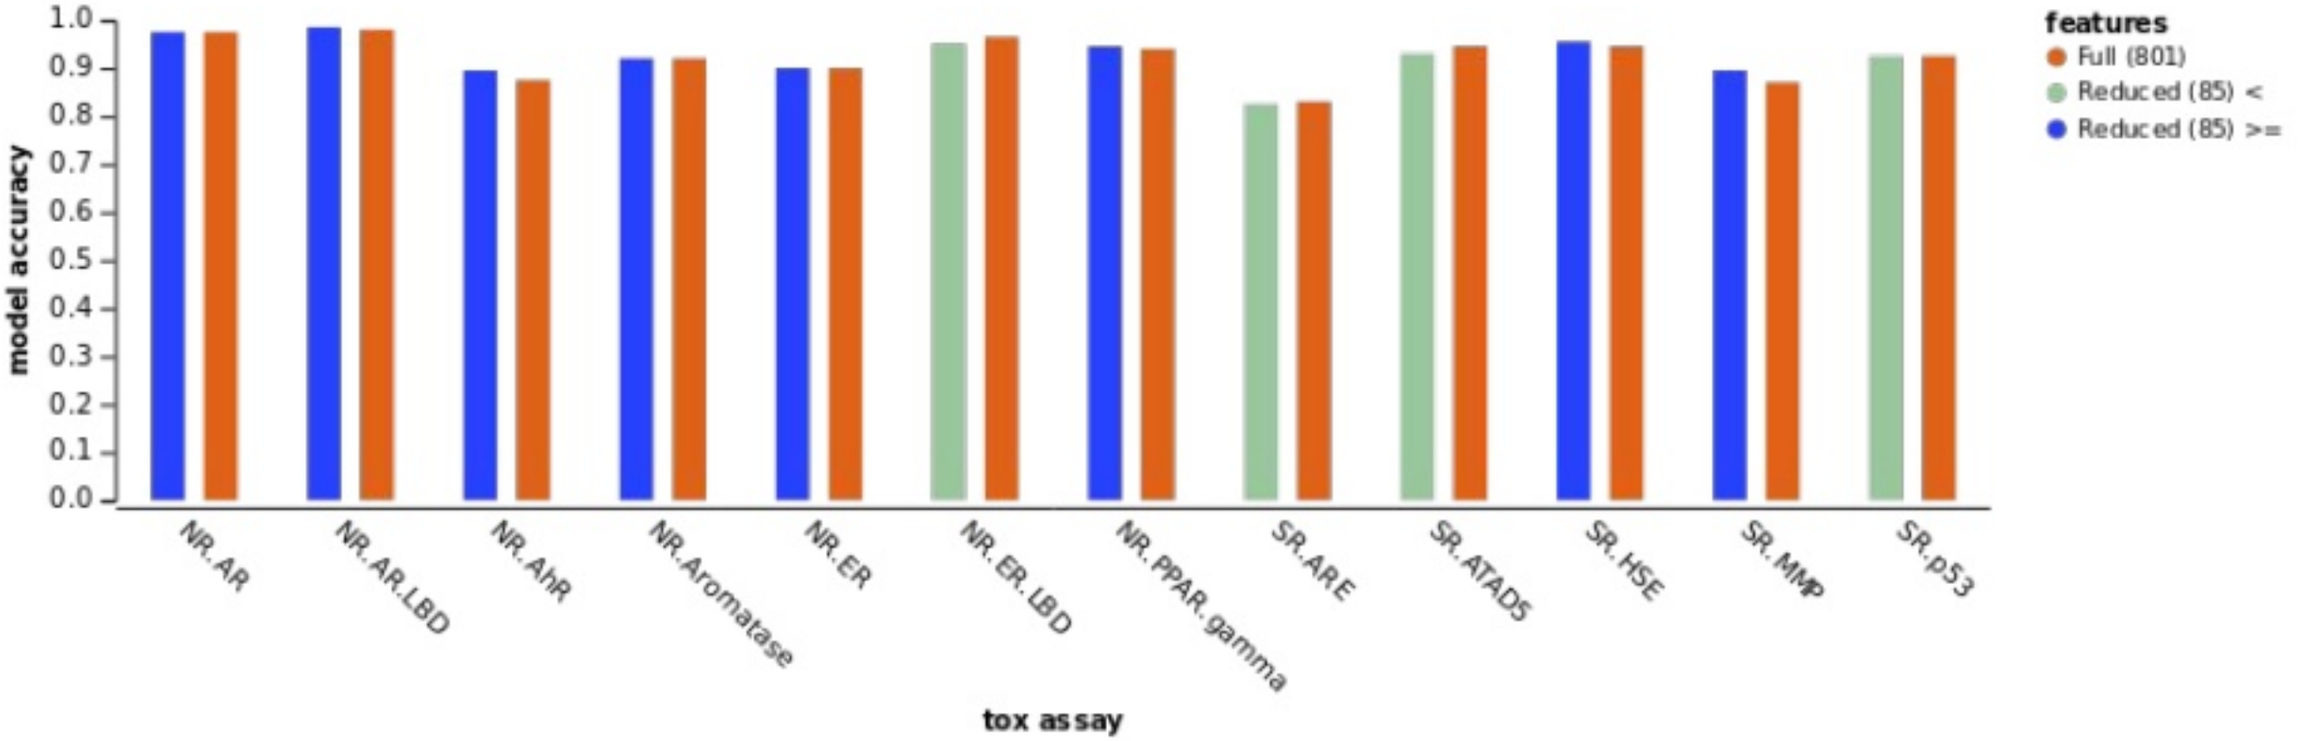

Full vs Reduced (SVM)

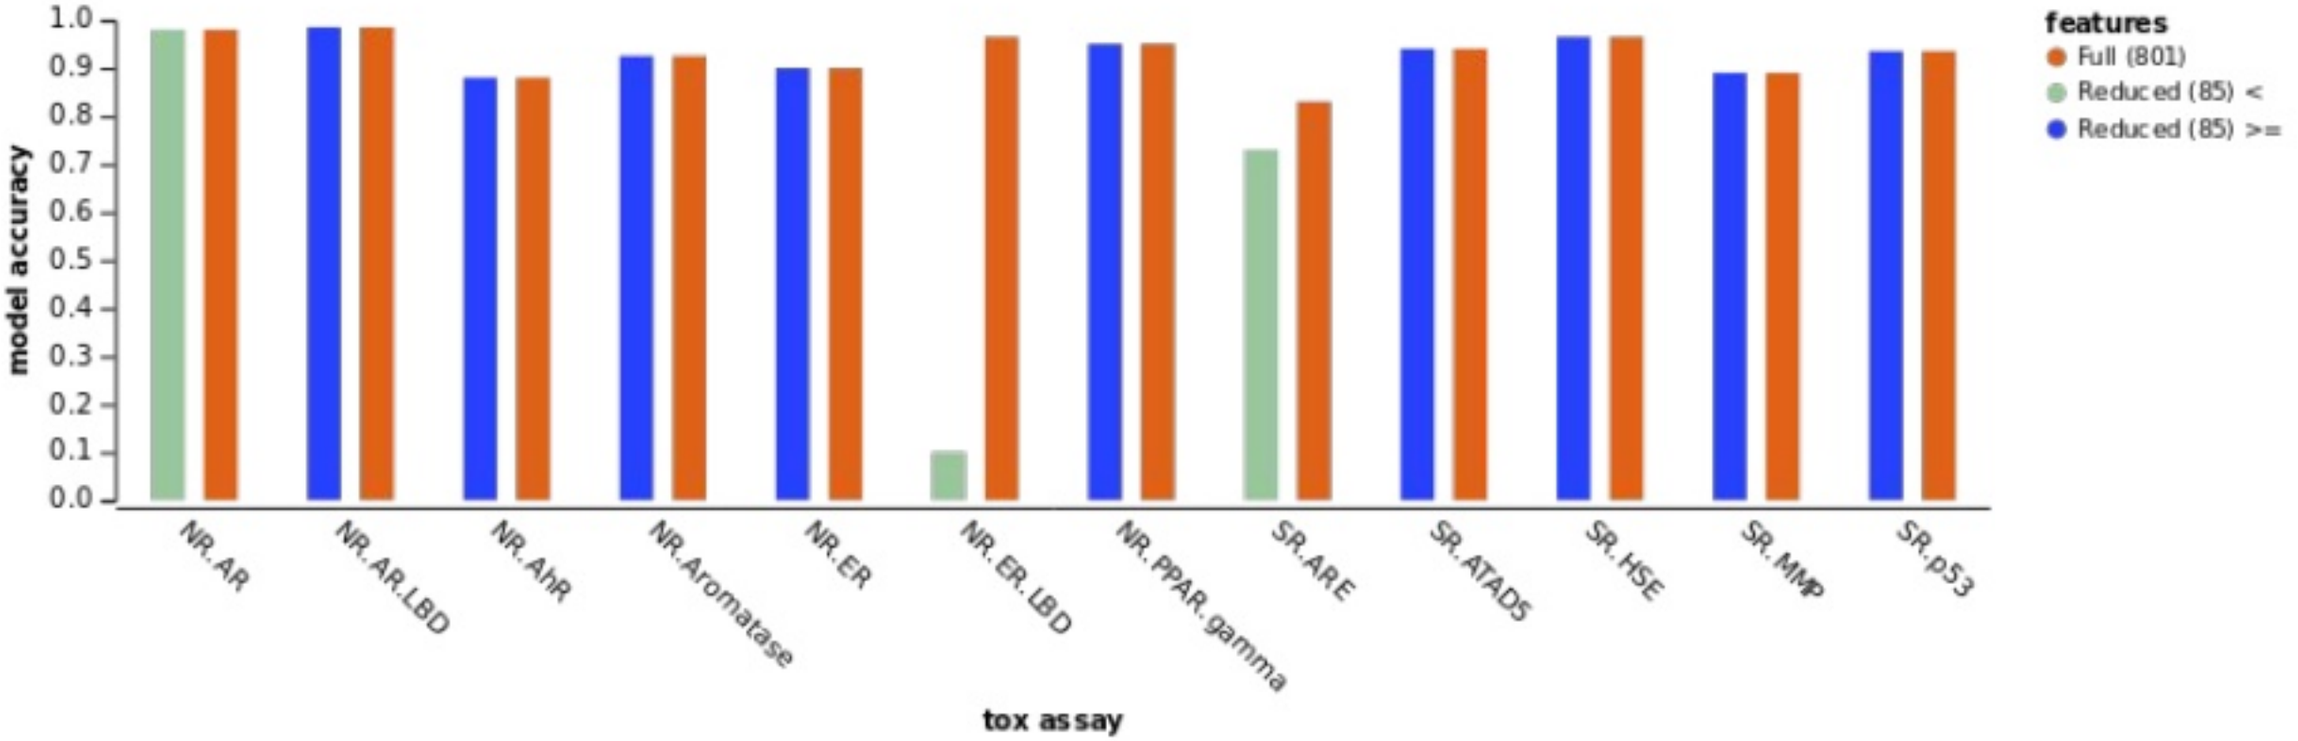

Training Times

(times in seconds)	Full (801)	Reduced (85)
Logistic Regression	68.5	46.8
SVM	35.3	33.8
GB Tree Ensemble	247	65.0

Evaluation Times

(times in seconds)	Full (801)	Reduced (85)
Logistic Regression	32.1	3.88
SVM	0.59	0.23
GB Tree Ensemble	1.33	0.88

Thank You

Erik Erlandson

eje@redhat.com

@manyangled

https://github.com/erikerlandson/feature-reduction-talk