



Spark: Data Science as a Service

Sridhar Alla, Kiran Muglurmamath
Comcast

Who we are

- **Sridhar Alla**
Director, Data Engineering, Comcast
Architecting and building solutions on a big data scale
sridhar_alla@cable.comcast.com
- **Kiran Muglurmath**
Executive Director, Data Science, Comcast
Data science on a big data scale.
kiran_muglurmath@cable.comcast.com



Agenda

- Why do this?
- Where are we now
- Real world challenges
- Introduction to Roadrunner – Our Solution to the real world challenges
- How we use Roadrunner in Comcast
- Q & A

Our Data

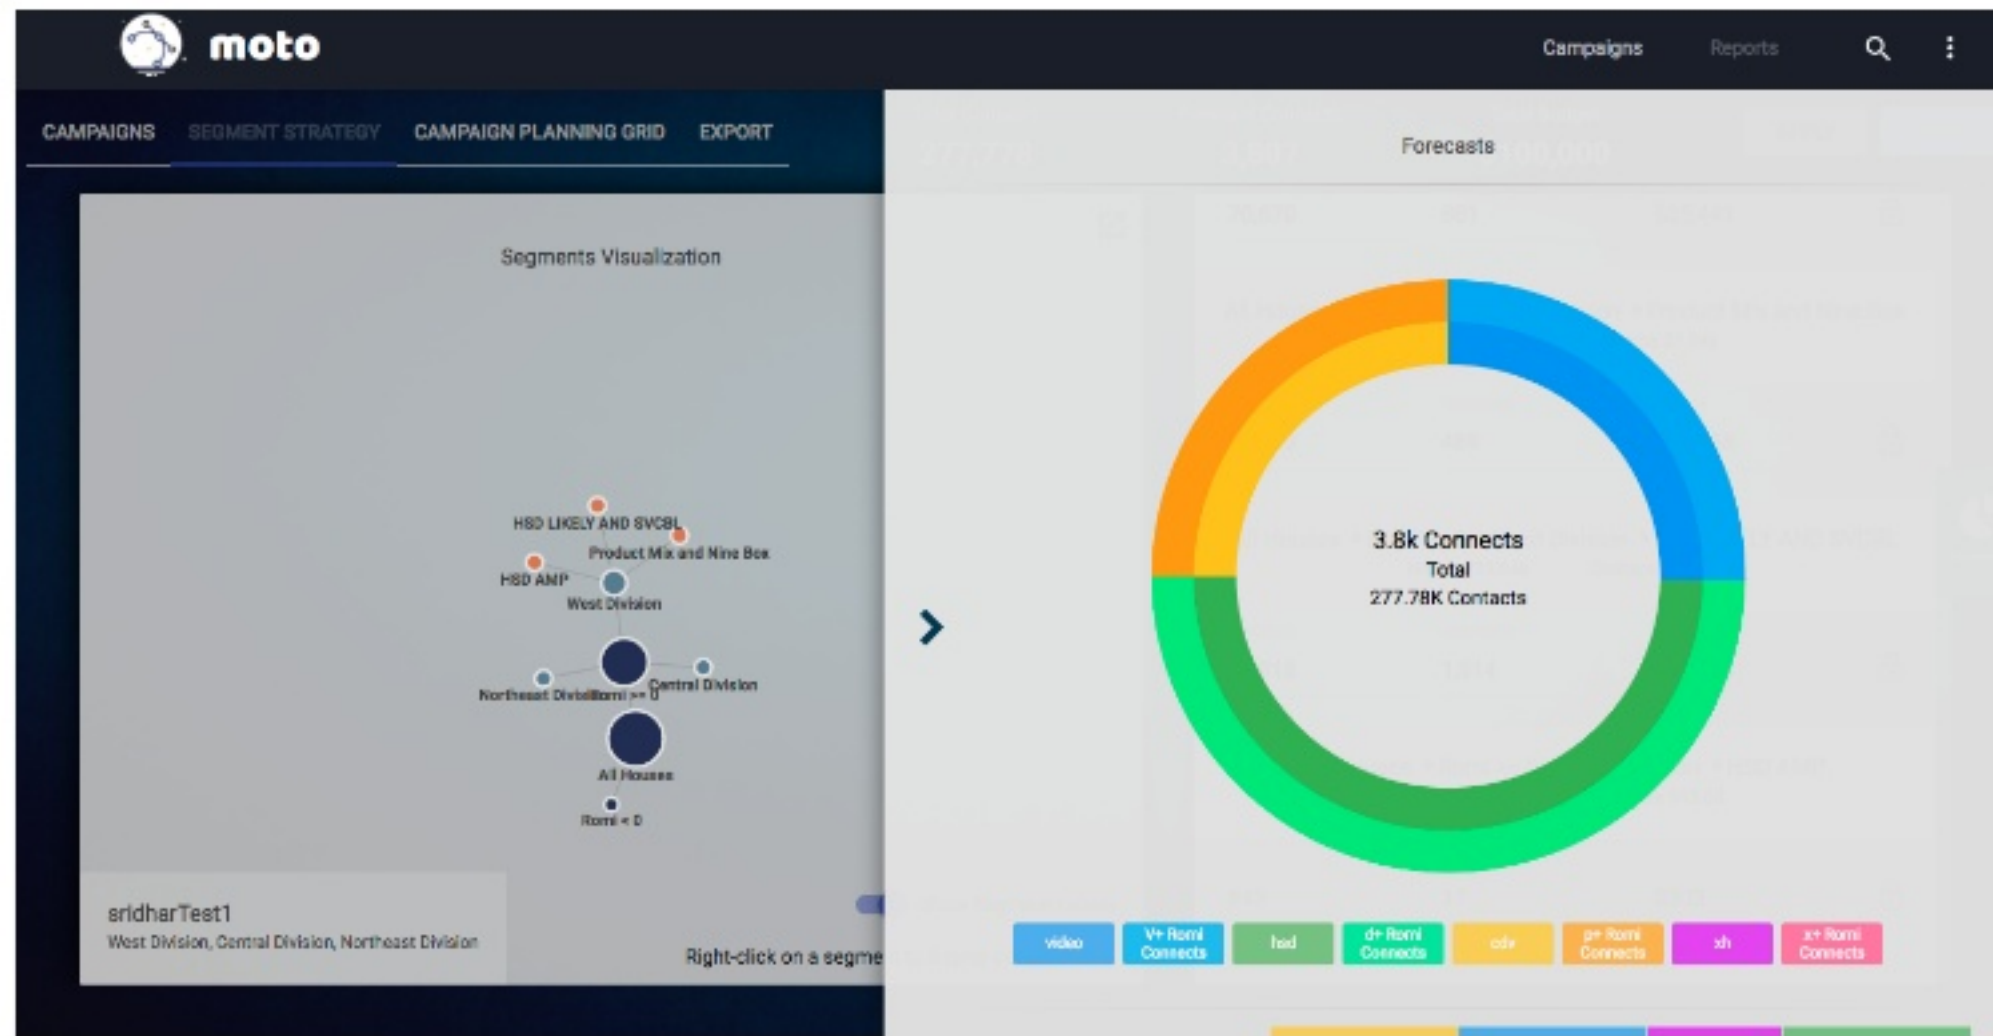
- 40PB in HDFS capacity and 100s of TBs in Teradata space
- ~1200 data nodes in total in Hadoop and Spark clusters
- Multiple 1Trillion+ row datasets
- Datasets with 12000+ columns
- 100s of models
 - Logistic regression, Neural Networks
 - LDA and other text analytics
 - Bayesian Networks
 - Clustering that includes kmeans, hierarchical, density
 - Geospatial

Data Science Use Cases

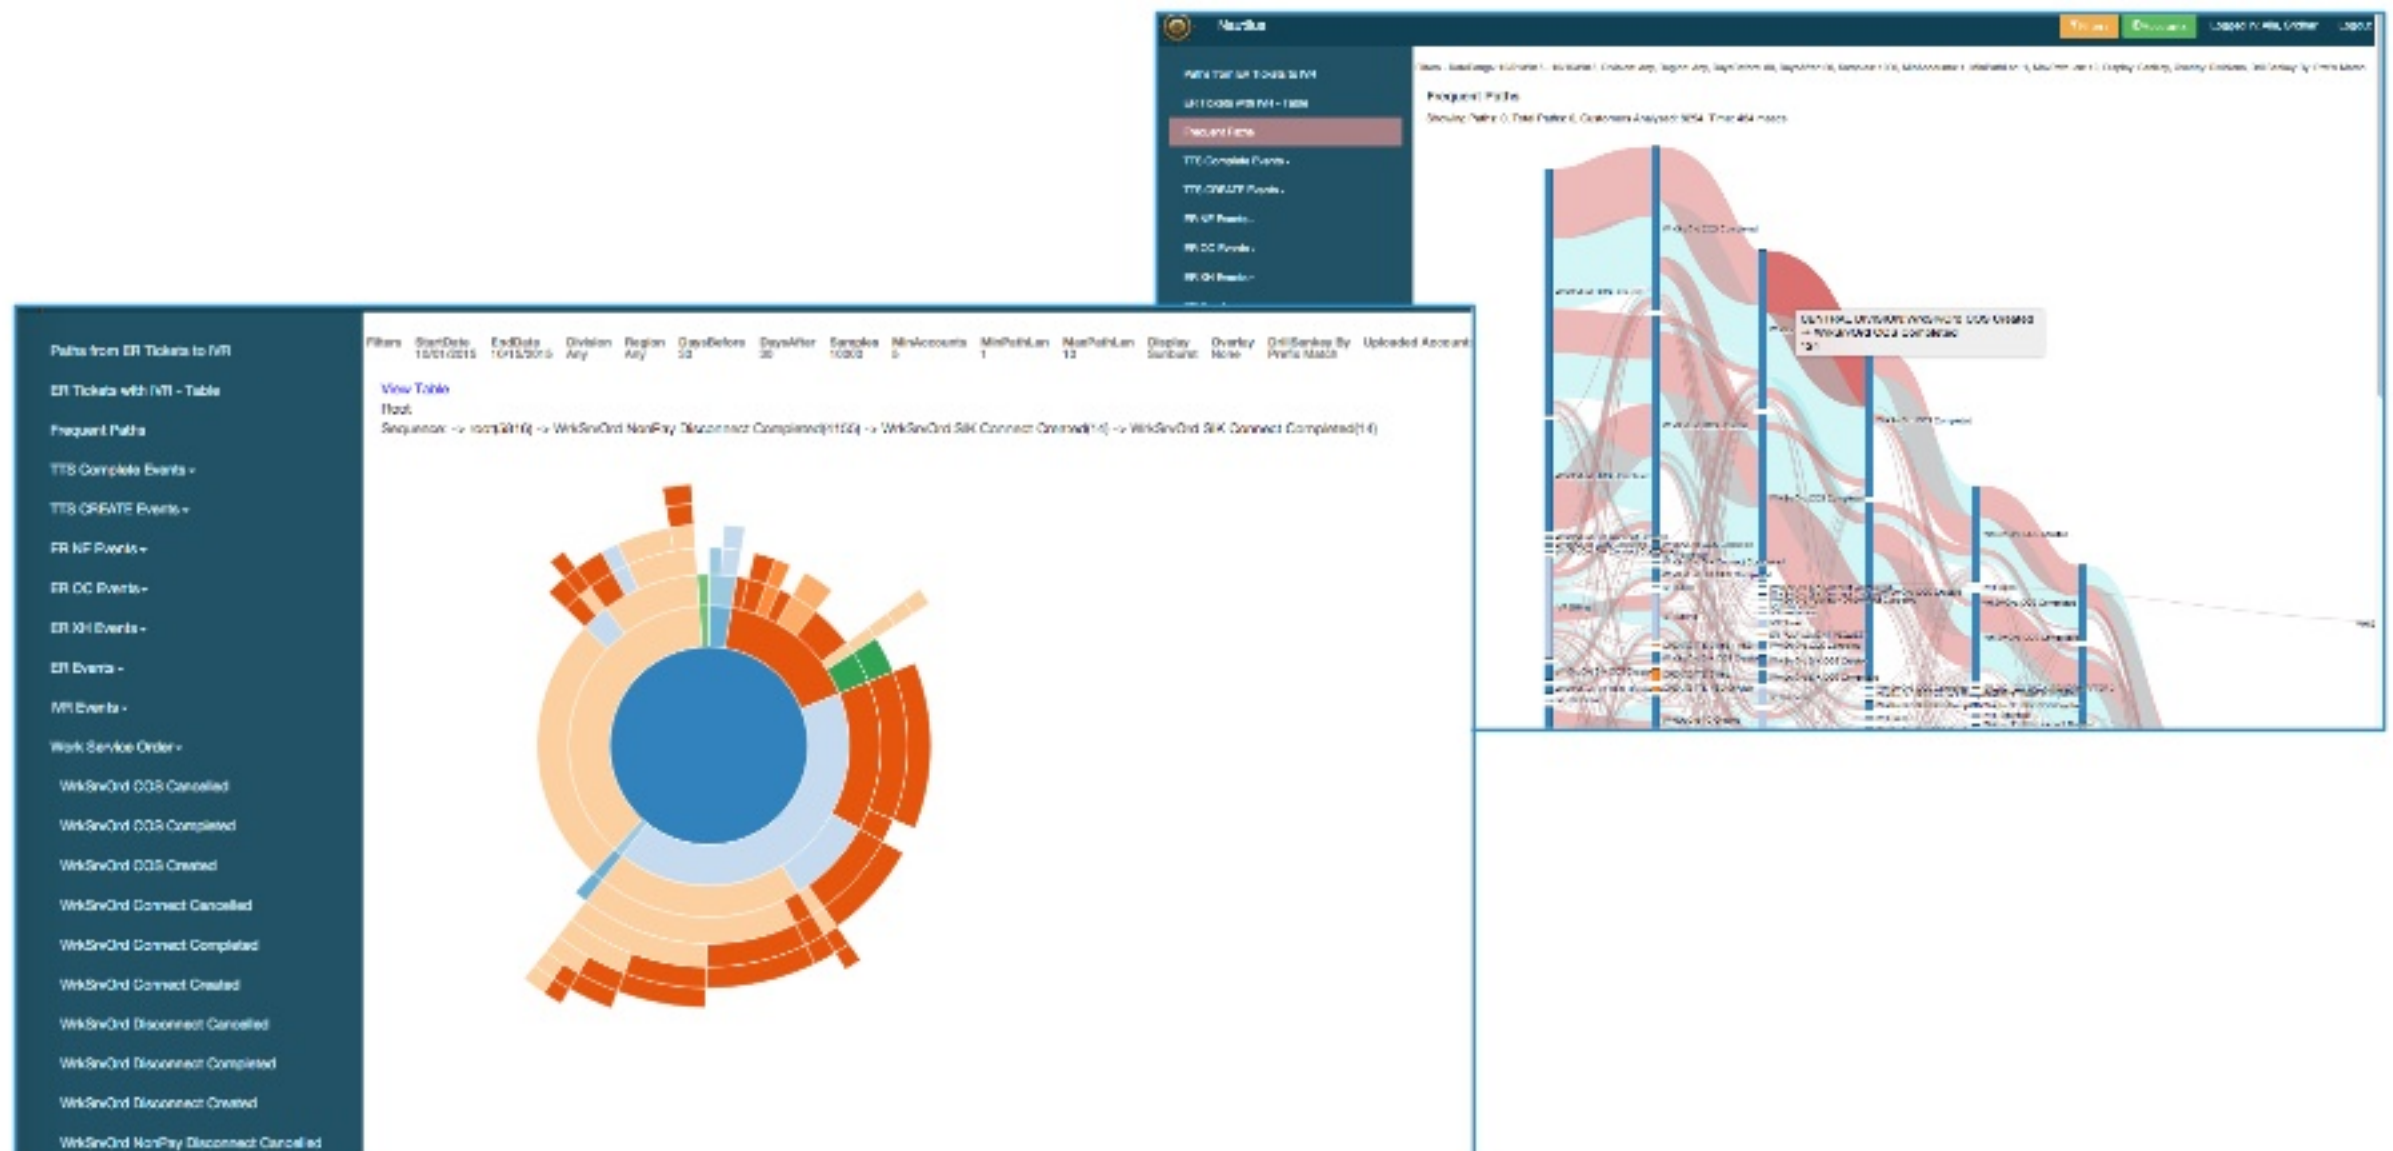
- Churn Models
- Price Elasticity
- Geo Spatial Route Optimization
- Direct Mail Campaign
- Customer call Analytics

many more

Direct Mail Campaign Optimization



Customer Journey Analytics



Main Challenges of Data Science

- Feature Engineering
 - Making sense of variety in data
- Model Scoring
 - Implementing ML algorithms
- Operational consumption for Business use cases

Main Challenges of Data Science

- Data ingestion, profiling and quality control
- We store and process massive amounts of data, still lack critical ability to stitch together pieces of data to make meaningful predictions. This is due to
 - Massive data size
 - Lack of service level architecture
- Multiple teams working on the same dataset
 - Increase development time because everyone has to process/feature engineer same dataset

What we needed

- A Central Processing System
 - Highly Scalable
 - Persisted and Cached
 - SQL capabilities and connection with multiple data sources and databases
 - Statistical Process Control methodology for data quality at every stage
 - Machine Learning capabilities and connection with multiple ML tools
 - Multi Tenancy
 - Access through APIs and programming languages
 - Fully automated workflow management for data science operations

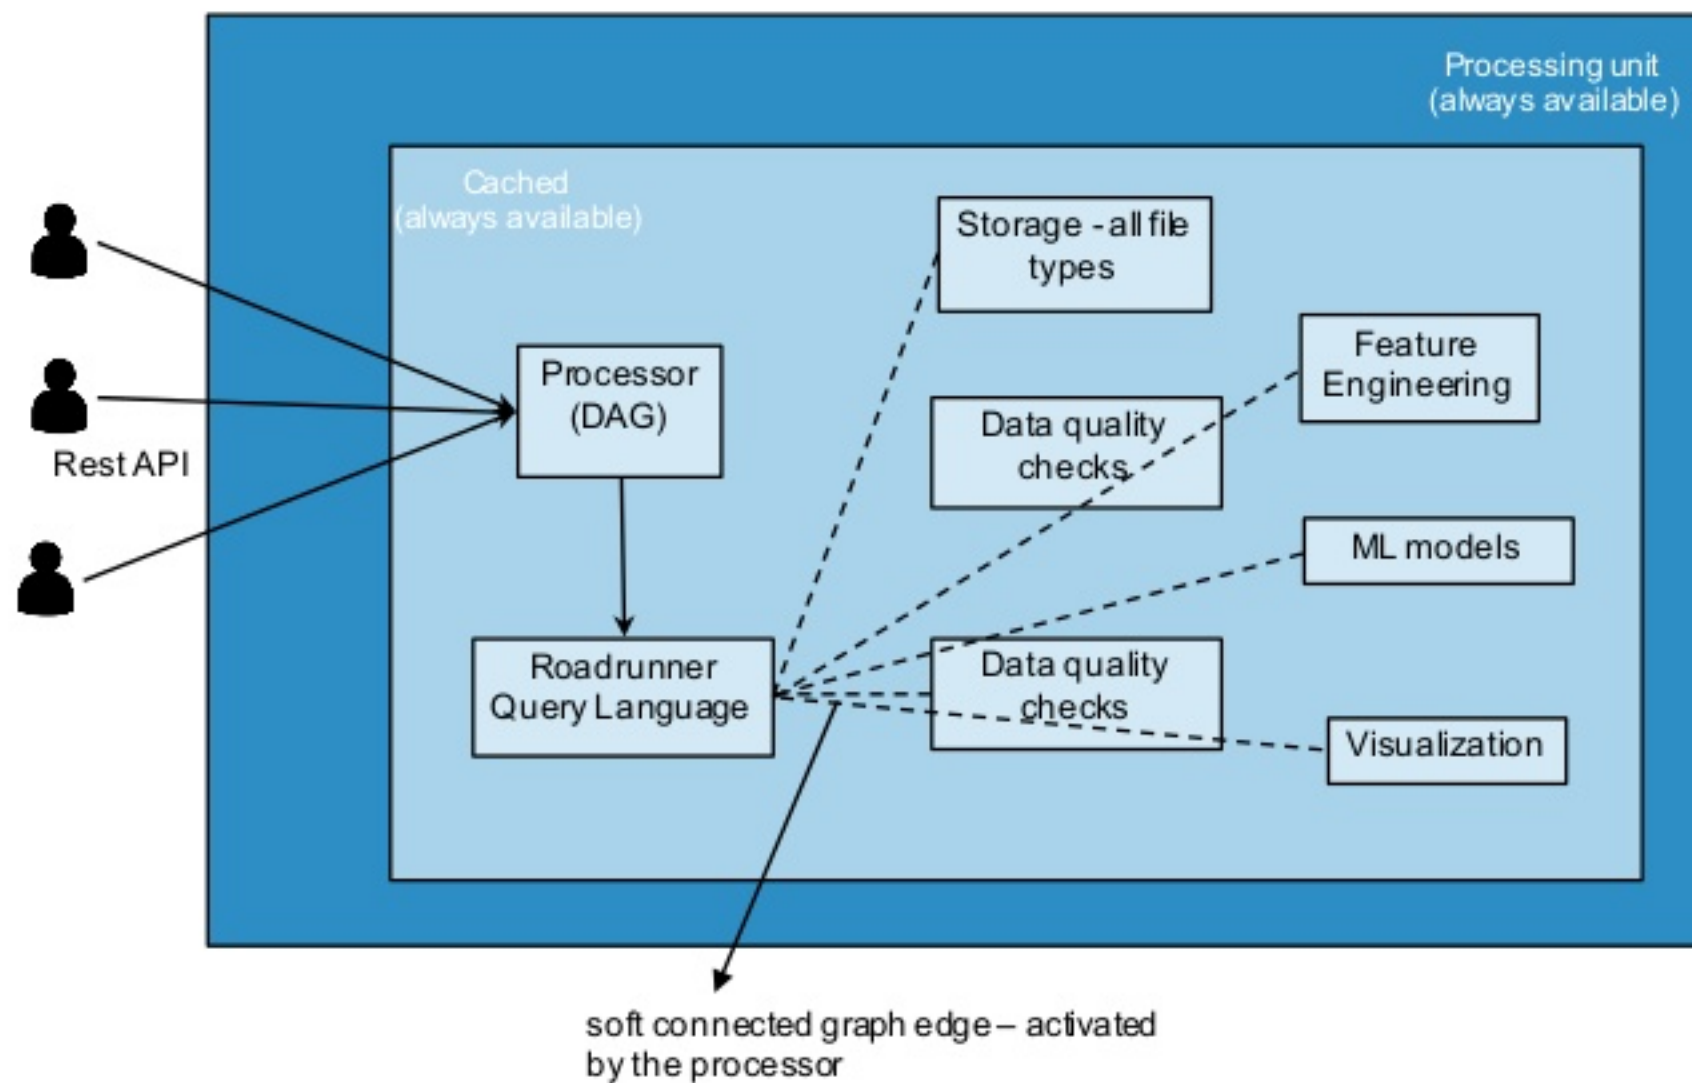
What we built

- Perpetual Spark Engine
- RESTful API to control all aspects
- Massively parallel quality control of petabyte scale datasets
 - Use Statistical Process Control methodology to check data at the record level
 - Parallelized data profilers on blind datasets
- Connectors to
 - Cassandra, Hbase, MongoDB, Teradata, MySQL, Hive, Elasticsearch, etc
 - Kafka, Storm for streaming data
 - ORC, Parquet, text files

What we built

- Role based control on who sees what
- Integration with modeling using Python, R, SAS, SparkML, H2O with language conversion tools
- Automated workflow management using graph methodology for data science

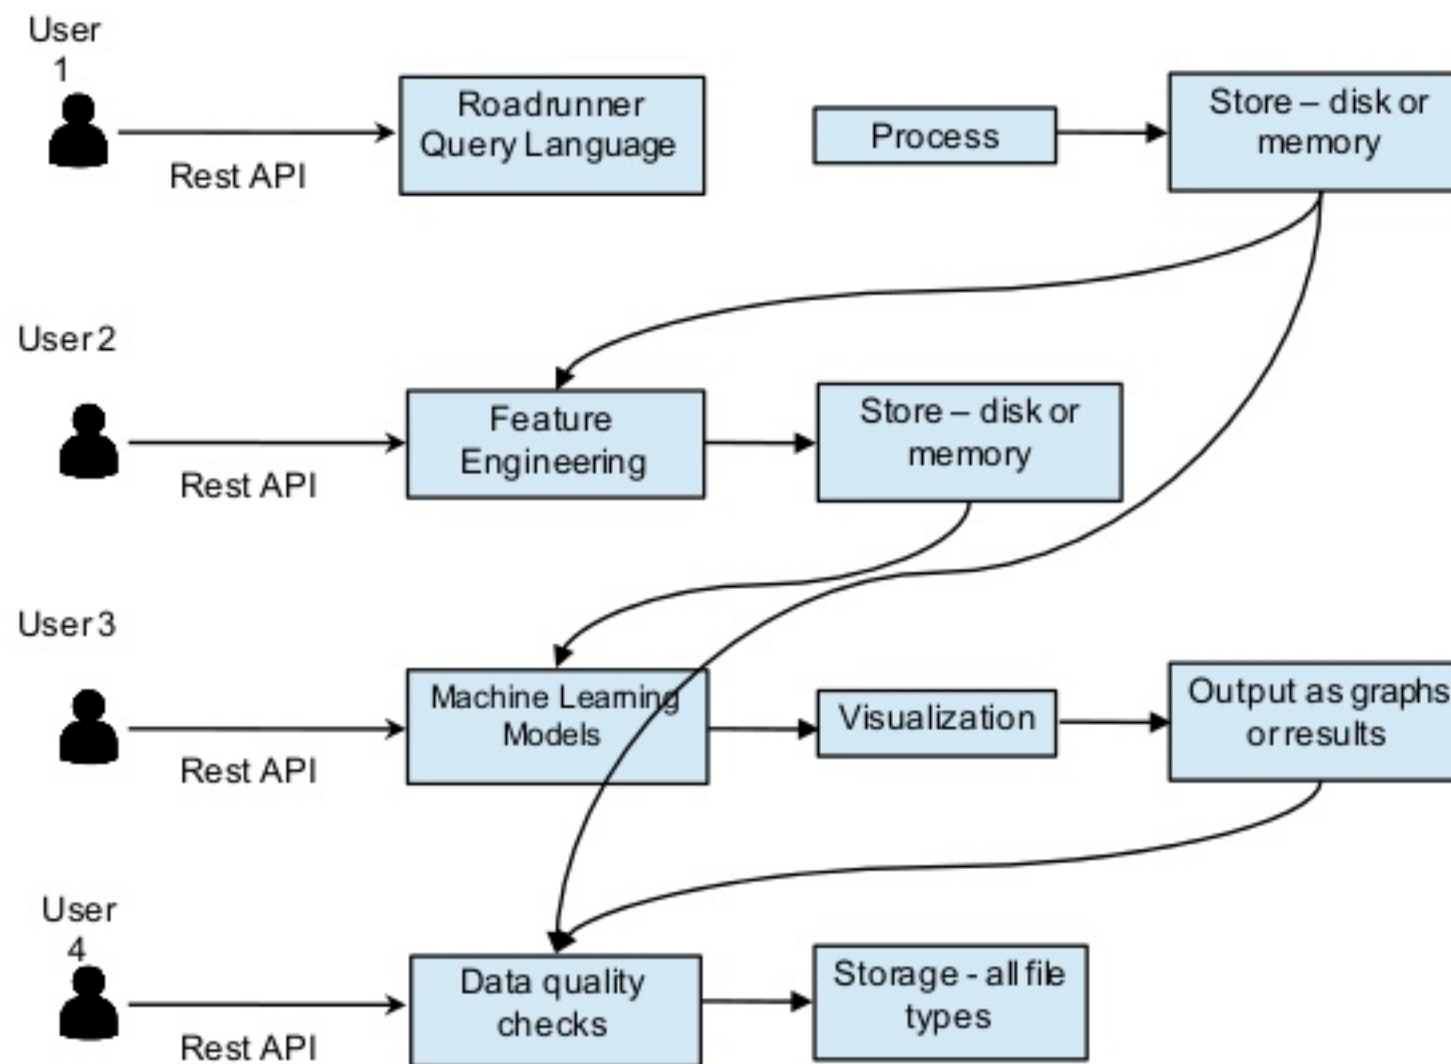
Roadrunner



Who can use Roadrunner

- Data Scientist
- DevOps
- Validation
- Modeler
- Engineer

How Roadrunner Works



Sample Rest API

Processing → Rest API

Feature Engineering → Rest API

Modeling → Rest API

```
{
  "jobType": "nautilusPathsJob",
  "jobId": "JobId4",
  "rosettaTableName": "base.adm_meld_201607",
  "startTime": "2016-01-01 00:00:00",
  "endTime": "2016-02-01 00:00:00",
  "eventId": "ANY",
  "appendToEventId": "",
  "minAccounts": 1,
  "accountFilters": "ALL",
  "eventRules": {
    "condition": "OR",
    "rules": [
      {
        "ruleType": 2,
        "firstEventId": "ER.*",
        "secondEventId": "IVR.*",
        "op": "gt",
        "threshold": 3,
        "timeGap": 166400,
        "generateRuleSequences": true,
        "overlappingSequences": true,
        "exactMatchingEvents": true
      }
    ]
  }
}
```

Sample Rest API

Processing → Rest API

Feature Engineering → Rest API

Modeling → Rest API

```
{  
  "jobType": "motoJob",  
  "jobId" : "moto1",  
  "campaignName" : "NE_Explore",  
  "campaignCondition":  
    {  
      "division": "NORTHEAST DIVISION",  
      "channels": "tbd_mailed"  
    },  
  "jobStage" : "updateNode",  
  "category": {"romi": ">=0", "connects": "2010"}  
}
```

Examples of Transformations

```
{
  "rules": [
    {"name": "tolowercase", "colname": "regionallowercase", "columns": ["SERLOC_CURRENT_REGION_NAME"]},
    {"name": "touppercase", "columns": ["SERLOC_CURRENT_DIVISION_NAME", "SERLOC_CURRENT_REGION_NAME"]},
    {"name": "concat", "columns": ["SERLOC_CURRENT_DIVISION_NAME", "SERLOC_CURRENT_REGION_NAME"]},
    {"name": "arithexpr", "expr": "(DAYSSINCE + (DAYSSINCE * CURRDAYSINCE))", "columns": ["DAYSSINCE", "CURRDAYSINCE"]},
    {"name": "zscore", "columns": ["DAYSSINCE"]},
    {"name": "filter", "filters": ["DAYSSINCE > 122"]},
    {"name": "decile", "colname": "DN", "scoreColumn": "DAYSSINCE", "columns": ["ACCOUNTSTATUS"]},
    {"name": "nothing", "columns": [""]}
  ]
}
```

Examples of Joins

```
{
  "rules": [
    {"name": "innerjoin", "joins": ["ACCOUNTSTATUS, astatusname"]},
    {"name": "leftouterjoin", "joins": ["ACCOUNTSTATUS, astatusname"]},
    {"name": "rightouterjoin", "joins": ["ACCOUNTSTATUS, astatusname"]},
    {"name": "outerjoin", "joins": ["ACCOUNTSTATUS, astatusname"]},
    {"name": "nothing", "joins": [""]}
  ]
}
```

Examples of Joins

```
{
  "results": [
    {
      "joinType": "inner",
      "joinTotal": 301,
      "leftTotal": 1000,
      "rightTotal": 2,
      "leftNulls": 0,
      "rightNulls": 0
    },
    {
      "joinType": "leftouter",
      "joinTotal": 1000,
      "leftTotal": 1000,
      "rightTotal": 2,
      "leftNulls": 383,
      "rightNulls": 699
    },
    {
      "joinType": "rightouter",
      "joinTotal": 302,
      "leftTotal": 1000,
      "rightTotal": 2,
      "leftNulls": 1,
      "rightNulls": 0
    },
    {
      "joinType": "outer",
      "joinTotal": 1001,
      "leftTotal": 1000,
      "rightTotal": 2,
      "leftNulls": 384,
      "rightNulls": 699
    }
  ]
}
```

Examples of Aggregations

```
{
  "rules": [ {
    "functions" : [ {
      "name" : "approx_count_distinct"
    }, {
      "name" : "histogram_string",
      "buckets" : 20
    }
  ],
  "groupBy": "SNAPSHOTDATE",
  "columns" : [ "SERLOC_CURRENT_DIVISION_NAME" ]
},
{
  "functions" : [ {
    "name" : "approx_count_distinct"
  }, {
    "name" : "min"
  }, {
    "name" : "var_pop"
  }, {
    "name" : "sum"
  }, {
    "name" : "percentile_approx",
    "percentiles" : [ 0.25, 0.5, 0.75, 0.9, 0.95, 0.99 ]
  }
  ],
  "groupBy": "SNAPSHOTDATE,SERLOC_CURRENT_DIVISION_NAME,SERLOC_CURRENT_REGION_NAME",
  "columns" : [ "DAYSSINCE" ]
}
]
```


Examples of Aggregations

```
{
  "results": {
    "approx_count_distinct(6)": 10,
    "min(6)": "174.56",
    "max(6)": "995.09",
    "avg(6)": 566.268,
    "count(6)": 10,
    "first(6)": "706.99",
    "last(6)": "995.09",
    "kurtosis(6)": -1.4016327462769833,
    "skewness(6)": 0.12851131741001862,
    "stddev(6)": 298.53423354040245,
    "stddev_pop(6)": 283.21444125609133,
    "variance(6)": 89122.68859555555,
    "var_pop(6)": 80210.419736,
    "sum(6)": 5662.68,
    "percentile_approx(CAST(6 AS DOUBLE), array(0.25, 0.5, 0.75, 0.9, 0.95, 0.99))": [
      272.63,
      506.58,
      885.62,
      911.8,
      995.09,
      995.09
    ]
  }
}
```

Examples of Aggregations

```
{
  "Aggregations" : [ {
    "SNAPSHOTDATE" : "20170201",
    "SERLOC_CURRENT_DIVISION_NAME" : "NORTHEAST DIVISION",
    "SERLOC_CURRENT_REGION_NAME" : "KEYSTONE REGION",
    "approx_count_distinct(SERLOC_CURRENT_DIVISION_NAME)" : 1,
    "approx_count_distinct(DAYSSINCE)" : 33,
    "stddev(DAYSSINCE)" : 3611.7773139718893,
    "stddev_pop(DAYSSINCE)" : 3565.171783171744,
    "variance(DAYSSINCE)" : 1.3044935365721995E7,
    "var_pop(DAYSSINCE)" : 1.2710449843523994E7,
    "sum(DAYSSINCE)" : 84277,
    "percentile_approx(CAST(DAYSSINCE AS DOUBLE), array(0.25, 0.5, 0.75, 0.9, 0.95, 0.99))" : [ 112.0, 552.0, 2218.0,
  ], {
    "SNAPSHOTDATE" : "20170101",
    "SERLOC_CURRENT_DIVISION_NAME" : "NORTHEAST DIVISION",
    "SERLOC_CURRENT_REGION_NAME" : "BELTWAY REGION",
    "approx_count_distinct(SERLOC_CURRENT_DIVISION_NAME)" : 1,
    "approx_count_distinct(DAYSSINCE)" : 28,
    "stddev(DAYSSINCE)" : 1183.8048117688786,
    "stddev_pop(DAYSSINCE)" : 1170.8666657885606,
    "variance(DAYSSINCE)" : 1401393.83236715,
    "var_pop(DAYSSINCE)" : 1370928.7490548207,
    "sum(DAYSSINCE)" : 28387,
    "percentile_approx(CAST(DAYSSINCE AS DOUBLE), array(0.25, 0.5, 0.75, 0.9, 0.95, 0.99))" : [ 0.0, 60.0, 810.0, 1738
  ] },
  "customAggregations" : [ {
    "histogram_string(SERLOC_CURRENT_DIVISION_NAME)" : {
      "NORTHEAST DIVISION" : 364,
      "CENTRAL DIVISION" : 435,
      "WEST DIVISION" : 435
    }
  } ]
}
```

Deciles – Spark + Scala

```
val filters =  
  dfTransformed  
    .groupBy(column)  
    .count  
    .distinct  
    .rdd  
    .map(r => r.getString(0))  
    .collect  
  
val rdds = for { f <- filters } yield {  
  val dfTmp = {  
    if (f == null)  
      dfTransformed.filter(col(column).isNull)  
    else  
      dfTransformed.filter(col(column) === f)  
  }  
  
  val bw = Window.partitionBy(column).orderBy(col(scoreColumn).desc)  
  val df2 = dfTmp.select(col("*"), ntile(10).over(bw).alias(colName.getOrElse("decile")))  
}
```

Deciles – the Roadrunner way...

```
{  
  "rules": [  
    {"name": "decile", "colname": "DN", "scoreColumn": "DAYSSINCE", "columns": ["ACCOUNTSTATUS"]}  
  ]  
}
```

Grouped Aggregations – easy?

```
val (aggColumnFunctions, toCalculate) = columnFunctions.partition(
  _.function.isDefined
) //if a columnFunction has a function=None then it is a custom function and cannot be handled by the `df.agg` call
val aggFunctions = aggColumnFunctions.flatMap(_.function)

def columnNames(funcs: Seq[ColumnFunction]) =
  funcs.flatMap(_.columnNames).distinct.map(col)

def groupByColumnsFunc(funcs: Seq[ColumnFunction]) =
  funcs.flatMap(_.groupBy.getOrElse("").split(",")).map(x => x.trim).distinct.map(col)

val groupByColumns = groupByColumnsFunc(aggColumnFunctions)

val resultsF = Future(blocking {
  aggFunctions.isEmpty.fold(
    { List() }, {
      logger.debug(s"Running df.agg(${aggFunctions.mkString(",")})")
      if (groupByColumns.isEmpty) {
        df.select(columnNames(aggColumnFunctions): _*)
          .agg(aggFunctions.head, aggFunctions.tail: _*)
          .toJSON
          .collectAsList()
          .toList
      } else {
        df.groupBy(groupByColumnsFunc(aggColumnFunctions): _*)
          .agg(aggFunctions.head, aggFunctions.tail: _*)
          .toJSON
          .collectAsList()
          .toList
      }
    }
  )
})
```


Grouped Aggregations – easy?

```
val (aggColumnFunctions, toCalculate) = columnFunctions.partition(
  _.function.isDefined
) //if a columnFunction has a function=None then it is a custom function and cannot be handled by the `df.agg` call
val aggFunctions = aggColumnFunctions.flatMap(_.function)

def columnNames(funcs: Seq[ColumnFunction]) =
  funcs.flatMap(_.columnNames).distinct.map(col)

def groupByColumnsFunc(funcs: Seq[ColumnFunction]) =
  funcs.flatMap(_.groupBy.getOrElse("").split(",")).map(x => x.trim).distinct.map(col)

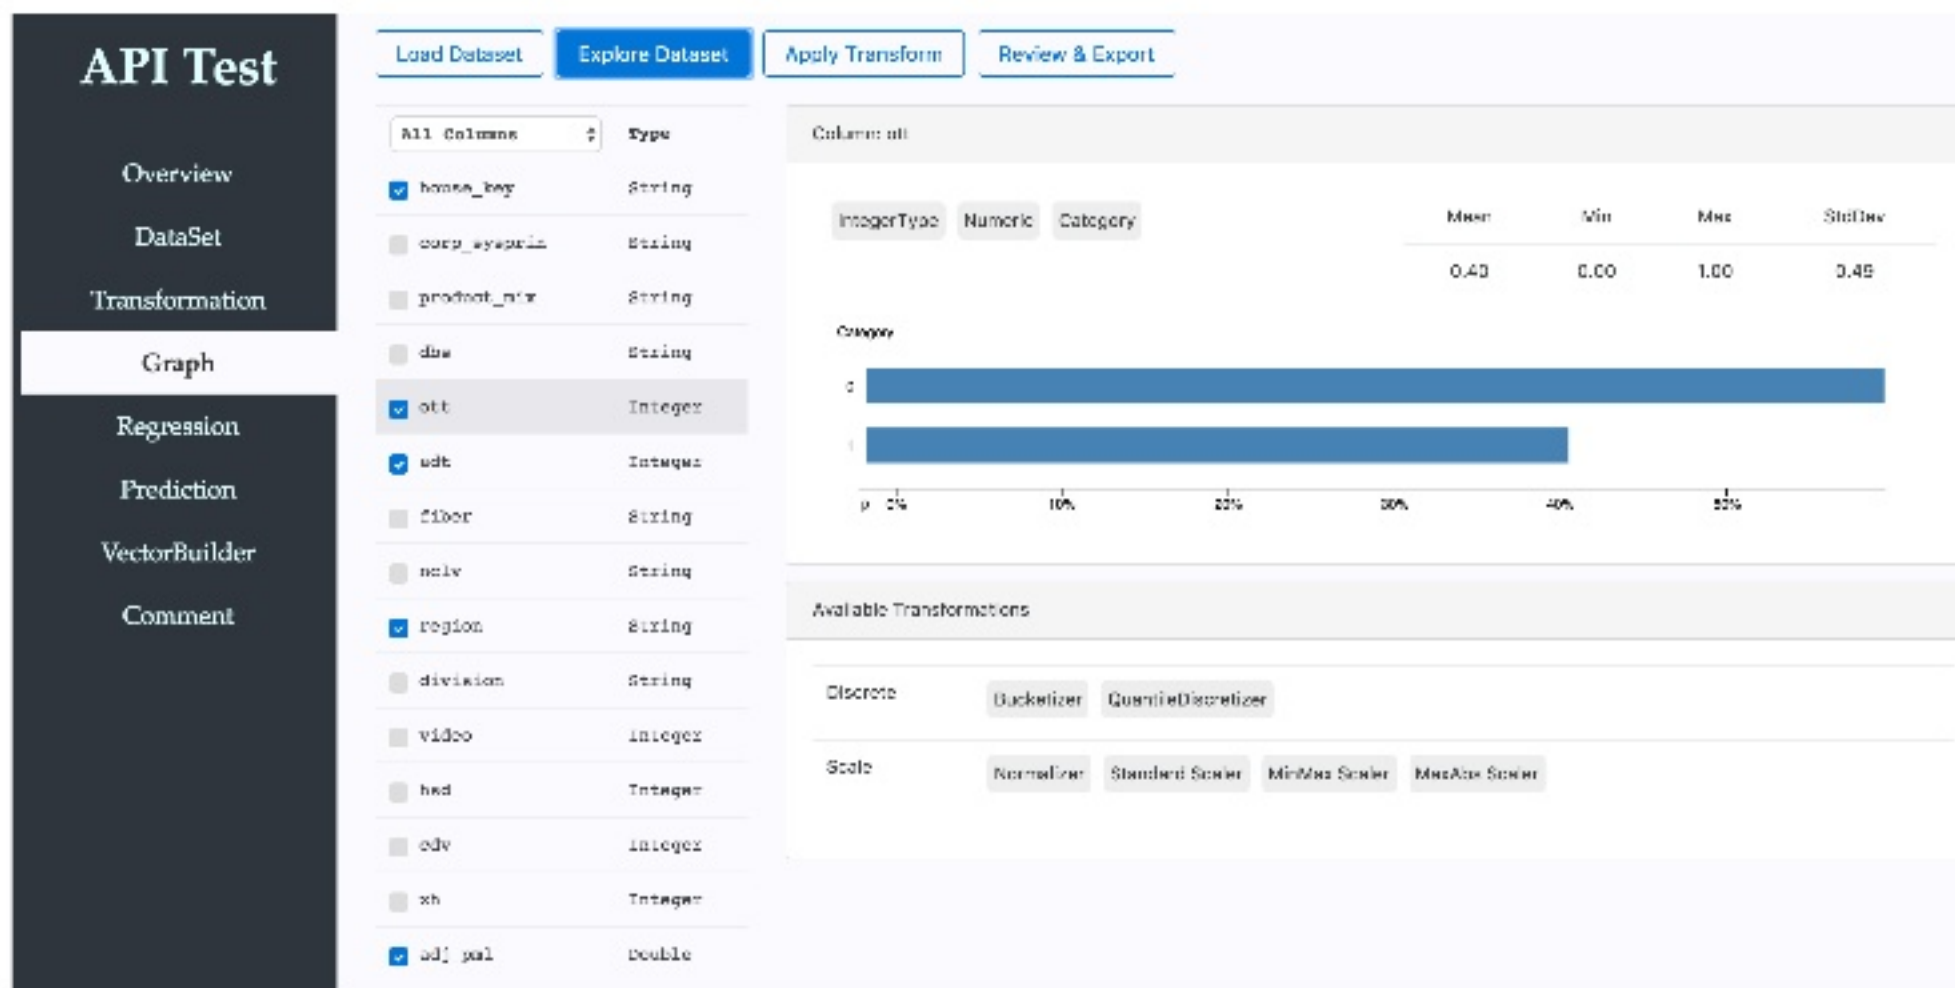
val groupByColumns = groupByColumnsFunc(aggColumnFunctions)

val resultsF = Future(blocking {
  aggFunctions.isEmpty.fold(
    { List() }, {
      logger.debug(s"Running df.agg(${aggFunctions.mkString(",")})")
      if (groupByColumns.isEmpty) {
        df.select(columnNames(aggColumnFunctions): _*)
          .agg(aggFunctions.head, aggFunctions.tail: _*)
          .toJSON
          .collectAsList()
          .toList
      } else {
        df.groupBy(groupByColumnsFunc(aggColumnFunctions): _*)
          .agg(aggFunctions.head, aggFunctions.tail: _*)
          .toJSON
          .collectAsList()
          .toList
      }
    }
  )
})
```


Grouped Aggregations - the Roadrunner way

```
{
  "rules": [ {
    "functions" : [ {
      "name" : "approx_count_distinct"
    }, {
      "name" : "histogram_string",
      "buckets" : 20
    }
  ],
  "groupBy": "SNAPSHOTDATE",
  "columns" : [ "SERLOC_CURRENT_DIVISION_NAME" ]
}, {
  "functions" : [ {
    "name" : "approx_count_distinct"
  }, {
    "name" : "min"
  }, {
    "name" : "max"
  }, {
    "name" : "avg"
  }, {
    "name" : "stddev"
  }, {
    "name" : "stddev_pop"
  }, {
    "name" : "var_pop"
  }, {
    "name" : "sum"
  }, {
    "name" : "percentile_approx",
    "percentiles" : [ 0.25, 0.5, 0.75, 0.9, 0.95, 0.99 ]
  }
  ],
  "groupBy": "SNAPSHOTDATE, SERLOC_CURRENT_DIVISION_NAME, SERLOC_CURRENT_REGION_NAME",
  "columns" : [ "DAYSSINCE" ]
}
]
```

Exploration



Features

API Test

Overview

DataSet

Transformation

Graph

Regression

Prediction

VectorBuilder

Comment

Vector Builder

Manually build a simple experimental dataset for training, testing and prototyping. This part is not implemented

Dataset

Enter Name

Load

Save

Vector Table

Add Col

Remove Col

Add Row

Remove Row

Label

Feature 1

Feature 2

0

0

0

1

1

1

0

0

0

```
{"data": [{"0", "0", "0"}, {"1", "1", "1"}, {"0", "0", "0"}]}
```

Converter



Saspark: A Trans-Compiler Experiment
from SAS Dataset Transformation to Spark, Scala, Python & perhaps R

Demo

Result

Documentation

API

SAS

```
IF hispanic_index=0 THEN DO;
  * calculate x using nclv & adj_pml scores;
  x = nclv * (1 + adj_pml);
  IF 0<x<1 THEN v=2;
  ELSE IF x in (1,2,3) OR nclv>0 THEN DO; v=3; END;
  ELSE v=5;
  END;
  ELSE v = 100;
```

Python

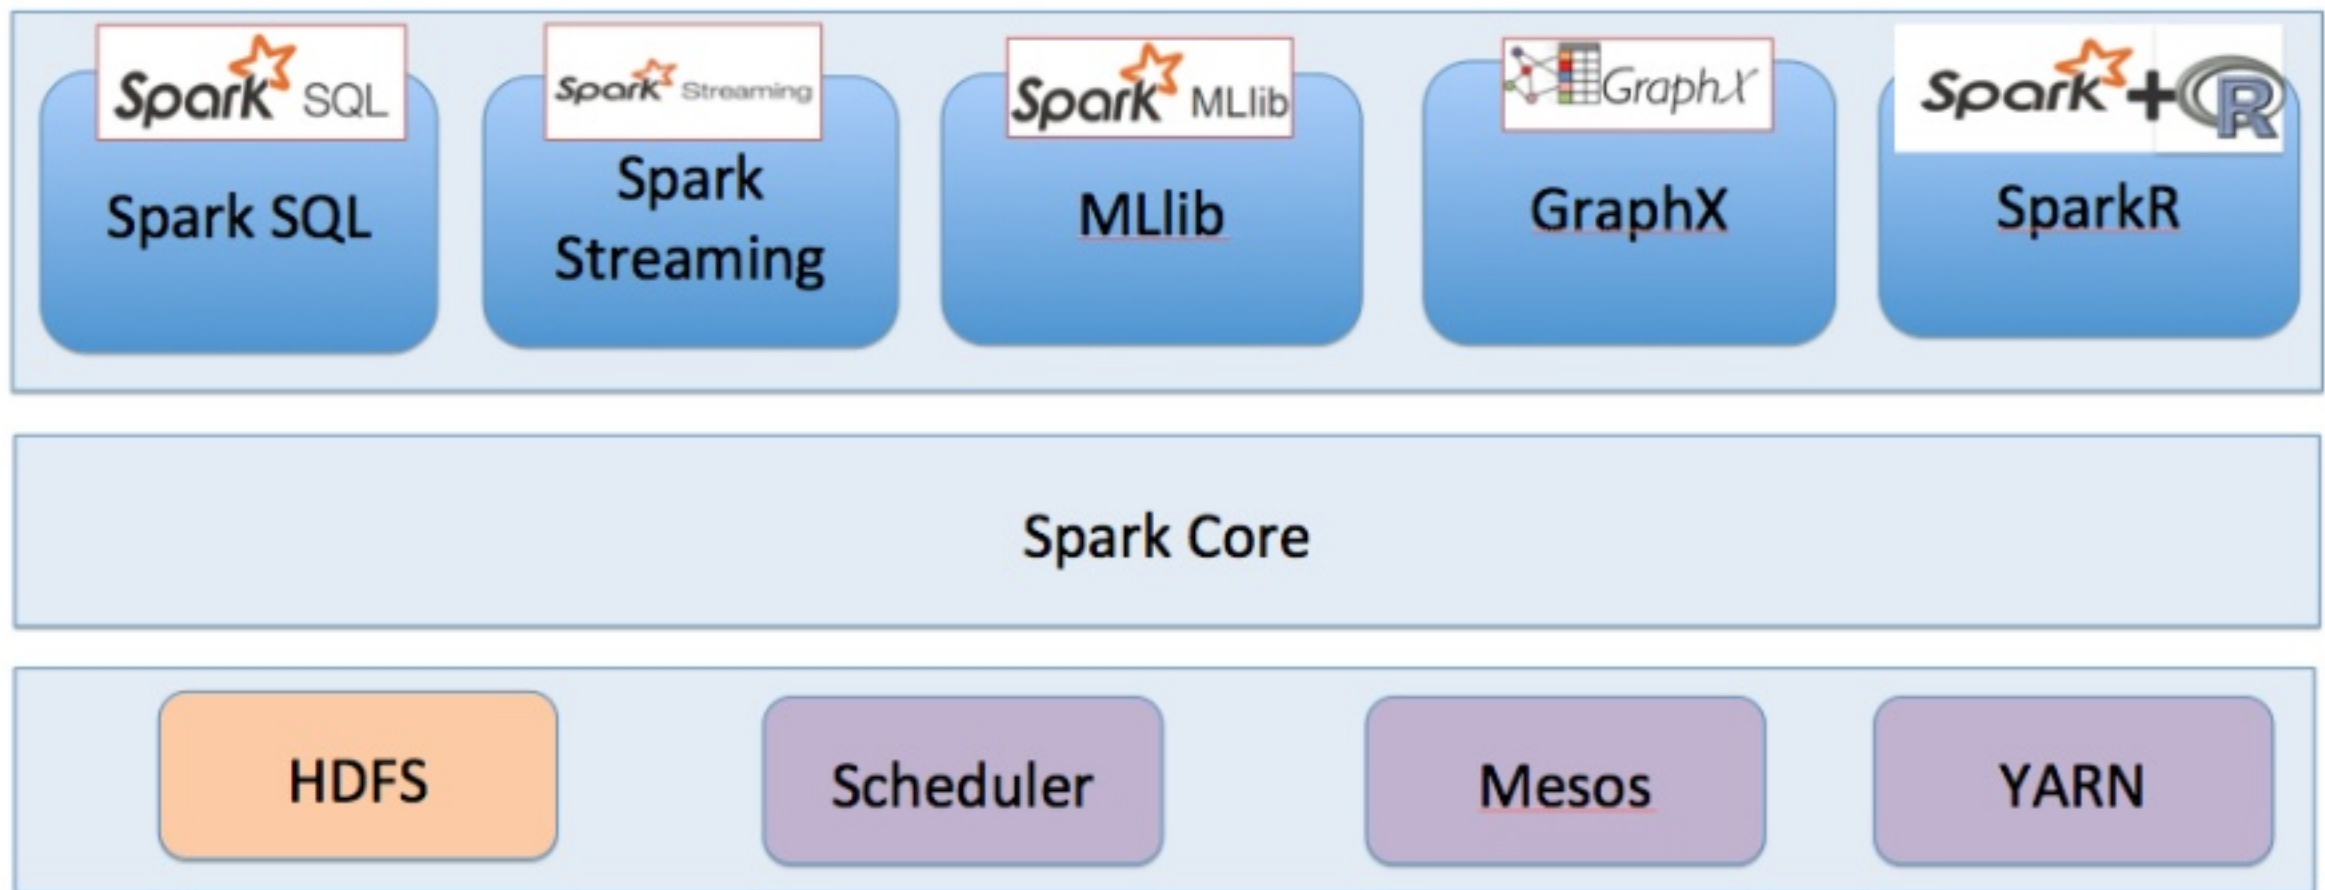
```
if row.hispanic_index == 0:
    # calculate x using nclv & adj_pml scores
    x = row.nclv * (1 + row.adj_pml)
    if 0 < x < 1:
        v = 2
    elif x in [1,2,3] or row.nclv > 0:
        v = 3
    else:
        v = 5
else:
    v = 100
```

Scala

```
//[Any] must be replaced with actual type for scala to compile
val adj_pml = row.getAs[Any]("adj_pml")
val nclv = row.getAs[Any]("nclv")
val hispanic_index = row.getAs[Any]("hispanic_index")
var x : Any
var v : Any

if (hispanic_index == 0) {
    // calculate x using nclv & adj_pml scores
    x = nclv * (1 + adj_pml)
    if (0 < x && x < 1) v = 2
    else if (List(1,2,3).contains(x) || nclv > 0) {
        v = 3
    }
    else v = 5
}
else v = 100
```

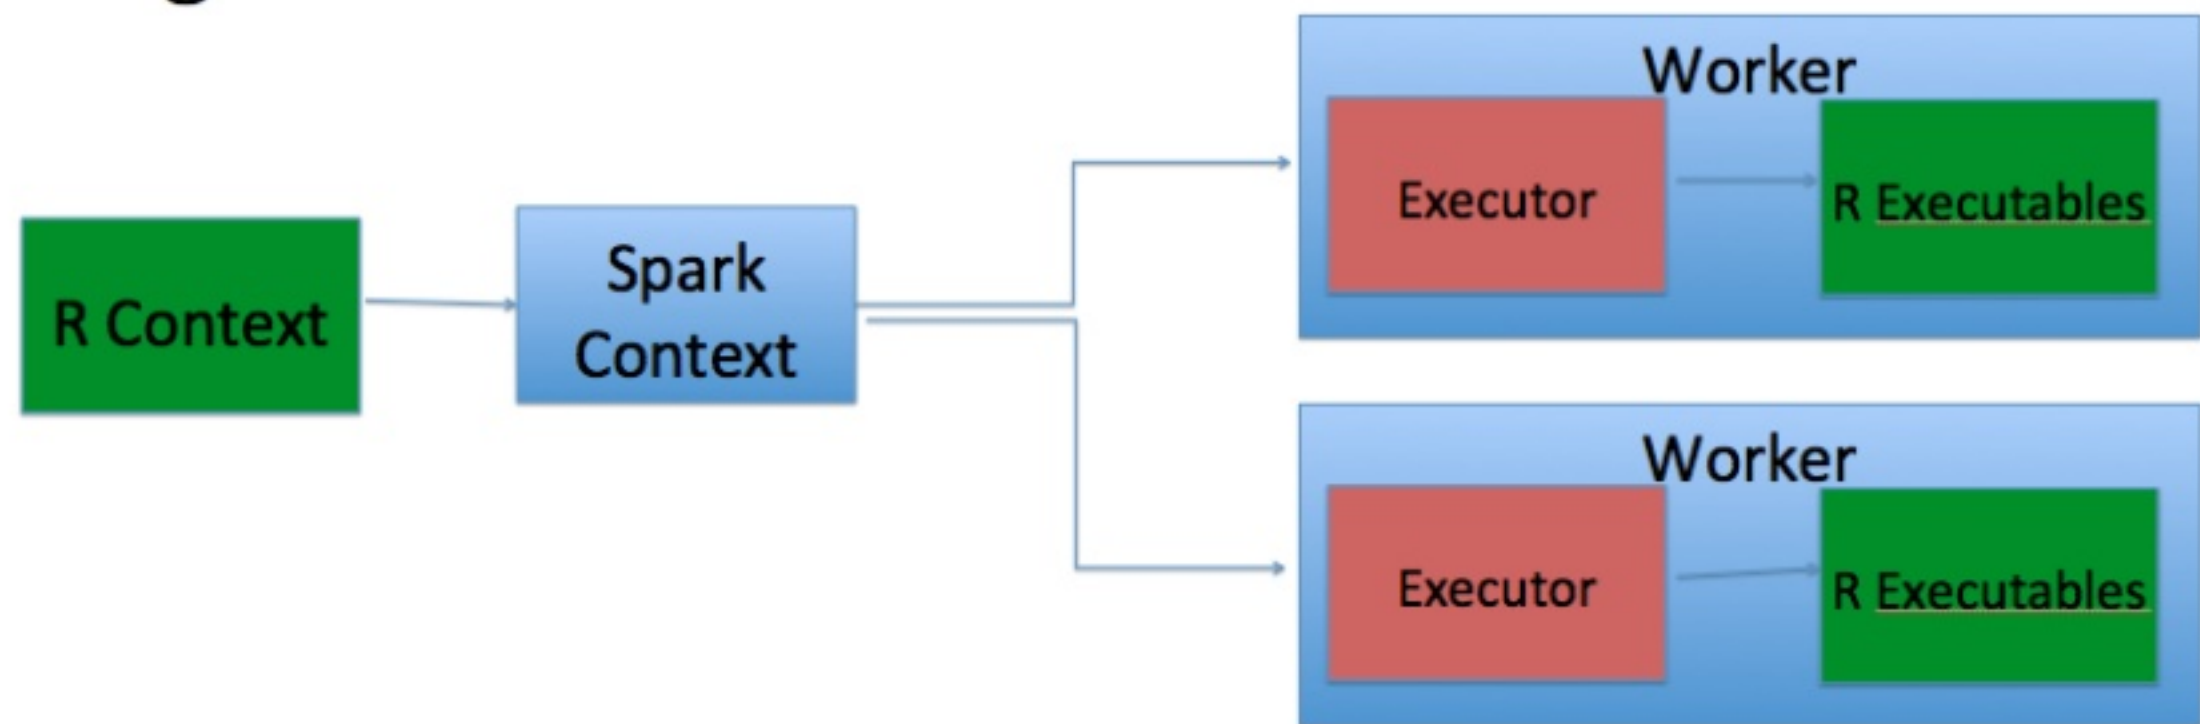
Spark Stack



SparkR





- Enables using R packages to process data
- Can run Machine Learning and Statistical Analysis



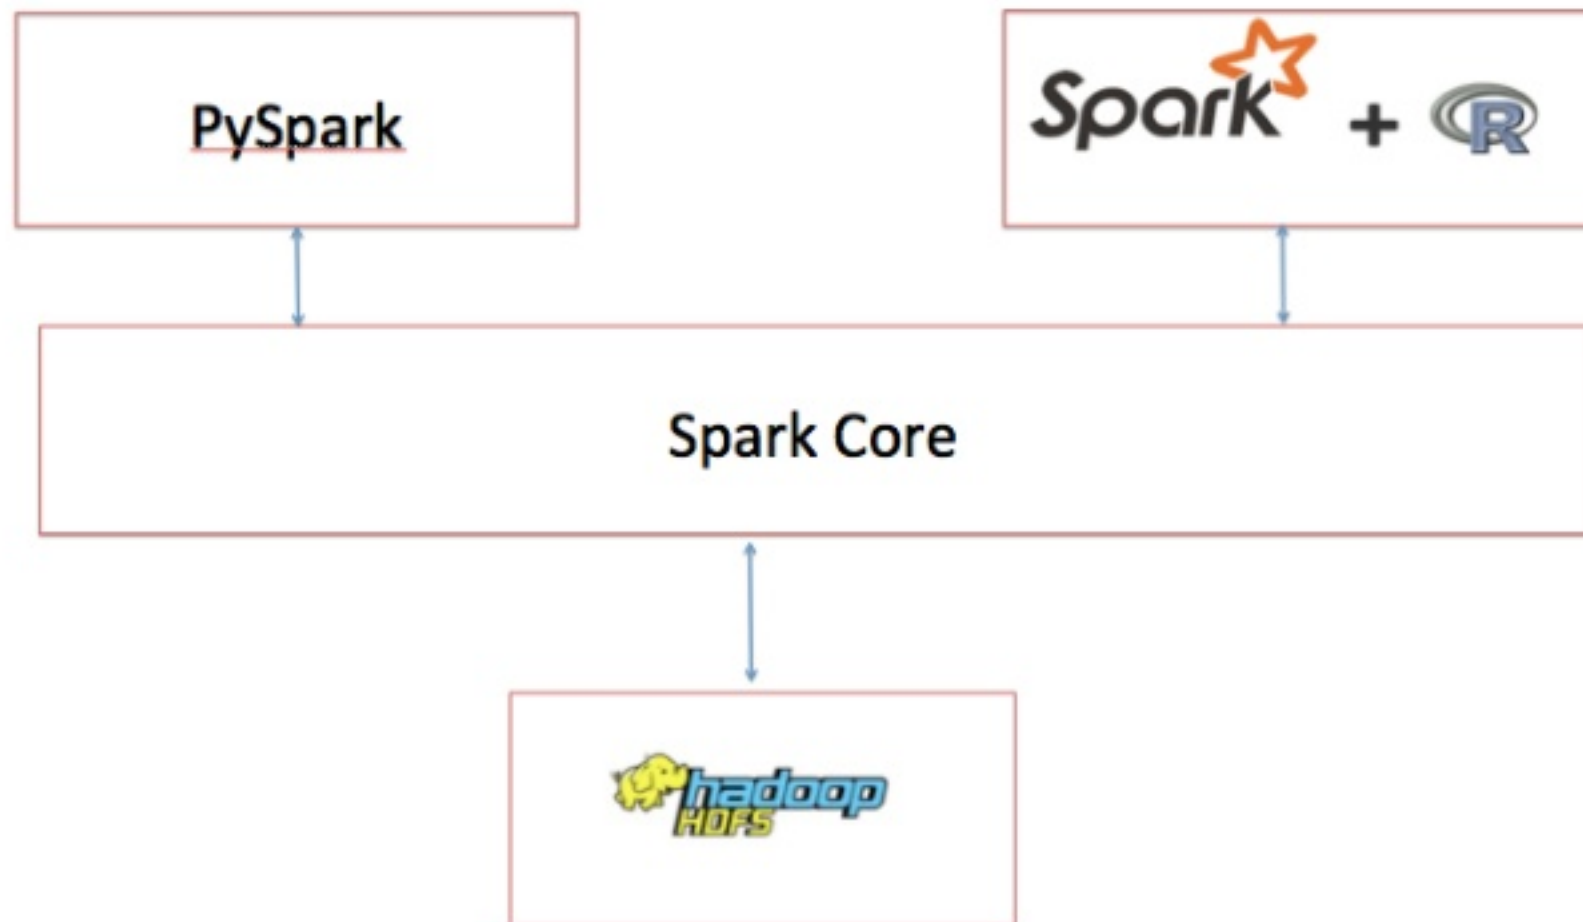
Spark MLlib



- Implements various Machine Learning Algorithms
- Classification, Regression, Collaborative Filtering, Clustering, Decomposition
- Works with Streaming, Spark SQL, GraphX or with SparkR.



PySpark, TensorFlow and SparkR



Scala and Spark for Big Data Analytics

Md. Rezaul Karim, Sridhar Alla

Tame Big Data with Scala and Apache Spark!

Preorder this eBook at 50% OFF and print book at 15% OFF

Packt>
www.packtpub.com



Discount Code:
eBook: **MLDROU50**
Print: **KVCUVW15**

Offer Expires
30th June 2017

We are hiring!

- Big Data Engineers (Hadoop, Spark, Kafka...)
- Data Analysts (R, SAS.....)
- Big Data Analysts (Hive, Pig)

jobs.comcast.com

Thank You.



Data Science Initiatives

- Customer Churn Prediction
- Click-thru Analytics
- Personalization
- Customer Journey
- Modeling
- Anomaly Detection
- GPU driven optimizations

Anomaly Detection

- Identification of observations which do not conform to an expected pattern.
- Ex: Network Intrusion Detection, Spikes in operational data, Unusual usage activity.

Popular Algorithms

- Unsupervised
 - KMeans
 - DBScan
- Supervised
 - HMM
 - Neural networks

KMeans Clustering

- Clustering is an unsupervised learning problem
- Groups subsets of entities with one another based on some notion of similarity.
- Easy to check if a new entity is falling outside known groups/clusters

Sample Code

```
import org.apache.spark.mllib.clustering.KMeans
import org.apache.spark.mllib.linalg.Vectors

val lines = sc.textFile("training.csv")
val data = lines.map(line => line.split(",").map(_.trim))
val inData = data.map{(x) => (x(3))}.map(_.toLong)
val inVector = inData.map{a => Vectors.dense(a)}.cache()
val numClusters = 3
val numIterations = 100
val kMeans = new KMeans().setMaxIterations(numIterations).setK(numClusters)
val kMeansModel = kMeans.run(inVector)

// Print cluster index for a given observation point
var ci = kMeansModel.predict(Vectors.dense(10000.0))
var ci = kMeansModel.predict(Vectors.dense(900008830.0))
```

Sample Code (R):

```
library('RHmm')
indata <- read.csv(file.choose(), header = FALSE, sep = ",", quote = "\"", dec = ".")
testdata <- read.csv(file.choose(), header = FALSE, sep = ",", quote = "\"", dec = ".")
dataSets <- c(as.numeric(indata$V4))
dataSetModel <- HMMFit(dataSets, nStates=3)
testdataSets <- c(as.numeric(testdata$V4))
tVitPath <- viterbi(dataSetModel, testdataSets)

#Forward-backward procedure, compute probabilities
tfb <- forwardBackward(dataSetModel, testdataSets)

# Plot implied states
layout(1:3) dataSet
plot(testdataSets[1:100], ylab="StateA", type="l", main="dataSet A")
plot(tVitPath$states[1:100], ylab="StateB", type="l", main="dataSet B")
```

Add Slides as Necessary

- Supporting points go here.





Thank You.

Contact information or call to action goes here.