

Why you should care about data layout in the file system

Cheng Lian, [@liancheng](#)

Vida Ha, [@femineer](#)

Spark Summit 2017



About Databricks

TEAM

Started Spark project (now Apache Spark) at UC Berkeley in 2009

MISSION

Making Big Data Simple

PRODUCT


Unified Analytics Platform



Apache Spark is a
powerful framework
with some temper



Just like
super mario



Serve him the
right ingredients



Powers up and
gets more efficient



Keep serving




He even knows
how to *Spark*!



However,
once served
a wrong dish...



Meh...



And sometimes...



It can be messy...

Secret sauces
we feed  **spark**TM



File Formats

Choosing a compression scheme



The obvious

- Compression ratio: the higher the better
- De/compression speed: the faster the better

Choosing a compression scheme



Splittable v.s. non-splittable

- Affects parallelism, crucial for big data
- Common splittable compression schemes
 - LZ4, Snappy, BZip2, LZO, and etc.
- *GZip is non-splittable*
 - Still common if file sizes are \ll 1GB
 - Still applicable for Parquet

Columnar formats



Smart, analytics friendly, optimized for big data

- Support for nested data types
- Efficient data skipping
 - Column pruning
 - Min/max statistics based predicate push-down
- Nice interoperability
- Examples:
 - Spark SQL built-in support: [Apache Parquet](#) and [Apache ORC](#)
 - Newly emerging: [Apache CarbonData](#) and Spinach

Columnar formats



Parquet

- Apache Spark default output format
- Usually the best practice for Spark SQL
- Relatively heavy write path
 - Worth the time to encode for repeated analytics scenario
- Does not support fine grained appending
 - Not ideal for, e.g., collecting logs
- Check out Parquet [presentations](#) for more details

Semi-structured text formats



Sort of structured but not self-describing

- Excellent write path performance but slow on the read path
 - Good candidates for collecting raw data (e.g., logs)
- Subject to inconsistent and/or malformed records
- Schema inference provided by Spark (for JSON and CSV)
 - Sampling-based
 - Handy for exploratory scenario but can be inaccurate
 - Always specify an accurate schema in production

Semi-structured text formats



JSON

- Supported by Apache Spark out of the box
- One JSON object per line for fast file splitting
- JSON object: map or struct?
 - Spark schema inference always treats JSON objects as structs
 - Watch out for arbitrary number of keys (may OOM executors)
 - Specify an accurate schema if you decide to stick with maps

Semi-structured text formats



JSON

- Malformed records
 - Bad records are collected into column `_corrupted_record`
 - All other columns are set to null

Semi-structured text formats



CSV

- Supported by Spark 2.x out of the box
 - Check out the [spark-csv](#) package for Spark 1.x
- Often used for handling legacy data providers & consumers
 - Lacks of a standard file specification
 - Separator, escaping, quoting, and etc.
 - Lacks of support for nested data types

Raw text files



Arbitrary line-based text files

- Splitting files into lines using `spark.read.text()`
 - Keep your lines a reasonable size
- Keep file size < 1GB if compressed with a non-splittable compression scheme (e.g., GZip)
- Handling inevitable malformed data
 - Use a `filter()` transformation to drop bad lines, or
 - Use a `map()` transformation to fix bad line




Directory layout

Partitioning



Overview

- Coarse-grained data skipping
- Available for both persisted tables and raw directories
- Automatically discovers Hive style partitioned directories



Partitioning



SQL

```
CREATE TABLE ratings
USING PARQUET
PARTITIONED BY (year, genre)
AS SELECT artist, rating, year, genre
FROM music
```

DataFrame API


```
spark
  .table("music")
  .select('artist', 'rating', 'year', 'genre')
  .write
  .format("parquet")
  .partitionBy('year', 'genre')
  .saveAsTable("ratings")
```

Partitioning



Filter predicates

Use simple filter predicates containing partition columns to leverage partition pruning



Partitioning



Filter predicates

- year = 2000 AND genre = 'folk'
- year > 2000 AND rating > 3
- year > 2000 OR genre <> 'rock'



Partitioning



Filter predicates

- year > 2000 OR rating = 5
- year > rating




Partitioning



Avoid excessive partitions

- Stress metastore for persisted tables
- Stress file system when reading directly from the file system
- Suggestions
 - Avoid using too many partition columns
 - Avoid using partition columns with too many distinct values
 - Try hashing the values
 - E.g., partition by first letter of first name rather than first name

Partitioning



Scalable partition handling

Using persisted partitioned tables with Spark 2.1+

- Per-partition metadata gets persisted into the metastore
- Avoids unnecessary partition discovery (esp. valuable for S3)


Check our [blog post](#) for more details

Bucketing



Overview

- *Pre-shuffles* and optionally *pre-sorts* the data while writing
- Layout information gets persisted in the metastore
- Avoids shuffling and sorting when joining large datasets
- Only available for persisted tables



Bucketing



SQL

```
CREATE TABLE ratings
USING PARQUET
PARTITIONED BY (year, genre)
CLUSTERED BY (rating) INTO 5 BUCKETS
SORTED BY (rating)
AS SELECT artist, rating, year, genre
FROM music
```

DataFrame

```
ratings
  .select('artist', 'rating', 'year', 'genre')
  .write
  .format("parquet")
  .partitionBy("year", "genre")
  .bucketBy(5, "rating")
  .sortBy("rating")
  .saveAsTable("ratings")
```

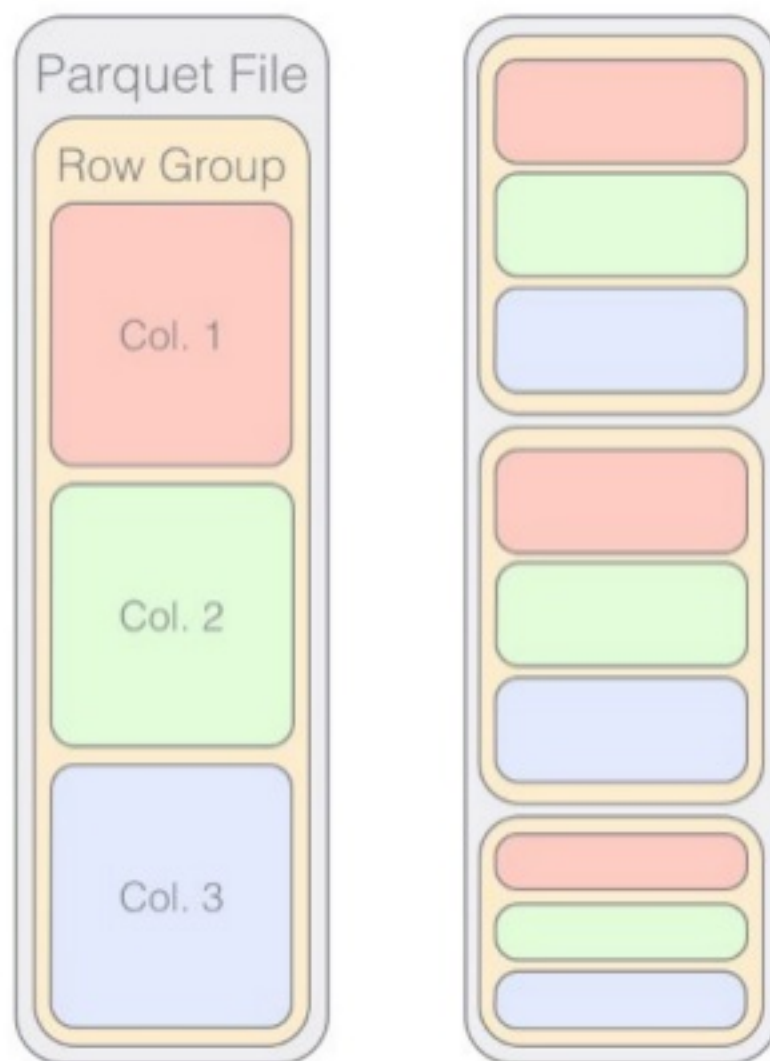
Bucketing



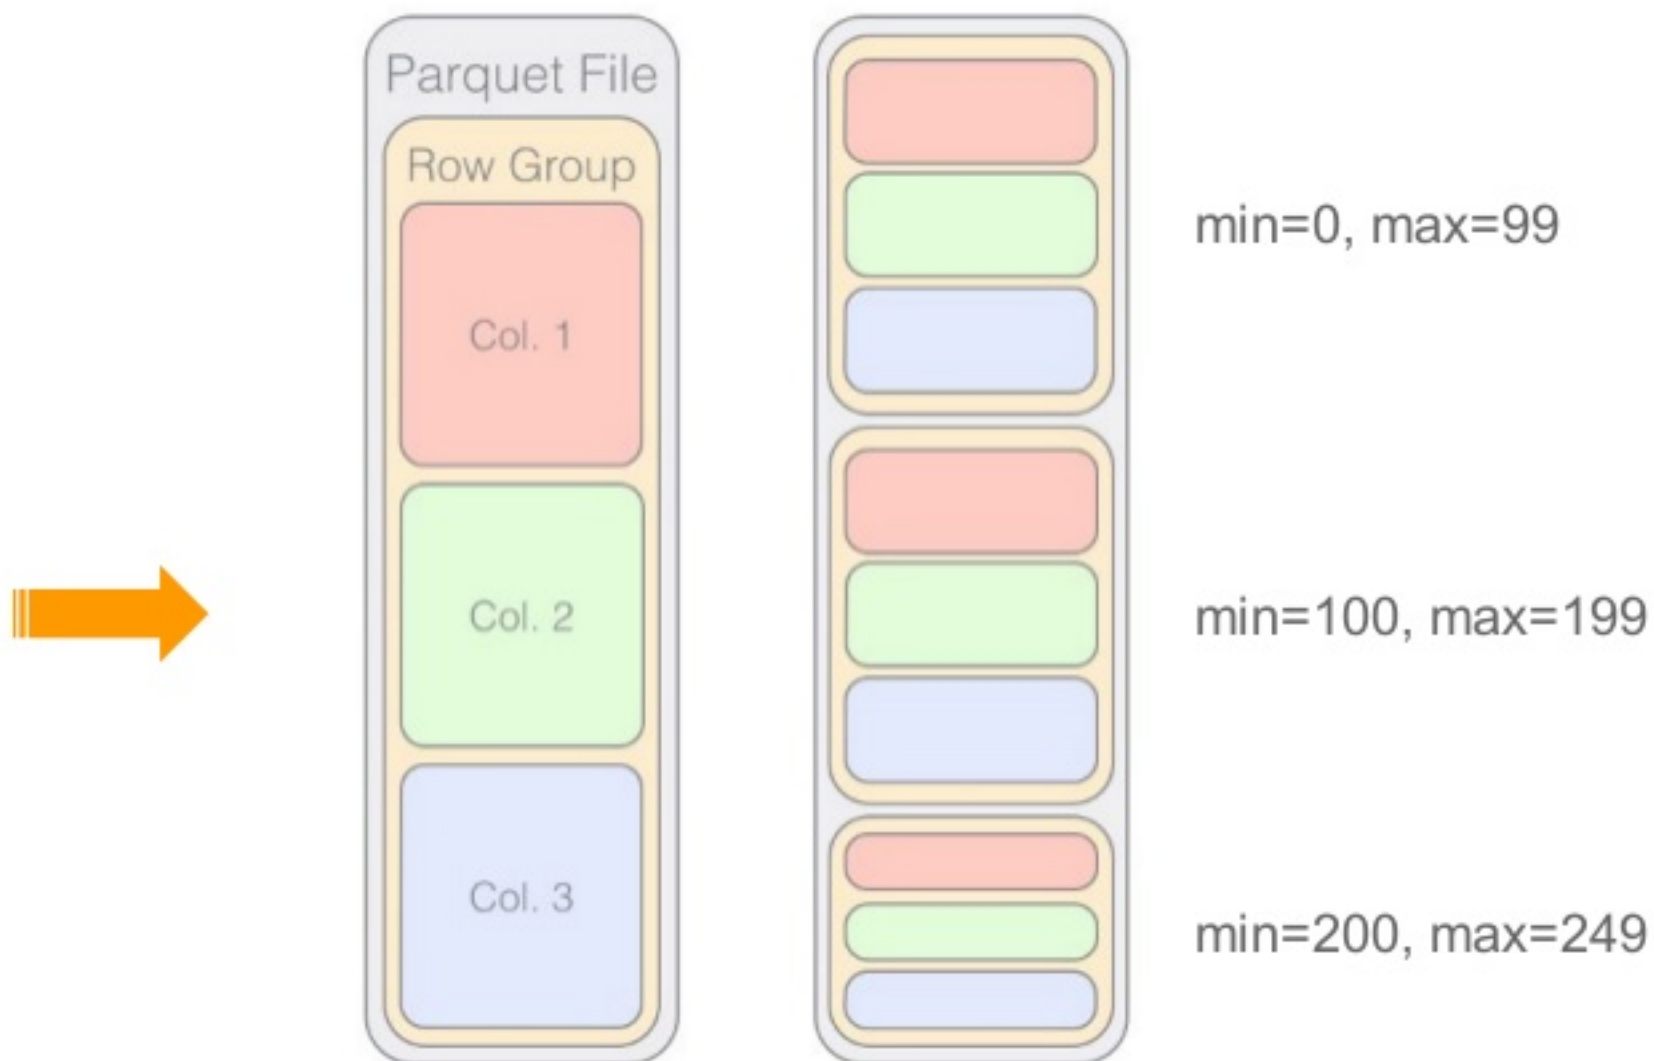
In combo with columnar formats

- Bucketing
 - Per-bucket sorting
- Columnar formats
 - Efficient *data skipping* based on min/max statistics
 - Works best when the searched columns are *sorted*

Bucketing



Bucketing



Bucketing



In combo with columnar formats

Perfect combination, makes your Spark jobs FLY!



★ More tips

File size and compaction



Avoid small files

- Cause excessive parallelism
 - Spark 2.x improves this by packing small files
- Cause extra file metadata operations
 - Particularly bad when hosted on S3

File size and compaction



How to control output file sizes

- In general, one task in the output stage writes one file
 - Tune parallelism of the output stage
- `coalesce(N)`, for
 - Reduces parallelism for small jobs
- `repartition(N)`, for
 - Increasing parallelism for all jobs, or
 - Reducing parallelism of final output stage for large jobs
 - Still preserves high parallelism for previous stages

True story



Customer

- Spark ORC Read Performance is much slower than Parquet
- The same query took
 - 3 seconds on a Parquet dataset
 - 4 minutes on an equivalent ORC dataset

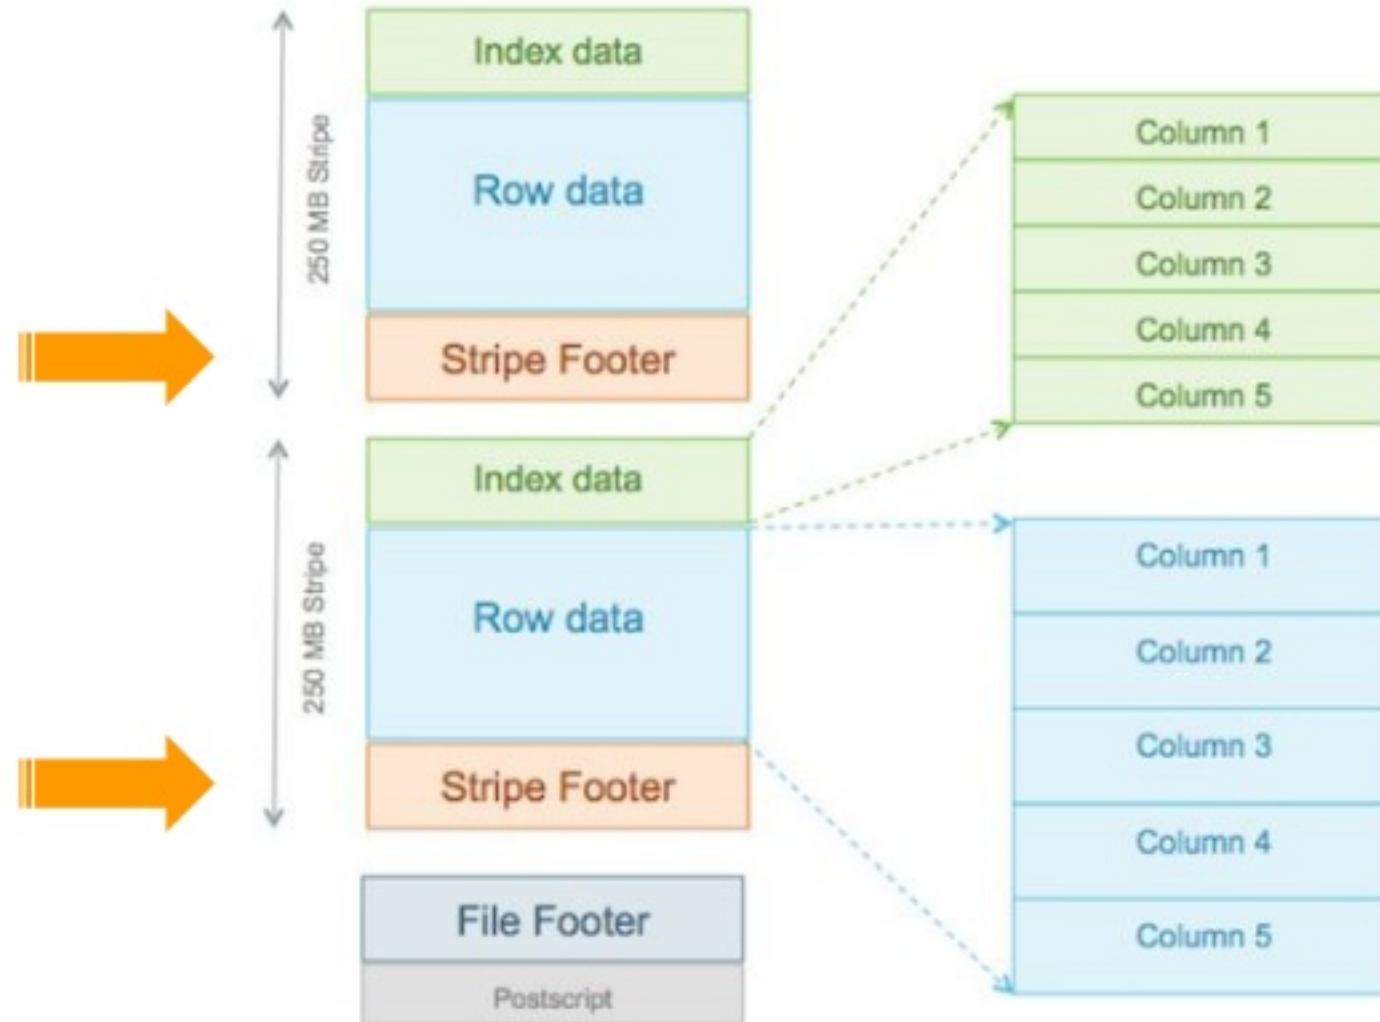
True story



Me

- Ran a simple `count(*)`, which took
 - **Seconds** on the Parquet dataset with a handful IO requests
 - **35 minutes** on the ORC dataset with 10,000s of IO requests
- Most task execution threads are *reading ORC stripe footers*

True story



True story



```
import org.apache.hadoop.hive.q1.io.orc._
import org.apache.hadoop.conf.Configuration
import org.apache.hadoop.fs.Path

val conf = new Configuration


def countStripes(file: String): Int = {
  val path = new Path(file)
  val reader = OrcFile.createReader(path, OrcFile.readerOptions(conf))
  val metadata = reader.getMetadata
  metadata.getStripeStatistics.size
}
```

True story



Maximum file size: ~15 MB

Maximum ORC stripe counts: ~1,400



True story



Root cause

Malformed (but not corrupted) ORC dataset

- ORC readers read the footer first before consuming a strip
- ~1,400 stripes within a single file as small as 15 MB
- ~1,400 x 2 read requests issued to S3 for merely 15 MB of data

True story



Root cause

Malformed (but not corrupted) ORC dataset

- ORC readers read the footer first before consuming a strip
- ~1,400 stripes within a single file as small as 15 MB
- ~1,400 x 2 read requests issued to S3 for merely 15 MB of data

Much worse than even CSV, not mention Parquet

True story



Why?

- Tiny ORC files (~10 KB) generated by Streaming jobs
 - Resulting one tiny ORC stripe inside each ORC file
 - The footers might take even more space than the actual data!

True story



Why?

Tiny files got compacted into larger ones using

```
ALTER TABLE ... PARTITION (...) CONCATENATE;
```

The **CONCATENATE** command just, well, *concatenated* those tiny stripes and produced larger (~15 MB) files with a huge number of tiny stripes.

True story



Lessons learned

Again, avoid writing small files in *columnar formats*

- Output files using CSV or JSON for Streaming jobs
 - For better write path performance
- Compact small files into large chunks of columnar files later
 - For better read path performance

True story



The cure

Simply read the ORC dataset and write it back using

```
spark.read.orc(input).write.orc(output)
```

So that stripes are adjusted into more reasonable sizes.

Schema evolution



Columns come and go

- Never ever change the data type of a published column
- Columns with the same name should have the same data type
- If you really dislike the data type of some column
 - Add a new column with a new name and the right data type
 - Deprecate the old one
 - Optionally, drop it after updating all downstream consumers

Schema evolution



Columns come and go

Spark built-in data sources that support schema evolution

- JSON
- Parquet
- ORC

Schema evolution



Common columnar formats are less tolerant of data type mismatch. E.g.:

- INT cannot be promoted to LONG
- FLOAT cannot be promoted to DOUBLE

JSON is more tolerating, though

- LONG → DOUBLE → STRING

True story



Customer

Parquet dataset corrupted!!! **HALP!!!**



True story

What happened?

Original schema

- {col1: DECIMAL(19, 4), col2: INT}

Accidentally appended data with schema

- {col1: DOUBLE, col2: DOUBLE}

All files written into the same directory

True story



What happened?

Common columnar formats are less tolerant of data type mismatch. E.g.:

- INT cannot be promoted to LONG
- FLOAT cannot be promoted to DOUBLE

Parquet considered these schemas as incompatible ones and refused to merge them.

True story



BTW

JSON schema inference is more tolerating

- LONG → DOUBLE → STRING

However

- JSON is NOT suitable for analytics scenario
- Schema inference is unreliable, not suitable for production

True story



The cure

Correct the schema

- Filter out all the files with the wrong schema
- Rewrite those files using the correct schema

Exhausting because all files are appended into a single directory

True story



Lessons learned

- Be very careful on the write path
- Consider partitioning when possible
 - Better read path performance
 - Easier to fix the data when something went wrong

Recap



File formats

- Compression schemes
- Columnar (Parquet, ORC)
- Semi-structured (JSON, CSV)
- Raw text format

Directory layout

- Partitioning
- Bucketing

Other tips

- File sizes and compaction
- Schema evolution

Try Apache Spark in Databricks!

UNIFIED ANALYTICS PLATFORM

- Collaborative cloud environment
- Free version (community edition)

DATABRICKS RUNTIME 3.0

- Apache Spark - optimized for the cloud
- Caching and optimization layer - DBIO
- Enterprise security - DBES

Try for free today.
databricks.com

Early draft available for free today!

go.databricks.com/book





Thank you

Q & A